
Elliptic Curves

— *EYNTKA* Series —

NATHANAEL CHWOJKO-SRAWLEY

March 24, 2026

Contents

Contents	3
Curves	i
0.1 Speedy Summary of Varieties	i
0.2 Curves	ii
0.3 Frobenius Map	vi
0.4 Divisors	viii
0.5 Differentials	x
0.6 Riemann-Roch Theorem	xv
 I Foundations: Geometry and Algebra of Elliptic Curves	 1
1 Basics of Elliptic Curves	5
1.1 Elliptic Curves and Weierstrass Equations	5
1.2 The Discriminant and the j -Invariant	14
1.3 The Group Law	24
1.4 The Invariant Differential	33
1.5 Automorphisms and Twists	38
1.6 Singular Cubics	45
 2 Isogenies and the Endomorphism Ring	 53
2.1 Isogenies — Definition and First Properties	53

2.2	The Dual Isogeny	62
2.3	The Multiplication-by- n Map and the Structure of $E[n]$	72
2.4	The Frobenius Endomorphism	79
2.5	The Tate Module and the Weil Pairing	85
2.6	The Endomorphism Ring	90
3	The Analytic Theory	97
3.1	Lattices and Elliptic Functions	97
3.2	The Weierstrass $\wp$ -Function	102
3.3	The Uniformization Theorem	109
3.4	Eisenstein Series and the j -Invariant	116
3.5	The Analytic Dictionary	122
3.6	The Modular Group and the Fundamental Domain	130
4	The Formal Group of an Elliptic Curve	137
4.1	Formal Group Laws	138
4.2	The Formal Group of an Elliptic Curve	144
4.3	The Formal Logarithm and Formal Exponential	152
4.4	Height and the $[p]$ -Series	159
4.5	Formal Groups over Local Fields	166
II	Arithmetic: Local and Global Theory	175
5	Elliptic Curves over Finite Fields	179
5.1	The Characteristic Polynomial of Frobenius	179
5.2	The Zeta Function and the Weil Conjectures	185
5.3	The Group Structure of $E(\mathbb{F}_q)$	190
5.4	Tate's Isogeny Theorem	195
5.5	Supersingular Curves and the Deuring Correspondence	200
6	Elliptic Curves over Local Fields	207
6.1	Minimal Weierstrass Equations	208
6.2	Reduction Types	212
6.3	Points of Finite Order	217
6.4	The Action of Inertia and the Néron–Ogg–Shafarevich Criterion	222
6.5	Tate's Algorithm and the Kodaira–Néron Classification	227
6.6	The Component Group and Tamagawa Numbers	231

7	Build-up: Galois Cohomology	237
7.1	Group Cohomology	238
7.2	Galois Cohomology of Fields	242
7.3	Cohomology of Elliptic Curves	247
7.4	Twists Revisited	253
7.5	Local and Global Duality	256
8	Elliptic Curves over Number Fields: The Mordell–Weil Theorem	263
8.1	The Weak Mordell–Weil Theorem	263
8.2	Heights	266
8.3	The Néron–Tate Height	274
8.4	Proof of the Mordell–Weil Theorem	280
8.5	The Rank, Torsion, and the Structure of $E(K)$	286
9	Computing the Mordell–Weil Group	293
9.1	The 2-Descent with Full 2-Torsion	293
9.2	Descent via 2-Isogeny	299
9.3	Higher Descent and the Limits of 2-Descent	303
9.4	Searching for Points and Height Bounds	310
9.5	Saturation and Verifying Generators	315
10	Elliptic Curves over Number Rings	321
10.1	Siegel’s Theorem	321
10.2	Baker’s Method and Effective Bounds	328
10.3	Computing Integral Points in Practice	334
10.4	The ABC Conjecture and Szpiro’s Conjecture	340
III	Advanced Topics: Representations, Modularity, and Beyond	345
	Index	348
	Bibliography	351

Curves

(here, will be a quick overview from other notes)

0.1 Speedy Summary of Varieties

This is a *very* speedy reminder of how varieties are defined in modern algebraic geometry.

Let $k = \bar{k}$ be an algebraically closed field. Then the set of points k^n will be denoted $\mathbb{A}_k^n$ (so that there is no confusing $\mathbb{A}_k^n$ with a vector-space). Any polynomial $f(x_1, \dots, x_n)$ in n -variables can be defined as a function $\mathbb{A}_k^n \rightarrow \mathbb{A}_k^1$. The zero-set of these polynomials are closed under arbitrary intersection and finite unions, and $\emptyset$ as well as $\mathbb{A}_k^n$ are in this collection (given by $0, 1 \in k[x_1, \dots, x_n]$) and hence forms a topology known as the *Zariski topology*. The closed sets in this space are called the *algebraic sets*. The Zariski topology has non-trivial irreducible sets. The irreducible sets are called *varieties*. See [8] for the important notions about varieties, including the global sections that correspond to them. By the Nullstellensatz, every maximal ideal of a global section ring of a variety corresponds to a point in the variety. This lays the foundation for considering more general rings as geometric objects. It was observed by Grothendieck that we should consider the collection of *prime ideals* as they operate better functorially (namely, the pre-image of a prime ideal is always prime). The collection of all prime ideals is the spectrum of a ring R . The Zariski topology is defined similarly. We then keep track of the global and local sections on open subsets via a sheaf known as the *structure sheaf*; a spectrum along with the structure sheaf is called an *affine scheme*. A *Scheme* is a (locally) ringed space $(X, \mathcal{O}_X)$ that is covered by open sets that are isomorphic to affine schemes. The dimension of a scheme is defined as the maximum chain of irreducible subsets (starting the indexing at 0). A curve and a surface is a one and two dimensional (irreducible) scheme respectively. A k -variety is a separated integral reduced k -scheme of finite type. Varieties defined in this way correspond to quasi-projective varieties in classical algebraic geometry. A proper variety is called a *complete variety*, where classical projective varieties (when interpreted as schemes) make up most examples of such varieties (though there do exist non-projective complete varieties, see for example the blow up example in [3]). This book focuses on 1-dimensional complete varieties with genus 1 that have at least one solution.

Let us now set some conventions, and for simplicity stick to classical algebraic geometry language. Let K be given (and recall assumed to be perfect). Then we shall call *affine n -space*:

$$\mathbb{A}^n = \mathbb{A}(\overline{K}) = \{P = (x_1, \dots, x_n) \mid x_i \in \overline{K}\}$$

Notice that we are taking its algebraic closure; we can think of this as the more *geometric setting*. The set of K -rational points of $\mathbb{A}^n$ will be denoted:

$$\mathbb{A}^n(K) = \{P = (x_1, \dots, x_n) \mid x_i \in K\}$$

There is a natural way of defining $\mathbb{A}^n(K)$: if $G_{\overline{K}/K} = \text{Gal}_K(\overline{K})$ is the absolute galois group with respect to K , then there is a natural $G_{\overline{K}/K}$ action on $\mathbb{A}^n$

$$G_{\overline{K}/K} \curvearrowright \mathbb{A}^n \quad \sigma \cdot P = P^\sigma = (x_1^\sigma, \dots, x_n^\sigma)$$

We shall let $V \subseteq \mathbb{A}^n$ be our varieties (in particular, the coefficients of the defining polynomials belong to $\overline{K}$). Let V/K be a variety defined over the field K , and $V(K)$ be the set:

$$V(K) = V \cap \mathbb{A}^n(K)$$

Note the difference between the two: V/K takes on coefficients in K , while $V(K)$ allows for “solutions” over K . We shall often be looking at $V/\mathbb{Q}$ and ask for $V(K)$ for various fields (ex. $\mathbb{C}$, local fields, global fields, etc.) Note that $V(K)$ could also be defined as the points of V which are invariant under the $G_{\overline{K}/K}$ action:

$$V(K) = \left\{ P \in V \mid P^\sigma = P, \text{ for all } \sigma \in G_{\overline{K}/K} \right\}$$

0.2 Curves

We shall assume k is a perfect field (that every algebraic extension is separable), unless stated otherwise.

We summarize the important concepts from [3] and [8], with some extra details for the material that is covered later in [3] to alleviate some of the burden of pre-requisite knowledge. We shall consider a curve to be a one-dimensional k -scheme that is

- Noetherian (the natural finiteness condition)
- Integral (hence it is irreducible and reduced)
- Separated (no points are arbitrarily close to each other)

It was shown in [3] that all such schemes *must* be projective, and so we can simplify our notion of curve by saying a curve is a projective variety of dimension 1 (we shall provide a proof in chapter 1 for elliptic curves). As C is irreducible, the rational field for C is well-defined, and is denoted $k(C)$. The notation C/k shall often be written to represent the curve over the field k , as this may be easier to reader as compact to $C(k)$ which is similar to the $k(C)$ notation.

If furthermore all local rings are regular, the curve is said to be a *regular curve* or *smooth curve*. In one dimension, regular local rings are DVRs. If C is a curve and $\mathcal{O}_p(C)$ is the local ring at a regular point p , we can define the discrete valuation ν or ord as:

$$\nu_p : \mathcal{O}_p(C) \rightarrow \mathbb{N} \cup \{\infty\} \quad \nu_p(f) = \sup \{d \in \mathbb{Z} \mid f \in \mathfrak{m}_p^d\}$$

This can be extended to $k(C)$ by defining $\nu_p(f/g) = \nu_p(f) - \nu_p(g)$, giving us:

$$\nu_p : k(C) \rightarrow \mathbb{Z} \cup \{\infty\}$$

Recall that the *uniformizer* for C at p is a function whose order of vanishing is 1, and it is unique up to multiplication by a unit. Using the discrete valuation, we can identify the zero's and poles of a function $f \in k(C)$, define the regular points when $\nu_p(f) \geq 0$, and prove that all but finitely many points are regular¹.

When working over 1-dimensional irreducible projective varieties, we have enough restriction to be able to define smooth curves by taking the Zariski topology on $k(C)$ on the set of DVR's that dominate k . Hence, smooth curves are characterized by the collection DVR's, and can even be defined in this way, with morphisms given by the map $\varphi^* : k(C_2) \rightarrow k(C_1)$ where $f \mapsto f \circ \varphi$ (²)

For the purposes of working with curves over finite characteristic, the following lemma will prove useful:

Lemma 0.2.1: Curves and Uniformizers

Let C/k be a curve, $p \in C$ a point, and $t \in k(C)$ a uniformizer for the in p . Then $k(C)/k(t)$ is a finite separable field extension.

Proof:

Let $p = \text{char}(k)$.

First, the extension is certainly finite by Noether's Normalization theorem as it is finitely generated over k and has transcendence degree 1 over k (as C is a curve) and $t \notin k$. What's left to show is that any $a \in k(C)$ is separable over $k(t)$.

As it is algebraic over $k(t)$, there is a polynomial equation $\varphi(t, x)$ such that $\varphi(t, a) = 0$. Without loss of generality we may assume φ is a minimal polynomial for a . Now, if φ contains a monomial $a_{ij}t^i x^j$ where $j \not\equiv 0 \pmod{p}$, then the x -derivative of φ is non-zero, hence x is separable over $k(t)$.

For the sake of contradiction, let's say this is not the case. Then we may re-write $\varphi(t, x) = \psi(t, x^p)$. Then as $\text{char}(k) = p$, we have that $\varphi(t, x)^p = \varphi(t^p, x^p)$. Note how it was important that k is a separable extension so that each coefficient of $\varphi(t, x)$ can be represented as the p th power of an element in an algebraic extension. Thus, if $a_{ij} = b_{ij}^p$, we can write:

$$\varphi(t, x) = \psi(t, x^p) = \sum_{k=0}^{p-1} \left(\sum_{i,j} b_{ijk} t^{ip} x^{jp} \right) t^k = \sum_{k=0}^{p-1} \varphi_k(t, x)^p t^k$$

Now, as φ is the minimal polynomial, $\varphi(t, a) = 0$. On the other hand, as t is the uniformizer at p , we have:

$$\nu_p(\varphi_k(t, a)^p t^k) = p\nu_p(\varphi_k(t, a)) + k\nu_p(t) \equiv k \pmod{p}$$

Thus, each term $\varphi_k(t, a)^p t^k$ must have a distinct order at p , and hence every term must vanish:

$$\varphi_0(t, a) = \varphi_1(t, a) = \cdots = \varphi_{p-1}(t, a) = 0$$

¹A very simple but elegant proof of this fact is that the stalk at the generic point is locally free in some open neighbourhood and is isomorphic to the differential of the generic point, which means it is regular on an open dense subset, and in dimension 1 such a set is cofinite

²Though this characterization does not generalize to higher dimensions as the local valuation do not form a scheme, see (this link – commented because not compiling pdf) for more details

However, at least one of the terms $\varphi_k(t, x)$ must have an x term, which will then be a polynomial of degree *smaller* than the minimal polynomial ($\deg_x \varphi_k(t, x) \leq \frac{1}{p} \deg_x \varphi(t, x)$), contradicting minimality of $\varphi(t, x)$, as we sought to show.

Let us now look at maps between curves. Recall that a rational map $\varphi : C_1 \rightarrow V \subseteq \mathbb{P}_k^n$ is of the form $[f_0, f_1, \dots, f_n]$ where the polynomials are subject to the restrictions given by the defining polynomials for V . Most of the time, we shall consider the case where the codomain is a curve C_2 . Let $\deg \varphi = [k(C_1) : \varphi^*(k(C_2))]$ be the *degree* of φ . If we must keep track of the separable and inseparable extensions and their corresponding degree, we shall write $\deg_s \varphi$ and $\deg_i \varphi$ respectively. One thing that we left as an exercise in [3] that is worth pointing out is that we can define a map $\varphi_* : k(C_1) \rightarrow k(C_2)$ by doing the following: let $L = k(C_1)$ and $K = k(C_2)$, and $N_{L/K} : L \rightarrow K$ be the field norm. Then:

$$\varphi_* = (\varphi^*)^{-1} \circ N_{L/K}$$

This is the *pushforward*, as compared to φ^* which is the *pullback*.

If $V = \mathbb{P}_k^1$, then we have $[f_0, f_1]$ with no defining restrictions. If $f_1 \neq 0$, then we can write this map as $[f_0/f_1 : 1]$. If $f_1 = 0$, then $[f_0 : 0] = [1 : 0]$ and hence it is a constant function at infinity. Thus, all rational maps $C \rightarrow \mathbb{P}_k^1$ are in bijective correspondence with $k(C) \cup \{\infty\}$. When working with rational maps over curves, most points will be regular; indeed show that if $p \in C$ is smooth then φ is regular at p (exercise).

Next, it was shown in [8] that maps to a projective variety is either surjective or constant, hence a regular morphism $C_1 \rightarrow C_2$ is either surjective or constant. When φ is non-constant, φ induce a map $\varphi^* : k(C_2) \rightarrow k(C_1)$. These maps between function fields and rational maps between polynomials have the following properties:

Proposition 0.2.2: Properties of Pullbacks of Morphisms

1. $k(C_1)$ is a finite extension of $\varphi^*(k(C_2))$
2. if $\iota : k(C_2) \rightarrow k(C_1)$ is a k -homomorphism, there is a unique non-constant rational map $\varphi : C_1 \rightarrow C_2$ such that $\varphi^* = \iota$
3. If $L \subseteq k(C_1)$ is a subfield of finite index containing k , then there exists a unique smooth curve C'/k up to k -isomorphism and a non-constant rational map $\varphi : C_1 \rightarrow C'$ such that $\varphi^*(k(C')) = L$
4. If $\varphi : C_1 \rightarrow C_2$ is a map of degree 1 (so $[k(C_1) : \varphi^*(k(C_2))] = 1$), then φ^* is an isomorphism

Proof:

see [8]

These results that the category of smooth curves over k with morphisms being non-constant rational morphisms (which as smooth curves must be regular, can equivalently be characterized as surjective regular morphisms) is dual to the category whose objects are finitely generated extensions K/k of transcendence degree 1 where $K \cap \bar{k} = k$ whose morphisms are k -field homomorphisms.

An important notion for curves is the ramification at a point:

Definition 0.2.3: Ramification Index

Let $\varphi : C_1 \rightarrow C_2$ be a non-constant map of smooth curves, and $p \in C_1$ a point. Then the *ramification index* of φ at p is given by:

$$e_\varphi(p) = \nu_p(\varphi^*(t_{\varphi(p)}))$$

where $t_{\varphi(p)} \in k(C_2)$ is the uniformizer at $\varphi(p)$. If $e_\varphi(p) = 1$, we say that p is unramified, and ramified otherwise (i.e. if $e_\varphi(p) > 1$ ^a). If φ is unramified at every point, we shall say that φ is *unramified*.

^aThe ramification will always be positive

The following facts about the ramification can be found in exercises in [8] or in [3, chapter 7]:

Proposition 0.2.4: Properties of Ramification Index

Let $\varphi : C_1 \rightarrow C_2$ be a non-constant map of smooth curves. Then:

- for every $q \in C_2$,

$$\sum_{p \in \varphi^{-1}(q)} e_\varphi(p) = \deg \varphi$$

- For all but finitely many $q \in C_2$,

$$\#\varphi^{-1}(q) = \deg_s \varphi$$

- If $\psi : C_2 \rightarrow C_2$ is another non-constant map of smooth curves, then for each $p \in C_1$

$$e_{\psi \circ \varphi}(p) = e_\varphi(p) \cdot e_\psi(\varphi(p))$$

Proof:

see [3]

By combining the first two bullet points, we get that a map $\varphi : C_1 \rightarrow C_2$ is unramified if and only if for all $q \in C_2$:

$$\#\varphi^{-1}(q) = \deg(\varphi)$$

Example 0.2.5: Ramification Index

Consider the map

$$\varphi : \mathbb{P}_k^1 \rightarrow \mathbb{P}_k^1 \quad [x : y] \mapsto [x^3(x - y)^2 : y^5]$$

Verify that φ is ramified at $[0 : 1]$ and $[1 : 1]$ where:

$$e_\varphi([0 : 1]) = 3 \quad \text{and} \quad e_\varphi([1 : 1]) = 2$$

We see that $\deg \varphi = 5$.

0.3 Frobenius Map

The Frobenius map is a natural map in finite characteristic (the galois group of finite extensions of finite fields is generated by the Frobenius automorphism, but this is not necessarily the case for non-finite fields! It is at least always injective). Let $\text{char}(k) > 0$ and $q = p^r$. Then for any $f \in k[x]$, let $f^{(q)} \in k[x]$ represent the polynomial where each coefficient of f was raised to the power of q (which is still in $k[x]$ as the Frobenius automorphism was applied to each element). Consider then $C^{(q)}/k$ be the curve generated by:

$$I(C^{(q)}) = \langle f^{(q)} : f \in I(C) \rangle$$

Then the natural map:

$$\varphi : C \rightarrow C^{(q)} \quad \varphi([x_0 : \cdots : x_n]) = [x_0^q : \cdots : x_n^q]$$

is a regular morphism, since for every $\varphi(P)$:

$$\begin{aligned} f^{(q)}(\varphi(P)) &= f^{(q)}(x_0^q, \dots, x_n^q) \\ &= (f(x_0, \dots, x_n))^q & \text{char}(k) = p \\ &= 0^q = 0 & f(P) = 0 \end{aligned}$$

Proposition 0.3.1: Properties of Frobenius Morphism

Let k be a (perfect) field of characteristic $p > 0$, $q = p^r$, and C/k a curve. Let $\varphi : C \rightarrow C^{(q)}$ be the q th power Frobenius morphism. Then^a

1. $\varphi^*(K^{(q)}) = K(C)^q = \{f^q \mid f \in K(C)\}$
2. φ is purely inseparable
3. $\deg \varphi = q$

^aNote that if k were not perfect, then (2) and (3) remain unchanged, but (1) requires modification

Proof:

1.

$$\varphi^* \left(\frac{f}{g} \right) = \frac{f(x_0^q, \dots, x_n^q)}{g(x_0^q, \dots, x_n^q)}$$

Similarly, elements of $K(C)^q \subseteq K(C)$ are of the form:


$$\frac{f(x_0, \dots, x_n)^q}{g(x_0, \dots, x_n)^q}$$

As k is perfect, every k is a q th power, hence:

$$(k[x_0, \dots, x_n])^q = k[x_0^q, \dots, x_n^q]$$

giving equality

2. immediate from above and field theory (see [8, chapter 17])
3. Taking a finite extension if necessary^a take a smooth point $P \in C(k)$. Then let $t \in K(C)$ be a uniformizer at P . Then $K(C)$ is separable over $K(t)$. This gives a tower:



Recall that if L/F is a separable extension over a field of characteristic q , then $L = F(L^q)$ ([8]). We thus get that $k(C) = k(C)^q(t)$, thus by (1):

$$\deg \varphi = [k(C)^q(t) : k(C)^q]$$

Now, $t^q \in k(C)^q$, so in order to prove $\deg(\varphi) = q$, we must show $t^{q/p} \notin k(C)^q$. If $t^{q/p} = f^q$ for some $f \in k(C)$, then

$$\frac{p}{q} = \text{ord}_P(t^{q/p}) = q \text{ord}_P(f)$$

which is impossible: $\text{ord}_P(f)$ must be an integer, completing the proof.

^aOr using the fact that smooth points form an open dense subset

Corollary 0.3.2: Factoring Morphisms Over Positive Characteristic

Let $\psi : C_1 \rightarrow C_2$ be a map of (smooth) curves over a field of characteristic $p > 0$. Then φ factors as

$$C_1 \xrightarrow{\varphi} C_1^{(q)} \xrightarrow{\lambda} C_2$$

where $q = \deg_i(\psi)$, φ is the q -th power Frobenius map, and λ is a separable map.

Proof:

let K be the separable closure of $\psi^*(k(C_2)) \subseteq k(C_1)$. Then $k(C_1)/K$ is purely inseparable of degree q , hence $k(C_1)^1 \subseteq K$. Then by proposition 0.3.1

$$k(C_1)^q = \varphi^*(k(C_1^{(q)})) \quad \text{and} \quad k[k(C_1) : \varphi^*(k(C_1)^{(q)})] = q$$

By comparing degrees, it is immediate that $K\varphi^*(C_1^{(q)})$. Thus, there is a tower of fields:

$$k(C_1)/\varphi^*(k(C_1)^{(q)})/\psi^*(k(C_2))$$

and hence the correspondence morphisms

$$\begin{array}{ccccc}
 & & \psi & & \\
 C_1 & \xrightarrow{\varphi} & C_1^{(q)} & \xrightarrow{\lambda} & C_2
 \end{array}$$

as we sought to show.

0.4 Divisors

Let C/k be a curve. Then its divisor group $\text{Div}(C)$ is the free-abelian group generated by the points of C , i.e. any $D \in \text{Div}(C)$ is a finite sum:

$$\sum_{p \in C} n_p(p)$$

for some finite choices of points $p \in C$ and integers $n_p \in \mathbb{Z}$. The degree of $D \in \text{Div}(C)$ is the sum of the n_p . The divisors of degree zero over a curve C , denoted $\text{Div}^0(C)$, are all divisors whose degree is 0. Note that $\text{Gal}_k(\bar{k})$ acts on $\text{Div}(C)$ (resp. $\text{Div}^0(C)$) in the natural way:

$$D^\sigma = \sum_{p \in C} n_p(p^\sigma)$$

We would say that D is defined over k if $D^\sigma = D$ for all $\sigma \in \text{Gal}_k(\bar{k})$. The group of divisors over k is defined to be $\text{Div}_k(C)$ (and $\text{Div}_k^0(C)$ for the degree 0 terms).

Assume now that C is smooth. Then for each $f \in \bar{k}(C)^*$, we can define at each point $p \in C$ the order $\text{ord}_p(f)$. We thus can define:

$$\text{div}(f) = \sum_{p \in C} \text{ord}_p(f)(p)$$

which is a finite sum as f would have zero on all but finitely many points. Such a divisor is called a *principal divisor*. It should be checked that if $\sigma \in \text{Gal}_k(\bar{k})$, then:

$$\text{div}(f^\sigma) = (\text{div}(f))^\sigma$$

and hence $\text{div}(f) \in \text{Div}_k(C)$. As ν_p is a valuation, we can define the group homomorphism:

$$\text{Div} : \bar{k}(C) \rightarrow \text{Div}(C)$$

This can be thought as being analogous as sending an element $a \in F/\mathbb{Q}$ of a number field to its principal fractional ideal $(a) \in \mathcal{I}_{F/\mathbb{Q}}$. The principal divisors form an abelian subgroup, label it $\text{CP}(C)$, and hence we can define the *divisor class group* or *Picard Group* as:

$$\text{Pic}(C) = \frac{\text{Div}(C)}{\text{CP}(C)}$$

We shall denote $\text{Pic}_K(C) \leq \text{Pic}(C)$ the subgroup fixed by $\text{Gal}_k(\bar{k})$. As we are assuming C is smooth, we have the following from [3]:

- $\text{div}(f) = 0$ if and only if $f \in \bar{k}^*$: if we assume $\text{div}(f) = 0$, it has no poles so $f : C \rightarrow \mathbb{P}_{\bar{k}}^1$ is not surjective, hence constant, hence $f \in \bar{k}$. The converse is immediate.
- $\deg(\text{div}(f)) = 0$: This is a classical result stating that the number of poles and roots must match.

- If C is a curve, $\text{Pic}(C) \cong \mathbb{Z}$ if and only if $C \rightarrow \mathbb{P}^1$ is an isomorphism. Indeed, the if $D = 0$, then it must be that $D = \text{div}(f)$ namely since $D = \sum n_p(p)$, writing $p = [a_p : b_p]$ we get that:

$$D = \text{div} \left(\prod (b_p x - a_p y)^{n_p} \right)$$

which from this we see the degree map gives $\deg : \text{Pic}(C) \cong \mathbb{Z}$. Conversely, if $\text{Pic}(C) \cong \mathbb{Z}$ (this has to do with the decomposition of $\text{Pic}(C) \cong \text{Pic}^0(C) \oplus \mathbb{Z}$ where we thus deduce $\text{Pic}^0(C) \cong 0$ and get that the genus must be zero giving us a rational equivalence, which we may raise to an isomorphism)

Example 0.4.1: Divisor of Elliptic Curve

For an elliptic curve $y^2 = (x - e_1)(x - e_2)(x - e_3)$, let $p_i = (e_i, 0)$, and let p_∞ be the point at infinity. Then:

$$\begin{aligned} \text{div}(x - e_i) &= 2(p_i) - 2(p_\infty) \\ \text{div}(y) &= (p_1) + (p_2) + (p_3) - 3(p_\infty) \end{aligned}$$

This shall soon tie to a natural group structure we can put on divisors, see chapter 1.

We would be remiss if the connection between the Picard group and Class group from number theory was not mentioned. Let $\text{Div}^0(C) \subseteq \text{Div}(C)$ be the collection of degree 0 divisors. Then define $\text{Pic}^0(C)$ to be the quotient of $\text{Div}^0(C)$ by the degree 0 principal divisors, which as pointed out earlier is given by $\bar{k}(C)$. We thus get a natural exact sequence:

$$1 \rightarrow \bar{k} \hookrightarrow \bar{k}(C) \xrightarrow{\text{Div}} \text{Div}^0(C) \twoheadrightarrow \text{Pic}^0(C) \rightarrow 0$$

which mirrors the following canonical exact sequence in number theory:

$$1 \rightarrow (\text{units of number ring}) \rightarrow k^* \rightarrow (\text{frac. ideals}) \rightarrow (\text{ideal class group}) \rightarrow 1$$

Moving onto morphisms, if $\varphi : C_1 \rightarrow C_2$ is a non-constant map of smooth curves, from φ^* and φ_* we can define:

$$\begin{aligned} \varphi^* : \text{Div}(C_2) &\rightarrow \text{Div}(C_1) & \varphi_* : \text{Div}(C_1) &\rightarrow \text{Div}(C_2) \\ (q) &\mapsto \sum_{p \in \varphi^{-1}(q)} (e_\varphi(p)) (p) & (p) &\mapsto (\varphi(p)) \end{aligned}$$

If $f \in \bar{k}(C)^*$ which is associated to $f : C \rightarrow \mathbb{P}^1$, then we get that:

$$\text{div}(f) = f^*((0) - (\infty))$$

The following basic properties from [3] ought to be known or reviewed:

Proposition 0.4.2: Properties of Divisors

1. $\deg(\varphi^* D) = (\deg \varphi)(\deg D)$ for all $D \in \text{Div}(C_2)$.
2. $\varphi^*(\text{div } f) = \text{div}(\varphi^* f)$ for all $f \in \overline{K}(C_2)^*$.
3. $\deg(\varphi_* D) = \deg D$ for all $D \in \text{Div}(C_1)$.
4. $\varphi_*(\text{div } f) = \text{div}(\varphi_* f)$ for all $f \in \overline{K}(C_1)^*$.
5. $\varphi_* \circ \varphi^*$ acts as multiplication by $\deg \varphi$ on $\text{Div}(C_2)$.
6. φ^* and φ_* induce maps on the degree 0 Picard Groups.
7. If $\psi : C_2 \rightarrow C_3$ is another such map, then

$$(\psi \circ \varphi)^* = \varphi^* \circ \psi^* \quad \text{and} \quad (\psi \circ \varphi)_* = \psi_* \circ \varphi_*.$$

0.5 Differentials

Differentials on curves shall play the same role as differential forms from calculus, further be useful for determining when an algebraic map is separable (recall that there is a derivative condition to check when a polynomial creates a separable field extension?). The main difference between the differential-geometric case and the current algebraic case is that we may define differentials on non-smooth curves. This shall be a review of [3, chapter 5.5]

Definition 0.5.1: Differential Forms On A Curve

Let C be a (not necessarily smooth) curve. Then the (meromorphic) *differential forms* on C , denoted Ω_C , is the $\overline{k}(C)$ -vector-space generated by symbols of the form dx for each $x \in \overline{k}(C)$ such that:

1. $d(x + y) = dx + dy$
2. $d(xy) = xdy + dx \cdot y$
3. $da = 0$ for all $a \in \overline{k}$

If $\varphi : C_1 \rightarrow C_2$ is a non-constant map of curves, then φ^* induces a map:

$$\varphi^* : \Omega_{C_2} \rightarrow \Omega_{C_1} \quad \varphi^* \left(\sum_i f_i dx_i \right) = \sum_i (\varphi^* f_i) d(\varphi^* x_i)$$

Proposition 0.5.2: Properties of Differentials

Let C be a curve.

1. Ω_C is a 1-dimensional $\overline{K}(C)$ -vector space.
2. Let $x \in \overline{K}(C)$. Then dx is a $\overline{K}(C)$ -basis for Ω_C if and only if $\overline{K}(C)/\overline{K}(x)$ is a finite separable extension.
3. Let $\varphi : C_1 \rightarrow C_2$ be a nonconstant map of curves. Then φ is separable if and only if the map

$$\varphi^* : \Omega_{C_2} \longrightarrow \Omega_{C_1}$$

is injective (equivalently, nonzero).

Proof:

1. See [3]
2. First, note that for dx to be a basis, x must be transcendental over $\overline{K}$ (otherwise $dx = 0$). For the $\Rightarrow$ direction, suppose dx is a $\overline{K}(C)$ -basis for Ω_C . This means every differential form can be written as $f \cdot dx$ for some $f \in \overline{K}(C)$. If $\overline{K}(C)/\overline{K}(x)$ were inseparable, then there would exist an element $y \in \overline{K}(C)$ such that $dy = 0$ (in characteristic $p > 0$, this happens when $y = z^p$ for some z). But then dy cannot be expressed as $f \cdot dx$ for any f , contradicting our assumption that dx is a basis. Therefore, $\overline{K}(C)/\overline{K}(x)$ must be separable. Conversely, suppose $\overline{K}(C)/\overline{K}(x)$ is a finite separable extension. Let $\overline{K}(C) = \overline{K}(x)[y]/(P(y))$ where P is the minimal polynomial of y over $\overline{K}(x)$. Since the extension is separable, $P'(y) \neq 0$. For any element $f \in \overline{K}(C)$, we can write $df = \frac{\partial f}{\partial x} dx + \frac{\partial f}{\partial y} dy$. But from the relation $P(y) = 0$, we get $P'(y)dy + \sum_i \frac{\partial P}{\partial a_i} da_i = 0$ where a_i are the coefficients of P . Since $P'(y) \neq 0$, we can solve for dy in terms of dx . This means any differential form can be expressed as $f \cdot dx$ for some $f \in \overline{K}(C)$, making dx a basis for Ω_C .
3. Using (1) and (2), choose $y \in \overline{K}(C_2)$ such that $\Omega_{C_2} = \overline{K}(C_2) dy$ and such that $\overline{K}(C_2)/\overline{K}(y)$ is a separable extension. Note that $\varphi^* \overline{K}(C_2)$ is then separable over $\varphi^* \overline{K}(y) = \overline{K}(\varphi^* y)$. Then:

$$\begin{aligned} \varphi^* \text{ is injective} &\iff d(\varphi^* y) \neq 0 \\ &\iff d(\varphi^* y) \text{ is a basis for } \Omega_{C_1} \text{ (from (a))}, \\ &\iff \overline{K}(C_1)/\overline{K}(\varphi^* y) \text{ is separable (from (b))}, \\ &\iff \overline{K}(C_1)/\varphi^* \overline{K}(C_2) \text{ is separable,} \end{aligned}$$

where the last equivalence follows because $\varphi^* \overline{K}(C_2)/\overline{K}(\varphi^* y)$ is separable.

Example 0.5.3: Differential On Curves

1. For the projective line $\mathbb{P}_k^1$, the function field is $k(t)$ where t is a coordinate on the affine line.

Any rational function on $\mathbb{P}_k^1$ can be written as $f(t) = \frac{p(t)}{q(t)}$ where p and q are polynomials. The differential forms on $\mathbb{P}_k^1$ are of the form $\omega = f(t)dt$ where $f(t) \in k(t)$.

Since every differential form can be expressed in terms of dt , it forms a basis for $\Omega_{\mathbb{P}_k^1}$. Therefore, $\Omega_{\mathbb{P}_k^1}$ is a 1-dimensional $k(t)$ -vector space.

Consider the differential form $\omega = \frac{t^2+1}{t-3}dt$ on $\mathbb{P}_k^1$. This is expressed in terms of the basis element dt multiplied by the rational function $\frac{t^2+1}{t-3} \in k(t)$.

For the circle defined by the equation $x^2 + y^2 = 1$, we examine its function field and differential forms.

2. From the equation $x^2 + y^2 = 1$, we can parameterize the circle using the rational parameterization:

$$\begin{aligned} x &= \frac{1-t^2}{1+t^2} \\ y &= \frac{2t}{1+t^2} \end{aligned}$$

where $t = \tan(\theta/2)$ and θ is the angle parameter.

The function field $k(C) = k(x, y)$ with the relation $x^2 + y^2 = 1$ is isomorphic to $k(t)$.

Taking differentials of the equation $x^2 + y^2 = 1$, we get:

$$\begin{aligned} 2x dx + 2y dy &= 0 \\ \Rightarrow dy &= -\frac{x}{y}dx \quad (\text{when } y \neq 0) \end{aligned}$$

This means any differential form can be expressed in terms of dx alone. Therefore, Ω_C is a 1-dimensional $k(C)$ -vector space with basis $\{dx\}$.

Using the parameterization, we can express the basic differential forms in terms of dt :

$$\begin{aligned} dx &= \frac{d}{dt} \left(\frac{1-t^2}{1+t^2} \right) dt = \frac{-4t}{(1+t^2)^2} dt \\ dy &= \frac{d}{dt} \left(\frac{2t}{1+t^2} \right) dt = \frac{2(1-t^2)}{(1+t^2)^2} dt \end{aligned}$$

Any differential form $\omega = f(x, y)dx + g(x, y)dy$ can be rewritten as $\omega = h(t)dt$ for some rational function $h(t)$, confirming that Ω_C is 1-dimensional with basis $\{dt\}$.

3. For the singular curve defined by the equation $x^2 - y^3 = 0$ (known as a cusp), the curve C has a singularity at the origin, but we can still analyze its function field. From the equation $x^2 = y^3$, we can parameterize the curve as:

$$\begin{aligned} x &= t^3 \\ y &= t^2 \end{aligned}$$

where t is a parameter.

Using our parameterization, we can express the basic differential forms in terms of dt :

$$\begin{aligned} dx &= \frac{d}{dt}(t^3)dt = 3t^2dt \\ dy &= \frac{d}{dt}(t^2)dt = 2tdt \end{aligned}$$

This shows that any differential form can be expressed in terms of dt , making $\{dt\}$ a basis for Ω_C over $k(t) \cong k(C)$.

Consider the differential form $\omega = xdx + ydy$. Using our parameterization and the relations above:

$$\begin{aligned} \omega &= t^3(3t^2dt) + t^2(2tdt) \\ &= (3t^5 + 2t^3)dt \end{aligned}$$

demonstrating how any differential form can be written as $h(t)dt$ for some rational function $h(t)$.

Proposition 0.5.4: Differentials and Order

Let C be a curve, let $p \in C$, and let $t \in \bar{k}(C)$ be a uniformizer at p .

1. For every $\omega \in \Omega_C$ there exists a unique function $g \in \bar{k}(C)$, depending on ω and t , satisfying

$$\omega = gdt.$$

We denote g by ω/dt .

2. Let $f \in \bar{k}(C)$ be regular at p . Then df/dt is also regular at p .
3. Let $\omega \in \Omega_C$ with $\omega \neq 0$. The quantity

$$\text{ord}_p(\omega/dt)$$

depends only on ω and p , independent of the choice of uniformizer t . We call this value the order of ω at p and denote it by $\text{ord}_p(\omega)$.

4. Let $x, f \in \bar{k}(C)$ with $x(p) = 0$, and let $P = \text{char}k$. Then

$$\begin{aligned} \text{ord}_p(fdx) &= \text{ord}_p(f) + \text{ord}_p(x) - 1, & \text{if } P = 0 \text{ or } P \nmid \text{ord}_p(x) \\ \text{ord}_p(fdx) &\geq \text{ord}_p(f) + \text{ord}_p(x), & \text{if } P > 0 \text{ and } P \mid \text{ord}_p(x) \end{aligned}$$

5. Let $\omega \in \Omega_C$ with $\omega \neq 0$. Then

$$\text{ord}_p(\omega) = 0 \quad \text{for all but finitely many } p \in C.$$

Proof:

see [21]

From this, we can define:

Definition 0.5.5: Divisor of Differential

Let $\omega \in \Omega_C$. Then the *divisor associated to ω* is:

$$\text{div} = \sum_{p \in C} \text{ord}_p(\omega)(p) \in \text{Div}(C)$$

the differential $\omega \in \Omega_C$ is *regular* (or *holomorphic*) if:

$$\text{ord}_p(\omega) \geq 0 \quad \forall p \in C$$

and *non-vanishing* if

$$\text{ord}_p(\omega) \leq 0 \quad \forall p \in C$$

and thus it is holomorphic and non-vanishing if $\text{ord}_p(\omega) = 0$ for all $p \in C$.

Note that if $\omega_1, \omega_2 \in \Omega_C$ are nonzero differentials, then proposition 0.5.2(1) implies that there must exist a $f \in \bar{k}(C)^*$ such that $\omega_1 = f\omega_2$. By proposition 0.5.4 this implies:

$$\text{div}(\omega_1) = \text{div}(f) + \text{div}(\omega_2)$$

In particular:

$$\text{div}(\omega_1) - \text{div}(\omega_2) = \text{div}(f)$$

showing that divisor differ up to a principal divisor! We thus have the following well-defined notion:

Definition 0.5.6: Canonical Divisor

Let C be a curve. Then the *canonical divisor* on C is the image $\text{div}(\omega)$ in $\text{Pic}(C)$ for any nonzero differential $\omega \in \Omega_C$. Any divisor in this class is called the *canonical divisor*.

Example 0.5.7: Canonical Divisors

1. Let $C = \mathbb{P}^1$, and let t be a coordinate function. Then for all $\alpha \in \bar{k}$ the function $t - \alpha$ is a uniformizer at α and so:

$$\text{ord}_\alpha(dt) = \text{ord}_\alpha(d(t - \alpha)) = 0$$

However, at $\infty \in \mathbb{P}^1$, we need to use a function like $1/t$ as our uniformizer, hence:

$$\text{ord}_\infty(dt) = \text{ord}_\infty\left(-t^2 d\left(\frac{1}{t}\right)\right) = -2$$

Thus,

$$\text{div}(dt) = -2(\infty)$$

Note that this shows that dt is *not* holomorphic. Then, for any $\omega \in \Omega_{\mathbb{P}^1}$ by proposition 0.5.4 we have:

$$\deg \text{div}(\omega) = \deg \text{div}(dt) = -2$$

and so no ω can be holomorphic

2. Let $C = y^2 = (x - e_1)(x - e_2)(x - e_3)$. Letting $p_i = (e_i, 0)$, we can compute to get:

$$\operatorname{div}(dx) = (p_1) + (p_2) + (p_3) - 3(p_\infty)$$

where the $(p_1), (p_2), (p_3)$ come from how $dx = d(x - e_i)$ and the $-3(p_\infty)$ comes from:

$$dx = -x^2 d\left(\frac{1}{x}\right) \Rightarrow \operatorname{ord}_{p_\infty}(dx) = -2 - 1 = -3$$

Thus, combining with $\operatorname{div}(y) = (p_1) + (p_2) + (p_3) - 3(p_\infty)$, we can construct a non-vanishing holomorphic form:

$$\operatorname{div}\left(\frac{dx}{y}\right) = 0$$

0.6 Riemann-Roch Theorem

So what is the invariant information that the canonical divisor provides? The *Riemann-Roch* theorem provides this information: its *degree* shall be invariant, and shall be deterministically be tied to the *genus* of the underlying curve. The Riemann Roch Theorem gives us a way to use algebraic tools to measure the genus of a curve.

Definition 0.6.1: Effective Divisor

Let C be a curve and $D = \sum_p n_p(p)$ be a divisor. Then D is an *effective divisor* or *positive divisor* if $n_p \geq 0$ for every $p \in C$. In this case, we write

$$D \geq 0$$

If D_1, D_2 are two divisors, we write $D_1 \geq D_2$ is the divisor $(D_1 - D_2)$ is positive.

Intuitively, every $f \in \mathcal{L}(D)$ can have poles only at points where D has negative coefficients (that is, there is a bound on the poles). As an example, let $f \in \bar{k}(C)^*$ be a function that is regular at all points except one, where it has a pole of degree at most n at p . Then we can write:

$$\operatorname{div}(f) \geq -n(p)$$

Definition 0.6.2: Riemann-Roch Space For Divisor

Let $D \in \operatorname{Div}(C)$ be a divisor. Then the *Riemann Roch space* associated to D is the space:

$$\mathcal{L}(D) = \{f \in \bar{k}(C)^* \mid \operatorname{div}(f) \geq -D\} \cup \{0\}$$

Furthermore, label:

$$\ell(D) = \dim_{\bar{k}} \mathcal{L}(D)$$

We can think of $\mathcal{L}(D)$ as being the collection of (meromorphic) functions that are allowed to have poles at the divisor D .

Proposition 0.6.3: Properties of Riemann-Roch Space

Let $D \in \text{Div}(C)$.

1. If $\deg D < 0$, then

$$\mathcal{L}(D) = \{0\} \quad \text{and} \quad \ell(D) = 0.$$

2. $\mathcal{L}(D)$ is a finite-dimensional $\bar{k}$ -vector space.

3. If $D' \in \text{Div}(C)$ is linearly equivalent to D , then

$$\mathcal{L}(D) \cong \mathcal{L}(D'), \quad \text{and so} \quad \ell(D) = \ell(D').$$

Proof:

see [21]

From these, we can interpret $\mathcal{L}(K_C)$ where $K_C \in \text{Div}(C)$ is the canonical divisor (so $K_C = \text{div}(\omega)$ for some appropriate choice of ω). By definition we have that $f \in \mathcal{L}(K_C)$ is:

$$\text{div}(f) \geq -\text{div}(\omega)$$

Thus $\text{div}(f\omega) \geq 0$, i.e. $f\omega$ is holomorphic. Then certainly by definition if $f\omega$ is holomorphic then $\text{div}(f\omega) \geq 0$ so $f \in \mathcal{L}(K_C)$. As K_C is canonical, every differential form on C can be represented as $f\omega$, hence we have a $\bar{k}$ -vector space isomorphism:

$$\mathcal{L}(K_C) \cong \{\omega \in \Omega_C \mid \omega \text{ is holomorphic}\}$$

The dimension $\ell(K_C)$ forms an important invariant of our curve in the *Riemann-Roch equation*:

Theorem 0.6.4: Riemann-Roch Theorem

let C be a curve and K_C the canonical divisor on C . Then there exists an integer $g \geq 0$, called the *genus* of C such that for every divisor $D \in \text{Div}(C)$:

$$\ell(D) - \ell(K_C - D) = \deg(D) - g + 1$$

Proof:

see [3, chapter 7.1]

Corollary 0.6.5: Canonical Divisor and Genus

Let K_C be the canonical divisor for a curve C . Then:

1. $\ell(K_C) = g$
2. $\deg(K_C) = 2g - 2$
3. if $\deg D > 2g - 2$ then $\ell(D) = \deg(D) - g + 1$

Proof:

immediate

Example 0.6.6: Riemann-Roch Application

1. Let $C = \mathbb{P}_k^1$. Then we saw in example 0.5.7 that there are no holomorphic differentials on C , and hence $\ell(K_C) = 0$. Thus, we get from the Riemann-Roch Theorem that:

$$\ell(D) - \ell(-2(\infty) - D) = \deg(D) + 1$$

if $D \geq -1$, then:

$$\ell(D) = \deg(D) + 1$$

2. Let $C = y^2 = (x - e_1)(x - e_2)(x - e_3)$. Then in example 0.5.7 we had:

$$\operatorname{div} \left(\frac{dx}{y} \right) = 0$$

and hence by proposition 0.5.4 we have that $K_C = 0$. Thus, we can compute:

$$g = \ell(K_C) = \ell(0) = 1$$

Hence, C has genus 1. Given $D \geq 1$, we further get that:

$$\ell(D) = \deg(D)$$

Let us use this to compute some special cases of D :

- (a) Let $p \in C$ be some point. Then by the above it must be that $\ell((p)) = 1$, however $\mathcal{L}((p))$ certainly contains the constant functions, but these have no poles, hence no functions on C have a single simple point
- (b) If we take the point at infinity p_∞ , then $\ell(2(p_\infty)) = 2$. The set $\{1, x\}$ forms a basis for this set.
- (c) Similarly, for $\ell(3(p_\infty))$ and $\ell(4(p_\infty))$, we have a basis $\{1, x, y\}$ and $\{1, x, y, x^2\}$.
- (d) For $\mathcal{L}(6(p_\infty))$, note that $\{1, x, y, x^2, xy, x^3, y^2\}$ is a subset, but by our above formula we have $\ell(6(p_\infty)) = 6$. Thus, that set is *linearly dependent*, which we can also see by using the defining elliptic equation.

Next, we should take a look at non-algebraically closed fields.

Proposition 0.6.7: Riemann-Roch over a field K

Let C/K be a smooth curve and let $D \in \operatorname{Div}_K(C)$. Then $\mathcal{L}(D)$ has a basis consisting of functions in $K(C)$.

Proof:

Since D is defined over K , we have

$$f^\sigma \in \mathcal{L}(D^\sigma) = \mathcal{L}(D) \quad \text{for all } f \in \mathcal{L}(D) \text{ and all } \sigma \in G_{\bar{K}/K}$$

Thus $G_{\bar{K}/K}$ acts on $\mathcal{L}(D)$. Then by the following lemma (lemma 0.6.8), we get the desired result, completing the proof.

Lemma 0.6.8: Galois Action and Invariant Basis

Let C/K be a smooth curve and let $D \in \text{Div}_k(C)$. Then $\mathcal{L}(D)$ has a basis consisting of functions in $k(C)$

Proof:

Finally, we shall take advantage of the following theorem from [3] for the relations between maps between curves and genus:

Theorem 0.6.9: Hurwitz Theorem

Let $\varphi : C_1 \rightarrow C_2$ be a non-constant separable map of smooth curves of genus g_1, g_2 . Then:

$$2g_1 - 2 \geq \deg(\varphi)(2g_2 - 2) + \sum_{p \in C_1} (e_\varphi(p) - 1)$$

where equality holds if and only if either:

1. $\text{char}(k) = 0$, or
2. $\text{char}(k) = p > 0$ and $p \nmid e_\varphi(p)$ for all $p \in C_1$

Proof:

see [3]. (or look at [21, p.37])

Exercise 0.6.1

1. Let C be a curve and $\varphi : C \rightarrow V \subseteq \mathbb{P}_k^n$. Show that if $p \in C$ is smooth then φ is regular at p . Thus if C is smooth, φ must be regular.

Part I

Foundations: Geometry and Algebra of Elliptic Curves

here

Basics of Elliptic Curves

The reader who has completed a study of algebraic geometry possesses a formidable toolkit: schemes, sheaves, cohomology, the Riemann–Roch theorem, divisor class groups, and the general theory of algebraic curves. This chapter puts all of that machinery to work on a single class of objects — smooth projective curves of genus 1 with a marked rational point. These are *elliptic curves*, and they are, in a precise sense, the simplest algebraic curves whose geometry forces an interesting arithmetic. A curve of genus 0 over a field K is either isomorphic to $\mathbb{P}_K^1$ or has no rational points at all, and in either case its arithmetic is elementary. A curve of genus ≥ 2 has only finitely many rational points over a number field (by Faltings’ theorem). It is genus 1 — and only genus 1 — where the set of rational points can carry the structure of a finitely generated abelian group, and where the interplay between geometry and arithmetic becomes truly rich.

Our goal in this chapter is to establish the basic objects and their structure. We will use the Riemann–Roch theorem to derive the Weierstrass equation from first principles, define the discriminant and j -invariant, construct the group law via the Picard group, classify automorphisms and twists, and analyze what happens when the curve degenerates. By the end, the reader will see that an elliptic curve is simultaneously a projective algebraic curve, a one-dimensional group scheme, and (when the singular fibers are included) a bridge to the classical algebraic groups $\mathbb{G}_m$ and $\mathbb{G}_a$. Every abstract theorem the reader has learned will produce something explicit and computable here.

1.1 Elliptic Curves and Weierstrass Equations

We begin with the definition and immediately demonstrate that the abstract theory of divisors and linear systems, far from being a formalism for its own sake, produces the single most important equation in the subject.

Definition 1.1.1: Elliptic Curve

An *elliptic curve* over a field K is a pair (E, O) where E is a smooth projective curve of genus 1 over K and $O \in E(K)$ is a distinguished K -rational point. When the choice of O is understood, we write simply E/K or E .

The requirement that O be a rational point is essential and not merely a convenience. There exist smooth projective genus-1 curves over $\mathbb{Q}$ that have no rational points at all — for instance, the curve $3X^3 + 4Y^3 + 5Z^3 = 0$ in $\mathbb{P}_{\mathbb{Q}}^2$, due to Selmer, has points over every completion $\mathbb{Q}_p$ and over $\mathbb{R}$ but no global rational point. Such a curve is a *genus-1 curve* but not an elliptic curve; it is a principal homogeneous space (or *torsor*) for its Jacobian, a distinction that will become central when we study the Tate–Shafarevich group in ???. The marked point O will serve as the identity element for the group law constructed in section 1.3.

The Riemann–Roch Derivation

The Weierstrass equation is not an *ad hoc* formula imposed on elliptic curves. It emerges inevitably from the Riemann–Roch theorem, and working through this derivation is the best way to see abstract cohomological machinery produce concrete equations.

Let E/K be an elliptic curve with marked point O , and let $g = 1$ denote the genus. Recall that for a divisor D on a smooth projective curve of genus g , the Riemann–Roch theorem [3] asserts

$$\ell(D) - \ell(K_E - D) = \deg(D) - g + 1,$$

where $\ell(D) = \dim_K H^0(E, \mathcal{O}_E(D))$ is the dimension of the space of rational functions whose divisor of poles is bounded by D , and K_E is the canonical divisor. For a genus-1 curve, $\deg(K_E) = 2g - 2 = 0$, and the global sections of the canonical bundle have dimension $\ell(K_E) = g = 1$. Since $\ell(K_E) = 1$ and $\deg(K_E) = 0$, the canonical class K_E is linearly equivalent to the zero divisor, i.e., $K_E \sim 0$.

The key consequence is immediate: for any divisor D with $\deg(D) \geq 1$, we have $\deg(K_E - D) = -\deg(D) < 0$, and a divisor of negative degree has no nonzero global sections, so $\ell(K_E - D) = 0$. Therefore, the Riemann–Roch theorem simplifies to

$$\ell(D) = \deg(D) \quad \text{for all divisors } D \text{ with } \deg(D) \geq 1. \quad (1.1)$$

We now apply this to the divisors $n \cdot (O)$ for $n = 1, 2, 3, \dots$ and systematically build a basis for each Riemann–Roch space $L(n \cdot (O)) := H^0(E, \mathcal{O}_E(n \cdot (O)))$, which consists of all rational functions $f \in K(E)^*$ satisfying $\operatorname{div}(f) + n \cdot (O) \geq 0$ (i.e., f has at worst a pole of order n at O and no other poles), together with 0.

- **$n = 1$:** We have $\ell(O) = 1$. A function in $L(O)$ has at most a simple pole at O and no other poles. On a curve of genus 1, such a function must be constant (a function with a single simple pole would give an isomorphism $E \xrightarrow{\sim} \mathbb{P}^1$, contradicting $g = 1$). Thus $L(O) = K \cdot 1$, spanned by the constant function 1.
- **$n = 2$:** We have $\ell(2 \cdot (O)) = 2$. The space $L(2 \cdot (O))$ contains $L(O)$ as a subspace, so it is spanned by 1 and a new function $x \in K(E)$ that has a pole of order *exactly* 2 at O and is regular elsewhere. Such a function exists by the dimension count, and we fix a choice. The pair $\{1, x\}$ is a basis for $L(2 \cdot (O))$.

- **$n = 3$:** We have $\ell(3 \cdot (O)) = 3$. We need one new function $y \in K(E)$ with a pole of order *exactly* 3 at O . (A pole of order ≤ 2 would place y in $L(2 \cdot (O))$, which is already accounted for.) We fix such a y ; then $\{1, x, y\}$ is a basis for $L(3 \cdot (O))$.

At this point we have the two fundamental coordinate functions x and y . We continue:

- **$n = 4$:** We have $\ell(4 \cdot (O)) = 4$. The products of our existing functions give candidates for new basis elements: x^2 has a pole of order 4 at O ($\text{ord}_O(x) = -4$). Since $x^2 \notin L(3 \cdot (O))$, the set $\{1, x, y, x^2\}$ is a basis.
- **$n = 5$:** We have $\ell(5 \cdot (O)) = 5$. The function xy has pole order $2+3 = 5$ at O , so $\{1, x, y, x^2, xy\}$ is a basis.
- **$n = 6$:** We have $\ell(6 \cdot (O)) = 6$. Now consider all monomials in x and y with pole order at most 6 at O :

$$\underbrace{1}_{\text{order } 0}, \quad \underbrace{x}_{\text{order } 2}, \quad \underbrace{y}_{\text{order } 3}, \quad \underbrace{x^2}_{\text{order } 4}, \quad \underbrace{xy}_{\text{order } 5}, \quad \underbrace{y^2}_{\text{order } 6}, \quad \underbrace{x^3}_{\text{order } 6}.$$

This gives *seven* elements of the six-dimensional space $L(6 \cdot (O))$. They must therefore satisfy a nontrivial linear relation over K .

The relation involves both y^2 and x^3 (since these are the only elements of pole order exactly 6, and removing either would leave six elements that are linearly independent by a pole-order argument: the remaining six elements have distinct pole orders 0, 2, 3, 4, 5, 6, and functions with distinct pole orders at a point are linearly independent). Therefore the relation has the form

$$\alpha y^2 + \beta x^3 = (\text{linear combination of } 1, x, y, x^2, xy)$$

with $\alpha, \beta \neq 0$. By rescaling x and y (specifically, substituting $x = -(\alpha/\beta)x'$ and $y = (\alpha/\beta)y'$, and dividing the entire relation by α^3/β^2 — note that such rescaling preserves the property that the functions have pole orders 2 and 3 respectively), we may assume $\alpha = 1$ and $\beta = -1$. The relation then takes the form

$$y^2 + a_1xy + a_3y = x^3 + a_2x^2 + a_4x + a_6 \tag{1.2}$$

for some $a_1, a_2, a_3, a_4, a_6 \in K$. (The numbering of the coefficients is traditional: a_i is assigned “weight” i , where x has weight 2 and y has weight 3. The terms y^2 and x^3 each have weight 6; the term $a_i \cdot (\text{monomial of weight } 6 - i)$ accounts for each coefficient. There is no a_5 because there is no monomial in x, y of weight 1.)

Definition 1.1.2: Weierstrass Equation

An equation of the form

$$y^2 + a_1xy + a_3y = x^3 + a_2x^2 + a_4x + a_6$$

with $a_1, a_2, a_3, a_4, a_6 \in K$ is called a (*generalized*) *Weierstrass equation*. If the corresponding projective curve in $\mathbb{P}_K^2$ (given by the homogenization $Y^2Z + a_1XYZ + a_3YZ^2 = X^3 + a_2X^2Z + a_4XZ^2 + a_6Z^3$) is smooth, we say the Weierstrass equation defines an elliptic curve, and we take $O = [0 : 1 : 0]$ as the distinguished point.

The point $O = [0 : 1 : 0]$ is the unique point at infinity on the Weierstrass model. In the affine chart $Z \neq 0$, the equation is (1.2). In the affine chart $Y \neq 0$ (setting $Y = 1$), the equation becomes $Z + a_1XZ + a_3Z^2 = X^3 + a_2X^2Z + a_4XZ^2 + a_6Z^3$, and one checks that the point $(X, Z) = (0, 0)$ (which is O) lies on this curve and is a smooth point (the partial derivative with respect to Z at the origin is $1 \neq 0$). Thus O is always a smooth point of any Weierstrass model.

The derivation above proves the following:

Proposition 1.1.3: Weierstrass Form from Riemann–Roch

Every elliptic curve E/K admits a Weierstrass equation: there exist functions $x, y \in K(E)$ such that the map

$$\varphi : E \longrightarrow \mathbb{P}_K^2, \quad P \longmapsto [x(P) : y(P) : 1]$$

(extended to $\varphi(O) = [0 : 1 : 0]$) is a closed immersion whose image is the projective closure of a Weierstrass equation (1.2).

Proof:

The functions x, y are constructed as above via Riemann–Roch. The map φ is the morphism associated to the complete linear system $|3 \cdot (O)|$, which has dimension $\ell(3 \cdot (O)) - 1 = 2$ and thus maps into $\mathbb{P}_K^2$. Because $\deg(3 \cdot (O)) = 3 > 2 = 2g$, the linear system $|3 \cdot (O)|$ is very ample [3], so φ is a closed immersion. The image is a curve of degree 3 in $\mathbb{P}_K^2$ (since $\deg(\varphi^* \mathcal{O}_{\mathbb{P}^2}(1)) = \deg(3 \cdot (O)) = 3$), and the relation among the seven elements of $L(6 \cdot (O))$ shows that the image satisfies the Weierstrass equation.

This is worth pausing on. The Riemann–Roch theorem, applied to nothing more than the divisor of a single point on a genus-1 curve, has produced for us a canonical embedding into the projective plane and the explicit cubic equation governing the image.

Admissible Changes of Variables

The Weierstrass equation for a given elliptic curve is not unique: different choices of the functions x and y in the Riemann–Roch spaces will yield different coefficients a_i . We now determine exactly how much freedom these choices allow.

Any other function $x' \in L(2 \cdot (O))$ with a pole of exact order 2 at O must satisfy $x' = u^2x + r$ for some $u \in K^*$ and $r \in K$ (since x' and u^2x differ by an element of $L(O) = K$, and the leading coefficient is adjusted by u^2 to account for the freedom in normalizing the pole). Similarly, $y' \in L(3 \cdot (O))$ with pole order exactly 3 must have the form $y' = u^3y + u^2sx + t$ for some $s, t \in K$ (the u^3y term fixes the leading behavior, and the correction terms lie in $L(2 \cdot (O))$).

Definition 1.1.4: Admissible Change of Variables

An *admissible change of variables* is a substitution of the form

$$x = u^2x' + r, \quad y = u^3y' + u^2sx' + t, \quad (1.3)$$

where $u \in K^*$ and $r, s, t \in K$. Two Weierstrass equations are *isomorphic (over K)* if they are related by such a substitution.

The set of admissible changes forms a group under composition: the identity is $(u, r, s, t) = (1, 0, 0, 0)$, and one can verify that the composition of two admissible changes and the inverse of an admissible change are again admissible (with explicit formulas for the composed parameters).

One can substitute (1.3) directly into a Weierstrass equation to determine how the coefficients a_i transform. We record the result:

Proposition 1.1.5: Transformation of Weierstrass Coefficients

Under the admissible change of variables (u, r, s, t) , the coefficients of the Weierstrass equation transform as:

$$\begin{aligned} u a'_1 &= a_1 + 2s, \\ u^2 a'_2 &= a_2 - sa_1 + 3r - s^2, \\ u^3 a'_3 &= a_3 + ra_1 + 2t, \\ u^4 a'_4 &= a_4 - sa_3 + 2ra_2 - (t + rs)a_1 + 3r^2 - 2st, \\ u^6 a'_6 &= a_6 + ra_4 + r^2a_2 + r^3 - ta_3 - t^2 - rta_1. \end{aligned}$$

Proof:

This is a direct (if tedious) substitution. Replace x by $u^2x' + r$ and y by $u^3y' + u^2sx' + t$ in

$$y^2 + a_1xy + a_3y = x^3 + a_2x^2 + a_4x + a_6,$$

expand, and collect terms to write the result in the standard Weierstrass form in the variables (x', y') . For instance, the left-hand side becomes

$$(u^3y' + u^2sx' + t)^2 + a_1(u^2x' + r)(u^3y' + u^2sx' + t) + a_3(u^3y' + u^2sx' + t)$$

and collecting the y'^2 term gives $u^6(y')^2$, the $x'y'$ term gives $u^6(a_1 + 2s)(x')(y')$, and so on. Dividing through by u^6 and reading off coefficients yields the stated formulas.

Simplification to Short Weierstrass Form

When the characteristic of K is not too small, we can use admissible changes to eliminate several coefficients. The result is the short Weierstrass form, which is the most commonly used equation for elliptic curves in practice.

Step 1: Eliminating a_1 and a_3 (requires $\text{char}K \neq 2$). The left-hand side of the Weierstrass

equation,

$$y^2 + a_1xy + a_3y,$$

is a quadratic in y . If $\text{char}K \neq 2$, we can complete the square. Set $s = 0$, $r = 0$, and $t = 0$ in the admissible change, and instead apply the substitution

$$y \mapsto \frac{1}{2}(y - a_1x - a_3), \quad (1.4)$$

which corresponds to the admissible change $(u, r, s, t) = (1, 0, 0, 0)$ followed by the replacement $y_{\text{old}} = \frac{1}{2}(y_{\text{new}} - a_1x - a_3)$, or equivalently $y_{\text{new}} = 2y_{\text{old}} + a_1x + a_3$.

Let us carry out this substitution explicitly. Writing $Y = 2y + a_1x + a_3$ (so that $y = (Y - a_1x - a_3)/2$), the left-hand side of the original equation becomes:

$$y^2 + a_1xy + a_3y = \left(\frac{Y - a_1x - a_3}{2}\right)^2 + a_1x \cdot \frac{Y - a_1x - a_3}{2} + a_3 \cdot \frac{Y - a_1x - a_3}{2}.$$

Expanding the square in the first term:

$$\left(\frac{Y - a_1x - a_3}{2}\right)^2 = \frac{Y^2 - 2Y(a_1x + a_3) + (a_1x + a_3)^2}{4}.$$

Adding the remaining two terms (with common denominator 4):

$$\begin{aligned} &= \frac{1}{4} \left[Y^2 - 2Y(a_1x + a_3) + (a_1x + a_3)^2 + 2a_1xY - 2a_1x(a_1x + a_3) \right. \\ &\quad \left. + 2a_3Y - 2a_3(a_1x + a_3) \right] \\ &= \frac{1}{4} \left[Y^2 + (a_1x + a_3)^2 - 2(a_1x + a_3)^2 \right] \\ &= \frac{1}{4} \left[Y^2 - (a_1x + a_3)^2 \right] \\ &= \frac{1}{4} \left[Y^2 - a_1^2x^2 - 2a_1a_3x - a_3^2 \right]. \end{aligned}$$

In the simplification above, the key cancellation is that $-2Y(a_1x + a_3) + 2a_1xY + 2a_3Y = 0$, which is precisely why completing the square works. Setting this equal to the right-hand side $x^3 + a_2x^2 + a_4x + a_6$ and multiplying through by 4:

$$Y^2 = 4x^3 + (a_1^2 + 4a_2)x^2 + (2a_1a_3 + 4a_4)x + (a_3^2 + 4a_6). \quad (1.5)$$

This motivates the standard auxiliary quantities.

Definition 1.1.6: The b -Invariants

Given a Weierstrass equation with coefficients a_1, a_2, a_3, a_4, a_6 , we define

$$\begin{aligned} b_2 &= a_1^2 + 4a_2, \\ b_4 &= a_1a_3 + 2a_4, \\ b_6 &= a_3^2 + 4a_6, \\ b_8 &= a_1^2a_6 - a_1a_3a_4 + 4a_2a_6 + a_2a_3^2 - a_4^2. \end{aligned}$$

The quantity b_8 does not appear in (1.5), but it satisfies the useful relation

$$4b_8 = b_2b_6 - b_4^2, \quad (1.6)$$

which one verifies by direct expansion. This relation means b_8 is determined by b_2, b_4, b_6 , but having a separate name for it simplifies many formulas that follow (particularly the discriminant).

In terms of the b -invariants, equation (1.5) reads

$$Y^2 = 4x^3 + b_2x^2 + 2b_4x + b_6. \quad (1.7)$$

Step 2: Eliminating the x^2 term (requires additionally $\text{char}K \neq 3$). Equation (1.7) is a depressed form in Y but not in x : the cubic on the right still has an x^2 term. If $\text{char}K \neq 3$, we eliminate it by completing the cube. We substitute

$$x \mapsto \frac{1}{12}(x - 3b_2), \quad (1.8)$$

or equivalently $x_{\text{new}} = 12x_{\text{old}} + 3b_2$, which shifts the variable to center the cubic. Simultaneously, to clear the factor of 4 on the leading cubic term, we substitute $Y \mapsto Y/108$, or equivalently $Y_{\text{new}} = 108Y_{\text{old}} + \dots$. It is cleaner to carry this out in two stages.

First, in (1.7), substitute $x = (X - 3b_2)/36$ (and leave Y unchanged for now):

$$\begin{aligned} 4x^3 + b_2x^2 &= 4\left(\frac{X - 3b_2}{36}\right)^3 + b_2\left(\frac{X - 3b_2}{36}\right)^2 \\ &= \frac{4(X - 3b_2)^3}{36^3} + \frac{b_2(X - 3b_2)^2}{36^2}. \end{aligned}$$

Expanding $(X - 3b_2)^3 = X^3 - 9b_2X^2 + 27b_2^2X - 27b_2^3$ and $(X - 3b_2)^2 = X^2 - 6b_2X + 9b_2^2$:

$$= \frac{4(X^3 - 9b_2X^2 + 27b_2^2X - 27b_2^3)}{46656} + \frac{b_2(X^2 - 6b_2X + 9b_2^2)}{1296}.$$

Converting to a common denominator of $46656 = 36^3$:

$$\begin{aligned} &= \frac{4X^3 - 36b_2X^2 + 108b_2^2X - 108b_2^3 + 36b_2X^2 - 216b_2^2X + 324b_2^3}{46656} \\ &= \frac{4X^3 - 108b_2^2X + 216b_2^3}{46656} \\ &= \frac{X^3}{11664} - \frac{b_2^2X}{432} + \frac{b_2^3}{216}. \end{aligned}$$

The remaining linear and constant terms in (1.7) transform as:

$$2b_4x + b_6 = \frac{2b_4(X - 3b_2)}{36} + b_6 = \frac{b_4X}{18} - \frac{b_2b_4}{6} + b_6.$$

Adding everything together, the right-hand side of (1.7) becomes:

$$\frac{X^3}{11664} - \frac{b_2^2}{432}X + \frac{b_2^3}{216} + \frac{b_4}{18}X - \frac{b_2b_4}{6} + b_6.$$

Collecting the X -coefficient:

$$-\frac{b_2^2}{432} + \frac{b_4}{18} = \frac{-b_2^2 + 24b_4}{432} = -\frac{b_2^2 - 24b_4}{432}.$$

Collecting the constant term:

$$\frac{b_2^3}{216} - \frac{b_2b_4}{6} + b_6 = \frac{b_2^3 - 36b_2b_4 + 216b_6}{216}.$$

This motivates the c -invariants:

Definition 1.1.7: The c -Invariants

$$\begin{aligned} c_4 &= b_2^2 - 24b_4, \\ c_6 &= -b_2^3 + 36b_2b_4 - 216b_6. \end{aligned}$$

With these definitions, the right-hand side equals

$$\frac{X^3}{11664} - \frac{c_4}{432}X - \frac{c_6}{216}.$$

The left-hand side is Y^2 , so we have

$$Y^2 = \frac{X^3}{11664} - \frac{c_4}{432}X - \frac{c_6}{216}.$$

Finally, substitute $Y = \tilde{Y}/108$ (i.e., $\tilde{Y} = 108Y$) to obtain:

$$\frac{\tilde{Y}^2}{11664} = \frac{X^3}{11664} - \frac{c_4}{432}X - \frac{c_6}{216}.$$

Multiplying through by $11664 = 108^2$:

$$\tilde{Y}^2 = X^3 - 27c_4X - 54c_6.$$

Renaming $(\tilde{Y}, X)$ back to (y, x) and writing $A = -27c_4$ and $B = -54c_6$, we arrive at the *short Weierstrass form*:

$$\boxed{y^2 = x^3 + Ax + B.} \tag{1.9}$$

To summarize the relationship between the original coefficients and the short form: when $\text{char}K \neq 2, 3$, every Weierstrass equation can be brought to the form $y^2 = x^3 + Ax + B$ via an admissible change of variables, with

$$A = -27c_4 = -27(b_2^2 - 24b_4), \quad B = -54c_6 = -54(-b_2^3 + 36b_2b_4 - 216b_6).$$

1.1.8 Remark: Conventions in the Literature Different authors use different normalizations. Silverman [21] writes the short form as $y^2 = x^3 - 27c_4x - 54c_6$, which matches our (1.9). Some authors instead normalize so that $y^2 = x^3 - (c_4/48)x - (c_6/864)$, which arises from a slightly different rescaling (avoiding the factors of 27 and 54 at the cost of denominators). For concrete arithmetic over $\mathbb{Q}$, the form $y^2 = x^3 + Ax + B$ with $A, B \in \mathbb{Z}$ is almost universally preferred, and any Weierstrass equation over $\mathbb{Q}$ with $a_1, \dots, a_6 \in \mathbb{Z}$ can be brought to such a form after clearing denominators.

1.1.9 Remark: Characteristic 2 and 3 In characteristics 2 and 3, the simplifications above fail: we cannot divide by 2 to complete the square, nor by 3 to complete the cube. The generalized Weierstrass equation (1.2) is the correct form to use over any field, and all of the theory we develop (the group law, the discriminant, isogeny theory) works with it. In characteristic 2, the best one can achieve is either $y^2 + xy = x^3 + a_2x^2 + a_6$ (if $a_1 \neq 0$) or $y^2 + a_3y = x^3 + a_4x + a_6$ (if $a_1 = 0$). In characteristic 3, one can reach $y^2 = x^3 + a_2x^2 + a_4x + a_6$ (i.e., $a_1 = a_3 = 0$ but the x^2 term remains). For simplicity, many of our explicit examples will use the short Weierstrass form over fields of characteristic $\neq 2, 3$, but we will state the general results for all characteristics.

Weierstrass Equations over Rings

The Weierstrass equation (1.2) makes sense with coefficients a_i in any commutative ring R , not just a field. By taking $R = \mathbb{Z}$, for example, we obtain *integral* Weierstrass models, which will be essential when we study reduction modulo primes in chapter 6. More generally, one can define Weierstrass models over an arbitrary base scheme S , obtaining a family of curves parametrized by S . We will not develop this formalism now, but the reader should keep in mind that the constructions of this section are not limited to the case of a base field.

Example 1.1.10: The Running Examples

The following elliptic curves will appear repeatedly throughout this book. We record their Weierstrass equations for reference.

(a) *The Fermat cubic.* The projective curve $X^3 + Y^3 = Z^3$ in $\mathbb{P}_{\mathbb{Q}}^2$ is a smooth cubic. In the affine chart $Z = 1$, it is $x^3 + y^3 = 1$. This is *not* in Weierstrass form as written, but a change of variables transforms it to one. Setting $u = \frac{12-36y}{x+y}$ and $v = \frac{36(1-y)}{x+y}$ (or equivalently, applying the standard algorithm to put a smooth plane cubic with a rational point into Weierstrass form), one obtains the Weierstrass equation

$$v^2 = u^3 - 432.$$

This is a short Weierstrass form with $A = 0$ and $B = -432$. The rational point $[1 : -1 : 0]$ on the Fermat cubic maps to the point at infinity $O = [0 : 1 : 0]$ on the Weierstrass model.

(b) *The curve $y^2 = x^3 - x$.* This is a short Weierstrass equation with $A = -1$, $B = 0$ over any field K with $\text{char} K \neq 2, 3$. Over $\mathbb{Q}$, it has the rational points O , $(0, 0)$, $(1, 0)$, and $(-1, 0)$ (all of which are 2-torsion points, as we shall see). This curve has complex multiplication by $\mathbb{Z}[i]$, the Gaussian integers, which will be explored in depth in ??.

(c) *The curve $y^2 = x^3 + 1$.* This is a short Weierstrass equation with $A = 0$, $B = 1$. It has complex multiplication by $\mathbb{Z}[\omega]$, the Eisenstein integers, where $\omega = e^{2\pi i/3}$. The presence of a third root of unity in the endomorphism ring gives this curve a richer automorphism group than a generic elliptic curve, as we will see in section 1.5.

(d) *The congruent number curves.* For a positive integer n , the curve

$$E_n : y^2 = x^3 - n^2x$$

has $A = -n^2$, $B = 0$. The name comes from the classical *congruent number problem*: n is a congruent number (i.e., the area of a right triangle with rational side lengths) if and only if $E_n(\mathbb{Q})$ has a point of infinite order. These curves form a family of quadratic twists of $E_1 : y^2 = x^3 - x$ (example (b)), and their arithmetic is a rich source of open questions.

Exercise 1.1.1

1. Let E/K be an elliptic curve with marked point O . Show that the map $\varphi_{|n \cdot (O)|} : E \rightarrow \mathbb{P}^{n-1}$ associated to the complete linear system $|n \cdot (O)|$ is a closed immersion for all $n \geq 3$, and determine the degree of the image. What happens for $n = 2$?
2. Show that the admissible changes of variables form a group under composition: given two admissible changes (u_1, r_1, s_1, t_1) and (u_2, r_2, s_2, t_2) , determine explicit formulas for the parameters (u_3, r_3, s_3, t_3) of their composition, and find the formula for the inverse of (u, r, s, t) .
3. Verify the relation $4b_8 = b_2b_6 - b_4^2$ by direct expansion from the definitions $b_2 = a_1^2 + 4a_2$, $b_4 = a_1a_3 + 2a_4$, $b_6 = a_3^2 + 4a_6$, $b_8 = a_1^2a_6 - a_1a_3a_4 + 4a_2a_6 + a_2a_3^2 - a_4^2$.
4. Carry out the change of variables transforming the Fermat cubic $x^3 + y^3 = 1$ into Weierstrass form in detail. (*Hint*: The curve has the rational point $(1, 0)$. Parametrize lines through this point as $y = t(x - 1)$, substitute to get a rational parametrization, then use this to build functions with prescribed pole orders at the point at infinity.)
5. Let K be a field of characteristic 2. Starting from the general Weierstrass equation $y^2 + a_1xy + a_3y = x^3 + a_2x^2 + a_4x + a_6$, show:
 - (a) If $a_1 \neq 0$, the substitution $x \mapsto a_1^2x + a_3/a_1$, $y \mapsto a_1^3y + (a_1^2a_4 + a_3^2)/a_1^3$ brings the equation to the form $y^2 + xy = x^3 + a'_2x^2 + a'_6$.
 - (b) If $a_1 = 0$, the equation can be brought to the form $y^2 + a_3y = x^3 + a_4x + a_6$.
6. Let C be a smooth projective curve over K with a rational point $P \in C(K)$. Suppose that the complete linear system $|3 \cdot (P)|$ embeds C as a smooth cubic in $\mathbb{P}^2$. Use the Riemann–Roch theorem to prove that C has genus 1. (*Thus the genus-1 condition in the definition of an elliptic curve is equivalent to asking that the Weierstrass model is a smooth cubic.*)

1.2 The Discriminant and the j -Invariant

The Weierstrass equation of an elliptic curve is not unique: different choices of coordinates yield different coefficients. Two natural questions arise. First, *when* does a Weierstrass equation define an elliptic curve — that is, when is the resulting projective cubic smooth? Second, *when* do two Weierstrass equations define isomorphic elliptic curves? The discriminant answers the first question, and the j -invariant answers the second (at least over an algebraically closed field). Note that genus 0 curves are all (over $\bar{k}$) rational, and hence this is new behaviour for curves of genus ≥ 1 .

Smoothness and the Discriminant

We begin with the short Weierstrass form, where the smoothness criterion is both geometrically transparent and algebraically clean.

Proposition 1.2.1: Smoothness for Short Weierstrass

Let $\text{char}K \neq 2, 3$. The curve $E : y^2 = x^3 + Ax + B$ in $\mathbb{P}_K^2$ is smooth if and only if $4A^3 + 27B^2 \neq 0$.

Proof:

The point at infinity $O = [0 : 1 : 0]$ is always smooth (as verified in section 1.1), so any singularity must be an affine point (x_0, y_0) with $y_0^2 = x_0^3 + Ax_0 + B$. Setting $f(x, y) = y^2 - x^3 - Ax - B$, the Jacobian criterion for smoothness requires that $\partial f/\partial x$ and $\partial f/\partial y$ do not simultaneously vanish at any point on the curve. We compute

$$\frac{\partial f}{\partial y} = 2y, \quad \frac{\partial f}{\partial x} = -3x^2 - A.$$

Since $\text{char}K \neq 2$, the condition $\partial f/\partial y = 0$ forces $y_0 = 0$. Then $\partial f/\partial x = 0$ gives $3x_0^2 + A = 0$, and the curve equation gives $x_0^3 + Ax_0 + B = 0$. A singularity exists if and only if the cubic polynomial $g(x) = x^3 + Ax + B$ has a root x_0 at which $g'(x_0) = 3x_0^2 + A = 0$ — that is, if and only if g has a repeated root.

The discriminant of the cubic $g(x) = x^3 + px + q$ (with no x^2 term) is classically

$$\text{disc}(g) = -4p^3 - 27q^2,$$

and a polynomial has a repeated root (over the algebraic closure) if and only if its discriminant vanishes. With $p = A$ and $q = B$, the curve is smooth if and only if $-4A^3 - 27B^2 \neq 0$, i.e., $4A^3 + 27B^2 \neq 0$.

We recall why the discriminant formula works: if the roots of $x^3 + px + q$ are $\alpha_1, \alpha_2, \alpha_3$, then $\text{disc}(g) = \prod_{i < j} (\alpha_i - \alpha_j)^2$. By Vieta's formulas, $\alpha_1 + \alpha_2 + \alpha_3 = 0$, $\alpha_1\alpha_2 + \alpha_1\alpha_3 + \alpha_2\alpha_3 = p$, and $\alpha_1\alpha_2\alpha_3 = -q$. A direct computation of the product of squared differences yields $-4p^3 - 27q^2$; see [5] for the general theory of discriminants and resultants.

We now define the discriminant for an arbitrary Weierstrass equation.

Definition 1.2.2: Discriminant

Given a Weierstrass equation with b -invariants b_2, b_4, b_6, b_8 as in definition 1.1.6, the *discriminant* is

$$\Delta = -b_2^2 b_8 - 8b_4^3 - 27b_6^2 + 9b_2 b_4 b_6. \quad (1.10)$$

Let us verify that this definition is consistent with proposition 1.2.1. For the short Weierstrass form $y^2 = x^3 + Ax + B$, we have $a_1 = a_2 = a_3 = 0$, $a_4 = A$, $a_6 = B$, so:

$$\begin{aligned} b_2 &= 0^2 + 4 \cdot 0 = 0, \\ b_4 &= 0 \cdot 0 + 2A = 2A, \\ b_6 &= 0^2 + 4B = 4B, \\ b_8 &= 0 - 0 + 0 + 0 - A^2 = -A^2. \end{aligned}$$

Substituting into (1.10):

$$\begin{aligned}\Delta &= -0 \cdot (-A^2) - 8(2A)^3 - 27(4B)^2 + 9 \cdot 0 \cdot 2A \cdot 4B \\ &= -64A^3 - 432B^2 \\ &= -16(4A^3 + 27B^2).\end{aligned}$$

Thus $\Delta \neq 0$ if and only if $4A^3 + 27B^2 \neq 0$, confirming that the general formula specializes correctly. The factor of -16 is a normalization convention; some authors define $\Delta' = 4A^3 + 27B^2$ as the discriminant for the short form, but the equation (1.10) has the advantage of transforming cleanly under admissible changes of variables, as we shall see shortly.

The following fundamental identity relates the discriminant to the c -invariants introduced in definition 1.1.7.

Proposition 1.2.3: The c - Δ Identity

For any Weierstrass equation,

$$1728 \Delta = c_4^3 - c_6^2. \quad (1.11)$$

Proof:

We verify this for the short Weierstrass form and then argue that the general case follows. For $y^2 = x^3 + Ax + B$, we computed $\Delta = -16(4A^3 + 27B^2)$, and from definition 1.1.7:

$$\begin{aligned}c_4 &= b_2^2 - 24b_4 = 0 - 24(2A) = -48A, \\ c_6 &= -b_2^3 + 36b_2b_4 - 216b_6 = 0 + 0 - 216(4B) = -864B.\end{aligned}$$

Then:

$$\begin{aligned}c_4^3 - c_6^2 &= (-48A)^3 - (-864B)^2 \\ &= -110592A^3 - 746496B^2. \quad = 1728(-64A^3 - 432B^2)\end{aligned}$$

On the other hand:

$$1728 \Delta = 1728 \cdot (-16)(4A^3 + 27B^2) = -27648(4A^3 + 27B^2) = -110592A^3 - 746496B^2.$$

The two expressions agree. For the general Weierstrass equation, the identity (1.11) is verified by expanding both sides in terms of b_2, b_4, b_6, b_8 (using the relation $4b_8 = b_2b_6 - b_4^2$ from (1.6)) and confirming that every monomial cancels. This is a tedious but straightforward computation that we leave to the reader — or to a computer algebra system.

The identity $1728\Delta = c_4^3 - c_6^2$ is not just a computational curiosity. It means that the pair (c_4, c_6) lies on the affine curve $X^3 - Y^2 = 1728\Delta$ in the (c_4, c_6) -plane, and the smoothness condition $\Delta \neq 0$ is equivalent to $(c_4, c_6) \neq (0, 0)$ on the curve $X^3 - Y^2 = 0$. We will see that this identity is the algebraic shadow of the analytic relation between Eisenstein series and the modular discriminant when we study the complex-analytic theory in chapter 3.

We now establish that $\Delta \neq 0$ characterizes smoothness for arbitrary Weierstrass equations, not just the short form.

Theorem 1.2.4: Smoothness Criterion

A Weierstrass equation defines an elliptic curve (i.e., the associated projective cubic in $\mathbb{P}_K^2$ is smooth) if and only if $\Delta \neq 0$.

Proof:

If $\text{char} K \neq 2, 3$, we may apply the admissible change of variables from section 1.1 to reach the short Weierstrass form $y^2 = x^3 + Ax + B$. Since admissible changes are isomorphisms of projective varieties, the original curve is smooth if and only if the short form is, which by proposition 1.2.1 is equivalent to $4A^3 + 27B^2 \neq 0$, i.e., $\Delta \neq 0$.

For the general case (including characteristics 2 and 3), we argue directly from the Jacobian criterion. The affine Weierstrass equation is $f(x, y) = y^2 + a_1xy + a_3y - x^3 - a_2x^2 - a_4x - a_6 = 0$, with partial derivatives

$$\frac{\partial f}{\partial y} = 2y + a_1x + a_3, \quad \frac{\partial f}{\partial x} = a_1y - 3x^2 - 2a_2x - a_4.$$

A singular point (x_0, y_0) satisfies $f = 0$, $\partial f/\partial x = 0$, and $\partial f/\partial y = 0$ simultaneously. These three polynomial equations in two unknowns can be combined (via resultants) into a single polynomial condition on the coefficients a_i , and one verifies that this condition is precisely $\Delta = 0$. The resultant computation can be organized by first solving $\partial f/\partial y = 0$ for y_0 (or x_0 in characteristic 2), substituting into the other two equations, and computing the discriminant of the resulting system. We refer to Silverman [21, Chapter III, Proposition 1.4] for the complete verification in all characteristics.

Next, the discriminant is not an invariant of the isomorphism class of an elliptic curve, but it transforms in a controlled way:

Proposition 1.2.5: Transformation Law for c_4 , c_6 , and Δ

Under the admissible change of variables (u, r, s, t) from definition 1.1.4, the quantities c_4 , c_6 , and Δ transform as

$$c'_4 = u^{-4} c_4, \quad c'_6 = u^{-6} c_6, \quad \Delta' = u^{-12} \Delta.$$

Proof:

The transformation laws for c'_4 and c'_6 follow from the transformation formulas for the a'_i in proposition 1.1.5: one substitutes the expressions for a'_i into the definitions of b'_i and then c'_i . This is a bookkeeping exercise. Once $c'_4 = u^{-4}c_4$ and $c'_6 = u^{-6}c_6$ are established, the transformation law for Δ follows immediately from the identity (1.11):

$$1728 \Delta' = c'^3_4 - c'^2_6 = u^{-12} c^3_4 - u^{-12} c^2_6 = u^{-12} (c^3_4 - c^2_6) = u^{-12} \cdot 1728 \Delta.$$

The weight -12 of Δ under rescaling is characteristic: it will reappear when we study the modular discriminant function $\Delta(\tau)$ in the analytic theory, which is a modular form of weight 12. The fact that c_4 has weight -4 and c_6 has weight -6 will similarly correspond to the Eisenstein series $E_4(\tau)$

and $E_6(\tau)$.

The j -Invariant

Since c_4 and Δ both change by powers of u under admissible changes, any ratio that cancels the u -dependence will be a genuine isomorphism invariant. The simplest such ratio is:

Definition 1.2.6: The j -Invariant

The j -invariant of an elliptic curve given by a Weierstrass equation with $\Delta \neq 0$ is

$$j = j(E) = \frac{c_4^3}{\Delta}.$$

By proposition 1.2.5, if we apply an admissible change of variables, the new j -invariant is

$$j' = \frac{c_4'^3}{\Delta'} = \frac{(u^{-4}c_4)^3}{u^{-12}\Delta} = \frac{u^{-12}c_4^3}{u^{-12}\Delta} = \frac{c_4^3}{\Delta} = j.$$

Thus j is invariant under all admissible changes of variables and depends only on the K -isomorphism class of the elliptic curve. The normalizing factor of 1728 that some authors include (writing $j = 1728 c_4^3/\Delta$) is a matter of convention related to the theory of modular forms; both conventions appear in the literature. Our convention matches Silverman [21], and the choice is pinned down by the requirement that the j -function $j(\tau)$ on the upper half-plane have residue 1 at the cusp $\tau \rightarrow i\infty$ (see chapter 3).

For the short Weierstrass form $y^2 = x^3 + Ax + B$ (with $\text{char}K \neq 2, 3$), we have $c_4 = -48A$ and $\Delta = -16(4A^3 + 27B^2)$, so:

$$j = \frac{(-48A)^3}{-16(4A^3 + 27B^2)} = \frac{-110592 A^3}{-16(4A^3 + 27B^2)} = \frac{6912 A^3}{4A^3 + 27B^2} = -\frac{1728 \cdot (4A)^3}{\Delta}. \quad (1.12)$$

Example 1.2.7: j -Invariants of the Running Examples

We compute the j -invariant for each of the curves introduced in example 1.1.10.

(a) The curve $y^2 = x^3 - x$. Here $A = -1$, $B = 0$. Then

$$\Delta = -16(4(-1)^3 + 27 \cdot 0^2) = -16 \cdot (-4) = 64,$$

and $c_4 = -48 \cdot (-1) = 48$. So

$$j = \frac{48^3}{64} = \frac{110592}{64} = 1728.$$

(b) The curve $y^2 = x^3 + 1$. Here $A = 0$, $B = 1$. Then

$$\Delta = -16(0 + 27) = -432,$$

and $c_4 = -48 \cdot 0 = 0$. So

$$j = \frac{0^3}{-432} = 0.$$

(c) The congruent number curves $E_n : y^2 = x^3 - n^2x$. Here $A = -n^2$, $B = 0$. Then

$$\Delta = -16(4(-n^2)^3 + 0) = -16 \cdot (-4n^6) = 64n^6,$$

and $c_4 = -48 \cdot (-n^2) = 48n^2$. So

$$j = \frac{(48n^2)^3}{64n^6} = \frac{110592n^6}{64n^6} = 1728.$$

Every congruent number curve E_n has $j = 1728$, regardless of n . This is not a coincidence: the curves E_n are all quadratic twists of $E_1 : y^2 = x^3 - x$, and we will see in section 1.5 that twists of a given curve always share the same j -invariant.

(d) The Fermat cubic. In its Weierstrass form $v^2 = u^3 - 432$ (example 1.1.10(a)), we have $A = 0$, $B = -432$. Since $A = 0$, we immediately get $c_4 = 0$ and $j = 0$. The Fermat cubic and the curve $y^2 = x^3 + 1$ both have $j = 0$, so they are isomorphic over $\overline{\mathbb{Q}}$ (though not over $\mathbb{Q}$).

The two special values $j = 0$ and $j = 1728$ are distinguished by the vanishing of c_6 and c_4 respectively. From the identity $1728\Delta = c_4^3 - c_6^2$: if $j = 0$ then $c_4 = 0$, so $1728\Delta = -c_6^2$, and the curve (in short form) has $A = 0$; if $j = 1728$ then $c_6 = 0$, so $1728\Delta = c_4^3$, and the curve has $B = 0$. These curves have extra automorphisms compared to generic elliptic curves, a phenomenon we analyze in section 1.5.

The Classification Theorem

The j -invariant is a complete isomorphism invariant over an algebraically closed field. This is the first major classification result for elliptic curves.

Theorem 1.2.8: Classification by the j -Invariant

Let $K = \overline{K}$ be an algebraically closed field, and let E, E' be elliptic curves over K . Then $E \cong E'$ (as elliptic curves over K) if and only if $j(E) = j(E')$.

Proof:

We have already shown that j is an isomorphism invariant, so $E \cong E'$ implies $j(E) = j(E')$. The content is the converse.

Assume $\text{char} K \neq 2, 3$ (the exceptional characteristics require similar arguments with the generalized Weierstrass form; see [21, Appendix A]). Write both curves in short Weierstrass form: $E : y^2 = x^3 + Ax + B$ and $E' : y'^2 = x'^3 + A'x' + B'$. By the discussion following proposition 1.1.5, an isomorphism between two short Weierstrass equations (preserving the condition $a_1 = a_2 = a_3 = 0$) must have the form $x = u^2x'$, $y = u^3y'$ for some $u \in K^*$, which transforms the coefficients as

$$A = u^4A', \quad B = u^6B'. \quad (1.13)$$

To see this, recall from proposition 1.1.5 that $ua'_1 = a_1 + 2s$, and since $a_1 = a'_1 = 0$ we get $s = 0$. Next, $u^2a'_2 = a_2 - sa_1 + 3r - s^2 = 3r$, and since $a_2 = a'_2 = 0$ and $\text{char} K \neq 3$, we get $r = 0$. Finally, $u^3a'_3 = a_3 + ra_1 + 2t = 2t$, and since $a_3 = a'_3 = 0$ and $\text{char} K \neq 2$, we get $t = 0$. The

remaining transformation laws give $u^4 a'_4 = a_4$ and $u^6 a'_6 = a_6$, yielding (1.13).

Thus we need to show: if $j(E) = j(E')$, there exists $u \in K^*$ with $A = u^4 A'$ and $B = u^6 B'$.

Case 1: $j = 0$. Then $A = A' = 0$ (since $c_4 = -48A = 0$), and $B, B' \neq 0$ (since $\Delta \neq 0$). We need $u^6 = B/B'$. Since K is algebraically closed, a sixth root exists.

Case 2: $j = 1728$. Then $B = B' = 0$ (since $c_6 = -864B = 0$), and $A, A' \neq 0$. We need $u^4 = A/A'$, which has a solution over K .

Case 3: $j \neq 0, 1728$. Then A, A', B, B' are all nonzero. From (1.13), if such a u exists, then $u^{12} = (u^4)^3 = (A/A')^3$ and also $u^{12} = (u^6)^2 = (B/B')^2$. Hence a necessary condition is

$$\left(\frac{A}{A'}\right)^3 = \left(\frac{B}{B'}\right)^2. \quad (1.14)$$

We claim that $j(E) = j(E')$ implies (1.14). Indeed, from the explicit formula (1.12), the j -invariant depends only on the ratio A^3/B^2 (and similarly for A', B'):

$$j = \frac{6912 A^3}{4A^3 + 27B^2} = \frac{6912}{4 + 27(B^2/A^3)},$$

so $j(E) = j(E')$ forces $B^2/A^3 = B'^2/A'^3$, which is precisely (1.14).

Now choose any $u \in K^*$ with $u^4 = A/A'$ (possible since $K = \overline{K}$). Then

$$u^{12} = \left(\frac{A}{A'}\right)^3 = \left(\frac{B}{B'}\right)^2,$$

so $u^6 = \pm B/B'$. If $u^6 = B/B'$, we are done. If $u^6 = -B/B'$, replace u by ζu where ζ is a primitive fourth root of unity in K (which exists since K is algebraically closed and $\text{char } K \neq 2$). Then $(\zeta u)^4 = \zeta^4 u^4 = u^4 = A/A'$ (unchanged) and $(\zeta u)^6 = \zeta^6 u^6 = \zeta^2 \cdot u^6 = (-1)(-B/B') = B/B'$ (using $\zeta^2 = -1$). Thus the admissible change $(u', 0, 0, 0)$ with $u' = \zeta u$ gives the desired isomorphism.

The proof reveals exactly *why* the theorem fails over non-algebraically closed fields: the required $u \in K^*$ may not exist in K . Two curves E and E' with $j(E) = j(E')$ are always isomorphic over $\overline{K}$, but the isomorphism may not descend to K . The curves related in this way are called *twists* of each other, and we will classify them in section 1.5.

Surjectivity: Every Element of K is a j -Invariant

The j -invariant defines a map from the set of K -isomorphism classes of elliptic curves over K to the field K . Over an algebraically closed field, this map is surjective: every element of K arises as the j -invariant of some elliptic curve. In fact, surjectivity holds over any field.

Proposition 1.2.9: Surjectivity of j

For every $j_0 \in K$, there exists an elliptic curve E/K with $j(E) = j_0$.

Proof:

We construct explicit Weierstrass equations. Assume first that $\text{char} K \neq 2, 3$.

Case $j_0 = 0$: Take $E : y^2 = x^3 + 1$. Then $A = 0$, $B = 1$, so $\Delta = -432 \neq 0$ (since $\text{char} K \neq 2, 3$) and $j = 0$.

Case $j_0 = 1728$: Take $E : y^2 = x^3 + x$. Then $A = 1$, $B = 0$, so $\Delta = -64 \neq 0$ and $j = (-48)^3/(-64) = 1728$.

Case $j_0 \neq 0, 1728$: Set

$$c = \frac{j_0}{j_0 - 1728}, \quad (1.15)$$

which is well-defined (the denominator is nonzero) and nonzero (the numerator is nonzero). Take the curve

$$E : y^2 = x^3 - 3cx - 2c.$$

We verify that $j(E) = j_0$. With $A = -3c$ and $B = -2c$:

$$4A^3 + 27B^2 = 4(-3c)^3 + 27(-2c)^2 = -108c^3 + 108c^2 = 108c^2(1 - c).$$

Now $1 - c = 1 - j_0/(j_0 - 1728) = -1728/(j_0 - 1728)$, so

$$4A^3 + 27B^2 = 108c^2 \cdot \frac{-1728}{j_0 - 1728} = \frac{-108 \cdot 1728 c^2}{j_0 - 1728}.$$

The discriminant is $\Delta = -16(4A^3 + 27B^2) = 16 \cdot 108 \cdot 1728 c^2/(j_0 - 1728)$, which is nonzero since $c \neq 0$ and $j_0 \neq 1728$, confirming the curve is smooth. For the j -invariant:

$$\begin{aligned} c_4 &= -48A = -48(-3c) = 144c, \\ j &= \frac{c_4^3}{\Delta} = \frac{(144c)^3}{16 \cdot 108 \cdot 1728 c^2/(j_0 - 1728)} = \frac{144^3 c^3 \cdot (j_0 - 1728)}{16 \cdot 108 \cdot 1728 c^2}. \end{aligned}$$

We compute $144^3 = 2985984$ and $16 \cdot 108 \cdot 1728 = 2985984$, so these factors cancel, leaving

$$j = c \cdot (j_0 - 1728) = \frac{j_0}{j_0 - 1728} \cdot (j_0 - 1728) = j_0.$$

1.2.10 Remark: j -Invariant as a Coarse Moduli Space Over an algebraically closed field K , the classification theorem says that the map

$$\{\text{elliptic curves over } K\}/\text{isomorphism} \longrightarrow K, \quad [E] \longmapsto j(E),$$

is a bijection. In the language of moduli theory, this identifies the j -line $\mathbb{A}_K^1$ as the *coarse moduli space* of elliptic curves. It is not a *fine* moduli space, because there is no universal family of elliptic curves over $\mathbb{A}_K^1$: the extra automorphisms at $j = 0$ and $j = 1728$ obstruct the existence of a universal family. The construction of fine moduli spaces requires rigidifying the moduli problem by adding level structure, a topic taken up in ??.

Detecting Special Values of j

We collect a useful observation about the vanishing of c_4 and c_6 .

Proposition 1.2.11: Special Values of the j -Invariant

Let E/K be an elliptic curve with $\text{char} K \neq 2, 3$ in short Weierstrass form $y^2 = x^3 + Ax + B$. Then:

1. $j(E) = 0$ if and only if $A = 0$ (equivalently, $c_4 = 0$).
2. $j(E) = 1728$ if and only if $B = 0$ (equivalently, $c_6 = 0$).
3. $j(E) \neq 0, 1728$ if and only if $A \neq 0$ and $B \neq 0$ (equivalently, $c_4 \neq 0$ and $c_6 \neq 0$).

Proof:

From $c_4 = -48A$ and $c_6 = -864B$ (with 48, 864 invertible since $\text{char} K \neq 2, 3$):

(1) $j = 0$ iff $c_4^3 = 0$ iff $c_4 = 0$ iff $A = 0$. In this case $B \neq 0$ (else $\Delta = 0$), so the curve has the form $y^2 = x^3 + B$.

(2) $j = 1728$ iff $c_4^3 = 1728\Delta$ iff $c_4^3 = c_4^3 - c_6^2$ iff $c_6 = 0$ iff $B = 0$. In this case $A \neq 0$, so the curve has the form $y^2 = x^3 + Ax$.

(3) follows from (1) and (2).

Example 1.2.12: A Non-Running Example: General j -Value

Consider the elliptic curve $E : y^2 = x^3 - 4x + 4$ over $\mathbb{Q}$. We have $A = -4$ and $B = 4$. Computing in prime factorizations throughout:

$$\begin{aligned} 4A^3 + 27B^2 &= 4(-64) + 27(16) = -256 + 432 = 176 = 2^4 \cdot 11, \\ \Delta &= -16 \cdot 176 = -2^4 \cdot 2^4 \cdot 11 = -2^8 \cdot 11, \\ c_4 &= -48(-4) = 192 = 2^6 \cdot 3, \\ c_4^3 &= 2^{18} \cdot 3^3. \end{aligned}$$

Therefore

$$j = \frac{c_4^3}{\Delta} = \frac{2^{18} \cdot 27}{-2^8 \cdot 11} = -\frac{2^{10} \cdot 27}{11} = -\frac{27648}{11}.$$

Since $j \neq 0, 1728$, this is a “generic” elliptic curve with no extra automorphisms (proposition 1.2.11). The non-integral value of j illustrates that j -invariants of elliptic curves over $\mathbb{Q}$ need not be integers.

1.2.13 Remark: Integrality of the j -Invariant The j -invariant of an elliptic curve over $\mathbb{Q}$ is a rational number, not necessarily an integer. However, a remarkable theorem in the theory of complex multiplication states that if E has CM by an order in an imaginary quadratic field, then $j(E)$ is an algebraic *integer*. In fact, $j(E)$ generates the Hilbert class field of the relevant imaginary quadratic field. The integrality of j -invariants for CM curves is a deep arithmetic phenomenon that we will prove in ??.

Exercise 1.2.1

1. Verify the discriminant formula for two generalized Weierstrass equations:

- (a) $y^2 + y = x^3 - x^2$ (Cremona label 11a3, conductor 11).
- (b) $y^2 + y = x^3 - x^2 - 10x - 20$ (Cremona label 11a1, a model of $X_0(11)$).

For each, compute b_2, b_4, b_6, b_8 , then c_4, c_6 , then Δ and j . Verify that $1728\Delta = c_4^3 - c_6^2$.

2. Verify the identity $1728\Delta = c_4^3 - c_6^2$ for an arbitrary Weierstrass equation by expanding both sides in terms of b_2, b_4, b_6 (using the relation $4b_8 = b_2b_6 - b_4^2$ to eliminate b_8). (*Hint:* A computer algebra system makes this painless. Alternatively, argue that both sides are polynomial expressions in $a_1, \dots, a_6$ of weighted degree 12; verify agreement for sufficiently many specializations.)
3. Verify directly (without using the identity $1728\Delta = c_4^3 - c_6^2$) that under the admissible change $(u, 0, 0, 0)$, the discriminant transforms as $\Delta' = u^{-12}\Delta$. (*Hint:* The change $x = u^2x', y = u^3y'$ sends $a_4 \mapsto u^{-4}a_4, a_6 \mapsto u^{-6}a_6$ in the short form.)
4. For each integer n with $1 \leq n \leq 10$, compute the j -invariant of the congruent number curve $E_n : y^2 = x^3 - n^2x$ and the discriminant $\Delta(E_n)$. What pattern do you observe for $\Delta(E_n)$ as a function of n ?
5. Consider the explicit construction in proposition 1.2.9: for $j_0 \neq 0, 1728$, the curve $y^2 = x^3 - 3cx - 2c$ with $c = j_0/(j_0 - 1728)$ has j -invariant j_0 .

- (a) When $K = \mathbb{Q}$, for which $j_0 \in \mathbb{Z}$ do the coefficients $A = -3c$ and $B = -2c$ lie in $\mathbb{Z}$?
 - (b) Compute c , A , B , and Δ explicitly for $j_0 = -2^{15} \cdot 3 \cdot 5^3$ (the j -invariant of the Frey curve associated to certain Diophantine equations, as we will see in ??).
6. Let $E : y^2 = x^3 + Ax + B$ be an elliptic curve over an algebraically closed field K with $\text{char} K \neq 2, 3$. An *automorphism* of E (as an elliptic curve, i.e., preserving O) corresponds to an admissible change $(u, 0, 0, 0)$ from E to itself. Show:
- (a) Such an automorphism requires $u^4 A = A$ and $u^6 B = B$.
 - (b) If $j \neq 0, 1728$ (i.e., $A \neq 0$ and $B \neq 0$), then $u^2 = 1$, so $u = \pm 1$, and $\text{Aut}(E) \cong \mathbb{Z}/2\mathbb{Z}$.
 - (c) If $j = 1728$ (i.e., $B = 0$), then $u^4 = 1$, so $\text{Aut}(E) \cong \mathbb{Z}/4\mathbb{Z}$.
 - (d) If $j = 0$ (i.e., $A = 0$), then $u^6 = 1$, so $\text{Aut}(E) \cong \mathbb{Z}/6\mathbb{Z}$.
- (This computation will be refined and proved in full generality in section 1.5.)

1.3 The Group Law

The single most remarkable feature of an elliptic curve is that its points form an abelian group. This is not a phenomenon one encounters for generic curves: a smooth projective curve of genus 0 has no natural group structure (its automorphism group acts, but the point set $\mathbb{P}^1(K)$ is not itself a group), and curves of genus ≥ 2 are far too rigid — Faltings' theorem ensures they have only finitely many rational points over a number field, which is incompatible with the structure of an interesting group. The group law on an elliptic curve is a consequence of the genus being exactly 1, and the cleanest way to see it is through the Picard group.

The Picard Group Isomorphism

Let E/K be an elliptic curve with marked point O . Recall [3] that the *Picard group* $\text{Pic } E$ is the group of divisor classes on E under linear equivalence, and $\text{Pic}^0 E$ is the subgroup of degree-zero divisor classes. For any point $P \in E(K)$, the divisor $(P) - (O)$ has degree 0, so its class lies in $\text{Pic}^0 E$.

Theorem 1.3.1: The Picard Group Isomorphism

Let E/K be an elliptic curve with marked point O . The map

$$\sigma : E(K) \longrightarrow \text{Pic}^0 E, \quad P \longmapsto [(P) - (O)],$$

is a bijection. (Here $[(P) - (O)]$ denotes the linear equivalence class of the divisor $(P) - (O)$.)

Proof:

Injectivity. Suppose $\sigma(P) = \sigma(Q)$, i.e., $(P) - (O) \sim (Q) - (O)$, which is to say $(P) \sim (Q)$. Then $[(P)]$ and $[(Q)]$ are the same class in $\text{Pic}^1 E$ (the divisor classes of degree 1). By the Riemann–Roch computation of section 1.1, $\ell(P) = 1$, so the complete linear system $|P|$ consists of the single effective divisor (P) . Similarly, $|Q| = \{(Q)\}$. Since $(P) \sim (Q)$ means $|P| = |Q|$, we conclude $(P) = (Q)$ as divisors, hence $P = Q$.

The key input here is precisely the genus-1 condition. If E had genus 0, then $\ell(P) = 2$, and two distinct points could be linearly equivalent (as they are on $\mathbb{P}^1$). If E had genus ≥ 2 , the map σ would still be injective, but it would fail to be surjective.

Surjectivity. Let D be a divisor of degree 0 on E ; we must show $D \sim (P) - (O)$ for some $P \in E(K)$. Consider $D + (O)$, which has degree 1. By the Riemann–Roch theorem (eq. (1.1)), $\ell(D + (O)) = 1$. A nonzero element $f \in L(D + (O))$ satisfies $\text{div}(f) + D + (O) \geq 0$, and since $\text{div}(f) + D + (O)$ has degree $0 + 1 = 1$ and is effective, it must be a single point: $\text{div}(f) + D + (O) = (P)$ for some $P \in E(K)$. Therefore $D \sim (P) - (O)$.

This shows σ maps an elliptic curves to groups. We can extend this to morphisms, showing it is a covariant functor, namely let $\varphi : E_1 \rightarrow E_2$ be a non-constant morphism between smooth, projective curves. Since φ is a finite surjective morphism, it induces a covariant *pushforward* on Weil divisors by mapping points directly to their images:

Definition 1.3.2: Pushforward of Divisors

Let $\varphi : E_1 \rightarrow E_2$ be a non-constant morphism between smooth, projective curves. Then the *pushforward* φ_* is:

$$\varphi_* : \text{WDiv } E_1 \longrightarrow \text{WDiv } E_2, \quad \sum n_i P_i \longmapsto \sum n_i (\varphi(P_i)).$$

Crucially, the pushforward of a principal divisor $\text{div}(f)$ is the principal divisor $\text{div}(\text{Nm}(f))$, where Nm is the field-theoretic norm map from $K(E_1)$ to $K(E_2)$. Consequently, φ_* respects linear equivalence and descends to a well-defined group homomorphism

$$\varphi_* : \text{Pic}^0 E_1 \longrightarrow \text{Pic}^0 E_2$$

on degree-zero divisor classes.

Corollary 1.3.3: Pushforward via Pic

Let (E_1, O_1) and (E_2, O_2) be elliptic curves, and let $\varphi : E_1 \rightarrow E_2$ be a non-constant morphism of algebraic varieties such that $\varphi(O_1) = O_2$. Then φ is a group homomorphism.

The pushforward assignment $(E, \varphi) \mapsto (\text{Pic}^0 E, \varphi_*)$ defines a covariant functor $\text{Pic}^0 : \mathbf{Ell}_K \rightarrow \mathbf{Ab}$ from the category of elliptic curves to the category of abelian groups. Because functors preserve isomorphisms, any isomorphism of curves φ induces a group isomorphism φ_* .

Proof:

By ??, the morphism φ induces a well-defined group homomorphism $\varphi_* : \text{Pic}^0 E_1 \rightarrow \text{Pic}^0 E_2$. For any point $P \in E_1(K)$, we evaluate the action of φ_* on the corresponding divisor class mapped by the bijection σ_1 :

$$\varphi_*([(P) - (O_1)]) = [(\varphi(P)) - (\varphi(O_1))] = [(\varphi(P)) - (O_2)],$$

where the second equality relies on the assumption $\varphi(O_1) = O_2$.

The group laws on E_1 and E_2 are defined precisely by transporting the abelian group structure from Pic^0 via the bijections σ_1 and σ_2 . Because φ_* is a homomorphism of Picard groups

that commutes with these bijections, the underlying map on points φ is forced to be a group homomorphism.

The reader can check that Pic^0 is a faithful, additive functor, but it is not full.

This bijection is the mechanism that produces the group law. Since $\text{Pic}^0 E$ is an abelian group (under addition of divisor classes), we can transport this group structure to $E(K)$ via σ .

Definition 1.3.4: The Group Law on an Elliptic Curve

Let E/K be an elliptic curve with marked point O . The *group law* on $E(K)$ is defined by declaring σ to be a group isomorphism. Concretely, for $P, Q \in E(K)$, their sum $P + Q$ is the unique point satisfying

$$[(P + Q) - (O)] = [(P) - (O)] + [(Q) - (O)] \quad \text{in } \text{Pic}^0 E,$$

or equivalently,

$$(P + Q) \sim (P) + (Q) - (O). \quad (1.16)$$

The identity element is O , and the inverse $-P$ of a point P is the unique point satisfying $(-P) \sim 2(O) - (P)$.

The elegance of this definition is that *associativity, commutativity, and the existence of inverses are all free*: they follow immediately from the corresponding properties of $\text{Pic}^0 E$, which is a group by the general theory of divisor class groups. We have not had to verify a single algebraic identity.

Let us record this as a formal statement.

Corollary 1.3.5: Associativity and Commutativity

The group law on $E(K)$ is associative and commutative, with identity O . Thus $(E(K), +)$ is an abelian group.

Proof:

The map σ is a bijection by theorem 1.3.1, and $\text{Pic}^0 E$ is an abelian group. By definition, $\sigma(P + Q) = \sigma(P) + \sigma(Q)$, so σ is a group isomorphism. All group axioms transfer from $\text{Pic}^0 E$ to $E(K)$.

This is worth lingering on. The classical approach to the group law on elliptic curves proceeds by defining addition geometrically (via the chord-and-tangent construction, which we will describe momentarily), deriving explicit rational formulas, and then *proving* associativity by a painful direct verification involving rational functions or intersections in $\mathbb{P}^2$. That verification is notoriously tedious — Cassels [2] calls it “an exercise in patience, not intelligence” — and sheds no light on *why* the law is associative. The Picard group approach shows that associativity is not a computational accident but a structural inevitability: we are simply reading off the group structure that was always present in $\text{Pic}^0 E$. The reader who has invested in learning scheme theory and sheaf cohomology earns this clarity.

The Chord-and-Tangent Construction

Having established the group law abstractly, we now describe it geometrically. The Weierstrass model embeds E as a cubic curve in $\mathbb{P}^2$, and in the projective plane, every line meets a cubic in exactly three points (counted with multiplicity, by Bézout's theorem). This intersection theory gives the group law a visual interpretation.

Let $P, Q \in E(K)$ with $P \neq Q$. The line ℓ_{PQ} through P and Q meets the cubic E in a third point, call it R . (If $P = Q$, we take ℓ_{PP} to be the tangent line to E at P .) By Bézout's theorem, the divisor of the line ℓ , viewed as a section of $\mathcal{O}_{\mathbb{P}^2}(1)$ restricted to E , satisfies

$$\operatorname{div}(\ell|_E) = (P) + (Q) + (R),$$

where all intersection multiplicities are accounted for. Since two lines in $\mathbb{P}^2$ are linearly equivalent as divisors (they are both sections of $\mathcal{O}_{\mathbb{P}^2}(1)$), the divisor $(P) + (Q) + (R)$ is linearly equivalent to the divisor cut out by *any* line. In particular, taking the vertical line through O (which meets E at O with some multiplicity), we obtain:

$$(P) + (Q) + (R) \sim 3(O). \quad (1.17)$$

Here we use the fact that $O = [0 : 1 : 0]$ is the unique point at infinity on the Weierstrass model, and a generic line through O meets E in O plus two affine points. More precisely, the line $Z = 0$ (the “line at infinity”) meets the Weierstrass cubic at O with multiplicity 3, so $\operatorname{div}(Z|_E) = 3(O)$.

Now, from (1.17), $(R) \sim 3(O) - (P) - (Q)$, and therefore

$$(P) + (Q) - (O) \sim 3(O) - (R) - (O) = 2(O) - (R).$$

The group law (eq. (1.16)) tells us $P + Q$ is the unique point satisfying $(P + Q) \sim (P) + (Q) - (O) \sim 2(O) - (R)$. Now, the vertical line through $R = (x_R, y_R)$ has the equation $x = x_R$, and it meets E at R , at the “reflected” point $R' = (x_R, -y_R - a_1x_R - a_3)$, and at O (the point at infinity on the vertical line). So:

$$(R) + (R') + (O) \sim 3(O),$$

hence $(R') \sim 2(O) - (R)$. Thus $P + Q = R'$: the sum of P and Q is obtained by finding the third intersection R of the line PQ with E , and then reflecting R across the x -axis (relative to the Weierstrass involution).

Proposition 1.3.6: Geometric Description of the Group Law

Let E/K be an elliptic curve in Weierstrass form and let $P, Q \in E(K)$.

1. (*Addition*) If $P \neq Q$, draw the line through P and Q ; it meets E in a third point R . Then $P + Q$ is the reflection of R across the x -axis.
2. (*Doubling*) If $P = Q$, draw the tangent line to E at P ; it meets E in a residual point R (with appropriate multiplicity). Then $2P = P + P$ is the reflection of R .
3. (*Identity*) $P + O = P$ for all P .
4. (*Inverse*) The inverse $-P$ of a point $P = (x_0, y_0)$ is its reflection: $-P = (x_0, -y_0 - a_1x_0 - a_3)$ in the general Weierstrass form, or simply $-P = (x_0, -y_0)$ in the short form.

Proof:

All four statements follow from the divisor-class computation above:

(1) and (2): The third intersection point R of the line $\overline{PQ}$ (or the tangent at P) with E satisfies $(P) + (Q) + (R) \sim 3(O)$. By the argument above, $P + Q$ is the vertical reflection of R .

(3): The “line through P and O ” is the vertical line $x = x_P$ (since O is the point at infinity on vertical lines). This line meets E at P , $-P$, and O , giving $(P) + (-P) + (O) \sim 3(O)$. The third intersection of the line through P and O is $-P$, and the reflection of $-P$ is P . So $P + O = P$.

(4): From $(P) + (-P) + (O) \sim 3(O)$, we get $(P) + (-P) \sim 2(O)$, i.e., $(-P) \sim 2(O) - (P)$, confirming that the vertical reflection is the group inverse.

Explicit Addition Formulas

We now translate the geometric construction into explicit rational formulas. We work first with the short Weierstrass form $y^2 = x^3 + Ax + B$ (assuming $\text{char} K \neq 2, 3$), and then state the general formulas.

Let $P = (x_1, y_1)$ and $Q = (x_2, y_2)$ be affine points on E with $P \neq -Q$ (so that $P + Q \neq O$).

Case 1: $P \neq Q$ (and $x_1 \neq x_2$). The line through P and Q has slope

$$\lambda = \frac{y_2 - y_1}{x_2 - x_1},$$

and equation $y = \lambda(x - x_1) + y_1$. To find the third intersection with E , substitute into $y^2 = x^3 + Ax + B$:

$$(\lambda(x - x_1) + y_1)^2 = x^3 + Ax + B.$$

Expanding the left side:

$$\lambda^2 x^2 - 2\lambda^2 x_1 x + \lambda^2 x_1^2 + 2\lambda y_1 x - 2\lambda y_1 x_1 + y_1^2 = x^3 + Ax + B.$$

Rearranging:

$$x^3 - \lambda^2 x^2 + (\text{lower order terms}) = 0. \quad (1.18)$$

By Bézout, this cubic in x has three roots: x_1 , x_2 , and x_3 (the x -coordinate of the third intersection point R). By Vieta's formulas applied to (1.18):

$$x_1 + x_2 + x_3 = \lambda^2,$$

so

$$x_3 = \lambda^2 - x_1 - x_2.$$

The y -coordinate of R is $y_3 = \lambda(x_3 - x_1) + y_1$. Since $P + Q$ is the reflection of R , we get $P + Q = (x_3, -y_3)$:

$$\boxed{\begin{aligned} x(P + Q) &= \lambda^2 - x_1 - x_2, \\ y(P + Q) &= \lambda(x_1 - x(P + Q)) - y_1, \end{aligned}} \quad \lambda = \frac{y_2 - y_1}{x_2 - x_1}. \quad (1.19)$$

Case 2: $P = Q$ (doubling, with $y_1 \neq 0$). We use the tangent line at P . Implicitly differentiating $y^2 = x^3 + Ax + B$ gives $2y dy = (3x^2 + A) dx$, so the slope of the tangent is

$$\lambda = \frac{3x_1^2 + A}{2y_1}.$$

The same Vieta argument applies: the tangent line at P meets E at P with multiplicity 2 and at a third point R . The cubic $x^3 - \lambda^2 x^2 + \dots = 0$ has x_1 as a double root and x_3 as the third root, so $2x_1 + x_3 = \lambda^2$, giving $x_3 = \lambda^2 - 2x_1$. The doubling formulas are:

$$\boxed{\begin{aligned} x(2P) &= \lambda^2 - 2x_1, \\ y(2P) &= \lambda(x_1 - x(2P)) - y_1, \end{aligned}} \quad \lambda = \frac{3x_1^2 + A}{2y_1}. \quad (1.20)$$

Special cases. If $P = -Q$ (i.e., $x_1 = x_2$ and $y_1 = -y_2$), the “line” through P and Q is vertical and meets E at O ; thus $P + Q = O$. If $P = Q$ and $y_1 = 0$, then P is a point of order 2 (see below), and $2P = O$.

Example 1.3.7: Addition on $y^2 = x^3 - x$

Consider $E : y^2 = x^3 - x$ over $\mathbb{Q}$, with $A = -1$, $B = 0$. The affine rational points include the three points with $y = 0$: namely $(0, 0)$, $(1, 0)$, and $(-1, 0)$, which are all 2-torsion points (since $-P = (x_P, -y_P) = (x_P, 0) = P$ implies $2P = O$). Are there other rational points?

Let us verify the group law by computing the sum of two of the 2-torsion points. Take $P = (0, 0)$ and $Q = (1, 0)$. The slope is $\lambda = (0 - 0)/(1 - 0) = 0$, so the line through P and Q is $y = 0$. Substituting into $y^2 = x^3 - x$ gives $0 = x^3 - x = x(x - 1)(x + 1)$, with roots $x = 0, 1, -1$. The third intersection is $R = (-1, 0)$. Reflecting: $P + Q = (-1, -0) = (-1, 0)$. Thus

$$(0, 0) + (1, 0) = (-1, 0).$$

This is consistent with the group structure: the four points $\{O, (0, 0), (1, 0), (-1, 0)\}$ form a group isomorphic to $\mathbb{Z}/2\mathbb{Z} \times \mathbb{Z}/2\mathbb{Z}$ (the Klein four-group), since each non-identity element has order 2, and the sum of any two distinct non-identity elements is the third.

Example 1.3.8: Doubling on $y^2 = x^3 + 1$

Consider $E : y^2 = x^3 + 1$ over $\mathbb{Q}$, with $A = 0$, $B = 1$. The point $P = (0, 1)$ lies on E . We compute $2P$: the slope of the tangent is

$$\lambda = \frac{3 \cdot 0^2 + 0}{2 \cdot 1} = 0.$$

So $x(2P) = 0 - 2 \cdot 0 = 0$ and $y(2P) = 0 \cdot (0 - 0) - 1 = -1$. Thus $2P = (0, -1)$, which is indeed $-P$ (the reflection of P). This means $3P = 2P + P = -P + P = O$, so P has order 3 in $E(\mathbb{Q})$.

Similarly, the point $Q = (-1, 0)$ satisfies $y_Q = 0$, so Q is a 2-torsion point. We also check: $P + Q$ has slope $\lambda = (0 - 1)/(-1 - 0) = 1$, so $x(P + Q) = 1 - 0 - (-1) = 2$, $y(P + Q) = 1 \cdot (0 - 2) - 1 = -3$. Thus $P + Q = (2, -3)$. One verifies: $(-3)^2 = 9 = 8 + 1 = 2^3 + 1$.

General Weierstrass Addition Formulas

For completeness, and because the general formulas are needed in characteristics 2 and 3, we record the addition law for the full Weierstrass equation $y^2 + a_1xy + a_3y = x^3 + a_2x^2 + a_4x + a_6$.

Given $P = (x_1, y_1)$ and $Q = (x_2, y_2)$ on E with $P + Q \neq O$, define:

$$\lambda = \begin{cases} \frac{y_2 - y_1}{x_2 - x_1} & \text{if } P \neq Q, \\ \frac{3x_1^2 + 2a_2x_1 + a_4 - a_1y_1}{2y_1 + a_1x_1 + a_3} & \text{if } P = Q, \end{cases} \quad (1.21)$$

$$\mu = \begin{cases} \frac{y_1x_2 - y_2x_1}{x_2 - x_1} & \text{if } P \neq Q, \\ \frac{-x_1^3 + a_4x_1 + 2a_6 - a_3y_1}{2y_1 + a_1x_1 + a_3} & \text{if } P = Q. \end{cases} \quad (1.22)$$

In both cases, the line ℓ through P and Q (or the tangent at P if $P = Q$) has equation $y = \lambda x + \mu$. The x -coordinates of the three intersection points of ℓ with E are the roots of a cubic; Vieta's formulas give:

$$\begin{aligned} x(P + Q) &= \lambda^2 + a_1\lambda - a_2 - x_1 - x_2, \\ y(P + Q) &= -(\lambda + a_1)x(P + Q) - \mu - a_3. \end{aligned} \quad (1.23)$$

The derivation is entirely parallel to the short Weierstrass case. Substituting $y = \lambda x + \mu$ into $y^2 + a_1xy + a_3y = x^3 + a_2x^2 + a_4x + a_6$, the left side contributes $(\lambda x + \mu)^2 + a_1x(\lambda x + \mu) + a_3(\lambda x + \mu) = (\lambda^2 + a_1\lambda)x^2 + \dots$. Rearranging to isolate the x^3 coefficient on the right:

$$x^3 + (a_2 - \lambda^2 - a_1\lambda)x^2 + \dots = 0.$$

The sum of roots gives $x_1 + x_2 + x_3 = \lambda^2 + a_1\lambda - a_2$, and the reflection formula for the general Weierstrass involution $(x, y) \mapsto (x, -y - a_1x - a_3)$ produces the y -coordinate.

1.3.9 Remark: The Addition Law as a Rational Map The formulas (1.23) express the coordinates of $P + Q$ as rational functions of the coordinates of P and Q . This is a crucial observation: the addition map

$$m : E \times E \longrightarrow E, \quad (P, Q) \longmapsto P + Q,$$

is given by rational functions, and thus defines a *morphism of varieties* (one must check regularity at the finitely many exceptional points where the denominators vanish, which is done by working in appropriate affine charts around O). Similarly, the inversion map $\iota : E \rightarrow E$ given by $P \mapsto -P$ is a morphism. These facts mean that the group law is not merely a set-theoretic structure but a *geometric* one: elliptic curves are group varieties.

The 2-Torsion Points

The explicit formulas make the 2-torsion structure immediately visible. A point $P = (x_0, y_0)$ has order 2 (i.e., $2P = O$, equivalently $P = -P$) if and only if $(x_0, y_0) = (x_0, -y_0 - a_1x_0 - a_3)$, i.e., $2y_0 + a_1x_0 + a_3 = 0$. In the short Weierstrass form, this simplifies to $y_0 = 0$.

Proposition 1.3.10: 2-Torsion Points

Let $E : y^2 = x^3 + Ax + B$ with $\text{char} K \neq 2$.

1. The non-identity 2-torsion points of E are exactly the points $(x_0, 0)$ where x_0 is a root of $x^3 + Ax + B$.
2. Over $\overline{K}$, the cubic $x^3 + Ax + B$ has three distinct roots e_1, e_2, e_3 (distinct because $\Delta \neq 0$), and

$$E[2](\overline{K}) = \{O, (e_1, 0), (e_2, 0), (e_3, 0)\} \cong \mathbb{Z}/2\mathbb{Z} \times \mathbb{Z}/2\mathbb{Z}.$$

3. Over K itself, the structure of $E[2](K)$ depends on how the cubic $x^3 + Ax + B$ factors: $E[2](K) \cong \mathbb{Z}/2\mathbb{Z} \times \mathbb{Z}/2\mathbb{Z}$ if all roots lie in K ; $E[2](K) \cong \mathbb{Z}/2\mathbb{Z}$ if exactly one root lies in K ; and $E[2](K) = \{O\}$ if no root lies in K .

Proof:

(1) follows from the observation that $P = -P$ iff $y_0 = 0$, and if $y_0 = 0$ then $x_0^3 + Ax_0 + B = 0$. (2) Over $\overline{K}$, $x^3 + Ax + B$ has three roots (distinct since $\Delta = -16(4A^3 + 27B^2) \neq 0$), yielding three non-identity 2-torsion points, plus O . The group $\{O, (e_1, 0), (e_2, 0), (e_3, 0)\}$ has exponent 2, so it is isomorphic to $(\mathbb{Z}/2\mathbb{Z})^k$ for some k ; since it has order 4, $k = 2$. (3) is immediate from (1) and (2).

The identification $E[2](\overline{K}) \cong (\mathbb{Z}/2\mathbb{Z})^2$ is the first instance of a general pattern: we will prove in chapter 2 that $E[n](\overline{K}) \cong (\mathbb{Z}/n\mathbb{Z})^2$ for all n coprime to $\text{char} K$. Over K itself, the Galois group $\text{Gal}(\overline{K}/K)$ acts on $E[n](\overline{K})$, and the structure of $E[n](K)$ is determined by this Galois action. This interplay between torsion and Galois theory is a central theme of the arithmetic of elliptic curves.

The Group Scheme Structure

For the reader who has studied algebraic groups (e.g., in [7]), we now formalize the group law as a group scheme structure. This is not merely a repackaging: the group scheme perspective is the one that generalizes to families of elliptic curves and to base schemes that are not fields.

Proposition 1.3.11: Elliptic Curves as Group Schemes

Let E/K be an elliptic curve. The addition morphism $m : E \times_K E \rightarrow E$, the inversion morphism $\iota : E \rightarrow E$, and the identity section $e : \text{Spec } K \rightarrow E$ (sending the unique point to O) make E into a group scheme over K .

Proof:

The addition formulas (1.23) show that m is defined by rational functions on an open subset of $E \times E$. The fact that these formulas extend to a well-defined morphism on all of $E \times E$ (including the points where the denominators appear to vanish) requires a calculation in local coordinates near O ; see [21] or [3]. Alternatively, the universal property of the Picard scheme gives m directly as a morphism, bypassing the need to check regularity of the formulas. The group

axioms hold because σ is a group isomorphism from the functor of points $E(S) \rightarrow \text{Pic}_{E/K}^0(S)$ for any K -scheme S .

As a group scheme over K , E is a smooth, projective, connected, one-dimensional group variety. By Chevalley's structure theorem [7], any connected algebraic group over a perfect field sits in a short exact sequence $1 \rightarrow H \rightarrow G \rightarrow A \rightarrow 1$ with H affine and A an abelian variety. For an elliptic curve, the affine part is trivial ($H = 1$) and $A = E$: an elliptic curve is an abelian variety. Indeed, elliptic curves are the simplest abelian varieties — they are exactly the abelian varieties of dimension 1.

The commutativity of the group law, which we obtained from the commutativity of $\text{Pic}^0 E$, now also follows from a general fact: abelian varieties are always commutative (a consequence of the rigidity lemma for proper group schemes). This is the phenomenon mentioned in chapter 1's introduction: the geometric rigidity of a projective variety forces the group operation to be abelian.

Example 1.3.12: The Formal Group Preview

The group scheme structure allows us to study the *local* behavior of the group law near the identity O . Working in the local parameter $t = -x/y$ at O (which is a uniformizer of the local ring $\mathcal{O}_{E,O}$), the addition law $m : E \times E \rightarrow E$ induces a formal power series

$$F(t_1, t_2) = t_1 + t_2 + (\text{higher order terms in } t_1, t_2) \in K[[t_1, t_2]].$$

This formal power series F is a *one-dimensional commutative formal group law*, and it encodes the p -adic behavior of the elliptic curve over local fields. The study of this formal group will be the subject of chapter 4.

Exercise 1.3.1

- Let $E : y^2 = x^3 + 1$ over $\mathbb{Q}$.
 - Find all 2-torsion points of E over $\mathbb{Q}$.
 - Verify that $P = (2, 3)$ lies on E , and compute $2P$, $3P$, and $4P$ using the explicit addition formulas.
 - Determine the order of P in $E(\mathbb{Q})$.
- On the congruent number curve $E_5 : y^2 = x^3 - 25x$, verify that $P = (-4, 6)$ lies on E_5 , and compute $2P$. Verify your answer satisfies the curve equation.
- Let E be given by the general Weierstrass equation $y^2 + a_1xy + a_3y = x^3 + a_2x^2 + a_4x + a_6$. Show that the negation map $P = (x_0, y_0) \mapsto -P = (x_0, -y_0 - a_1x_0 - a_3)$ preserves the curve equation (i.e., if (x_0, y_0) lies on E , so does $(x_0, -y_0 - a_1x_0 - a_3)$).
- On the Fermat cubic $x^3 + y^3 = 1$, consider the group law obtained by taking $O = [1 : -1 : 0]$ as the identity. The “line through P and Q ” construction works directly on any smooth plane cubic. Compute the sum $P + Q$ where $P = (1, 0)$ and $Q = (0, 1)$. (*Hint:* The line through P and Q is $x + y = 1$. Find the third intersection with $x^3 + y^3 = 1$, then reflect through O .)
- Show that three points $P, Q, R \in E(K)$ (on a Weierstrass model) are collinear in $\mathbb{P}^2$ if and only if $P + Q + R = O$ in the group law. (*Here collinear means they are the intersection of a single line with E , counted with appropriate multiplicity.*)

6. (Associativity without Riemann–Roch.) Let E be a smooth cubic in $\mathbb{P}_K^2$ with a rational point O . Assuming only Bézout’s theorem and the Cayley–Bacharach theorem (which states: if two cubics C_1, C_2 in $\mathbb{P}^2$ meet in 9 points and a third cubic C_3 passes through 8 of these 9 points, then C_3 also passes through the ninth), prove that the chord-and-tangent group law is associative. (*This is the classical proof that avoids the Picard group. It works, but compare the effort required to the one-line argument of corollary 1.3.5.*)

1.4 The Invariant Differential

By the Riemann–Roch theorem, a smooth projective curve E of genus $g = 1$ has a canonical divisor K_E of degree $2g - 2 = 0$. In section 1.1, we saw that $\ell(K_E) = 1$, meaning $K_E \sim 0$. In the language of sheaves, the canonical bundle Ω_E^1 of regular 1-forms is trivial, and the space of global sections $H^0(E, \Omega_E^1)$ is a one-dimensional vector space over K .

Geometrically, this means an elliptic curve admits a globally regular differential form that vanishes nowhere. Because the space of such forms is one-dimensional, any two such forms are scalar multiples of each other. The Weierstrass equation provides us with a canonical choice of generator for this space. Fixing this generator provides an explicit trivialization of the canonical bundle, meaning it defines an isomorphism $\mathcal{O}_E \xrightarrow{\sim} \Omega_E^1$ via $f \mapsto f\omega$.

Definition and Regularity

Definition 1.4.1: The Invariant Differential

Let E/K be an elliptic curve given by a Weierstrass equation $y^2 + a_1xy + a_3y = x^3 + a_2x^2 + a_4x + a_6$. The *invariant differential* on E is the 1-form

$$\omega = \frac{dx}{2y + a_1x + a_3}. \quad (1.24)$$

For the short Weierstrass form $y^2 = x^3 + Ax + B$, this simplifies to $\omega = \frac{dx}{2y}$.

Let us formally verify that this explicit rational formula yields the globally non-vanishing regular differential promised by the abstract theory.

Proposition 1.4.2: Regularity of ω

The differential ω is regular and nowhere vanishing on E , including at the point at infinity O .

Proof:

Let $f(x, y) = y^2 + a_1xy + a_3y - x^3 - a_2x^2 - a_4x - a_6 = 0$ be the affine equation. Taking the exterior derivative yields

$$\frac{\partial f}{\partial x} dx + \frac{\partial f}{\partial y} dy = 0.$$

Since $\partial f/\partial y = 2y + a_1x + a_3$, we can rewrite ω as

$$\omega = \frac{dx}{\partial f/\partial y} = -\frac{dy}{\partial f/\partial x}.$$

At any affine point $P = (x_0, y_0) \in E$, the Jacobian criterion for smoothness guarantees that the partial derivatives $\partial f/\partial x$ and $\partial f/\partial y$ cannot simultaneously vanish. If $\partial f/\partial y \neq 0$, the expression $dx/(\partial f/\partial y)$ shows that ω is regular and non-vanishing at P . If $\partial f/\partial y = 0$, then $\partial f/\partial x \neq 0$, and the alternative expression $-dy/(\partial f/\partial x)$ establishes regularity and non-vanishing at P .

It remains to check the point at infinity O . We work in the short Weierstrass form $y^2 = x^3 + Ax + B$ (assuming $\text{char} K \neq 2, 3$) for algebraic simplicity; the calculation for the general form is identical in spirit. The local parameter at O is $t = -x/y$ (as previewed in example 1.3.12). We wish to express x and y as Laurent series in t .

Dividing the curve equation by y^3 gives $y^{-1} = (x/y)^3 + Ay^{-3} + By^{-3}$. Substituting $w = -1/y$ and $t = -x/y$, we obtain the relation $w = t^3 - Atw^2 + Bw^3$. By successive approximation (or Hensel's lemma), the lowest-order term of w is t^3 , so $y = -1/w = -t^{-3} + \dots$. Since $x = -ty$, we have $x = t^{-2} + \dots$.

Now we simply compute the expansion of ω :

$$\begin{aligned} dx &= d(t^{-2} + \dots) = -2t^{-3}dt + \dots, \\ 2y &= -2t^{-3} + \dots. \end{aligned}$$

Taking the quotient gives

$$\omega = \frac{dx}{2y} = \frac{-2t^{-3}dt + \dots}{-2t^{-3} + \dots} = (1 + (\text{higher order terms})) dt.$$

Evaluated at $t = 0$ (which is the point O), the coefficient of dt is $1 \neq 0$. Thus ω is regular and non-vanishing at the point at infinity.

1.4.3 Geometric and Formal Group Connection

1. **Poincaré Residue:** The equality $\frac{dx}{f_y} = -\frac{dy}{f_x}$ is a specific instance of the Poincaré residue map. By the adjunction formula, a nowhere-vanishing top-form on a smooth hypersurface $X \subset \mathbb{P}^n$ given by $F = 0$ can always be constructed locally by taking the standard volume form on the ambient space, dividing by F , and taking the residue.
2. **Formal Groups:** The calculation at infinity reveals that $\omega = (1 + \dots)dt$. This shows that ω is the *normalized* invariant differential of the formal group $\hat{E}$ associated to E . The choice of scaling in the definition of ω ensures the leading coefficient is exactly 1, linking the global curve directly to its local formal group law.

Translation Invariance

The name “invariant differential” stems from its interaction with the group law. Because E is an algebraic group, it acts on itself by translation. Just as Lie groups admit left-invariant Haar measures, algebraic groups admit translation-invariant differential forms. For any $P \in E(K)$, the translation map

$$\tau_P : E \longrightarrow E, \quad Q \longmapsto P + Q$$

is an isomorphism of algebraic varieties. It therefore pulls back regular differentials to regular differentials.

Theorem 1.4.4: Translation Invariance

For any point $P \in E(K)$, we have $\tau_P^* \omega = \omega$.

Proof:

Since τ_P is an isomorphism, the pullback $\tau_P^* \omega$ is another everywhere-regular, non-vanishing differential form. Since $\dim_K H^0(E, \Omega_E^1) = 1$, any two global regular differentials are proportional. Thus, there exists a constant $a_P \in K^*$ such that

$$\tau_P^* \omega = a_P \omega.$$

We must show $a_P = 1$. Consider the assignment $P \mapsto a_P$. Because the group addition morphism $m : E \times E \rightarrow E$ is given by rational functions, the map $\Phi : E \rightarrow \mathbb{G}_m$ defined by $\Phi(P) = a_P$ is a morphism of algebraic varieties.

But E is a projective variety and $\mathbb{G}_m$ is affine. The only globally regular functions on a smooth projective connected variety are constants. Thus Φ must map all of E to a single point in $\mathbb{G}_m$. Since τ_O is the identity map, $\tau_O^* \omega = \omega$, meaning $\Phi(O) = 1$. It follows that $\Phi(P) = 1$ for all P , and therefore $\tau_P^* \omega = \omega$.

This proof is a classic example of the power of viewing elliptic curves as proper group varieties. An alternative algebraic proof involves plugging the explicit rational formulas for addition (eq. (1.23)) into the differential and forcing the denominators to clear. While possible, such a computation requires heroic persistence and obscures the structural necessity of the result. Furthermore, evaluating ω at the identity O provides a canonical isomorphism between the global sections of the canonical bundle $H^0(E, \Omega_E^1)$ and the cotangent space at the identity, $T_O^* E$.

An immediate useful corollary of the explicit formula is the action of the negation map on ω .

Proposition 1.4.5: Action of the Hyperelliptic Involution

Let $[-1] : E \rightarrow E$ be the negation map $P \mapsto -P$. Then $[-1]^* \omega = -\omega$.

Proof:

Recall from eq. (1.26) that $[-1](x, y) = (x, -y - a_1 x - a_3)$. Pulling back $\omega = dx/(2y + a_1 x + a_3)$, the numerator dx is unchanged because the x -coordinate is preserved. The denominator transforms to

$$2(-y - a_1 x - a_3) + a_1 x + a_3 = -2y - 2a_1 x - 2a_3 + a_1 x + a_3 = -(2y + a_1 x + a_3).$$

Taking the quotient introduces an overall minus sign, yielding $-\omega$.

Transformation Under Admissible Changes

Since ω is defined in terms of a specific Weierstrass equation, it is important to track how it transforms under an admissible change of variables. This transformation law will be critical when we classify twists and complex multiplication.

Proposition 1.4.6: Transformation Law for ω

Under the admissible change of variables (u, r, s, t) defining $x = u^2x' + r$ and $y = u^3y' + u^2sx' + t$, the invariant differential transforms as

$$\omega = u^{-1}\omega', \quad \text{where } \omega' = \frac{dx'}{2y' + a'_1x' + a'_3}. \quad (1.25)$$

Proof:

Taking the derivative of $x = u^2x' + r$ gives $dx = u^2dx'$. For the denominator, we substitute the expressions for x and y :

$$\begin{aligned} 2y + a_1x + a_3 &= 2(u^3y' + u^2sx' + t) + a_1(u^2x' + r) + a_3 \\ &= 2u^3y' + u^2(2s + a_1)x' + (2t + a_1r + a_3). \end{aligned}$$

Recall the transformation formulas for the coefficients from proposition 1.1.5: we have $ua'_1 = a_1 + 2s$, so $u^2(a_1 + 2s) = u^3a'_1$. Similarly, $u^3a'_3 = a_3 + ra_1 + 2t$. Substituting these into the expression above gives:

$$2y + a_1x + a_3 = 2u^3y' + u^3a'_1x' + u^3a'_3 = u^3(2y' + a'_1x' + a'_3).$$

Forming the quotient for ω :

$$\omega = \frac{dx}{2y + a_1x + a_3} = \frac{u^2dx'}{u^3(2y' + a'_1x' + a'_3)} = u^{-1}\omega'.$$

1.4.7 Remark: Elliptic Integrals and the Origin of the Name The terms “elliptic curve” and “elliptic integral” are historical artifacts. When early modern mathematicians tried to calculate the arc length of an ellipse with semi-major axis a and eccentricity e , they encountered integrals of the form

$$\int \sqrt{\frac{1-e^2t^2}{1-t^2}} dt,$$

which, after algebraic substitutions, can be transformed into integrals of the form $\int \frac{dx}{\sqrt{x^3+Ax+B}}$. Such integrals could not be expressed in terms of elementary functions and were dubbed *elliptic integrals*.

In the 19th century, Abel, Jacobi, and Weierstrass noticed you can invert the function. To understand what this means, consider the arc length of a circle, which gives the integral $z = \int_0^x \frac{dt}{\sqrt{1-t^2}}$. Instead of struggling to evaluate this, we define the inverse function $x = \sin(z)$. By the fundamental theorem of calculus, $\frac{dz}{dx} = \frac{1}{\sqrt{1-x^2}}$, which means the derivative of the inverse is $\frac{dx}{dz} = \sqrt{1-x^2}$. We call this derivative $\cos(z)$, instantly yielding the parametrization of the circle $x^2 + y^2 = 1$ via $(x, y) = (\sin z, \cos z)$.

Weierstrass applied the exact same logic to elliptic integrals. Consider the complex integral in the normal form:

$$z = \int_{\infty}^x \frac{dt}{\sqrt{4t^3 - g_2t - g_3}}.$$

Weierstrass defined the inverse function to be the *Weierstrass $\wp$ -function*, so $x = \wp(z)$.

What is the y -coordinate? In our algebraic curve, $y = \sqrt{4x^3 - g_2x - g_3}$. Looking at the integral, we see that $\frac{dz}{dx} = \frac{1}{y}$. By the chain rule, the derivative of our new inverse function must be:

$$\wp'(z) = \frac{dx}{dz} = y.$$

Since the points (x, y) must satisfy $y^2 = 4x^3 - g_2x - g_3$, substituting our analytic functions yields the famous differential equation:

$$(\wp'(z))^2 = 4\wp(z)^3 - g_2\wp(z) - g_3.$$

Just as sine and cosine parametrize the circle, the functions $z \mapsto (\wp(z), \wp'(z))$ provide a complete complex analytic parametrization of the elliptic curve. Furthermore, because the complex curve has the topology of a torus (it contains two independent loops we can integrate around), the integral z accumulates two independent periods, making $\wp(z)$ and $\wp'(z)$ *doubly periodic* complex functions.

Notice that from $z = \int \frac{dx}{y}$, the differential element is exactly $dz = \frac{dx}{y}$. The invariant differential ω is literally the differential element dz of these historic analytic functions! We will see this analytic theory fully developed in chapter 3.

Exercise 1.4.1

1. Let E be an elliptic curve over a field K of characteristic 2.
 - (a) If $a_1 \neq 0$, show that the invariant differential can be written as $\omega = \frac{dx}{a_1x+a_3}$.
 - (b) What is the invariant differential if $a_1 = 0$? Explain why it still nowhere-vanishing.

2. Consider the Fermat cubic curve $X^3 + Y^3 = Z^3$. In the affine chart $Z = 1$, the equation is $x^3 + y^3 = 1$. The invariant differential on this curve (which translates up to scaling to the Weierstrass model) is given by $\omega = \frac{dx}{y^2}$. Show that $\omega = -\frac{dy}{x^2}$.
3. Let $E : y^2 = x^3 - x$ and $E' : y^2 = x^3 + 4x$. The map $\varphi : E \rightarrow E'$ given by

$$(x, y) \mapsto \left(\frac{x^2 + 1}{x}, y \frac{x^2 - 1}{x^2} \right)$$

is a non-constant morphism mapping $O \rightarrow O$ (an *isogeny*). Let $\omega' = \frac{dx'}{2y'}$ be the invariant differential on E' . Compute the pullback $\varphi^*\omega'$ and express it in terms of the invariant differential $\omega = \frac{dx}{2y}$ on E .

1.5 Automorphisms and Twists

The classification theorem (theorem 1.2.8) tells us that over an algebraically closed field, j -invariants biject with isomorphism classes of elliptic curves. Over a non-algebraically closed field K , the situation is more interesting: multiple non-isomorphic curves over K can share the same j -invariant. These are related by *twisting*, and the twist classification is governed by the *automorphism group* of the curve. This section determines $\text{Aut}(E)$ completely and uses it to classify all twists.

Automorphisms of Elliptic Curves

An *automorphism* of an elliptic curve (E, O) is an isomorphism $\varphi : E \xrightarrow{\sim} E$ of varieties that fixes the marked point: $\varphi(O) = O$. Since φ preserves O and the group law is determined by O (via Pic^0), any automorphism is automatically a group homomorphism: $\varphi(P + Q) = \varphi(P) + \varphi(Q)$ for all $P, Q \in E(K)$. (This follows from the functoriality of the Picard group construction, or can be checked directly from the geometric description of addition.)

Example 1.5.1: Hyperelliptic Involution

Every elliptic curve has at least one nontrivial automorphism: the negation map $[-1] : E \rightarrow E$ defined by $P \mapsto -P$. In terms of coordinates on a Weierstrass model, this is

$$[-1] : (x, y) \mapsto (x, -y - a_1x - a_3), \quad (1.26)$$

which in short Weierstrass form ($a_1 = a_3 = 0$) simplifies to $(x, y) \mapsto (x, -y)$. This is the *hyperelliptic involution*: it is the unique involution of E that fixes the marked point O and acts as -1 on the tangent space $T_O E$. The quotient $E/\langle [-1] \rangle$ is isomorphic to $\mathbb{P}^1$ (the x -line, via the degree-2 map $E \rightarrow \mathbb{P}^1$ given by $(x, y) \mapsto x$), and the branch points of this double cover are the four fixed points of $[-1]$: the identity O and the three 2-torsion points $E[2](\bar{K}) \setminus \{O\}$ (when $\text{char} K \neq 2$).

Thus $\text{Aut}(E)$ always contains the subgroup $\langle [-1] \rangle \cong \mathbb{Z}/2\mathbb{Z}$. The question is whether there are any *additional* automorphisms.

For a Weierstrass model, an automorphism is an admissible change of variables (u, r, s, t) that sends the equation back to itself. In the short Weierstrass form $y^2 = x^3 + Ax + B$ (with $\text{char} K \neq 2, 3$), the

analysis in exercise 1.2.1 (6) showed that such an automorphism must have $r = s = t = 0$, and so it takes the form

$$(x, y) \mapsto (u^2x, u^3y) \quad (1.27)$$

where $u \in K^*$ satisfies $u^4A = A$ and $u^6B = B$. We now give the complete classification.

Theorem 1.5.2: Automorphism Group of an Elliptic Curve

Let E/K be an elliptic curve with $\text{char}K \neq 2, 3$, given in short Weierstrass form $y^2 = x^3 + Ax + B$. Then:

1. If $j(E) \neq 0, 1728$ (i.e., $AB \neq 0$), then $\text{Aut}(E) = \{\pm 1\} \cong \mathbb{Z}/2\mathbb{Z}$, generated by $[-1] : (x, y) \mapsto (x, -y)$.
2. If $j(E) = 1728$ (i.e., $B = 0, A \neq 0$), then $\text{Aut}(E) \cong \mathbb{Z}/4\mathbb{Z}$, generated by $\iota : (x, y) \mapsto (-x, iy)$, where $i \in K$ is a primitive fourth root of unity.
3. If $j(E) = 0$ (i.e., $A = 0, B \neq 0$), then $\text{Aut}(E) \cong \mathbb{Z}/6\mathbb{Z}$, generated by $\zeta : (x, y) \mapsto (\omega x, -y)$, where $\omega \in K$ is a primitive cube root of unity.

Here (2) requires that K contains a primitive fourth root of unity, and (3) requires a primitive cube root; otherwise $\text{Aut}_K(E)$ is the subgroup of the stated group that is defined over K .

Proof:

From (1.27), an automorphism corresponds to $u \in K^*$ with $u^4A = A$ and $u^6B = B$.

Case 1: $A \neq 0$ and $B \neq 0$. Then $u^4 = 1$ and $u^6 = 1$, which together give $u^2 = u^{\gcd(4,6)} = 1$, so $u = \pm 1$. The automorphism $u = 1$ is the identity, and $u = -1$ gives $(x, y) \mapsto ((-1)^2x, (-1)^3y) = (x, -y)$, the hyperelliptic involution. Thus $\text{Aut}(E) \cong \mathbb{Z}/2\mathbb{Z}$.

Case 2: $B = 0, A \neq 0$ (so $j = 1728$). The equation is $y^2 = x^3 + Ax$. The condition $u^6B = B$ is automatic, and $u^4A = A$ forces $u^4 = 1$. Over $\overline{K}$, the fourth roots of unity are $\{1, -1, i, -i\}$. The automorphism corresponding to $u = i$ is

$$(x, y) \mapsto (i^2x, i^3y) = (-x, -iy).$$

Write $\varphi(x, y) = (-x, -iy)$. We compute its iterates:

$$\varphi^2(x, y) = \varphi(-x, -iy) = (-(-x), -i(-iy)) = (x, i^2y) = (x, -y),$$

$$\varphi^3(x, y) = \varphi(x, -y) = (-x, -i(-y)) = (-x, iy),$$

$$\varphi^4(x, y) = \varphi(-x, iy) = (x, -i \cdot iy) = (x, -i^2y) = (x, y).$$

So φ has order 4, $\varphi^2 = [-1]$, and $\text{Aut}(E) = \langle \varphi \rangle \cong \mathbb{Z}/4\mathbb{Z}$.

One verifies that φ preserves the equation: if $y^2 = x^3 + Ax$, then $(-iy)^2 = -y^2$ and $(-x)^3 + A(-x) = -x^3 - Ax = -(x^3 + Ax) = -y^2$, so $(-iy)^2 = -y^2 = (-x)^3 + A(-x)$.

Case 3: $A = 0, B \neq 0$ (so $j = 0$). The equation is $y^2 = x^3 + B$. The condition $u^4A = A$ is automatic, and $u^6B = B$ forces $u^6 = 1$. Over $\overline{K}$, the sixth roots of unity are $\{1, -1, \omega, -\omega, \omega^2, -\omega^2\}$

where $\omega = e^{2\pi i/3}$. The automorphism corresponding to $u = -\omega^2$ (a primitive sixth root of unity, since $(-\omega^2)^6 = \omega^{12} = 1$ and $(-\omega^2)^3 = -\omega^6 = -1 \neq 1$, $(-\omega^2)^2 = \omega^4 = \omega \neq 1$) is

$$\psi(x, y) = ((-\omega^2)^2 x, (-\omega^2)^3 y) = (\omega^4 x, -\omega^6 y) = (\omega x, -y).$$

Let us verify this has order 6:

$$\begin{aligned}\psi^1(x, y) &= (\omega x, -y), \\ \psi^2(x, y) &= (\omega \cdot \omega x, -(-y)) = (\omega^2 x, y), \\ \psi^3(x, y) &= (\omega \cdot \omega^2 x, -y) = (x, -y) = [-1](x, y), \\ \psi^4(x, y) &= (\omega x, y), \\ \psi^5(x, y) &= (\omega^2 x, -y), \\ \psi^6(x, y) &= (x, y).\end{aligned}$$

So ψ has order 6, $\psi^3 = [-1]$, ψ^2 has order 3, and $\text{Aut}(E) = \langle \psi \rangle \cong \mathbb{Z}/6\mathbb{Z}$.

Verification: if $y^2 = x^3 + B$, then $(-y)^2 = y^2 = x^3 + B$ and $(\omega x)^3 = \omega^3 x^3 = x^3$ (since $\omega^3 = 1$). So $(-y)^2 = (\omega x)^3 + B$.

The automorphism groups are cyclic of orders 2, 4, and 6 — precisely the finite subgroups of K^* that arise from the condition $u^{\text{lcm}(4,6)} = u^{12} = 1$ combined with either $u^4 = 1$, $u^6 = 1$, or both. The geometric content is that the extra automorphisms at $j = 1728$ and $j = 0$ come from the extra symmetry of the cubic $x^3 + Ax$ (invariant under $x \mapsto -x$) and $x^3 + B$ (invariant under $x \mapsto \omega x$), respectively.

Example 1.5.3: Automorphisms of the Running Examples

(a) $E : y^2 = x^3 - x$ ($j = 1728$). The automorphism $\varphi : (x, y) \mapsto (-x, -iy)$ has order 4, with $\varphi^2 = [-1]$. Over $\mathbb{Q}$, the element $i = \sqrt{-1}$ does not lie in $\mathbb{Q}$, so φ is not defined over $\mathbb{Q}$. However, φ is defined over $\mathbb{Q}(i)$, and $\text{Aut}_{\mathbb{Q}}(E) = \{\pm 1\} \cong \mathbb{Z}/2\mathbb{Z}$, while $\text{Aut}_{\mathbb{Q}(i)}(E) \cong \mathbb{Z}/4\mathbb{Z}$. The map φ is the complex multiplication by i in the endomorphism ring $\text{End}(E) \cong \mathbb{Z}[i]$.

(b) $E : y^2 = x^3 + 1$ ($j = 0$). The automorphism $\psi : (x, y) \mapsto (\omega x, -y)$ has order 6. Over $\mathbb{Q}$, neither ω nor $-\omega^2$ lies in $\mathbb{Q}$, but $\psi^3 = [-1]$ is defined over $\mathbb{Q}$, and $\psi^2 : (x, y) \mapsto (\omega^2 x, y)$ is defined over $\mathbb{Q}(\omega)$. Thus $\text{Aut}_{\mathbb{Q}}(E) \cong \mathbb{Z}/2\mathbb{Z}$ and $\text{Aut}_{\mathbb{Q}(\omega)}(E) \cong \mathbb{Z}/6\mathbb{Z}$. The map ψ^2 is the complex multiplication by ω in the endomorphism ring $\text{End}(E) \cong \mathbb{Z}[\omega]$.

Twists of Elliptic Curves

The classification theorem tells us that over $\overline{K}$, the j -invariant is a complete invariant. Over K itself, two elliptic curves with the same j -invariant need not be isomorphic; the obstruction is measured by the *twists*.

Definition 1.5.4: Twist

Let E/K be an elliptic curve. A *twist* of E is an elliptic curve E'/K such that $E'_K \cong E_{\overline{K}}$ (i.e., E and E' become isomorphic after base change to the algebraic closure). Equivalently, E' is a twist of E if and only if $j(E') = j(E)$.

The set of K -isomorphism classes of twists of E is denoted $\text{Twist}(E/K)$.

We now classify twists in all three cases. Recall from theorem 1.2.8 that an isomorphism between short Weierstrass models $y^2 = x^3 + Ax + B$ and $y^2 = x^3 + A'x + B'$ over $\overline{K}$ has the form $(x, y) \mapsto (u^2x, u^3y)$ with $A' = u^{-4}A$ and $B' = u^{-6}B$. Such an isomorphism is defined over K if and only if $u \in K^*$.

The simplest and most important type of twist is the quadratic twist. This is the generic case: it accounts for all twists when $j \neq 0, 1728$.

Definition 1.5.5: Quadratic Twist

Let $E : y^2 = x^3 + Ax + B$ be an elliptic curve over K (with $\text{char} K \neq 2$) and let $d \in K^*$. The *quadratic twist of E by d* is the elliptic curve

$$E^{(d)} : dy^2 = x^3 + Ax + B$$

(which is isomorphic to E over $K(\sqrt{d})$ via $y \mapsto y/\sqrt{d}$). Equivalently, applying the change of variables $(x, y) \mapsto (x/d, y/d^2)$ gives the standard Weierstrass form:

$$E^{(d)} : y^2 = x^3 + d^2Ax + d^3B. \quad (1.28)$$

The two curves E and $E^{(d)}$ become isomorphic over $K(\sqrt{d})$. Concretely, starting from $E^{(d)} : y^2 = x^3 + d^2Ax + d^3B$ and $E : y^2 = x^3 + Ax + B$, an admissible change with $u = 1/\sqrt{d}$ gives $u^{-4} = d^2$ and $u^{-6} = d^3$, so the map $(x, y) \mapsto (u^2x, u^3y) = (x/d, y/d\sqrt{d})$ sends a point on E to a point on $E^{(d)}$. This isomorphism requires $\sqrt{d}$ and so is defined over $K(\sqrt{d})$ but not over K (unless $d \in (K^*)^2$).

Proposition 1.5.6: Quadratic Twists and j -Invariants

Let $E : y^2 = x^3 + Ax + B$ and let $d \in K^*$.

1. $j(E^{(d)}) = j(E)$ for all d .
2. $E^{(d)} \cong_K E$ if and only if $d \in (K^*)^2$.
3. $E^{(d)} \cong_K E^{(d')}$ if and only if $d/d' \in (K^*)^2$.
4. If $\text{char} K \neq 2$, the map $d \mapsto E^{(d)}$ induces a bijection

$$K^*/(K^*)^2 \xrightarrow{\sim} \text{Twist}(E/K) \quad \text{when } j(E) \neq 0, 1728.$$

Proof:

1. From (1.28), $E^{(d)}$ has coefficients $A' = d^2A$ and $B' = d^3B$, so

$$\Delta(E^{(d)}) = -16(4(d^2A)^3 + 27(d^3B)^2) = -16d^6(4A^3 + 27B^2) = d^6\Delta(E)$$

and $c_4(E^{(d)}) = -48d^2A = d^2c_4(E)$. Therefore

$$j(E^{(d)}) = \frac{c_4(E^{(d)})^3}{\Delta(E^{(d)})} = \frac{d^6 c_4(E)^3}{d^6 \Delta(E)} = j(E).$$

2. The curves E and $E^{(d)}$ are isomorphic over K iff there exists $u \in K^*$ with $d^2A = u^{-4}A$ and $d^3B = u^{-6}B$. If $A \neq 0$, the first equation gives $u^4 = d^{-2}$, so $u^2 = \pm d^{-1}$. If also $B \neq 0$ (which holds when $j \neq 0, 1728$), the second gives $u^6 = d^{-3}$, so $u^2 = d^{-1}$ (since $u^6 = (u^2)^3$ must be compatible with $u^4 = (u^2)^2$). Thus $d = u^{-2} \in (K^*)^2$.
3. $E^{(d)} \cong_K E^{(d')}$ iff $E^{(d/d')} \cong_K E$ (by composing twists), iff $d/d' \in (K^*)^2$.
4. When $j \neq 0, 1728$, every twist of E is a quadratic twist. Indeed, if E' is a twist with $j(E') = j(E)$, write both in short Weierstrass form. The isomorphism $E'_K \cong E_K$ gives $A' = u^{-4}A$ and $B' = u^{-6}B$ for some $u \in \overline{K}^*$. Since $AB \neq 0$, we get $u^2 = (A/A')^{1/2}$ from the first equation and $u^2 = (B/B')^{1/3}$ from the second. Set $d = u^{-2} \in \overline{K}^*$; then $A' = d^2A$ and $B' = d^3B$. But $A', B', A, B \in K$, so $d^2 = A'/A \in K$ and $d^3 = B'/B \in K$, hence $d = d^3/d^2 = (B'/B)/(A'/A) \in K$. Thus $E' = E^{(d)}$ with $d \in K^*$, and the twist is quadratic. The classification by $K^*/(K^*)^2$ follows from (3).

Example 1.5.7: Quadratic Twists of the Congruent Number Curves

Let $E : y^2 = x^3 - x + 1$ (with $A = -1$, $B = 1$, so $j \neq 0, 1728$). The twist by $d = 5$ is

$$E^{(5)} : y^2 = x^3 - 25x + 125.$$

We have $j(E^{(5)}) = j(E)$, and E and $E^{(5)}$ become isomorphic over $\mathbb{Q}(\sqrt{5})$ via $u = 1/\sqrt{5}$. Over $\mathbb{Q}$ they are non-isomorphic, since $5 \notin (\mathbb{Q}^*)^2$.

When $j = 1728$ or $j = 0$, the automorphism group is larger than $\mathbb{Z}/2\mathbb{Z}$, and there exist twists that are not quadratic.

Proposition 1.5.8: Quartic Twists ($j = 1728$)

Let $E : y^2 = x^3 + Ax$ over K with $\text{char} K \neq 2, 3$, $A \neq 0$, and $j(E) = 1728$. Then $\text{Twist}(E/K) \cong K^*/(K^*)^4$, via the map

$$d \mapsto E^{[d]} : y^2 = x^3 + dAx.$$

In particular, two curves $y^2 = x^3 + Ax$ and $y^2 = x^3 + A'x$ (with $A, A' \neq 0$) are isomorphic over K if and only if $A'/A \in (K^*)^4$.

Proof:

An isomorphism between $y^2 = x^3 + Ax$ and $y^2 = x^3 + A'x$ over K requires $u \in K^*$ with $u^4 A = A'$ (the condition on $B = B' = 0$ is automatic). Thus $A'/A = u^4 \in (K^*)^4$. For the twist by $d \in K^*$: the curve $y^2 = x^3 + dAx$ becomes isomorphic to E over $K(d^{1/4})$ via $u = d^{-1/4}$.

Example 1.5.9: Quartic Twist

The congruent number curves $E_n : y^2 = x^3 - n^2x$ are related to $E_1 : y^2 = x^3 - x$ by the formula (1.28) with $A = -1$, $B = 0$, $d = n$: the twist has $A' = n^2 \cdot (-1) = -n^2$ and $B' = n^3 \cdot 0 = 0$. As $j(E_1) = 1728$ (since $B = 0$), these are *quartic* twist, and correspond to the class of n in $\mathbb{Q}^*/(\mathbb{Q}^*)^4$.

Proposition 1.5.10: Sextic Twists ($j = 0$)

Let $E : y^2 = x^3 + B$ over K with $\text{char} K \neq 2, 3$, $B \neq 0$, and $j(E) = 1728$. Then $\text{Twist}(E/K) \cong K^*/(K^*)^6$, via the map

$$d \mapsto E^{\{d\}} : y^2 = x^3 + dB.$$

In particular, two curves $y^2 = x^3 + B$ and $y^2 = x^3 + B'$ (with $B, B' \neq 0$) are isomorphic over K if and only if $B'/B \in (K^*)^6$.

Proof:

An isomorphism requires $u \in K^*$ with $u^6 B = B'$, giving $B'/B = u^6 \in (K^*)^6$.

Let us record the full picture as a summary.

Theorem 1.5.11: Classification of Twists

Let E/K be an elliptic curve over a field with $\text{char} K \neq 2, 3$. The set of K -isomorphism classes of twists of E is classified as follows:

$j(E)$	$\text{Aut}_{\overline{K}}(E)$	$\text{Twist}(E/K)$
$\neq 0, 1728$	$\mathbb{Z}/2\mathbb{Z}$	$K^*/(K^*)^2$
$= 1728$	$\mathbb{Z}/4\mathbb{Z}$	$K^*/(K^*)^4$
$= 0$	$\mathbb{Z}/6\mathbb{Z}$	$K^*/(K^*)^6$

The pattern is clear: the twist set is $K^*/(K^*)^n$ where $n = |\text{Aut}_{\overline{K}}(E)|$. This is not a coincidence. The general theory of Galois cohomology (which we will develop in chapter 7) shows that for any algebraic object X over K , the set of twists (forms) of X is classified by the cohomology set $H^1(\text{Gal}(\overline{K}/K), \text{Aut}_{\overline{K}}(X))$. For an elliptic curve with $\text{Aut}_{\overline{K}}(E) = \mu_n$ (the group of n -th roots of unity), this cohomology set is isomorphic to $K^*/(K^*)^n$ by Kummer theory. We state this now for future reference but defer the full proof:

Proposition 1.5.12: Twists via Galois Cohomology

For any elliptic curve E/K (with $\text{char } K \neq 2, 3$),

$$\text{Twist}(E/K) \cong H^1(\text{Gal}(\overline{K}/K), \text{Aut}_{\overline{K}}(E)).$$

The bijection sends a 1-cocycle $\sigma \mapsto a_\sigma \in \text{Aut}_{\overline{K}}(E)$ to the twist E' whose $\overline{K}$ -isomorphism $\varphi : E_{\overline{K}} \xrightarrow{\sim} E'_{\overline{K}}$ satisfies $\varphi^{-1} \circ \sigma \varphi = a_\sigma$ for all $\sigma \in \text{Gal}(\overline{K}/K)$.

This is the same mechanism as the twisting of finite group schemes described in [7] and is closely related to the twisting of varieties in [6]. The explicit description via $K^*/(K^*)^n$ that we have derived in this section is the “unrolled” version of this cohomological statement.

Example 1.5.13: Twists over $\mathbb{Q}$

Over $\mathbb{Q}$, the group $\mathbb{Q}^*/(\mathbb{Q}^*)^2$ is enormous (it is a vector space over $\mathbb{F}_2$ with a basis indexed by the primes, together with the class of -1). Thus there are infinitely many quadratic twists of any elliptic curve over $\mathbb{Q}$, one for each squarefree integer d . For instance, the curves $y^2 = x^3 - d^2x$ for squarefree $d \geq 1$ give a complete set of representatives for the quadratic twists of $y^2 = x^3 - x$.

The group $\mathbb{Q}^*/(\mathbb{Q}^*)^4$ is even larger: the class of $d \in \mathbb{Q}^*$ is trivial iff d is a fourth power of a rational number. For $j = 1728$, the curves $y^2 = x^3 + dx$ for d ranging over representatives of $\mathbb{Q}^*/(\mathbb{Q}^*)^4$ give all twists.

Similarly, $\mathbb{Q}^*/(\mathbb{Q}^*)^6$ classifies sextic twists of $j = 0$ curves.

Example 1.5.14: Twists over Finite Fields

Over a finite field $\mathbb{F}_q$ with q odd, the group $\mathbb{F}_q^*/(\mathbb{F}_q^*)^2$ has order 2 (every element of $\mathbb{F}_q^*$ is either a square or a nonsquare). Thus a generic elliptic curve over $\mathbb{F}_q$ has exactly one nontrivial quadratic twist.

Concretely, let $E : y^2 = x^3 + Ax + B$ over $\mathbb{F}_q$ and let $\delta \in \mathbb{F}_q^*$ be a fixed nonsquare. The unique nontrivial quadratic twist is $E^{(\delta)} : y^2 = x^3 + \delta^2 Ax + \delta^3 B$. The point counts of E and $E^{(\delta)}$ are related by the elegant formula

$$|E(\mathbb{F}_q)| + |E^{(\delta)}(\mathbb{F}_q)| = 2(q + 1). \quad (1.29)$$

This is because a point (x_0, y_0) with $y_0 \neq 0$ lies on E iff $x_0^3 + Ax_0 + B$ is a nonzero square, and it lies on $E^{(\delta)}$ iff $x_0^3 + Ax_0 + B$ is δ times a square — so exactly one of E and $E^{(\delta)}$ gets the point (with both $\pm y_0$), while points with $y_0 = 0$ lie on both curves. Summing over all $x_0 \in \mathbb{F}_q$ and adding the points at infinity gives (1.29). We will revisit this formula in chapter 5 when we study the zeta function.

Exercise 1.5.1

1. Let $E : y^2 = x^3 + 1$ over $\mathbb{Q}(\omega)$, where $\omega = (-1 + \sqrt{-3})/2$. Write down all six automorphisms

- of E explicitly as maps $(x, y) \mapsto (\cdot, \cdot)$, verify that they form a cyclic group of order 6, and identify the generator.
2. Let K be an algebraically closed field of characteristic 2, and let $E : y^2 + y = x^3$ over K (this is a supersingular elliptic curve with $j = 0$). Show that $\text{Aut}(E)$ is nonabelian, and determine its order. (*Hint:* In characteristic 2, the automorphism group can be larger than $\mathbb{Z}/6\mathbb{Z}$; this curve admits automorphisms of the form $(x, y) \mapsto (x + s^2, y + sx + t)$.)
 3. Let $E : y^2 = x^3 + 3x + 7$ over $\mathbb{Q}$. Show that $j(E) \neq 0, 1728$. Determine the quadratic twist $E^{(-1)}$ of E by $d = -1$, and compute $\Delta(E)$ and $\Delta(E^{(-1)})$. How do the discriminants compare?
 4. Let $E : y^2 = x^3 + 1$ over $\mathbb{F}_7$.
 - (a) Count $|E(\mathbb{F}_7)|$ directly.
 - (b) Let $\delta = 3$ (a nonsquare in $\mathbb{F}_7^*$). Write down the quadratic twist $E^{(\delta)}$ and count $|E^{(\delta)}(\mathbb{F}_7)|$.
 - (c) Verify the formula $|E(\mathbb{F}_7)| + |E^{(\delta)}(\mathbb{F}_7)| = 2(7 + 1) = 16$.
 5. Over $\mathbb{Q}$, the curves $y^2 = x^3 + 1$ and $y^2 = x^3 + 4$ both have $j = 0$. Are they isomorphic over $\mathbb{Q}$? Over $\mathbb{Q}$? If they are twists, what element of $\mathbb{Q}^*/(\mathbb{Q}^*)^6$ classifies the twist?
 6. Let $E : y^2 = x^3 - x$ over $\mathbb{Q}$ and $E^{(d)} : y^2 = x^3 - d^2x$ for squarefree $d > 0$.
 - (a) Show that $E^{(d)}[2](\mathbb{Q}) \cong \mathbb{Z}/2\mathbb{Z} \times \mathbb{Z}/2\mathbb{Z}$ for every d .
 - (b) Look up (e.g., in the LMFDB database) the ranks of $E^{(d)}$ for $d = 1, 2, 3, 5, 6, 7$. What pattern, if any, do you observe? (*The connection to the congruent number problem will be explored in chapter 8.*)

1.6 Singular Cubics

Throughout this chapter we have required $\Delta \neq 0$, ensuring that the Weierstrass equation defines a smooth curve — an elliptic curve. What happens when $\Delta = 0$? The projective cubic still exists as a variety, but it has a singular point, and it no longer carries the structure of an elliptic curve. Nevertheless, the *nonsingular locus* of a singular Weierstrass cubic retains a group structure from the chord-and-tangent law, and the resulting group turns out to be one of two familiar objects: the multiplicative group $\mathbb{G}_m$ or the additive group $\mathbb{G}_a$. This classification closes the chapter by connecting the theory of elliptic curves to the classical algebraic groups the reader already knows.

The results of this section also serve a concrete purpose: when we study reduction of elliptic curves modulo a prime in chapter 6, the reduced curve will sometimes be singular, and the type of singularity (node or cusp) will determine the *reduction type* (multiplicative or additive) of the original curve. Understanding singular Weierstrass cubics here is an investment that pays off immediately in the local theory.

The Singular Point

Assume $\text{char} K \neq 2, 3$ and consider the short Weierstrass equation $E : y^2 = x^3 + Ax + B$ with $\Delta = -16(4A^3 + 27B^2) = 0$, so $4A^3 + 27B^2 = 0$. By the analysis in section 1.2, a singular point

(x_0, y_0) must satisfy $y_0 = 0$ and $3x_0^2 + A = 0$, as well as $x_0^3 + Ax_0 + B = 0$. From $3x_0^2 = -A$, we get $x_0^3 = -Ax_0/3$, and substituting into the curve equation:

$$0 = x_0^3 + Ax_0 + B = -\frac{Ax_0}{3} + Ax_0 + B = \frac{2Ax_0}{3} + B,$$

so $x_0 = -3B/(2A)$ (assuming $A \neq 0$; the case $A = 0$ forces $B = 0$ as well, which we handle separately). Thus the singular point is unique:

$$P_{\text{sing}} = \left(-\frac{3B}{2A}, 0 \right). \quad (1.30)$$

We can simplify the equation by translating the singular point to the origin. Starting from $y^2 = x^3 + Ax + B$ and substituting $x \mapsto x + x_0$ where $x_0 = -3B/(2A)$:

$$\begin{aligned} y^2 &= (x + x_0)^3 + A(x + x_0) + B \\ &= x^3 + 3x_0x^2 + 3x_0^2x + x_0^3 + Ax + Ax_0 + B \\ &= x^3 + 3x_0x^2 + (3x_0^2 + A)x + (x_0^3 + Ax_0 + B). \end{aligned}$$

Since $P_{\text{sing}} = (x_0, 0)$ lies on the curve, $x_0^3 + Ax_0 + B = 0$, so the constant term vanishes. Since $3x_0^2 + A = 0$ (the singularity condition), the linear term vanishes. Writing $\alpha = 3x_0$, we obtain:

$$y^2 = x^3 + \alpha x^2, \quad \text{where } \alpha = 3x_0 = -\frac{9B}{2A}. \quad (1.31)$$

The right-hand side factors as $x^2(x + \alpha)$, so we can write:

$$y^2 = x^2(x + \alpha). \quad (1.32)$$

Nodes and Cusps

The type of singularity at the origin is determined by the *tangent cone*: the lowest-degree homogeneous part of the defining equation $f(x, y) = y^2 - x^3 - \alpha x^2$. Dropping the cubic term x^3 (which has higher order), the tangent cone is

$$y^2 - \alpha x^2 = 0, \quad \text{i.e.,} \quad y^2 = \alpha x^2.$$

Definition 1.6.1: Node and Cusp

Let C be a Weierstrass cubic with a singular point at the origin, given by $y^2 = x^2(x + \alpha)$.

1. If $\alpha \neq 0$, the tangent cone $y^2 = \alpha x^2$ defines two distinct tangent directions $y = \pm\sqrt{\alpha}x$ (over $\overline{K}$). The singularity is a *node* (also called an *ordinary double point*).
2. If $\alpha = 0$, the equation becomes $y^2 = x^3$ and the tangent cone degenerates to $y^2 = 0$, i.e., the single line $y = 0$ with multiplicity 2. The singularity is a *cusp*.

The distinction between nodes and cusps can be read off from the original invariants. Translating back: $\alpha = 0$ occurs when $x_0 = 0$, which by $3x_0^2 + A = 0$ forces $A = 0$, and then $x_0^3 + B = 0$ gives $B = 0$. Both A and B vanish, so $c_4 = -48A = 0$ and the curve is $y^2 = x^3$. In general:

Proposition 1.6.2: Node vs. Cusp via c_4

Let $\Delta = 0$ and $\text{char} K \neq 2, 3$. The singular Weierstrass cubic has:

1. a *node* if $c_4 \neq 0$ (equivalently, the cubic $x^3 + Ax + B$ has a double root but not a triple root);
2. a *cusp* if $c_4 = 0$ (equivalently, $A = B = 0$, and the cubic has a triple root at $x = 0$, giving $y^2 = x^3$).

Proof:

From $c_4 = -48A$: $c_4 = 0$ iff $A = 0$. If $A = 0$ and $\Delta = 0$, then $27B^2 = 0$, so $B = 0$. The equation reduces to $y^2 = x^3$: cusp. If $c_4 \neq 0$ (i.e., $A \neq 0$), then $\alpha = -9B/(2A) \neq 0$ (since $B \neq 0$, as $4A^3 + 27B^2 = 0$ with $A \neq 0$ forces $B \neq 0$), so the tangent cone has two distinct lines: node.

Example 1.6.3: A Node and a Cusp

(a) *Node*. The curve $y^2 = x^3 + x^2$ over $\mathbb{Q}$ has $A_{\text{long}} = 0$, $a_2 = 1$ in the translated form (it is already in the form (1.31) with $\alpha = 1$). Let us work with the short Weierstrass perspective: it has a singularity at $(0, 0)$ with tangent cone $y^2 = x^2$, i.e., two tangent lines $y = x$ and $y = -x$. These lines are defined over $\mathbb{Q}$: the node is *split*.

(b) *Node (non-split)*. The curve $y^2 = x^3 - x^2$ has $\alpha = -1$ and tangent cone $y^2 = -x^2$. Over $\mathbb{Q}$, this has no real solutions other than $x = y = 0$; the tangent lines $y = \pm ix$ are defined over $\mathbb{Q}(i)$ but not over $\mathbb{Q}$. The node is *non-split* over $\mathbb{Q}$.

(c) *Cusp*. The curve $y^2 = x^3$ has $A = B = 0$ and $\alpha = 0$. The tangent cone is $y^2 = 0$, a double line. This is a cusp.

The Group Structure of the Nonsingular Locus

The singular point must be excluded from any group law (it has no well-defined tangent line in the sense needed for the chord-and-tangent construction). Let E_{ns} denote the set of nonsingular points of the Weierstrass cubic C , including the point at infinity O . The chord-and-tangent construction restricts to a well-defined group law on E_{ns} : if $P, Q \in E_{\text{ns}}$ and the line $\overline{PQ}$ meets C at a third point R , then R is also nonsingular (?? bellow), and the reflection of R remains nonsingular. Thus E_{ns} is a subgroup of the “would-be” group on C .

Proposition 1.6.4: Closure of the Nonsingular Locus

Let C be a singular Weierstrass cubic over a field K , and let P_{sing} be its singular point. If $P, Q \in E_{\text{ns}}$, then the third point of intersection R of the line $L = \overline{PQ}$ with C is also in E_{ns} . Furthermore, the reflection $-R$ is in E_{ns} .

Proof:

By Bézout's Theorem, the line L and the cubic C intersect at exactly three points in the projective plane $\mathbb{P}^2(\overline{K})$, counting intersection multiplicities.

Since P_{sing} is a singularity of C , the intersection multiplicity of C with *any* line passing through P_{sing} is at least 2 at P_{sing} .

Suppose for contradiction that the third intersection point R is the singular point P_{sing} .

- **Case 1** ($P \neq Q$): The line L is the secant line passing through P , Q , and P_{sing} . The sum of the intersection multiplicities would be at least 1 (for P) + 1 (for Q) + 2 (for P_{sing}) = 4. This exceeds the total of 3 allowed by Bézout's Theorem, a contradiction.
- **Case 2** ($P = Q$): The line L is the tangent line to C at the nonsingular point P . Thus, L intersects C at P with multiplicity at least 2. If L also passes through P_{sing} , the total intersection multiplicity would be at least 2 (for P) + 2 (for P_{sing}) = 4, again contradicting Bézout's Theorem.

Therefore, R cannot be the singular point, meaning $R \in E_{\text{ns}}$.

Finally, we verify that the reflection $-R$ is nonsingular. For a singular Weierstrass cubic $y^2 = f(x)$, the singular point must be a root of both the curve equation and its partial derivatives. Since $\partial/\partial y(y^2 - f(x)) = 2y$, the singular point must lie on the x -axis. Thus $P_{\text{sing}} = (x_0, 0)$.

Because $R \in E_{\text{ns}}$, we know $R = (x_R, y_R) \neq (x_0, 0)$. If x_R were equal to x_0 , then we would have $y_R^2 = f(x_0) = 0$, forcing $R = (x_0, 0) = P_{\text{sing}}$, a contradiction. Thus $x_R \neq x_0$. Since the reflection $-R = (x_R, -y_R)$ shares the same x -coordinate as R , it also avoids x_0 . Therefore $-R \neq P_{\text{sing}}$, meaning $-R \in E_{\text{ns}}$.

The key theorem identifies this group explicitly.

Theorem 1.6.5: Group of the Nonsingular Locus

Let C be a singular Weierstrass cubic over K with $\text{char} K \neq 2, 3$, and let $E_{\text{ns}} = C \setminus \{P_{\text{sing}}\}$ be its nonsingular locus. Then:

1. (*Node — multiplicative type.*) If $c_4 \neq 0$, then

$$E_{\text{ns}}(\overline{K}) \cong \overline{K}^* \cong \mathbb{G}_m(\overline{K}).$$

Over K itself, two cases arise:

- (a) (*Split node.*) If $\alpha \in (K^*)^2$ (i.e., the tangent directions at the node are defined over K), then $E_{\text{ns}}(K) \cong K^* \cong \mathbb{G}_m(K)$.
- (b) (*Non-split node.*) If $\alpha \notin (K^*)^2$, then $E_{\text{ns}}(K) \cong \ker(\text{Nm}_{L/K} : L^* \rightarrow K^*)$, where $L = K(\sqrt{\alpha})$.

2. (*Cusp — additive type.*) If $c_4 = 0$, then $E_{\text{ns}}(K) \cong K \cong \mathbb{G}_a(K)$.

Proof:

We give explicit parametrizations in each case, using the translated form $y^2 = x^2(x + \alpha)$.

Cusp ($\alpha = 0$): the equation is $y^2 = x^3$. A nonsingular affine point (x_0, y_0) with $y_0^2 = x_0^3$ and $(x_0, y_0) \neq (0, 0)$ can be parametrized by the slope $t = y_0/x_0$. From $y_0 = tx_0$ and $y_0^2 = x_0^3$:

$$t^2 x_0^2 = x_0^3, \quad \text{so} \quad x_0 = t^2, \quad y_0 = t^3$$

(dividing by $x_0 \neq 0$). This defines a bijection

$$\varphi : K \longrightarrow E_{\text{ns}}(K), \quad t \longmapsto \begin{cases} (t^2, t^3) & \text{if } t \neq 0, \\ O & \text{if } t = 0, \end{cases} \quad (1.33)$$

with inverse $(x, y) \mapsto y/x$ for affine points and $O \mapsto 0$. We verify that φ is a group isomorphism from $(\mathbb{G}_a(K), +)$ to $(E_{\text{ns}}(K), +)$ by checking the chord-and-tangent law directly.

Take $P_1 = (t_1^2, t_1^3)$ and $P_2 = (t_2^2, t_2^3)$ with $t_1 \neq t_2$ and both nonzero. The slope of the line $\overline{P_1 P_2}$ is

$$\lambda = \frac{t_2^3 - t_1^3}{t_2^2 - t_1^2} = \frac{(t_2 - t_1)(t_2^2 + t_1 t_2 + t_1^2)}{(t_2 - t_1)(t_2 + t_1)} = \frac{t_1^2 + t_1 t_2 + t_2^2}{t_1 + t_2}.$$

By the addition formula for $y^2 = x^3$ (short Weierstrass with $A = B = 0$): $x_3 = \lambda^2 - t_1^2 - t_2^2$. We compute

$$\lambda^2 = \frac{(t_1^2 + t_1 t_2 + t_2^2)^2}{(t_1 + t_2)^2},$$

so that after simplification, the third intersection point $R = (x_3, y_3)$ has t -parameter $t_3 = y_3/x_3 = -(t_1 + t_2)$. The reflection (negation) changes the sign of y , hence $t \mapsto -t$, so the sum has parameter $t_1 + t_2$. Thus $\varphi(t_1) + \varphi(t_2) = \varphi(t_1 + t_2)$, confirming φ is a group homomorphism from $(\mathbb{G}_a, +)$.

Node ($\alpha \neq 0$): the equation is $y^2 = x^2(x + \alpha)$. We parametrize by $t = y/x$ (the slope from the node). From $y = tx$:

$$t^2 x^2 = x^2(x + \alpha), \quad \text{so} \quad x = t^2 - \alpha, \quad y = t(t^2 - \alpha).$$

This parametrizes all nonsingular affine points except those with $x = 0$ (which would require $t^2 = \alpha$). We define a map

$$\psi : \overline{K}^* \longrightarrow E_{\text{ns}}(\overline{K}), \quad t \longmapsto (t^2 - \alpha, t(t^2 - \alpha)) \quad (1.34)$$

for $t^2 \neq \alpha$ (i.e., $t \neq \pm\sqrt{\alpha}$), with the two “missing” slopes $t = \pm\sqrt{\alpha}$ corresponding to the tangent directions at the node (which are the “limits” as one approaches the singular point along each branch).

To make this a genuine parametrization by $\mathbb{G}_m$, we change coordinates. Over $\overline{K}$, we can factor $y^2 = x^2(x + \alpha)$ by writing $y = tx$ and working with the parameter

$$u = \frac{y + \sqrt{\alpha}x}{y - \sqrt{\alpha}x} = \frac{t + \sqrt{\alpha}}{t - \sqrt{\alpha}} \in \overline{K}^*.$$

This is well-defined for $t \neq \pm\sqrt{\alpha}$, and the inverse is $t = \sqrt{\alpha}(u+1)/(u-1)$. One verifies (by a calculation analogous to the cuspidal case) that the chord-and-tangent group law on $E_{\text{ns}}(\bar{K})$ corresponds to *multiplication* in $\bar{K}^*$ under this parametrization:

$$\psi(u_1) + \psi(u_2) = \psi(u_1 u_2).$$

Thus $E_{\text{ns}}(\bar{K}) \cong \mathbb{G}_m(\bar{K})$. The identity element O corresponds to $u = 1$ (i.e., $t = \infty$, which is the “vertical” direction).

Over K itself, the parameter $u = (t + \sqrt{\alpha})/(t - \sqrt{\alpha})$ is defined over K if and only if $\sqrt{\alpha} \in K$:

- If $\alpha \in (K^*)^2$ (split node), then $u \in K^*$, and $E_{\text{ns}}(K) \cong K^* = \mathbb{G}_m(K)$.
- If $\alpha \notin (K^*)^2$ (non-split node), set $L = K(\sqrt{\alpha})$ and $\beta = \sqrt{\alpha}$. Then $u = (t + \beta)/(t - \beta) \in L^*$, and the Galois conjugate is $\bar{u} = (t - \beta)/(t + \beta) = 1/u$. A K -rational point on E_{ns} corresponds to a parameter $u \in L^*$ satisfying $\bar{u} = 1/u$, i.e., $u\bar{u} = 1$, which is exactly the condition $\text{Nm}_{L/K}(u) = 1$. Thus $E_{\text{ns}}(K) \cong \ker(\text{Nm}_{L/K} : L^* \rightarrow K^*) = L^{*, \text{Nm}=1}$, the group of norm-1 elements.

The appearance of $\mathbb{G}_m$ and $\mathbb{G}_a$ is deeply satisfying from the perspective of algebraic groups. Over an algebraically closed field, there are exactly three connected one-dimensional algebraic groups (up to isomorphism): the multiplicative group $\mathbb{G}_m$, the additive group $\mathbb{G}_a$, and the elliptic curves (one for each j -invariant). All three arise from a single family — the Weierstrass cubics — as the discriminant varies:

Condition	Singularity	Group
$\Delta \neq 0$	smooth	elliptic curve E
$\Delta = 0, c_4 \neq 0$	node	$\mathbb{G}_m$ (or $\ker \text{Nm}$ if non-split)
$\Delta = 0, c_4 = 0$	cuspidal	$\mathbb{G}_a$

For a reader who has seen Chevalley’s structure theorem [7]: every connected algebraic group over a perfect field decomposes as $1 \rightarrow H \rightarrow G \rightarrow A \rightarrow 1$ with H affine and A an abelian variety. In dimension 1, the abelian varieties are the elliptic curves, and the affine groups are $\mathbb{G}_m$ and $\mathbb{G}_a$. The degeneration of an elliptic curve to a nodal or cuspidal cubic is the degeneration from the “abelian variety” side of Chevalley’s theorem to the “affine” side.

The Split/Non-Split Distinction

The split vs. non-split distinction for nodes deserves emphasis, as it will play a central role in reduction theory. Over a local field K with residue field k , if an elliptic curve E/K has bad reduction at a prime $\mathfrak{p}$, the reduced curve $\tilde{E}/k$ is a singular cubic. The type of singularity determines the *reduction type*:

- *Multiplicative reduction*: $\tilde{E}$ has a node. The reduction is *split* if the tangent lines at the node are defined over the residue field k , and *non-split* otherwise.
- *Additive reduction*: $\tilde{E}$ has a cusp.

The terminology is now self-explanatory: “multiplicative reduction” means the nonsingular locus of the reduced curve is $\mathbb{G}_m$, and “additive reduction” means it is $\mathbb{G}_a$. We will develop this in full detail in chapter 6.

Example 1.6.6: Reduction of $y^2 = x^3 - x$ modulo Small Primes

Consider $E : y^2 = x^3 - x$ over $\mathbb{Q}$, with $\Delta = 64 = 2^6$.

Modulo $p = 3$: The reduced curve is $y^2 = x^3 - x \equiv x^3 + 2x \pmod{3}$. The discriminant of the cubic $x^3 + 2x$ modulo 3 is $-4(2)^3 = -32 \equiv 1 \pmod{3}$, which is nonzero. So $\tilde{E}$ is smooth over $\mathbb{F}_3$: the curve has *good reduction* at 3.

Modulo $p = 2$: The discriminant $\Delta = 64 \equiv 0 \pmod{2}$, so $\tilde{E}$ is singular over $\mathbb{F}_2$. In characteristic 2 we need the generalized Weierstrass form, and a more careful analysis (via Tate’s algorithm, chapter 6) is required. For now, note that $2|\Delta$ indicates bad reduction.

This curve has good reduction at all odd primes (since $\Delta = 2^6$), and the only prime of bad reduction is $p = 2$.

Example 1.6.7: A Nodal Cubic over $\mathbb{F}_5$

Consider the curve $y^2 = x^3 + x^2$ over $\mathbb{F}_5$. This is in the form (1.31) with $\alpha = 1$. Since $1 \in (\mathbb{F}_5^*)^2$ (as $1 = 1^2$), the node is split, and $E_{\text{ns}}(\mathbb{F}_5) \cong \mathbb{F}_5^*$. We verify: $|\mathbb{F}_5^*| = 4$, so $|E_{\text{ns}}(\mathbb{F}_5)| = 4$ (plus the singular point $(0, 0)$, which is excluded).

The nonsingular affine points are found by computing $x^3 + x^2 = x^2(x + 1)$ for $x = 0, 1, 2, 3, 4$ modulo 5: $x = 0 : 0$ (singular point, excluded); $x = 1 : 2$ (nonsquare in $\mathbb{F}_5^*$); $x = 2 : 12 \equiv 2$ (nonsquare); $x = 3 : 36 \equiv 1$ (square, $y = \pm 1$); $x = 4 : 80 \equiv 0$ ($y = 0$, one point). This gives $2 + 1 = 3$ affine points, plus O , for a total of 4.

Exercise 1.6.1

1. Verify directly that the parametrization $\varphi : (\mathbb{G}_a(K), +) \rightarrow (E_{\text{ns}}(K), +)$ given by $t \mapsto (t^2, t^3)$ is a group homomorphism for the cuspidal cubic $y^2 = x^3$, by computing the sum $(t_1^2, t_1^3) + (t_2^2, t_2^3)$ using the chord-and-tangent formulas and showing it equals $((t_1 + t_2)^2, (t_1 + t_2)^3)$.
2. For the nodal cubic $y^2 = x^2(x + 1)$ over $\mathbb{Q}$, the node is split ($\alpha = 1$ is a square). Write down the explicit parametrization $\mathbb{G}_m(\mathbb{Q}) \xrightarrow{\sim} E_{\text{ns}}(\mathbb{Q})$ using the parameter $u = (t + 1)/(t - 1)$ where $t = y/x$, and verify the multiplicative group law for two specific points.
3. Consider the curve $y^2 = x^2(x - 1)$ over $\mathbb{Q}$ (note the sign: $\alpha = -1$). Show that the node is non-split over $\mathbb{Q}$ (the tangent directions are $y = \pm ix$), and identify $E_{\text{ns}}(\mathbb{Q})$ with a subgroup of $\mathbb{Q}(i)^*$.
4. Show that the cuspidal cubic $y^2 = x^3$ over $\mathbb{F}_p$ (with $p \geq 5$) has $|E_{\text{ns}}(\mathbb{F}_p)| = p$. Conclude that every element of $E_{\text{ns}}(\mathbb{F}_p)$ has order dividing p , and contrast this with the Hasse bound $|E(\mathbb{F}_p)| \in [p + 1 - 2\sqrt{p}, p + 1 + 2\sqrt{p}]$ for smooth curves.

5. Let E be a Weierstrass cubic over K in the general form $y^2 + a_1xy + a_3y = x^3 + a_2x^2 + a_4x + a_6$ with $\Delta = 0$ and a singular point at (x_0, y_0) . After translating the singular point to the origin, show that the tangent cone is

$$(2y_0 + a_1x_0 + a_3)Y = (3x_0^2 + 2a_2x_0 + a_4 - a_1y_0)X$$

when the singularity is a cusp, and determine the tangent cone in the nodal case. (*Hint:* expand $f(x_0 + X, y_0 + Y) = 0$ and collect the lowest-degree terms.)

6. Using the results of this section and the classification of one-dimensional connected algebraic groups over $\overline{K}$ (cf. [7]), explain why an elliptic curve cannot be isomorphic (as an algebraic group) to $\mathbb{G}_m$ or $\mathbb{G}_a$, even though the degenerate Weierstrass cubics produce these groups. (*Hint:* one approach uses properness.)

I do not talk about the invariant differential. I think I should in the first chapter.

In general, look at the summary table at the end, and make sure you have that content in the first chapter!

2

Isogenies and the Endomorphism Ring

Chapter 1 established the group law on an elliptic curve. In this chapter, we study the natural maps between these groups: *isogenies*. We will develop this algebraic theory from scratch, valid over any field. We define isogenies as non-constant morphisms of elliptic curves, prove the key structural results (the dual isogeny, the degree formula $\deg([n]) = n^2$, the torsion structure $E[n] \cong (\mathbb{Z}/n\mathbb{Z})^2$), classify endomorphism rings, and introduce the Tate module $T_\ell(E)$.

Over the complex numbers, these concepts possess an analytic description in terms of lattices, which we will explore in detail in chapter 3. However, over a number field, a finite field, or a p -adic field, we have no complex lattice to work with, which is why the algebraic theory must stand on its own. Along the way, we encounter new phenomena that have no analytic counterpart — the Frobenius endomorphism and the distinction between ordinary and supersingular curves — which arise from the arithmetic of characteristic p .

2.1 Isogenies — Definition and First Properties

We begin by defining isogenies and establishing their fundamental properties. The key surprise for the reader coming from general algebraic geometry is the *rigidity*: a non-constant morphism between elliptic curves that sends the identity to the identity is automatically a group homomorphism. This is a consequence of the properness of elliptic curves and distinguishes them sharply from affine algebraic groups. Building on this, we develop the Galois theory of elliptic function fields, which gives us precise control over the kernel, the fibers, and the ramification of an isogeny, and which enables us to construct quotient curves E/Φ for any finite subgroup Φ .

Definition and Basic Properties

Definition 2.1.1: Isogeny

Let E_1/K and E_2/K be elliptic curves. An *isogeny* from E_1 to E_2 is a morphism $\varphi : E_1 \rightarrow E_2$ of varieties that satisfies $\varphi(O_{E_1}) = O_{E_2}$. If φ is non-constant (equivalently, surjective — since a morphism of projective curves is either surjective or constant), we say φ is a *non-trivial isogeny*. Two elliptic curves are *isogenous* if there exists a non-trivial isogeny between them.

Unless specified otherwise, all isogenies will be considered non-trivial unless we write $\text{Hom}(E_1, E_2)$ which will contain the trivial isogeny.

2.1.2 Isogenies on Abelian Varieties When working over curves, all non-trivial morphisms are finite maps. When working over general abelian varieties, it shall be necessary to specify that the *kernel* is finite.

The condition $\varphi(O) = O$ is essential: without it, φ would be an “affine morphism” (in the linear algebra sense, like an affine linear map), not a morphism of “pointed curves.” Since the group structure on an elliptic curve is determined by the marked point O (section 1.3), preserving O is the natural condition for a morphism in the category of elliptic curves.

The following rigidity result is the foundational fact of the theory:

Theorem 2.1.3: Isogenies are Group Homomorphisms

Let $\varphi : E_1 \rightarrow E_2$ be a non-constant isogeny. Then φ is a group homomorphism: $\varphi(P + Q) = \varphi(P) + \varphi(Q)$ for all $P, Q \in E_1(\bar{K})$.

Proof:

We give two arguments, one using the rigidity lemma and one relying on the Picard group.

Approach 1 (Rigidity lemma). The rigidity lemma for abelian varieties [3] states: if A is an abelian variety, V is a variety, and $f : A \times V \rightarrow B$ is a morphism to an abelian variety such that $f(0, v_0) = 0$ for some $v_0 \in V$, then $f(a, v) = f(a, v_0) + f(0, v)$ for all (a, v) . Taking $A = V = E_1$, $B = E_2$, and $f(P, Q) = \varphi(P + Q) - \varphi(P) - \varphi(Q)$: we have $f(O, O) = \varphi(O) - \varphi(O) - \varphi(O) = O$, so f satisfies the hypothesis. By rigidity, $f(P, Q) = f(P, O) + f(O, Q)$ for all P, Q . But $f(P, O) = \varphi(P) - \varphi(P) - \varphi(O) = O$ and $f(O, Q) = \varphi(Q) - \varphi(O) - \varphi(Q) = O$. So $f \equiv O$, meaning $\varphi(P + Q) = \varphi(P) + \varphi(Q)$.

Approach 2 (Via the Picard Group). We actually already proved this during our construction of the group law (corollary 1.3.3), but it is worth recalling the mechanism. The group law on an elliptic curve is defined such that the map $\sigma : P \mapsto [(P) - (O)]$ is an isomorphism to $\text{Pic}^0(E)$.

Since an isogeny φ is a finite morphism of curves, it induces a group homomorphism on divisor classes $\varphi_* : \text{Pic}^0(E_1) \rightarrow \text{Pic}^0(E_2)$ via pushforward (??). Evaluating this homomorphism on the

divisor class associated to a point P gives:

$$\varphi_*([(P) - (O_{E_1})]) = [(\varphi(P)) - (\varphi(O_{E_1}))] = [(\varphi(P)) - (O_{E_2})],$$

where the crucial step is that an isogeny satisfies $\varphi(O_{E_1}) = O_{E_2}$. This exactly matches the divisor class of $\varphi(P)$ in $\text{Pic}^0(E_2)$. Since φ_* is a homomorphism *on the points of E_1* , the underlying map φ must also be a homomorphism.

This theorem has no analogue for affine algebraic groups: a polynomial map $\mathbb{G}_m \rightarrow \mathbb{G}_m$ sending 1 to 1 need not be a group homomorphism (consider $t \mapsto t^2 + t - 1$, which sends 1 to 1 but is not multiplicative). The properness of elliptic curves is the essential ingredient.

An immediate consequence of the group homomorphism property is the following:

Corollary 2.1.4: Isogenies Preserve Torsion

If $\varphi : E_1 \rightarrow E_2$ is an isogeny and $P \in E_1[n]$ (i.e., $[n]P = O$), then $\varphi(P) \in E_2[n]$. In particular, φ restricts to a group homomorphism $\varphi : E_1[n] \rightarrow E_2[n]$ for every $n \geq 1$.

Proof:

$[n]\varphi(P) = \varphi([n]P) = \varphi(O) = O$, so $\varphi(P) \in E_2[n]$.

Another consequence that we will use repeatedly:

Lemma 2.1.5: Translation Absorbs the Kernel

Let $\varphi : E_1 \rightarrow E_2$ be a non-constant isogeny and $T \in \ker(\varphi)$. Then $\varphi \circ \tau_T = \varphi$, where $\tau_T : E_1 \rightarrow E_1$ is the translation $P \mapsto P + T$.

Proof:

For any $P \in E_1$: $\varphi(\tau_T(P)) = \varphi(P + T) = \varphi(P) + \varphi(T) = \varphi(P) + O = \varphi(P)$.

Just as with any non-constant morphism of smooth projective curves, an isogeny $\varphi : E_1 \rightarrow E_2$ induces an injective K -algebra homomorphism on the fields of rational functions via pullback:

$$\varphi^* : K(E_2) \hookrightarrow K(E_1), \quad f \mapsto f \circ \varphi.$$

Because elliptic curves are curves, this makes $K(E_1)$ a finite extension of the subfield $\varphi^*K(E_2)$. This field extension encodes the geometric and arithmetic properties of the isogeny.

Definition 2.1.6: Degree, Separable and Inseparable Degrees

Let $\varphi : E_1 \rightarrow E_2$ be a non-constant isogeny. The *degree* of φ is

$$\deg(\varphi) = [K(E_1) : \varphi^*K(E_2)],$$

the degree of the induced field extension. We define $\deg([0]) = 0$.

The extension $K(E_1)/\varphi^*K(E_2)$ is a finite extension of function fields and decomposes as a separable extension followed by a purely inseparable extension. The *separable degree* $\deg_s(\varphi)$ and the *inseparable degree* $\deg_i(\varphi)$ are the degrees of these two parts, satisfying

$$\deg(\varphi) = \deg_s(\varphi) \cdot \deg_i(\varphi).$$

We say φ is *separable* if $\deg_i(\varphi) = 1$, *inseparable* if $\deg_i(\varphi) > 1$, and *purely inseparable* if $\deg_s(\varphi) = 1$.

In characteristic 0, every finite extension of fields is separable, so every isogeny is separable. In characteristic p , inseparable isogenies arise naturally from the Frobenius endomorphism, as we shall see in section 2.4.

Proposition 2.1.7: Multiplicativity of Degree

For isogenies $\varphi : E_1 \rightarrow E_2$ and $\psi : E_2 \rightarrow E_3$,

$$\deg(\psi \circ \varphi) = \deg(\psi) \cdot \deg(\varphi).$$

The same holds for the separable and inseparable degrees separately.

Proof:

This is the tower law for field extensions: $[K(E_1) : (\psi \circ \varphi)^*K(E_3)] = [K(E_1) : \varphi^*K(E_2)] \cdot [\varphi^*K(E_2) : (\psi \circ \varphi)^*K(E_3)]$. Since $(\psi \circ \varphi)^* = \varphi^* \circ \psi^*$ and φ^* is an isomorphism on its image, the second factor is $[K(E_2) : \psi^*K(E_3)] = \deg(\psi)$. The separable and inseparable degrees multiply by the standard tower law for separable/inseparable extensions.

The Galois Theory of Isogenies

The group homomorphism property of isogenies, combined with lemma 2.1.5, gives remarkable control over the fiber structure and the function field extension of an isogeny. This is the key structural result that underlies much of the theory.

Theorem 2.1.8: Galois Theory of Elliptic Function Fields

Let $\varphi : E_1 \rightarrow E_2$ be a non-constant isogeny.

1. (*Uniform fibers and ramification.*) For every $Q \in E_2(\overline{K})$,

$$|\varphi^{-1}(Q)| = \deg_s(\varphi), \quad \text{and} \quad e_\varphi(P) = \deg_i(\varphi) \text{ for every } P \in \varphi^{-1}(Q),$$

where $e_\varphi(P)$ is the ramification index of φ at P .

2. (*Kernel as Galois group.*) The map

$$\ker(\varphi) \longrightarrow \text{Gal}(\overline{K}(E_1) / \varphi^* \overline{K}(E_2)), \quad T \longmapsto \tau_T^*,$$

is an isomorphism of groups, where τ_T^* is the automorphism of $\overline{K}(E_1)$ induced by the translation $\tau_T : P \mapsto P + T$.

3. (*Separable $\Rightarrow$ unramified and Galois.*) If φ is separable, then φ is unramified, $|\ker(\varphi)| = \deg(\varphi)$, and $\overline{K}(E_1)/\varphi^* \overline{K}(E_2)$ is a Galois extension.

Proof:

1. From the general theory of morphisms of curves [3], we know that $|\varphi^{-1}(Q)| = \deg_s(\varphi)$ holds for all but finitely many $Q \in E_2$. We show it holds for *all* Q .

Let $Q, Q' \in E_2(\overline{K})$ be arbitrary, and choose $R \in E_1(\overline{K})$ with $\varphi(R) = Q' - Q$ (possible since φ is surjective). The translation $\tau_R : P \mapsto P + R$ is an isomorphism $E_1 \rightarrow E_1$ that restricts to an injection $\varphi^{-1}(Q) \hookrightarrow \varphi^{-1}(Q')$: if $\varphi(P) = Q$, then $\varphi(P + R) = \varphi(P) + \varphi(R) = Q + (Q' - Q) = Q'$. Since τ_R is a bijection, $|\varphi^{-1}(Q)| = |\varphi^{-1}(Q')|$. As Q, Q' were arbitrary, the fiber size is constant, so the “generic” value $\deg_s(\varphi)$ holds everywhere.

For the ramification: let $P, P' \in \varphi^{-1}(Q)$ and set $T = P' - P$. Then $\varphi(T) = \varphi(P') - \varphi(P) = O$, so $T \in \ker(\varphi)$, and $\varphi \circ \tau_T = \varphi$ (lemma 2.1.5). Since τ_T is an isomorphism:

$$e_\varphi(P') = e_{\varphi \circ \tau_T}(P) = e_\varphi(\tau_T(P)) \cdot e_{\tau_T}(P) = e_\varphi(P') \cdot 1,$$

and symmetrically $e_\varphi(P) = e_\varphi(P')$. Thus all points in $\varphi^{-1}(Q)$ have the same ramification index. Now:

$$\deg(\varphi) = \sum_{P \in \varphi^{-1}(Q)} e_\varphi(P) = |\varphi^{-1}(Q)| \cdot e_\varphi(P) = \deg_s(\varphi) \cdot e_\varphi(P),$$

so $e_\varphi(P) = \deg(\varphi) / \deg_s(\varphi) = \deg_i(\varphi)$.

2. We show the map $T \mapsto \tau_T^*$ is a well-defined injective group homomorphism into the Galois group, and then count to conclude it is an isomorphism.

Well-defined: For $T \in \ker(\varphi)$, the translation τ_T^* sends $f \in \overline{K}(E_1)$ to f^T defined by $f^T(P) = f(P + T)$. We must show τ_T^* fixes $\varphi^* \overline{K}(E_2)$. For $g \in \overline{K}(E_2)$:

$$\tau_T^*(\varphi^* g) = (\varphi \circ \tau_T)^* g = \varphi^* g,$$

using $\varphi \circ \tau_T = \varphi$ (lemma 2.1.5). So τ_T^* fixes $\varphi^* \overline{K}(E_2)$.

Homomorphism: $\tau_{S+T}^* = (\tau_S \circ \tau_T)^* = \tau_T^* \circ \tau_S^*$. (The order reversal is the standard contravariance of pullbacks.)

Injectivity: If τ_T^* is the identity on $\overline{K}(E_1)$, then every rational function f satisfies $f(P+T) = f(P)$ for all P . In particular, taking $f = x$ (the Weierstrass coordinate, which separates points on E_1 up to the involution): if $x(P+T) = x(P)$ and $y(P+T) = y(P)$ for all P , then $P+T = P$ for all P , hence $T = O$.

Surjectivity (by counting): The group $\ker(\varphi)$ has order $\deg_s(\varphi)$ (by part (1) applied to $Q = O$). The Galois group $\text{Gal}(\overline{K}(E_1)/\varphi^* \overline{K}(E_2))$ has order $\leq [\overline{K}(E_1) : \varphi^* \overline{K}(E_2)]_s = \deg_s(\varphi)$ (with equality iff the extension is Galois). Since we have an injection from a group of size $\deg_s(\varphi)$ into a group of size $\leq \deg_s(\varphi)$, the map is an isomorphism and the upper bound is achieved.

3. If φ is separable, then $\deg_i(\varphi) = 1$, so $e_\varphi(P) = 1$ for all P by part (1): φ is unramified. The fiber $\varphi^{-1}(O) = \ker(\varphi)$ has $|\ker(\varphi)| = \deg_s(\varphi) = \deg(\varphi)$. And from part (2), $|\text{Gal}(\overline{K}(E_1)/\varphi^* \overline{K}(E_2))| = \deg_s(\varphi) = \deg(\varphi) = [\overline{K}(E_1) : \varphi^* \overline{K}(E_2)]$, which is the criterion for the extension to be Galois.

This theorem is the algebraic engine of the theory. Part (1) says that isogenies between elliptic curves are “as uniform as possible”: the fiber size and ramification are constant across all points (contrast this with a generic finite morphism of curves, where both can vary). Part (2) says the kernel of an isogeny *is* the Galois group of the corresponding function field extension, with the translation automorphisms τ_T^* playing the role of the Galois elements. Part (3) says that separable isogenies are always étale covers.

An important consequence is the “factoring” criterion:

Corollary 2.1.9: Extending Isogenies

Let $\varphi : E_1 \rightarrow E_2$ and $\psi : E_1 \rightarrow E_3$ be non-constant isogenies, with φ separable. If $\ker(\varphi) \subseteq \ker(\psi)$, then there exists a unique isogeny $\lambda : E_2 \rightarrow E_3$ such that $\psi = \lambda \circ \varphi$.

$$\begin{array}{ccc}
 & E_2 & \\
 \varphi \nearrow & & \downarrow \lambda \\
 E_1 & & E_3 \\
 \psi \searrow & & \\
 & &
 \end{array}$$

Proof:

Since φ is separable, theorem 2.1.8(3) says $\overline{K}(E_1)/\varphi^* \overline{K}(E_2)$ is Galois with Galois group $\ker(\varphi)$ (acting by $T \mapsto \tau_T^*$). The inclusion $\ker(\varphi) \subseteq \ker(\psi)$ means every $T \in \ker(\varphi)$ also satisfies $\psi(P+T) = \psi(P)$, i.e., τ_T^* fixes $\psi^* \overline{K}(E_3)$. By Galois theory (a subfield fixed by the entire Galois

group lies in the base field):

$$\psi^* \overline{K}(E_3) \subseteq \varphi^* \overline{K}(E_2) \subseteq \overline{K}(E_1).$$

The inclusion $\psi^* \overline{K}(E_3) \subseteq \varphi^* \overline{K}(E_2)$ gives a morphism $\lambda : E_2 \rightarrow E_3$ at the level of function fields: $\varphi^*(\lambda^* \overline{K}(E_3)) = \psi^* \overline{K}(E_3)$, which geometrically means $\lambda \circ \varphi = \psi$. Finally, $\lambda(O_{E_2}) = \lambda(\varphi(O_{E_1})) = \psi(O_{E_1}) = O_{E_3}$, so λ is an isogeny. Uniqueness follows from the injectivity of φ^* .

Isogenies from Finite Subgroups

The extension corollary above gives us “downward” factoring. We now address the fundamental “upward” question: given a finite subgroup $\Phi \subseteq E$, does there exist an isogeny $\varphi : E \rightarrow E'$ with $\ker(\varphi) = \Phi$? Over $\mathbb{C}$, the answer is immediate: $E' = \mathbb{C}/\Lambda'$ where $\Lambda' \supset \Lambda$ with $\Lambda'/\Lambda = \Phi$. Algebraically, we use the Galois theory just developed to construct E' from the fixed field.

Theorem 2.1.10: Isogenies from Finite Subgroups

Let E/K be an elliptic curve and let $\Phi \subseteq E(\overline{K})$ be a finite subgroup. Then there exists a unique (up to isomorphism) smooth projective curve $C/\overline{K}$ and a separable isogeny $\varphi : E \rightarrow C$ with $\ker(\varphi) = \Phi$. Moreover, C is an elliptic curve (of genus 1).

If Φ is $\text{Gal}(\overline{K}/K)$ -stable (i.e., defined over K), then C and φ can be taken to be defined over K . We write $C = E/\Phi$.

Proof:

Each $T \in \Phi$ gives a translation automorphism $\tau_T : E \rightarrow E$, which induces an automorphism $\tau_T^* : \overline{K}(E) \rightarrow \overline{K}(E)$. Since Φ is a finite group of automorphisms of the function field $\overline{K}(E)$, Artin’s theorem on fixed fields ensures that

$$L := \overline{K}(E)^\Phi = \{f \in \overline{K}(E) \mid \tau_T^* f = f \text{ for all } T \in \Phi\}$$

is a subfield of $\overline{K}(E)$ with $[\overline{K}(E) : L] = |\Phi|$ and $\overline{K}(E)/L$ is a Galois extension with Galois group isomorphic to Φ .

The field L has transcendence degree 1 over $\overline{K}$ (since $\overline{K}(E)$ does and L is an intermediate field of a finite extension). By the correspondence between smooth projective curves and their function fields [3], there exists a unique smooth projective curve $C/\overline{K}$ with $\overline{K}(C) = L$, and the inclusion $L \hookrightarrow \overline{K}(E)$ corresponds to a finite morphism

$$\varphi : E \longrightarrow C, \quad \varphi^*(\overline{K}(C)) = L.$$

Claim: φ is unramified. Let $P \in E$ and $T \in \Phi$. For every $f \in \overline{K}(C) = L$:

$$f(\varphi(P + T)) = (\varphi^* f)(P + T) = (\tau_T^* \circ \varphi^*)(f)(P) = (\varphi^* f)(P) = f(\varphi(P)),$$

where the key step uses $\tau_T^*(L) = L$ (since L is the fixed field of Φ and $T \in \Phi$). Since this holds for every rational function f on C , we conclude $\varphi(P + T) = \varphi(P)$. In particular, choosing

$Q = \varphi(P) \in C$ and varying $T \in \Phi$:

$$\{P + T \mid T \in \Phi\} \subseteq \varphi^{-1}(Q).$$

Since the points $P + T$ are pairwise distinct (as T ranges over Φ), we have $|\varphi^{-1}(Q)| \geq |\Phi|$. But $|\varphi^{-1}(Q)| \leq \deg(\varphi) = |\Phi|$ by the general theory of morphisms of curves, with equality if and only if φ is unramified at all points of $\varphi^{-1}(Q)$. Therefore φ is unramified at every point above Q , and as Q was arbitrary, φ is unramified everywhere.

Claim: φ is separable and $\ker(\varphi) = \Phi$. Since φ is unramified, $e_\varphi(P) = 1$ for all P , hence $\deg_i(\varphi) = 1$ and φ is separable. The fiber $\varphi^{-1}(\varphi(O)) = \varphi^{-1}(O_C)$ (where $O_C = \varphi(O_E)$) contains $\{O + T : T \in \Phi\} = \Phi$, and has exactly $\deg(\varphi) = |\Phi|$ elements, so $\ker(\varphi) = \Phi$.

Claim: C has genus 1. Since φ is unramified, the Hurwitz formula [3] gives

$$2g(E) - 2 = \deg(\varphi) (2g(C) - 2) + \underbrace{\sum_{P \in E} (e_\varphi(P) - 1)}_{=0}.$$

With $g(E) = 1$ and $\deg(\varphi) = |\Phi| \neq 0$: $0 = |\Phi|(2g(C) - 2)$, so $g(C) = 1$.

Taking $O_C = \varphi(O_E)$ as the marked point makes C an elliptic curve, φ an isogeny, and $\varphi(O_E) = O_C$ by construction. Uniqueness: if $\psi : E \rightarrow C'$ is another separable isogeny with $\ker(\psi) = \Phi$, then $\ker(\varphi) = \ker(\psi) = \Phi$ and corollary 2.1.9 (applied in both directions) gives isogenies $\lambda : C \rightarrow C'$ and $\mu : C' \rightarrow C$ with $\psi = \lambda \circ \varphi$ and $\varphi = \mu \circ \psi$. Then $\varphi = \mu \circ \lambda \circ \varphi$, so $\mu \circ \lambda = [\deg(\mu \circ \lambda) / \deg(\varphi)] \cdots$ — more simply, $\deg(\lambda) \deg(\mu) = \deg(\varphi) / \deg(\varphi) = 1$... in fact, λ and μ have $\deg(\lambda) \deg(\mu) = 1$ since $\mu \circ \lambda \circ \varphi = \varphi$ gives $\deg(\mu) \deg(\lambda) = 1$, so both are isomorphisms.

For the statement over K : if Φ is $\text{Gal}(\overline{K}/K)$ -stable, then $L = \overline{K}(E)^\Phi$ is stable under $\text{Gal}(\overline{K}/K)$ (since the Galois action commutes with the translation automorphisms for a K -rational subgroup). By Galois descent, C and φ descend to K .

Vélu's Formulas

theorem 2.1.10 constructs E/Φ abstractly via function fields. In practice, we need explicit Weierstrass equations. Vélu [25] gave formulas for the isogeny and its image curve in terms of the Weierstrass data and the subgroup Φ . The derivation amounts to explicitly computing Φ -invariant functions with prescribed pole orders at O ; the details are a (long but elementary) exercise in the addition formulas.

Proposition 2.1.11: Vélu's Formulas

Let $E : y^2 = x^3 + Ax + B$ (with $\text{char} K \neq 2, 3$) and let $\Phi \subseteq E(\overline{K})$ be a finite subgroup. For each $Q = (x_Q, y_Q) \in \Phi \setminus \{O\}$, define

$$t_Q = 3x_Q^2 + A, \quad u_Q = 2y_Q^2.$$

Let $\Phi^+ \subseteq \Phi \setminus \{O\}$ be a set of representatives for the equivalence relation $Q \sim -Q$ (choosing one from each pair $\{Q, -Q\}$), and set

$$t = \sum_{Q \in \Phi^+} t_Q, \quad w = \sum_{Q \in \Phi^+} (u_Q + x_Q t_Q).$$

Then the quotient curve $E' = E/\Phi$ has Weierstrass equation

$$E' : Y^2 = X^3 + (A - 5t)X + (B - 7w), \quad (2.1)$$

and the isogeny $\varphi : E \rightarrow E'$ is given on affine points $P = (x_P, y_P) \notin \Phi$ by

$$\varphi(P) = \left(x_P + \sum_{Q \in \Phi^+} \left(\frac{t_Q}{x_P - x_Q} + \frac{u_Q}{(x_P - x_Q)^2} \right), y_P \left(1 - \sum_{Q \in \Phi^+} \frac{u_Q}{(x_P - x_Q)^3} \right) + \dots \right) \quad (2.2)$$

where the y -coordinate formula involves additional terms from the Taylor expansion of the addition law around each Q .^a

^aThe complete y -coordinate formula and the generalization to arbitrary Weierstrass equations are given in Silverman [21, Chapter III, Section 4].

The key idea behind Vélu's formulas is that the x -coordinate of the isogeny, $X(P) = x_P + \sum(\dots)$, is the unique Φ -invariant function on E that has a double pole at O and no other poles modulo Φ : it is the natural “ x -coordinate on E/Φ .” The terms $t_Q/(x_P - x_Q) + u_Q/(x_P - x_Q)^2$ arise from absorbing the poles at Q and $-Q$ into the sum by subtracting off the principal part of x_{P+Q} near $P = Q$.

Example 2.1.12: A Degree-2 Isogeny via Vélu

Let $E : y^2 = x^3 - x$ over $\mathbb{Q}$ and $\Phi = \{O, (0, 0)\}$. Then $\Phi^+ = \{(0, 0)\}$ (a single element, since $(0, 0) = -(0, 0)$ is its own inverse). We compute:

$$t_{(0,0)} = 3 \cdot 0^2 + (-1) = -1, \quad u_{(0,0)} = 2 \cdot 0^2 = 0.$$

So $t = -1$ and $w = 0 + 0 \cdot (-1) = 0$. The image curve is

$$E' : Y^2 = X^3 + (-1 - 5(-1))X + (0 - 7 \cdot 0) = X^3 + 4X.$$

The x -coordinate of the isogeny simplifies to $\varphi_x(x) = x - 1/x = (x^2 - 1)/x$ (using $t_Q/(x - 0) = -1/x$ and $u_Q = 0$). The y -coordinate works out to $\varphi_y(x, y) = y(x^2 + 1)/x^2$.

Verification: $\varphi(1, 0) = (0, 0)$ on E' ; $\varphi(-1, 0) = (0, 0)$ on E' ; $\varphi(0, 0) = O$ (the point at infinity, since φ_x has a pole). So $\ker(\varphi) = \{O, (0, 0)\}$ as required.

Example 2.1.13: An Isogeny Between Running Examples

The degree-2 isogeny $\varphi : E \rightarrow E'$ with $E : y^2 = x^3 - x$ and $E' : y^2 = x^3 + 4x$ has a dual isogeny $\hat{\varphi} : E' \rightarrow E$ (section 2.2). Both curves have $j = 1728$, confirming that isogenous curves can share the same j -invariant. More precisely, $j(E') = (-48 \cdot 4)^3 / (-16 \cdot 4 \cdot 64) = (-192)^3 / (-4096) = 1728$.

We conclude this section by showing that isogeny (certainly) need not preserve j -invariance:

Example 2.1.14: Isogeny and J-invariance

Take from example 2.1.12 the degree-2 isogeny from $E : y^2 = x^3 + 1$ (with $j = 0$). Using the kernel $\{O, (-1, 0)\}$ produces, via Vélu, the curve $E' : Y^2 = X^3 + (-1 + 5)X + (1 - 0) = X^3 + 4X + 1$, which has $j(E') \neq 0$ in general.

Exercise 2.1.1

1. (The rigidity lemma for elliptic curves.) Let $\varphi, \psi : E_1 \rightarrow E_2$ be two morphisms of elliptic curves with $\varphi(P_0) = \psi(P_0)$ for some $P_0 \in E_1(\bar{K})$. Show that $\varphi = \psi$ or both are constant. (*Hint:* Consider the morphism $\varphi - \psi : E_1 \rightarrow E_2$, which sends P_0 to O , and use the fact that a morphism of projective curves is either constant or surjective.)
2. Let $\varphi : E_1 \rightarrow E_2$ be a non-constant isogeny and $T \in \ker(\varphi)$. Show that for any $f \in \bar{K}(E_2)$, the pullback φ^*f satisfies $(\varphi^*f)(P + T) = (\varphi^*f)(P)$ for all $P \in E_1$. (*This is the function-field version of lemma 2.1.5.*)
3. For the isogeny $\varphi : E \rightarrow E'$ of example 2.1.12 (with $\varphi(x, y) = ((x^2 - 1)/x, y(x^2 + 1)/x^2)$), verify theorem 2.1.8(2) directly: show that $\bar{K}(E)/\varphi^*\bar{K}(E')$ has degree 2, identify the nontrivial Galois automorphism as $\tau_{(0,0)}$, and verify it acts as the identity on $\varphi^*\bar{K}(E')$.
4. Use Vélu's formulas to compute the degree-3 isogeny $\varphi : E \rightarrow E'$ on $E : y^2 = x^3 + 1$ with kernel $\Phi = \{O, (0, 1), (0, -1)\}$ (the 3-torsion subgroup generated by $(0, 1)$). Determine the Weierstrass equation of E' and verify that $\deg(\varphi) = 3$.
5. Let $\varphi : E_1 \rightarrow E_2$ and $\psi : E_2 \rightarrow E_3$ be separable isogenies. Show that $\ker(\psi \circ \varphi)$ fits in the exact sequence $0 \rightarrow \ker(\varphi) \rightarrow \ker(\psi \circ \varphi) \xrightarrow{\varphi} \ker(\psi) \rightarrow 0$.
6. Let $E : y^2 = x^3 - x$ and consider the 4-torsion subgroup $\Phi = E[2] \cong (\mathbb{Z}/2\mathbb{Z})^2$. Use corollary 2.1.9 to show that the isogeny $\varphi : E \rightarrow E/\Phi$ factors through any degree-2 isogeny $E \rightarrow E/\langle T \rangle$ for $T \in E[2] \setminus \{O\}$.

2.2 The Dual Isogeny

Over $\mathbb{C}$, an isogeny $\varphi : E_1 \rightarrow E_2$ corresponding to $\alpha \in \mathbb{C}$ with $\alpha\Lambda_1 \subseteq \Lambda_2$ has a dual $\hat{\varphi}$ corresponding to $\hat{\alpha} = \deg(\varphi)/\alpha$ (proposition 3.5.7). The algebraic construction of the dual isogeny cannot use this formula directly — there is no ambient field $\mathbb{C}$ to divide in. Instead, we use the Picard group: the dual isogeny is the map obtained by pulling back divisors under φ and translating via the isomorphism $E \cong \text{Pic}^0(E)$ from theorem 1.3.1. The resulting construction works over any field and gives us the fundamental identities $\hat{\varphi} \circ \varphi = [\deg \varphi]$ and $\varphi \circ \hat{\varphi} = [\deg \varphi]$ that are the backbone of the theory.

Construction via the Picard Group

Let $\varphi : E_1 \rightarrow E_2$ be a non-constant isogeny. The pullback of divisors under φ defines a map $\varphi^* : \text{Div}(E_2) \rightarrow \text{Div}(E_1)$ by $\varphi^*(\sum n_i(Q_i)) = \sum n_i \varphi^*(Q_i)$, where $\varphi^*(Q) = \sum_{P \in \varphi^{-1}(Q)} e_\varphi(P) (P)$ (summing over the preimage with ramification multiplicities). This map preserves degree (since $\sum_{P \in \varphi^{-1}(Q)} e_\varphi(P) = \deg(\varphi)$ for each Q) and preserves linear equivalence (since pullback of principal divisors is principal: $\varphi^* \text{div}(g) = \text{div}(g \circ \varphi)$ for $g \in \overline{K}(E_2)^*$). Therefore φ^* descends to a group homomorphism

$$\varphi^* : \text{Pic}^0(E_2) \longrightarrow \text{Pic}^0(E_1).$$

Composing with the Picard isomorphisms $\sigma_i : E_i \xrightarrow{\sim} \text{Pic}^0(E_i)$ from theorem 1.3.1:

Definition 2.2.1: The Dual Isogeny

Let $\varphi : E_1 \rightarrow E_2$ be a non-constant isogeny. The *dual isogeny* $\hat{\varphi} : E_2 \rightarrow E_1$ is the unique isogeny making the diagram

$$\begin{array}{ccc} E_2 & \xrightarrow{\sigma_2} & \text{Pic}^0(E_2) \\ \hat{\varphi} \downarrow & & \downarrow \varphi^* \\ E_1 & \xleftarrow{\sigma_1^{-1}} & \text{Pic}^0(E_1) \end{array} \quad (2.3)$$

commute: $\hat{\varphi} = \sigma_1^{-1} \circ \varphi^* \circ \sigma_2$. We also define $[\hat{0}] = [0]$.

Unpacking the definition: for $Q \in E_2(\overline{K})$,

$$\hat{\varphi}(Q) = \sigma_1^{-1}(\varphi^*[(Q) - (O_{E_2})]) = \sigma_1^{-1}[(\varphi^*(Q)) - \deg(\varphi) \cdot (O_{E_1})]. \quad (2.4)$$

Here $\varphi^*(Q) = \sum_{P \in \varphi^{-1}(Q)} e_\varphi(P) (P)$ is a divisor of degree $\deg(\varphi)$ on E_1 , and we subtract $\deg(\varphi) \cdot (O_{E_1})$ to get a degree-zero divisor. The map σ_1^{-1} then converts this divisor class to a point on E_1 .

Since σ_2 , φ^* , and σ_1^{-1} are all group homomorphisms, $\hat{\varphi}$ is automatically a group homomorphism. The remaining question is whether it is an isogeny (i.e., non-constant when φ is non-constant). This will follow from the fundamental identity below.

Note that the existence of a dual isogeny is not necessarily the case for general curves: by the Riemann-Hurwitz formula a degree $n > 1$ map from C to D forces the genus of C to be strictly greater than the genus of D (assuming $g(D) \geq 1$). Because you cannot map a lower-genus curve onto a higher-genus curve, a reverse map $D \rightarrow C$ is topologically and geometrically impossible.

Proposition 2.2.2: Isogenies Form An Equivalence Relationship

The relation “there is a (nontrivial) isogeny between elliptic curves” forms an equivalence relation on elliptic curves over $\overline{K}$.

Proof:
exercise.

2.2.3 The Arithmetic and Geometric Meaning of Isogeny The fact that isogeny forms an equivalence relation—symmetry being guaranteed by the existence of the dual isogeny—is a foundational pillar of modern arithmetic geometry. Grouping abelian varieties into “isogeny classes” implies they share identical core structures, differing only by finite, “negligible” kernels. Depending on the setting, this equivalence manifests in different ways:

- *Rational Isomorphism (Algebraic)*: If E_1 and E_2 are isogenous, then $E_2 \cong E_1/\Phi$ for some finite group Φ . If the category of abelian varieties is tensored with $\mathbb{Q}$ (creating the “isogeny category”), isogenies become exact isomorphisms, with the dual isogeny acting as the inverse map.
- *Point Counts (Finite Fields)*: Tate showed in [23] that two elliptic curves (or abelian varieties) over a finite field $\mathbb{F}_q$ are isogenous if and only if they have the exact same number of points over $\mathbb{F}_q$ and all its finite extensions; they share the same characteristic polynomial of Frobenius. This is an important criterion that we will prove in chapter 5.
- *L-functions and Rank (Number Fields)*: Faltings’ Isogeny Theorem establishes that two abelian varieties over a number field are isogenous if and only if they have the exact same L-function. Consequently, isogenous elliptic curves over $\mathbb{Q}$ share the exact same algebraic rank, even if their finite torsion subgroups differ.
- *Commensurable Lattices (Complex Geometry)*: Over $\mathbb{C}$, isogenous curves $E_1 = \mathbb{C}/\Lambda_1$ and $E_2 = \mathbb{C}/\Lambda_2$ possess commensurable lattices. They share a common sub-lattice of finite index, meaning their rational structures are identical: $\Lambda_1 \otimes \mathbb{Q} = \Lambda_2 \otimes \mathbb{Q}$.
- *Poincaré Reducibility*: Up to isogeny, every abelian variety decomposes uniquely into a product of “simple” abelian varieties ($A \sim A_1 \times \cdots \times A_k$). This geometric analog to prime factorization is only possible because we allow equivalence up to isogeny rather than strict isomorphism.

As isogenies are dual, their composition should yield some form of natural canceling out. The following makes this precise:

Proposition 2.2.4: Properties of the Dual Isogeny

Let $\varphi : E_1 \rightarrow E_2$ be a non-constant isogeny of degree n . Then:

1. $\hat{\varphi} \circ \varphi = [n]$ on E_1 .
2. $\varphi \circ \hat{\varphi} = [n]$ on E_2 .

Proof:

1. We need to show $\hat{\varphi}(\varphi(P)) = [n]P$ for all $P \in E_1(\overline{K})$. By the definition of $\hat{\varphi}$:

$$\hat{\varphi}(\varphi(P)) = \sigma_1^{-1}(\varphi^*[(\varphi(P)) - (O_{E_2})]).$$

The pullback $\varphi^*(\varphi(P))$ is the divisor $\sum_{R \in \varphi^{-1}(\varphi(P))} e_\varphi(R)(R)$. By theorem 2.1.8(1), the fiber $\varphi^{-1}(\varphi(P))$ consists of the points $P + T$ for $T \in \ker(\varphi)$ (each with multiplicity

$e_\varphi = \deg_i(\varphi)$, so

$$\varphi^*(\varphi(P)) = \deg_i(\varphi) \sum_{T \in \ker(\varphi)} (P + T).$$

Also $\varphi^*(O_{E_2}) = \deg_i(\varphi) \sum_{T \in \ker(\varphi)} (T)$. Thus:

$$\begin{aligned} \varphi^*[(\varphi(P)) - (O_{E_2})] &= \deg_i(\varphi) \left[\sum_{T \in \ker(\varphi)} (P + T) - \sum_{T \in \ker(\varphi)} (T) \right] \\ &= \deg_i(\varphi) \sum_{T \in \ker(\varphi)} [(P + T) - (T)]. \end{aligned}$$

In $\text{Pic}^0(E_1)$, using the group isomorphism $\sigma_1(P) = [(P) - (O_{E_1})]$, we can rewrite this sum of divisors:

$$\begin{aligned} \sum_{T \in \ker(\varphi)} [(P + T) - (T)] &= \sum_{T \in \ker(\varphi)} \left([(P + T) - (O_{E_1})] - [(T) - (O_{E_1})] \right) \\ &= \sum_{T \in \ker(\varphi)} (\sigma_1(P + T) - \sigma_1(T)) \\ &= \sum_{T \in \ker(\varphi)} \sigma_1(P) \quad (\text{since } \sigma_1 \text{ is a homomorphism}) \\ &= |\ker(\varphi)| \cdot \sigma_1(P). \end{aligned}$$

Therefore, $\varphi^*[(\varphi(P)) - (O_{E_2})] = \deg_i(\varphi) \cdot |\ker(\varphi)| \cdot \sigma_1(P) = n \cdot \sigma_1(P) = \sigma_1([n]P)$. Applying σ_1^{-1} to both sides gives $\hat{\varphi}(\varphi(P)) = [n]P$.

2. We need $\varphi(\hat{\varphi}(Q)) = [n]Q$ for all $Q \in E_2(\bar{K})$. Consider the isogeny $\varphi \circ \hat{\varphi} : E_2 \rightarrow E_2$. From (1), $\hat{\varphi} \circ \varphi = [n]$, so $\varphi \circ \hat{\varphi} \circ \varphi = \varphi \circ [n] = [n] \circ \varphi$ (since $[n]$ commutes with every endomorphism). Thus $(\varphi \circ \hat{\varphi}) \circ \varphi = [n] \circ \varphi$.

Since φ is surjective (non-constant), we can “cancel” it on the right. To justify the cancellation rigorously: let $\psi = \varphi \circ \hat{\varphi} - [n]$, which is a morphism $E_2 \rightarrow E_2$ with $\psi \circ \varphi = 0$. If ψ is non-constant, then ψ is surjective, making $\psi \circ \varphi$ surjective — contradicting $\psi \circ \varphi = 0$. Hence ψ is constant, and since $\psi(O) = \varphi(\hat{\varphi}(O)) - [n]O = O$, we get $\psi = 0$, meaning $\varphi \circ \hat{\varphi} = [n]$.

The identities $\hat{\varphi} \circ \varphi = [\deg \varphi]$ and $\varphi \circ \hat{\varphi} = [\deg \varphi]$ are the most important properties of the dual isogeny. They say that every isogeny has an “approximate inverse” — not an actual inverse (since $\deg(\varphi)$ may be > 1), but an isogeny that composes with φ to give multiplication by the degree. In the analytic theory, this was the statement $\hat{\alpha} \cdot \alpha = n$ and $\alpha \cdot \hat{\alpha} = n$.

A result that is worth singling out is that the dualizing operator $\widehat{(-)}$ is a *homomorphism*. This is much harder to show, technically requiring a special case of the theory of abelian varieties as $E \times E$ is no longer an elliptic curve). We thus single it out and box it separately:

Theorem 2.2.5: Dualizing is A Homomorphism

$$\widehat{\varphi + \psi} = \hat{\varphi} + \hat{\psi}$$

Proof:

The proof shall be presented for characteristic 0 (in [21, exercise 3.31] it goes over for arbitrary characteristic).

Let $x_1, y_1 \in K(E_1)$ and $x_2, y_2 \in K(E_2)$ be Weierstrass coordinates. Start by looking at E_2 considered as an elliptic curve defined over the field $K(E_1) = K(x_1, y_1)^a$. Then another way of saying that $\varphi : E_1 \rightarrow E_2$ is an isogeny is to note that $\varphi(x_1, y_1) \in E_2(K(x_1, y_1))$, and similarly for $\psi(x_1, y_1)$ and $(\varphi + \psi)(x_1, y_1)$. Now consider the divisor

$$D = ((\varphi + \psi)(x_1, y_1)) - (\varphi(x_1, y_1)) - (\psi(x_1, y_1)) + (O_{E_2}) \in \text{Div}_{K(x_1, y_1)}(E_2).$$

The definition of $\varphi + \psi$ implies that D sums to O_{E_2} , so by ?? D is linearly equivalent to 0. Thus there is a function

$$f \in K(x_1, y_1)(E_2) = K(x_1, y_1, x_2, y_2)$$

that, when considered as a function of x_2 and y_2 , has divisor D .

Now, consider f as a function of x_1 and y_1 , that is, treat f as a function on E_1 considered as an elliptic curve defined over $K(x_2, y_2)$. Suppose that $P_1 \in E_1(K(x_2, y_2))$ is a point satisfying $\varphi(P_1) = (x_2, y_2)$. Then examining D , specifically the term $-(\varphi(x_1, y_1))$, notice that f has a pole at P_1 , namely the function $f(x_1, y_1; x_2, y_2)$ has a pole when x_1, y_1, x_2, y_2 satisfies $(x_2, y_2) = \varphi(x_1, y_1)$. Furthermore, by [3]:

$$\nu_{P_1}(f) = -e_\varphi(P_1).$$

Similarly, f has a pole at P_1 if $(x_2, y_2) = \psi(P_1)$, and it has a zero at P_1 if $(x_2, y_2) = (\varphi + \psi)(P_1)$. It follows that as a function of x_1 and y_1 , the divisor of f has the form

$$(\varphi + \psi)^*((x_2, y_2)) - \varphi^*((x_2, y_2)) - \psi^*((x_2, y_2)) + \sum n_i(P_i) \in \text{Div}_{K(x_2, y_2)}(E_1),$$

where the P_i 's are points in $E_1(\bar{K})$ that do not depend on (x_2, y_2) , i.e., $\sum n_i(P_i) \in \text{Div}_{\bar{K}}(E_1)$.

Since this is the divisor of a function, its points sum to O_{E_1} in the group law. By the definition of the dual isogeny, summing the points in a pullback divisor $\alpha^*(Q)$ yields exactly $\hat{\alpha}(Q)$. Therefore, summing the points in the divisor above shows that the point

$$\widehat{(\varphi + \psi)}(x_2, y_2) - \hat{\varphi}(x_2, y_2) - \hat{\psi}(x_2, y_2)$$

is equal to the constant point $-\sum n_i P_i$, meaning it does not depend on (x_2, y_2) . Evaluating at $(x_2, y_2) = O_{E_2}$ (and recalling that isogenies always map $O \mapsto O$) shows that this constant point must be O_{E_1} , which completes the proof that

$$\widehat{\varphi + \psi} = \hat{\varphi} + \hat{\psi}.$$

^aThis is where we use the characteristic 0 assumption, since all of our results on elliptic curves have assumed that the base field is perfect, but the function field $K(x_1, y_1)$ is imperfect in characteristic $p > 0$.

Proposition 2.2.6: Further Properties of Dual

1. $\deg \varphi = \deg \hat{\varphi}$
2. $\hat{\hat{\varphi}} = \varphi$.
3. For isogenies $\varphi : E_1 \rightarrow E_2$ and $\psi : E_2 \rightarrow E_3$: $\widehat{\psi \circ \varphi} = \hat{\varphi} \circ \hat{\psi}$.
4. $\widehat{[n]} = [n]$ for all $n \in \mathbb{Z}$ (the multiplication-by- n map is self-dual).

Proof:

1. I realized I proved $\deg[m] = m^2$ in the next section!! Fix this!!

$$m^2 = \deg[m] = \deg(\varphi \circ \hat{\varphi}) = (\deg(\varphi))(\deg(\hat{\varphi})) = m(\deg(\hat{\varphi}))$$

2. By definition, the dual of $\hat{\varphi}$ satisfies $\hat{\hat{\varphi}} \circ \hat{\varphi} = [\deg(\hat{\varphi})] = [n]$. From (2), we also know that $\varphi \circ \hat{\varphi} = [n]$. Equating these gives $\hat{\hat{\varphi}} \circ \hat{\varphi} = \varphi \circ \hat{\varphi}$. Since $\hat{\varphi}$ is a non-constant isogeny, it is surjective, allowing us to cancel it on the right to conclude $\hat{\hat{\varphi}} = \varphi$.

3. For $\psi : E_2 \rightarrow E_3$ and $\varphi : E_1 \rightarrow E_2$: the pullback on Picard groups satisfies $(\psi \circ \varphi)^* = \varphi^* \circ \psi^*$. Therefore:

$$\begin{aligned} \widehat{\psi \circ \varphi} &= \sigma_1^{-1} \circ (\psi \circ \varphi)^* \circ \sigma_3 = \sigma_1^{-1} \circ \varphi^* \circ \psi^* \circ \sigma_3 \\ &= \sigma_1^{-1} \circ \varphi^* \circ \sigma_2 \circ \sigma_2^{-1} \circ \psi^* \circ \sigma_3 = \hat{\varphi} \circ \hat{\psi}. \end{aligned}$$

4. For $n \geq 1$: we compute $\widehat{[n]}$ using the definition. For $Q \in E(\overline{K})$:

$$\widehat{[n]}(Q) = \sigma^{-1}([n]^*([Q] - (O))).$$

Rather than computing the pullback explicitly, we use the identity from (1): $\widehat{[n]} \circ [n] = [\deg([n])] = [n^2] = [n] \circ [n]$. Cancelling $[n]$ on the right (which is surjective for $n \neq 0$) yields $\widehat{[n]} = [n]$.

Example 2.2.7: The Dual of a Degree-2 Isogeny

For the isogeny $\varphi : E \rightarrow E'$ of example 2.1.12 with $E : y^2 = x^3 - x$, $E' : y^2 = x^3 + 4x$, and $\varphi(x, y) = ((x^2 - 1)/x, y(x^2 + 1)/x^2)$, the dual $\hat{\varphi} : E' \rightarrow E$ satisfies $\hat{\varphi} \circ \varphi = [2]$ on E and $\varphi \circ \hat{\varphi} = [2]$ on E' .

The dual can be computed explicitly: $\hat{\varphi}$ is also a degree-2 isogeny from E' to E . Its kernel is precisely the image of the 2-torsion under φ : $\ker(\hat{\varphi}) = \varphi(E[2])$. Since $E[2]$ has order 4 and $\ker(\varphi)$ has order 2, the image $\varphi(E[2])$ is a subgroup of $E'[2]$ of order 2. Evaluating φ on the 2-torsion points of E reveals that $\ker(\hat{\varphi}) = \{O, (0, 0)\}$. One can use Vélú's formulas on E' with this kernel

to recover E as the target.

Example 2.2.8: Self-Duality of $[n]$

The map $[2] : E \rightarrow E$ on any elliptic curve satisfies $\widehat{[2]} = [2]$: the dual of “double” is “double.” This means $[2] \circ [2] = [4] = [\deg([2])]$, which is of course obvious ($[2]^2 = [4]$). But the self-duality is a non-trivial structural fact: it says that pulling back the divisor class $[(Q) - (O)]$ under $[2]$ gives the divisor class $\sigma([2]Q)$, not some other class.

The Dual and the Degree Form

The dual isogeny turns $\text{Hom}(E_1, E_2)$ into a module with a rich structure.

Proposition 2.2.9: The Degree as a Quadratic Form

For $\varphi, \psi \in \text{Hom}(E_1, E_2)$, $\deg$ satisfies the parallelogram identity^a:

$$\deg(\varphi + \psi) + \deg(\varphi - \psi) = 2\deg(\varphi) + 2\deg(\psi). \quad (2.5)$$

In particular, the map $\deg : \text{Hom}(E_1, E_2) \rightarrow \mathbb{Z}$ is a positive definite quadratic form (on $\text{Hom}(E_1, E_2) \setminus \{0\}$), and the associated bilinear form is

$$\langle \varphi, \psi \rangle = \deg(\varphi + \psi) - \deg(\varphi) - \deg(\psi) = \frac{1}{2}(\deg(\varphi + \psi) - \deg(\varphi) - \deg(\psi)) \cdot 2 \cdots \quad (2.6)$$

More precisely, the “dual pairing” $\langle \varphi, \psi \rangle = \widehat{\varphi + \psi} - \widehat{\varphi} - \widehat{\psi}$ relates the degree to the dual via $\widehat{\varphi} \circ \psi + \widehat{\psi} \circ \varphi = [\deg(\varphi + \psi) - \deg(\varphi) - \deg(\psi)]$ in $\text{End}(E_1)$.

^aRecall that a norm is an inner product if it satisfies the $\|x + y\| + \|x - y\| = 2\|x\| + 2\|y\|$

Proof:

The identity (2.5) follows from the theory of the dual: for any $\varphi, \psi : E_1 \rightarrow E_2$, the dual satisfies $\widehat{\varphi + \psi} = \widehat{\varphi} + \widehat{\psi}$ (since pullback on Picard groups is additive). Therefore:

$$\begin{aligned} (\widehat{\varphi} + \widehat{\psi})(\varphi + \psi) &= \widehat{\varphi} \circ \varphi + \widehat{\varphi} \circ \psi + \widehat{\psi} \circ \varphi + \widehat{\psi} \circ \psi \\ &= [\deg \varphi] + \widehat{\varphi} \circ \psi + \widehat{\psi} \circ \varphi + [\deg \psi]. \end{aligned}$$

Also $(\widehat{\varphi} + \widehat{\psi})(\varphi + \psi) = \widehat{\varphi + \psi} \circ (\varphi + \psi) = [\deg(\varphi + \psi)]$. (Here we used the fact that the dual of $\varphi + \psi$ is $\widehat{\varphi} + \widehat{\psi}$, which follows from the linearity of pullback on Pic^0 .) So $\widehat{\varphi} \circ \psi + \widehat{\psi} \circ \varphi = [\deg(\varphi + \psi)] - [\deg \varphi] - [\deg \psi]$.

Replacing ψ by $-\psi$ (and using $\widehat{-\psi} = -\widehat{\psi}$, $\deg(-\psi) = \deg(\psi)$): $-\widehat{\varphi} \circ \psi - \widehat{\psi} \circ \varphi = [\deg(\varphi - \psi)] - [\deg \varphi] - [\deg \psi]$. Adding the two identities: $0 = [\deg(\varphi + \psi)] + [\deg(\varphi - \psi)] - 2[\deg \varphi] - 2[\deg \psi]$, which gives (2.5).

The parallelogram identity (2.5) says $\deg$ is a quadratic form; if we were working over a field this would be an inner product. Its positive-definiteness ($\deg(\varphi) > 0$ for $\varphi \neq 0$) means $\text{Hom}(E_1, E_2)$ is a

torsion-free $\mathbb{Z}$ -module of finite rank as we'll soon show (the rank is at most 4, as we will prove in section 2.6). This quadratic form is the algebraic counterpart of the norm form $|\alpha|^2$ on the lattice of complex multipliers from section 3.5.

Corollary 2.2.10: Absence of Zero Divisors and Torsion

1. The abelian group $\text{Hom}(E_1, E_2)$ is torsion-free, in other words $\text{Hom}(E_1, E_2)$ is a torsion-free $\mathbb{Z}$ -module, and so free (as $\mathbb{Z}$ is a PID).
2. The endomorphism ring $\text{End}(E)$ has no zero divisors (it is a domain).

Proof:

1. Suppose $\varphi \in \text{Hom}(E_1, E_2)$ is a torsion element, meaning $[m]\varphi = 0$ for some integer $m \neq 0$. Taking the degree of both sides gives $\deg([m]\varphi) = \deg(0) = 0$. Because degree is multiplicative, we have $\deg([m]) \deg(\varphi) = m^2 \deg(\varphi) = 0$. Since $m \neq 0$, we must have $\deg(\varphi) = 0$, which implies $\varphi = 0$.
2. Suppose $\varphi, \psi \in \text{End}(E)$ satisfy $\varphi \circ \psi = 0$. Taking the degree gives $\deg(\varphi \circ \psi) = \deg(\varphi) \deg(\psi) = \deg(0) = 0$. Since the degrees are integers, either $\deg(\varphi) = 0$ or $\deg(\psi) = 0$, meaning $\varphi = 0$ or $\psi = 0$.

An important consequence of the dual pairing is that it allows us to define the “trace” of an endomorphism without ever writing down a matrix, and forces every endomorphism to satisfy a quadratic equation. When we introduce the Tate-module in section 2.5 this shall be an actual trace. For now, the trace can be thought of as measuring the angle of the inner product induced by the degree map:

Proposition 2.2.11: The Trace and the Characteristic Equation

Let $\varphi \in \text{End}(E)$ be an endomorphism. Define the *trace* of φ as the integer

$$\text{tr}(\varphi) = \deg(\varphi + [1]) - \deg(\varphi) - 1. \quad (2.7)$$

Then $\varphi + \hat{\varphi} = [\text{tr}(\varphi)]$, and φ satisfies the quadratic characteristic equation in $\text{End}(E)$:

$$\varphi^2 - [\text{tr}(\varphi)]\varphi + [\deg(\varphi)] = [0]. \quad (2.8)$$

Proof:

By the definition of the dual pairing $\langle \varphi, \psi \rangle = \hat{\varphi} \circ \psi + \hat{\psi} \circ \varphi$ applied to φ and $[1]$, we have:

$$\hat{\varphi} \circ [1] + \widehat{[1]} \circ \varphi = [\deg(\varphi + [1]) - \deg(\varphi) - \deg([1])].$$

Since $[1]$ is the identity map and $\widehat{[1]} = [1]$ (proposition 2.2.6), the left side simplifies to $\hat{\varphi} + \varphi$. The right side, by our definition of the trace, is exactly $[\text{tr}(\varphi)]$. Thus, $\varphi + \hat{\varphi} = [\text{tr}(\varphi)]$.

To find the characteristic equation, compose this identity with φ on the right:

$$\varphi^2 + \hat{\varphi} \circ \varphi = [\text{tr}(\varphi)] \circ \varphi.$$

We know from the fundamental property of the dual that $\hat{\varphi} \circ \varphi = [\deg(\varphi)]$. Substituting this in and rearranging terms (noting that integer multiplication commutes with φ) yields:

$$\varphi^2 - [\text{tr}(\varphi)]\varphi + [\deg(\varphi)] = [0].$$

Because $\deg$ is a positive definite quadratic form, the associated bilinear form satisfies the Cauchy-Schwarz inequality. This yields a universal bound on the trace that is one of the key components of the Riemann Hypothesis for elliptic curves over finite fields:

Corollary 2.2.12: Cauchy-Schwarz and the Hasse Bound

For any $\varphi \in \text{End}(E)$, the trace satisfies the bound:

$$|\text{tr}(\varphi)| \leq 2\sqrt{\deg(\varphi)}. \quad (2.9)$$

Proof:

For any rational numbers $m, n \in \mathbb{Q}$, the degree of $m\varphi + n[1]$ is a non-negative real number (extending the quadratic form to $\text{Hom}(E, E) \otimes \mathbb{Q}$). Expanding this using the bilinear pairing gives a quadratic polynomial in m and n :

$$\deg(m\varphi + n[1]) = m^2 \deg(\varphi) + mn \text{tr}(\varphi) + n^2 \geq 0.$$

For this real quadratic polynomial in the variable (m/n) to be strictly non-negative, its discriminant must be less than or equal to zero. The discriminant is $b^2 - 4ac = \text{tr}(\varphi)^2 - 4 \deg(\varphi)(1) \leq 0$. Rearranging this gives $\text{tr}(\varphi)^2 \leq 4 \deg(\varphi)$, which implies $|\text{tr}(\varphi)| \leq 2\sqrt{\deg(\varphi)}$.

These ideas shall be explored at length in chapter 5.

2.2.13 The Roots of the Characteristic Equation Because every endomorphism $\varphi \in \text{End}(E)$ satisfies the polynomial $P(X) = X^2 - \text{tr}(\varphi)X + \deg(\varphi)$, we can ask what the formal roots of this polynomial actually represent.

By the quadratic formula, the roots are:

$$\alpha, \bar{\alpha} = \frac{\text{tr}(\varphi) \pm \sqrt{\text{tr}(\varphi)^2 - 4\deg(\varphi)}}{2}$$

Crucially, the Cauchy-Schwarz bound (corollary 2.2.12) dictates that the discriminant $\text{tr}(\varphi)^2 - 4\deg(\varphi)$ is always less than or equal to zero. This physically forces the roots to be a pair of *complex conjugates* (which are only purely real when the discriminant is zero, corresponding to the standard integer multiplication map $\varphi = [n]$).

These complex conjugate roots carry immediate, profound meaning depending on the setting:

- *Analytically (Over $\mathbb{C}$):* If $E = \mathbb{C}/\Lambda$, the root α is literally the complex number that scales the lattice to induce the isogeny. Its modulus $|\alpha| = \sqrt{\deg(\varphi)}$ is the geometric stretch factor, and its argument $\arg(\alpha)$ is the angle of rotation.
- *Algebraically (The Tate Module):* When we later translate φ into a 2×2 matrix acting on the Tate module $T_\ell(E)$, the roots α and $\bar{\alpha}$ are exactly the *eigenvalues* of that matrix. The characteristic polynomial defined here purely via the degree form will perfectly match the characteristic polynomial of that linear algebraic matrix.

Exercise 2.2.1

1. For a non-constant isogeny $\varphi : E_1 \rightarrow E_2$, show that the pullback $\varphi^* : \text{Pic}^0(E_2) \rightarrow \text{Pic}^0(E_1)$ is a group homomorphism by verifying: (a) φ^* preserves degree-zero divisors, and (b) φ^* preserves linear equivalence (i.e., $\varphi^* \text{div}(g) = \text{div}(g \circ \varphi)$ for $g \in \overline{K}(E_2)^*$).
2. Verify $\hat{\varphi} \circ \varphi = [2]$ directly for the degree-2 isogeny $\varphi(x, y) = ((x^2 - 1)/x, y(x^2 + 1)/x^2)$ from $E : y^2 = x^3 - x$ to $E' : y^2 = x^3 + 4x$, by computing $\hat{\varphi}(\varphi(P))$ for a generic point $P = (x, y)$ and checking the result equals $[2](P)$. (You will need Vélu's formulas for $\hat{\varphi}$ on E' with kernel $\{O, (0, 0)\}$.)
3. Let $\varphi : E_1 \rightarrow E_2$ and $\psi : E_2 \rightarrow E_3$ be non-constant isogenies. Verify the contravariance identity $\widehat{\psi \circ \varphi} = \hat{\varphi} \circ \hat{\psi}$ by checking that both sides, when composed with $\psi \circ \varphi$, give $[\deg(\psi) \deg(\varphi)]$ on E_1 .
4. Show that $\widehat{[n]} = [n]$ using only the identity $\hat{\varphi} \circ \varphi = [\deg \varphi]$ and the multiplicativity of duals. (Hint: $[n] = [1] \circ [1] \circ \cdots \circ [1]$, and $\widehat{[1]} = [1]$.)
5. Let $\varphi : E_1 \rightarrow E_2$ be a separable isogeny of degree n . Show that $\ker(\hat{\varphi})$ is isomorphic to $E_2[n]/\varphi(E_1[n])$ (the quotient of the n -torsion by the image of the source's n -torsion under φ). (Hint: use $\varphi \circ \hat{\varphi} = [n]$ to show $\hat{\varphi}(E_2[n]) \subseteq E_1[n]$, then count.)
6. On $\text{End}(E)$ for $E : y^2 = x^3 - x$ over $\overline{\mathbb{Q}}$ (which has $\text{End}(E) = \mathbb{Z}[i]$), the degree form is $\deg(\alpha) = |\alpha|^2 = a^2 + b^2$ for $\alpha = a + bi$. Verify the parallelogram identity (2.5) for $\varphi = [2]$ and $\psi = [i]$ (the CM endomorphism $(x, y) \mapsto (-x, iy)$).

2.3 The Multiplication-by- n Map and the Structure of $E[n]$

In the previous section, we established that $\text{End}(E)$ is a torsion-free $\mathbb{Z}$ -module. Since the automorphism group always contains at least the identity $[1]$ and its negation $[-1]$, the endomorphism ring must have rank at least 1 over $\mathbb{Z}$. Therefore, the most foundational family of isogenies which must exist on absolutely every elliptic curve are the scalar multiplication maps $[n]$. In this section, we study these maps and their kernels.

As there is a natural group structure, all these maps are immediately definable as the multiplication-by- n map $[n] : E \rightarrow E$, defined by

$$P \mapsto \underbrace{P + P + \cdots + P}_{n \text{ times}}$$

for $n > 0$, $[0] = 0$, and $[-n] = -[n]$. In this section we prove the key facts: $\deg([n]) = n^2$, and the kernel $E[n]$ has the structure $(\mathbb{Z}/n\mathbb{Z})^2$ when n is coprime to the characteristic. The main tool is the theory of *division polynomials*, which give explicit formulas for $[n]$ in terms of the Weierstrass coordinates.

Division Polynomials

The idea is to express the coordinates of $[n](P)$ as rational functions of the coordinates of P . The denominators of these rational functions are the *division polynomials*, whose zeros are exactly the n -torsion points.

Definition 2.3.1: Division Polynomials

Let $E : y^2 = x^3 + Ax + B$ be an elliptic curve over K with $\text{char} K \neq 2, 3$. The *division polynomials* $\psi_n \in K[x, y]$ are defined recursively by:

$$\begin{aligned} \psi_0 &= 0, \\ \psi_1 &= 1, \\ \psi_2 &= 2y, \\ \psi_3 &= 3x^4 + 6Ax^2 + 12Bx - A^2, \\ \psi_4 &= 4y(x^6 + 5Ax^4 + 20Bx^3 - 5A^2x^2 - 4ABx - 8B^2 - A^3), \end{aligned}$$

with the general recurrence (for $n \geq 2$):

$$\psi_{2n+1} = \psi_{n+2}\psi_n^3 - \psi_{n-1}\psi_{n+1}^3, \quad (2.10)$$

$$2y \cdot \psi_{2n} = \psi_n(\psi_{n+2}\psi_{n-1}^2 - \psi_{n-2}\psi_{n+1}^2). \quad (2.11)$$

The division polynomials satisfy a parity structure: ψ_n is a polynomial in x alone when n is odd (after substituting $y^2 = x^3 + Ax + B$), and $\psi_n/(2y)$ is a polynomial in x when n is even. More precisely, define

$$\varphi_n = x\psi_n^2 - \psi_{n+1}\psi_{n-1}, \quad \omega_n = \frac{1}{4y}(\psi_{n+2}\psi_{n-1}^2 - \psi_{n-2}\psi_{n+1}^2). \quad (2.12)$$

Then the multiplication-by- n map is given by:

Proposition 2.3.2: Formulas for $[n]$

For $P = (x, y) \in E(\overline{K})$ with $P \notin E[n]$:

$$[n](P) = \left(\frac{\varphi_n(x)}{\psi_n(x)^2}, \frac{\omega_n(x, y)}{\psi_n(x, y)^3} \right). \quad (2.13)$$

The x -coordinate of $[n](P)$ depends only on x (not on y) and can be written as

$$x([n]P) = \frac{\varphi_n(x)}{\psi_n(x)^2} = x - \frac{\psi_{n-1}(x)\psi_{n+1}(x)}{\psi_n(x)^2}, \quad (2.14)$$

a rational function of x alone (after substituting $y^2 = x^3 + Ax + B$ wherever y appears).

Proof:

By induction on n , using the addition formulas of section 1.3. The base cases: for $n = 1$, $[1](x, y) = (x, y)$, and $\varphi_1 = x$, $\psi_1 = 1$. For $n = 2$ (doubling): the doubling formula gives $x([2]P) = (x^4 - 2Ax^2 - 8Bx + A^2)/(4(x^3 + Ax + B))$, which matches φ_2/ψ_2^2 with $\psi_2 = 2y$ (so $\psi_2^2 = 4y^2 = 4(x^3 + Ax + B)$). For the inductive step: the recurrences (2.10)–(2.11) are derived from applying the addition law to $[n+1]P = [n]P + P$ (for odd indices) and $[2n]P = [n]P + [n]P$ (for even indices), expressing the numerators and denominators of $x([n+1]P)$ and $x([2n]P)$ in terms of ψ_k and φ_k for $k \leq n$. The computation is lengthy but entirely mechanical: at each step one clears denominators in the addition formula, substitutes $y^2 = x^3 + Ax + B$, and identifies the resulting expressions with the recurrence. The interested reader can verify the $n = 3$ case in exercise 2.3.1 (1).

The Degree of $[n]$

The rational function $x([n]P) = \varphi_n(x)/\psi_n(x)^2$ determines the degree of $[n]$ via the degree of this rational function.

Lemma 2.3.3: Degree of ψ_n

As a polynomial in x (after substituting $y^2 = x^3 + Ax + B$), ψ_n^2 has degree $n^2 - 1$ in x for $n \geq 1$. More precisely:

1. For n odd, $\psi_n \in K[x]$ is a polynomial of degree $(n^2 - 1)/2$ with leading coefficient n .
2. For n even, $\psi_n/(2y) \in K[x]$ is a polynomial of degree $(n^2 - 4)/2$ with leading coefficient $n/2$, so $\psi_n^2/(4y^2) = \psi_n^2/(4(x^3 + Ax + B))$ has degree $n^2 - 4$ in x , and ψ_n^2 has degree $n^2 - 1$ in x .

Proof:

By induction on n , using the recurrences (2.10)–(2.11). The base cases: $\psi_1 = 1$ (degree $0 = (1^2 - 1)/2$, leading coefficient 1); $\psi_3 = 3x^4 + \dots$ (degree $4 = (9 - 1)/2$, leading coefficient 3); $\psi_2 = 2y$ (so $\psi_2/(2y) = 1$, degree $0 = (4 - 4)/2$, leading coefficient $1 = 2/2$); $\psi_4 = 4y(x^6 + \dots)$

(so $\psi_4/(2y) = 2x^6 + \dots$, degree $6 = (16 - 4)/2$, leading coefficient $2 = 4/2$). All match the claimed formulas.

For the inductive step, we treat the odd-index recurrence; the even-index case is analogous. Assume the result holds for all ψ_k with $k \leq 2n$. By (2.10), $\psi_{2n+1} = \psi_{n+2}\psi_n^3 - \psi_{n-1}\psi_{n+1}^3$. In x -degree (after substituting $y^2 = x^3 + Ax + B$ to eliminate y): both products $\psi_{n+2}\psi_n^3$ and $\psi_{n-1}\psi_{n+1}^3$ have the same total degree in x , namely $((2n+1)^2 - 1)/2 = 2n^2 + 2n$ (one verifies this by checking both parity cases for n , using the degree formulas in (1) and (2)). The leading coefficients are $(n+2) \cdot n^3$ and $(n-1) \cdot (n+1)^3$ respectively. Their difference is

$$(n+2)n^3 - (n-1)(n+1)^3 = n^4 + 2n^3 - (n^4 + 2n^3 - 2n - 1) = 2n + 1,$$

which is nonzero, confirming that the leading terms do not cancel. Thus ψ_{2n+1} has degree $2n^2 + 2n = ((2n+1)^2 - 1)/2$ and leading coefficient $2n + 1$, as claimed.

Since $\varphi_n = x\psi_n^2 - \psi_{n+1}\psi_{n-1}$, the numerator φ_n has degree n^2 in x (the term $x\psi_n^2$ dominates), and the denominator ψ_n^2 has degree $n^2 - 1$. Thus $x([n]P)$ is a rational function of degree n^2 in x . For the map $[n] : E \rightarrow E$ viewed as a morphism of curves, this means:

Theorem 2.3.4: Degree of the Multiplication Map

For any elliptic curve E/K and any integer $n \neq 0$,

$$\deg([n]) = n^2. \quad (2.15)$$

Proof:

The degree of the isogeny $[n] : E \rightarrow E$ is defined algebraically as the degree of the corresponding extension of function fields: $\deg([n]) = [K(E) : [n]^*K(E)]$.

The function field of E is $K(E) = K(x, y)$, and the pullback $[n]^*$ acts by substituting the coordinates of $[n]P$, meaning the subfield is $[n]^*K(E) = K(x([n]P), y([n]P))$. To compute the degree, we analyze the tower of field extensions over the base field $K(x([n]P))$.

On one hand, we can build $K(x, y)$ by passing through $K(x)$:

$$[K(x, y) : K(x([n]P))] = [K(x, y) : K(x)] \cdot [K(x) : K(x([n]P))].$$

The top extension $K(x, y)/K(x)$ has degree 2, as it is generated by y satisfying $y^2 = x^3 + Ax + B$. The bottom extension $K(x)/K(x([n]P))$ is generated by the rational function $x([n]P) = \varphi_n(x)/\psi_n(x)^2$. Because the numerator and denominator do not share roots (if they did, $x([n]P)$ would evaluate to $0/0$ rather than the required pole at n -torsion points), the degree of this extension is exactly the maximum of their degrees. By the preceding lemma, $\deg(\varphi_n) = n^2$ and $\deg(\psi_n^2) = n^2 - 1$, so this degree is $\max(n^2, n^2 - 1) = n^2$. Thus, the total degree via this path is $2n^2$.

On the other hand, we can build $K(x, y)$ by passing through the pullback field:

$$[K(x, y) : K(x([n]P))] = [K(x, y) : K(x([n]P), y([n]P))] \cdot [K(x([n]P), y([n]P)) : K(x([n]P))].$$

By definition, the first factor is exactly $\deg([n])$. The second factor is the degree of the extension generated by $y([n]P)$ over $K(x([n]P))$. Since $y([n]P)$ satisfies the Weierstrass equation $Y^2 = X^3 + AX + B$ (which is an irreducible quadratic over any rational function field in X), this degree is 2.

Equating the total degrees from the two paths yields $2 \cdot \deg([n]) = 2n^2$, which immediately gives $\deg([n]) = n^2$. Because this derivation relies only on algebraic identities of the division polynomials, it holds universally over any field, regardless of characteristic.

If $\text{char} K = 0$, there is a simpler more elegant proof worth presenting using theorem 2.2.5

Proof alternative:

If $m = 0$ it is true by definition, and if $m = 1$ it is true by proposition 2.2.4. For $m \neq 0, 1$ use induction with:

$$[\widehat{m+1}] = [\widehat{m}] + [\widehat{1}]$$

Then considering by definition $[\widehat{m}] = [\deg[m]]$; letting $d = \deg[m]$ we get, by definition of dual and of what we just proved:

$$[d] = [\widehat{m}] \circ [m] = [m] \circ [m] = [m^2]$$

And as $\text{End}(E)$ is torsion free, $d = m^2$.

The result $\deg([n]) = n^2$ is the algebraic counterpart of the analytic computation $[\Lambda : n\Lambda] = n^2$ (as shall be done in example 3.5.2). It holds in all characteristics, including the case $p|n$ in characteristic p (where $[n]$ is inseparable).

Separability of $[n]$

Proposition 2.3.5: Separability of the Multiplication Map

Let E/K be an elliptic curve and $n \geq 1$ an integer.

1. If $\text{char} K = 0$, or $\text{char} K = p > 0$ and $p \nmid n$, then $[n]$ is separable.
2. If $\text{char} K = p > 0$ and $p|n$, then $[n]$ is inseparable.
3. In particular, $[p]$ in characteristic p has inseparable degree either p or p^2 :
 - (a) $\deg_i([p]) = p$ and $\deg_s([p]) = p$ (*ordinary* case).
 - (b) $\deg_i([p]) = p^2$ and $\deg_s([p]) = 1$ (*supersingular* case).

Proof:

A morphism $\varphi : C_1 \rightarrow C_2$ of smooth curves over a field of characteristic p is inseparable if and only if the pullback $\varphi^* : \Omega_{C_2}^1 \rightarrow \Omega_{C_1}^1$ on differentials is the zero map. The invariant differential

on E is $\omega = dx/(2y)$. We compute:

$$[n]^*\omega = [n]^*\left(\frac{dx}{2y}\right) = \frac{d(x \circ [n])}{2(y \circ [n])}.$$

For a group scheme endomorphism, the pullback of the invariant differential satisfies $[n]^*\omega = n\omega$ (this can be checked by expanding $[n] = [n-1] + [1]$ and using the translation-invariance of ω , or by computing directly from the formal group in a neighborhood of O ; see chapter 4 for the full development).

Thus $[n]^*\omega = n\omega$, which vanishes if and only if $n = 0$ in K , i.e., $\text{char}K \mid n$. So $[n]$ is separable iff $\text{char}K \nmid n$. This proves (1) and (2).

For (3): since $\deg([p]) = p^2$ and $[p]$ is inseparable, $\deg_i([p]) \geq p$ (the inseparable degree is always a power of p). The two possibilities are $\deg_i = p$ (giving $\deg_s = p$, so $|E[p]| = p$) or $\deg_i = p^2$ (giving $\deg_s = 1$, so $|E[p]| = 1$). Both cases occur: the first is the ordinary case, the second is the supersingular case. See section 2.4 for the detailed analysis.

The formula $[n]^*\omega = n\omega$ is worth emphasizing. It says the invariant differential is an “eigenvector” for every multiplication map, with eigenvalue n . This is the algebraic incarnation of the fact that the analytic invariant differential dz on $\mathbb{C}/\Lambda$ satisfies $[n]^*(dz) = d(nz) = n dz$.

The Structure of $E[n]$

We now prove the torsion structure theorem, fulfilling the promise made in section 1.3.

Theorem 2.3.6: Structure of the n -Torsion

Let E/K be an elliptic curve and $n \geq 1$ an integer with $\gcd(n, \text{char}K) = 1$ (with the convention $\text{char}K = 0$ means no restriction). Then:

$$E[n](\overline{K}) \cong \mathbb{Z}/n\mathbb{Z} \times \mathbb{Z}/n\mathbb{Z}. \quad (2.16)$$

Proof:

Since $\gcd(n, \text{char}K) = 1$, the map $[n]$ is separable (proposition 2.3.5), so $|E[n](\overline{K})| = \deg_s([n]) = \deg([n]) = n^2$.

The group $E[n]$ is a finite abelian group killed by n (since $[n]P = O$ for every $P \in E[n]$). By the structure theorem for finite abelian groups, $E[n] \cong \mathbb{Z}/d_1\mathbb{Z} \times \mathbb{Z}/d_2\mathbb{Z}$ for some $d_1 \mid d_2 \mid n$ with $d_1 d_2 = n^2$. We need to show $d_1 = d_2 = n$.

Step 1: Reduce to prime powers. For $\gcd(m, n) = 1$, the Chinese Remainder Theorem gives $E[mn] \cong E[m] \times E[n]$ (since $E[mn]$ is killed by mn and has order $m^2 n^2$, and the natural map $E[mn] \rightarrow E[m] \times E[n]$ given by $P \mapsto ([m]P, [n]P)$ is an isomorphism). Thus it suffices to prove $E[\ell^k] \cong (\mathbb{Z}/\ell^k\mathbb{Z})^2$ for each prime power ℓ^k with $\ell \neq \text{char}K$.

Step 2: Prove $E[\ell] \cong (\mathbb{Z}/\ell\mathbb{Z})^2$ for a prime $\ell \neq \text{char}K$. The group $E[\ell]$ is a finite abelian group of order ℓ^2 killed by ℓ . Any finite abelian group of exponent ℓ is an $\mathbb{F}_\ell$ -vector space, so $E[\ell] \cong (\mathbb{Z}/\ell\mathbb{Z})^a$ for some $a \geq 0$. The order condition $\ell^a = \ell^2$ forces $a = 2$.

Step 3: Induction on k for $E[\ell^k] \cong (\mathbb{Z}/\ell^k\mathbb{Z})^2$. The multiplication map $[\ell] : E[\ell^{k+1}] \rightarrow E[\ell^k]$ is surjective (since $[\ell] : E \rightarrow E$ is surjective and $[\ell](E[\ell^{k+1}]) = E[\ell^k]$), and its kernel is $E[\ell]$. By the exact sequence

$$0 \longrightarrow E[\ell] \longrightarrow E[\ell^{k+1}] \xrightarrow{[\ell]} E[\ell^k] \longrightarrow 0,$$

we have $|E[\ell^{k+1}]| = |E[\ell]| \cdot |E[\ell^k]| = \ell^2 \cdot \ell^{2k} = \ell^{2(k+1)}$. Since $E[\ell^{k+1}]$ is killed by ℓ^{k+1} , it is isomorphic to $(\mathbb{Z}/\ell^{k+1}\mathbb{Z})^{a_1} \times (\mathbb{Z}/\ell^k\mathbb{Z})^{a_2} \times \cdots$ for some partition of $2(k+1)$. The surjectivity of $[\ell] : E[\ell^{k+1}] \rightarrow E[\ell^k]$ and the inductive hypothesis $E[\ell^k] \cong (\mathbb{Z}/\ell^k\mathbb{Z})^2$ force $E[\ell^{k+1}]$ to have at least two cyclic summands of order ℓ^{k+1} (otherwise the image of $[\ell]$ could not surject onto $(\mathbb{Z}/\ell^k\mathbb{Z})^2$). With $|E[\ell^{k+1}]| = \ell^{2(k+1)}$, the only possibility is $E[\ell^{k+1}] \cong (\mathbb{Z}/\ell^{k+1}\mathbb{Z})^2$.

The proof illustrates a general principle: the structure of $E[n]$ is determined by its order and the surjectivity of the multiplication maps. The key input is $|E[\ell]| = \ell^2$ (from $\deg([\ell]) = \ell^2$), which forces the rank to be 2. The induction then propagates this to all ℓ -power torsion.

Example 2.3.7: $E[3]$ for $y^2 = x^3 + 1$

Consider $E : y^2 = x^3 + 1$ over $\mathbb{Q}$ (or over $\overline{\mathbb{Q}}$, to see all torsion). We seek all points $P \in E(\overline{\mathbb{Q}})$ with $[3]P = O$. The division polynomial $\psi_3 = 3x^4 + 6 \cdot 0 \cdot x^2 + 12 \cdot 1 \cdot x - 0^2 = 3x^4 + 12x = 3x(x^3 + 4)$. Setting $\psi_3 = 0$: $x = 0$ or $x^3 = -4$. The roots $x = 0$ give $y^2 = 1$, so $y = \pm 1$: the points $(0, \pm 1)$, which we already know have order 3 (example 1.3.8). The roots $x^3 = -4$ give $x = -\sqrt[3]{4}$ (and its conjugates $-\sqrt[3]{4}\omega, -\sqrt[3]{4}\omega^2$ for $\omega = e^{2\pi i/3}$), with $y^2 = -4 + 1 = -3$, so $y = \pm\sqrt{-3}$. This gives 6 more points.

In total: $E[3]$ consists of O plus the 8 points above, for $|E[3]| = 9 = 3^2$. The group structure is $E[3] \cong (\mathbb{Z}/3\mathbb{Z})^2$, with $(0, 1)$ and $(-\sqrt[3]{4}, \sqrt{-3})$ as generators (one can verify they are not multiples of each other).

Over $\mathbb{Q}$ itself, only $(0, \pm 1)$ are rational, so $E[3](\mathbb{Q}) = \{O, (0, 1), (0, -1)\} \cong \mathbb{Z}/3\mathbb{Z}$. The remaining 6 points are defined over $\mathbb{Q}(\sqrt[3]{4}, \sqrt{-3})$, and the Galois group $\text{Gal}(\overline{\mathbb{Q}}/\mathbb{Q})$ permutes them — this Galois action on $E[3]$ is the beginning of the Galois representation theory that will be developed in ??.

Torsion in Characteristic p

When $\text{char}K = p > 0$ and $p \nmid n$, the map $[n]$ is inseparable and $|E[n](\overline{K})| < n^2$. The structure depends on whether E is ordinary or supersingular.

Proposition 2.3.8: p -Torsion Structure

Let E/K be an elliptic curve over a field of characteristic $p > 0$.

1. (*Ordinary case.*) If E is ordinary (i.e., $E[p](\overline{K}) \neq \{O\}$), then $E[p^r](\overline{K}) \cong \mathbb{Z}/p^r\mathbb{Z}$ for all $r \geq 1$.
2. (*Supersingular case.*) If E is supersingular (i.e., $E[p](\overline{K}) = \{O\}$), then $E[p^r](\overline{K}) = \{O\}$ for all $r \geq 1$.

For mixed torsion $E[n]$ with $n = p^r m$ and $\gcd(m, p) = 1$:

$$E[n](\overline{K}) \cong E[p^r](\overline{K}) \times E[m](\overline{K}) \cong \begin{cases} \mathbb{Z}/p^r\mathbb{Z} \times (\mathbb{Z}/m\mathbb{Z})^2 & \text{(ordinary),} \\ (\mathbb{Z}/m\mathbb{Z})^2 & \text{(supersingular).} \end{cases}$$

Proof:

The key input is proposition 2.3.5(3): in the ordinary case, $\deg_s([p]) = p$, so $|E[p](\overline{K})| = p$, hence $E[p] \cong \mathbb{Z}/p\mathbb{Z}$ (the only group of order p and exponent p). For the induction: the exact sequence $0 \rightarrow E[p] \rightarrow E[p^{r+1}] \xrightarrow{[p]} E[p^r] \rightarrow 0$ (where the map $[\ell] : E[p^{r+1}] \rightarrow E[p^r]$ is surjective since $[p] : E \rightarrow E$ is surjective) gives $|E[p^{r+1}]| = p \cdot |E[p^r]| = p^{r+1}$. Since $E[p^{r+1}]$ is killed by p^{r+1} and the surjection onto $E[p^r] \cong \mathbb{Z}/p^r\mathbb{Z}$ requires at least one summand of order p^{r+1} , we get $E[p^{r+1}] \cong \mathbb{Z}/p^{r+1}\mathbb{Z}$.

In the supersingular case, $\deg_s([p]) = 1$, so $|E[p](\overline{K})| = 1$, and $[p]$ is purely inseparable. Then $[p^r] = [p]^r$ is also purely inseparable, giving $E[p^r](\overline{K}) = \{O\}$.

The decomposition $E[n] \cong E[p^r] \times E[m]$ for $n = p^r m$ follows from $\gcd(p^r, m) = 1$ and the Chinese Remainder Theorem for abelian groups.

The ordinary/supersingular dichotomy will be analyzed in depth in section 2.4, where we characterize it in terms of the Frobenius endomorphism, and again in chapter 5, where we relate it to the trace of Frobenius and the point count $|E(\mathbb{F}_q)|$.

Example 2.3.9: p -Torsion on $y^2 = x^3 - x$ over $\mathbb{F}_5$

Consider $E : y^2 = x^3 - x$ over $\mathbb{F}_5$. We count: $|E(\mathbb{F}_5)| = (\text{compute } x^3 - x \pmod{5} \text{ for } x = 0, 1, 2, 3, 4)$:
 $x = 0 : 0$ ($y = 0$, one point); $x = 1 : 0$ ($y = 0$); $x = 2 : 6 \equiv 1$ (square, $y = \pm 1$); $x = 3 : 24 \equiv 4$ (square, $y = \pm 2$); $x = 4 : 60 \equiv 0$ ($y = 0$). Affine points: $3 + 2 + 2 = 7$, plus O : $|E(\mathbb{F}_5)| = 8$.

The trace of Frobenius is $a_5 = 5 + 1 - 8 = -2$, which is nonzero modulo 5, so E is ordinary over $\mathbb{F}_5$. We expect $E[5](\overline{\mathbb{F}_5}) \cong \mathbb{Z}/5\mathbb{Z}$: it has 5 elements, not 25.

On the other hand, the 2-torsion is full: $E[2](\mathbb{F}_5) = \{O, (0, 0), (1, 0), (4, 0)\} \cong (\mathbb{Z}/2\mathbb{Z})^2$ (since $x^3 - x = x(x-1)(x+1)$ splits completely over $\mathbb{F}_5$).

Exercise 2.3.1

1. Verify the formula $\psi_3 = 3x^4 + 6Ax^2 + 12Bx - A^2$ by computing $[3](x, y) = [2](x, y) + (x, y)$ using the doubling and addition formulas from section 1.3, and identifying the denominator of the x -coordinate.
2. From the doubling formula $x([2]P) = \frac{(3x^2+A)^2}{4(x^3+Ax+B)} - 2x$ (after simplification), verify that the numerator of $x([2]P)$ (written over a common denominator) is a polynomial of degree 4 in x , confirming $\deg([2]) = 4 = 2^2$.
3. Compute $E[4]$ for the curve $E : y^2 = x^3 - x$ over $\overline{\mathbb{Q}}$. (*Hint*: $E[4]$ consists of the points P with $[2]P \in E[2]$. Find the preimages of each 2-torsion point under the doubling map.) Verify that $|E[4]| = 16$ and $E[4] \cong (\mathbb{Z}/4\mathbb{Z})^2$.
4. Show that $[n]^*\omega = n\omega$ for the invariant differential $\omega = dx/(2y)$ by computing explicitly for $n = 2$ (using the doubling formulas). (*Hint*: compute $d(x([2]P))$ in terms of dx and compare with $2y([2]P) \cdot \omega$.)
5. Let $E : y^2 = x^3 + 1$ over $\mathbb{F}_2$. This curve has $|E(\mathbb{F}_2)| = 3$ (verify!), so the trace of Frobenius is $a_2 = 2 + 1 - 3 = 0$, and E is supersingular. Show directly that $E[2](\overline{\mathbb{F}_2}) = \{O\}$ by verifying that the equation $y^2 + y = x^3$ (the correct equation in characteristic 2, after a change of variables) has no affine solutions with $2P = O$.
6. For $E : y^2 = x^3 - x$ over $\mathbb{Q}$, determine the smallest field extension $\mathbb{Q} \subseteq L$ over which all of $E[3]$ is defined. (*Hint*: compute $\psi_3 = 3x^4 - 6x^2 + 12x - 1$ (with $A = -1$, $B = 0$... actually $-A^2 = -1$), find its splitting field, and adjoin the necessary y -coordinates.)

2.4 The Frobenius Endomorphism

In characteristic 0, every isogeny is separable and the theory of the previous two sections tells the whole story. In characteristic p , a fundamentally new actor enters: the *Frobenius endomorphism*, which raises coordinates to the p -th power. It is the prototypical inseparable isogeny — purely inseparable of degree p — and it controls the arithmetic of elliptic curves over finite fields. The factorization $[p] = \hat{\varphi}_p \circ \varphi_p$ (Frobenius composed with its dual, the Verschiebung) is the key to the ordinary/supersingular dichotomy that was stated in section 2.3 and is now proved in full.

The Frobenius Morphism

Let K be a field of characteristic $p > 0$.

Definition 2.4.1: The q -Frobenius Morphism

Let $E/\mathbb{F}_q$ be an elliptic curve defined over $\mathbb{F}_q$ (where $q = p^r$), given by a Weierstrass equation with coefficients in $\mathbb{F}_q$. The q -power *Frobenius morphism* is the map

$$\varphi_q : E \longrightarrow E, \quad (x, y) \longmapsto (x^q, y^q), \quad O \longmapsto O. \quad (2.17)$$

Since the Weierstrass coefficients lie in $\mathbb{F}_q$ (hence satisfy $a^q = a$), the point (x^q, y^q) satisfies the Weierstrass equation whenever (x, y) does: if $y^2 = x^3 + Ax + B$, then $(y^q)^2 = (y^2)^q = (x^3 + Ax + B)^q = x^{3q} + A^q x^q + B^q = (x^q)^3 + A(x^q) + B$. So φ_q is indeed a morphism $E \rightarrow E$.

More generally, if E is defined over a subfield $\mathbb{F}_{p^s} \subseteq \mathbb{F}_q$ but we work with the q -Frobenius, we can also consider the p -power Frobenius $\varphi_p : (x, y) \mapsto (x^p, y^p)$, which is a morphism $E \rightarrow E^{(p)}$, where $E^{(p)}$ denotes the curve obtained by raising the coefficients of E to the p -th power. When E is defined over $\mathbb{F}_p$, $E^{(p)} = E$ and φ_p is an endomorphism. The q -Frobenius is the r -fold iterate: $\varphi_q = \varphi_p^r$.

Proposition 2.4.2: Basic Properties of Frobenius

Let $E/\mathbb{F}_q$ be an elliptic curve.

1. φ_q is a non-constant isogeny (hence a group homomorphism, by theorem 2.1.3).
2. φ_q is purely inseparable: $\deg_s(\varphi_q) = 1$ and $\deg_i(\varphi_q) = q$.
3. $\deg(\varphi_q) = q$.
4. $\ker(\varphi_q) = \{O\}$ (the kernel consists of only the identity).
5. The fixed points of φ_q are exactly the $\mathbb{F}_q$ -rational points: $\{P \in E(\overline{\mathbb{F}_q}) : \varphi_q(P) = P\} = E(\mathbb{F}_q)$.

Proof:

1. The map φ_q is non-constant since $(x, y) \mapsto (x^q, y^q)$ is not the constant map to O (it preserves affine points).
- 2-3. The pullback on function fields is $\varphi_q^* : \overline{\mathbb{F}_q}(E) \rightarrow \overline{\mathbb{F}_q}(E)$, sending $f(x, y) \mapsto f(x^q, y^q)$. The image is $\varphi_q^*(\overline{\mathbb{F}_q}(E)) = \overline{\mathbb{F}_q}(x^q, y^q) = \overline{\mathbb{F}_q}(E)^q$ (the subfield of q -th powers). The extension $\overline{\mathbb{F}_q}(E)/\overline{\mathbb{F}_q}(E)^q$ is purely inseparable of degree q : every element $f \in \overline{\mathbb{F}_q}(E)$ satisfies $f^q \in \overline{\mathbb{F}_q}(E)^q$, so f is a root of $T^q - f^q \in \overline{\mathbb{F}_q}(E)^q[T]$, which is inseparable. The degree is q because $[\overline{\mathbb{F}_q}(x, y) : \overline{\mathbb{F}_q}(x^q, y^q)] = q$ (the element x satisfies $x^q - x^q = 0$ over $\overline{\mathbb{F}_q}(x^q, y^q)$, and the minimal polynomial of x over $\overline{\mathbb{F}_q}(x^q)$ is $T^q - x^q$ of degree q ; since y is algebraic of degree ≤ 2 over $\overline{\mathbb{F}_q}(x)$ and also over $\overline{\mathbb{F}_q}(x^q)$, the total degree is q).
4. Since φ_q is purely inseparable, $\deg_s(\varphi_q) = 1$, so by theorem 2.1.8(1), $|\varphi_q^{-1}(O)| = \deg_s(\varphi_q) = 1$. Hence $\ker(\varphi_q) = \{O\}$.
5. A point $P = (x_0, y_0) \in E(\overline{\mathbb{F}_q})$ satisfies $\varphi_q(P) = P$ iff $(x_0^q, y_0^q) = (x_0, y_0)$, iff $x_0^q = x_0$ and $y_0^q = y_0$, iff $x_0, y_0 \in \mathbb{F}_q$, iff $P \in E(\mathbb{F}_q)$. (The point O satisfies $\varphi_q(O) = O$ always.)

Property (5) is the reason Frobenius is so important for arithmetic: it detects rational points. The number of $\mathbb{F}_q$ -rational points on E is the number of fixed points of φ_q , which is $|E(\mathbb{F}_q)| = \deg(\varphi_q - [1]) = \deg(1 - \varphi_q)$ (we will explore further in chapter 5).

The multiplication-by- p map is inseparable in characteristic p (proposition 2.3.5), and the Frobenius endomorphism is the “inseparable part.” The precise relationship is:

Theorem 2.4.3: Frobenius–Verschiebung Factorization

Let $E/\mathbb{F}_q$ be an elliptic curve with $q = p^r$. There exists a unique isogeny $V_q : E \rightarrow E$, called the q -Verschiebung (or q -transfer), such that

$$V_q \circ \varphi_q = [q] \quad \text{and} \quad \varphi_q \circ V_q = [q]. \quad (2.18)$$

The Verschiebung V_q is the dual isogeny $\hat{\varphi}_q$ of Frobenius (definition 2.2.1).

Proof:

We prove the first identity $V_q \circ \varphi_q = [q]$ directly; the second identity $\varphi_q \circ V_q = [q]$ then follows immediately from proposition 2.2.4(2), which gives $\varphi \circ \hat{\varphi} = [\deg \varphi]$ for any isogeny φ .

Note that corollary 2.1.9 does not apply directly (it requires the first map to be separable, but φ_q is purely inseparable). We argue instead at the level of function fields.

The map $[q] : E \rightarrow E$ has $\deg([q]) = q^2$ and $[q]^*\omega = q\omega = 0$ (since $q = p^r = 0$ in $\mathbb{F}_q$), so $[q]$ is inseparable. The pullback on function fields gives $[q]^*\overline{\mathbb{F}_q}(E) \subseteq \overline{\mathbb{F}_q}(E)$ with $[\overline{\mathbb{F}_q}(E) : [q]^*\overline{\mathbb{F}_q}(E)] = q^2$. The Frobenius gives $\varphi_q^*\overline{\mathbb{F}_q}(E) = \overline{\mathbb{F}_q}(E)^q \subseteq \overline{\mathbb{F}_q}(E)$ with index q .

We claim $\varphi_q^*\overline{\mathbb{F}_q}(E) \supseteq [q]^*\overline{\mathbb{F}_q}(E)$. Indeed, for any $f \in \overline{\mathbb{F}_q}(E)$, $[q]^*f = f \circ [q]$. We need to show $f \circ [q] \in \overline{\mathbb{F}_q}(E)^q$, i.e., that $f \circ [q] = g^q$ for some $g \in \overline{\mathbb{F}_q}(E)$. Write $q = p^r$ and $[q] = [p]^r$. Since $[p]^*\omega = p\omega = 0$, the map $[p]$ factors through the Frobenius φ_p : any inseparable morphism in characteristic p factors through the p -Frobenius at the level of function fields (this is the universal property of the Frobenius — the p -power map is the “maximal inseparable” part of any inseparable morphism). Concretely: $[p]^*\overline{\mathbb{F}_q}(E) \subseteq \varphi_p^*\overline{\mathbb{F}_q}(E) = \overline{\mathbb{F}_q}(E)^p$. Iterating r times: $[q]^*\overline{\mathbb{F}_q}(E) = [p]^r\overline{\mathbb{F}_q}(E) \subseteq \overline{\mathbb{F}_q}(E)^{p^r} = \overline{\mathbb{F}_q}(E)^q = \varphi_q^*\overline{\mathbb{F}_q}(E)$.

The inclusion $[q]^*\overline{\mathbb{F}_q}(E) \subseteq \varphi_q^*\overline{\mathbb{F}_q}(E)$ gives a morphism $V_q : E \rightarrow E$ at the level of function fields: $\varphi_q^*(V_q^*\overline{\mathbb{F}_q}(E)) = [q]^*\overline{\mathbb{F}_q}(E)$, meaning $V_q \circ \varphi_q = [q]$ geometrically. Since $V_q(\varphi_q(O)) = [q](O) = O$ and $\varphi_q(O) = O$, we get $V_q(O) = O$, so V_q is an isogeny.

The degree of V_q is $\deg(V_q) = \deg([q])/\deg(\varphi_q) = q^2/q = q$ (using $V_q \circ \varphi_q = [q]$ and multiplicativity of degrees). Thus V_q is an isogeny of degree q , and the factorization $[q] = V_q \circ \varphi_q$ is established.

The factorization $[p] = V_p \circ \varphi_p$ (for $q = p$, $r = 1$) is the most important special case. It says that the multiplication-by- p map decomposes as: first apply Frobenius (purely inseparable, degree p), then apply the Verschiebung (which carries the separable part). The separability of V_p determines the ordinary/supersingular dichotomy:

Theorem 2.4.4: Ordinary and Supersingular Elliptic Curves

Let E be an elliptic curve over a field K of characteristic $p > 0$. The following are equivalent:

- (a) $E[p](\overline{K}) \cong \mathbb{Z}/p\mathbb{Z}$ (“ E is *ordinary*”).
- (b) $\deg_i([p]) = p$ (i.e., $[p]$ has inseparable degree exactly p , not p^2).
- (c) The Verschiebung V_p is separable.
- (d) The formal group $\hat{E}$ associated to E has height 1 (see chapter 4).

And the following are equivalent:

- (a') $E[p](\overline{K}) = \{O\}$ (“ E is *supersingular*”).
- (b') $\deg_i([p]) = p^2$ (i.e., $[p]$ is purely inseparable).
- (c') The Verschiebung V_p is inseparable.
- (d') The formal group $\hat{E}$ associated to E has height 2.

Every elliptic curve over $\overline{K}$ is either ordinary or supersingular.

Proof:

The equivalences (a) $\Leftrightarrow$ (b) and (a') $\Leftrightarrow$ (b') were established in proposition 2.3.5: since $\deg([p]) = p^2$ and $|E[p](\overline{K})| = \deg_s([p]) = p^2 / \deg_i([p])$, the two possibilities are $\deg_i = p$ (giving $|E[p]| = p$, hence $E[p] \cong \mathbb{Z}/p\mathbb{Z}$) or $\deg_i = p^2$ (giving $|E[p]| = 1$). There is no third option.

(b) $\Leftrightarrow$ (c): From the factorization $[p] = V_p \circ \varphi_p$, we have $\deg_i([p]) = \deg_i(V_p) \cdot \deg_i(\varphi_p) = \deg_i(V_p) \cdot p$ (since φ_p is purely inseparable of degree p). So:

- $\deg_i([p]) = p$ iff $\deg_i(V_p) = 1$ iff V_p is separable.
- $\deg_i([p]) = p^2$ iff $\deg_i(V_p) = p$ iff V_p is purely inseparable.

(b) $\Leftrightarrow$ (d): This will be proved in chapter 4, where the height of the formal group is defined and related to the inseparable degree of $[p]$ in a neighborhood of the identity. We include condition (d) here for completeness and forward reference.

That every curve is one or the other follows from the fact that $\deg_i([p])$ is a power of p dividing $\deg([p]) = p^2$, so $\deg_i([p]) \in \{1, p, p^2\}$. But $\deg_i([p]) = 1$ is impossible (since $[p]^*\omega = p\omega = 0$, the map $[p]$ is inseparable, hence $\deg_i([p]) > 1$). This leaves exactly two cases.

The terminology “supersingular” is somewhat misleading: supersingular curves are not singular (they are smooth). The term was introduced by Deuring and refers to the fact that these curves have “more endomorphisms than expected” — their endomorphism ring is an order in a quaternion algebra, rather than an order in an imaginary quadratic field or simply $\mathbb{Z}$ (see section 2.6). In fact, supersingular curves are quite special: there are only finitely many supersingular j -invariants in characteristic p (approximately $p/12$ of them), while the remaining j -invariants correspond to

ordinary curves.

Example 2.4.5: Ordinary and Supersingular in Small Characteristic

1. Let $E : y^2 = x^3 - x$ over $\mathbb{F}_5$. From example 2.3.9, $|E(\mathbb{F}_5)| = 8$, so $a_5 = 5 + 1 - 8 = -2 \not\equiv 0 \pmod{5}$. Thus E is ordinary over $\mathbb{F}_5$. The p -torsion is $E[5](\overline{\mathbb{F}_5}) \cong \mathbb{Z}/5\mathbb{Z}$.
2. Let $E : y^2 = x^3 + 1$ over $\mathbb{F}_2$. We have $|E(\mathbb{F}_2)| = 3$ (from exercise 2.3.1 (5)), so $a_2 = 2 + 1 - 3 = 0$. Thus E is supersingular over $\mathbb{F}_2$, and $E[2](\overline{\mathbb{F}_2}) = \{O\}$.
3. Let $E : y^2 = x^3 + x$ over $\mathbb{F}_3$. We compute $|E(\mathbb{F}_3)|$: for $x = 0$, $y^2 = 0$, $y = 0$ (one point); for $x = 1$, $y^2 = 2$ (non-square mod 3: no points); for $x = 2$, $y^2 = 2 + 2 = 1$, $y = \pm 1$ (two points). Total: $3 + 1 = 4$, so $a_3 = 3 + 1 - 4 = 0$. Thus E is supersingular over $\mathbb{F}_3$. In fact, for $p = 2, 3$, the curve $y^2 = x^3 + 1$ (resp. $y^2 = x^3 + x$) is the unique supersingular curve (up to isomorphism) in that characteristic.
4. Let the $j = 0$ and $j = 1728$ curves. Over $\overline{\mathbb{F}_p}$: the curve $y^2 = x^3 + 1$ (with $j = 0$) is supersingular iff $p \equiv 2 \pmod{3}$, and the curve $y^2 = x^3 - x$ (with $j = 1728$) is supersingular iff $p \equiv 3 \pmod{4}$. (These follow from the explicit point counts using quadratic residue arguments; see chapter 5.)

The Frobenius endomorphism φ_q satisfies a degree-2 polynomial relation, which encodes the point count $|E(\mathbb{F}_q)|$.

Theorem 2.4.6: Characteristic Polynomial of Frobenius

Let $E/\mathbb{F}_q$ be an elliptic curve. The Frobenius endomorphism $\varphi_q \in \text{End}(E)$ satisfies

$$\varphi_q^2 - a_q \varphi_q + [q] = [0] \quad \text{in } \text{End}(E), \quad (2.19)$$

where $a_q = q + 1 - |E(\mathbb{F}_q)|$ is the *trace of Frobenius*. Equivalently, φ_q is a root of the polynomial $T^2 - a_q T + q \in \mathbb{Z}[T]$.

I actually decided to prove this earlier so this is no longer valid, and maybe I don't even need this here anymore.

Proof:

We defer the full proof to chapter 5, where it will be established using the theory of the dual isogeny and the Weil pairing: $\deg(\varphi_q - [n]) = \varphi_q^2 - a_q n + qn^2$ as a quadratic form in n , which forces $\varphi_q^2 - a_q \varphi_q + q = 0$ in $\text{End}(E)$. The key input is the Hasse bound $|a_q| \leq 2\sqrt{q}$ (see chapter 5).

For now, we record the consequence: the Frobenius and Verschiebung satisfy $\varphi_q + V_q = [a_q]$ in

$\text{End}(E)$. This follows from (2.19): $V_q = [q]/\varphi_q = ([a_q]\varphi_q - \varphi_q^2)/\varphi_q = [a_q] - \varphi_q$ (dividing the relation $\varphi_q^2 - a_q\varphi_q + q = 0$ by φ_q , using $[q] = V_q \circ \varphi_q$).

The identity $\varphi_q + V_q = [a_q]$ is the algebraic incarnation of the point count: it says that the “sum” of Frobenius and Verschiebung acts as multiplication by the trace. In particular, E is supersingular iff $a_p \equiv 0 \pmod{p}$ (for $q = p$), iff $V_p = -\varphi_p$ (up to a p -torsion correction), iff the Frobenius and Verschiebung “collide.”

Frobenius and the Function Field

For later use (especially in the construction of the dual isogeny), we record the function-field behavior of Frobenius.

Proposition 2.4.7: Frobenius and the Function Field

Let $E/\mathbb{F}_q$ be an elliptic curve.

1. The pullback $\varphi_q^* : \overline{\mathbb{F}_q}(E) \rightarrow \overline{\mathbb{F}_q}(E)$ sends $f \mapsto f^q$ (the q -th power map on functions). Its image is $\overline{\mathbb{F}_q}(E)^q$, the subfield of q -th powers.
2. Any inseparable isogeny $\psi : E_1 \rightarrow E_2$ in characteristic p factors through the Frobenius: $\psi = \lambda \circ \varphi_p^s$ for some separable isogeny λ and some $s \geq 1$ with $p^s = \deg_i(\psi)$.
3. The invariant differential satisfies $\varphi_q^*\omega = 0$ (since $\varphi_q^*\omega = d(x^q)/(2y^q) = qx^{q-1}dx/(2y^q) = 0$ in characteristic p dividing q).

Proof:

1. this is immediate from the definition: $\varphi_q^*f(x, y) = f(x^q, y^q) = f(x, y)^q$ (using the Frobenius endomorphism of $\overline{\mathbb{F}_q}$).
2. The inseparable part of a field extension L/M in characteristic p is generated by the p -th power map. More precisely: if $[L : M]_i = p^s$, then $M \subseteq L^{p^s} \subseteq L$ and L/L^{p^s} is purely inseparable of degree p^s . For an inseparable morphism $\psi : E_1 \rightarrow E_2$, the inclusion $\psi^*\overline{\mathbb{F}_q}(E_2) \subseteq \overline{\mathbb{F}_q}(E_1)^{p^s} = (\varphi_p^s)^*\overline{\mathbb{F}_q}(E_1)$ gives the factorization $\psi = \lambda \circ \varphi_p^s$ at the level of function fields, hence geometrically.
3. Follows from direct computation: $\varphi_q^*(dx/(2y)) = d(x^q)/(2y^q) = qx^{q-1}dx/(2y^q)$. Since $\text{char} K = p$ and $p|q$, we have $q = 0$ in K , so $\varphi_q^*\omega = 0$.

Exercise 2.4.1

1. Let $E : y^2 = x^3 + x + 1$ over $\mathbb{F}_7$. Compute $|E(\mathbb{F}_7)|$ by evaluating $x^3 + x + 1$ for each $x \in \mathbb{F}_7$ and counting the number of y -values. Determine a_7 , and decide whether E is ordinary or supersingular over $\mathbb{F}_7$.
2. Show directly from the definition that $\varphi_q^2 = \varphi_{q^2}$: the q -Frobenius composed with itself is the q^2 -Frobenius. Conclude $\deg(\varphi_q^n) = q^n$ for all $n \geq 1$.

3. Using the factorization $[p] = V_p \circ \varphi_p$ and the degrees $\deg([p]) = p^2$, $\deg(\varphi_p) = p$, determine $\deg(V_p) = p$. Is V_p separable, inseparable, or purely inseparable? (*Answer: it depends on whether E is ordinary or supersingular.*)
4. Over $\overline{\mathbb{F}_5}$, there is a unique supersingular j -invariant. Show it is $j = 0$ by verifying that $y^2 = x^3 + 1$ has $|E(\mathbb{F}_5)| = 6$ (hence $a_5 = 0$), and that the other special curve $y^2 = x^3 - x$ has $a_5 = -2 \neq 0$ (hence is ordinary). (*For $p = 5$, there is exactly one supersingular j -invariant, consistent with the formula $[p/12] + \epsilon$ for the number of supersingular j -invariants.*)
5. Let $E/\mathbb{F}_p$ be an ordinary elliptic curve. Show that $\ker(V_p) \cong \mathbb{Z}/p\mathbb{Z}$ (a cyclic group of order p), and that $\ker(\varphi_p) = \{O\}$. Conclude that the exact sequence $0 \rightarrow \ker(\varphi_p) \rightarrow E[p] \xrightarrow{\varphi_p} \ker(V_p) \rightarrow 0$ degenerates to $E[p] \cong \ker(V_p) \cong \mathbb{Z}/p\mathbb{Z}$.
6. Let $\psi : E_1 \rightarrow E_2$ be an isogeny in characteristic p with $\deg(\psi) = p^2$ and $\deg_i(\psi) = p$. Show that ψ factors as $\psi = \lambda \circ \varphi_p$ where $\lambda : E_1^{(p)} \rightarrow E_2$ is a separable isogeny of degree p . (*This illustrates proposition 2.4.7(2).*)

2.5 The Tate Module and the Weil Pairing

As will be shown in chapter 3, there is a lattice Λ representing an elliptic curve as a free $\mathbb{Z}$ -module of rank 2, and all the key constructions (isogenies, the dual, the endomorphism ring) were expressed as linear algebra on this module. Over a general field, we have no lattice — but we have the next best thing: the *Tate module* $T_\ell(E)$, which packages all ℓ -power torsion into a single free $\mathbb{Z}_\ell$ -module of rank 2. This is the algebraic substitute for Λ , and it carries a natural symplectic form — the *Weil pairing* — that is the algebraic substitute for the intersection pairing on Λ .

Fix a prime $\ell \neq \text{char} K$.

Definition 2.5.1: The Tate Module

The ℓ -adic Tate module of E is the limit

$$T_\ell(E) = \varprojlim_n E[\ell^n], \quad (2.20)$$

where the transition maps are the multiplication-by- ℓ maps $[\ell] : E[\ell^{n+1}] \rightarrow E[\ell^n]$. The *rationalized Tate module* is $V_\ell(E) = T_\ell(E) \otimes_{\mathbb{Z}_\ell} \mathbb{Q}_\ell$.

An element of $T_\ell(E)$ is a sequence $(P_1, P_2, P_3, \dots)$ with $P_n \in E[\ell^n]$ and $[\ell]P_{n+1} = P_n$ for all $n \geq 1$. Such a sequence is a “compatible system of ℓ -power torsion points” — it is the algebraic analogue of choosing a generator of the ℓ -primary part of the lattice Λ .

Proposition 2.5.2: Structure of the Tate Module

$T_\ell(E)$ is a free $\mathbb{Z}_\ell$ -module of rank 2. Equivalently, $V_\ell(E)$ is a $\mathbb{Q}_\ell$ -vector space of dimension 2.

Proof:

By theorem 2.3.6, $E[\ell^n] \cong (\mathbb{Z}/\ell^n\mathbb{Z})^2$ for all $n \geq 1$ (since $\ell \neq \text{char}K$). The multiplication map $[\ell] : E[\ell^{n+1}] \rightarrow E[\ell^n]$ corresponds to the natural surjection $(\mathbb{Z}/\ell^{n+1}\mathbb{Z})^2 \twoheadrightarrow (\mathbb{Z}/\ell^n\mathbb{Z})^2$. Taking the limit:

$$T_\ell(E) = \varprojlim (\mathbb{Z}/\ell^n\mathbb{Z})^2 = \left(\varprojlim \mathbb{Z}/\ell^n\mathbb{Z} \right)^2 = \mathbb{Z}_\ell^2.$$

The Tate module linearizes the theory of isogenies:

Proposition 2.5.3: Isogenies Induce Linear Maps on Tate Modules

Let $\varphi : E_1 \rightarrow E_2$ be an isogeny. Then φ induces a $\mathbb{Z}_\ell$ -linear map

$$T_\ell(\varphi) : T_\ell(E_1) \longrightarrow T_\ell(E_2), \quad (P_1, P_2, \dots) \longmapsto (\varphi(P_1), \varphi(P_2), \dots).$$

This map is injective if $\varphi \neq 0$, and the assignment $\varphi \mapsto T_\ell(\varphi)$ is a ring homomorphism $\text{End}(E) \rightarrow \text{End}_{\mathbb{Z}_\ell}(T_\ell(E)) \cong M_2(\mathbb{Z}_\ell)$.

Proof:

The map is well-defined: φ restricts to $\varphi : E_1[\ell^n] \rightarrow E_2[\ell^n]$ (corollary 2.1.4), and the compatibility $[\ell]\varphi(P_{n+1}) = \varphi([\ell]P_{n+1}) = \varphi(P_n)$ holds since φ commutes with $[\ell]$.

$\mathbb{Z}_\ell$ -linearity: $T_\ell(\varphi)(aP + bQ) = \varphi(aP + bQ) = a\varphi(P) + b\varphi(Q) = aT_\ell(\varphi)(P) + bT_\ell(\varphi)(Q)$ for $a, b \in \mathbb{Z}_\ell$ (interpreted via the inverse limit).

Injectivity for $\varphi \neq 0$: if $T_\ell(\varphi) = 0$, then $\varphi(P_n) = O$ for all $P_n \in E_1[\ell^n]$, so $E_1[\ell^n] \subseteq \ker(\varphi)$ for all n . But $|E_1[\ell^n]| = \ell^{2n} \rightarrow \infty$, while $|\ker(\varphi)| < \infty$, a contradiction.

Ring homomorphism: $T_\ell(\varphi \circ \psi) = T_\ell(\varphi) \circ T_\ell(\psi)$ and $T_\ell(\varphi + \psi) = T_\ell(\varphi) + T_\ell(\psi)$ (both follow from the pointwise definitions).

Corollary 2.5.4: The Rank Bound for Hom

For elliptic curves E_1, E_2 over $\overline{K}$ with $\ell \neq \text{char}K$:

$$\text{Hom}(E_1, E_2) \otimes_{\mathbb{Z}} \mathbb{Z}_\ell \hookrightarrow \text{Hom}_{\mathbb{Z}_\ell}(T_\ell(E_1), T_\ell(E_2)) \cong M_2(\mathbb{Z}_\ell).$$

In particular, $\text{rank}_{\mathbb{Z}} \text{Hom}(E_1, E_2) \leq 4$.

Proof:

The map $\varphi \mapsto T_\ell(\varphi)$ is a $\mathbb{Z}$ -linear injection $\text{Hom}(E_1, E_2) \hookrightarrow \text{Hom}_{\mathbb{Z}_\ell}(T_\ell(E_1), T_\ell(E_2))$ (injectivity was shown in proposition 2.5.3). Because $\mathbb{Z}_\ell$ is a flat $\mathbb{Z}$ -module, extending scalars to $\mathbb{Z}_\ell$ preserves this injectivity, giving $\text{Hom}(E_1, E_2) \otimes_{\mathbb{Z}} \mathbb{Z}_\ell \hookrightarrow M_2(\mathbb{Z}_\ell)$. Since $M_2(\mathbb{Z}_\ell)$ is a free $\mathbb{Z}_\ell$ -module of rank 4, the left side has rank ≤ 4 as a $\mathbb{Z}_\ell$ -module, hence $\text{rank}_{\mathbb{Z}} \text{Hom}(E_1, E_2) \leq 4$.

Over finite fields, a deep theorem of Tate [23] says the injection is an *isomorphism*: $\text{Hom}(E_1, E_2) \otimes \mathbb{Z}_\ell \cong \text{Hom}_{\mathbb{Z}_\ell}(T_\ell(E_1), T_\ell(E_2))$. This will be proved in chapter 5.

The Algebraic Weil Pairing

more words here giving intuition! The pairing is super interesting!

Over $\mathbb{C}$, the Weil pairing was defined by the explicit formula $e_n(P, Q) = \zeta_n^{ad-bc}$ (definition 3.5.9). Algebraically, we construct it using functions with prescribed divisors and the Weil reciprocity law.

Theorem 2.5.5: The Weil Pairing

Let E/K be an elliptic curve and $n \geq 1$ an integer with $\gcd(n, \text{char}K) = 1$. There exists a pairing

$$e_n : E[n] \times E[n] \longrightarrow \mu_n \quad (2.21)$$

(where μ_n is the group of n -th roots of unity in $\overline{K}^*$) satisfying:

1. (*Bilinearity.*) $e_n(P_1 + P_2, Q) = e_n(P_1, Q) e_n(P_2, Q)$, and similarly in the second variable.
2. (*Alternating.*) $e_n(P, P) = 1$ for all P , and $e_n(P, Q) = e_n(Q, P)^{-1}$.
3. (*Nondegeneracy.*) If $e_n(P, Q) = 1$ for all $Q \in E[n]$, then $P = O$.
4. (*Galois equivariance.*) For $\sigma \in \text{Gal}(\overline{K}/K)$: $e_n(\sigma P, \sigma Q) = \sigma(e_n(P, Q))$.
5. (*Compatibility with isogenies.*) For an isogeny $\varphi : E_1 \rightarrow E_2$: $e_n(\varphi(P), \varphi(Q)) = e_n(P, Q)^{\deg \varphi}$.
6. (*Compatibility with $[\ell]$.*) For $P, Q \in E[\ell^{n+1}]$: $e_{\ell^{n+1}}(P, Q)^\ell = e_{\ell^n}([\ell]P, [\ell]Q)$.

Proof:

Construction. Let $P, Q \in E[n]$. We construct $e_n(P, Q)$ using the function field of E .

Step 1: Choose a function for P . Since $P \in E[n]$, we have $[n]P = O$. By the definition of the group law, the divisor class of $(P) - (O)$ in the Picard group satisfies $n[(P) - (O)] = [([n]P) - (O)] = [(O) - (O)] = 0$. This means the divisor $n(P) - n(O)$ is principal. Thus, there exists a function $f_P \in \overline{K}(E)^*$ such that

$$\text{div}(f_P) = n(P) - n(O).$$

Step 2: Pull back by $[n]$. Since the isogeny $[n]$ is unramified (as $\gcd(n, \text{char}K) = 1$), pulling back the divisor of f_P simply multiplies the points in the preimage:

$$[n]^* \text{div}(f_P) = n \cdot [n]^*(P) - n \cdot [n]^*(O).$$

Because we also know $[n]^* \text{div}(f_P) = \text{div}(f_P \circ [n])$, this tells us that $\text{div}(f_P \circ [n])$ is exactly n times the divisor $[n]^*(P) - [n]^*(O)$. Therefore, the divisor $[n]^*(P) - [n]^*(O)$ itself must be principal. We can choose a function $g_P \in \overline{K}(E)^*$ such that

$$\text{div}(g_P) = [n]^*(P) - [n]^*(O) \quad \text{and} \quad g_P^n = f_P \circ [n].$$

Step 3: Define the pairing. For any $X \in E$, consider the value $g_P(X + Q)$. Because $Q \in E[n]$, we know $[n](X + Q) = [n]X + [n]Q = [n]X$. Raising $g_P(X + Q)$ to the n -th power yields:

$$g_P(X + Q)^n = f_P([n](X + Q)) = f_P([n]X) = g_P(X)^n.$$

This implies that the ratio $g_P(X + Q)/g_P(X)$ raised to the n -th power is 1. Therefore, the ratio is a constant (an n -th root of unity) that is independent of the choice of X . We define the Weil pairing as this constant:

$$e_n(P, Q) = \frac{g_P(X + Q)}{g_P(X)}. \quad (2.22)$$

Properties. (1) Bilinearity in Q is immediate from the definition:

$$e_n(P, Q_1 + Q_2) = \frac{g_P(X + Q_1 + Q_2)}{g_P(X)} = \frac{g_P(X + Q_1 + Q_2)}{g_P(X + Q_2)} \frac{g_P(X + Q_2)}{g_P(X)} = e_n(P, Q_1) \cdot e_n(P, Q_2).$$

Bilinearity in P requires a separate argument verifying that $g_{P_1+P_2}$ is proportional to $g_{P_1}g_{P_2}$; we refer to [21, p. III.8.1].

(2) Alternating: Proving $e_n(P, P) = 1$ involves a careful application of the Weil reciprocity law to functions with disjoint supports; see [21, p. III.8.1]. Once $e_n(P, P) = 1$ is known, bilinearity implies $1 = e_n(P + Q, P + Q) = e_n(P, P)e_n(P, Q)e_n(Q, P)e_n(Q, Q) = e_n(P, Q)e_n(Q, P)$, so $e_n(Q, P) = e_n(P, Q)^{-1}$.

(3) Nondegeneracy: Suppose $e_n(P, Q) = 1$ for all $Q \in E[n]$. Then $g_P(X + Q) = g_P(X)$ for all $Q \in \ker([n])$, meaning g_P is invariant under translation by the kernel of the separable isogeny $[n]$. By Galois theory for function fields, g_P must descend to the quotient curve, meaning there is a function $h \in \overline{K}(E)^*$ such that $g_P = h \circ [n]$.

Then $f_P \circ [n] = g_P^n = h^n \circ [n]$, which implies $f_P = h^n$. Taking divisors, $n(P) - n(O) = \operatorname{div}(f_P) = n\operatorname{div}(h)$, so $\operatorname{div}(h) = (P) - (O)$. However, a function with a single simple pole provides an isomorphism from E to $\mathbb{P}^1$, which is impossible for a genus 1 curve unless the function is constant, meaning $P = O$.

(4) Galois equivariance: For $\sigma \in \operatorname{Gal}(\overline{K}/K)$, applying σ to $\operatorname{div}(f_P)$ gives $\operatorname{div}(\sigma f_P) = n(\sigma P) - n(O)$, so we may choose $f_{\sigma P} = \sigma f_P$. This identically leads to choosing $g_{\sigma P} = \sigma g_P$. Then:

$$e_n(\sigma P, \sigma Q) = \frac{(\sigma g_P)(X + \sigma Q)}{(\sigma g_P)(X)} = \sigma \left(\frac{g_P(\sigma^{-1}X + Q)}{g_P(\sigma^{-1}X)} \right) = \sigma(e_n(P, Q)).$$

(5) and (6) follow by similar algebraic manipulations of the pullback divisors; see the exercises.

The compatibility property (6) allows us to assemble the Weil pairings e_{ℓ^n} into a single pairing on the Tate module.

Definition 2.5.6: The ℓ -adic Weil Pairing

The ℓ -adic Weil pairing is the $\mathbb{Z}_\ell$ -bilinear map

$$e_\ell : T_\ell(E) \times T_\ell(E) \longrightarrow T_\ell(\mu) := \varprojlim \mu_{\ell^n} \cong \mathbb{Z}_\ell(1), \quad (2.23)$$

defined by $e_\ell((P_n)_n, (Q_n)_n) = (e_{\ell^n}(P_n, Q_n))_n$, where the compatibility $e_{\ell^{n+1}}(P_{n+1}, Q_{n+1})^\ell = e_{\ell^n}(P_n, Q_n)$ (property (6)) ensures that the sequence is compatible in the inverse limit.

The ℓ -adic Weil pairing inherits all properties of e_n : it is $\mathbb{Z}_\ell$ -bilinear, alternating, nondegenerate, and Galois-equivariant. The target $\mathbb{Z}_\ell(1)$ is the Tate twist: as a $\mathbb{Z}_\ell$ -module it is free of rank 1, but $\text{Gal}(\overline{K}/K)$ acts via the ℓ -adic cyclotomic character $\chi_\ell : \text{Gal}(\overline{K}/K) \rightarrow \mathbb{Z}_\ell^*$.

The Weil pairing gives a fundamental relationship between isogenies and determinants.

Theorem 2.5.7: Determinant of an Endomorphism

For any isogeny $\varphi : E \rightarrow E$ (i.e., $\varphi \in \text{End}(E)$), the determinant of the induced map $T_\ell(\varphi) : T_\ell(E) \rightarrow T_\ell(E)$ satisfies

$$\det(T_\ell(\varphi)) = \deg(\varphi). \quad (2.24)$$

More generally, for an isogeny $\varphi : E_1 \rightarrow E_2$: $e_n(\varphi(P), \varphi(Q)) = e_n(P, Q)^{\deg \varphi}$ for all $P, Q \in E_1[n]$.

Proof:

The “more generally” statement is property (5) of the Weil pairing. For the determinant formula: choose a $\mathbb{Z}_\ell$ -basis (e_1, e_2) of $T_\ell(E)$. Then $e_\ell(e_1, e_2) \in \mathbb{Z}_\ell(1)$ is a $\mathbb{Z}_\ell$ -generator (by nondegeneracy). Writing $T_\ell(\varphi)(e_i) = \sum a_{ij}e_j$:

$$e_\ell(T_\ell(\varphi)(e_1), T_\ell(\varphi)(e_2)) = e_\ell(e_1, e_2)^{\deg \varphi} \quad (\text{by property (5)}).$$

On the other hand, by bilinearity and the alternating property:

$$\begin{aligned} e_\ell(T_\ell(\varphi)(e_1), T_\ell(\varphi)(e_2)) &= e_\ell(a_{11}e_1 + a_{21}e_2, a_{12}e_1 + a_{22}e_2) \\ &= e_\ell(e_1, e_2)^{a_{11}a_{22} - a_{12}a_{21}} = e_\ell(e_1, e_2)^{\det(T_\ell(\varphi))}. \end{aligned}$$

Since $e_\ell(e_1, e_2)$ is a generator of the free rank-1 module $\mathbb{Z}_\ell(1)$, comparing exponents gives $\det(T_\ell(\varphi)) = \deg(\varphi)$.

This is a powerful result: it says the degree of an endomorphism (a geometric quantity, defined via function fields) equals the determinant of its action on the Tate module (a linear-algebraic quantity). Combined with the trace formula $\text{tr}(T_\ell(\varphi)) = \deg(\varphi + [1]) - \deg(\varphi) - 1$ (from expanding the degree of $\varphi + [1]$), one recovers the full characteristic polynomial of $T_\ell(\varphi)$.

For the Frobenius endomorphism φ_q on $E/\mathbb{F}_q$: $\det(T_\ell(\varphi_q)) = \deg(\varphi_q) = q$, and $\text{tr}(T_\ell(\varphi_q)) = a_q = q + 1 - |E(\mathbb{F}_q)|$. The matrix $T_\ell(\varphi_q) \in M_2(\mathbb{Z}_\ell)$ is determined (up to conjugacy) by its characteristic polynomial $T^2 - a_q T + q$, recovering the statement of theorem 2.4.6. **fix this reference once I delete this theorem**

With the Tate module and the Weil pairing in hand, we shall establish an analytic-algebraic dictionary (the full analytic dictionary will be explored in section 3.5). For reference and interest of the reader, the following is what we shall build up to:

Analytic	Algebraic	Type
Lattice $\Lambda \cong \mathbb{Z}^2$	Tate module $T_\ell(E) \cong \mathbb{Z}_\ell^2$	Free module of rank 2
Intersection form on Λ	Weil pairing e_ℓ on $T_\ell(E)$	Alternating, nondegenerate
$\alpha \in \mathbb{C}$ with $\alpha\Lambda \subseteq \Lambda$	$T_\ell(\varphi) \in M_2(\mathbb{Z}_\ell)$	Linear map
$ \alpha ^2 = \deg(\varphi)$	$\det(T_\ell(\varphi)) = \deg(\varphi)$	Norm = determinant

The Tate module will return in ??, where the Galois action $\text{Gal}(\overline{K}/K) \rightarrow \text{Aut}(T_\ell(E)) \cong \text{GL}_2(\mathbb{Z}_\ell)$ defines the ℓ -adic Galois representation attached to E . The Weil pairing will ensure $\det(\rho_{E,\ell}) = \chi_\ell$ (the cyclotomic character), and the classification of $\text{End}(E)$ from section 2.6 will constrain the image of this representation.

Exercise 2.5.1

1. For $E : y^2 = x^3 - x$ over $\overline{\mathbb{Q}}$ with $\ell = 3$, choose compatible systems of 3-power torsion to exhibit an explicit $\mathbb{Z}_3$ -basis of $T_3(E)$. (*Hint*: start with a basis (P_1, Q_1) of $E[3] \cong (\mathbb{Z}/3\mathbb{Z})^2$ and extend to $E[9]$ by choosing P_2, Q_2 with $[3]P_2 = P_1$, $[3]Q_2 = Q_1$.)
2. For $E : y^2 = x^3 - x$ over $\mathbb{F}_5$ (ordinary, $a_5 = -2$), the Frobenius φ_5 acts on $T_\ell(E)$ for any $\ell \neq 5$. What is the characteristic polynomial of $T_\ell(\varphi_5)$? Verify that $\det(T_\ell(\varphi_5)) = 5$ and $\text{tr}(T_\ell(\varphi_5)) = -2$.
3. Verify the bilinearity of e_n in the second variable directly from the construction (2.22): show $e_n(P, Q_1 + Q_2) = e_n(P, Q_1) \cdot e_n(P, Q_2)$ using the base-point shifting argument.
4. Prove property (5) of the Weil pairing: $e_n(\varphi(P), \varphi(Q)) = e_n(P, Q)^{\deg \varphi}$ for a separable isogeny φ . (*Hint*: if $\text{div}(f_P) = n(P) - n(O)$, what is the divisor of $f_P \circ \hat{\varphi}$? Use $\hat{\varphi} \circ \varphi = [\deg \varphi]$.)
5. Using the Galois equivariance of e_n and the determinant formula, show that for $\sigma \in \text{Gal}(\overline{K}/K)$:

$$\det(\rho_{E,\ell}(\sigma)) = \chi_\ell(\sigma),$$

where $\rho_{E,\ell} : \text{Gal}(\overline{K}/K) \rightarrow \text{GL}_2(\mathbb{Z}_\ell)$ is the representation on $T_\ell(E)$ and χ_ℓ is the ℓ -adic cyclotomic character.

6. Show that the injection $\text{End}(E) \otimes \mathbb{Z}_\ell \hookrightarrow M_2(\mathbb{Z}_\ell)$ is *not* always surjective. (*Hint*: for $\text{End}(E) = \mathbb{Z}$, the image is $\mathbb{Z}_\ell \cdot I_2$, a rank-1 submodule of the rank-4 module $M_2(\mathbb{Z}_\ell)$.)

2.6 The Endomorphism Ring

The endomorphism ring $\text{End}(E) = \text{Hom}(E, E)$ is a ring under composition and addition: $(\varphi + \psi)(P) = \varphi(P) + \psi(P)$ and $(\varphi \circ \psi)(P) = \varphi(\psi(P))$. It always contains $\mathbb{Z}$ (via the multiplication maps $[n]$), and the question is how much larger it can be. Over $\mathbb{C}$, we shall see in theorem 3.5.5 that $\text{End}(E)$ is either $\mathbb{Z}$ (generic) or an order in an imaginary quadratic field (the CM case). Over fields of characteristic p , a third possibility appears — an order in a *quaternion algebra* — corresponding to supersingular

curves. In this section we give the complete algebraic classification, valid in all characteristics, using the positive definite degree form from proposition 2.2.9 and the rank bound provided by the Tate module.

For completeness, we box the fact that $\text{End}(Z)$ is a ring with a natural faithful $\mathbb{Z}$ -action:

Proposition 2.6.1: $\text{End}(E)$ is a Ring

Let E/K be an elliptic curve. Then $\text{End}(E)$ is a (possibly non-commutative) ring with the following properties:

1. (*Addition.*) $(\varphi + \psi)(P) = \varphi(P) + \psi(P)$ makes $\text{End}(E)$ an abelian group.
2. (*Multiplication.*) $\varphi \circ \psi$ is the ring product.
3. (*Identity.*) $[1] = \text{id}_E$ is the multiplicative identity.
4. (*Characteristic map.*) The map $\mathbb{Z} \rightarrow \text{End}(E)$, $n \mapsto [n]$, is an injective ring homomorphism.

Proof:

exercise for those with insomnia.

Recall from corollary 2.2.10 that $\text{Hom}(E_1, E_2)$ and $\text{End}(E)$ are both free $\mathbb{Z}$ -modules. Using the Tate-module, we can bound the ranks:

Proposition 2.6.2: Finiteness of Hom

Let E_1, E_2 be elliptic curves over $\overline{K}$. Then

$$\text{rank}_{\mathbb{Z}} \text{Hom}(E_1, E_2) \leq 4 \quad \text{and} \quad \text{rank}_{\mathbb{Z}} \text{End}(E) \leq 4$$

Proof:

As established in corollary 2.5.4, for any prime $\ell \neq \text{char} K$, the ℓ -adic Tate module provides a natural injection

$$\text{Hom}(E_1, E_2) \otimes_{\mathbb{Z}} \mathbb{Z}_{\ell} \hookrightarrow \text{Hom}_{\mathbb{Z}_{\ell}}(T_{\ell}(E_1), T_{\ell}(E_2)) \cong M_2(\mathbb{Z}_{\ell}).$$

Since $M_2(\mathbb{Z}_{\ell})$ is a free $\mathbb{Z}_{\ell}$ -module of rank 4, the $\mathbb{Z}$ -rank of $\text{Hom}(E_1, E_2)$ can be at most 4. By taking $E_1 = E_2$, we get the same result for $\text{End}(E)$.

For $\text{End}(E)$, we can make the bound sharp:

Theorem 2.6.3: Classification of Endomorphism Rings

Let E be an elliptic curve over $\overline{K}$. Then $\text{End}(E)$ is one of the following:

- I. $\text{End}(E) \cong \mathbb{Z}$. (*The generic case.*)
- II. $\text{End}(E)$ is an order $\mathcal{O}$ in an imaginary quadratic field K . (*The CM case.*)
- III. $\text{End}(E)$ is an order in a definite quaternion algebra B over $\mathbb{Q}$. (*The supersingular case; only occurs when $\text{char} K = p > 0$.*)

These three cases are mutually exclusive and exhaustive. In cases II and III, we say E has *complex multiplication* (CM).

The proof occupies the remainder of this subsection. The strategy is:

1. Show $R := \text{End}(E)$ is an integral domain of characteristic 0.
2. Show $D := R \otimes_{\mathbb{Z}} \mathbb{Q}$ is a division algebra of finite dimension over $\mathbb{Q}$.
3. Use the degree form to show D admits a positive-definite “norm form.”
4. Classify the possibilities for D using Hurwitz’s Theorem (see [1]).

Proof:

Step 1: $\text{End}(E)$ is an integral domain. If $\varphi \circ \psi = 0$ with $\varphi, \psi \in \text{End}(E)$, then $\deg(\varphi) \deg(\psi) = \deg(\varphi \circ \psi) = 0$. Since the degree of a non-zero isogeny is positive, either $\varphi = 0$ or $\psi = 0$. The characteristic of $\text{End}(E)$ is 0: $[n] \neq 0$ for $n \neq 0$ (since $\deg([n]) = n^2 \neq 0$).

Step 2: $D = \text{End}(E) \otimes \mathbb{Q}$ is a finite-dimensional division algebra over $\mathbb{Q}$. Since $\text{End}(E)$ is a torsion-free $\mathbb{Z}$ -module of finite rank (proposition 2.6.2), $D = \text{End}(E) \otimes_{\mathbb{Z}} \mathbb{Q}$ is a finite-dimensional $\mathbb{Q}$ -vector space. It is a domain (inherited from $\text{End}(E)$). A finite-dimensional domain over a field is always a division algebra: if $\alpha \neq 0$ in D , the map $x \mapsto \alpha x$ is an injective $\mathbb{Q}$ -linear map on a finite-dimensional space, hence surjective, so α has an inverse.

Step 3: The degree extends to a positive definite norm form on D . The degree map $\deg : \text{End}(E) \rightarrow \mathbb{Z}_{\geq 0}$ is a quadratic form (proposition 2.2.9) that is positive definite on $\text{End}(E) \setminus \{0\}$, and it satisfies $\deg(\varphi \circ \psi) = \deg(\varphi) \cdot \deg(\psi)$ (multiplicativity, proposition 2.1.7). Extending to D by $\deg(\varphi/n) = \deg(\varphi)/n^2$, we get a $\mathbb{Q}$ -valued positive definite quadratic form $N : D \rightarrow \mathbb{Q}$ satisfying $N(\alpha\beta) = N(\alpha)N(\beta)$: a *norm form* on the division algebra D .

The associated bilinear form $\langle \alpha, \beta \rangle = N(\alpha + \beta) - N(\alpha) - N(\beta)$ extends to a positive definite bilinear form on $D \otimes_{\mathbb{Q}} \mathbb{R}$, making $D \otimes \mathbb{R}$ a real inner product space.

Step 4: Classify the possibilities. A finite-dimensional division algebra D over $\mathbb{Q}$ equipped with a positive definite norm form N satisfying $N(\alpha\beta) = N(\alpha)N(\beta)$ is severely constrained. We consider three cases based on $\dim_{\mathbb{Q}} D$.

Case $\dim_{\mathbb{Q}} D = 1$: $D = \mathbb{Q}$, and $\text{End}(E) = \mathbb{Z}$. This is case I.

Case $\dim_{\mathbb{Q}} D = 2$: D is a quadratic field extension $K/\mathbb{Q}$. The norm form N must be positive definite on $K \otimes_{\mathbb{Q}} \mathbb{R}$. If $K = \mathbb{Q}(\sqrt{d})$ for squarefree d , then $K \otimes \mathbb{R} = \mathbb{R}(\sqrt{d})$, which is $\mathbb{R} \times \mathbb{R}$ (if $d > 0$) or $\mathbb{C}$ (if $d < 0$). The norm form on $\mathbb{Q}(\sqrt{d})$ is $N(a + b\sqrt{d}) = a^2 - db^2$ (the field norm $N_{K/\mathbb{Q}}$). For this to be positive definite, we need $a^2 - db^2 > 0$ for all $(a, b) \neq (0, 0)$, which requires $d < 0$ (if $d > 0$, take $a = 0, b = 1$ to get $-d < 0$). So K is an *imaginary quadratic* field.

That the norm form $N = \deg$ agrees with the field norm $N_{K/\mathbb{Q}}$ follows from the multiplicativity $N(\alpha\beta) = N(\alpha)N(\beta)$ and the fact that both are quadratic forms on a 2-dimensional $\mathbb{Q}$ -vector space that are multiplicative: such a form is uniquely determined (up to scaling) by the algebra structure, and the scaling is fixed by $N([1]) = \deg([1]) = 1$.

In this case $\text{End}(E) \subseteq \mathcal{O}_K$ is an order in K (a subring that is a free $\mathbb{Z}$ -module of rank 2), and we are in case II.

Case $\dim_{\mathbb{Q}} D = 4$: By the classification of finite-dimensional division algebras over $\mathbb{Q}$, D is a *quaternion algebra* over $\mathbb{Q}$: $D = \mathbb{Q} + \mathbb{Q}\alpha + \mathbb{Q}\beta + \mathbb{Q}\alpha\beta$ with $\alpha^2, \beta^2 \in \mathbb{Q}^*$ and $\alpha\beta = -\beta\alpha$. The reduced norm $\text{Nrd} : D \rightarrow \mathbb{Q}$ is a multiplicative quadratic form of rank 4. As in case II, the positive-definiteness of N forces D to be *definite*: $D \otimes_{\mathbb{Q}} \mathbb{R} \cong \mathbb{H}$ (the Hamilton quaternions), not $M_2(\mathbb{R})$.

We claim this case occurs only in characteristic p . Indeed, if $\text{char} K = 0$, then by the analytic theory (theorem 3.5.5), $\text{End}(E_{\bar{K}})$ is either $\mathbb{Z}$ or an order in an imaginary quadratic field — the quaternion case cannot occur over $\bar{\mathbb{Q}}$ (or any algebraically closed field of characteristic 0). The underlying reason is that a quaternion algebra over $\mathbb{Q}$ cannot embed in $\mathbb{C}$: such an embedding would make $D \otimes \mathbb{R}$ a subalgebra of $\mathbb{C}$, but $\dim_{\mathbb{R}}(D \otimes \mathbb{R}) = 4 > 2 = \dim_{\mathbb{R}} \mathbb{C}$.

In characteristic p , supersingular curves realize case III. The Frobenius endomorphism φ_p generates an imaginary quadratic field $\mathbb{Q}(\varphi_p)$ inside D , but for a supersingular curve, one can construct additional endomorphisms that do not commute with φ_p . This forces the dimension of D to be strictly greater than 2, leaving the quaternion algebra (dimension 4) as the only possibility. We will explore the explicit structure of this quaternion algebra in ??.

Why other dimensions cannot occur: We have shown that $D \otimes_{\mathbb{Q}} \mathbb{R}$ is a finite-dimensional associative $\mathbb{R}$ -algebra equipped with a positive-definite multiplicative quadratic form. By Hurwitz's Theorem [1], the only such algebras over $\mathbb{R}$ are $\mathbb{R}$ (dimension 1), $\mathbb{C}$ (dimension 2), and the Hamilton quaternions $\mathbb{H}$ (dimension 4). (The octonions $\mathbb{O}$ also admit such a form, but they are non-associative, whereas $\text{End}(E)$ is associative under composition). Therefore, $\dim_{\mathbb{Q}} D = \dim_{\mathbb{R}}(D \otimes_{\mathbb{Q}} \mathbb{R})$ must be 1, 2, or 4.

The Three Cases in Detail

Case I: $\text{End}(E) = \mathbb{Z}$. This is the “generic” case: for a “random” elliptic curve over $\overline{K}$, the only endomorphisms are the multiplication maps $[n]$. Over $\overline{\mathbb{Q}}$ (or any algebraically closed field of characteristic 0), this holds for all but countably many j -invariants (those corresponding to CM curves). Over $\overline{\mathbb{F}_p}$, this holds for ordinary curves that are not CM.

Example 2.6.4: A Curve with $\text{End} = \mathbb{Z}$

The curve $E : y^2 = x^3 + x + 1$ over $\overline{\mathbb{Q}}$ has j -invariant $j(E) = 1728 \frac{4(1)^3}{4(1)^3 + 27(1)^2} = \frac{6912}{31}$. Because CM j -invariants are always algebraic integers, and $6912/31$ is strictly rational (not an integer), E cannot have complex multiplication. So $\text{End}(E) = \mathbb{Z}$. One can verify computationally that no extra endomorphisms exist by checking that the division polynomials ψ_n for small n have no “unexpected” factorizations.

Case II: $\text{End}(E)$ is an order in an imaginary quadratic field. The curve E is said to have *complex multiplication* (CM) by the order $\mathcal{O} = \text{End}(E)$. The imaginary quadratic field $K = \mathcal{O} \otimes \mathbb{Q}$ is called the *CM field* of E .

Example 2.6.5: CM by the Gaussian Integers

For $E : y^2 = x^3 - x$ over $\overline{\mathbb{Q}}$, the map $\iota : (x, y) \mapsto (-x, iy)$ is an automorphism of E satisfying $\iota^2 = [-1]$ (the negation map). Thus $\iota^2 + [1] = [0]$, and ι satisfies the minimal polynomial $T^2 + 1$ over $\mathbb{Z}$. The endomorphism ring is $\text{End}(E) = \mathbb{Z}[\iota] \cong \mathbb{Z}[i]$, the ring of Gaussian integers, which is the maximal order in $K = \mathbb{Q}(i) = \mathbb{Q}(\sqrt{-1})$.

The degree form is $\deg(a + b\iota) = a^2 + b^2 = |a + bi|^2$: the ordinary absolute value squared on $\mathbb{Z}[i]$. This matches the analytic computation of example 3.5.6(a).

Example 2.6.6: CM by the Eisenstein Integers

For $E : y^2 = x^3 + 1$ over $\overline{\mathbb{Q}}$, let $\omega = e^{2\pi i/3} = (-1 + \sqrt{-3})/2$. The map $\psi : (x, y) \mapsto (\omega x, y)$ is an endomorphism of E with $\psi^3 = [1]$ (since $\omega^3 = 1$ applied to x , and y is unchanged). More precisely, $\psi^2 + \psi + [1] = [0]$, so ψ satisfies $T^2 + T + 1$. The endomorphism ring is $\text{End}(E) = \mathbb{Z}[\psi] \cong \mathbb{Z}[\omega]$, the ring of Eisenstein integers, which is the maximal order in $K = \mathbb{Q}(\omega) = \mathbb{Q}(\sqrt{-3})$.

Case III: $\text{End}(E)$ is an order in a quaternion algebra. This occurs only for supersingular curves in characteristic $p > 0$. The precise statement is:

Theorem 2.6.7: Endomorphism Ring of Supersingular Curves

Let $E/\overline{\mathbb{F}_p}$ be a supersingular elliptic curve. Then $\text{End}(E) \otimes \mathbb{Q}$ is a definite quaternion algebra $B_{p,\infty}$ over $\mathbb{Q}$, ramified at p and ∞ (and unramified at all other places). The endomorphism ring $\text{End}(E)$ is a maximal order in $B_{p,\infty}$.

Proof:

From the classification (theorem 2.6.3), $\text{End}(E) \otimes \mathbb{Q}$ is a division algebra D of dimension 1, 2, or 4 over $\mathbb{Q}$. We show $\dim_{\mathbb{Q}} D = 4$ by exhibiting “enough” endomorphisms.

Since E is supersingular, $E[p](\overline{\mathbb{F}}_p) = \{O\}$ (theorem 2.4.4). We may assume up to isogeny that E is defined over $\mathbb{F}_p$ and has trace $a_p = 0$. The Frobenius endomorphism φ_p then satisfies $\varphi_p^2 + [p] = [0]$. The subring $\mathbb{Z}[\varphi_p] \subseteq \text{End}(E)$ is isomorphic to the imaginary quadratic order $\mathbb{Z}[\sqrt{-p}]$, which has rank 2 over $\mathbb{Z}$.

To show $\dim_{\mathbb{Q}} D = 4$ (not 2), one must show $\text{End}(E)$ is strictly larger than $\mathbb{Z}[\varphi_p]$: there exist endomorphisms not in $\mathbb{Q}(\varphi_p)$. This is the deepest part of the theorem. The proof uses the theory of the formal group (chapter 4) and Dieudonné modules, or alternatively the Deuring lifting theorem, which we will develop in chapter 5. The key idea is that over $\overline{\mathbb{F}}_p$, the group scheme $E[p^\infty]$ (the p -divisible group of E) is isomorphic to the unique 1-dimensional formal group of height 2 over $\overline{\mathbb{F}}_p$, whose endomorphism ring is the maximal order in the quaternion algebra ramified at p and ∞ .

That B is ramified at p (i.e., $B \otimes \mathbb{Q}_p$ is a division algebra, not $M_2(\mathbb{Q}_p)$) follows from the fact that $E[p](\overline{\mathbb{F}}_p) = \{O\}$: if $B \otimes \mathbb{Q}_p \cong M_2(\mathbb{Q}_p)$, then $\text{End}(E) \otimes \mathbb{Z}_p$ would contain idempotents producing a nontrivial p -torsion point, contradicting supersingularity. That B is ramified at ∞ follows from the positive-definiteness of the degree form (a quaternion algebra carrying a positive definite norm form is definite, hence ramified at ∞). By the classification of quaternion algebras over $\mathbb{Q}$ (they are ramified at an even number of places), $B = B_{p,\infty}$ is uniquely determined.

Example 2.6.8: A Supersingular Endomorphism Ring

For $E : y^2 = x^3 + 1$ over $\overline{\mathbb{F}}_2$ (which is supersingular: $a_2 = 0$, see example 2.4.5(b)), the endomorphism ring $\text{End}(E)$ is a maximal order in the quaternion algebra $B_{2,\infty}$ ramified at 2 and ∞ . Concretely, $B_{2,\infty}$ can be presented as $\mathbb{Q}\langle i, j \rangle$ with $i^2 = -1$, $j^2 = -2$, $ij = -ji$. The maximal order is $\mathbb{Z}[1, i, j, (1 + ij)/2]$. Identifying these abstract generators with specific endomorphisms of E requires the Deuring correspondence, which we develop in chapter 5.

Summary and Comparison with the Analytic Theory

The algebraic classification (theorem 2.6.3) refines the analytic classification (theorem 3.5.5) in two ways. First, it works over any field, not just $\mathbb{C}$. Second, it reveals the supersingular case (III), which is invisible over $\mathbb{C}$ because a quaternion algebra over $\mathbb{Q}$ cannot embed in $\mathbb{C}$ (as observed in Step 4 of the proof).

The hierarchy of endomorphism rings reflects a hierarchy of “symmetry”:

Case	$\text{rank}_{\mathbb{Z}} \text{End}(E)$	$\dim_{\mathbb{Q}} D$	Occurs when
I: $\text{End} = \mathbb{Z}$	1	1	Generic (all characteristics)
II: $\text{End} = \mathcal{O} \subseteq K$	2	2	CM (all characteristics)
III: $\text{End} = \mathcal{O} \subseteq B$	4	4	Supersingular ($\text{char} K = p$ only)

This trichotomy governs the entire arithmetic of the curve: the Galois representation on $T_\ell(E)$ (to be constructed in the next section) is “as large as possible” (GL_2 -type) in case I, factors through a Cartan subgroup in case II, and is “as small as possible” (scalar on the p -adic side) in case III. The deeper one goes into the theory, the more this classification determines.

Exercise 2.6.1

1. Show directly that $\text{End}(E)$ has no zero divisors. (*Hint*: if $\varphi \circ \psi = 0$ with $\varphi \neq 0$, apply $\deg$.)
2. Let R be a finitely generated torsion-free $\mathbb{Z}$ -algebra with no zero divisors. Show that $D = R \otimes_{\mathbb{Z}} \mathbb{Q}$ is a division algebra over $\mathbb{Q}$.
3. For $E : y^2 = x^3 - x$ over $\overline{\mathbb{Q}}$, verify that the endomorphism $\iota : (x, y) \mapsto (-x, iy)$ satisfies $\deg(\iota) = 1$ and $\iota^2 = [-1]$. Conclude $\text{End}(E) \supseteq \mathbb{Z}[i]$, and explain why $\text{End}(E) = \mathbb{Z}[i]$ (not strictly larger).
4. Show that if D is a division algebra over $\mathbb{Q}$ carrying a positive definite quadratic form $N : D \rightarrow \mathbb{Q}$ with $N(\alpha\beta) = N(\alpha)N(\beta)$, then $\dim_{\mathbb{Q}} D \in \{1, 2, 4\}$. (*Hint*: use the Frobenius theorem on real division algebras: the only finite-dimensional associative division algebras over $\mathbb{R}$ are $\mathbb{R}, \mathbb{C}, \mathbb{H}$.)
5. Let $E : y^2 = x^3 - x$ over $\mathbb{Q}$, so $\text{End}(E_{\overline{\mathbb{Q}}}) = \mathbb{Z}[i]$. Reducing modulo 5: the reduced curve $\tilde{E} : y^2 = x^3 - x$ over $\mathbb{F}_5$ has $|\tilde{E}(\mathbb{F}_5)| = 8$ and is ordinary (example 2.4.5(a)). What is $\text{End}(\tilde{E}_{\overline{\mathbb{F}_5}})$? (*Hint*: the endomorphism ι reduces to an endomorphism of $\tilde{E}$; does End grow or stay the same?)
6. Show that the number of supersingular j -invariants over $\overline{\mathbb{F}_p}$ is $\lfloor p/12 \rfloor + \epsilon$, where $\epsilon \in \{0, 1, 2\}$ depends on $p \pmod{12}$. (*Hint*: use the mass formula: $\sum_{j \text{ s.s.}} 1/|\text{Aut}(E_j)| = (p-1)/24$, where the sum is over supersingular j -invariants and $|\text{Aut}(E_j)|$ accounts for extra automorphisms at $j = 0$ and $j = 1728$.)

The Analytic Theory

The preceding chapters built the algebraic theory of elliptic curves and isogenies over an arbitrary field. Over the complex numbers, every elliptic curve $E/\mathbb{C}$ is analytically isomorphic to a complex torus $\mathbb{C}/\Lambda$ for a lattice $\Lambda \subset \mathbb{C}$, and conversely every such torus carries the structure of an elliptic curve. This *uniformization* is the structural result in the classical theory. It converts the algebraic questions from chapter 2 — the group law, isogenies, torsion, and the endomorphism ring — into questions about lattices, which are discrete subgroups of $\mathbb{C}$ and therefore amenable to direct computation.

Many of the algebraic constructions that required algebraic proofs in the previous chapter become transparent analytically: an isogeny is simply the multiplication-by- α map for a complex number α satisfying $\alpha\Lambda_1 \subseteq \Lambda_2$; the n -torsion consists of the classes $\frac{1}{n}\Lambda/\Lambda$; and complex multiplication is the existence of an $\alpha \notin \mathbb{Z}$ with $\alpha\Lambda \subseteq \Lambda$.

The drawback of this clarity is that it is specific to $\mathbb{C}$: the uniformization does not exist over $\mathbb{Q}$ or over finite fields. Nevertheless, the analytic theory is indispensable, both for the intuition it provides and for the concrete results it yields. The Eisenstein series G_4 and G_6 , the modular discriminant $\Delta(\tau)$, and the j -function $j(\tau)$ that we construct in this chapter will reappear in the theory of complex multiplication (??) and modular forms (??). The reader who has studied complex analysis [4] has already encountered elliptic functions and Riemann surfaces; this chapter organizes that theory into a framework needed for arithmetic applications.

3.1 Lattices and Elliptic Functions

We begin by fixing the notation for lattices that will persist throughout this chapter and the rest of the book, and by recalling the basic structural facts about elliptic functions. The reader who has studied [4] will find this section largely familiar; its purpose is to establish conventions and to identify precisely the properties we will need.

Lattices in $\mathbb{C}$

Definition 3.1.1: Lattice

A *lattice* in $\mathbb{C}$ is a discrete subgroup $\Lambda \subset \mathbb{C}$ of rank 2 over $\mathbb{Z}$. Equivalently, $\Lambda = \mathbb{Z}\omega_1 + \mathbb{Z}\omega_2$ for two complex numbers ω_1, ω_2 that are linearly independent over $\mathbb{R}$ (i.e., $\omega_1/\omega_2 \notin \mathbb{R}$).

The pair (ω_1, ω_2) is called a *basis* for Λ ; it is far from unique. The set of all bases for a given lattice Λ is related by the action of $\mathrm{SL}_2(\mathbb{Z})$: if (ω_1, ω_2) is a basis, then $(\omega'_1, \omega'_2) = (a\omega_1 + b\omega_2, c\omega_1 + d\omega_2)$ is also a basis if and only if $\begin{pmatrix} a & b \\ c & d \end{pmatrix} \in \mathrm{SL}_2(\mathbb{Z})$. We adopt the following normalization.

Definition 3.1.2: Normalized Basis and the Period Ratio

A basis (ω_1, ω_2) for a lattice Λ is *positively oriented* if $\mathrm{Im}(\omega_1/\omega_2) > 0$. The ratio $\tau = \omega_1/\omega_2$ then lies in the *upper half-plane*

$$\mathbb{H} = \{\tau \in \mathbb{C} \mid \mathrm{Im}(\tau) > 0\}.$$

We call τ a *period ratio* for Λ .

Any lattice admits a positively oriented basis (interchange ω_1 and ω_2 if necessary). By rescaling Λ — replacing Λ by $\omega_2^{-1}\Lambda$ — we can always assume $\omega_2 = 1$, so that $\Lambda = \mathbb{Z}\tau + \mathbb{Z}$ for some $\tau \in \mathbb{H}$. We give this a name.

Definition 3.1.3: Standard Lattice

For $\tau \in \mathbb{H}$, the *standard lattice* is $\Lambda_\tau = \mathbb{Z}\tau + \mathbb{Z}$. Two lattices Λ and Λ' are *homothetic*, written $\Lambda \sim \Lambda'$, if $\Lambda' = \alpha\Lambda$ for some $\alpha \in \mathbb{C}^*$.

Proposition 3.1.4: Homothety Classes and the Upper Half-Plane

Every lattice is homothetic to Λ_τ for some $\tau \in \mathbb{H}$, and $\Lambda_\tau \sim \Lambda_{\tau'}$ if and only if $\tau' = \gamma \cdot \tau := \frac{a\tau+b}{c\tau+d}$ for some $\gamma = \begin{pmatrix} a & b \\ c & d \end{pmatrix} \in \mathrm{SL}_2(\mathbb{Z})$.

Proof:

The first claim is the rescaling argument above. For the second: $\Lambda_\tau = \mathbb{Z}\tau + \mathbb{Z}$ and $\Lambda_{\tau'} = \mathbb{Z}\tau' + \mathbb{Z}$ are homothetic iff $\Lambda_{\tau'} = \alpha\Lambda_\tau$ for some $\alpha \in \mathbb{C}^*$, which means $\tau' = \alpha(a\tau+b)$ and $1 = \alpha(c\tau+d)$ for integers a, b, c, d with $ad - bc = 1$ (the determinant condition ensures $(\tau', 1)$ and $(\tau, 1)$ generate the same lattice up to scaling). From the second equation, $\alpha = (c\tau+d)^{-1}$, and substituting into the first gives $\tau' = (a\tau+b)/(c\tau+d)$.

Thus the set of homothety classes of lattices in $\mathbb{C}$ is identified with the quotient $\mathrm{SL}_2(\mathbb{Z}) \backslash \mathbb{H}$. This is the space that parametrizes elliptic curves over $\mathbb{C}$ up to isomorphism — the first incarnation of the *moduli space* of elliptic curves. We will study this quotient in detail in section 3.6.

Complex Tori

Given a lattice Λ , the quotient $\mathbb{C}/\Lambda$ inherits from $\mathbb{C}$ the structure of a compact Riemann surface. The projection $\pi : \mathbb{C} \rightarrow \mathbb{C}/\Lambda$ is a holomorphic covering map with deck transformation group Λ , acting by $z \mapsto z + \omega$ for $\omega \in \Lambda$. The topology of $\mathbb{C}/\Lambda$ is that of a torus: a fundamental domain for Λ in $\mathbb{C}$ is the parallelogram

$$\mathcal{P} = \{s\omega_1 + t\omega_2 \mid 0 \leq s, t < 1\},$$

and $\mathbb{C}/\Lambda$ is obtained by identifying opposite edges.

Proposition 3.1.5: Genus of a Complex Torus

The Riemann surface $\mathbb{C}/\Lambda$ is compact, connected, and of genus 1.

Proof:

Compactness: $\mathbb{C}/\Lambda$ is the continuous image of the closed parallelogram $\overline{\mathcal{P}}$, which is compact. Connectedness: $\mathbb{C}$ is connected and π is surjective. For the genus, compute the Euler characteristic: $\mathbb{C}/\Lambda$ is topologically a torus $S^1 \times S^1$, so $\chi(\mathbb{C}/\Lambda) = 0$, giving $g = 1$ from $\chi = 2 - 2g$.

Alternatively, the holomorphic 1-form dz on $\mathbb{C}$ is Λ -invariant and thus descends to a nowhere-vanishing holomorphic 1-form on $\mathbb{C}/\Lambda$. A compact Riemann surface of genus g has exactly g linearly independent holomorphic 1-forms (by Riemann–Roch or Hodge theory); here the space is spanned by the single form dz , confirming $g = 1$.

Since $\mathbb{C}/\Lambda$ is a compact Riemann surface of genus 1 and has the distinguished point $0 + \Lambda$, it is (analytically) an elliptic curve. Moreover, it is a *group*: the addition in $\mathbb{C}$ descends to a well-defined, holomorphic group operation on $\mathbb{C}/\Lambda$, with identity $0 + \Lambda$ and inverse $-z + \Lambda$. This makes $\mathbb{C}/\Lambda$ simultaneously a Riemann surface and a complex Lie group.

Elliptic Functions

The meromorphic functions on $\mathbb{C}/\Lambda$ are precisely the Λ -periodic meromorphic functions on $\mathbb{C}$: these are the *elliptic functions* with respect to Λ .

Definition 3.1.6: Elliptic Function

An *elliptic function* with respect to a lattice Λ is a meromorphic function $f : \mathbb{C} \rightarrow \mathbb{C} \cup \{\infty\}$ satisfying $f(z + \omega) = f(z)$ for all $z \in \mathbb{C}$ and all $\omega \in \Lambda$. The set of all elliptic functions for Λ forms a field, denoted $\mathbb{C}(\Lambda)$.

We recall the fundamental structural results about elliptic functions. All of these are proved in [4]; we collect them here for convenient reference.

Proposition 3.1.7: Basic Properties of Elliptic Functions

Let $f \in \mathbb{C}(\Lambda)$ be a nonconstant elliptic function.

1. (*Liouville.*) An elliptic function with no poles is constant. Equivalently, a holomorphic elliptic function is constant.
2. (*Residue theorem.*) The sum of the residues of f in any fundamental parallelogram is zero: $\sum_P \text{Res}_P(f) = 0$.
3. (*Valence.*) The number of zeros of f in a fundamental parallelogram (counted with multiplicity) equals the number of poles (counted with multiplicity). This common value is the *order* of f , denoted $\text{ord}(f)$.
4. (*Minimum order.*) $\text{ord}(f) \geq 2$. In other words, there is no elliptic function with a single simple pole in a fundamental parallelogram. (A simple pole would violate the residue condition.)
5. (*Zeros minus poles.*) If $a_1, \dots, a_n$ are the zeros of f and $b_1, \dots, b_n$ are the poles (in a fundamental parallelogram, counted with multiplicity), then $\sum a_i - \sum b_i \in \Lambda$.

Proof:

See [4]. Briefly: (1) follows from the maximum modulus principle on the compact surface $\mathbb{C}/\Lambda$. (2)–(5) follow from integrating f'/f , zf'/f , and f itself around the boundary of a fundamental parallelogram, using the periodicity of f to cancel opposite edges.

Property (4) is the reason the $\wp$ -function, which we construct in the next section, has a pole of order 2 rather than 1: there is simply no elliptic function of order 1.

Property (5) has a clean interpretation in terms of divisors. For an elliptic function f of order n , the divisor $\text{div}(f) = \sum(a_i) - \sum(b_i)$ is a principal divisor on $\mathbb{C}/\Lambda$, and the condition $\sum a_i \equiv \sum b_i \pmod{\Lambda}$ is the “Abel–Jacobi” constraint on the complex torus. More precisely: on a genus-1 curve, a degree-zero divisor $D = \sum n_i(P_i)$ is principal if and only if $\deg(D) = 0$ and the “Abel sum” $\sum n_i P_i = 0$ in the group law of $\mathbb{C}/\Lambda$. This is the analytic incarnation of the theorem that $\text{Pic}^0(E) \cong E$ from theorem 1.3.1.

The Field of Elliptic Functions

The field $\mathbb{C}(\Lambda)$ is much smaller than the field $\mathbb{C}(z)$ of all meromorphic functions on $\mathbb{C}$. Its structure is determined by a single function of order 2 — the Weierstrass $\wp$ -function, which we construct in the next section.

Proposition 3.1.8: Structure of the Function Field

The field $\mathbb{C}(\Lambda)$ of elliptic functions for Λ is generated over $\mathbb{C}$ by any elliptic function of order 2 and its derivative. Concretely, if $\wp$ is an even elliptic function of order 2, then:

1. Every even elliptic function is a rational function of $\wp$: $\mathbb{C}(\Lambda)^{\text{even}} = \mathbb{C}(\wp)$.
2. Every elliptic function f can be written uniquely as $f = R(\wp) + S(\wp) \cdot \wp'$ for rational functions $R, S \in \mathbb{C}(\wp)$.
3. $\mathbb{C}(\Lambda) = \mathbb{C}(\wp, \wp')$, and $[\mathbb{C}(\Lambda) : \mathbb{C}(\wp)] = 2$.

Proof Key ideas:

(1) If f is even and of order $2n$, then f has n zeros $\{a_1, \dots, a_n\}$ modulo Λ (where each a_i accounts for both a_i and $-a_i$, except for the half-lattice points which are their own negatives). The function $\prod_{i=1}^n (\wp(z) - \wp(a_i))$ has the same zeros as f (with the same multiplicities, by a careful case analysis at the half-lattice points) and poles only at the lattice points. Since f is also polar only at the lattice points (its poles occur in $\pm$ -pairs by evenness, hence at the lattice points modulo Λ), the ratio $f / \prod (\wp - \wp(a_i))$ is an elliptic function with no poles, hence constant by Liouville.

(2) Any elliptic function f decomposes as $f = f_{\text{even}} + f_{\text{odd}}$ where $f_{\text{even}}(z) = (f(z) + f(-z))/2$ and $f_{\text{odd}}(z) = (f(z) - f(-z))/2$. The even part is a rational function of $\wp$ by (1). The odd part satisfies $f_{\text{odd}}/\wp'$ is even (since $\wp'$ is odd and both are elliptic), hence is a rational function of $\wp$ by (1). Thus $f = R(\wp) + S(\wp) \cdot \wp'$.

(3) follows from (2): $\mathbb{C}(\Lambda) = \mathbb{C}(\wp)[\wp']$, and $\wp'$ satisfies a quadratic relation over $\mathbb{C}(\wp)$ (namely, $(\wp')^2 = 4\wp^3 - g_2\wp - g_3$, as we shall prove in section 3.2), so $[\mathbb{C}(\Lambda) : \mathbb{C}(\wp)] = 2$.

The field extension $\mathbb{C}(\Lambda)/\mathbb{C}(\wp)$ is a degree-2 extension generated by $\wp'$, with $(\wp')^2 \in \mathbb{C}(\wp)$. Geometrically, this reflects the degree-2 map $\mathbb{C}/\Lambda \rightarrow \mathbb{P}^1$ given by $z \mapsto \wp(z)$, which is a ramified double cover — the hyperelliptic structure of the elliptic curve. The Galois group of this cover is $\{1, [-1]\}$, acting by $z \mapsto -z$, and $\wp'$ is the function that distinguishes the two sheets (it changes sign under $z \mapsto -z$).

This result tells us that once we have the $\wp$ -function in hand, we know the entire function field of $\mathbb{C}/\Lambda$, and hence (by the equivalence between smooth projective curves and their function fields) we know $\mathbb{C}/\Lambda$ as an algebraic curve. The next section constructs $\wp$ and derives the algebraic relation it satisfies.

Exercise 3.1.1

1. Let $\Lambda = \mathbb{Z}\omega_1 + \mathbb{Z}\omega_2$ be a lattice with positively oriented basis (ω_1, ω_2) .
 - (a) Show that $\gamma = \begin{pmatrix} a & b \\ c & d \end{pmatrix} \in \text{GL}_2(\mathbb{Z})$ sends (ω_1, ω_2) to another basis of Λ if and only if $\det(\gamma) = \pm 1$, and to a positively oriented basis if and only if $\det(\gamma) = +1$.
 - (b) Conclude that the set of positively oriented bases of Λ is a principal homogeneous space (torsor) for $\text{SL}_2(\mathbb{Z})$.

2. Show that the area of the fundamental parallelogram for $\Lambda_\tau = \mathbb{Z}\tau + \mathbb{Z}$ is $\text{Im}(\tau)$. More generally, for $\Lambda = \mathbb{Z}\omega_1 + \mathbb{Z}\omega_2$, the area is $|\text{Im}(\omega_1\overline{\omega_2})|$. Deduce that homothetic lattices have proportional areas: $\text{area}(\mathbb{C}/\alpha\Lambda) = |\alpha|^2 \cdot \text{area}(\mathbb{C}/\Lambda)$.
3. Determine which of the following lattices are homothetic:
 - (a) $\Lambda_1 = \mathbb{Z} + \mathbb{Z}i$ and $\Lambda_2 = \mathbb{Z}(1+i) + \mathbb{Z}(1-i)$.
 - (b) $\Lambda_1 = \mathbb{Z} + \mathbb{Z}i$ and $\Lambda_3 = \mathbb{Z} + \mathbb{Z}(2i)$.
 - (c) $\Lambda_4 = \mathbb{Z} + \mathbb{Z}\omega$ (where $\omega = e^{2\pi i/3}$) and $\Lambda_5 = \mathbb{Z} + \mathbb{Z}(2\omega + 1)$.
4. Let $\Lambda = \mathbb{Z} + \mathbb{Z}\tau$ and define $f(z) = \wp(z) - \wp(z + \tau/3)$, where $\wp$ is the Weierstrass $\wp$ -function for Λ (to be constructed in section 3.2).
 - (a) Show that f is an elliptic function for Λ .
 - (b) Determine the order of f , and find its zeros and poles in a fundamental parallelogram.

(Hint: $\wp(z)$ has a unique double pole at $z = 0$ in $\mathbb{C}/\Lambda$.)
5. (The Abel–Jacobi constraint.) Let f be an elliptic function of order 3 for the lattice Λ , with simple zeros at a_1, a_2, a_3 and a triple pole at $z = 0$ (all modulo Λ). Show that $a_1 + a_2 + a_3 \in \Lambda$. (This is the analytic counterpart of the collinearity result exercise 1.3.1 (5): three points on E sum to O in the group law iff they are collinear.)
6. Let Λ be a lattice and let $f \in \mathbb{C}(\Lambda)$ be an elliptic function of order n .
 - (a) Show that $f : \mathbb{C}/\Lambda \rightarrow \mathbb{P}^1$ is a holomorphic map of degree n .
 - (b) Conclude that $[\mathbb{C}(\Lambda) : \mathbb{C}(f)] = n$.

(Hint for (b): use the fact that for a non-constant morphism $\varphi : C_1 \rightarrow C_2$ of smooth projective curves, the degree of φ equals the degree of the field extension $[K(C_1) : \varphi^*K(C_2)]$.)

3.2 The Weierstrass $\wp$ -Function

We now construct the function that generates the entire field of elliptic functions: the Weierstrass $\wp$ -function. It is the simplest nonconstant elliptic function (having order 2, the minimum possible by proposition 3.1.7(4)), and its Laurent expansion at the origin produces the Eisenstein series $G_{2k}(\Lambda)$ as coefficients. Most importantly, $\wp$ and its derivative $\wp'$ satisfy an algebraic relation that is precisely a Weierstrass equation — thereby connecting the analytic theory to the algebraic objects of chapter 1.

Construction and Convergence

The idea is simple: we want an even elliptic function with a double pole at each lattice point and no other poles. The most naive attempt — $\sum_{\omega \in \Lambda} (z - \omega)^{-2}$ — does not converge (the terms decay too slowly). The trick is to subtract the “constant part” of each term to improve convergence.

Definition 3.2.1: The Weierstrass $\wp$ -Function

Let Λ be a lattice in $\mathbb{C}$. The *Weierstrass $\wp$ -function* associated to Λ is

$$\wp(z; \Lambda) = \frac{1}{z^2} + \sum_{\omega \in \Lambda \setminus \{0\}} \left(\frac{1}{(z - \omega)^2} - \frac{1}{\omega^2} \right). \quad (3.1)$$

When the lattice is understood, we write simply $\wp(z)$.

The subtraction of $1/\omega^2$ is the key: it ensures that each summand is $O(|\omega|^{-3})$ for $|z|$ bounded and $|\omega| \rightarrow \infty$, making the series converge.

Proposition 3.2.2: Convergence and Basic Properties of $\wp$

1. The series (3.1) converges absolutely and uniformly on every compact subset of $\mathbb{C} \setminus \Lambda$. In particular, $\wp(z)$ is meromorphic on $\mathbb{C}$ with double poles at each $\omega \in \Lambda$ and no other poles.
2. $\wp(z)$ is an even function: $\wp(-z) = \wp(z)$.
3. $\wp(z)$ is Λ -periodic: $\wp(z + \omega) = \wp(z)$ for all $\omega \in \Lambda$.

Thus $\wp \in \mathbb{C}(\Lambda)$ is an even elliptic function of order 2.

Proof:

(1) For $|z| \leq R$ and $|\omega| \geq 2R$, we estimate:

$$\frac{1}{(z - \omega)^2} - \frac{1}{\omega^2} = \frac{\omega^2 - (z - \omega)^2}{\omega^2(z - \omega)^2} = \frac{2\omega z - z^2}{\omega^2(z - \omega)^2}.$$

Since $|z - \omega| \geq |\omega| - |z| \geq |\omega|/2$, we have $|z - \omega|^2 \geq |\omega|^2/4$, so

$$\left| \frac{1}{(z - \omega)^2} - \frac{1}{\omega^2} \right| \leq \frac{|2\omega z - z^2|}{|\omega|^2 \cdot |\omega|^2/4} \leq \frac{2|\omega||z| + |z|^2}{|\omega|^4/4} \leq \frac{C_R}{|\omega|^3}$$

for a constant C_R depending only on R . The series $\sum_{\omega \neq 0} |\omega|^{-3}$ converges for any lattice Λ (the number of lattice points with $|\omega| \approx N$ is $O(N)$, so the sum behaves like $\sum N \cdot N^{-3} = \sum N^{-2} < \infty$). More precisely, the convergence of lattice sums $\sum_{\omega \neq 0} |\omega|^{-s}$ for $s > 2$ is a standard fact; see lemma 3.2.3 below.

(2) Replacing z by $-z$ in (3.1):

$$\wp(-z) = \frac{1}{z^2} + \sum_{\omega \neq 0} \left(\frac{1}{(-z - \omega)^2} - \frac{1}{\omega^2} \right) = \frac{1}{z^2} + \sum_{\omega \neq 0} \left(\frac{1}{(z + \omega)^2} - \frac{1}{\omega^2} \right).$$

As ω ranges over $\Lambda \setminus \{0\}$, so does $-\omega$, so the sum is unchanged: $\wp(-z) = \wp(z)$.

(3) Differentiating (3.1) term by term (justified by the uniform convergence on compact sets

away from Λ):

$$\wp'(z) = -2 \sum_{\omega \in \Lambda} \frac{1}{(z - \omega)^3}, \quad (3.2)$$

where the sum now includes $\omega = 0$. This series converges absolutely (the terms are $O(|\omega|^{-3})$), and its Λ -periodicity is immediate: replacing z by $z + \omega_0$ for $\omega_0 \in \Lambda$ just reindexes the sum. Therefore $\wp'(z + \omega_0) = \wp'(z)$ for all $\omega_0 \in \Lambda$.

Now, $\wp(z + \omega_0) - \wp(z)$ has derivative $\wp'(z + \omega_0) - \wp'(z) = 0$, so $\wp(z + \omega_0) - \wp(z) = c$ is a constant. Setting $z = -\omega_0/2$ (assuming $\omega_0/2 \notin \Lambda$, which holds for at least one generator ω_0): $\wp(\omega_0/2) - \wp(-\omega_0/2) = c$. Since $\wp$ is even, both sides are equal, so $c = 0$.

We record the convergence fact we used:

Lemma 3.2.3: Convergence of Lattice Sums

For a lattice $\Lambda \subset \mathbb{C}$ and a real number $s > 2$, the sum $\sum_{\omega \in \Lambda \setminus \{0\}} |\omega|^{-s}$ converges. For $s \leq 2$, it diverges.

Proof:

Let (ω_1, ω_2) be a basis for Λ and let $\delta = \min\{|\operatorname{Im}(\omega_1/\omega_2)|, 1\} \cdot |\omega_2| > 0$. Then $|m\omega_1 + n\omega_2| \geq \delta \cdot \max(|m|, |n|)$ for all $(m, n) \neq (0, 0)$ (a consequence of the $\mathbb{R}$ -linear independence of ω_1 and ω_2). The number of lattice points with $\max(|m|, |n|) = N$ is $O(N)$, so

$$\sum_{\omega \neq 0} |\omega|^{-s} \leq C \sum_{N=1}^{\infty} N \cdot N^{-s} = C \sum_{N=1}^{\infty} N^{1-s},$$

which converges for $s > 2$ and diverges for $s \leq 2$.

This lemma explains the threshold: $\sum \omega^{-2}$ diverges (hence the need for the correction in (3.1)), while $\sum \omega^{-k}$ converges for $k \geq 3$, which is why $\wp'$ has an absolutely convergent series without any correction.

The Laurent Expansion and Eisenstein Series

We now compute the Laurent expansion of $\wp(z)$ at $z = 0$. This expansion is the source of the Eisenstein series.

For $|z| < |\omega|$, we expand each summand in (3.1) using the geometric series:

$$\begin{aligned} \frac{1}{(z-\omega)^2} - \frac{1}{\omega^2} &= \frac{1}{\omega^2} \left(\frac{1}{(1-z/\omega)^2} - 1 \right) \\ &= \frac{1}{\omega^2} \left(\sum_{n=0}^{\infty} (n+1) \left(\frac{z}{\omega} \right)^n - 1 \right) \\ &= \frac{1}{\omega^2} \sum_{n=1}^{\infty} (n+1) \frac{z^n}{\omega^n} \\ &= \sum_{n=1}^{\infty} (n+1) \frac{z^n}{\omega^{n+2}}. \end{aligned}$$

In the second line, we used the identity $\frac{1}{(1-w)^2} = \sum_{n=0}^{\infty} (n+1)w^n$ for $|w| < 1$, which is obtained by differentiating the geometric series $\frac{1}{1-w} = \sum_{n=0}^{\infty} w^n$.

Summing over all $\omega \in \Lambda \setminus \{0\}$ and interchanging the order of summation (justified by absolute convergence for $|z| < \min_{\omega \neq 0} |\omega|$):

$$\begin{aligned} \wp(z) &= \frac{1}{z^2} + \sum_{\omega \neq 0} \sum_{n=1}^{\infty} (n+1) \frac{z^n}{\omega^{n+2}} \\ &= \frac{1}{z^2} + \sum_{n=1}^{\infty} (n+1) \left(\sum_{\omega \neq 0} \frac{1}{\omega^{n+2}} \right) z^n. \end{aligned} \tag{3.3}$$

Definition 3.2.4: Eisenstein Series

For an integer $k \geq 1$, the *Eisenstein series of weight $2k$* for the lattice Λ is

$$G_{2k}(\Lambda) = \sum_{\omega \in \Lambda \setminus \{0\}} \frac{1}{\omega^{2k}}.$$

By lemma 3.2.3, this converges absolutely for $k \geq 2$. (The case $k = 1$ requires special treatment; see remark 3.2.9.)

Since $\wp$ is even, all odd-power terms in its Laurent expansion must vanish. Indeed, the inner sum $\sum_{\omega \neq 0} \omega^{-(n+2)}$ vanishes when n is odd (because ω and $-\omega$ contribute terms that cancel). When $n = 2m$ is even, the inner sum is $G_{2m+2}(\Lambda)$. Substituting into (3.3):

Proposition 3.2.5: Laurent Expansion of $\wp$

For $0 < |z| < \min_{\omega \neq 0} |\omega|$,

$$\wp(z) = \frac{1}{z^2} + \sum_{m=1}^{\infty} (2m+1) G_{2m+2}(\Lambda) z^{2m}. \quad (3.4)$$

Explicitly, writing out the first few terms:

$$\wp(z) = \frac{1}{z^2} + 3 G_4 z^2 + 5 G_6 z^4 + 7 G_8 z^6 + 9 G_{10} z^8 + \cdots \quad (3.5)$$

The Eisenstein series G_4 and G_6 will play the starring roles: they determine the Weierstrass equation of the associated elliptic curve. The higher Eisenstein series $G_8, G_{10}, \dots$ are not independent — they can be expressed as polynomials in G_4 and G_6 (a consequence of the fact that the ring of modular forms for $\mathrm{SL}_2(\mathbb{Z})$ is a polynomial ring in G_4 and G_6 ; see ??).

Differentiating (3.4) term by term, we also obtain the Laurent expansion of $\wp'$:

$$\wp'(z) = -\frac{2}{z^3} + 6 G_4 z + 20 G_6 z^3 + 42 G_8 z^5 + \cdots \quad (3.6)$$

The Differential Equation

The Laurent expansions of $\wp$ and $\wp'$ make it possible to find the algebraic relation between them by a direct comparison of poles and coefficients.

Theorem 3.2.6: Weierstrass Differential Equation

For any lattice Λ , the $\wp$ -function satisfies the differential equation

$$\wp'(z)^2 = 4 \wp(z)^3 - g_2 \wp(z) - g_3, \quad (3.7)$$

where

$$g_2 = g_2(\Lambda) = 60 G_4(\Lambda), \quad g_3 = g_3(\Lambda) = 140 G_6(\Lambda). \quad (3.8)$$

Proof:

The strategy is to show that $h(z) = \wp'(z)^2 - 4\wp(z)^3 + g_2\wp(z) + g_3$ is an elliptic function with no poles, hence constant by Liouville's theorem, and then that the constant is zero. We carry this out by comparing Laurent series at $z = 0$.

Write $\wp(z) = z^{-2} + c_2 z^2 + c_4 z^4 + \cdots$ with $c_{2k} = (2k+1)G_{2k+2}$, so $c_2 = 3G_4$ and $c_4 = 5G_6$. The derivative is $\wp'(z) = -2z^{-3} + 2c_2 z + 4c_4 z^3 + \cdots$.

Step 1: Compute $\wp'(z)^2$ through order z^0 . Squaring $\wp' = -2z^{-3} + 2c_2 z + 4c_4 z^3 + \cdots$: the cross terms through z^0 are $2(-2z^{-3})(2c_2 z) = -8c_2 z^{-2}$ and $2(-2z^{-3})(4c_4 z^3) = -16c_4$. The

$(2c_2z)^2 = 4c_2^2z^2$ term is $O(z^2)$. Thus:

$$(\wp')^2 = 4z^{-6} - 8c_2z^{-2} - 16c_4 + O(z^2) = 4z^{-6} - 24G_4z^{-2} - 80G_6 + O(z^2). \quad (3.9)$$

Step 2: Compute $4\wp(z)^3$ through order z^0 . We first compute $\wp^2$ through order z^2 . The product $(z^{-2} + c_2z^2 + c_4z^4 + \dots)^2$ gives: z^{-4} from $z^{-2} \cdot z^{-2}$; $2c_2$ from $z^{-2} \cdot c_2z^2$ (twice); and $2c_4z^2$ from $z^{-2} \cdot c_4z^4$ (twice). So $\wp^2 = z^{-4} + 2c_2 + 2c_4z^2 + O(z^4)$.

Now $\wp^3 = \wp \cdot \wp^2 = (z^{-2} + c_2z^2 + c_4z^4 + \dots)(z^{-4} + 2c_2 + 2c_4z^2 + \dots)$. Reading off:

- z^{-6} : $z^{-2} \cdot z^{-4} = 1$.
- z^{-2} : $z^{-2} \cdot 2c_2 + c_2z^2 \cdot z^{-4} = 3c_2$.
- z^0 : $z^{-2} \cdot 2c_4z^2 + c_4z^4 \cdot z^{-4} = 3c_4$.

(All other cross terms contributing at z^0 — such as $c_2z^2 \cdot 2c_2 = 2c_2^2z^2$ — are $O(z^2)$.) Therefore:

$$4\wp^3 = 4z^{-6} + 12c_2z^{-2} + 12c_4 + O(z^2) = 4z^{-6} + 36G_4z^{-2} + 60G_6 + O(z^2). \quad (3.10)$$

Step 3: Assemble $h(z)$. Subtracting (3.10) from (3.9):

$$(\wp')^2 - 4\wp^3 = -60G_4z^{-2} - 140G_6 + O(z^2).$$

Adding $g_2\wp + g_3 = 60G_4(z^{-2} + O(z^2)) + 140G_6$:

$$h(z) = (-60G_4 + 60G_4)z^{-2} + (-140G_6 + 140G_6) + O(z^2) = O(z^2).$$

In particular, h has no pole at $z = 0$: the z^{-6} terms cancel between $(\wp')^2$ and $4\wp^3$, the z^{-2} terms cancel after adding $g_2\wp$, and the constant terms cancel after adding g_3 .

Since h is a linear combination of elliptic functions, it is elliptic. Its only possible poles are at $z \equiv 0 \pmod{\Lambda}$, and we have just shown these are removable. By Liouville's theorem, h is constant. Since $h(z) = O(z^2)$ as $z \rightarrow 0$, the constant must be zero.

The computation above is bookkeeping-heavy but the underlying idea is simple: $(\wp')^2$ and $4\wp^3$ are both elliptic functions with a pole of order 6 at the origin, and subtracting them cancels the leading term $4z^{-6}$. The remaining polar term (z^{-2}) is canceled by adding $g_2\wp$, and the constant is canceled by g_3 . Liouville does the rest.

The relation (3.7) is a *Weierstrass equation*: setting $x = \wp(z)$ and $y = \wp'(z)$, we obtain

$$y^2 = 4x^3 - g_2x - g_3. \quad (3.11)$$

Comparing with the short Weierstrass form $Y^2 = X^3 + AX + B$ from section 1.1: the substitution $X = \wp(z)$, $Y = \wp'(z)/2$ transforms (3.11) into

$$Y^2 = X^3 - \frac{g_2}{4}X - \frac{g_3}{4},$$

so $A = -g_2/4$ and $B = -g_3/4$. The Weierstrass equation of the elliptic curve $\mathbb{C}/\Lambda$ is determined by the Eisenstein series of the lattice.

Example 3.2.7: The Square Lattice

For $\Lambda = \mathbb{Z} + \mathbb{Z}i$ (the Gaussian lattice), the symmetry $\omega \mapsto i\omega$ preserves Λ (since $i \cdot \mathbb{Z} + i \cdot \mathbb{Z}i = \mathbb{Z}i + \mathbb{Z}(-1) = \Lambda$). This implies $G_{2k}(i\Lambda) = i^{-2k}G_{2k}(\Lambda)$. But $i\Lambda = \Lambda$, so $G_{2k}(\Lambda) = i^{-2k}G_{2k}(\Lambda)$, which forces $G_{2k} = 0$ whenever $4 \nmid 2k$. In particular, $G_6 = 0$ (since $i^{-6} = -1 \neq 1$). Thus $g_3 = 140G_6 = 0$, and the Weierstrass equation is

$$y^2 = 4x^3 - g_2x, \quad \text{i.e., } Y^2 = X^3 - \frac{g_2}{4}X.$$

This is the curve $Y^2 = X^3 + AX$ with $A = -g_2/4 < 0$ (one can show $g_2 > 0$ for the square lattice) and $B = 0$. By proposition 1.2.11, this curve has $j = 1728$, matching our computation in example 1.2.7(a). The extra symmetry $\omega \mapsto i\omega$ of the lattice is the complex multiplication $[i] : z \mapsto iz$ on the torus $\mathbb{C}/\Lambda$, which descends to the automorphism $(X, Y) \mapsto (-X, -iY)$ of the elliptic curve — exactly the order-4 automorphism from theorem 1.5.2(2).

Example 3.2.8: The Hexagonal Lattice

For $\Lambda = \mathbb{Z} + \mathbb{Z}\omega$ where $\omega = e^{2\pi i/3} = (-1 + \sqrt{-3})/2$, the symmetry $\omega \mapsto \omega \cdot \Lambda$ preserves Λ (since $\omega \cdot 1 = \omega$ and $\omega \cdot \omega = \omega^2 = -1 - \omega$). This gives $G_{2k}(\Lambda) = \omega^{-2k}G_{2k}(\Lambda)$, forcing $G_{2k} = 0$ unless $3 \mid k$. In particular, $G_4 = 0$ (since $\omega^{-4} = \omega^{-1} = \omega^2 \neq 1$). Thus $g_2 = 60G_4 = 0$, and the Weierstrass equation is

$$y^2 = 4x^3 - g_3, \quad \text{i.e., } Y^2 = X^3 - \frac{g_3}{4}.$$

This is a curve with $A = 0$ and $j = 0$, matching example 1.2.7(b). The lattice symmetry $z \mapsto \omega z$ gives the order-6 automorphism ψ from theorem 1.5.2(3).

3.2.9 Remark: The Eisenstein Series G_2 The sum $G_2(\Lambda) = \sum_{\omega \neq 0} \omega^{-2}$ does *not* converge absolutely (lemma 3.2.3 gives divergence for $s = 2$). One can assign it a value via a specific ordering of the summation (Eisenstein summation), but the resulting function $G_2(\tau)$ is *not* modular of weight 2 — it satisfies a modified transformation law with a correction term. This failure is related to the fact that the space of modular forms of weight 2 for $\text{SL}_2(\mathbb{Z})$ is zero-dimensional. The “non-holomorphic Eisenstein series” $G_2^*(\tau) = G_2(\tau) - \pi/\text{Im}(\tau)$ transforms correctly but is not holomorphic. This subtlety will not affect us until ??; for now, all the Eisenstein series we use ($G_4, G_6, G_8, \dots$) converge absolutely and transform as expected.

Exercise 3.2.1

1. Verify the Laurent expansion (3.6) by differentiating (3.4) term by term. Use this to confirm that $\wp'(z)$ is an odd function of order 3, with a triple pole at $z = 0$ and zeros at the three half-lattice points $\omega_1/2, \omega_2/2, (\omega_1 + \omega_2)/2$.
2. Carry out the Laurent expansion comparison in the proof of theorem 3.2.6 to one higher order: compute the coefficient of z^2 in both $(\wp')^2$ and $4\wp^3 - g_2\wp - g_3$, and verify they agree. (*The coefficient of z^2 in both sides should be 0, reflecting the fact that $h(z)$ is identically zero, not just zero through z^0 .*)

3. Using the fact that $(\wp')^2 = 4\wp^3 - g_2\wp - g_3$ and the Laurent expansion of both sides, show that

$$G_8 = \frac{3}{7} G_4^2.$$

(*Hint:* Compare the coefficients of z^2 in the differential equation. This is the first instance of a “recurrence” for Eisenstein series.)

4. By the same method, show that $G_{10} = \frac{5}{11} G_4 G_6$.
5. (The addition formula for $\wp$.) Let $a, b \in \mathbb{C}$ with $a, b, a \pm b \notin \Lambda$. Show that

$$\wp(a+b) = -\wp(a) - \wp(b) + \frac{1}{4} \left(\frac{\wp'(a) - \wp'(b)}{\wp(a) - \wp(b)} \right)^2.$$

(*Hint:* Consider the elliptic function $f(z) = \wp'(z) - \lambda\wp(z) - \mu$ where λ and μ are chosen so that $f(a) = f(b) = 0$. Then f has a triple pole at $z = 0$ and zeros at a, b , and a third point c . By proposition 3.1.7(5), $a + b + c \equiv 0 \pmod{\Lambda}$, so $c \equiv -(a + b)$. Now use $f(c) = 0$ and the explicit formula for λ .)

6. Show that $\wp(z; \alpha\Lambda) = \alpha^{-2} \wp(z/\alpha; \Lambda)$ for any $\alpha \in \mathbb{C}^*$. Conclude that $g_2(\alpha\Lambda) = \alpha^{-4} g_2(\Lambda)$ and $g_3(\alpha\Lambda) = \alpha^{-6} g_3(\Lambda)$.

3.3 The Uniformization Theorem

We have seen that the Weierstrass $\wp$ -function for a lattice Λ satisfies the differential equation $(\wp')^2 = 4\wp^3 - g_2\wp - g_3$, which is a Weierstrass equation. The map $z \mapsto (\wp(z), \wp'(z))$ therefore sends points of the complex torus $\mathbb{C}/\Lambda$ to points of the corresponding elliptic curve. In this section we prove that this map is an isomorphism — of Riemann surfaces, of algebraic curves, and of groups — and that every elliptic curve over $\mathbb{C}$ arises this way.

The Uniformization Map

Theorem 3.3.1: Uniformization Theorem

Let Λ be a lattice in $\mathbb{C}$, and let E be the elliptic curve defined by $Y^2 = 4X^3 - g_2(\Lambda)X - g_3(\Lambda)$, with $g_2 = 60G_4(\Lambda)$ and $g_3 = 140G_6(\Lambda)$. The map

$$\varphi : \mathbb{C}/\Lambda \longrightarrow E(\mathbb{C}), \quad z \longmapsto \begin{cases} [\wp(z) : \wp'(z) : 1] & \text{if } z \notin \Lambda, \\ O = [0 : 1 : 0] & \text{if } z \in \Lambda, \end{cases} \quad (3.12)$$

is an isomorphism of complex Lie groups. In particular, it is a biholomorphism of Riemann surfaces.

Before proving this, let us note the key consequence: the complex torus $\mathbb{C}/\Lambda$ and the elliptic curve $E(\mathbb{C})$ are the same object, viewed from different perspectives. Passing from $\mathbb{C}/\Lambda$ to $E(\mathbb{C})$ via φ means “writing down the Weierstrass embedding”; passing from $E(\mathbb{C})$ to $\mathbb{C}/\Lambda$ via φ^{-1} means “integrating the invariant differential.” The group law on $E(\mathbb{C})$ corresponds to simple addition in $\mathbb{C}$ modulo Λ .

Proof:

We prove the theorem in several steps.

Step 1: $\wp$ is well-defined. For $z \notin \Lambda$, the point $(\wp(z), \wp'(z))$ lies on E by theorem 3.2.6. Since $\wp$ and $\wp'$ are Λ -periodic, $\wp$ is well-defined on $\mathbb{C}/\Lambda$. At $z = 0$ (a lattice point), $\wp$ has a double pole and $\wp'$ has a triple pole, so in projective coordinates $[\wp(z) : \wp'(z) : 1] = [1/\wp'(z) : 1 : 1/\wp'(z) \cdot 1/\wp(z) \cdot \wp(z)] \dots$ it is cleaner to use the local parameter $t = 1/z$ near $z = 0$. We have $\wp(z) \sim z^{-2}$ and $\wp'(z) \sim -2z^{-3}$, so

$$[\wp(z) : \wp'(z) : 1] = [z^{-2} : -2z^{-3} : 1] = [z : -2 : z^3] \rightarrow [0 : -2 : 0] = [0 : 1 : 0]$$

as $z \rightarrow 0$. Thus $\wp$ extends continuously (and holomorphically) to $z = 0$ with $\wp(0) = O$.

Step 2: $\wp$ is holomorphic. Away from $z = 0$, $\wp$ is a composition of meromorphic functions and the embedding $E \hookrightarrow \mathbb{P}^2$, hence holomorphic. At $z = 0$, the extension in Step 1 is holomorphic because $\wp$ and $\wp'$ have poles of known order, and the transition to projective coordinates removes the singularity (as the computation above shows).

Step 3: $\wp$ is injective. Suppose $\wp(z_1) = \wp(z_2)$ with $z_1, z_2 \notin \Lambda$. Then $\wp(z_1) = \wp(z_2)$ and $\wp'(z_1) = \wp'(z_2)$. Since $\wp$ has order 2, the equation $\wp(z) = c$ has exactly two solutions in $\mathbb{C}/\Lambda$: they are $z = a$ and $z = -a$ for some a (because $\wp$ is even, so $\wp(a) = \wp(-a)$). Thus $z_2 \equiv \pm z_1 \pmod{\Lambda}$.

If $z_2 \equiv -z_1$ and $z_2 \not\equiv z_1$, then $\wp'(z_1) = \wp'(z_2) = \wp'(-z_1) = -\wp'(z_1)$ (since $\wp'$ is odd), giving $\wp'(z_1) = 0$. But this means z_1 is a zero of $\wp'$, i.e., z_1 is a half-lattice point (exercise 3.2.1 (1)). At such a point, $z_1 \equiv -z_1 \pmod{\Lambda}$, so in fact $z_2 \equiv z_1$ after all. Thus $z_1 \equiv z_2 \pmod{\Lambda}$ in all cases: $\wp$ is injective.

Step 4: $\wp$ is surjective. Let $(x_0, y_0) \in E(\mathbb{C})$ be an affine point, so $y_0^2 = 4x_0^3 - g_2x_0 - g_3$. Since $\wp : \mathbb{C}/\Lambda \rightarrow \mathbb{P}^1$ has degree 2, every value in $\mathbb{C}$ is attained by $\wp$: there exists $z_0 \in \mathbb{C}/\Lambda$ with $\wp(z_0) = x_0$. By the differential equation, $\wp'(z_0)^2 = y_0^2$, so $\wp'(z_0) = \pm y_0$. If $\wp'(z_0) = y_0$, we are done. If $\wp'(z_0) = -y_0$, replace z_0 by $-z_0$: then $\wp(-z_0) = \wp(z_0) = x_0$ and $\wp'(-z_0) = -\wp'(z_0) = y_0$. In either case, $\wp(z_0) = (x_0, y_0)$.

Step 5: $\wp$ is biholomorphic. A bijective holomorphic map between compact Riemann surfaces is automatically biholomorphic (the inverse is holomorphic by the open mapping theorem or by the fact that a bijective morphism of smooth curves is an isomorphism). Since $\mathbb{C}/\Lambda$ and $E(\mathbb{C})$ are both compact Riemann surfaces and $\wp$ is a bijective holomorphic map, $\wp$ is a biholomorphism.

3.3.2 Remark: The Invariant Differential The inverse map $\varphi^{-1} : E(\mathbb{C}) \rightarrow \mathbb{C}/\Lambda$ can be described explicitly using integration. The *invariant differential* on E (in the form $Y^2 = 4X^3 - g_2X - g_3$) is $\omega = dX/Y = dX/(2Y)$ in the standard short form. Since $X = \wp(z)$ and $Y = \wp'(z)$, we have $dX = \wp'(z) dz$ and $Y = \wp'(z)$, so $\omega = dz$. Thus $\varphi^*\omega = dz$, and the inverse map is given by integration:

$$\varphi^{-1}(P) = \int_O^P \omega \pmod{\Lambda},$$

where the integral is taken along any path from O to P on the Riemann surface $E(\mathbb{C})$. Different paths yield values differing by elements of Λ (the periods of ω), hence the result is well-defined in $\mathbb{C}/\Lambda$. This is the classical *Abel–Jacobi map* for genus-1 curves.

The Uniformization as a Group Isomorphism

The map φ is not merely a biholomorphism of Riemann surfaces — it is a group homomorphism, and therefore a group isomorphism. This means the group law on $E(\mathbb{C})$ (defined via Pic^0 in section 1.3) corresponds to plain addition in $\mathbb{C}$ modulo Λ .

Theorem 3.3.3: φ is a Group Isomorphism

The map $\varphi : (\mathbb{C}/\Lambda, +) \rightarrow (E(\mathbb{C}), +)$ is an isomorphism of abelian groups.

Proof:

We must show $\varphi(z_1 + z_2) = \varphi(z_1) + \varphi(z_2)$ for all $z_1, z_2 \in \mathbb{C}/\Lambda$, where the left side is addition in $\mathbb{C}/\Lambda$ and the right side is the group law on E . Since φ is already a bijection, it suffices to prove it is a homomorphism.

Recall from definition 1.3.4 that the group law on E is characterized by: $P + Q + R = O$ if and only if $(P) + (Q) + (R) \sim 3(O)$, which happens if and only if P, Q, R are the three intersection points (counted with multiplicity) of a line with E . Equivalently, $P + Q = S$ means the line through P and Q meets E in a third point R , and the vertical line through R meets E at S (the “reflection” of R).

We use the analytic characterization. Let $P = \varphi(a)$ and $Q = \varphi(b)$ with $a, b \notin \Lambda$ and $a \not\equiv \pm b \pmod{\Lambda}$ (the generic case). Consider the elliptic function

$$f(z) = \wp'(z) - \lambda \wp(z) - \mu,$$

where λ and μ are chosen so that $f(a) = f(b) = 0$:

$$\lambda = \frac{\wp'(a) - \wp'(b)}{\wp(a) - \wp(b)}, \quad \mu = \frac{\wp'(b)\wp(a) - \wp'(a)\wp(b)}{\wp(a) - \wp(b)}.$$

Note that λ is the slope of the line ℓ through P and Q in the (X, Y) -plane, and $Y = \lambda X + \mu$ is the equation of ℓ .

The function f has a triple pole at $z = 0$ (from $\wp'$; the double pole of $\wp$ is of lower order). Since

f is elliptic of order 3, it has exactly three zeros in a fundamental parallelogram: a , b , and a third zero c . By proposition 3.1.7(5) (the Abel–Jacobi constraint), $a + b + c \equiv 0 \pmod{\Lambda}$, so $c \equiv -(a + b)$.

The zeros of f correspond to the intersection points of ℓ with E : these are $P = \varphi(a)$, $Q = \varphi(b)$, and $R = \varphi(c) = \varphi(-(a + b))$. The three intersection points of a line with E satisfy $P + Q + R = O$ in the group law (exercise 1.3.1 (5)). So:

$$\varphi(a) + \varphi(b) + \varphi(-(a + b)) = O.$$

Since $-(a + b) \equiv -(a + b) \pmod{\Lambda}$, we have $\varphi(-(a + b)) = (\wp(-(a + b)), \wp'(-(a + b))) = (\wp(a + b), -\wp'(a + b))$ (using evenness of $\wp$ and oddness of $\wp'$). This point is exactly $-\varphi(a + b)$ (the group inverse is $(X, Y) \mapsto (X, -Y)$ for the form $Y^2 = 4X^3 - \dots$). Therefore:

$$\varphi(a) + \varphi(b) - \varphi(a + b) = O, \quad \text{i.e.,} \quad \varphi(a) + \varphi(b) = \varphi(a + b).$$

The special cases ($a \equiv \pm b$, or a or b in Λ) are handled by continuity or by direct verification with the doubling formula. Thus φ is a group homomorphism, and since it is bijective, it is a group isomorphism.

This result is the fundamental bridge between the analytic and algebraic theories. It tells us, for instance, that the addition formula for $\wp$ (exercise 3.2.1 (5)) is not a curiosity but the *analytic form of the group law*: the formula $\wp(a + b) = -\wp(a) - \wp(b) + \frac{1}{4} \left(\frac{\wp'(a) - \wp'(b)}{\wp(a) - \wp(b)} \right)^2$ is the chord-and-tangent law in disguise, with the slope $\lambda = (\wp'(a) - \wp'(b))/(\wp(a) - \wp(b))$ playing exactly the same role as in the algebraic addition formulas of section 1.3.

The Converse: Every Elliptic Curve over $\mathbb{C}$ is a Torus

The uniformization theorem as stated constructs, for each lattice, an elliptic curve. The converse asserts that every elliptic curve over $\mathbb{C}$ arises from some lattice.

Theorem 3.3.4: Converse Uniformization

Let $E/\mathbb{C}$ be an elliptic curve. Then there exists a lattice $\Lambda \subset \mathbb{C}$ such that $E(\mathbb{C}) \cong \mathbb{C}/\Lambda$ as complex Lie groups. The lattice Λ is unique up to homothety.

Proof Key ideas:

There are two approaches: one via Riemann surface theory and one via explicit construction.

Approach 1 (Universal cover). The elliptic curve $E(\mathbb{C})$ is a compact, connected Riemann surface of genus 1. By the uniformization theorem for Riemann surfaces [4], the universal cover of any Riemann surface of genus ≥ 1 is either $\mathbb{C}$ or the unit disk $\mathbb{D}$. Since $E(\mathbb{C})$ admits a nowhere-vanishing holomorphic 1-form (namely dX/Y), it cannot be covered by $\mathbb{D}$ (the disk has negative Euler characteristic, and holomorphic 1-forms on $\mathbb{D}/\Gamma$ would have to come from automorphic forms with nontrivial transformation behavior). Thus the universal cover is $\mathbb{C}$, and $E(\mathbb{C}) \cong \mathbb{C}/\Gamma$ for some discrete subgroup $\Gamma \subset \mathbb{C}$ acting by translations. Since $E(\mathbb{C})$ is compact, Γ must have

rank 2, hence Γ is a lattice.

Approach 2 (Integration of the invariant differential). Put E in short Weierstrass form $Y^2 = X^3 + AX + B$, and consider the invariant differential $\omega = dX/(2Y)$. For a closed loop γ on $E(\mathbb{C})$, the period $\int_\gamma \omega$ depends only on the homology class $[\gamma] \in H_1(E(\mathbb{C}), \mathbb{Z}) \cong \mathbb{Z}^2$. The two periods $\omega_1 = \int_{\gamma_1} \omega$ and $\omega_2 = \int_{\gamma_2} \omega$ (where γ_1, γ_2 form a basis for H_1) are $\mathbb{R}$ -linearly independent (this follows from the Riemann bilinear relations), so they generate a lattice $\Lambda = \mathbb{Z}\omega_1 + \mathbb{Z}\omega_2$.

The Abel–Jacobi map $\Phi : E(\mathbb{C}) \rightarrow \mathbb{C}/\Lambda$ defined by $P \mapsto \int_O^P \omega \pmod{\Lambda}$ is a holomorphic group homomorphism. It is injective (because if $\Phi(P) = 0$ then P is a zero of a holomorphic differential on E , forcing $P = O$) and surjective (because $\mathbb{C}/\Lambda$ is compact and Φ is a non-constant holomorphic map between compact Riemann surfaces of the same genus). Hence Φ is an isomorphism.

For uniqueness: if $E(\mathbb{C}) \cong \mathbb{C}/\Lambda \cong \mathbb{C}/\Lambda'$, then Λ and Λ' are related by the automorphism $z \mapsto \alpha z$ of $\mathbb{C}$ (the only biholomorphisms of $\mathbb{C}$ that fix the origin are the linear maps), so $\Lambda' = \alpha\Lambda$.

The converse uniformization theorem, combined with theorem 3.3.1, establishes a bijection:

$$\left\{ \begin{array}{c} \text{Lattices in } \mathbb{C} \\ \text{up to homothety} \end{array} \right\} \xleftrightarrow{1:1} \left\{ \begin{array}{c} \text{Elliptic curves over } \mathbb{C} \\ \text{up to isomorphism} \end{array} \right\}. \quad (3.13)$$

By proposition 3.1.4, the left side is $\mathrm{SL}_2(\mathbb{Z}) \backslash \mathbb{H}$, and by theorem 1.2.8, the right side is $\mathbb{C}$ (via the j -invariant). The j -function $j : \mathbb{H} \rightarrow \mathbb{C}$ (to be constructed in section 3.4) mediates between the two: $j(\tau) = j(E_{\Lambda_\tau})$.

The Discriminant and 2-Torsion

The uniformization gives us immediate access to the discriminant and the 2-torsion of the associated curve.

Proposition 3.3.5: The Roots of the Weierstrass Cubic

Let $\Lambda = \mathbb{Z}\omega_1 + \mathbb{Z}\omega_2$. Define the half-lattice points

$$e_1 = \wp(\omega_1/2), \quad e_2 = \wp(\omega_2/2), \quad e_3 = \wp((\omega_1 + \omega_2)/2).$$

Then e_1, e_2, e_3 are the three roots of the cubic $4X^3 - g_2X - g_3 = 0$. They are pairwise distinct, and:

1. $e_1 + e_2 + e_3 = 0$.
2. $e_1e_2 + e_1e_3 + e_2e_3 = -g_2/4$.
3. $e_1e_2e_3 = g_3/4$.
4. The discriminant satisfies

$$\Delta(\Lambda) = 16[(e_1 - e_2)(e_1 - e_3)(e_2 - e_3)]^2, \quad (3.14)$$

and in particular $\Delta(\Lambda) \neq 0$.

Proof:

The half-lattice points $\omega_1/2, \omega_2/2, (\omega_1 + \omega_2)/2$ are exactly the points $z \in \mathbb{C}/\Lambda$ satisfying $z \equiv -z \pmod{\Lambda}$ but $z \not\equiv 0$. At these points, $\wp'(z) = 0$ (since $\wp'(-z) = -\wp'(z)$ and $z \equiv -z$ gives $\wp'(z) = -\wp'(z)$, hence $\wp'(z) = 0$). The differential equation $\wp'(z)^2 = 4\wp(z)^3 - g_2\wp(z) - g_3$ then gives $4e_i^3 - g_2e_i - g_3 = 0$ for $i = 1, 2, 3$: the e_i are roots of $4X^3 - g_2X - g_3$.

Since $\wp$ has order 2 and is even, the equation $\wp(z) = c$ has two solutions $\{a, -a\}$ for generic c . For $c = e_i$, the two solutions $\{z, -z\}$ coincide (since z is a half-lattice point), so e_i is a value of $\wp$ of multiplicity 2. This means e_1, e_2, e_3 are three *distinct* roots of a cubic (each with multiplicity 1 as a root of the polynomial): they are distinct because $\wp$ takes each value exactly twice (counted with multiplicity), and the three half-lattice points are distinct elements of $\mathbb{C}/\Lambda$.

Properties (1)–(3) follow from Vieta's formulas applied to $4X^3 - g_2X - g_3 = 4(X - e_1)(X - e_2)(X - e_3)$. For (4), the discriminant of $4X^3 - g_2X - g_3$ is $\Delta = -16(4A^3 + 27B^2)$ where $A = -g_2/4$ and $B = -g_3/4$ (from section 1.2), and the classical formula $\Delta = 16 \prod_{i < j} (e_i - e_j)^2$ follows. Since e_1, e_2, e_3 are distinct, $\Delta \neq 0$.

The non-vanishing of $\Delta(\Lambda)$ is reassuring: it confirms that the Weierstrass equation $Y^2 = 4X^3 - g_2X - g_3$ always defines a smooth curve (an elliptic curve) when g_2 and g_3 come from a lattice. The identity (3.14) also gives the discriminant an analytic meaning in terms of the half-periods.

For the 2-torsion, the uniformization immediately yields:

Corollary 3.3.6: 2-Torsion of $E(\mathbb{C})$

Under the uniformization $\varphi : \mathbb{C}/\Lambda \xrightarrow{\sim} E(\mathbb{C})$, the 2-torsion subgroup is

$$E[2](\mathbb{C}) = \varphi(\tfrac{1}{2}\Lambda/\Lambda) = \{O, (e_1, 0), (e_2, 0), (e_3, 0)\} \cong (\mathbb{Z}/2\mathbb{Z})^2.$$

More generally, $E[n](\mathbb{C}) = \varphi(\tfrac{1}{n}\Lambda/\Lambda) \cong (\mathbb{Z}/n\mathbb{Z})^2$ for all $n \geq 1$.

Proof:

Under φ , the n -torsion points of $E(\mathbb{C})$ correspond to elements $z \in \mathbb{C}/\Lambda$ with $nz \in \Lambda$, i.e., $z \in \tfrac{1}{n}\Lambda/\Lambda$. This group is $(\mathbb{Z}\omega_1/n + \mathbb{Z}\omega_2/n)/(\mathbb{Z}\omega_1 + \mathbb{Z}\omega_2) \cong (\mathbb{Z}/n\mathbb{Z})^2$. For $n = 2$, the nontrivial elements are $\omega_1/2, \omega_2/2, (\omega_1 + \omega_2)/2$, mapping to $(e_1, 0), (e_2, 0), (e_3, 0)$.

This gives a clean proof of the fact $E[n](\overline{K}) \cong (\mathbb{Z}/n\mathbb{Z})^2$ that was stated (for $K = \mathbb{C}$) but not proved in section 1.3. The algebraic proof over arbitrary algebraically closed fields of characteristic coprime to n requires more machinery (the theory of isogenies, chapter 2); the analytic proof is immediate.

Example 3.3.7: The 3-Torsion of $y^2 = x^3 + 1$

For $E : Y^2 = X^3 + 1$ over $\mathbb{C}$ (with $g_2 = 0$ and $g_3 = -4$), the uniformizing lattice is $\Lambda = \mathbb{Z} + \mathbb{Z}\omega$ with $\omega = e^{2\pi i/3}$ (up to scaling). The 3-torsion consists of the nine points $\tfrac{1}{3}\Lambda/\Lambda = \{(a + b\omega)/3 : 0 \leq a, b \leq 2\}$. The identity O corresponds to $(a, b) = (0, 0)$. The point $P = (0, 1)$ in $E(\mathbb{C})$ corresponds to $z = \omega/3$ (or $z = (1 + \omega)/3$) in $\mathbb{C}/\Lambda$, which is a 3-torsion point: this is consistent with the computation in example 1.3.8 that $3P = O$ on this curve.

The remaining eight non-identity 3-torsion points include the six points $\varphi(a/3 + b\omega/3)$ for $(a, b) \neq (0, 0)$. Their coordinates involve the values $\wp(1/3), \wp(\omega/3), \wp((1 + \omega)/3), \dots$, which are algebraic numbers (they lie in the field generated by the 3-division values of $\wp$). The Galois action on these coordinates will be a key ingredient in the theory of complex multiplication (??).

Example 3.3.8: Periods of Running Examples

(a) $y^2 = x^3 - x$ ($j = 1728$). The invariant differential is $\omega = dx/(2y)$. The periods of E are

$$\omega_1 = \int_{-1}^0 \frac{dx}{2\sqrt{x^3 - x}} = \int_{-1}^0 \frac{dx}{2\sqrt{-x(1-x)(1+x)}},$$

and $\omega_2 = i\omega_1$ (by the symmetry $x \mapsto -x, y \mapsto iy$). These are related to the lemniscate constant: $\omega_1 = \Gamma(1/4)^2/(4\sqrt{\pi})$. The period ratio is $\tau = \omega_1/\omega_2 = i$, confirming that the uniformizing lattice is $\Lambda_i = \mathbb{Z} + \mathbb{Z}i$ (up to scaling).

(b) $y^2 = x^3 + 1$ ($j = 0$). Similarly, the period ratio is $\tau = e^{2\pi i/3}$, and the lattice is $\Lambda_\omega = \mathbb{Z} + \mathbb{Z}\omega$ (up to scaling). The real period involves the integral $\int_{-1}^{\infty} dx/(2\sqrt{x^3 + 1})$, which is expressible in terms of $\Gamma(1/3)$.

Exercise 3.3.1

1. Let $E : y^2 = x^3 + Ax + B$ be an elliptic curve over $\mathbb{C}$ in standard short Weierstrass form, and let $\omega = dx/(2y)$ be the invariant differential.
 - (a) Show that ω is holomorphic and nowhere-vanishing on E (including at the point at infinity O).
 - (b) Conclude that the space of holomorphic differentials on E is one-dimensional, spanned by ω . (*This is consistent with $g = 1$.*)
2. In the proof of theorem 3.3.3, we used the Abel–Jacobi constraint to show $\varphi(a) + \varphi(b) = \varphi(a+b)$ for generic a, b . Handle the remaining cases:
 - (a) $a \equiv -b \pmod{\Lambda}$ but $a \not\equiv 0$ (so $P = -Q$).
 - (b) $a \equiv b \pmod{\Lambda}$ (the doubling case $P = Q$).
 - (c) $a \equiv 0$ or $b \equiv 0$ (one point is O).
3. Let $E : y^2 = 4x^3 - g_2x - g_3$ with roots $e_1 > e_2 > e_3$ (assumed real). Show that the real period of E is

$$\omega_1 = 2 \int_{e_1}^{\infty} \frac{dx}{\sqrt{4x^3 - g_2x - g_3}}$$

and compute the imaginary period ω_2 as a similar integral along a path from e_3 to e_2 .

4. Using $\Delta = 16(e_1 - e_2)^2(e_1 - e_3)^2(e_2 - e_3)^2$ and the Vieta relations $e_1 + e_2 + e_3 = 0$, $e_1e_2 + e_1e_3 + e_2e_3 = -g_2/4$, $e_1e_2e_3 = g_3/4$, verify directly that $\Delta = g_2^3 - 27g_3^2$.
5. Show that the 5-torsion $E[5](\mathbb{C})$ of any elliptic curve $E/\mathbb{C}$ has exactly 25 elements, and determine its group structure.
6. For the congruent number curve $E_n : y^2 = x^3 - n^2x$, the uniformizing lattice Λ_n is related to Λ_1 (the lattice for $y^2 = x^3 - x$) by a scaling. Determine this scaling factor in terms of n and the real period ω_1 of E_1 .

3.4 Eisenstein Series and the j -Invariant

The Eisenstein series $G_{2k}(\Lambda)$ first appeared as coefficients in the Laurent expansion of $\wp$. In this section we promote them to functions of the period ratio $\tau \in \mathbb{H}$, study their transformation behavior under $\mathrm{SL}_2(\mathbb{Z})$ (which is the origin of the theory of modular forms), and use them to construct the analytic j -function $j(\tau)$. The payoff is that the analytic j -function agrees with the algebraic j -invariant of section 1.2, closing the circle: lattice invariants = curve invariants.

Eisenstein Series as Functions of τ

For a lattice $\Lambda = \mathbb{Z}\omega_1 + \mathbb{Z}\omega_2$ and an integer $k \geq 2$, the Eisenstein series is $G_{2k}(\Lambda) = \sum_{\omega \in \Lambda \setminus \{0\}} \omega^{-2k}$. Passing to the standard lattice $\Lambda_\tau = \mathbb{Z}\tau + \mathbb{Z}$ (with $\tau = \omega_1/\omega_2 \in \mathbb{H}$), we define:

Definition 3.4.1: Eisenstein Series on $\mathbb{H}$

For $\tau \in \mathbb{H}$ and an integer $k \geq 2$, the *Eisenstein series of weight $2k$* is

$$G_{2k}(\tau) = G_{2k}(\Lambda_\tau) = \sum_{\substack{(m,n) \in \mathbb{Z}^2 \\ (m,n) \neq (0,0)}} \frac{1}{(m\tau + n)^{2k}}. \quad (3.15)$$

The sum converges absolutely for $k \geq 2$ by lemma 3.2.3.

We also define the normalized invariants

$$g_2(\tau) = 60 G_4(\tau), \quad g_3(\tau) = 140 G_6(\tau), \quad (3.16)$$

so that the elliptic curve associated to Λ_τ has the Weierstrass equation $y^2 = 4x^3 - g_2(\tau)x - g_3(\tau)$.

Modularity: The Transformation Law

The key property of Eisenstein series is their behavior under the action of $\mathrm{SL}_2(\mathbb{Z})$ on $\mathbb{H}$. This is the entry point to the theory of modular forms.

Proposition 3.4.2: Modular Transformation of G_{2k}

For $\gamma = \begin{pmatrix} a & b \\ c & d \end{pmatrix} \in \mathrm{SL}_2(\mathbb{Z})$ and $k \geq 2$,

$$G_{2k}(\gamma \cdot \tau) = (c\tau + d)^{2k} G_{2k}(\tau), \quad (3.17)$$

where $\gamma \cdot \tau = (a\tau + b)/(c\tau + d)$.

Proof:

The lattice $\Lambda_{\gamma \cdot \tau}$ is homothetic to Λ_τ with scaling factor $(c\tau + d)^{-1}$. Indeed, a basis for Λ_τ is $(\tau, 1)$, and a basis for $\Lambda_{\gamma \cdot \tau}$ is $(\gamma \cdot \tau, 1)$. Since

$$\gamma \cdot \tau = \frac{a\tau + b}{c\tau + d}, \quad 1 = \frac{c\tau + d}{c\tau + d},$$

we see that $(\gamma \cdot \tau, 1) = (c\tau + d)^{-1}(a\tau + b, c\tau + d)$. Since $\begin{pmatrix} a & b \\ c & d \end{pmatrix} \in \mathrm{SL}_2(\mathbb{Z})$, the pair $(a\tau + b, c\tau + d)$ is another basis for Λ_τ . Thus $\Lambda_{\gamma \cdot \tau} = (c\tau + d)^{-1}\Lambda_\tau$.

The Eisenstein series transforms under homothety as $G_{2k}(\alpha\Lambda) = \alpha^{-2k}G_{2k}(\Lambda)$ (exercise 3.2.1 (6)), so:

$$G_{2k}(\gamma \cdot \tau) = G_{2k}(\Lambda_{\gamma \cdot \tau}) = G_{2k}((c\tau + d)^{-1}\Lambda_\tau) = (c\tau + d)^{2k} G_{2k}(\Lambda_\tau) = (c\tau + d)^{2k} G_{2k}(\tau).$$

A function $f : \mathbb{H} \rightarrow \mathbb{C}$ satisfying $f(\gamma \cdot \tau) = (c\tau + d)^k f(\tau)$ for all $\gamma \in \mathrm{SL}_2(\mathbb{Z})$ is called a *modular form of weight k* (provided it is holomorphic on $\mathbb{H}$ and satisfies a growth condition at $\tau \rightarrow i\infty$). The Eisenstein series G_{2k} is thus a modular form of weight $2k$ for $\mathrm{SL}_2(\mathbb{Z})$. The systematic theory of modular forms — including the structure of the graded ring they form, the role of Hecke operators, and the connection

to L -functions — will be developed in ???. For now, we need only the transformation law and its immediate consequences.

The transformation law for g_2 and g_3 follows:

$$g_2(\gamma \cdot \tau) = (c\tau + d)^4 g_2(\tau), \quad g_3(\gamma \cdot \tau) = (c\tau + d)^6 g_3(\tau). \quad (3.18)$$

Compare this with the algebraic transformation law of proposition 1.2.5: $c'_4 = u^{-4}c_4$ and $c'_6 = u^{-6}c_6$. The weights match. The analytic invariants g_2, g_3 and the algebraic invariants c_4, c_6 are the same objects (up to explicit numerical factors), viewed from different sides of the uniformization.

The Modular Discriminant

Definition 3.4.3: Modular Discriminant

The *modular discriminant* is the function $\Delta : \mathbb{H} \rightarrow \mathbb{C}$ defined by

$$\Delta(\tau) = g_2(\tau)^3 - 27 g_3(\tau)^2. \quad (3.19)$$

By (3.18), Δ transforms as

$$\Delta(\gamma \cdot \tau) = (c\tau + d)^{12} \Delta(\tau), \quad (3.20)$$

so Δ is a modular form of weight 12. This is the analytic incarnation of the weight -12 transformation of Δ under admissible changes in proposition 1.2.5: the sign difference ($+12$ vs. -12) reflects the duality between “transforming the lattice” (analytic) and “transforming the coordinates” (algebraic).

Proposition 3.4.4: Non-Vanishing of Δ on $\mathbb{H}$

$\Delta(\tau) \neq 0$ for all $\tau \in \mathbb{H}$.

Proof:

By proposition 3.3.5, $\Delta(\Lambda_\tau) = 16(e_1 - e_2)^2(e_1 - e_3)^2(e_2 - e_3)^2$, where e_1, e_2, e_3 are the distinct values of $\wp$ at the three half-lattice points. Since these are distinct (as proved in proposition 3.3.5), $\Delta(\tau) \neq 0$.

The modular discriminant admits a remarkable product expansion, which we state here and prove in ???:

$$\Delta(\tau) = (2\pi)^{12} q \prod_{n=1}^{\infty} (1 - q^n)^{24}, \quad q = e^{2\pi i \tau}. \quad (3.21)$$

This is the *Jacobi product formula*. The non-vanishing of Δ on $\mathbb{H}$ is transparent from this formula: $|q| = e^{-2\pi \text{Im}(\tau)} < 1$ for $\tau \in \mathbb{H}$, so neither q nor any factor $(1 - q^n)$ vanishes. The product formula also shows that $\Delta(\tau) \rightarrow 0$ as $\text{Im}(\tau) \rightarrow \infty$ (since $q \rightarrow 0$), and that this zero is simple. These properties will be important for the theory of modular forms.

The Analytic j -Function

Since g_2^3 has weight 12 and Δ has weight 12 (and is nonvanishing), their ratio is a modular function of weight 0:

Definition 3.4.5: The j -Function

The j -function (or modular invariant) is

$$j(\tau) = 1728 \frac{g_2(\tau)^3}{\Delta(\tau)}. \quad (3.22)$$

The factor of $1728 = 12^3$ is a normalization convention, chosen so that the q -expansion (see below) has leading coefficient 1. Some authors omit this factor; our convention matches the analytic tradition and ensures that the analytic j -function agrees with the algebraic j -invariant of definition 1.2.6.

Proposition 3.4.6: Properties of the j -Function

1. $j(\tau)$ is $\mathrm{SL}_2(\mathbb{Z})$ -invariant: $j(\gamma \cdot \tau) = j(\tau)$ for all $\gamma \in \mathrm{SL}_2(\mathbb{Z})$.
2. $j(\tau)$ is holomorphic on $\mathbb{H}$.
3. $j(\tau)$ has a simple pole at the cusp $\tau \rightarrow i\infty$ (i.e., as $q \rightarrow 0$).
4. j induces a bijection $\mathrm{SL}_2(\mathbb{Z}) \backslash \mathbb{H}^* \xrightarrow{\sim} \mathbb{P}^1(\mathbb{C})$, where $\mathbb{H}^* = \mathbb{H} \cup \{\text{cusps}\}$.
5. The special values are $j(i) = 1728$ and $j(e^{2\pi i/3}) = 0$.

Proof:

(1) Immediate from (3.18) and (3.20): $j(\gamma \cdot \tau) = 1728 g_2(\gamma \cdot \tau)^3 / \Delta(\gamma \cdot \tau) = 1728 (c\tau + d)^{12} g_2(\tau)^3 / ((c\tau + d)^{12} \Delta(\tau)) = j(\tau)$.

(2) Both $g_2(\tau)^3$ and $\Delta(\tau)$ are holomorphic on $\mathbb{H}$ (as Eisenstein series are holomorphic, and sums and products of holomorphic functions are holomorphic). Since $\Delta(\tau) \neq 0$ on $\mathbb{H}$ (proposition 3.4.4), the ratio is holomorphic.

(3) As $\tau \rightarrow i\infty$, $q = e^{2\pi i\tau} \rightarrow 0$. From the product formula (3.21), $\Delta(\tau) \sim (2\pi)^{12} q$ as $q \rightarrow 0$ (the infinite product approaches 1). Meanwhile, $g_2(\tau) \rightarrow (2\pi)^4 / 12 \cdot 2\zeta(4) \cdot (2\pi)^{-4} \dots$ — more precisely, $G_4(\tau) = 2\zeta(4) + O(q) = \pi^4/45 + O(q)$, so $g_2(\tau) = 60G_4(\tau) = 4\pi^4/3 + O(q)$, and $g_2(\tau)^3 = 64\pi^{12}/27 + O(q)$. Therefore

$$j(\tau) = 1728 \cdot \frac{64\pi^{12}/27 + O(q)}{(2\pi)^{12}q + O(q^2)} = \frac{1728 \cdot 64\pi^{12}/27}{(2\pi)^{12}} \cdot q^{-1} + O(1) = q^{-1} + O(1),$$

where the numerical prefactor is $1728 \cdot 64 / (27 \cdot 2^{12}) = 1728 \cdot 64 / 110592 = 1$. So j has a simple pole at the cusp with residue 1 in the q -expansion.

(4) By (1), j descends to a map $\mathrm{SL}_2(\mathbb{Z}) \backslash \mathbb{H}^* \rightarrow \mathbb{P}^1(\mathbb{C})$. The quotient $\mathrm{SL}_2(\mathbb{Z}) \backslash \mathbb{H}^*$ is a compact

Riemann surface (of genus 0, as we will see in section 3.6). By (3), j has exactly one pole on this surface (at the cusp), hence is a meromorphic function of degree 1. A degree-1 meromorphic function on a compact Riemann surface is a biholomorphism to $\mathbb{P}^1(\mathbb{C})$.

(5) For $\tau = i$: the square lattice has $G_6(i) = 0$ (example 3.2.7), so $g_3(i) = 0$ and $\Delta(i) = g_2(i)^3$. Thus $j(i) = 1728 \cdot g_2(i)^3 / g_2(i)^3 = 1728$.

For $\tau = e^{2\pi i/3}$: the hexagonal lattice has $G_4(e^{2\pi i/3}) = 0$ (example 3.2.8), so $g_2(e^{2\pi i/3}) = 0$ and $j(e^{2\pi i/3}) = 0$.

The q -Expansion

The behavior of $j(\tau)$ near the cusp is captured by its q -*expansion*: the Laurent series in $q = e^{2\pi i\tau}$.

Proposition 3.4.7: q -Expansion of j

$$j(\tau) = q^{-1} + 744 + 196884q + 21493760q^2 + 864299970q^3 + \cdots \quad (3.23)$$

where $q = e^{2\pi i\tau}$.

Proof Sketch:

From the q -expansions of G_4 and G_6 (derived by Poisson summation applied to the lattice sums):

$$G_4(\tau) = \frac{\pi^4}{45} \left(1 + 240 \sum_{n=1}^{\infty} \sigma_3(n) q^n \right), \quad (3.24)$$

$$G_6(\tau) = \frac{\pi^6}{945} \left(1 - 504 \sum_{n=1}^{\infty} \sigma_5(n) q^n \right), \quad (3.25)$$

where $\sigma_k(n) = \sum_{d|n} d^k$ is the k -th divisor sum. The normalized Eisenstein series $E_4 = G_4/(2\zeta(4)) = 1 + 240q + 2160q^2 + \cdots$ and $E_6 = G_6/(2\zeta(6)) = 1 - 504q - 16632q^2 - \cdots$ are the standard objects. Computing $g_2 = 60G_4$, $g_3 = 140G_6$, and forming $\Delta = g_2^3 - 27g_3^2$, the product formula (3.21) emerges after substantial simplification. The q -expansion of j then follows from $j = 1728g_2^3/\Delta$.

The full derivation of (3.24)–(3.25) via Poisson summation (or equivalently, via the Lipschitz summation formula) will be given in ???. For now, we record the result and note that the first few terms can be verified numerically.

The coefficient $744 = 3 \cdot 248$ and $196884 = 196883 + 1$ have remarkable interpretations: 196883 is the dimension of the smallest nontrivial representation of the Monster group, and the near-miss $196884 = 196883 + 1$ was the first hint of *Monstrous Moonshine*, a deep connection between modular functions and the representation theory of the Monster. This is far beyond our scope, but it illustrates the extraordinary richness of the j -function.

Agreement of Analytic and Algebraic j -Invariants

We now close the circle promised since section 1.2.

Theorem 3.4.8: Analytic = Algebraic

Let $E/\mathbb{C}$ be an elliptic curve with uniformizing lattice Λ , and let $\tau \in \mathbb{H}$ be a period ratio for Λ (so that $\Lambda \sim \Lambda_\tau$). Then the algebraic j -invariant $j(E)$ (from definition 1.2.6) equals the analytic j -function $j(\tau)$ (from definition 3.4.5).

Proof:

The uniformization gives E the Weierstrass equation $Y^2 = 4X^3 - g_2(\Lambda)X - g_3(\Lambda)$, or in the standard short form $Y^2 = X^3 + AX + B$ with $A = -g_2/4$ and $B = -g_3/4$. By section 1.2, the algebraic invariants are $c_4 = -48A = 12g_2$ and $\Delta_{\text{alg}} = -16(4A^3 + 27B^2) = -16(-g_2^3/16 + 27g_3^2/16) = g_2^3 - 27g_3^2 = \Delta(\Lambda)$.

The algebraic j -invariant is $j(E) = c_4^3/\Delta_{\text{alg}} = (12g_2)^3/(g_2^3 - 27g_3^2) = 1728 g_2^3/\Delta$, which is exactly $j(\tau)$ by definition.

For the dependence on τ : if $\Lambda \sim \Lambda_\tau$, then $\Lambda = \alpha\Lambda_\tau$ for some $\alpha \in \mathbb{C}^*$, and $g_2(\Lambda) = \alpha^{-4}g_2(\Lambda_\tau) = \alpha^{-4}g_2(\tau)$ (exercise 3.2.1 (6)), $g_3(\Lambda) = \alpha^{-6}g_3(\tau)$. The ratio g_2^3/Δ is weight-0 and hence independent of α : $j(E) = 1728 g_2(\tau)^3/\Delta(\tau) = j(\tau)$.

This theorem says we can compute the j -invariant of an elliptic curve $E/\mathbb{C}$ in two completely independent ways:

- *Algebraically*: from the Weierstrass coefficients $a_1, \dots, a_6$ via the formulas of section 1.2.
- *Analytically*: by finding the period ratio τ of the uniformizing lattice and evaluating $j(\tau) = q^{-1} + 744 + \dots$.

Both methods give the same answer. This agreement is a manifestation of a deeper principle: the algebraic structure of an elliptic curve over $\mathbb{C}$ is completely determined by its analytic structure (the underlying lattice), and vice versa. This is an instance of the “GAGA” principle (Géométrie Algébrique et Géométrie Analytique) of Serre.

Example 3.4.9: Computing j Both Ways

For $E : y^2 = x^3 - x$ over $\mathbb{C}$:

Algebraically (from example 1.2.7(a)): $A = -1$, $B = 0$, $\Delta = 64$, $c_4 = 48$, $j = 48^3/64 = 1728$.

Analytically: The period ratio is $\tau = i$. By proposition 3.4.6(5), $j(i) = 1728$.

Similarly, for $E : y^2 = x^3 + 1$: algebraically $j = 0$, and analytically $j(e^{2\pi i/3}) = 0$.

3.4.10 Remark: The j -Function and the Moduli Space The j -function $j : \mathbb{H} \rightarrow \mathbb{C}$ is the composition of two maps: the projection $\mathbb{H} \rightarrow \mathrm{SL}_2(\mathbb{Z}) \backslash \mathbb{H}$ (which sends τ to the isomorphism class of Λ_τ , hence to the isomorphism class of E_{Λ_τ}), followed by the bijection $\mathrm{SL}_2(\mathbb{Z}) \backslash \mathbb{H} \rightarrow \mathbb{C}$ given by j . The first map forgets the “marking” of the lattice (i.e., the choice of basis); the second map assigns the isomorphism invariant. The extended map $j : \mathrm{SL}_2(\mathbb{Z}) \backslash \mathbb{H}^* \xrightarrow{\sim} \mathbb{P}^1(\mathbb{C})$ identifies the compactified moduli space with the Riemann sphere. Adding level structure (replacing $\mathrm{SL}_2(\mathbb{Z})$ by a congruence subgroup $\Gamma_0(N)$, $\Gamma_1(N)$, or $\Gamma(N)$) yields the *modular curves* that will be central to ??.

Exercise 3.4.1

1. Verify the transformation law $G_4(\gamma \cdot \tau) = (c\tau + d)^4 G_4(\tau)$ directly for the two generators of $\mathrm{SL}_2(\mathbb{Z})$: $T = \begin{pmatrix} 1 & 1 \\ 0 & 1 \end{pmatrix}$ (translation $\tau \mapsto \tau + 1$) and $S = \begin{pmatrix} 0 & -1 \\ 1 & 0 \end{pmatrix}$ (inversion $\tau \mapsto -1/\tau$).
2. Using the q -expansions $E_4(\tau) = 1 + 240q + 2160q^2 + \cdots$ and $E_6(\tau) = 1 - 504q - 16632q^2 - \cdots$:
 - (a) Compute $E_4(\tau)^3$ through order q^2 .
 - (b) Compute $E_6(\tau)^2$ through order q^2 .
 - (c) Verify that $E_4^3 - E_6^2 = 1728 \tilde{\Delta}$ where $\tilde{\Delta} = q - 24q^2 + \cdots$, confirming the first terms of the product formula for Δ .
3. Compute $j(\tau)$ for $\tau = (1+i)/2$. (*Hint*: First show that $\Lambda_{(1+i)/2}$ is homothetic to Λ_i by finding the $\mathrm{SL}_2(\mathbb{Z})$ -transformation.)
4. Show that the only zero of $\Delta(\tau)$ on $\mathrm{SL}_2(\mathbb{Z}) \backslash \mathbb{H}^*$ is at the cusp $\tau = i\infty$, and that this zero is simple. Conclude that the modular discriminant is the unique (up to scalar) modular form of weight 12 that vanishes at the cusp.
5. Using the fact that $j : \mathrm{SL}_2(\mathbb{Z}) \backslash \mathbb{H}^* \rightarrow \mathbb{P}^1(\mathbb{C})$ is a bijection, show that for every $j_0 \in \mathbb{C}$, there exists a unique $\tau_0 \in \mathrm{SL}_2(\mathbb{Z}) \backslash \mathbb{H}$ with $j(\tau_0) = j_0$. Compare this with the algebraic surjectivity result proposition 1.2.9.
6. Show that $j(\tau) \in \mathbb{R}$ whenever τ lies on the boundary of the standard fundamental domain $\mathcal{F}$: that is, when $|\tau| = 1$ or $\mathrm{Re}(\tau) = \pm 1/2$. (*Hint*: On $|\tau| = 1$, the map $\tau \mapsto -\bar{\tau} = -1/\bar{\tau}$ preserves Λ_τ up to conjugation, forcing g_2, g_3 to be real.)

3.5 The Analytic Dictionary

The uniformization $\mathbb{C}/\Lambda \cong E(\mathbb{C})$ is not merely a classification tool — it is a dictionary that translates every algebraic notion about elliptic curves into concrete lattice arithmetic. In this section we build the main entries of this dictionary: isogenies become lattice inclusions, the endomorphism ring becomes a subring of $\mathbb{C}$, complex multiplication is the existence of non-integer complex multipliers, the dual isogeny has a clean description, and the Weil pairing is computed by a determinant. The results of this section will be the analytic foundation for the algebraic theory of isogenies in chapter 2 and for the theory of complex multiplication in ??.

Isogenies as Lattice Inclusions

The key principle is that every holomorphic group homomorphism between complex tori lifts to a $\mathbb{C}$ -linear map on the universal covers.

Theorem 3.5.1: Analytic Theory of Isogenies

Let $E_1 = \mathbb{C}/\Lambda_1$ and $E_2 = \mathbb{C}/\Lambda_2$ be elliptic curves over $\mathbb{C}$. There is a bijection

$$\mathrm{Hom}(E_1, E_2) \xleftarrow{1:1} \{\alpha \in \mathbb{C} \mid \alpha\Lambda_1 \subseteq \Lambda_2\},$$

given by sending a holomorphic group homomorphism $\varphi : E_1 \rightarrow E_2$ to the unique $\alpha \in \mathbb{C}$ such that the diagram

$$\begin{array}{ccc} \mathbb{C} & \xrightarrow{\cdot\alpha} & \mathbb{C} \\ \pi_1 \downarrow & & \downarrow \pi_2 \\ \mathbb{C}/\Lambda_1 & \xrightarrow{\varphi} & \mathbb{C}/\Lambda_2 \end{array}$$

commutes, where π_i are the projections. In particular:

1. φ is an isogeny (i.e., surjective with finite kernel) if and only if $\alpha \neq 0$.
2. If φ is an isogeny, $\deg(\varphi) = [\Lambda_2 : \alpha\Lambda_1]$, the index of $\alpha\Lambda_1$ in Λ_2 .
3. $\ker(\varphi) = \alpha^{-1}\Lambda_2/\Lambda_1$ as a subgroup of $\mathbb{C}/\Lambda_1$.

Proof:

Since $\mathbb{C}$ is the universal cover of both $E_i = \mathbb{C}/\Lambda_i$, any holomorphic map $\varphi : E_1 \rightarrow E_2$ with $\varphi(0) = 0$ lifts uniquely to a holomorphic map $\tilde{\varphi} : \mathbb{C} \rightarrow \mathbb{C}$ with $\tilde{\varphi}(0) = 0$ and $\pi_2 \circ \tilde{\varphi} = \varphi \circ \pi_1$. If φ is a group homomorphism, then $\tilde{\varphi}$ satisfies $\tilde{\varphi}(z+w) = \tilde{\varphi}(z) + \tilde{\varphi}(w)$ for all $z, w \in \mathbb{C}$ (lift the identity $\varphi(z+w) = \varphi(z) + \varphi(w)$ to universal covers). A holomorphic function $\tilde{\varphi} : \mathbb{C} \rightarrow \mathbb{C}$ satisfying $\tilde{\varphi}(z+w) = \tilde{\varphi}(z) + \tilde{\varphi}(w)$ and $\tilde{\varphi}(0) = 0$ is $\mathbb{C}$ -linear: differentiating with respect to w and setting $w = 0$ gives $\tilde{\varphi}'(z) = \tilde{\varphi}'(0) = \alpha$ for all z , so $\tilde{\varphi}(z) = \alpha z$.

The condition that $\tilde{\varphi}$ descend to a map $E_1 \rightarrow E_2$ is exactly $\tilde{\varphi}(\Lambda_1) \subseteq \Lambda_2$, i.e., $\alpha\Lambda_1 \subseteq \Lambda_2$.

(1) If $\alpha = 0$, then φ is the zero map (not surjective). If $\alpha \neq 0$, then $\tilde{\varphi}(z) = \alpha z$ is surjective on $\mathbb{C}$, hence φ is surjective on $\mathbb{C}/\Lambda_2$. The kernel $\ker(\varphi) = \{z \in \mathbb{C}/\Lambda_1 : \alpha z \in \Lambda_2\} = \alpha^{-1}\Lambda_2/\Lambda_1$ is a finite group (since $\alpha\Lambda_1 \subseteq \Lambda_2$ implies $\Lambda_1 \subseteq \alpha^{-1}\Lambda_2$, and both are lattices in $\mathbb{C}$, so the quotient is finite).

(2) The degree of φ equals $|\ker(\varphi)| = |\alpha^{-1}\Lambda_2/\Lambda_1| = [\alpha^{-1}\Lambda_2 : \Lambda_1]$. Since $\alpha^{-1}\Lambda_2 \supset \Lambda_1$ with finite index, and multiplication by α gives an isomorphism $\alpha^{-1}\Lambda_2/\Lambda_1 \cong \Lambda_2/\alpha\Lambda_1$, we get $\deg(\varphi) = [\Lambda_2 : \alpha\Lambda_1]$.

(3) is immediate from the construction: $\ker(\varphi) = \{z + \Lambda_1 : \alpha z \in \Lambda_2\} = \alpha^{-1}\Lambda_2/\Lambda_1$.

The beauty of this result is that an isogeny — a priori a complicated algebraic morphism defined

by rational functions on the curve — reduces to multiplication by a single complex number. The algebraic complexity is hidden in the lattice inclusion $\alpha\Lambda_1 \subseteq \Lambda_2$.

Example 3.5.2: The Multiplication-by- n Map

For any lattice Λ and positive integer n , the map $[n] : \mathbb{C}/\Lambda \rightarrow \mathbb{C}/\Lambda$ given by $z \mapsto nz$ corresponds to $\alpha = n$. Since $n\Lambda \subseteq \Lambda$, this is an endomorphism of $E = \mathbb{C}/\Lambda$. Its kernel is

$$\ker([n]) = n^{-1}\Lambda/\Lambda = \frac{1}{n}\Lambda/\Lambda \cong (\mathbb{Z}/n\mathbb{Z})^2,$$

confirming $\deg([n]) = [\Lambda : n\Lambda] = n^2$ (the index of $n\Lambda$ in $\Lambda = \mathbb{Z}\omega_1 + \mathbb{Z}\omega_2$ is n^2 since each generator is scaled by n).

Example 3.5.3: An Isogeny of Degree 2

Consider the lattice $\Lambda = \mathbb{Z} + \mathbb{Z}\tau$ and the sublattice $\Lambda' = \mathbb{Z} + \mathbb{Z}(2\tau)$. Then $\Lambda' \subseteq \Lambda$ with index $[\Lambda : \Lambda'] = 2$. The inclusion $\Lambda' \hookrightarrow \Lambda$ corresponds to an isogeny $\varphi : \mathbb{C}/\Lambda' \rightarrow \mathbb{C}/\Lambda$ with $\alpha = 1$ (the identity on $\mathbb{C}$), kernel $\Lambda/\Lambda' \cong \mathbb{Z}/2\mathbb{Z}$ (generated by $\tau + \Lambda'$), and degree 2. The curves $E' = \mathbb{C}/\Lambda'$ and $E = \mathbb{C}/\Lambda$ have the same x -coordinate function $\wp(z; \Lambda)$ restricted to $\mathbb{C}/\Lambda'$... but the period ratio for Λ' is 2τ , not τ , so E' is a different elliptic curve with $j(E') = j(2\tau) \neq j(\tau)$ in general.

Alternatively, we can take $\alpha = 2$ and $\Lambda_1 = \Lambda_2 = \Lambda$, giving the multiplication-by-2 map, which has degree 4 and kernel $(\mathbb{Z}/2\mathbb{Z})^2$. The degree-2 isogeny above is genuinely different: it sends one elliptic curve to a *different* one.

The Endomorphism Ring

Definition 3.5.4: Endomorphism Ring (Analytic)

The *endomorphism ring* of $E = \mathbb{C}/\Lambda$ is

$$\text{End}(E) = \{\alpha \in \mathbb{C} \mid \alpha\Lambda \subseteq \Lambda\}.$$

This is a subring of $\mathbb{C}$ (closed under addition and multiplication, and containing 1).

Every lattice satisfies $n\Lambda \subseteq \Lambda$ for all $n \in \mathbb{Z}$, so $\mathbb{Z} \subseteq \text{End}(E)$ always. The question is whether $\text{End}(E)$ is strictly larger.

Theorem 3.5.5: Classification of Endomorphism Rings

Let $E = \mathbb{C}/\Lambda$ be an elliptic curve over $\mathbb{C}$. Then exactly one of the following holds:

1. (*Generic case.*) $\text{End}(E) = \mathbb{Z}$.
2. (*CM case.*) $\text{End}(E)$ is an order in an imaginary quadratic field $K = \mathbb{Q}(\sqrt{-d})$ for some squarefree positive integer d . In this case, we say E has *complex multiplication* (CM) by $\text{End}(E)$.

Case (2) occurs if and only if the period ratio τ of Λ is a quadratic irrationality (i.e., τ satisfies a quadratic equation $a\tau^2 + b\tau + c = 0$ with $a, b, c \in \mathbb{Z}$ and $a \neq 0$).

Proof:

Suppose $\alpha \in \text{End}(E) \setminus \mathbb{Z}$, so $\alpha\Lambda \subseteq \Lambda$ with $\alpha \notin \mathbb{Z}$. Writing $\Lambda = \mathbb{Z}\omega_1 + \mathbb{Z}\omega_2$ with $\tau = \omega_1/\omega_2$:

$$\alpha\omega_1 = a\omega_1 + b\omega_2$$

$$\alpha\omega_2 = c\omega_1 + d\omega_2$$

for some $a, b, c, d \in \mathbb{Z}$. Dividing the first equation by ω_2 : $\alpha\tau = a\tau + b$. Dividing the second by ω_2 : $\alpha = c\tau + d$. Substituting:

$$(c\tau + d)\tau = a\tau + b, \quad \text{i.e.,} \quad c\tau^2 + (d - a)\tau - b = 0.$$

Since $\alpha = c\tau + d \notin \mathbb{Z}$ and $\tau \notin \mathbb{R}$, we must have $c \neq 0$ (else $\alpha = d \in \mathbb{Z}$). So τ satisfies a quadratic equation with integer coefficients and $c \neq 0$. Since $\tau \in \mathbb{H}$ (hence $\tau \notin \mathbb{R}$), the discriminant $(d - a)^2 + 4bc < 0$, so $\mathbb{Q}(\tau)$ is an imaginary quadratic field.

Conversely, if τ is a quadratic irrationality, say $c\tau^2 + (d - a)\tau - b = 0$, then $\alpha = c\tau + d$ satisfies $\alpha\Lambda \subseteq \Lambda$ (as the computation above reverses), so $\alpha \in \text{End}(E) \setminus \mathbb{Z}$.

Now, $\text{End}(E)$ is a subring of $\mathbb{C}$ containing $\mathbb{Z}$. If $\alpha \in \text{End}(E) \setminus \mathbb{Z}$, then $\alpha = c\tau + d$ for some $c \neq 0$, $d \in \mathbb{Z}$, and $\mathbb{Q}(\alpha) = \mathbb{Q}(\tau) = K$ is an imaginary quadratic field. Every element of $\text{End}(E)$ lies in K (by the same argument), so $\text{End}(E) \subseteq K$. Since $\text{End}(E)$ is a subring of K containing $\mathbb{Z}$ and at least one element not in $\mathbb{Z}$, it is an *order* in K : a subring of the ring of integers $\mathcal{O}_K$ that is a free $\mathbb{Z}$ -module of rank 2.

The full ring of integers $\mathcal{O}_K$ is the maximal order. Every order has the form $\mathbb{Z} + f\mathcal{O}_K$ for a positive integer f (the *conductor* of the order). The endomorphism ring $\text{End}(E)$ is an order of conductor f in $K = \mathbb{Q}(\sqrt{-d})$.

Example 3.5.6: CM of the Running Examples

(a) $E : y^2 = x^3 - x$ ($\Lambda = \mathbb{Z} + \mathbb{Z}i$, $\tau = i$). The element $\alpha = i$ satisfies $i\Lambda = i\mathbb{Z} + i \cdot \mathbb{Z}i = \mathbb{Z}i - \mathbb{Z} = \Lambda$. So $i \in \text{End}(E)$, and $\text{End}(E) \supseteq \mathbb{Z}[i]$. Since $\mathbb{Z}[i] = \mathcal{O}_{\mathbb{Q}(i)}$ is the full ring of integers of $\mathbb{Q}(i) = \mathbb{Q}(\sqrt{-1})$, the endomorphism ring is $\text{End}(E) = \mathbb{Z}[i]$, the Gaussian integers. The conductor is $f = 1$.

(b) $E : y^2 = x^3 + 1$ ($\Lambda = \mathbb{Z} + \mathbb{Z}\omega$, $\tau = \omega = e^{2\pi i/3}$). The element $\alpha = \omega$ satisfies $\omega\Lambda = \mathbb{Z}\omega + \mathbb{Z}\omega^2 = \mathbb{Z}\omega + \mathbb{Z}(-1 - \omega) = \Lambda$ (using $\omega^2 = -1 - \omega$). So $\text{End}(E) = \mathbb{Z}[\omega]$, the Eisenstein integers, which is $\mathcal{O}_{\mathbb{Q}(\omega)} = \mathcal{O}_{\mathbb{Q}(\sqrt{-3})}$.

(c) A generic curve. For $\tau = 1 + \sqrt{2}i$ (an element of $\mathbb{H}$ that is not a quadratic irrationality over $\mathbb{Q}$), the lattice Λ_τ has $\text{End}(\mathbb{C}/\Lambda_\tau) = \mathbb{Z}$: no extra endomorphisms.

The CM curves are exceptional: they form a countable subset of the moduli space (one for each order in each imaginary quadratic field). Yet they are arithmetically the most interesting curves, and the theory of complex multiplication — which describes how the j -invariants and torsion points of CM curves generate class fields of imaginary quadratic fields — is one of the crowning achievements of algebraic number theory. We will develop this in ??.

The Dual Isogeny

Given an isogeny $\varphi : E_1 \rightarrow E_2$ corresponding to $\alpha \in \mathbb{C}$ with $\alpha\Lambda_1 \subseteq \Lambda_2$, there is a canonical isogeny in the reverse direction.

Proposition 3.5.7: The Dual Isogeny (Analytic)

Let $\varphi : E_1 \rightarrow E_2$ be an isogeny of degree n , corresponding to $\alpha \in \mathbb{C}^*$ with $\alpha\Lambda_1 \subseteq \Lambda_2$. The dual isogeny $\hat{\varphi} : E_2 \rightarrow E_1$ is the isogeny corresponding to $\hat{\alpha} = n/\alpha$. It satisfies:

1. $\hat{\alpha}\Lambda_2 \subseteq \Lambda_1$, so $\hat{\varphi}$ is a well-defined isogeny.
2. $\hat{\varphi} \circ \varphi = [n]$ on E_1 (since $\hat{\alpha} \cdot \alpha = n$).
3. $\varphi \circ \hat{\varphi} = [n]$ on E_2 .
4. $\deg(\hat{\varphi}) = \deg(\varphi) = n$.
5. $\widehat{\hat{\varphi}} = \varphi$.

Proof:

(1) We need $(n/\alpha)\Lambda_2 \subseteq \Lambda_1$. Since $\alpha\Lambda_1 \subseteq \Lambda_2$, we have $\Lambda_2 \supseteq \alpha\Lambda_1$, so $(n/\alpha)\Lambda_2 \supseteq (n/\alpha)(\alpha\Lambda_1) = n\Lambda_1 \subseteq \Lambda_1$. But we need the inclusion in the other direction. Consider the chain $\alpha\Lambda_1 \subseteq \Lambda_2$, which gives $n\Lambda_1 = \alpha \cdot (n/\alpha)\Lambda_1 \subseteq \alpha\Lambda_1$.?... let us argue differently.

Since $n = \deg(\varphi) = [\Lambda_2 : \alpha\Lambda_1]$, we have $n\Lambda_2 \subseteq \alpha\Lambda_1$ (because n annihilates the quotient $\Lambda_2/\alpha\Lambda_1$). Thus $(n/\alpha)\Lambda_2 \subseteq \Lambda_1$.

(2) The composition $\hat{\varphi} \circ \varphi$ corresponds to $\hat{\alpha} \cdot \alpha = (n/\alpha) \cdot \alpha = n$, which is the multiplication-by- n map.

(3) Similarly, $\varphi \circ \hat{\varphi}$ corresponds to $\alpha \cdot (n/\alpha) = n$.

(4) $\deg(\hat{\varphi}) = [\Lambda_1 : \hat{\alpha}\Lambda_2] = [\Lambda_1 : (n/\alpha)\Lambda_2]$. From $\hat{\varphi} \circ \varphi = [n]$ of degree n^2 and $\deg(\hat{\varphi}) \deg(\varphi) = \deg([n]) = n^2$, we get $\deg(\hat{\varphi}) = n^2/n = n$.

(5) The dual of $\hat{\varphi}$ corresponds to $n/\hat{\alpha} = n/(n/\alpha) = \alpha$, which is φ .

The dual isogeny will be constructed algebraically (without the crutch of uniformization) in chapter 2. The algebraic construction is considerably more involved — it requires working with the pushforward of divisors and the theory of function fields. The analytic construction, by contrast, is a one-liner: $\hat{\alpha} = n/\alpha$.

Example 3.5.8: Dual of the Multiplication-by- n Map

For $[n] : E \rightarrow E$ with $\alpha = n$, the dual is $\widehat{[n]}$ with $\hat{\alpha} = n/n = 1 \cdot n/n \dots$ wait, $\hat{\alpha} = \deg([n])/\alpha = n^2/n = n$. So $\widehat{[n]} = [n]$: the multiplication-by- n map is self-dual.

The Weil Pairing

The Weil pairing is a bilinear form on the n -torsion of an elliptic curve, taking values in the n -th roots of unity. Its analytic construction is particularly transparent.

Definition 3.5.9: The Weil Pairing (Analytic)

Let $E = \mathbb{C}/\Lambda$ with $\Lambda = \mathbb{Z}\omega_1 + \mathbb{Z}\omega_2$, and let $n \geq 1$. The n -torsion is $E[n] = \frac{1}{n}\Lambda/\Lambda$, with generators $P = \omega_1/n + \Lambda$ and $Q = \omega_2/n + \Lambda$. The *Weil pairing* is the map

$$e_n : E[n] \times E[n] \longrightarrow \mu_n, \quad e_n(P, Q) = e^{2\pi i/n}, \quad (3.26)$$

extended bilinearly. Explicitly, for $P = (a\omega_1 + b\omega_2)/n$ and $Q = (c\omega_1 + d\omega_2)/n$ with $a, b, c, d \in \mathbb{Z}$:

$$e_n(P, Q) = e^{2\pi i(ad-bc)/n} = \zeta_n^{ad-bc}, \quad (3.27)$$

where $\zeta_n = e^{2\pi i/n}$ is a primitive n -th root of unity.

The formula (3.27) shows that e_n is the exponential of the standard symplectic form on $\Lambda/n\Lambda \cong (\mathbb{Z}/n\mathbb{Z})^2$: the “determinant” $ad - bc$ is the symplectic intersection pairing.

Proposition 3.5.10: Properties of the Weil Pairing

The Weil pairing satisfies:

1. (*Bilinearity.*) $e_n(P_1 + P_2, Q) = e_n(P_1, Q) \cdot e_n(P_2, Q)$, and similarly in the second variable.
2. (*Alternating.*) $e_n(P, P) = 1$ for all $P \in E[n]$, and $e_n(P, Q) = e_n(Q, P)^{-1}$.
3. (*Nondegeneracy.*) If $e_n(P, Q) = 1$ for all $Q \in E[n]$, then $P = O$.
4. (*Galois equivariance.*) For any isogeny $\varphi : E_1 \rightarrow E_2$ and $P \in E_1[n]$, $Q \in E_1[n]$: $e_n(\varphi(P), \varphi(Q)) = e_n(P, Q)^{\deg \varphi}$. In particular, for $\sigma \in \text{Aut}(E)$: $e_n(\sigma P, \sigma Q) = e_n(P, Q)$.
5. (*Compatibility.*) For $m|n$: $e_n(P, Q)^{n/m} = e_m([n/m]P, [n/m]Q)$ for $P, Q \in E[n]$.

Proof:

All properties follow directly from the explicit formula (3.27).

- (1) If $P_1 = (a_1\omega_1 + b_1\omega_2)/n$ and $P_2 = (a_2\omega_1 + b_2\omega_2)/n$, then $P_1 + P_2 = ((a_1 + a_2)\omega_1 + (b_1 + b_2)\omega_2)/n$, and

$$e_n(P_1 + P_2, Q) = \zeta_n^{(a_1 + a_2)d - (b_1 + b_2)c} = \zeta_n^{a_1d - b_1c} \cdot \zeta_n^{a_2d - b_2c} = e_n(P_1, Q) \cdot e_n(P_2, Q).$$

- (2) $e_n(P, P) = \zeta_n^{ad - ba} = \zeta_n^0 = 1$ (writing $P = (a\omega_1 + b\omega_2)/n$), since $ad - ba = 0$. And $e_n(Q, P) = \zeta_n^{ca - db} = \zeta_n^{-(ad - bc)} = e_n(P, Q)^{-1}$.

- (3) If $e_n(P, Q) = 1$ for all Q , then $ad - bc \equiv 0 \pmod{n}$ for all $c, d \in \mathbb{Z}$. Taking $Q = \omega_1/n$ (so $c = 1, d = 0$): $-b \equiv 0 \pmod{n}$, hence $n|b$. Taking $Q = \omega_2/n$ (so $c = 0, d = 1$): $a \equiv 0 \pmod{n}$, hence $n|a$. So $P = (a\omega_1 + b\omega_2)/n \in \Lambda$, i.e., $P = O$ in $E[n]$.

- (4) For an isogeny φ corresponding to $\alpha \in \mathbb{C}$: φ maps $E_1[n] \rightarrow E_2[n]$, and the lattice inclusion $\alpha\Lambda_1 \subseteq \Lambda_2$ induces a map on $\frac{1}{n}\Lambda_1/\Lambda_1 \rightarrow \frac{1}{n}\Lambda_2/\Lambda_2$. Tracking the effect of α on the symplectic form: if $\Lambda_1 = \mathbb{Z}\omega_1 + \mathbb{Z}\omega_2$, then $\alpha\omega_1 = a'\omega'_1 + b'\omega'_2$ in terms of a basis (ω'_1, ω'_2) for Λ_2 , and the determinant of the change-of-basis matrix is $\deg(\varphi)$, giving the stated formula.

- (5) follows from $[n/m]P = (a\omega_1 + b\omega_2)/m$ when $P = (a\omega_1 + b\omega_2)/n$, and $e_m([n/m]P, [n/m]Q) = \zeta_m^{ad - bc} = (\zeta_n^{n/m})^{ad - bc} = e_n(P, Q)^{n/m}$.

The Weil pairing will be constructed algebraically in chapter 2, where its Galois-equivariance has deep arithmetic consequences: it implies that the determinant of the Galois representation $\rho_{E, \ell} : \text{Gal}(\overline{\mathbb{Q}}/\mathbb{Q}) \rightarrow \text{GL}_2(\mathbb{Z}_\ell)$ (attached to the Tate module) is the cyclotomic character. The analytic construction above makes this transparent: the Galois group acts on the “coordinates” (a, b) of a torsion point, and the Weil pairing is the determinant of the resulting matrix.

Example 3.5.11: The Weil Pairing on $E[2]$

For $n = 2$ and $E = \mathbb{C}/(\mathbb{Z}\omega_1 + \mathbb{Z}\omega_2)$, the 2-torsion is $E[2] = \{O, \omega_1/2, \omega_2/2, (\omega_1 + \omega_2)/2\}$. With $P = \omega_1/2$ and $Q = \omega_2/2$:

$$e_2(P, Q) = e^{2\pi i(1 \cdot 1 - 0 \cdot 0)/2} = e^{\pi i} = -1.$$

Since $\mu_2 = \{1, -1\}$, the Weil pairing is nontrivial. In the standard basis (P, Q) , the pairing matrix is $e_2(P, Q) = -1$, confirming nondegeneracy.

Summary: The Analytic–Algebraic Dictionary

We collect the main entries of the dictionary established in this chapter.

Algebraic (Chapter 1)	Analytic (Chapter 2)
Elliptic curve $E/\mathbb{C}$	Complex torus $\mathbb{C}/\Lambda$
Group law $P + Q$	Addition $z_1 + z_2 \pmod{\Lambda}$
n -torsion $E[n] \cong (\mathbb{Z}/n\mathbb{Z})^2$	$\frac{1}{n}\Lambda/\Lambda$
Isogeny $\varphi : E_1 \rightarrow E_2$	$\alpha \in \mathbb{C}$ with $\alpha\Lambda_1 \subseteq \Lambda_2$
Degree of φ	$[\Lambda_2 : \alpha\Lambda_1]$
Dual isogeny $\hat{\varphi}$	$\hat{\alpha} = \deg(\varphi)/\alpha$
$\text{End}(E) = \mathbb{Z}$ (generic)	τ is not a quadratic irrationality
$\text{End}(E) = \mathcal{O}$ (CM by order $\mathcal{O}$)	$\tau \in K = \mathbb{Q}(\sqrt{-d})$, $\mathcal{O} = \mathbb{Z} + f\mathcal{O}_K$
j -invariant $j(E)$	$j(\tau) = 1728 g_2^3/\Delta$
Discriminant Δ	$g_2^3 - 27g_3^2$; product formula
Weil pairing e_n	ζ_n^{ad-bc} (symplectic form)
Weierstrass equation	$(\wp')^2 = 4\wp^3 - g_2\wp - g_3$
Invariant differential $dx/(2y)$	dz

This dictionary will be used freely in subsequent chapters. When a concept is clearer from the analytic perspective, we will compute there and translate back; when the algebraic formulation is needed (over fields other than $\mathbb{C}$), we will develop it independently.

Exercise 3.5.1

- Let $\Lambda = \mathbb{Z} + \mathbb{Z}\tau$ and $\alpha = 2 + i$ (with $\tau = i$, so $\Lambda = \mathbb{Z} + \mathbb{Z}i$).
 - Verify that $\alpha\Lambda \subseteq \Lambda$.
 - Compute $[\Lambda : \alpha\Lambda]$.
 - Describe the kernel $\ker(\varphi)$ as a subgroup of $E[5]$.
- Determine $\text{End}(E)$ for the elliptic curve $E = \mathbb{C}/\Lambda_\tau$ when:
 - $\tau = \sqrt{-2}$.
 - $\tau = (1 + \sqrt{-7})/2$.

(c) $\tau = 2i$.

In each case, identify the imaginary quadratic field and the conductor of the order.

3. Let $\varphi : E_1 \rightarrow E_2$ be the isogeny corresponding to $\alpha = 1 + i$ with $\Lambda_1 = \Lambda_2 = \mathbb{Z} + \mathbb{Z}i$.

(a) Compute $\deg(\varphi)$ and $\ker(\varphi)$.

(b) Write down $\hat{\alpha}$ for the dual isogeny, and verify $\hat{\alpha} \cdot \alpha = \deg(\varphi)$.

(c) Verify that $\hat{\alpha}\Lambda_2 \subseteq \Lambda_1$.

4. Compute the Weil pairing $e_3(P, Q)$ for $P = \tau/3 + \Lambda$ and $Q = 1/3 + \Lambda$ on the curve $E = \mathbb{C}/(\mathbb{Z}\tau + \mathbb{Z})$. Verify that $e_3(P, Q)$ is a primitive cube root of unity.

5. Let $\varphi : E_1 \rightarrow E_2$ and $\psi : E_2 \rightarrow E_3$ be isogenies corresponding to α and β respectively. Show that $\psi \circ \varphi$ corresponds to $\beta\alpha$, and that $\deg(\psi \circ \varphi) = \deg(\psi) \cdot \deg(\varphi)$.

6. Let $E = \mathbb{C}/\Lambda$ have CM by the order $\mathcal{O} = \mathbb{Z}[\sqrt{-n}]$ for a squarefree $n \geq 1$.

(a) Show that Λ is homothetic to the lattice $\mathbb{Z} + \mathbb{Z}\sqrt{-n}$.

(b) Determine the conductor of $\mathcal{O}$ in $\mathcal{O}_{\mathbb{Q}(\sqrt{-n})}$ for $n = 1, 2, 3, 5, 7$.

(c) For which values of n is $\mathcal{O}$ the maximal order?

3.6 The Modular Group and the Fundamental Domain

We have repeatedly used the fact that the modular group $\mathrm{SL}_2(\mathbb{Z})$ acts on $\mathbb{H}$, and that the quotient $\mathrm{SL}_2(\mathbb{Z}) \backslash \mathbb{H}$ parametrizes elliptic curves up to isomorphism. In this final section of the chapter, we study the group $\mathrm{SL}_2(\mathbb{Z})$ itself — its generators, its fundamental domain, and the structure of the quotient as a Riemann surface. The punchline is that $\mathrm{SL}_2(\mathbb{Z}) \backslash \mathbb{H}^*$ has genus 0 (confirming that the j -function is an isomorphism to $\mathbb{P}^1$), and the two orbifold points in the quotient correspond precisely to the two exceptional j -values ($j = 1728$ and $j = 0$) where elliptic curves have extra automorphisms.

The Modular Group

Definition 3.6.1: The Modular Group

The *modular group* is $\mathrm{SL}_2(\mathbb{Z}) = \left\{ \begin{pmatrix} a & b \\ c & d \end{pmatrix} \mid a, b, c, d \in \mathbb{Z}, ad - bc = 1 \right\}$. Its quotient by the center $\{\pm I\}$ is the *projective modular group* $\mathrm{PSL}_2(\mathbb{Z}) = \mathrm{SL}_2(\mathbb{Z})/\{\pm I\}$.

The action of $\mathrm{SL}_2(\mathbb{Z})$ on $\mathbb{H}$ by Möbius transformations $\gamma \cdot \tau = (a\tau + b)/(c\tau + d)$ has kernel $\{\pm I\}$ (since $(-I) \cdot \tau = (-\tau - 0)/(0 - 1 \cdot \tau) = \tau$ — more directly, γ and $-\gamma$ induce the same transformation). The effective action is therefore by $\mathrm{PSL}_2(\mathbb{Z})$.

The group $\mathrm{SL}_2(\mathbb{Z})$ is generated by two elements with simple geometric meanings:

Proposition 3.6.2: Generators of $\mathrm{SL}_2(\mathbb{Z})$

The modular group $\mathrm{SL}_2(\mathbb{Z})$ is generated by

$$T = \begin{pmatrix} 1 & 1 \\ 0 & 1 \end{pmatrix} \quad (\text{translation : } \tau \mapsto \tau+1) \quad \text{and} \quad S = \begin{pmatrix} 0 & -1 \\ 1 & 0 \end{pmatrix} \quad (\text{inversion : } \tau \mapsto -1/\tau).$$

They satisfy the relations $S^2 = -I$ and $(ST)^3 = -I$, so in $\mathrm{PSL}_2(\mathbb{Z})$ the images $\bar{S}$ and $\bar{T}$ satisfy $\bar{S}^2 = 1$ and $(\bar{S}\bar{T})^3 = 1$. In fact, $\mathrm{PSL}_2(\mathbb{Z}) \cong \mathbb{Z}/2\mathbb{Z} * \mathbb{Z}/3\mathbb{Z}$, the free product of a cyclic group of order 2 (generated by $\bar{S}$) and a cyclic group of order 3 (generated by $\bar{S}\bar{T}$).

Proof Sketch:

The relations are verified directly: $S^2 = \begin{pmatrix} 0 & -1 \\ 1 & 0 \end{pmatrix}^2 = \begin{pmatrix} -1 & 0 \\ 0 & -1 \end{pmatrix} = -I$, and

$$ST = \begin{pmatrix} 0 & -1 \\ 1 & 0 \end{pmatrix} \begin{pmatrix} 1 & 1 \\ 0 & 1 \end{pmatrix} = \begin{pmatrix} 0 & -1 \\ 1 & 1 \end{pmatrix}, \quad (ST)^3 = -I$$

(the last can be checked by computing $(ST)^2 = \begin{pmatrix} -1 & -1 \\ 1 & 0 \end{pmatrix}$ and $(ST)^3 = \begin{pmatrix} -1 & 0 \\ 0 & -1 \end{pmatrix} = -I$).

That S and T generate $\mathrm{SL}_2(\mathbb{Z})$ will follow from the construction of the fundamental domain below: the algorithm that moves any $\tau \in \mathbb{H}$ into the fundamental domain uses only T (and T^{-1}) and S , and the proof that it terminates shows that the corresponding $\mathrm{SL}_2(\mathbb{Z})$ -element is a word in S and T . The free product structure $\mathrm{PSL}_2(\mathbb{Z}) \cong \mathbb{Z}/2\mathbb{Z} * \mathbb{Z}/3\mathbb{Z}$ follows from the theory of groups acting on trees (Bass–Serre theory) applied to the action of $\mathrm{PSL}_2(\mathbb{Z})$ on the Farey graph; see Serre [17].

The generator T acts by horizontal translation ($\tau \mapsto \tau + 1$), while S acts by inversion in the unit circle composed with reflection ($\tau \mapsto -1/\tau$, which maps τ to the point diametrically opposite on the unit circle and then reflects). The composition $ST : \tau \mapsto -1/(\tau + 1)$ is a rotation of order 3 around the fixed point $\tau = e^{2\pi i/3}$.

The Fundamental Domain

Theorem 3.6.3: Fundamental Domain for $\mathrm{SL}_2(\mathbb{Z})$

Define the region

$$\mathcal{F} = \{\tau \in \mathbb{H} \mid |\tau| \geq 1 \text{ and } |\mathrm{Re}(\tau)| \leq \tfrac{1}{2}\}. \quad (3.28)$$

Then:

1. Every $\tau \in \mathbb{H}$ is $\mathrm{SL}_2(\mathbb{Z})$ -equivalent to some $\tau' \in \mathcal{F}$.
2. If $\tau, \tau' \in \mathcal{F}$ satisfy $\tau' = \gamma \cdot \tau$ for some $\gamma \in \mathrm{SL}_2(\mathbb{Z})$ with $\gamma \neq \pm I$, then τ and τ' lie on the boundary of $\mathcal{F}$. Specifically, either:
 - (a) $\mathrm{Re}(\tau) = -1/2$ and $\tau' = \tau + 1$ (identified by T), or
 - (b) $|\tau| = 1$ and $\tau' = -1/\tau$ (identified by S).

Thus $\mathcal{F}$ is a fundamental domain for the action of $\mathrm{SL}_2(\mathbb{Z})$ on $\mathbb{H}$: it contains exactly one representative from each orbit, except that boundary points are identified pairwise.

Proof:

(1) Given $\tau \in \mathbb{H}$, we produce $\tau' \in \mathcal{F}$ by the following algorithm.

Step A. Apply T^n for an appropriate $n \in \mathbb{Z}$ to ensure $|\mathrm{Re}(\tau)| \leq 1/2$. (Choose $n = -\lfloor \mathrm{Re}(\tau) + 1/2 \rfloor$.)

Step B. If $|\tau| < 1$, apply S : replace τ by $-1/\tau$. Since $\mathrm{Im}(-1/\tau) = \mathrm{Im}(\tau)/|\tau|^2 > \mathrm{Im}(\tau)$ (because $|\tau| < 1$), this strictly increases the imaginary part.

Step C. Return to Step A.

The key observation is that $\mathrm{Im}(\gamma \cdot \tau) = \mathrm{Im}(\tau)/|c\tau + d|^2$ for $\gamma = \begin{pmatrix} a & b \\ c & d \end{pmatrix}$. For a fixed τ , the set of values $|c\tau + d|^2$ for $(c, d) \in \mathbb{Z}^2 \setminus \{(0, 0)\}$ is a discrete subset of $\mathbb{R}_{>0}$ (it is the set of values of a positive definite quadratic form in (c, d)), so $\mathrm{Im}(\gamma \cdot \tau)$ achieves a maximum over all $\gamma \in \mathrm{SL}_2(\mathbb{Z})$. The algorithm terminates because Step B strictly increases $\mathrm{Im}(\tau)$ whenever it is applied, and the imaginary part is bounded above.

When the algorithm terminates, we have $|\mathrm{Re}(\tau)| \leq 1/2$ (from Step A) and $|\tau| \geq 1$ (since Step B was not triggered), so $\tau \in \mathcal{F}$.

(2) The uniqueness statement requires showing that if $\tau, \tau' \in \mathcal{F}^\circ$ (the interior) satisfy $\tau' = \gamma \cdot \tau$ with $\gamma \neq \pm I$, we reach a contradiction. Writing $\gamma = \begin{pmatrix} a & b \\ c & d \end{pmatrix}$, the condition $\tau' \in \mathcal{F}^\circ$ gives $\mathrm{Im}(\tau') > \sqrt{3}/2$. Since $\mathrm{Im}(\tau') = \mathrm{Im}(\tau)/|c\tau + d|^2$ and $\mathrm{Im}(\tau) > \sqrt{3}/2$, we need $|c\tau + d|^2 < 2$. For τ in the interior of $\mathcal{F}$, one checks that $|c\tau + d| \geq 1$ for all $(c, d) \neq (0, 0)$, with equality only on the boundary arcs. A case analysis on c (the key cases being $c = 0$, $|c| = 1$) shows that interior points have no nontrivial identifications. The boundary identifications are exactly those

described in (2a) and (2b).

The fundamental domain $\mathcal{F}$ is the hyperbolic triangle with vertices at $e^{\pi i/3}$, $e^{2\pi i/3}$, and $i\infty$, with vertical sides $\operatorname{Re}(\tau) = \pm 1/2$ identified by T , and the circular arc $|\tau| = 1$ identified by S . After identification, the boundary of $\mathcal{F}$ folds into a closed surface: the quotient $\operatorname{SL}_2(\mathbb{Z}) \backslash \mathbb{H}$.

Stabilizers and Orbifold Points

At a generic point $\tau \in \mathcal{F}^\circ$, the stabilizer $\operatorname{Stab}_{\operatorname{PSL}_2(\mathbb{Z})}(\tau) = \{1\}$ is trivial. At two special points, the stabilizer is larger:

Proposition 3.6.4: Orbifold Points

The points of $\mathcal{F}$ with nontrivial stabilizer in $\operatorname{PSL}_2(\mathbb{Z})$ are:

1. $\tau = i$, with stabilizer $\langle \bar{S} \rangle \cong \mathbb{Z}/2\mathbb{Z}$. The generator S acts as $\tau \mapsto -1/\tau$, which fixes i (since $-1/i = i$).
2. $\tau = \rho = e^{2\pi i/3} = (-1 + \sqrt{-3})/2$, with stabilizer $\langle \overline{ST} \rangle \cong \mathbb{Z}/3\mathbb{Z}$. The generator ST acts as $\tau \mapsto -1/(\tau + 1) = \tau$ (one verifies $-1/(\rho + 1) = -1/(-\rho^2) = 1/\rho^2 = \rho$, using $\rho^2 + \rho + 1 = 0$).

These are the only orbifold points of $\operatorname{SL}_2(\mathbb{Z}) \backslash \mathbb{H}$.

Proof:

A point $\tau \in \mathbb{H}$ has nontrivial stabilizer iff $\gamma \cdot \tau = \tau$ for some $\gamma = \begin{pmatrix} a & b \\ c & d \end{pmatrix} \in \operatorname{SL}_2(\mathbb{Z})$ with $\gamma \neq \pm I$. This gives $(a\tau + b)/(c\tau + d) = \tau$, i.e., $c\tau^2 + (d - a)\tau - b = 0$. For $c \neq 0$, this is a quadratic with discriminant $(d - a)^2 + 4bc = (d + a)^2 - 4(ad - bc) = (a + d)^2 - 4$. Since $\tau \in \mathbb{H}$, the discriminant must be negative: $(a + d)^2 < 4$, so $|a + d| \leq 1$. The trace $a + d$ determines the type of fixed point:

- $a + d = 0$: discriminant -4 , $\tau = (a - d \pm 2i)/(2c)$. With $d = -a$ and $ad - bc = 1$: $-a^2 - bc = 1$. For example, $\gamma = S$ gives $a = d = 0$, $b = -1$, $c = 1$, $\tau = \pm i/(2 \cdot 1)$... more carefully: $c\tau^2 - b = 0$ gives $\tau^2 = b/c = -1$, so $\tau = i$. The stabilizer element has order 2 in $\operatorname{PSL}_2(\mathbb{Z})$.
- $a + d = \pm 1$: discriminant -3 , $\tau = ((a - d) \pm \sqrt{-3})/(2c)$. For $\gamma = ST = \begin{pmatrix} 0 & -1 \\ 1 & 1 \end{pmatrix}$: $a + d = 1$, $c = 1$, and $\tau = (-1 + \sqrt{-3})/2 = \rho$. The stabilizer element has order 3 in $\operatorname{PSL}_2(\mathbb{Z})$.

No other traces are possible (since $|a + d| \leq 1$), and any τ with nontrivial stabilizer is $\operatorname{SL}_2(\mathbb{Z})$ -equivalent to i or ρ .

Connection to Automorphisms of Elliptic Curves

The orbifold structure has a beautiful interpretation via the uniformization:

Proposition 3.6.5: Stabilizers = Automorphism Groups

For $\tau \in \mathbb{H}$, the stabilizer $\text{Stab}_{\text{PSL}_2(\mathbb{Z})}(\tau)$ is isomorphic to $\text{Aut}(E_\tau)/\langle [-1] \rangle$, where $E_\tau = \mathbb{C}/\Lambda_\tau$. In particular:

1. If τ is a generic point (trivial stabilizer), then $\text{Aut}(E_\tau) = \{\pm 1\} \cong \mathbb{Z}/2\mathbb{Z}$.
2. If $\tau \equiv i$ (stabilizer $\mathbb{Z}/2\mathbb{Z}$), then $\text{Aut}(E_\tau) \cong \mathbb{Z}/4\mathbb{Z}$.
3. If $\tau \equiv \rho$ (stabilizer $\mathbb{Z}/3\mathbb{Z}$), then $\text{Aut}(E_\tau) \cong \mathbb{Z}/6\mathbb{Z}$.

Proof:

An automorphism of $E_\tau = \mathbb{C}/\Lambda_\tau$ (fixing the origin) corresponds to $\alpha \in \mathbb{C}^*$ with $\alpha\Lambda_\tau = \Lambda_\tau$ (theorem 3.5.1). The condition $\alpha\Lambda_\tau = \Lambda_\tau$ means $(\alpha\tau, \alpha \cdot 1)$ is another basis for Λ_τ , so there exists $\gamma = \begin{pmatrix} a & b \\ c & d \end{pmatrix} \in \text{SL}_2(\mathbb{Z})$ with $\alpha\tau = a\tau + b$ and $\alpha = c\tau + d$. Eliminating α : $(c\tau + d)\tau = a\tau + b$, i.e., $\gamma \cdot \tau = \tau$. Conversely, $\gamma \cdot \tau = \tau$ gives $\alpha = c\tau + d$ with $\alpha\Lambda_\tau = \Lambda_\tau$.

The map $\text{Aut}(E_\tau) \rightarrow \text{Stab}_{\text{SL}_2(\mathbb{Z})}(\tau)$ sending α to γ has kernel $\{+1, -1\}$ (the elements $\gamma = \pm I$), so $\text{Aut}(E_\tau)/\langle [-1] \rangle \cong \text{Stab}_{\text{PSL}_2(\mathbb{Z})}(\tau)$.

For $\tau = i$: the stabilizer is $\langle S \rangle \cong \mathbb{Z}/2\mathbb{Z}$ in $\text{PSL}_2(\mathbb{Z})$, giving $\text{Aut}(E_i)/\langle [-1] \rangle \cong \mathbb{Z}/2\mathbb{Z}$, hence $|\text{Aut}(E_i)| = 4$. The automorphisms are $\alpha \in \{1, -1, i, -i\}$ (corresponding to $\gamma = I, -I, S, -S$).

For $\tau = \rho$: the stabilizer is $\langle \overline{ST} \rangle \cong \mathbb{Z}/3\mathbb{Z}$, giving $|\text{Aut}(E_\rho)| = 6$. The automorphisms are $\alpha \in \{1, -1, \rho, -\rho, \rho^2, -\rho^2\}$, the sixth roots of unity.

This closes a circle that has been open since section 1.5: the extra automorphisms at $j = 0$ and $j = 1728$ are not merely algebraic curiosities but reflect the *geometric symmetry* of the corresponding points in moduli space. The orbifold points of $\text{SL}_2(\mathbb{Z}) \backslash \mathbb{H}$ are precisely the points where the elliptic curve has extra symmetry.

The Quotient as a Riemann Surface

The quotient $\text{SL}_2(\mathbb{Z}) \backslash \mathbb{H}$ has a natural structure of a (non-compact) Riemann surface, with the orbifold points i and ρ requiring special local coordinates. It is compactified by adding a single point, the *cusp* ∞ , corresponding to the “degenerate” lattice $\tau \rightarrow i\infty$ (i.e., $\Lambda_\tau \rightarrow \mathbb{Z} \cdot i\infty + \mathbb{Z}$, which ceases to be a lattice).

Theorem 3.6.6: The Modular Curve $\text{SL}_2(\mathbb{Z}) \backslash \mathbb{H}^*$

Define $\mathbb{H}^* = \mathbb{H} \cup \mathbb{Q} \cup \{i\infty\}$ (the extended upper half-plane). The quotient

$$X(1) = \text{SL}_2(\mathbb{Z}) \backslash \mathbb{H}^*$$

is a compact Riemann surface of genus 0, and the j -function gives an isomorphism $j : X(1) \xrightarrow{\sim} \mathbb{P}^1(\mathbb{C})$.

Proof Sketch:

Compactness. The space $\mathbb{H}^*$ is given the topology in which a neighborhood of $i\infty$ consists of sets $\{\tau \in \mathbb{H} : \text{Im}(\tau) > M\} \cup \{i\infty\}$ for $M > 0$. The rational cusps $\mathbb{Q}$ are identified with $i\infty$ by $\text{SL}_2(\mathbb{Z})$ (since $\gamma \cdot (i\infty) = a/c$ for $\gamma = \begin{pmatrix} a & b \\ c & d \end{pmatrix}$, and conversely any a/c maps back to $i\infty$). Thus $\text{SL}_2(\mathbb{Z}) \backslash \mathbb{H}^*$ has a single cusp. The quotient is compact because the fundamental domain $\mathcal{F} \cup \{i\infty\}$ is compact in the topology of $\mathbb{H}^*$.

Riemann surface structure. At a generic point $\tau_0 \in \mathbb{H}$ with trivial stabilizer, a local coordinate on the quotient is simply $\tau - \tau_0$. At the orbifold point $\tau = i$ (stabilizer of order 2), a local coordinate is $(\tau - i)^2$ (or any quadratic function vanishing at i that is invariant under S). At $\tau = \rho$ (stabilizer of order 3), a local coordinate is $(\tau - \rho)^3$. At the cusp, a local coordinate is $q = e^{2\pi i \tau}$ (which uniformizes the neighborhood of ∞ in $\text{SL}_2(\mathbb{Z}) \backslash \mathbb{H}^*$, since $T : \tau \mapsto \tau + 1$ acts trivially on q).

Genus computation. By the Gauss–Bonnet theorem (or the Riemann–Hurwitz formula applied to $j : X(1) \rightarrow \mathbb{P}^1$): the Euler characteristic of $X(1)$ is computed from the fundamental domain. The hyperbolic area of $\mathcal{F}$ is $\pi/3$ (a classical computation using the hyperbolic metric $ds^2 = |d\tau|^2 / \text{Im}(\tau)^2$). By the Gauss–Bonnet formula for orbifold surfaces:

$$\chi_{\text{orb}}(X(1)) = -\frac{\text{Area}(\mathcal{F})}{2\pi} + \left(1 - \frac{1}{2}\right) + \left(1 - \frac{1}{3}\right) + 1 = -\frac{1}{6} + \frac{1}{2} + \frac{2}{3} + 1 = 2,$$

where the three correction terms account for the orbifold points of orders 2 and 3 and the cusp. Since $\chi = 2 - 2g$, we get $g = 0$.

Isomorphism via j . Since $X(1)$ has genus 0, it is isomorphic to $\mathbb{P}^1(\mathbb{C})$. The j -function is a meromorphic function on $X(1)$ with a single simple pole (at the cusp), hence is a degree-1 map to $\mathbb{P}^1(\mathbb{C})$: an isomorphism.

The notation $X(1)$ for $\text{SL}_2(\mathbb{Z}) \backslash \mathbb{H}^*$ is standard: it is the *modular curve of level 1*. In ??, we will study modular curves of higher level — the quotients $X_0(N) = \Gamma_0(N) \backslash \mathbb{H}^*$, $X_1(N) = \Gamma_1(N) \backslash \mathbb{H}^*$, and $X(N) = \Gamma(N) \backslash \mathbb{H}^*$ — obtained by replacing $\text{SL}_2(\mathbb{Z})$ with congruence subgroups. These will have genus ≥ 1 for sufficiently large N , and the modular curves $X_0(N)$ will be central to the modularity theorem.

3.6.7 Remark: From j -Line to Moduli Stack The fact that $X(1) \cong \mathbb{P}^1$ via j makes $\mathbb{A}^1 = \mathbb{P}^1 \setminus \{\infty\}$ the “ j -line,” parametrizing isomorphism classes of elliptic curves. This is a *coarse* moduli space: it classifies isomorphism classes but does not carry a universal family (because of the non-trivial automorphisms at $j = 0$ and $j = 1728$). In the language of algebraic geometry, the *moduli stack* $\mathcal{M}_{1,1}$ of elliptic curves is a Deligne–Mumford stack with coarse moduli space $\mathbb{A}^1$. The stacky structure at $j = 0$ and $j = 1728$ is precisely the orbifold structure of $\text{SL}_2(\mathbb{Z}) \backslash \mathbb{H}$ at ρ and i . For most of this book, the coarse moduli space suffices; the stack-theoretic perspective becomes necessary when dealing with level structures and representability of moduli problems in ??.

Example 3.6.8: The Fundamental Domain in Action

Let us use the algorithm from theorem 3.6.3 to find the $\mathrm{SL}_2(\mathbb{Z})$ -representative of $\tau = -5/7 + i/7$.

Step A: $\mathrm{Re}(\tau) = -5/7$. Apply T^1 : $\tau \mapsto \tau + 1 = 2/7 + i/7$. Now $|\mathrm{Re}(\tau)| = 2/7 \leq 1/2$.

Step B: $|\tau|^2 = 4/49 + 1/49 = 5/49 < 1$. Apply S : $\tau \mapsto -1/\tau = -\bar{\tau}/|\tau|^2 = -(2/7 - i/7)/(5/49) = -(2/7 - i/7) \cdot 49/5 = -14/5 + 7i/5$. Now $\mathrm{Im}(\tau) = 7/5 > 1/7$, confirming the imaginary part increased.

Step A: $\mathrm{Re}(\tau) = -14/5 = -2.8$. Apply T^3 : $\tau \mapsto \tau + 3 = 1/5 + 7i/5$. Now $|\mathrm{Re}(\tau)| = 1/5 \leq 1/2$.

Step B: $|\tau|^2 = 1/25 + 49/25 = 50/25 = 2 > 1$. No action needed.

The result is $\tau' = 1/5 + 7i/5 \in \mathcal{F}$, obtained by the $\mathrm{SL}_2(\mathbb{Z})$ -transformation $T^3ST = \begin{pmatrix} 0 & -1 \\ 1 & 1 \end{pmatrix} \begin{pmatrix} 1 & 1 \\ 0 & 1 \end{pmatrix} \cdots$ applied to the original τ . The j -invariant $j(\tau') = j(-5/7 + i/7)$ can now be computed from τ' .

Exercise 3.6.1

1. Verify that $(ST)^3 = -I$ by direct matrix multiplication. Then show that $U = ST$ has order 3 in $\mathrm{PSL}_2(\mathbb{Z})$ and that U fixes $\rho = e^{2\pi i/3}$.
2. Compute the hyperbolic area of $\mathcal{F}$ using the formula $\mathrm{Area}(\mathcal{F}) = \int_{\mathcal{F}} \frac{dx dy}{y^2}$ where $\tau = x + iy$. (*Hint:* integrate x from $-1/2$ to $1/2$ and y from $\sqrt{1-x^2}$ to ∞ , handling the lower boundary $y = \sqrt{1-x^2}$ from the unit circle.)
3. Apply the fundamental domain algorithm to $\tau = 3/5 + i/5$. Determine the $\mathrm{SL}_2(\mathbb{Z})$ -element that maps τ into $\mathcal{F}$, and verify that the result lies in $\mathcal{F}$.
4. Show that the stabilizer of $\tau = i$ in $\mathrm{SL}_2(\mathbb{Z})$ (not $\mathrm{PSL}_2(\mathbb{Z})$) is the cyclic group of order 4 generated by S . List all four elements explicitly.
5. Use the Riemann–Hurwitz formula to compute the genus of $X(1)$ directly from the properties of $j : X(1) \rightarrow \mathbb{P}^1(\mathbb{C})$, which is a degree-1 map with ramification at $j = 0$ (of type $1/3$), $j = 1728$ (of type $1/2$), and $j = \infty$ (a simple pole). Verify that the formula gives $g = 0$.
6. The *principal congruence subgroup* of level N is $\Gamma(N) = \ker(\mathrm{SL}_2(\mathbb{Z}) \rightarrow \mathrm{SL}_2(\mathbb{Z}/N\mathbb{Z}))$. Show that $[\mathrm{SL}_2(\mathbb{Z}) : \Gamma(N)] = N^3 \prod_p |N|_p (1 - p^{-2})$ for $N \geq 1$, and compute $[\mathrm{SL}_2(\mathbb{Z}) : \Gamma(2)]$ and $[\mathrm{SL}_2(\mathbb{Z}) : \Gamma(3)]$ explicitly. (*These indices control the genus of higher-level modular curves in ??.*)

The Formal Group of an Elliptic Curve

The first three chapters studied elliptic curves as global objects: projective curves defined by Weierstrass equations, the morphisms between them, and complex tori uniformized by lattices. All of this was *global* — we worked with the entire curve, or at least with all of its torsion at once. In this chapter we change perspective and zoom in to the infinitesimal neighborhood of the identity point O . The group law on E , which is globally a morphism $m : E \times E \rightarrow E$ of projective varieties, becomes locally a formal power series $F(X, Y) \in R[[X, Y]]$ in a local parameter at O . This power series — the *formal group law* $\hat{E}$ associated to E — is the p -adic engine of the theory.

Why should one care about the local behavior at O ? The answer is that the formal group captures exactly the information that is invisible to the global theory but essential over local fields. Over a local field K with residue field k , the group $E(K)$ fits into a short exact sequence $0 \rightarrow \hat{E}(\mathfrak{m}_K) \rightarrow E(K) \rightarrow \tilde{E}(k) \rightarrow 0$, where $\hat{E}(\mathfrak{m}_K)$ is the kernel of reduction — the subgroup of points that “reduce to the identity.” This kernel is a pro- p group whose structure (its torsion, its rank as a $\mathbb{Z}_p$ -module, its filtration) is completely determined by the formal group law. The reader who has studied Lie groups in [7] will recognize the analogy: $\hat{E}(\mathfrak{m}_K)$ plays the role of a “ p -adic neighborhood of the identity” in $E(K)$, just as a Lie group is locally determined by its Lie algebra. But because elliptic curves are projective (not affine), there is no global coordinate ring — the formal group is new local information not derived by completing/localizing a global ring, and it encodes arithmetic that the function field alone cannot see. Its key invariant, the *height*, governs the ordinary/supersingular dichotomy from section 2.4 and will control the p -part of the torsion over local fields.

4.1 Formal Group Laws

Before specializing to elliptic curves, we set up the general theory of formal group laws. The idea is simple: a formal group law is a power series that satisfies the axioms of a group operation, without any convergence requirements. The “group” has no underlying topological space — it is a purely algebraic object, encoded by a formal power series. The payoff comes when we evaluate the power series on elements of the maximal ideal of a complete local ring, where convergence is guaranteed.

Definition and First Properties

Definition 4.1.1: Formal Group Law

Let R be a commutative ring with identity. A *(one-dimensional, commutative) formal group law* over R is a power series $F(X, Y) \in R[[X, Y]]$ satisfying:

1. (*Identity.*) $F(X, 0) = X$ and $F(0, Y) = Y$.
2. (*Associativity.*) $F(X, F(Y, Z)) = F(F(X, Y), Z)$ in $R[[X, Y, Z]]$.
3. (*Commutativity.*) $F(X, Y) = F(Y, X)$.

We sometimes write $X +_F Y$ for $F(X, Y)$ to emphasize the group-like nature of the operation.

The identity axiom forces $F(X, Y) = X + Y + (\text{terms of degree } \geq 2)$: the linear part of F is always the ordinary addition of formal variables. The higher-order terms encode the “curvature” of the group law near the identity.

Condition (3) (commutativity) is not part of the general definition of a formal group law — one can study non-commutative formal group laws — but it holds for all formal groups arising from abelian group schemes (hence for all elliptic curves), and we impose it throughout.

Example 4.1.2: The Additive Formal Group

The *additive formal group law* is

$$\hat{\mathbb{G}}_a : \quad F(X, Y) = X + Y.$$

All axioms are immediate. This is the formal group associated to the additive algebraic group $\mathbb{G}_a$: the group law is literally addition, with no higher-order corrections.

Example 4.1.3: The Multiplicative Formal Group

The *multiplicative formal group law* is

$$\hat{\mathbb{G}}_m : \quad F(X, Y) = X + Y + XY.$$

This arises from the substitution $t = g - 1$ in the multiplicative group $\mathbb{G}_m$ so that the (formal group law) identity axiom holds. If $g_1, g_2 \in \mathbb{G}_m$ correspond to $t_1 = g_1 - 1$ and $t_2 = g_2 - 1$, then

$$g_1 g_2 - 1 = (1 + t_1)(1 + t_2) - 1 = t_1 + t_2 + t_1 t_2,$$

which is $F(t_1, t_2)$. The identity axiom: $F(X, 0) = X + 0 + 0 = X$. Associativity:

$$\begin{aligned} F(X, F(Y, Z)) &= X + (Y + Z + YZ) + X(Y + Z + YZ) \\ &= X + Y + Z + XY + XZ + YZ + XYZ, \end{aligned}$$

which is symmetric in the three variables, so $F(X, F(Y, Z)) = F(F(X, Y), Z)$.

These two examples are the “trivial” formal groups — they come from the two classical 1-dimensional algebraic groups $\mathbb{G}_a$ and $\mathbb{G}_m$ that appeared as the nonsingular loci of degenerate Weierstrass equations in section 1.6. The formal groups of elliptic curves will be more interesting: their power series have infinitely many nonzero terms, and the complexity of these terms encodes arithmetic information.

The Formal Inverse

The definition of a formal group law does not explicitly include an inverse. This is because the inverse exists automatically — a fact that requires a short argument.

Proposition 4.1.4: Existence of the Formal Inverse

Let $F(X, Y)$ be a formal group law over R . There exists a unique power series $\iota(X) \in R[[X]]$ satisfying

$$F(X, \iota(X)) = 0. \quad (4.1)$$

Moreover, $\iota(X) = -X + (\text{terms of degree } \geq 2)$.

Proof:

We construct $\iota(X)$ by successive approximation. Write $\iota(X) = \sum_{n \geq 1} c_n X^n$ with $c_n \in R$ to be determined. The identity axiom gives $F(X, Y) = X + Y + \sum_{i+j \geq 2} a_{ij} X^i Y^j$ for some $a_{ij} \in R$. Substituting $Y = \iota(X)$:

$$F(X, \iota(X)) = X + \iota(X) + \sum_{i+j \geq 2} a_{ij} X^i \iota(X)^j.$$

Setting this equal to 0 and examining the coefficient of X : we get $1 + c_1 = 0$, so $c_1 = -1$. This confirms $\iota(X) = -X + O(X^2)$.

For $n \geq 2$: the coefficient of X^n in $F(X, \iota(X))$ involves c_n linearly (from the term $X + \iota(X)$, which contributes c_n at degree n) plus terms depending only on $c_1, \dots, c_{n-1}$ and the a_{ij} (from the higher-order terms in F , which contribute products of lower-degree coefficients). Thus the equation $[X^n] F(X, \iota(X)) = 0$ uniquely determines c_n in terms of the previous coefficients. By induction, $\iota(X)$ exists and is unique.

Example 4.1.5: Inverses in the Classical Cases

For $\hat{\mathbb{G}}_a$: $\iota(X) = -X$ (since $X + (-X) = 0$).

For $\hat{\mathbb{G}}_m$: we need $F(X, \iota(X)) = X + \iota(X) + X\iota(X) = 0$, i.e., $\iota(X)(1 + X) = -X$, so $\iota(X) =$

$-X/(1+X) = \sum_{n=1}^{\infty} (-1)^n X^n = -X + X^2 - X^3 + \cdots$. This is the formal power series for $1/(1+X) - 1$, corresponding to the multiplicative inverse $g^{-1} = 1/g$ under the substitution $t = g - 1$.

Formal Group Laws as Functors

The formal group law F defines a group-valued functor: for any commutative R -algebra S that is complete with respect to an ideal $I \subseteq S$, the set I carries a group structure via

$$a +_F b := F(a, b) \quad \text{for } a, b \in I.$$

The power series $F(a, b)$ converges in S because $a, b \in I$ and F has no constant term, so each homogeneous component of $F(a, b)$ lies in I^n for increasingly large n , and the I -adic completeness of S ensures the infinite sum converges. The formal inverse $\iota(a)$ converges similarly, and the group axioms hold because they hold as identities of formal power series.

This functorial perspective is the bridge between formal algebra and p -adic analysis. When we take $R = \mathbb{Z}_p$, $S = \mathbb{Z}_p$, and $I = p\mathbb{Z}_p$, the formal group $\hat{E}$ evaluated on $p\mathbb{Z}_p$ will produce the kernel of reduction — a genuine subgroup of $E(\mathbb{Q}_p)$.

Hensel's Lemma

Many arguments in the theory of formal groups (and more broadly in the study of elliptic curves over local fields) require lifting approximate solutions to exact ones. The following version of Hensel's lemma will serve as our principal tool.

Theorem 4.1.6: Hensel's Lemma – Augmented

Let R be a ring, complete with respect to an ideal $I \subseteq R$, and let $F(w) \in R[w]$. Suppose that there exists an integer $n \geq 1$ and an element $a \in R$ such that

$$F(a) \in I^n \quad \text{and} \quad F'(a) \in R^*.$$

Then for any $\alpha \in R^*$ satisfying $\alpha \equiv F'(a) \pmod{I}$, the sequence

$$w_0 = a, \quad w_{m+1} = w_m - \frac{F(w_m)}{\alpha}$$

converges to an element $b \in R$ satisfying

$$F(b) = 0 \quad \text{and} \quad b \equiv a \pmod{I^n}.$$

If R is an integral domain, these conditions determine b uniquely.

Proof:

First, note that $\alpha \in F'(a) + I \subseteq R^* + I = R^*(1 + I) = R^*$, so α is a unit. To make the recurrence more transparent, replace $F(w)$ by $F(w + a)/\alpha$, which gives the normalized recurrence

$$w_0 = 0, \quad F(0) \in I^n, \quad F'(0) \equiv 1 \pmod{I}, \quad w_{m+1} = w_m - F(w_m).$$

Since $F(0) \in I^n$, the relation $w_m \in I^n \implies w_m - F(w_m) \in I^n$ shows that $w_m \in I^n$ for all $m \geq 0$.

We claim by induction that $w_m \equiv w_{m+1} \pmod{I^{m+n}}$ for all $m \geq 0$. The base case $m = 0$ is the assumption $F(0) \in I^n$. For the inductive step, factor

$$F(X) - F(Y) = (X - Y)(F'(0) + XG(X, Y) + YH(X, Y))$$

for some polynomials $G, H \in R[X, Y]$. Then:

$$\begin{aligned} w_{m+1} - w_m &= (w_m - F(w_m)) - (w_{m-1} - F(w_{m-1})) \\ &= (w_m - w_{m-1}) - (F(w_m) - F(w_{m-1})) \\ &= (w_m - w_{m-1})(1 - F'(0) - w_m G(w_m, w_{m-1}) - w_{m-1} H(w_m, w_{m-1})). \end{aligned}$$

By the induction hypothesis, $w_m - w_{m-1} \in I^{m-1+n}$. The factor in parentheses lies in I (since $F'(0) \equiv 1 \pmod{I}$ and $w_m, w_{m-1} \in I^n \subseteq I$). Therefore $w_{m+1} - w_m \in I^{m+n}$.

Since R is I -adically complete, the sequence (w_m) converges to an element $b \in R$. As every $w_m \in I^n$, we have $b \in I^n$. Taking the limit of $w_{m+1} = w_m - F(w_m)$ gives $b = b - F(b)$, so $F(b) = 0$.

For uniqueness when R is an integral domain: suppose $c \in I^n$ also satisfies $F(c) = 0$. Then

$$0 = F(b) - F(c) = (b - c)(F'(0) + bG(b, c) + cH(b, c)).$$

If $b \neq c$, then $F'(0) + bG(b, c) + cH(b, c) = 0$, which would give $F'(0) = -bG(b, c) - cH(b, c) \in I$, contradicting $F'(0) \equiv 1 \pmod{I}$. Hence $b = c$.

The augmented form of Hensel's lemma is more flexible than the classical version (which requires $F(a) \in I$ and $F'(a) \in R^*$): by allowing $F(a) \in I^n$ for $n \geq 1$, we can start with a cruder approximation and still converge. The freedom to choose $\alpha \neq F'(a)$ (merely congruent modulo I) will be useful when the exact derivative is difficult to compute but a suitable approximation is at hand.

Homomorphisms of Formal Groups

Definition 4.1.7: Homomorphism of Formal Group Laws

Let F and G be formal group laws over R . A *homomorphism* from F to G is a power series $f(X) \in R[[X]]$ with $f(0) = 0$ satisfying

$$f(F(X, Y)) = G(f(X), f(Y)). \quad (4.2)$$

A homomorphism is an *isomorphism* if $f'(0) \in R^*$ (equivalently, if f has a composition inverse $f^{-1} \in R[[X]]$). We write $\text{Hom}(F, G)$ for the set of homomorphisms from F to G , and $\text{End}(F) = \text{Hom}(F, F)$.

The condition $f(0) = 0$ says that f “preserves the identity.” The functional equation (4.2) says f

intertwines the two group operations. Writing $f(X) = \sum_{n \geq 1} c_n X^n$, the leading term $c_1 = f'(0)$ plays the role of the “tangent map” or “derivative at the identity” — the formal group analogue of the Lie algebra map.

Proposition 4.1.8: Structure of $\text{Hom}(F, G)$

Let F and G be formal group laws over R .

1. The set $\text{Hom}(F, G)$ is an abelian group under the operation $(f +_G g)(X) := G(f(X), g(X))$.
2. $\text{End}(F)$ is a ring, with addition as in (1) and multiplication given by composition: $(f \circ g)(X) = f(g(X))$.
3. The map $\mathbb{Z} \rightarrow \text{End}(F)$ sending $n \mapsto [n]_F$ (where $[n]_F$ is the n -fold formal sum: $[1]_F(X) = X$, $[n+1]_F(X) = F([n]_F(X), X)$) is a ring homomorphism.

Proof:

(1) We verify that $f +_G g$ is a homomorphism:

$$\begin{aligned} (f +_G g)(F(X, Y)) &= G(f(F(X, Y)), g(F(X, Y))) \\ &= G(G(f(X), f(Y)), G(g(X), g(Y))). \end{aligned}$$

On the other hand:

$$G((f +_G g)(X), (f +_G g)(Y)) = G(G(f(X), g(X)), G(f(Y), g(Y))).$$

These are equal by the commutativity and associativity of G : both expressions equal the four-fold sum $f(X) +_G f(Y) +_G g(X) +_G g(Y)$ in G . The identity element is the zero homomorphism $0(X) = 0$, and the inverse of f is $\iota_G \circ f$ (where ι_G is the formal inverse of G).

(2) Composition distributes over addition because G is a group operation: $f \circ (g +_F h) = f(F(g, h))$, and since f is a homomorphism, this equals $G(f(g), f(h)) = (f \circ g) +_G (f \circ h)$. Similarly for right distributivity using the fact that $g +_F h$ is defined via F .

(3) $[n]_F$ is a homomorphism because $[n]_F(F(X, Y)) = F([n]_F(X), [n]_F(Y))$ (which follows by induction from the associativity of F). That $n \mapsto [n]_F$ respects addition and multiplication is verified directly.

The $[n]$ -Series

The endomorphism $[n]_F$ is central to the theory, just as the multiplication-by- n map was central to the global theory of elliptic curves.

Proposition 4.1.9: The $[n]$ -Series

Let F be a formal group law over R .

1. $[0]_F(X) = 0$ and $[1]_F(X) = X$.
2. $[n]_F(X) = nX + (\text{terms of degree } \geq 2)$ for all $n \in \mathbb{Z}$.
3. $[-1]_F(X) = \iota(X)$, the formal inverse.
4. $[n + m]_F = [n]_F +_F [m]_F$ and $[nm]_F = [n]_F \circ [m]_F$.

Proof:

(1) and (3) are immediate from the definitions. For (2): $[1]_F(X) = X = 1 \cdot X$, and inductively $[n + 1]_F(X) = F([n]_F(X), X) = [n]_F(X) + X + O(X^2) = (n + 1)X + O(X^2)$ (using the identity axiom of F). For negative n , $[n]_F = [-1]_F \circ [-n]_F$, and $[-1]_F(X) = -X + O(X^2)$, so $[n]_F'(0) = (-1)(-n) = n$. (4) is part of proposition 4.1.8(3).

The key content of (2) is the leading coefficient: $[n]_F'(0) = n$. Over a ring where n is a unit, $[n]_F$ is an isomorphism (it has an invertible leading coefficient). Over a ring of characteristic p , the series $[p]_F$ has $[p]_F'(0) = p = 0$, so it is not an isomorphism — the behavior of $[p]_F$ beyond its vanishing leading term is the main invariant of the formal group in characteristic p .

Endomorphisms of the Classical Formal Groups**Example 4.1.10: Endomorphisms of $\hat{\mathbb{G}}_a$ and $\hat{\mathbb{G}}_m$**

(a) An endomorphism of $\hat{\mathbb{G}}_a$ is a power series $f(X) \in R[[X]]$ with $f(0) = 0$ and $f(X + Y) = f(X) + f(Y)$. If R is a $\mathbb{Q}$ -algebra, the only such power series are $f(X) = cX$ for $c \in R$: expanding $f(X + Y) = f(X) + f(Y)$ and comparing coefficients of $X^i Y^j$ shows all coefficients of degree ≥ 2 must vanish. So $\text{End}(\hat{\mathbb{G}}_a) \cong R$ (as a ring under addition and multiplication).

If R has characteristic $p > 0$, however, the Frobenius map $f(X) = X^p$ satisfies $f(X + Y) = (X + Y)^p = X^p + Y^p = f(X) + f(Y)$ (using the freshman's dream identity in characteristic p). More generally, X^{p^n} is an endomorphism for every $n \geq 0$, and $\text{End}(\hat{\mathbb{G}}_a)$ is the ring of p -linearized power series — those of the form $\sum c_n X^{p^n}$. This is a non-commutative ring (under composition) when R is not a prime field.

The $[p]$ -series is $[p]_{\hat{\mathbb{G}}_a}(X) = pX$, which is 0 in characteristic p . Thus $\hat{\mathbb{G}}_a$ has “height ∞ ” in the sense that the $[p]$ -series vanishes identically.

(b) An endomorphism of $\hat{\mathbb{G}}_m$ is a power series $f(X)$ with $f(0) = 0$ and $f(X + Y + XY) = f(X) + f(Y) + f(X)f(Y)$. Under the substitution $T = 1 + X$ (so $f(X) = g(1 + X) - 1$), this becomes $g(ST) = g(S)g(T)$, i.e., g is a multiplicative homomorphism of $1 + \mathfrak{m}$. The power series $g(T) = T^n$ satisfy this for all $n \in \mathbb{Z}$, giving $f(X) = (1 + X)^n - 1 = \sum_{k=1}^{\infty} \binom{n}{k} X^k$. Over $\mathbb{Z}$, these are all the endomorphisms: $\text{End}(\hat{\mathbb{G}}_m) \cong \mathbb{Z}$.

The $[p]$ -series is $[p]_{\hat{\mathbb{G}}_m}(X) = (1+X)^p - 1$. In characteristic p , $(1+X)^p - 1 = X^p$: the first nonzero term beyond the leading (vanishing) linear term has degree $p = p^1$. This is the hallmark of a formal group of *height* 1.

The contrast between $\hat{\mathbb{G}}_a$ (height ∞ , with $[p] = 0$) and $\hat{\mathbb{G}}_m$ (height 1, with $[p](X) = X^p$) foreshadows the key invariant of formal groups in characteristic p . For the formal group of an elliptic curve, we will find height 1 (ordinary) or height 2 (supersingular) — but never height ∞ (since $[p]$ on an elliptic curve is never the zero isogeny). The precise definition and analysis of height will be given in section 4.4.

Exercise 4.1.1

1. Show that if $F(X, Y) \in R[[X, Y]]$ satisfies only the identity axiom ($F(X, 0) = X$) and the associativity axiom ($F(X, F(Y, Z)) = F(F(X, Y), Z)$), then F automatically satisfies $F(0, Y) = Y$. (*Hint:* Apply associativity with $Z = \iota(Y)$, using the existence of the formal inverse proved in proposition 4.1.4. Alternatively, argue directly from the power series coefficients.)
2. Compute the formal inverse $\iota(X)$ of the formal group law $F(X, Y) = X + Y + aXY$ (for $a \in R$) up to degree 4. Verify that your answer agrees with the closed form $\iota(X) = -X/(1 + aX)$.
3. Over a ring R of characteristic $p > 0$, show that the composition $X^p \circ (cX) = c^p X^p$ while $(cX) \circ X^p = cX^p$, so the endomorphism ring of $\hat{\mathbb{G}}_a$ is non-commutative when R contains elements c with $c \neq c^p$.
4. Compute $[3]_{\hat{\mathbb{G}}_m}(X) = (1+X)^3 - 1$ explicitly. Verify that $[3]_{\hat{\mathbb{G}}_m}([2]_{\hat{\mathbb{G}}_m}(X)) = [6]_{\hat{\mathbb{G}}_m}(X) = (1+X)^6 - 1$.
5. Over a ring of characteristic 2, compute $[2]_{\hat{\mathbb{G}}_m}(X)$ and $[4]_{\hat{\mathbb{G}}_m}(X)$. Show that $[2]_{\hat{\mathbb{G}}_m}(X) = X^2$ and $[4]_{\hat{\mathbb{G}}_m}(X) = X^4$. What is $[2^n]_{\hat{\mathbb{G}}_m}(X)$ for general n ?
6. Let $R = \mathbb{Z}_p$ and $F(w) = w^p - w \in \mathbb{Z}_p[w]$. Use theorem 4.1.6 to show that for any $a \in \mathbb{Z}_p$ with $a^p \equiv a \pmod{p}$ and $p \nmid (pa^{p-1} - 1)$, there exists $b \in \mathbb{Z}_p$ with $b^p - b = 0$ and $b \equiv a \pmod{p}$. What are the solutions for $p = 5$?

4.2 The Formal Group of an Elliptic Curve

We now specialize from abstract formal group laws to the one that matters: the formal group $\hat{E}$ associated to a Weierstrass equation. The construction is entirely explicit — we choose a local parameter at the identity O , express the coordinate functions x and y as Laurent series in this parameter, and then write the group law as a power series. The result is a formal group law whose coefficients are polynomial expressions in the Weierstrass coefficients $a_1, \dots, a_6$.

Local Coordinates at O

Let E/K be an elliptic curve given by a Weierstrass equation

$$y^2 + a_1xy + a_3y = x^3 + a_2x^2 + a_4x + a_6 \quad (4.3)$$

with distinguished point $O = [0 : 1 : 0] \in \mathbb{P}_K^2$. To study the group law near O , we work in the affine chart $Y \neq 0$ of $\mathbb{P}^2$. Following standard sign conventions that simplify later formulas, we set $z = -X/Y$ and $w = -Z/Y$. This implies $x = X/Z = z/w$ and $y = Y/Z = -1/w$, and the point $O = [0 : 1 : 0]$ corresponds to the origin $(z, w) = (0, 0)$.

Proposition 4.2.1: Local Parameter at O

In the (z, w) -coordinates, the point $O = (0, 0)$ is a smooth point of E , and z is a uniformizer of the local ring $\mathcal{O}_{E,O}$. Equivalently, $z = -x/y$ is a local parameter at O .

Proof:

Substituting $x = z/w$ and $y = -1/w$ into (4.3) gives:

$$\left(-\frac{1}{w}\right)^2 + a_1\left(\frac{z}{w}\right)\left(-\frac{1}{w}\right) + a_3\left(-\frac{1}{w}\right) = \left(\frac{z}{w}\right)^3 + a_2\left(\frac{z}{w}\right)^2 + a_4\left(\frac{z}{w}\right) + a_6.$$

Multiplying through by w^3 and rearranging to isolate w yields:

$$w = z^3 + a_1zw + a_2z^2w + a_3w^2 + a_4zw^2 + a_6w^3. \quad (4.4)$$

Denote the right-hand side by $\Phi(z, w)$, so the curve is defined by $g(z, w) := w - \Phi(z, w) = 0$ in the (z, w) -plane. At the origin: $g(0, 0) = 0$ (so O lies on the curve), and $\partial g / \partial w|_{(0,0)} = 1$ (since Φ has no constant or linear-in- w term). By the implicit function theorem for formal power series (or equivalently, by theorem 4.1.6 applied to the polynomial $g(z, w)$ viewed as a polynomial in w over the complete ring $K[[z]]$, with the approximate solution $w = 0$ satisfying $g(z, 0) = -z^3 \in (z)^3$ and $g'_w(z, 0)|_{z=0} = 1 \in K[[z]]^*$), there exists a unique power series $w(z) \in K[[z]]$ with $w(0) = 0$ satisfying (4.4). This shows w is a function of z near O , so z generates the maximal ideal $\mathfrak{m}_O$ of $\mathcal{O}_{E,O}$: it is a uniformizer.

The relation $z = -x/y$ is easy to remember: x has a pole of order 2 at O and y has a pole of order 3, so x/y has a pole of order $2 - 3 = -1$, i.e., a *zero* of order 1 at O .

We now solve (4.4) for w as a power series in z by successive substitution. Write $w = \sum_{n \geq 3} c_n z^n$ (the series starts at z^3 because $\Phi(z, 0) = z^3$) and determine the coefficients c_n by substituting into (4.4) and matching powers of z .

The equation $w = z^3 + a_1zw + a_2z^2w + a_3w^2 + a_4zw^2 + a_6w^3$ gives a recurrence: the coefficient c_n is determined by the a_i and the previously computed $c_3, \dots, c_{n-1}$. The first several terms are:

Proposition 4.2.2: $w(z)$ Expansion

$$w(z) = z^3 + a_1z^4 + (a_1^2 + a_2)z^5 + (a_1^3 + 2a_1a_2 + a_3)z^6 + \dots \quad (4.5)$$

Proof:

Setting $w = z^3 + c_4z^4 + c_5z^5 + c_6z^6 + \dots$ and substituting into (4.4):

Coefficient of z^4 : The only contribution at degree 4 comes from $a_1z \cdot z^3 = a_1z^4$ on the right-hand

side. Thus $c_4 = a_1$.

Coefficient of z^5 : The contributions are $a_1 z \cdot c_4 z^4 = a_1^2 z^5$ and $a_2 z^2 \cdot z^3 = a_2 z^5$. Thus $c_5 = a_1^2 + a_2$.

Coefficient of z^6 : The contributions are $a_1 z \cdot c_5 z^5 = a_1(a_1^2 + a_2)z^6$, $a_2 z^2 \cdot c_4 z^4 = a_1 a_2 z^6$, and $a_3 \cdot (z^3)^2 = a_3 z^6$. Thus $c_6 = a_1^3 + 2a_1 a_2 + a_3$.

The pattern continues: each coefficient c_n is a polynomial in the a_i with integer coefficients, weighted so that c_n has weight $n - 3$ (where a_i has weight i).

From $w(z)$ we recover the coordinate functions $x = z/w$ and $y = -1/w$ as Laurent series in z :

Proposition 4.2.3: Laurent Series for x and y

As elements of $K((z))$:

$$x = z^{-2} - a_1 z^{-1} - a_2 - a_3 z - (a_4 + a_1 a_3) z^2 - \cdots, \quad (4.6)$$

$$y = -z^{-3} + a_1 z^{-2} + a_2 z^{-1} + a_3 + (a_4 + a_1 a_3) z + \cdots. \quad (4.7)$$

Proof:

We have $w(z) = z^3(1 + c_4 z + c_5 z^2 + c_6 z^3 + \cdots)$ with $c_4 = a_1$, $c_5 = a_1^2 + a_2$, $c_6 = a_1^3 + 2a_1 a_2 + a_3$. Using the geometric series:

$$\frac{1}{w(z)} = z^{-3} \cdot \frac{1}{1 + c_4 z + c_5 z^2 + \cdots} = z^{-3} (1 - c_4 z + (c_4^2 - c_5) z^2 - (c_4^3 - 2c_4 c_5 + c_6) z^3 + \cdots).$$

For $x = z/w = z \cdot w^{-1} = z^{-2} (1 - c_4 z + (c_4^2 - c_5) z^2 + \cdots)$:

$$x = z^{-2} - c_4 z^{-1} + (c_4^2 - c_5) - (c_4^3 - 2c_4 c_5 + c_6) z + \cdots.$$

Substituting: $c_4 = a_1$, $c_4^2 - c_5 = a_1^2 - (a_1^2 + a_2) = -a_2$, and $c_4^3 - 2c_4 c_5 + c_6 = a_1^3 - 2a_1(a_1^2 + a_2) + (a_1^3 + 2a_1 a_2 + a_3) = a_3$. This gives (4.6).

For $y = -1/w = -z^{-3} (1 - c_4 z + (c_4^2 - c_5) z^2 - (c_4^3 - 2c_4 c_5 + c_6) z^3 + \cdots)$:

$$y = -z^{-3} + a_1 z^{-2} + a_2 z^{-1} + a_3 + \cdots.$$

These Laurent series invert the assignment $P \mapsto z(P)$: given a small value of the local parameter z , we can recover the affine coordinates (x, y) of the corresponding point on E near O . The poles at $z = 0$ reflect the fact that $O = [0 : 1 : 0]$ is the point at infinity on the Weierstrass model, where the affine coordinates x and y are not defined.

The Formal Group Law

We now come to the main construction. Given two points P_1, P_2 near O with local parameters $z_1 = z(P_1)$ and $z_2 = z(P_2)$, the sum $P_3 = P_1 + P_2$ has a local parameter $z_3 = z(P_1 + P_2)$. We claim that z_3 is a formal power series in z_1 and z_2 .

Theorem 4.2.4: The Formal Group of an Elliptic Curve

Let E/K be an elliptic curve given by a Weierstrass equation (4.3). The group law on E , expressed in the local parameter $z = -x/y$ at O , defines a formal group law

$$\hat{E}(z_1, z_2) = z_1 + z_2 - a_1 z_1 z_2 - a_2 (z_1^2 z_2 + z_1 z_2^2) + \cdots \in \mathbb{Z}[a_1, a_2, a_3, a_4, a_6][[z_1, z_2]]. \quad (4.8)$$

This power series satisfies all the axioms of definition 4.1.1, and its coefficients are universal polynomials in $a_1, \dots, a_6$ with integer coefficients.

Proof:

Step 1: $\hat{E}(z_1, z_2)$ is a power series. The addition map $m : E \times E \rightarrow E$ is a morphism of smooth varieties sending $(O, O) \mapsto O$. In local coordinates at the source (using (z_1, z_2) as local parameters at $(O, O) \in E \times E$) and at the target (using z_3 as local parameter at $O \in E$), the morphism m is represented by a formal power series $\hat{E}(z_1, z_2) \in K[[z_1, z_2]]$ with $\hat{E}(0, 0) = 0$. This is because a morphism of smooth varieties induces a local ring homomorphism $m^* : \mathcal{O}_{E,O} \rightarrow \mathcal{O}_{E \times E, (O,O)}$, and since z generates $\mathfrak{m}_O$, the image $m^*(z) = \hat{E}(z_1, z_2)$ lies in $\mathfrak{m}_{(O,O)} = (z_1, z_2) \cdot K[[z_1, z_2]]$.

This is the key point: although the addition formulas for E involve rational functions (with denominators), the group law is a *morphism* of the underlying scheme, and a morphism is locally given by regular functions — hence by a power series in the completed local ring. The denominators that appear in the explicit chord-and-tangent formulas all cancel when one passes to the local parameter z at O .

Step 2: The axioms. The identity axiom $\hat{E}(z_1, 0) = z_1$ holds because $P + O = P$ for all P , which in local coordinates says $z(P + O) = z(P)$. Associativity $\hat{E}(z_1, \hat{E}(z_2, z_3)) = \hat{E}(\hat{E}(z_1, z_2), z_3)$ holds because addition on E is associative (??). Commutativity $\hat{E}(z_1, z_2) = \hat{E}(z_2, z_1)$ holds because E is an abelian group.

Step 3: Integrality. The coefficients of $\hat{E}$ are polynomials in $a_1, \dots, a_6$ with coefficients in $\mathbb{Z}$, because the addition formulas (??) involve only the a_i and the ring operations. More precisely, the power series $x(z)$ and $y(z)$ from proposition 4.2.3 have coefficients in $\mathbb{Z}[a_1, \dots, a_6]$, the addition formulas produce rational functions of x_1, y_1, x_2, y_2 with coefficients in $\mathbb{Z}[a_1, \dots, a_6]$, and the resulting expression for $z_3 = -x_3/y_3$ is a power series (by Step 1) with the same integrality.

Step 4: Explicit computation of the first terms. To compute $\hat{E}(z_1, z_2)$, we need the group law $P_1 + P_2 = P_3$ expressed in local coordinates. Given z_1 and z_2 , we compute $x_i = x(z_i)$ and $y_i = y(z_i)$ from (4.6)–(4.7), then apply the addition formulas to get (x_3, y_3) , and finally extract

$$z_3 = -x_3/y_3.$$

The slope of the line through P_1 and P_2 (for $P_1 \neq P_2$) is

$$\lambda = \frac{y_1 - y_2}{x_1 - x_2}.$$

Using the Laurent series, we factor $z_1 - z_2$ from both numerator and denominator. For instance:

$$\begin{aligned} x_1 - x_2 &= (z_1^{-2} - z_2^{-2}) - a_1(z_1^{-1} - z_2^{-1}) - a_3(z_1 - z_2) - \cdots \\ &= \frac{z_2^2 - z_1^2}{z_1^2 z_2^2} - a_1 \cdot \frac{z_2 - z_1}{z_1 z_2} - a_3(z_1 - z_2) - \cdots \\ &= (z_1 - z_2) \left[\frac{-(z_1 + z_2)}{z_1^2 z_2^2} + \frac{a_1}{z_1 z_2} - a_3 - \cdots \right]. \end{aligned}$$

Similarly for $y_1 - y_2$. The ratio λ is therefore a Laurent series in z_1, z_2 (the factor $z_1 - z_2$ cancels). After computing $x_3 = \lambda^2 + a_1\lambda - a_2 - x_1 - x_2$ and y_3 from the line equation, one extracts $z_3 = -x_3/y_3$. The computation is lengthy but completely mechanical; we record the result:

$$\begin{aligned} \hat{E}(z_1, z_2) &= z_1 + z_2 - a_1 z_1 z_2 - a_2 (z_1^2 z_2 + z_1 z_2^2) \\ &\quad - (2a_3 (z_1^3 z_2 + z_1 z_2^3) + (a_1 a_2 - 3a_3) z_1^2 z_2^2) + \cdots \end{aligned}$$

where the omitted terms have total degree ≥ 5 in (z_1, z_2) . One can verify: $\hat{E}(z_1, 0) = z_1$ (identity), $\hat{E}(z_1, z_2) = \hat{E}(z_2, z_1)$ (commutativity), and the degree-2 term is $-a_1 z_1 z_2$ (reflecting the $a_1 xy$ term in the Weierstrass equation).

The formal group $\hat{E}$ depends on the choice of Weierstrass equation, not just on the isomorphism class of E . An admissible change of variables (u, r, s, t) replaces the local parameter z by $z' = u^{-1}z + \cdots$ (since $z = -x/y$ and x, y transform under (1.3)), giving an isomorphic but not identical formal group law. The isomorphism class of $\hat{E}$ as a formal group is an invariant of the curve, even though the specific power series is not.

The Formal Inverse

The formal inverse $\iota(z)$ of $\hat{E}$ corresponds to the negation map $P \mapsto -P$ on E . Since $-(x, y) = (x, -y - a_1 x - a_3)$ on the Weierstrass model (recall the hyperelliptic involution from section 1.5), we can compute $\iota(z)$ directly from the Laurent series.

Proposition 4.2.5: The Formal Inverse of $\hat{E}$

The formal inverse of $\hat{E}$ is

$$\iota(z) = -z - a_1 z^2 - a_1^2 z^3 + \cdots. \quad (4.9)$$

Proof:

The negation map sends $(x, y) \mapsto (x, -y - a_1x - a_3)$. In the local parameter $z = -x/y$:

$$\iota(z) = \frac{-x(-P)}{y(-P)} = \frac{-x}{-y - a_1x - a_3} = \frac{x}{y + a_1x + a_3}.$$

Using the Laurent series (4.6)–(4.7):

$$\begin{aligned} y + a_1x + a_3 &= (-z^{-3} + a_1z^{-2} + a_2z^{-1} + a_3 + \cdots) \\ &\quad + a_1(z^{-2} - a_1z^{-1} - a_2 - \cdots) + a_3 \\ &= -z^{-3} + 2a_1z^{-2} + (a_2 - a_1^2)z^{-1} + (2a_3 - a_1a_2) + \cdots \end{aligned}$$

Therefore, dividing the series:

$$\iota(z) = \frac{z^{-2}(1 - a_1z - a_2z^2 - \cdots)}{-z^{-3}(1 - 2a_1z - (a_2 - a_1^2)z^2 - \cdots)} = -z \cdot \frac{1 - a_1z - a_2z^2 - \cdots}{1 - 2a_1z - (a_2 - a_1^2)z^2 - \cdots}.$$

Expanding $(1 - 2a_1z - (a_2 - a_1^2)z^2 - \cdots)^{-1} = 1 + 2a_1z + (a_2 + 3a_1^2)z^2 + \cdots$ and multiplying by the numerator:

$$\begin{aligned} \iota(z) &= -z \cdot (1 - a_1z - a_2z^2 + \cdots)(1 + 2a_1z + (a_2 + 3a_1^2)z^2 + \cdots) \\ &= -z(1 + a_1z + a_1^2z^2 + \cdots) \\ &= -z - a_1z^2 - a_1^2z^3 - \cdots. \end{aligned}$$

In the case where $a_1 = a_3 = 0$ (which includes the short Weierstrass form), the formal inverse is exactly $\iota(z) = -z$. This reflects the fact that $-(x, y) = (x, -y)$ in this case, meaning $z(-P) = -x/(-y) = x/y = -(-x/y) = -z(P)$.

The Invariant Differential

The invariant differential $\omega = dx/(2y + a_1x + a_3)$ on E played a crucial role in section 2.3, where we used $[n]^*\omega = n\omega$ to detect separability. We now express ω in terms of the local parameter z .

Proposition 4.2.6: The Invariant Differential in Local Coordinates

In the local parameter z at O :

$$\omega = \frac{dx}{2y + a_1x + a_3} = (1 + a_1z + (a_1^2 + a_2)z^2 + \cdots) dz. \quad (4.10)$$

In particular, ω is a regular (nowhere-vanishing) differential at O , and its expansion in z has leading coefficient 1.

Proof:

From (4.6): $x = z^{-2} - a_1 z^{-1} - a_2 - a_3 z - \dots$, so differentiating gives:

$$\frac{dx}{dz} = -2z^{-3} + a_1 z^{-2} - a_3 - \dots = -2z^{-3} \left(1 - \frac{a_1}{2} z + \frac{a_3}{2} z^3 + \dots \right).$$

From the preceding computation of $y + a_1 x + a_3$:

$$2y + a_1 x + a_3 = y + (y + a_1 x + a_3) = -2z^{-3} + 3a_1 z^{-2} + (2a_2 - a_1^2) z^{-1} + \dots.$$

Forming the ratio and factoring $-2z^{-3}$ from both:

$$\begin{aligned} \omega &= \frac{dx/dz}{2y + a_1 x + a_3} dz = \frac{-2z^{-3} \left(1 - \frac{a_1}{2} z + \dots \right)}{-2z^{-3} \left(1 - \frac{3a_1}{2} z - \frac{2a_2 - a_1^2}{2} z^2 + \dots \right)} dz \\ &= \frac{1 - \frac{a_1}{2} z + \dots}{1 - \frac{3a_1}{2} z - \frac{2a_2 - a_1^2}{2} z^2 + \dots} dz. \end{aligned}$$

Expanding the denominator: $\left(1 - \frac{3a_1}{2} z - \frac{2a_2 - a_1^2}{2} z^2 + \dots \right)^{-1} = 1 + \frac{3a_1}{2} z + \left(\frac{2a_2 - a_1^2}{2} + \frac{9a_1^2}{4} \right) z^2 + \dots = 1 + \frac{3a_1}{2} z + \frac{4a_2 + 7a_1^2}{4} z^2 + \dots$. Multiplying by the numerator gives:

$$\begin{aligned} \omega &= \left(1 - \frac{a_1}{2} z + \dots \right) \left(1 + \frac{3a_1}{2} z + \frac{4a_2 + 7a_1^2}{4} z^2 + \dots \right) dz \\ &= \left(1 + \left(\frac{3a_1}{2} - \frac{a_1}{2} \right) z + \left(\frac{4a_2 + 7a_1^2}{4} - \frac{3a_1^2}{4} \right) z^2 + \dots \right) dz \\ &= (1 + a_1 z + (a_1^2 + a_2) z^2 + \dots) dz. \end{aligned}$$

The fact that $\omega = (1 + \text{higher order}) dz$ means ω is a generator of $\Omega_{E/K}^1$ at O (since dz is a generator and the coefficient $1 + \dots$ is a unit in $K[[z]]$). This is consistent with the fact that the canonical bundle of a genus-1 curve is trivial ($K_E \sim 0$, from section 1.1), so ω generates $H^0(E, \Omega_E^1)$ and is everywhere nonvanishing.

We write

$$\omega(z) = \sum_{n=0}^{\infty} d_n z^n dz \quad \text{with } d_0 = 1, d_1 = a_1, d_2 = a_1^2 + a_2, \dots \quad (4.11)$$

These coefficients d_n match the c_{n+3} in the $w(z)$ -expansion: $d_0 = c_3 = 1$, $d_1 = c_4 = a_1$, $d_2 = c_5 = a_1^2 + a_2$. This is not a coincidence — it reflects the relation $\omega = d(w/z^2) \cdot (\text{correction terms})$, ultimately rooted in the fact that both the invariant differential and the formal group are intrinsic to the local structure at O .

The $[n]$ -Endomorphism

On the formal group $\hat{E}$, the $[n]$ -series $[n]_{\hat{E}}(z) = \underbrace{\hat{E}(z, \hat{E}(z, \dots))}_n$ corresponds to the multiplication-by- n map $[n] : E \rightarrow E$ expressed in local coordinates. The key property, promised in section 2.3, follows immediately:

Proposition 4.2.7: The $[n]$ -Series and the Invariant Differential

The $[n]$ -series of $\hat{E}$ satisfies $[n]'_{\hat{E}}(0) = n$. Equivalently, $[n]^*\omega = n\omega$.

Proof:

From proposition 4.1.9(2), $[n]_F(z) = nz + O(z^2)$ for any formal group law F . For the equivalence with $[n]^*\omega = n\omega$: the pullback of $\omega = (1 + \cdots) dz$ under $[n]_{\hat{E}}(z) = nz + \cdots$ is

$$[n]^*\omega = (1 + \cdots)|_{z=[n](z)} \cdot [n]'(z) dz.$$

Evaluating the leading term: $(1 + \cdots)|_{z=0} \cdot [n]'(0) = 1 \cdot n = n$. So $[n]^*\omega = n \cdot (1 + \cdots) dz = n\omega$, as desired. Alternatively, this follows from translation-invariance: $\tau_P^*\omega = \omega$ for all $P \in E$, and $[n] = \sum_{i=1}^n \tau_{[i-1]P}$, so $[n]^*\omega = n\omega$.

This gives the clean proof of $[n]^*\omega = n\omega$ that was deferred from section 2.3: it is simply the statement that the linear coefficient of $[n]_{\hat{E}}$ is n .

Running Examples**Example 4.2.8: Formal Group of $y^2 = x^3 - x$**

For $E : y^2 = x^3 - x$ over K with $\text{char} K \neq 2, 3$, all $a_i = 0$ except $a_4 = -1$. The w -expansion simplifies dramatically: $c_n = 0$ for $3 \leq n \leq 6$, and $c_7 = a_4 = -1$. Thus:

$$w(z) = z^3 - z^7 + \cdots$$

The Laurent series become:

$$\begin{aligned} x &= z^{-2} + z^2 + \cdots, \\ y &= -z^{-3} - z + \cdots. \end{aligned}$$

One can verify: $x^3 - x = z^{-6} + 3z^{-2} + \cdots - z^{-2} - z^2 - \cdots = z^{-6} + 2z^{-2} + \cdots$ and $y^2 = z^{-6} + 2z^{-2} + \cdots$, confirming the Weierstrass equation is satisfied.

Since $a_1 = a_2 = a_3 = 0$, the formal group law starts as $\hat{E}(z_1, z_2) = z_1 + z_2 + (\text{terms of degree } \geq 5)$: the group law is “more additive” than for a general curve, with the first corrections arising from a_4 at total degree 5.

The invariant differential is $\omega = dz + O(z^4) dz = dx/(2y)$.

Example 4.2.9: Formal Group of $y^2 = x^3 + 1$

For $E : y^2 = x^3 + 1$ over K with $\text{char} K \neq 2, 3$, all $a_i = 0$ except $a_6 = 1$. The first correction to $w(z) = z^3$ occurs even later: $c_n = 0$ for $3 \leq n \leq 8$, and $c_9 = a_6 = 1$. So:

$$w(z) = z^3 + z^9 + \cdots$$

The formal group law is $\hat{E}(z_1, z_2) = z_1 + z_2 + (\text{terms of degree } \geq 7)$ — the group law is even closer to the additive formal group than in the previous example.

Example 4.2.10: A Curve with $a_1 \neq 0$

For $E : y^2 + xy + y = x^3$ (so $a_1 = 1, a_3 = 1, a_2 = a_4 = a_6 = 0$), the expansion is

$$w(z) = z^3 + z^4 + z^5 + 2z^6 + \cdots$$

and the formal group law begins

$$\hat{E}(z_1, z_2) = z_1 + z_2 - z_1 z_2 + \cdots$$

Here the $a_1 = 1$ makes the degree-2 term $-z_1 z_2$ nonzero, and the full complexity of the formal group law is visible from the start. This curve is the elliptic curve E_{11a} of conductor 11 (the smallest conductor of any elliptic curve over $\mathbb{Q}$), and its arithmetic will reappear in ??.

Exercise 4.2.1

1. Verify that $w(z) = z^3 + a_1 z^4 + (a_1^2 + a_2)z^5 + (a_1^3 + 2a_1 a_2 + a_3)z^6 + \cdots$ satisfies (4.4) through degree 6 by direct substitution.
2. The Weierstrass coefficients a_i have weight i , and z has weight -1 (since $z = -x/y$ and x has weight -2 , y has weight -3). Show that c_n (the coefficient of z^n in $w(z)$) has weight $n - 3$ in the a_i . (*Hint*: verify that equation (4.4) is weight-homogeneous of weight -3 , assigning $\text{wt}(z) = -1$ and $\text{wt}(w) = -3$.)
3. Compute c_7 in the expansion of $w(z)$, and verify the answer $c_7 = a_1^4 + 3a_1^2 a_2 + 3a_1 a_3 + a_2^2 + a_4$. Check that this has weight $4 = 7 - 3$ in the a_i .
4. For $E : y^2 = x^3 + Ax + B$ (short Weierstrass), verify directly that $\omega = dx/(2y) = (1 + O(z^4)) dz$ by computing dx/dz and $2y$ from the Laurent series $x = z^{-2} - Az^2 - Bz^4 - \cdots$ and $y = -z^{-3} + Az + Bz^3 + \cdots$.
5. Let (u, r, s, t) be an admissible change of variables (definition 1.1.4) transforming a Weierstrass equation E to E' . Show that the local parameters $z = -x/y$ and $z' = -x'/y'$ are related by $z = u^{-1}z' + (\text{higher order})$. Conclude that $\hat{E}$ and $\hat{E}'$ are isomorphic as formal group laws, with isomorphism given by a power series with leading coefficient u^{-1} .
6. For $E : y^2 = x^3 - x$, verify that $y^2 = x^3 - x$ holds as an identity in $K((z))$ through z^{-2} , using $x = z^{-2} + z^2 + \cdots$ and $y = -z^{-3} - z + \cdots$.

4.3 The Formal Logarithm and Formal Exponential

Over $\mathbb{C}$, the uniformization $\mathbb{C}/\Lambda \cong E(\mathbb{C})$ provides a logarithm for the group law on E : the local inverse of the exponential map $\mathbb{C} \rightarrow \mathbb{C}/\Lambda$ linearizes the addition on E near the identity. In the formal group setting, the analogous construction is the *formal logarithm*: a power series $\log_F(X) \in (R \otimes \mathbb{Q})[[X]]$ that converts the formal group law F into ordinary addition. Over $\mathbb{Q}$ -algebras, this trivializes the formal group completely — every one-dimensional commutative formal group is isomorphic to $\hat{\mathbb{G}}_a$. In

characteristic p , the logarithm breaks down (its coefficients develop denominators divisible by p), and the failure is governed by the $[p]$ -series, whose structure will lead us to the height in the next section.

The Formal Logarithm

Definition 4.3.1: The Formal Logarithm

Let F be a formal group law over a ring R , and let $\omega_F(T) = (1 + d_1T + d_2T^2 + \cdots) dT$ be its invariant differential (i.e., the unique normalized differential satisfying $F(X, Y)^*\omega_F = \omega_F$; see below). The *formal logarithm* of F is the power series

$$\log_F(X) = \int_0^X \omega_F(T) = X + \frac{d_1}{2}X^2 + \frac{d_2}{3}X^3 + \cdots \in (R \otimes_{\mathbb{Z}} \mathbb{Q})[[X]]. \quad (4.12)$$

The formal logarithm lives in $(R \otimes \mathbb{Q})[[X]]$, not $R[[X]]$: the integration introduces denominators. This is the source of all the subtlety in characteristic p .

Before proceeding, we need the invariant differential of a general formal group law. For the formal group $\hat{E}$ of an elliptic curve, we computed this in proposition 4.2.6: $\omega = (1 + a_1z + (a_1^2 + a_2)z^2 + \cdots) dz$. For an abstract formal group law $F(X, Y)$, the construction is:

Proposition 4.3.2: The Invariant Differential of a Formal Group

Let $F(X, Y)$ be a formal group law over R . Define

$$\omega_F(X) = P(X) dX \quad \text{where} \quad P(X) = \frac{1}{F_Y(X, 0)}, \quad (4.13)$$

and $F_Y = \partial F / \partial Y$ denotes the formal partial derivative with respect to the second variable. Then:

1. $\omega_F(X) = (1 + c_1X + c_2X^2 + \cdots) dX$ with $c_i \in R$.
2. ω_F is *translation-invariant*: $F(X, Y)^*\omega_F = \omega_F$. In terms of power series, this means $P(F(X, Y)) \cdot F_X(X, Y) = P(X)$.

Proof:

(1) Since $F(X, Y) = X + Y + \sum_{i+j \geq 2} a_{ij}X^iY^j$, we have $F_Y(X, 0) = 1 + \sum_{i \geq 1} a_{i1}X^i$. This is a unit in $R[[X]]$ (since the constant term is 1), so $P(X) = 1/F_Y(X, 0) = 1 + c_1X + c_2X^2 + \cdots$ for some $c_i \in R$.

(2) The invariance is a direct consequence of the associativity and commutativity axioms. Differentiating the associativity equation $F(F(X, Y), Z) = F(X, F(Y, Z))$ with respect to Z using the chain rule yields:

$$F_Y(F(X, Y), Z) = F_Y(X, F(Y, Z)) \cdot F_Y(Y, Z).$$

Evaluating this at $Z = 0$ and using $F(Y, 0) = Y$ gives:

$$F_Y(F(X, Y), 0) = F_Y(X, Y) \cdot F_Y(Y, 0).$$

Recalling that $P(T) = 1/F_Y(T, 0)$, we can rewrite this as:

$$P(F(X, Y))^{-1} = F_Y(X, Y) \cdot P(Y)^{-1} \implies F_Y(X, Y) = \frac{P(Y)}{P(F(X, Y))}.$$

By commutativity, $F(X, Y) = F(Y, X)$, so the partial derivative with respect to the first variable is $F_X(X, Y) = F_Y(Y, X)$. Swapping X and Y in our formula gives:

$$F_X(X, Y) = \frac{P(X)}{P(F(Y, X))} = \frac{P(X)}{P(F(X, Y))}.$$

Multiplying by $P(F(X, Y))$ yields exactly the invariance condition $P(F(X, Y)) \cdot F_X(X, Y) = P(X)$.

For the formal group $\hat{E}$ of an elliptic curve, the invariant differential ω_F from (4.13) agrees with $\omega = dx/(2y + a_1x + a_3)$ expressed in the local parameter z , as computed in proposition 4.2.6.

The Linearization Theorem

Theorem 4.3.3: The Formal Logarithm Linearizes

Let F be a formal group law over R , and let $\log_F(X) = \int_0^X \omega_F$ be the formal logarithm. Then:

$$\log_F(F(X, Y)) = \log_F(X) + \log_F(Y) \quad (4.14)$$

as an identity in $(R \otimes \mathbb{Q})[[X, Y]]$. In other words, $\log_F$ is a homomorphism from F to the additive formal group $\hat{\mathbb{G}}_a$.

Proof:

Define $\Phi(X, Y) = \log_F(F(X, Y)) - \log_F(X)$. We want to show $\Phi(X, Y) = \log_F(Y)$. We do this by showing Φ is independent of X , and then evaluating at $X = 0$.

Differentiating with respect to X using the chain rule gives:

$$\frac{\partial \Phi}{\partial X} = \log'_F(F(X, Y)) \cdot F_X(X, Y) - \log'_F(X).$$

By the definition of the formal logarithm, its derivative is exactly the power series of the invariant differential $\omega_F(T) = P(T) dT$, meaning $\log'_F(T) = P(T) = 1/F_Y(T, 0)$. Substituting this in yields:

$$\frac{\partial \Phi}{\partial X} = P(F(X, Y)) \cdot F_X(X, Y) - P(X).$$

But by the translation invariance of ω_F established in proposition 4.3.2(2), we know that $P(F(X, Y)) \cdot F_X(X, Y) = P(X)$. Therefore, $\partial \Phi / \partial X = 0$, meaning $\Phi(X, Y)$ is independent of X .

Evaluating Φ at $X = 0$ gives:

$$\Phi(0, Y) = \log_F(F(0, Y)) - \log_F(0) = \log_F(Y) - 0 = \log_F(Y).$$

Since Φ is independent of X , we conclude $\Phi(X, Y) = \Phi(0, Y) = \log_F(Y)$ for all X, Y . Rearranging the definition of Φ gives $\log_F(F(X, Y)) = \log_F(X) + \log_F(Y)$.

The linearization theorem says $\log_F : F \rightarrow \hat{\mathbb{G}}_a$ is a formal group homomorphism. Since $\log'_F(0) = 1 \neq 0$, it is an isomorphism over any ring where $\log_F$ has coefficients (i.e., over any $\mathbb{Q}$ -algebra). This gives:

Corollary 4.3.4: Classification over $\mathbb{Q}$ -Algebras

If R is a $\mathbb{Q}$ -algebra, every one-dimensional commutative formal group law over R is isomorphic to $\hat{\mathbb{G}}_a$, via the formal logarithm.

Proof:

The logarithm $\log_F(X) = X + \frac{d_1}{2}X^2 + \dots$ has coefficients in R (since R is a $\mathbb{Q}$ -algebra, the denominators $2, 3, \dots$ are invertible). Its leading coefficient is $1 \in R^*$, so it possesses a compositional inverse $\exp_F(X) \in R[[X]]$, making it an isomorphism. By theorem 4.3.3, $\log_F : F \xrightarrow{\sim} \hat{\mathbb{G}}_a$.

Over $\mathbb{Q}$ -algebras, there is no interesting theory of formal groups: they are all the same. The richness of the theory is entirely a characteristic- p (or mixed-characteristic) phenomenon.

The Formal Exponential

Definition 4.3.5: The Formal Exponential

Let F be a formal group law over a $\mathbb{Q}$ -algebra R . The *formal exponential* is the composition inverse of the formal logarithm:

$$\exp_F = \log_F^{-1} \in R[[X]], \quad \exp_F(\log_F(X)) = X = \log_F(\exp_F(X)).$$

It satisfies $F(X, Y) = \exp_F(\log_F(X) + \log_F(Y))$.

The formal exponential “recovers” the formal group law from the additive group: F is the conjugate of ordinary addition by the coordinate change $\exp_F$. This mirrors the Lie-theoretic picture: the exponential map converts the Lie algebra (with its additive structure) into the group (with its possibly non-commutative, non-additive operation). For one-dimensional commutative formal groups, the Lie algebra is always $\hat{\mathbb{G}}_a$, so the exponential always exists over $\mathbb{Q}$ -algebras.

The $[n]$ -Series via the Logarithm

The formal logarithm gives an elegant formula for the multiplication-by- n endomorphism.

Proposition 4.3.6: $[n]_F$ via the Logarithm

For any formal group law F over R and any integer n :

$$[n]_F(X) = \exp_F(n \cdot \log_F(X)) \quad (4.15)$$

as an identity in $(R \otimes \mathbb{Q})[[X]]$. Equivalently, $\log_F([n]_F(X)) = n \cdot \log_F(X)$.

Proof:

The logarithm is a homomorphism: $\log_F(F(X, Y)) = \log_F(X) + \log_F(Y)$. Applying this n times (with $Y = X$ at each step): $\log_F([n]_F(X)) = n \cdot \log_F(X)$. Applying $\exp_F$ to both sides gives (4.15).

This formula is simultaneously very useful (it makes $[n]_F$ computable) and very dangerous (it requires working over a $\mathbb{Q}$ -algebra). Over $\mathbb{Z}_p$, for instance, $\log_F(X)$ has coefficients with p -adic denominators, and $\exp_F(X)$ has even worse denominators. The formula (4.15) holds as an identity of power series with coefficients in $\mathbb{Q}_p$, but the result $[n]_F(X)$ has coefficients in $\mathbb{Z}_p$ — a miraculous cancellation of denominators that is not at all obvious from the right-hand side.

Example 4.3.7: The Logarithm of $\hat{\mathbb{G}}_m$

For the multiplicative formal group $F(X, Y) = X + Y + XY$, the invariant differential is

$$\omega_F = \frac{dX}{F_Y(X, 0)} = \frac{dX}{1 + X}.$$

The formal logarithm is

$$\log_{\hat{\mathbb{G}}_m}(X) = \int_0^X \frac{dT}{1+T} = X - \frac{X^2}{2} + \frac{X^3}{3} - \frac{X^4}{4} + \cdots = \log(1 + X),$$

the familiar Taylor series for the natural logarithm. The formal exponential is

$$\exp_{\hat{\mathbb{G}}_m}(X) = e^X - 1 = X + \frac{X^2}{2!} + \frac{X^3}{3!} + \cdots,$$

and indeed $F(X, Y) = (1 + X)(1 + Y) - 1 = \exp_{\hat{\mathbb{G}}_m}(\log(1 + X) + \log(1 + Y))$.

The $[n]$ -series: $[n]_{\hat{\mathbb{G}}_m}(X) = \exp_{\hat{\mathbb{G}}_m}(n \log_{\hat{\mathbb{G}}_m}(X)) = (1 + X)^n - 1$, matching the formula from example 4.1.10(b). In particular, $[p]_{\hat{\mathbb{G}}_m}(X) = (1 + X)^p - 1$, which in characteristic p reduces to X^p (by the binomial theorem).

Example 4.3.8: The Logarithm of $\hat{E}$ for $y^2 = x^3 - x$

For $E : y^2 = x^3 - x$ (short Weierstrass, $a_1 = a_2 = a_3 = 0$, $a_4 = -1$), the invariant differential is $\omega = (1 + O(z^4)) dz$ (example 4.2.8). The formal logarithm is

$$\log_{\hat{E}}(z) = z + O(z^5),$$

i.e., the logarithm agrees with the identity map through degree 4. The first correction term comes

from the coefficient of z^4 in ω . The formal group law $\hat{E}$ is “close to additive” through low degrees, and the logarithm reflects this: the correction terms that separate $\hat{E}$ from $\hat{\mathbb{G}}_a$ appear at high degree, arising from the $a_4 = -1$ coefficient.

The $[p]$ -Series in Characteristic p

The logarithm reveals its deepest information when we examine the $[p]$ -series in characteristic p . Over a ring R of characteristic p , the formula $[p]_F(X) = \exp_F(p \cdot \log_F(X))$ involves multiplying by $p = 0$, so naively $[p]_F = 0$. But this is wrong: $[p]_F$ is a power series with coefficients in R (where $p = 0$), and it is generally *not* zero. The formula (4.15) only holds over $\mathbb{Q}$ -algebras; over $\mathbb{F}_p$ -algebras, $[p]_F$ must be computed directly.

What the logarithm *does* tell us is the structure of $[p]_F$ modulo p , provided we can lift the formal group to characteristic zero. Suppose we lift F to a formal group law $\tilde{F}$ over a torsion-free ring like $\mathbb{Z}_p$, and write its formal logarithm as:

$$\log_{\tilde{F}}(X) = X + \sum_{n=2}^{\infty} \frac{b_n}{n} X^n \in \mathbb{Q}_p[[X]], \quad (4.16)$$

where the b_n are the coefficients of the invariant differential $\omega_{\tilde{F}}$. The denominators that cause trouble when reducing modulo p are exactly the multiples of p . The first problematic term is $\frac{b_p}{p} X^p$. Because $[p]_{\tilde{F}}(X) = \exp_{\tilde{F}}(p \log_{\tilde{F}}(X))$, the multiplication by p clears this first denominator, yielding an internal term $p \cdot \frac{b_p}{p} X^p = b_p X^p$. When we reduce back to characteristic p , this coefficient survives to become the leading term of the $[p]$ -series!

The key geometric observation is:

Proposition 4.3.9: Structure of the $[p]$ -Series

Let F be a formal group law over a ring R of characteristic $p > 0$. Then $[p]_F(X)$ has the form

$$[p]_F(X) = pX + (\text{terms involving both } p \text{ and higher powers of } X), \quad (4.17)$$

and since $p = 0$ in R , the linear term vanishes. The first possibly nonzero term of $[p]_F(X)$ occurs at degree p or higher. In fact, modulo (X^{p+1}) :

$$[p]_F(X) \equiv v_1 X^p \pmod{X^{p+1}} \quad (4.18)$$

for some $v_1 \in R$. If F is the reduction of a formal group $\tilde{F}$ in characteristic zero, v_1 is exactly the reduction modulo p of the invariant differential coefficient b_p .

Proof:

Since $[p]_F(X) = F(X, F(X, \dots))$ (p times), we have $[p]_F(X) = pX + O(X^2)$, and in characteristic p the linear term vanishes. For the degree- p claim: the formal group law $F(X, Y) = X + Y + (\deg \geq 2)$ involves only additions and multiplications. When we iterate F to sum p copies of X , the coefficients of the terms of degree $k < p$ arise from evaluating symmetric polynomials in p variables at $(X, \dots, X)$. These combinatorial evaluations always yield integer multiples of binomial coefficients $\binom{p}{l}$ for $1 \leq l \leq k < p$. Because $k < p$, these binomial coefficients are all

divisible by p , and hence vanish in characteristic p . The first term that does not necessarily vanish has degree p , corresponding to the unmixed p -th powers.

The coefficient v_1 is the crucial datum. If $v_1 \neq 0$ (i.e., $v_1 \in R^*$, which happens when R is a field), then $[p]_F(X) = v_1 X^p + (\text{higher})$, and the formal group “looks like $\hat{G}_m$ at first order” (recall $[p]_{\hat{G}_m}(X) = X^p$ in characteristic p). If $v_1 = 0$, we must look at higher powers of p — and this is the beginning of the height filtration that will be the subject of the next section.

Example 4.3.10: The $[p]$ -Series for $\hat{G}_a$ and $\hat{G}_m$

(a) $\hat{G}_a$: $[p]_{\hat{G}_a}(X) = pX = 0$ in characteristic p . Here $v_1 = 0$ and in fact $[p] = 0$ identically.

(b) $\hat{G}_m$: $[p]_{\hat{G}_m}(X) = (1 + X)^p - 1 = X^p$ in characteristic p . Here $v_1 = 1 \neq 0$.

For the formal group of an elliptic curve, the $[p]$ -series is never identically zero (since $[p] : E \rightarrow E$ is a nonzero isogeny with $\deg([p]) = p^2$), so we are always in the case $[p]_E \neq 0$. The coefficient v_1 encodes whether E is ordinary or supersingular — but we defer this analysis to section 4.4, where we define the height and prove the equivalence.

Convergence over Local Rings

The formal logarithm $\log_F(X) = X + d_1 X^2/2 + d_2 X^3/3 + \cdots$ has denominators, so it does not converge on all of $\mathfrak{m}$ for a local ring $(R, \mathfrak{m})$ of mixed characteristic $(0, p)$. However, it converges on a neighborhood of zero:

Proposition 4.3.11: Convergence of the Logarithm

Let $(R, \mathfrak{m}, k)$ be a complete DVR of mixed characteristic $(0, p)$ with fraction field K and valuation v , and let F be a formal group law over R . Then:

1. The formal logarithm $\log_F : \mathfrak{m} \rightarrow K$ converges on all of $\mathfrak{m}$ (i.e., the series $\log_F(a)$ converges in K for all $a \in \mathfrak{m}$).
2. The formal exponential $\exp_F$ converges on a smaller neighborhood of zero, specifically on the domain $\{a \in \mathfrak{m} \mid v(a) > \frac{v(p)}{p-1}\}$.
3. On the overlap, $\log_F$ and $\exp_F$ are inverse to each other, and $\log_F$ is a group isomorphism from $(\mathfrak{m}^n, +_F)$ to $(\mathfrak{m}^n, +)$ for sufficiently large n (depending on the ramification).

Proof Sketch:

(1) The coefficients of $\log_F$ are $d_n/(n+1)$, where $d_n \in R$. The p -adic valuation of $1/(n+1)$ grows at most like $\log_p(n)$, while $v(a^{n+1}) = (n+1)v(a) \rightarrow \infty$ for $a \in \mathfrak{m}$. So the terms $d_n a^{n+1}/(n+1)$ tend to zero p -adically.

(2) The coefficients of $\exp_F$ involve factorials: the n -th coefficient has p -adic valuation roughly $-n/(p-1)$ (from $v_p(n!) \approx n/(p-1)$). For convergence, we need $n \cdot v(a) - n/(p-1) \rightarrow \infty$, which requires $v(a) > v(p)/(p-1)$.

- (3) Follows from (1), (2), and the formal identity $\exp_F \circ \log_F = \text{id}$.

This convergence result is essential for the applications in section 4.5: it guarantees that the formal group $\hat{E}(\mathfrak{m}_K)$ is a genuine topological group, and that the logarithm isomorphism $\hat{E}(\mathfrak{m}_K^n) \cong \mathfrak{m}_K^n$ (as additive groups) holds for n sufficiently large.

4.3.12 Remark: The Logarithm and p -adic Analysis Over $\mathbb{Q}_p$, the formal logarithm of $\hat{\mathbb{G}}_m$ is $\log(1 + X) = X - X^2/2 + X^3/3 - \cdots$, which converges for $|X|_p < 1$ (i.e., $X \in p\mathbb{Z}_p$). The exponential $e^X - 1$ converges for $|X|_p < p^{-1/(p-1)}$, a smaller disc. The mismatch between the convergence radii of $\log$ and $\exp$ is a fundamental feature of p -adic analysis, and it persists for all formal groups. It is this mismatch that makes the formal group structure on $\hat{E}(\mathfrak{m}_K)$ nontrivial: the logarithm linearizes the group law on a large disc, but the exponential only recovers the group structure on a smaller disc.

Exercise 4.3.1

1. Verify directly that $\log_{\hat{\mathbb{G}}_m}(F(X, Y)) = \log_{\hat{\mathbb{G}}_m}(X) + \log_{\hat{\mathbb{G}}_m}(Y)$ for $F(X, Y) = X + Y + XY$, using $\log_{\hat{\mathbb{G}}_m}(X) = \log(1 + X)$. (*Hint*: $\log(1 + X + Y + XY) = \log((1 + X)(1 + Y))$.)
2. Let $F(X, Y) = X + Y + cXY$ for $c \in R$. Compute the invariant differential ω_F and the formal logarithm $\log_F(X)$. Verify that $\log_F(F(X, Y)) = \log_F(X) + \log_F(Y)$ through degree 3.
3. For $\hat{\mathbb{G}}_m$ in characteristic $p = 3$, compute $[3]_{\hat{\mathbb{G}}_m}(X) = (1 + X)^3 - 1$ directly and verify that the result is X^3 (since $3X + 3X^2 = 0$ in characteristic 3). What is $[9]_{\hat{\mathbb{G}}_m}(X)$ in characteristic 3?
4. For the formal group $\hat{E}$ of $E : y^2 = x^3 - x$, the invariant differential is $\omega = (1 + 0 \cdot z + 0 \cdot z^2 + 0 \cdot z^3 + Az^4 + \cdots) dz$ for some integer A . Compute the formal logarithm $\log_{\hat{E}}(z)$ through degree 5, and identify the first term with a denominator not in $\mathbb{Z}$.
5. Let E be an elliptic curve over $\mathbb{F}_p$. Explain why $[p]_{\hat{E}} \neq 0$ (i.e., $[p]_{\hat{E}}$ is not the zero power series), using the fact that $\deg([p]) = p^2$ on E . (*Hint*: if $[p]_{\hat{E}} = 0$, what would this say about the kernel of $[p]$ near O ?)
6. Let F be a formal group law over $\mathbb{Z}$. Using $\log_F$ (which lives in $\mathbb{Q}[[X]]$), show that $\text{End}_{\mathbb{Z}}(F) \cong \mathbb{Z}$ (i.e., the only endomorphisms with integer coefficients are the multiplication-by- n maps). (*Hint*: over $\mathbb{Q}$, any endomorphism f satisfies $\log_F \circ f = c \cdot \log_F$ for some $c \in \mathbb{Q}$; determine what c must be if $f(X)$ has coefficients in $\mathbb{Z}$.)

4.4 Height and the $[p]$ -Series

The formal logarithm showed that over $\mathbb{Q}$ -algebras, all formal groups are isomorphic to $\hat{\mathbb{G}}_a$. In characteristic p , the logarithm breaks down, and the formal group carries a fundamental invariant — its *height* — that measures how far it is from the additive formal group. For the formal group of an elliptic curve, the height turns out to be 1 or 2, and this dichotomy is exactly the ordinary/supersingular classification from theorem 2.4.4. This section defines the height, proves it is well-defined, and establishes the equivalence with the characterizations given in section 2.4.

Definition of Height

Let R be a commutative ring of characteristic $p > 0$, and let F be a formal group law over R .

By proposition 4.3.9, the $[p]$ -series $[p]_F(X) \in R[[X]]$ has vanishing linear term ($[p]_F'(0) = p = 0$), so the first nonzero coefficient occurs at some degree ≥ 2 — or $[p]_F = 0$ identically. The key observation is that this first nonzero term always occurs at a p -power degree.

Proposition 4.4.1: The $[p]$ -Series Has p -Power Leading Term

Let F be a formal group law over a ring R of characteristic p . If $[p]_F \neq 0$, then there exists a unique integer $h \geq 1$ and a power series $g(X) \in R[[X]]$ with $g'(0) \neq 0$ such that

$$[p]_F(X) = g(X^{p^h}). \quad (4.19)$$

Equivalently: $[p]_F(X) = uX^{p^h} + (\text{terms of degree } > p^h)$ with $u \in R$, $u \neq 0$, and all terms of degree $< p^h$ vanish.

Proof:

Define $f = [p]_F$. We prove: if $f(X) \neq 0$ and $f'(X) = 0$ (as a power series — meaning every coefficient $nc_n = 0$, which in characteristic p means $c_n = 0$ whenever $p \nmid n$), then $f(X) = f_1(X^p)$ for some $f_1 \in R[[X]]$. Indeed, $f'(X) = 0$ means $f(X) = \sum_{n \geq 0} c_{pn} X^{pn} = \sum_{n \geq 0} c_{pn} (X^p)^n = f_1(X^p)$.

Now apply this iteratively. We have $[p]_F'(0) = p = 0$, but we need to analyze the full derivative $[p]_F'(X)$. The argument proceeds as follows. The endomorphism $[p]_F$ satisfies the functional equation $[p]_F(F(X, Y)) = F([p]_F(X), [p]_F(Y))$. Differentiating with respect to Y and setting $Y = 0$:

$$[p]_F'(F(X, 0)) \cdot F_Y(X, 0) = F_Y([p]_F(X), [p]_F(0)) \cdot [p]_F'(0).$$

Using $F(X, 0) = X$, $[p]_F(0) = 0$, and $[p]_F'(0) = 0$:

$$[p]_F'(X) \cdot F_Y(X, 0) = F_Y(0, 0) \cdot 0 = 0.$$

Since $F_Y(X, 0) = 1 + \dots$ is a unit in $R[[X]]$, we conclude $[p]_F'(X) = 0$ as a power series.

By the observation above, $[p]_F(X) = f_1(X^p)$ for some $f_1 \in R[[X]]$. Now, f_1 is also an endomorphism of a formal group (the “Frobenius twist” of F), and we can ask whether $f_1'(X) = 0$. If $f_1'(0) \neq 0$, we stop: $h = 1$ and $g = f_1$. If $f_1'(0) = 0$, we apply the same argument (using the functional equation for f_1 , which inherits a suitable functional equation from $[p]_F$) to write $f_1(X) = f_2(X^p)$, giving $[p]_F(X) = f_2(X^{p^2})$. Continuing: either $f_k'(0) \neq 0$ at some stage (giving $h = k$), or $[p]_F$ involves only p^k -th powers for all k , forcing $[p]_F = 0$.

Since we assumed $[p]_F \neq 0$, the process must terminate at some $h \geq 1$, giving $[p]_F(X) = g(X^{p^h})$ with $g'(0) \neq 0$.

Definition 4.4.2: Height of a Formal Group

Let F be a formal group law over a ring R of characteristic p .

1. If $[p]_F \neq 0$, the *height* of F is the unique integer $h \geq 1$ such that $[p]_F(X) = g(X^{p^h})$ with $g'(0) \neq 0$ (as in proposition 4.4.1).
2. If $[p]_F = 0$, the *height* of F is $h = \infty$.

We write $\text{ht}(F) = h$.

The height measures the “depth” of vanishing of $[p]_F$ at the origin. At height 1, the leading term is uX^p (like $\hat{\mathbb{G}}_m$); at height 2, it is uX^{p^2} ; and at height ∞ , the series vanishes entirely (like $\hat{\mathbb{G}}_a$). The height is an isomorphism invariant: if $F \cong G$ (i.e., there exists an isomorphism $f : F \xrightarrow{\sim} G$), then $\text{ht}(F) = \text{ht}(G)$, since $[p]_G = f \circ [p]_F \circ f^{-1}$ and conjugation by an isomorphism does not change the p -power structure.

Example 4.4.3: Heights of the Classical Formal Groups

(a) $\hat{\mathbb{G}}_a$: $[p](X) = pX = 0$. Height $= \infty$.

(b) $\hat{\mathbb{G}}_m$: $[p](X) = (1 + X)^p - 1 = X^p$ in characteristic p . Here $g(X) = X$ with $g'(0) = 1 \neq 0$. Height $= 1$.

Height of the Formal Group of an Elliptic Curve

We now prove the main theorem of this section: for the formal group $\hat{E}$ of an elliptic curve E over a field of characteristic $p > 0$, the height is either 1 or 2, and the two cases correspond precisely to ordinary and supersingular.

Theorem 4.4.4: Height and the Ordinary/Supersingular Dichotomy

Let E/k be an elliptic curve over a field k of characteristic $p > 0$. Then:

1. $\text{ht}(\hat{E}) = 1$ if and only if E is ordinary (i.e., $E[p](\bar{k}) \cong \mathbb{Z}/p\mathbb{Z}$).
2. $\text{ht}(\hat{E}) = 2$ if and only if E is supersingular (i.e., $E[p](\bar{k}) = \{O\}$).

In particular, $\text{ht}(\hat{E}) \in \{1, 2\}$: it is never ∞ and never > 2 .

Proof:

Step 1: $\text{ht}(\hat{E}) \leq 2$. The map $[p] : E \rightarrow E$ is a nonzero isogeny of degree p^2 . In the formal group, $[p]_{\hat{E}}(z)$ is a nonzero power series (as argued in proposition 4.3.9: if $[p]_{\hat{E}} = 0$, then $[p]$ would kill an infinite neighborhood of O , contradicting the finiteness of $\ker([p])$). So $\text{ht}(\hat{E}) \neq \infty$.

Moreover, $[p]_{\hat{E}}(z) = g(z^{p^h})$ with $g'(0) \neq 0$ means the “local degree” of $[p]_{\hat{E}}$ at $z = 0$ is p^h (the power series vanishes to order p^h at the origin). Since $[p] : E \rightarrow E$ has degree p^2 , and the local

degree at O cannot exceed the global degree, we have $p^h \leq p^2$, giving $h \leq 2$.

We make this precise. The inseparable degree of $[p]$ equals p^h : the map $[p]_{\hat{E}}(z) = g(z^{p^h})$ factors through the p^h -th power map $z \mapsto z^{p^h}$ (which is the local expression of the iterated Frobenius $\varphi_{p^h} : E \rightarrow E^{(p^h)}$), followed by a power series g with $g'(0) \neq 0$ (which is a local isomorphism). Therefore $\deg_i([p]) = p^h$. Since $\deg_i([p])$ divides $\deg([p]) = p^2$, we have $p^h | p^2$, so $h \leq 2$.

Step 2: $h \neq 0$. We need $h \geq 1$. This follows from $[p]^*\omega = p\omega = 0$ (since $p = 0$ in k), which means $[p]$ is inseparable (proposition 2.3.5), hence $\deg_i([p]) \geq p$, so $h \geq 1$.

Step 3: $h = 1 \iff$ **ordinary**, $h = 2 \iff$ **supersingular**. From Step 1, $\deg_i([p]) = p^h$. By theorem 2.4.4:

$$\begin{aligned} E \text{ ordinary} &\iff \deg_i([p]) = p \iff h = 1, \\ E \text{ supersingular} &\iff \deg_i([p]) = p^2 \iff h = 2. \end{aligned}$$

This theorem completes the proof of the equivalence (b) $\Leftrightarrow$ (d) in theorem 2.4.4, which was deferred to this chapter. The height is the formal-group theoretic incarnation of the ordinary/supersingular distinction: ordinary curves have a formal group that “looks like $\hat{\mathbb{G}}_m$ at first order” (height 1, with $[p](z) = uz^p + \cdots$ for $u \neq 0$), while supersingular curves have a formal group that is “more inseparable” (height 2, with $[p](z)$ vanishing to order p^2).

The Inseparable Factorization

The proof above used the relationship between height and inseparable degree. Let us make this precise.

Proposition 4.4.5: Height Equals Inseparable Exponent

Let F be a formal group law of finite height h over a perfect field k of characteristic p . Then the endomorphism $[p]_F$ factors as

$$[p]_F = g \circ \varphi^h, \tag{4.20}$$

where $\varphi(X) = X^p$ is the Frobenius endomorphism on F (more precisely, φ is a homomorphism $F \rightarrow F^{(p)}$ where $F^{(p)}$ is the formal group obtained by raising the coefficients of F to the p -th power), and g is an isomorphism from $F^{(p^h)}$ to F (over $\bar{k}$).

For the formal group of an elliptic curve, this factorization is the local version of the global factorization $[p] = V_p \circ \varphi_p$ from theorem 2.4.3 (or its iterated version). The Frobenius φ_p contributes one factor of X^p to the inseparable degree, and the Verschiebung V_p contributes either 1 (ordinary: V_p is separable, so its local expression has $g'(0) \neq 0$) or X^p (supersingular: V_p is inseparable). The total is p^1 or p^2 , giving height 1 or 2.

Explicit $[p]$ -Series for Running Examples

Example 4.4.6: $[p]$ -Series for $y^2 = x^3 - x$ over $\mathbb{F}_p$

Consider $E : y^2 = x^3 - x$ (with $j = 1728$) over $\mathbb{F}_p$ for various primes p .

$p = 5$: From example 2.4.5(a), $a_5 = -2 \neq 0$, so E is ordinary over $\mathbb{F}_5$. The formal group has height 1, so $[5]_{\hat{E}}(z) = uz^5 + (\text{higher})$ with $u \neq 0$. If we were to compute the $[5]$ -series by iterating the formal group law five times ($[5](z) = \hat{E}(z, [4](z))$), we would observe exactly this. Since all $a_i = 0$ except $a_4 = -1$, the formal group law starts as $\hat{E}(z_1, z_2) = z_1 + z_2 + O(\deg 5)$, and in characteristic 5:

$$[5]_{\hat{E}}(z) = 5z + \cdots = 0 + \cdots,$$

with the first nonzero term appearing exactly at z^5 (height 1).

$p = 3$: We have $|E(\mathbb{F}_3)| = 4$ (points: $O, (0, 0), (1, 0), (-1, 0)$), so $a_3 = 3 + 1 - 4 = 0$. Since $p \mid a_p$, E is supersingular over $\mathbb{F}_3$. The formal group has height 2, so $[3]_{\hat{E}}(z) = uz^9 + (\text{higher})$ with $u \neq 0$: the $[3]$ -series vanishes to order $9 = 3^2$.

Example 4.4.7: $[p]$ -Series for $y^2 = x^3 + 1$ over $\mathbb{F}_p$

Consider $E : y^2 = x^3 + 1$ (with $j = 0$) over $\mathbb{F}_p$.

$p = 5$: From example 2.4.5(b), $|E(\mathbb{F}_5)| = 6$, so $a_5 = 0$ and E is supersingular. Height = 2: $[5]_{\hat{E}}(z) = uz^{25} + \cdots$ with $u \neq 0$.

$p = 7$: We compute $|E(\mathbb{F}_7)|$: for $x = 0, 1, \dots, 6$, evaluate $x^3 + 1$ and count quadratic residues. One finds $|E(\mathbb{F}_7)| = 7 + 1 - a_7$ with $a_7 \neq 0$ and $7 \nmid a_7$, so E is ordinary. Height = 1.

The p -Divisible Group

The formal group $\hat{E}$ captures the “connected component” of the p -torsion of E , but the full p -torsion has both a connected and an étale part. To see this, we introduce:

Definition 4.4.8: The p -Divisible Group

Let E/k be an elliptic curve over a field of characteristic p . The p -divisible group (or Barsotti–Tate group) of E is the inductive system

$$E[p^\infty] = \varinjlim E[p^n],$$

where the transition maps are the inclusions $E[p^n] \hookrightarrow E[p^{n+1}]$.

Over an algebraically closed field, $E[p^n]$ is a finite group scheme of order p^{2n} , but in characteristic p it is never an étale group scheme, meaning it is not reduced (its affine coordinate ring contains nonzero nilpotents). The structure of $E[p^n]$ as a group scheme decomposes into two pieces:

Proposition 4.4.9: Connected-Étale Decomposition

Let E/k be an elliptic curve over a perfect field of characteristic p . The group scheme $E[p^n]$ has a canonical short exact sequence

$$0 \longrightarrow E[p^n]^0 \longrightarrow E[p^n] \longrightarrow E[p^n]^{\text{ét}} \longrightarrow 0, \quad (4.21)$$

where $E[p^n]^0$ is the connected component of the identity (a local group scheme, corresponding to the formal group) and $E[p^n]^{\text{ét}}$ is an étale group scheme.

The two cases are:

1. (*Ordinary.*) Over $\bar{k}$, $E[p^n]^0 \cong \mu_{p^n}$ (a local group scheme of order p^n) and $E[p^n]^{\text{ét}} \cong \mathbb{Z}/p^n\mathbb{Z}$ (a constant group scheme of order p^n). The connected part and the étale part both have order p^n .
2. (*Supersingular.*) $E[p^n]^0 = E[p^n]$ (the entire group scheme is connected) and $E[p^n]^{\text{ét}} = 0$. The étale part is trivial: there are no “geometric” p^n -torsion points other than the identity.

Proof Key ideas:

The connected-étale decomposition is a general fact about finite group schemes over a perfect field [24]: every finite commutative group scheme G over a perfect field admits a unique short exact sequence $0 \rightarrow G^0 \rightarrow G \rightarrow G^{\text{ét}} \rightarrow 0$ where G^0 is connected (local) and $G^{\text{ét}}$ is étale. Over an algebraically closed field, $G^{\text{ét}}(\bar{k}) = G(\bar{k})$ (the geometric points), and G^0 has no geometric points other than the identity.

For $G = E[p^n]$: the geometric points $E[p^n](\bar{k})$ form a group of order p^n (ordinary) or 1 (supersingular), by theorem 2.4.4 and induction on n . In the ordinary case, $|E[p^n](\bar{k})| = p^n$, so $E[p^n]^{\text{ét}}$ has order p^n and $E[p^n]^0$ has order $p^{2n}/p^n = p^n$. In the supersingular case, $|E[p^n](\bar{k})| = 1$, so $E[p^n]^{\text{ét}} = 0$ and $E[p^n]^0 = E[p^n]$ has order p^{2n} .

The identification of $E[p^n]^0$ with the formal group is the content of the kernel-of-reduction isomorphism that will be established in section 4.5: over a complete local ring, the connected part of $E[p^n]$ is precisely the subgroup of p^n -torsion in $\hat{E}(\mathfrak{m})$.

The connected-étale decomposition reflects the height: a formal group of height h has “connected p^n -torsion” of order p^{hn} . For $\hat{E}$, this is p^n (height 1, ordinary) or p^{2n} (height 2, supersingular). The étale part accounts for the remaining $p^{2n}/p^{hn} = p^{(2-h)n}$ points: this gives p^n geometric p^n -torsion points when $h = 1$, and 1 geometric point when $h = 2$.

The Cartier–Dieudonné Classification

Over an algebraically closed field, formal groups of a given height are essentially unique:

Theorem 4.4.10: Classification of Formal Groups by Height

Let $k = \overline{\mathbb{F}_p}$. For each $h \in \{1, 2, 3, \dots\} \cup \{\infty\}$, there exists (up to isomorphism) exactly one one-dimensional commutative formal group law of height h over k .

Proof Key ideas:

This is a deep result in the theory of formal groups, proved using the Dieudonné module classification. We sketch the idea.

To each formal group F over k , one associates its *Dieudonné module* $\mathbb{D}(F)$, which is a module over the non-commutative ring $W(k)\langle \mathcal{F}, \mathcal{V} \rangle / (\mathcal{F}\mathcal{V} = \mathcal{V}\mathcal{F} = p)$ (where $W(k)$ is the ring of Witt vectors, $\mathcal{F}$ is the Frobenius operator, and $\mathcal{V}$ is the Verschiebung operator). For a one-dimensional formal group of height h , the Dieudonné module is a free $W(k)$ -module of rank h , and the classification reduces to showing that all such modules of a given rank are isomorphic over $\overline{\mathbb{F}_p}$.

For height 1: the unique formal group is $\hat{\mathbb{G}}_m$ (up to isomorphism). For height 2: the unique formal group is the one arising from any supersingular elliptic curve over $\overline{\mathbb{F}_p}$. For the full proof, see Lazard [12] for the existence and Honda [11] for the classification via Honda systems. See also Hazewinkel [10] for a comprehensive treatment.

For elliptic curves, theorem 4.4.10 has the following striking consequence:

Corollary 4.4.11: Formal Groups of Elliptic Curves over $\overline{\mathbb{F}_p}$

1. All ordinary elliptic curves over $\overline{\mathbb{F}_p}$ have the same formal group: $\hat{E} \cong \hat{\mathbb{G}}_m$ over $\overline{\mathbb{F}_p}$.
2. All supersingular elliptic curves over $\overline{\mathbb{F}_p}$ have isomorphic formal groups (a unique height-2 formal group, sometimes denoted $\hat{\mathbb{G}}_{1/2}$ or $\hat{H}_2$).

This is a remarkable rigidity result: the formal group, which a priori depends on the specific Weierstrass equation, is determined (up to isomorphism over $\overline{\mathbb{F}_p}$) solely by the height, which in turn is determined by the single bit of information “ordinary or supersingular.” Over non-algebraically-closed fields, formal groups of the same height need not be isomorphic (they can differ by “twists”), but over $\overline{\mathbb{F}_p}$ the height is a complete invariant.

Summary Table

Type	$\text{ht}(\hat{E})$	$[p]_{\hat{E}}(z)$	$ E[p](\bar{k}) $	$\hat{E} \cong$ (over $\overline{\mathbb{F}_p}$)
Ordinary	1	$uz^p + \dots$ ($u \neq 0$)	p	$\hat{\mathbb{G}}_m$
Supersingular	2	$uz^{p^2} + \dots$ ($u \neq 0$)	1	unique height-2
$\hat{\mathbb{G}}_a$ (not from EC)	∞	0	—	$\hat{\mathbb{G}}_a$

The row for $\hat{\mathbb{G}}_a$ is included for comparison: it is the formal group of the additive group $\mathbb{G}_a$, which arises as the formal group of a *cuspidal* singular cubic (from section 1.6), not of any elliptic curve.

The fact that $\text{ht}(\hat{E}) \neq \infty$ for any elliptic curve reflects the fact that $[p]$ on an elliptic curve is always a nonzero isogeny.

Exercise 4.4.1

1. Show that the height is an isomorphism invariant: if $f : F \xrightarrow{\sim} G$ is an isomorphism of formal group laws over a field of characteristic p , then $\text{ht}(F) = \text{ht}(G)$. (*Hint:* $[p]_G = f \circ [p]_F \circ f^{-1}$, and f has $f'(0) \neq 0$.)
2. Let $\varphi : E_1 \rightarrow E_2$ be an isogeny of degree d between elliptic curves over a field of characteristic p . Show that φ induces a homomorphism $\hat{\varphi} : \hat{E}_1 \rightarrow \hat{E}_2$ of formal groups, and that $\hat{\varphi}(z) = cz^{p^e} + (\text{higher})$ with $c \neq 0$ and $p^e = \deg_i(\varphi)$. In particular, if $\varphi = [p]$, then $e = h = \text{ht}(\hat{E})$.
3. For $E : y^2 = x^3 - x$ over $\mathbb{F}_3$, verify directly that $|E(\mathbb{F}_3)| = 4$ and $a_3 = 0$, confirming supersingularity. Using $\text{ht}(\hat{E}) = 2$, what is the order of vanishing of $[3]_{\hat{E}}(z)$ at $z = 0$?
4. For $E : y^2 + y = x^3$ over $\mathbb{F}_2$, compute $|E(\mathbb{F}_2)|$ and determine whether E is ordinary or supersingular. What is $\text{ht}(\hat{E})$?
5. Show that for formal groups F, G of finite heights h_1, h_2 over a field of characteristic p , the product formal group $F \times G$ (defined by $(F \times G)((X_1, Y_1), (X_2, Y_2)) = (F(X_1, X_2), G(Y_1, Y_2))$, a two-dimensional formal group) has height $h_1 + h_2$. (*Hint:* For higher-dimensional formal groups, the height h is defined such that the local degree of the $[p]$ map at the origin is p^h . Show that the local degree of the multi-variable map $[p]_{F \times G}(X, Y) = ([p]_F(X), [p]_G(Y))$ is the product $p^{h_1} \cdot p^{h_2} = p^{h_1+h_2}$.)
6. Let $E/\overline{\mathbb{F}_p}$ be a supersingular elliptic curve. By theorem 4.4.10, $\hat{E}$ is the unique height-2 formal group over $\overline{\mathbb{F}_p}$. Show that $\text{End}(\hat{E})$ is strictly larger than $\mathbb{Z}_p$ (the p -adic completion of $\mathbb{Z}$). (*Hint:* $\text{End}(\hat{E})$ contains $\mathbb{Z}_p$ via the $[n]$ -maps. If $\text{End}(\hat{E}) = \mathbb{Z}_p$, then $\hat{E}$ would be determined by its $[p]$ -series alone, but a height-2 formal group has “more room” for endomorphisms than $\hat{\mathbb{G}}_m$ does.)

4.5 Formal Groups over Local Fields

We now arrive at the arithmetic payoff of the formal group machinery. Throughout this section, K denotes a local field: K is the fraction field of a complete discrete valuation ring R with maximal ideal $\mathfrak{m} = (\pi)$ and residue field $k = R/\mathfrak{m}$ of characteristic $p > 0$. The standard examples are $K = \mathbb{Q}_p$ (with $R = \mathbb{Z}_p$, $\pi = p$, $k = \mathbb{F}_p$) and $K = \mathbb{F}_q((t))$ (with $R = \mathbb{F}_q[[t]]$, $\pi = t$, $k = \mathbb{F}_q$). We write $v : K^* \rightarrow \mathbb{Z}$ for the normalized valuation.

Let E/K be an elliptic curve given by a Weierstrass equation with coefficients in R (an *integral* Weierstrass model). Reducing the coefficients modulo $\mathfrak{m}$ gives a curve $\tilde{E}/k$, which may or may not be smooth. When $\tilde{E}$ is smooth (i.e., $\Delta \not\equiv 0 \pmod{\pi}$, the case of *good reduction*), the reduction map $E(K) \rightarrow \tilde{E}(k)$ is a group homomorphism, and its kernel is precisely the group of points in the formal group. This is the content of the main theorem of this section. The full theory of reduction types (good, multiplicative, additive) and minimal models will be developed in chapter 6; here we focus on good reduction, where the formal group tells the whole story.

The Reduction Map

Let E/K be given by a Weierstrass equation (1.2) with $a_i \in R$. A point $P = (x_0, y_0) \in E(K)$ with $x_0, y_0 \in R$ reduces to $\tilde{P} = (\tilde{x}_0, \tilde{y}_0) \in \tilde{E}(k)$ (where $\tilde{x}_0 = x_0 \bmod \mathfrak{m}$). But not every K -rational point has integral coordinates: the coordinates x_0, y_0 may have negative valuations. To handle this, we pass to projective coordinates.

Write $P = [X_0 : Y_0 : Z_0]$ with $X_0, Y_0, Z_0 \in R$, not all in $\mathfrak{m}$ (we can always scale to achieve this). The reduction $\tilde{P} = [\tilde{X}_0 : \tilde{Y}_0 : \tilde{Z}_0] \in \mathbb{P}^2(k)$ is well-defined. This gives a map $\text{red} : E(K) \rightarrow \tilde{E}(k)$ on points.

Definition 4.5.1: The Kernel of Reduction

Let E/K have an integral Weierstrass model with good reduction. The subgroups $E_0(K)$ and $E_1(K)$ are defined by:

$$\begin{aligned} E_0(K) &= \{P \in E(K) \mid \tilde{P} \in \tilde{E}_{\text{ns}}(k)\} \quad (\text{points reducing to the nonsingular locus}), \\ E_1(K) &= \ker(\text{red} : E(K) \rightarrow \tilde{E}(k)) = \{P \in E(K) \mid \tilde{P} = \tilde{O}\}. \end{aligned}$$

When E has good reduction, $\tilde{E}_{\text{ns}} = \tilde{E}$, so $E_0(K) = E(K)$ and $E_1(K) = \ker(\text{red})$.

Thus, for good reduction, we have an exact sequence

$$0 \longrightarrow E_1(K) \longrightarrow E(K) \xrightarrow{\text{red}} \tilde{E}(k) \longrightarrow 0, \quad (4.22)$$

where the surjectivity of the reduction map follows from Hensel's lemma (theorem 4.1.6): a smooth point $\tilde{Q} \in \tilde{E}(k)$ lifts to a point $Q \in E(K)$ because the Jacobian criterion for $\tilde{E}$ at $\tilde{Q}$ provides a non-degenerate derivative, giving the initial approximation needed for Hensel's lifting.

The group $E_1(K)$ consists of points whose projective coordinates satisfy $[\tilde{X}_0 : \tilde{Y}_0 : \tilde{Z}_0] = [0 : 1 : 0] = \tilde{O}$ in $\mathbb{P}^2(k)$. In affine coordinates, this means $P = (x_0, y_0)$ with $v(x_0) \leq -2$ and $v(y_0) \leq -3$ (the coordinates “blow up” as the point approaches O), or more precisely, P lies in the domain where the local parameter $z = -x/y$ has positive valuation.

The Main Isomorphism

Theorem 4.5.2: Kernel of Reduction $\cong$ Formal Group

Let E/K have an integral Weierstrass model with good reduction. There is a canonical group isomorphism

$$E_1(K) \xrightarrow{\sim} \hat{E}(\mathfrak{m}), \quad (4.23)$$

where $\hat{E}(\mathfrak{m})$ denotes the group $(\mathfrak{m}, +_{\hat{E}})$: the maximal ideal $\mathfrak{m}$ equipped with the group operation given by the formal group law $\hat{E}(z_1, z_2)$ of theorem 4.2.4.

Proof:

The isomorphism is given by the local parameter: $P \mapsto z(P) = -x(P)/y(P)$.

Step 1: The map $P \mapsto z(P)$ lands in $\mathfrak{m}$. For $P \in E_1(K)$, we have $\tilde{P} = \tilde{O}$, which means

$(z(P), w(P))$ reduces to $(0, 0)$ modulo $\mathfrak{m}$ in the (z, w) -coordinates of proposition 4.2.1. In particular, $z(P) \in \mathfrak{m}$.

Step 2: The map is a group homomorphism. If $P_1, P_2 \in E_1(K)$, then $z(P_1 + P_2) = \hat{E}(z(P_1), z(P_2))$ by definition of the formal group law (theorem 4.2.4). The power series $\hat{E}(z_1, z_2)$ converges for $z_1, z_2 \in \mathfrak{m}$ because the terms of $\hat{E}$ have no constant term, and each homogeneous component lands in a higher power of $\mathfrak{m}$.

Step 3: Injectivity. If $z(P) = 0$, then P has $(z, w) = (0, 0)$ in the local coordinates, which corresponds to $P = O$.

Step 4: Surjectivity. Given $z_0 \in \mathfrak{m}$, we need to produce a point $P \in E_1(K)$ with $z(P) = z_0$. The w -expansion (proposition 4.2.2) gives $w_0 = w(z_0) = z_0^3 + a_1 z_0^4 + \cdots \in \mathfrak{m}^3$. Since $z_0 \in \mathfrak{m}$ and the series converges in R (all coefficients lie in R and the terms tend to zero $\mathfrak{m}$ -adically), we get $w_0 \in R$. Then $x_0 = z_0/w_0$ and $y_0 = -1/w_0$ give a point $P = (x_0, y_0) \in E(K)$ (it satisfies the Weierstrass equation because $w(z)$ was constructed as the solution of (4.4)).

We must verify $P \in E_1(K)$, i.e., $\tilde{P} = \tilde{O}$. In projective coordinates, clearing denominators gives $P = [x_0 : y_0 : 1] = [z_0/w_0 : -1/w_0 : 1] = [-z_0 : 1 : -w_0]$. Since $z_0 \in \mathfrak{m}$ and $w_0 \in \mathfrak{m}^3$, these integral projective coordinates satisfy $v(-z_0) \geq 1$, $v(1) = 0$, $v(-w_0) \geq 3$. Reducing modulo $\mathfrak{m}$ gives $[\tilde{X}_0 : \tilde{Y}_0 : \tilde{Z}_0] = [0 : 1 : 0] = \tilde{O}$. So $P \in E_1(K)$.

This theorem is the fundamental link between the formal group (an algebraic object defined by a power series) and the arithmetic of $E(K)$ (a group of rational points). It says the kernel of reduction is *completely controlled* by the formal group: every question about $E_1(K)$ can be answered by analyzing power series in one variable over R .

The Filtration

The formal group $\hat{E}(\mathfrak{m})$ carries a natural filtration by powers of $\mathfrak{m}$:

$$\hat{E}(\mathfrak{m}) \supset \hat{E}(\mathfrak{m}^2) \supset \hat{E}(\mathfrak{m}^3) \supset \cdots,$$

where $\hat{E}(\mathfrak{m}^n) = \left\{ z \in \mathfrak{m}^n \mid z \text{ is an element of } \hat{E}(\mathfrak{m}) \right\}$. Under the isomorphism of theorem 4.5.2, this corresponds to the subgroups $E_n(K) = \{P \in E_1(K) \mid v(z(P)) \geq n\}$.

Proposition 4.5.3: The Filtration and its Quotients

Let E/K have good reduction, with residue field k of characteristic p .

1. For each $n \geq 1$, $\hat{E}(\mathfrak{m}^n)$ is a subgroup of $\hat{E}(\mathfrak{m})$.
2. The successive quotients are:

$$\hat{E}(\mathfrak{m}^n)/\hat{E}(\mathfrak{m}^{n+1}) \cong k^+ \quad (\text{the additive group of } k) \quad (4.24)$$

for all $n \geq 1$. The isomorphism sends $z \in \mathfrak{m}^n$ to $z/\pi^n \bmod \mathfrak{m} \in k$.

3. $\hat{E}(\mathfrak{m}) = \varprojlim \hat{E}(\mathfrak{m})/\hat{E}(\mathfrak{m}^n)$, so $\hat{E}(\mathfrak{m})$ is a pro- p group (when $\text{char } k = p$).

Proof:

(1) If $z_1, z_2 \in \mathfrak{m}^n$, then $\hat{E}(z_1, z_2) = z_1 + z_2 + (\text{terms of degree } \geq 2 \text{ in } z_1, z_2)$. Each degree- ≥ 2 term involves $z_1^i z_2^j$ with $i + j \geq 2$ and $z_1^i z_2^j \in \mathfrak{m}^{n(i+j)} \subseteq \mathfrak{m}^{2n}$. So $\hat{E}(z_1, z_2) \equiv z_1 + z_2 \pmod{\mathfrak{m}^{2n}} \in \mathfrak{m}^n$. Similarly, $\iota(z) = -z + (\text{higher}) \in \mathfrak{m}^n$ when $z \in \mathfrak{m}^n$.

(2) The map $\hat{E}(\mathfrak{m}^n) \rightarrow k$ sending $z \mapsto z\pi^{-n} \bmod \mathfrak{m}$ is a group homomorphism: $\hat{E}(z_1, z_2) = z_1 + z_2 + (\text{terms in } \mathfrak{m}^{2n})$, and $\mathfrak{m}^{2n} \subseteq \mathfrak{m}^{n+1}$ (since $2n \geq n+1$ for $n \geq 1$), so modulo $\mathfrak{m}^{n+1}$:

$$\hat{E}(z_1, z_2) \equiv z_1 + z_2 \pmod{\mathfrak{m}^{n+1}}.$$

Dividing by π^n : $(z_1 + z_2)/\pi^n \equiv z_1/\pi^n + z_2/\pi^n \pmod{\mathfrak{m}}$. The kernel is $\mathfrak{m}^{n+1}$, so $\hat{E}(\mathfrak{m}^n)/\hat{E}(\mathfrak{m}^{n+1}) \cong k^+$.

(3) Since $\bigcap_{n \geq 1} \mathfrak{m}^n = (0)$ (by completeness), $\hat{E}(\mathfrak{m}) = \varprojlim \hat{E}(\mathfrak{m})/\hat{E}(\mathfrak{m}^n)$. Each quotient is a k -vector space of dimension 1, hence a group of order $|k| = q$ (a power of p). So $\hat{E}(\mathfrak{m})$ is an inverse limit of finite p -groups, i.e., a pro- p group.

The filtration is the formal group analogue of the p -adic filtration on $\mathbb{Z}_p$: each quotient is a copy of the residue field, and the group law becomes “more additive” at deeper levels. At level n , the group operation on $\hat{E}(\mathfrak{m}^n)$ is addition modulo $\mathfrak{m}^{n+1}$ — the higher-order terms of the formal group law contribute only to deeper levels.

The Logarithm Isomorphism

For sufficiently deep levels of the filtration, the formal logarithm from section 4.3 converges and gives an explicit isomorphism to the additive group.

Proposition 4.5.4: The Logarithm on $\hat{E}(\mathfrak{m}^n)$

Let K be a local field of characteristic zero (i.e., a finite extension of $\mathbb{Q}_p$) with absolute ramification index e (so $v(p) = e$). Then for $n > e/(p-1)$, the formal logarithm

$$\log_{\hat{E}} : \hat{E}(\mathfrak{m}^n) \xrightarrow{\sim} (\mathfrak{m}^n, +)$$

is an isomorphism of groups (where the target has ordinary addition).

Proof:

By proposition 4.3.11, $\log_{\hat{E}}$ converges on $\mathfrak{m}$ and takes values in K . For $z \in \mathfrak{m}^n$:

$$\log_{\hat{E}}(z) = z + \frac{d_1}{2}z^2 + \frac{d_2}{3}z^3 + \cdots = \sum_{m=1}^{\infty} \frac{d_{m-1}}{m}z^m.$$

The term of degree m has valuation $\geq mn - v(m)$. Since $v(m) \leq e \cdot \log_p(m)$ (the p -adic valuation of an integer m is at most $\log_p(m)$, multiplied by e for the ramification), the terms converge to zero for $z \in \mathfrak{m}$, and $\log_{\hat{E}}(z) \in \mathfrak{m}^n$ when n is large enough that the correction terms lie in $\mathfrak{m}^n$.

More precisely, $\log_{\hat{E}}(z) \equiv z \pmod{\mathfrak{m}^{n+1}}$ when $2n - v(2) \geq n + 1$, i.e., $n \geq v(2) + 1 = e + 1$ (for $p = 2$). The general condition is $n > e/(p-1)$: the term $d_{m-1}z^m/m$ has valuation $\geq mn - v(m)$, and $mn - v(m) \geq n + 1$ requires $n(m-1) \geq 1 + v(m)$. The worst case is $m = p$, giving $n(p-1) \geq 1 + e$, i.e., $n > e/(p-1)$.

Under this condition, $\log_{\hat{E}}$ maps $\mathfrak{m}^n$ into $\mathfrak{m}^n$, and the inverse (the formal exponential) also converges on $\mathfrak{m}^n$. Since $\log_{\hat{E}}$ is a homomorphism (theorem 4.3.3) and both $\log$ and $\exp$ converge, the map is a group isomorphism $\hat{E}(\mathfrak{m}^n) \xrightarrow{\sim} (\mathfrak{m}^n, +)$.

In particular, for $K = \mathbb{Q}_p$ (where $e = 1$): $\hat{E}(p\mathbb{Z}_p) \cong (p\mathbb{Z}_p, +) \cong \mathbb{Z}_p$ as an abelian group when $p \geq 3$ (taking $n = 1$). For $p = 2$, we need $n > 1$, so $\hat{E}(4\mathbb{Z}_2) \cong (4\mathbb{Z}_2, +) \cong \mathbb{Z}_2$.

Torsion in the Formal Group

The filtration and the logarithm isomorphism combine to give precise control over the torsion in $\hat{E}(\mathfrak{m})$.

Theorem 4.5.5: Torsion in the Formal Group

Let E/K have good reduction, with residue characteristic p .

1. If $\gcd(n, p) = 1$, then $\hat{E}(\mathfrak{m})$ has no n -torsion: $\hat{E}(\mathfrak{m})[n] = \{0\}$.
2. The torsion subgroup $\hat{E}(\mathfrak{m})_{\text{tors}}$ is a p -group.
3. If K has characteristic zero and $n > e/(p-1)$, then $\hat{E}(\mathfrak{m}^n)$ is torsion-free.

Proof:

(1) The $[n]$ -series on the formal group is $[n]_{\hat{E}}(z) = nz + O(z^2)$. If $\gcd(n, p) = 1$, then n is a unit in R (since $v(n) = 0$), so $[n]_{\hat{E}}$ has $[n]_{\hat{E}}'(0) = n \in R^*$. By theorem 4.1.6 (or by the formal implicit function theorem), $[n]_{\hat{E}}$ is an isomorphism on $\mathfrak{m}$: for any $w \in \mathfrak{m}$, the equation $[n]_{\hat{E}}(z) = w$ has a unique solution $z \in \mathfrak{m}$. In particular, $[n]_{\hat{E}}(z) = 0$ has only the solution $z = 0$.

Alternatively: $[n]_{\hat{E}}(z) = nz(1 + c_2z + \cdots) = nz \cdot u(z)$ where $u(z) \in R[[z]]$ with $u(0) = 1$. For $z \in \mathfrak{m}$, the series $u(z)$ evaluates to an element of $1 + \mathfrak{m}$, which is a unit in the local ring R . If $[n]_{\hat{E}}(z) = 0$, then $nz \cdot u(z) = 0$. Since $n \in R^*$ and $u(z) \in R^*$, we can cancel them in the integral domain R to conclude $z = 0$.

(2) Immediate from (1): every torsion element has order a power of p .

(3) For $n > e/(p-1)$, the logarithm gives $\hat{E}(\mathfrak{m}^n) \cong (\mathfrak{m}^n, +)$ (proposition 4.5.4). The additive group $\mathfrak{m}^n$ is torsion-free (since K has characteristic zero: if $m \cdot z = 0$ for $z \in \mathfrak{m}^n$ and $m \neq 0$, then $z = 0$). So $\hat{E}(\mathfrak{m}^n)$ is torsion-free.

The Torsion Injection Theorem

Combining the exact sequence (4.22) with the torsion result gives one of the most useful elementary results in the arithmetic of elliptic curves.

Theorem 4.5.6: Torsion Injects under Good Reduction

Let E/K have good reduction, with residue field k of characteristic p . Then the reduction map induces an injection on prime-to- p torsion:

$$E(K)[n] \hookrightarrow \tilde{E}(k) \quad \text{for all } n \text{ with } \gcd(n, p) = 1. \quad (4.25)$$

More precisely, $E(K)[n] \cap E_1(K) = \{O\}$ for $\gcd(n, p) = 1$.

Proof:

If $P \in E(K)[n] \cap E_1(K)$, then $z(P) \in \hat{E}(\mathfrak{m})$ and $[n](z(P)) = z([n]P) = z(O) = 0$. By theorem 4.5.5(1), $z(P) = 0$, so $P = O$.

Corollary 4.5.7: Bounding Torsion over Number Fields

Let $E/\mathbb{Q}$ be an elliptic curve with good reduction at a prime p , and let $\tilde{E}/\mathbb{F}_p$ be the reduction. Then for any prime $\ell \neq p$:

$$|E(\mathbb{Q})[\ell^\infty]| \mid |\tilde{E}(\mathbb{F}_p)|.$$

In particular, $E(\mathbb{Q})_{\text{tors}}$ injects into $\tilde{E}(\mathbb{F}_p)$ after removing the p -part.

Proof:

The inclusion $\mathbb{Q} \hookrightarrow \mathbb{Q}_p$ gives $E(\mathbb{Q}) \subseteq E(\mathbb{Q}_p)$, and E has good reduction over $\mathbb{Q}_p$. By theorem 4.5.6, $E(\mathbb{Q})[\ell^n] \hookrightarrow \tilde{E}(\mathbb{F}_p)$ for all n . Since $\tilde{E}(\mathbb{F}_p)$ is finite, so is $E(\mathbb{Q})[\ell^\infty]$, and its order divides $|\tilde{E}(\mathbb{F}_p)|$.

This is a remarkably effective computational tool. To determine the torsion subgroup of $E(\mathbb{Q})$, one reduces modulo several primes of good reduction and takes the intersection of the constraints. For instance:

Example 4.5.8: Bounding Torsion on Running Examples

(a) $E : y^2 = x^3 - x$ over $\mathbb{Q}$. The discriminant is $\Delta = -64$, so E has good reduction at all primes $p \neq 2$. Over $\mathbb{F}_3$: $|E(\mathbb{F}_3)| = 4$ (example 4.4.6). Over $\mathbb{F}_5$: $|\tilde{E}(\mathbb{F}_5)| = 8$ (from $a_5 = -2$, so $|\tilde{E}| = 5 + 1 + 2 = 8$). The torsion must inject into both $\tilde{E}(\mathbb{F}_3)$ and $\tilde{E}(\mathbb{F}_5)$ away from the respective residue characteristics.

From $p = 3$: $E(\mathbb{Q})[\ell^\infty] \hookrightarrow \tilde{E}(\mathbb{F}_3)$ for $\ell \neq 3$, so $|E(\mathbb{Q})[2^\infty]|$ divides 4. From $p = 5$: $|E(\mathbb{Q})[3^\infty]|$ divides $|\tilde{E}(\mathbb{F}_5)| = 8$, but $\gcd(3^n, 8) = 1$, so $E(\mathbb{Q})$ has no 3-torsion. Combining: $|E(\mathbb{Q})_{\text{tors}}|$ divides $\gcd(4, 8) = 4$ (away from 2 and 3, the torsion is trivial). In fact, $E(\mathbb{Q})_{\text{tors}} = \{O, (0, 0), (1, 0), (-1, 0)\} \cong (\mathbb{Z}/2\mathbb{Z})^2$.

(b) $E : y^2 = x^3 + 1$ over $\mathbb{Q}$. The discriminant is $\Delta = -432$, so good reduction away from 2 and 3. Over $\mathbb{F}_5$: $|\tilde{E}(\mathbb{F}_5)| = 6$. Over $\mathbb{F}_7$: one computes $|\tilde{E}(\mathbb{F}_7)| = 12$. From $p = 5$: $|E(\mathbb{Q})_{\text{tors}}|$ divides 6 (away from 5). From $p = 7$: $|E(\mathbb{Q})_{\text{tors}}|$ divides 12 (away from 7). So $|E(\mathbb{Q})_{\text{tors}}|$ divides $\gcd(6, 12) = 6$. In fact, $E(\mathbb{Q})_{\text{tors}} = \{O, (0, \pm 1), (-1, 0), (2, \pm 3)\} \cong \mathbb{Z}/6\mathbb{Z}$.

The Structure of $\hat{E}(\mathfrak{m})$ as a $\mathbb{Z}_p$ -Module

When K has characteristic zero, the formal group $\hat{E}(\mathfrak{m})$ is not just a pro- p group — it has the structure of a $\mathbb{Z}_p$ -module, and we can determine its rank.

Proposition 4.5.9: $\hat{E}(\mathfrak{m})$ as a $\mathbb{Z}_p$ -Module

Let $K/\mathbb{Q}_p$ be a finite extension of degree $[K : \mathbb{Q}_p] = d$. Then:

1. $\hat{E}(\mathfrak{m}) \cong \mathbb{Z}_p^d \oplus (\text{finite } p\text{-group})$ as a $\mathbb{Z}_p$ -module.
2. The finite part is trivial (i.e., $\hat{E}(\mathfrak{m}) \cong \mathbb{Z}_p^d$) when p is odd and $e < p - 1$ (where e is the absolute ramification index of $K/\mathbb{Q}_p$).

Proof:

(1) By proposition 4.5.4, for $n > e/(p-1)$ the logarithm gives $\hat{E}(\mathfrak{m}^n) \cong (\mathfrak{m}^n, +) \cong R \cong \mathbb{Z}_p^d$ as a $\mathbb{Z}_p$ -module. The quotient $\hat{E}(\mathfrak{m})/\hat{E}(\mathfrak{m}^n)$ is finite (it is filtered by the quotients $\hat{E}(\mathfrak{m}^i)/\hat{E}(\mathfrak{m}^{i+1}) \cong k$, and there are finitely many layers). So $\hat{E}(\mathfrak{m})$ is an extension of the finite group $\hat{E}(\mathfrak{m})/\hat{E}(\mathfrak{m}^n)$ by the free module $\hat{E}(\mathfrak{m}^n) \cong \mathbb{Z}_p^d$. By the structure theory of finitely generated $\mathbb{Z}_p$ -modules, $\hat{E}(\mathfrak{m}) \cong \mathbb{Z}_p^d \oplus T$ with T finite.

- (2) The finite part T is the p -torsion of $\hat{E}(\mathfrak{m})$. When $e < p - 1$, we have $e/(p - 1) < 1$, so the condition $n > e/(p - 1)$ is satisfied for $n = 1$. Thus $\hat{E}(\mathfrak{m}) = \hat{E}(\mathfrak{m}^1)$ is torsion-free, and $T = 0$.

For $K = \mathbb{Q}_p$ with $p \geq 3$: $e = 1$, $d = 1$, and $e < p - 1$ holds for $p \geq 3$, so $\hat{E}(p\mathbb{Z}_p) \cong \mathbb{Z}_p$. This says the kernel of reduction over $\mathbb{Q}_p$ is a free $\mathbb{Z}_p$ -module of rank 1 — topologically, it looks like the p -adic integers under addition. This rank-1 structure will contribute to the proof of the Mordell–Weil theorem in chapter 8, where it enters through the theory of heights and descent.

Looking Ahead

The results of this section will be generalized in two directions in chapter 6. First, for curves with *bad* reduction (multiplicative or additive), the kernel of reduction still exists and is still described by the formal group, but the exact sequence (4.22) involves the nonsingular part $\tilde{E}_{\text{ns}}(k)$ of the special fiber (which is $\mathbb{G}_m(k) = k^*$ for multiplicative reduction and $\mathbb{G}_a(k) = k$ for additive reduction, by section 1.6). Second, over a number field K , the reduction at each prime gives independent constraints, and combining them with the formal group structure at each prime is the engine of the weak Mordell–Weil theorem.

For the arithmetic of the p -torsion (which the torsion injection theorem does *not* control), the formal group is still the right tool: the height determines whether $E[p](K)$ can be nontrivial. Specifically, if E has good ordinary reduction at p , then $E[p](K) \hookrightarrow E[p](\bar{K})^{I_p}$ where I_p is the inertia group, and the formal group of height 1 contributes a connected piece $\hat{E}[p] = \ker([p]_{\hat{E}} : \hat{E}(\mathfrak{m}) \rightarrow \hat{E}(\mathfrak{m}))$ of order p .

Exercise 4.5.1

1. Verify that the reduction map $E(K) \rightarrow \tilde{E}(k)$ is surjective for E with good reduction, by using Hensel's lemma (theorem 4.1.6). (*Hint*: a smooth point $\tilde{Q} = (\tilde{x}_0, \tilde{y}_0) \in \tilde{E}(k)$ gives an approximate solution (x_0, y_0) to the Weierstrass equation over R , with $F(x_0, y_0) \in \mathfrak{m}$ and $\partial F / \partial y(x_0, y_0) = 2\tilde{y}_0 + a_1\tilde{x}_0 + a_3 \in k^*$ by smoothness.)
2. Let $E : y^2 = x^3 + 3x + 7$ over $\mathbb{Q}$, which has $\Delta = -2^4 \cdot 3 \cdot 1399$. Use reduction modulo 5 and modulo 11 to bound $E(\mathbb{Q})_{\text{tors}}$.
3. For $E : y^2 + y = x^3$ over $\mathbb{Q}_2$, compute $E_1(\mathbb{Q}_2) = \hat{E}(2\mathbb{Z}_2)$, and describe the first few levels of the filtration $\hat{E}(2\mathbb{Z}_2) \supset \hat{E}(4\mathbb{Z}_2) \supset \hat{E}(8\mathbb{Z}_2) \supset \dots$.
4. Show that $\hat{E}(p\mathbb{Z}_p) \cong \mathbb{Z}_p$ for any elliptic curve $E/\mathbb{Q}_p$ with good reduction and $p \geq 5$. (*Hint*: use $e = 1$, $e/(p - 1) < 1$, and the logarithm isomorphism.)
5. Let $E/\mathbb{Q}_p$ have good ordinary reduction (so $\text{ht}(\hat{E}) = 1$). Show that $|\hat{E}(\mathfrak{m})[p]| = p$, and identify this with the connected part of $E[p]$. (*Hint*: $[p]_{\hat{E}}(z) = uz^p + \dots$ with $u \neq 0$; the equation $[p](z) = 0$ has p solutions in $\bar{\mathfrak{m}}$ by counting the degree.)
6. Let $E/\mathbb{Q}_p$ have good supersingular reduction (so $\text{ht}(\hat{E}) = 2$). Show that $|\hat{E}(\mathfrak{m})[p]| = p^2$, and explain why this means $E[p](\bar{K}) \subseteq \hat{E}(\bar{\mathfrak{m}})$ (all p -torsion is “in the formal group”).

Part II

Arithmetic: Local and Global Theory

here

5

Elliptic Curves over Finite Fields

We now bring the above machinery of elliptic curves to finite fields $\mathbb{F}_q$, where $q = p^r$ is a prime power. Over finite fields, the Frobenius endomorphism φ_q — introduced in section 2.4 as the map $(x, y) \mapsto (x^q, y^q)$ — is the dominant actor. Its characteristic polynomial $T^2 - a_q T + q$ (stated in theorem 2.4.6 but with proof deferred to this chapter) determines everything: the point count $|E(\mathbb{F}_q)| = q + 1 - a_q$, the Hasse bound $|a_q| \leq 2\sqrt{q}$, the zeta function of $E/\mathbb{F}_q$ and the Weil conjectures, the group structure $E(\mathbb{F}_q) \cong \mathbb{Z}/d_1\mathbb{Z} \times \mathbb{Z}/d_2\mathbb{Z}$, and the isogeny class of E . This chapter proves the characteristic polynomial formula, derives its consequences, and settles the remaining debts from chapter 2: Tate’s isogeny theorem (the Tate module functor is fully faithful over finite fields) and the structure of the endomorphism ring for supersingular curves. By the end, the reader will see that the arithmetic of an elliptic curve over a finite field is, in a precise sense, completely determined by a single integer a_q .

5.1 The Characteristic Polynomial of Frobenius

The characteristic polynomial of the Frobenius endomorphism was stated in theorem 2.4.6:

$$\varphi_q^2 - a_q \varphi_q + [q] = [0] \quad \text{in } \text{End}(E),$$

where $a_q = q + 1 - |E(\mathbb{F}_q)|$. The proof was deferred because it requires the dual isogeny and the degree quadratic form, which had not yet been developed. Both tools are now available (from section 2.2), and the proof is surprisingly clean: it is a three-line consequence of the identity $\hat{\varphi}_q \circ \psi + \hat{\psi} \circ \varphi_q = [\deg(\varphi_q + \psi) - \deg(\varphi_q) - \deg(\psi)]$ applied to the correct choice of ψ .

Point Counting via the Degree

We begin by making precise the relationship between $|E(\mathbb{F}_q)|$ and the degree of $\varphi_q - [1]$.

Proposition 5.1.1: Point Count as a Degree

Let $E/\mathbb{F}_q$ be an elliptic curve. The isogeny $\varphi_q - [1] : E \rightarrow E$ is separable, and

$$|E(\mathbb{F}_q)| = \deg(\varphi_q - [1]). \quad (5.1)$$

Proof:

A point $P \in E(\overline{\mathbb{F}_q})$ lies in $E(\mathbb{F}_q)$ if and only if $\varphi_q(P) = P$, i.e., $(\varphi_q - [1])(P) = O$. Thus $E(\mathbb{F}_q) = \ker(\varphi_q - [1])$, and $|E(\mathbb{F}_q)| = |\ker(\varphi_q - [1])| = \deg_s(\varphi_q - [1])$ (the separable degree equals the kernel size).

It remains to show $\varphi_q - [1]$ is separable, so that $\deg_s = \deg$. The pullback on differentials is $(\varphi_q - [1])^*\omega = \varphi_q^*\omega - [1]^*\omega$. Since φ_q is purely inseparable, $\varphi_q^*\omega = 0$ (proposition 2.4.2). And $[1]^*\omega = \omega$. So $(\varphi_q - [1])^*\omega = -\omega \neq 0$, which means $\varphi_q - [1]$ is separable. Therefore $|E(\mathbb{F}_q)| = \deg(\varphi_q - [1])$.

More generally, $|E(\mathbb{F}_{q^n})| = \deg(\varphi_q^n - [1])$ for all $n \geq 1$, by the same argument with $\varphi_{q^n} = \varphi_q^n$.

The Trace Identity

The key step is to show that $\varphi_q + V_q$ (the sum of Frobenius and Verschiebung in $\text{End}(E)$) is a multiplication map.

Theorem 5.1.2: The Trace of Frobenius

Let $E/\mathbb{F}_q$ be an elliptic curve with Frobenius φ_q and Verschiebung $V_q = \hat{\varphi}_q$. Then

$$\varphi_q + V_q = [a_q] \quad \text{in } \text{End}(E), \quad (5.2)$$

where $a_q = q + 1 - |E(\mathbb{F}_q)|$.

Proof:

We use the bilinear form associated to the degree quadratic form (proposition 2.2.9). For any $\varphi, \psi \in \text{End}(E)$:

$$\hat{\varphi} \circ \psi + \hat{\psi} \circ \varphi = [\deg(\varphi + \psi) - \deg(\varphi) - \deg(\psi)]. \quad (5.3)$$

Apply this with $\varphi = \varphi_q$ and $\psi = -[1]$:

$$\begin{aligned} \hat{\varphi}_q \circ (-[1]) + \widehat{(-[1])} \circ \varphi_q &= [\deg(\varphi_q - [1]) - \deg(\varphi_q) - \deg([1])] \\ -V_q + (-[1]) \circ \varphi_q &= [|E(\mathbb{F}_q)| - q - 1] \\ -V_q - \varphi_q &= [-a_q], \end{aligned}$$

where we used $\hat{\varphi}_q = V_q$ (theorem 2.4.3), $\widehat{(-[1])} = \widehat{[-1]} = [-1]$ (proposition 2.2.4(6)), $[-1] \circ \varphi_q = -\varphi_q$ (since $[-1]$ is the negation map), $\deg(\varphi_q) = q$, $\deg([1]) = 1$, and $\deg(\varphi_q - [1]) = |E(\mathbb{F}_q)|$ (proposition 5.1.1). Multiplying both sides by -1 : $\varphi_q + V_q = [a_q]$.

The identity $\varphi_q + V_q = [a_q]$ says that the “sum” of Frobenius and its dual acts as multiplication by the trace. This is the elliptic curve analogue of the linear algebra identity $\alpha + \bar{\alpha} = \text{tr}(\alpha)$ for an element α of a quadratic extension: the Frobenius φ_q plays the role of α , the Verschiebung $V_q = \hat{\varphi}_q$ plays the role of the conjugate $\bar{\alpha}$, and the trace a_q is the sum.

The Characteristic Polynomial

Theorem 5.1.3: Characteristic Polynomial of Frobenius

Let $E/\mathbb{F}_q$ be an elliptic curve. Then:

$$\varphi_q^2 - [a_q] \circ \varphi_q + [q] = [0] \quad \text{in } \text{End}(E), \quad (5.4)$$

where $a_q = q + 1 - |E(\mathbb{F}_q)|$. Equivalently, φ_q is a root of $P(T) = T^2 - a_q T + q \in \mathbb{Z}[T]$.

Proof:

From theorem 5.1.2: $V_q = [a_q] - \varphi_q$. Composing both sides on the right with φ_q :

$$V_q \circ \varphi_q = [a_q] \circ \varphi_q - \varphi_q^2. \quad (5.5)$$

The left side is $[q]$ (from $V_q \circ \varphi_q = \hat{\varphi}_q \circ \varphi_q = [\deg(\varphi_q)] = [q]$, by proposition 2.2.4(1)). Therefore:

$$[q] = a_q \varphi_q - \varphi_q^2,$$

which rearranges to $\varphi_q^2 - a_q \varphi_q + [q] = [0]$.

The proof is strikingly short: two lines, using only the trace identity and the fundamental property $\hat{\varphi} \circ \varphi = [\deg \varphi]$ of the dual isogeny. This illustrates the power of the dual isogeny formalism developed in section 2.2: what would be a substantial computation by other means becomes an exercise in the algebra of endomorphisms.

The Degree Formula

The characteristic polynomial yields a closed formula for $\deg(\varphi_q - [n])$ for all $n \in \mathbb{Z}$.

Proposition 5.1.4: Degree of $\varphi_q - [n]$

For all $n \in \mathbb{Z}$:

$$\deg(\varphi_q - [n]) = n^2 - a_q n + q. \quad (5.6)$$

In particular, $\deg(\varphi_q - [1]) = 1 - a_q + q = |E(\mathbb{F}_q)|$ and $\deg(\varphi_q - [0]) = q = \deg(\varphi_q)$.

Proof:

The dual of $\varphi_q - [n]$ is $\widehat{\varphi_q - [n]} = V_q - [n]$ (using linearity of the dual). Composing:

$$\begin{aligned} [\deg(\varphi_q - [n])] &= (V_q - [n]) \circ (\varphi_q - [n]) \\ &= V_q \circ \varphi_q - [n] \circ V_q - V_q \circ [n] + [n]^2 \dots \end{aligned}$$

Wait — we need to be careful with the order of terms since V_q and φ_q do not commute in general. However, $[n]$ commutes with every endomorphism (since $[n]$ is defined by the group law), so:

$$\begin{aligned} (V_q - [n])(\varphi_q - [n]) &= V_q \varphi_q - nV_q - n\varphi_q + [n]^2 \\ &= [q] - n(V_q + \varphi_q) + [n]^2 \\ &= [q] - n[a_q] + [n]^2 = [q - na_q + n^2]. \end{aligned}$$

Here we used $V_q \varphi_q = [q]$, $V_q + \varphi_q = [a_q]$ (theorem 5.1.2), and $[n] \circ V_q = n \cdot V_q$ (i.e., $[n](V_q(P)) = V_q(P) + \dots + V_q(P)$ (n times) $= V_q([n]P) = V_q \circ [n](P)$, using that V_q is a group homomorphism; but $n \cdot V_q = [n] \circ V_q$, which requires care). More directly: $(V_q - [n])(\varphi_q - [n]) = V_q \varphi_q - V_q[n] - [n]\varphi_q + [n][n]$. Since $[n]$ is central: $V_q[n] = [n]V_q$ and $[n]\varphi_q = \varphi_q[n]$... actually $[n]$ commutes with φ_q because φ_q is a group homomorphism: $\varphi_q([n]P) = [n]\varphi_q(P)$. Similarly V_q commutes with $[n]$. So: $V_q[n] + [n]\varphi_q = [n](V_q + \varphi_q) = [n][a_q] = [na_q]$. Therefore:

$$[\deg(\varphi_q - [n])] = [q] - [na_q] + [n]^2 = [n^2 - a_q n + q].$$

Hasse's Theorem

The characteristic polynomial has real roots if and only if $a_q^2 - 4q \geq 0$, i.e., $|a_q| \geq 2\sqrt{q}$. But the degree formula imposes the opposite constraint:

Theorem 5.1.5: Hasse's Theorem

Let $E/\mathbb{F}_q$ be an elliptic curve. Then

$$|a_q| \leq 2\sqrt{q}, \quad \text{i.e.,} \quad |q + 1 - |E(\mathbb{F}_q)|| \leq 2\sqrt{q}. \quad (5.7)$$

Equivalently, $(\sqrt{q} - 1)^2 \leq |E(\mathbb{F}_q)| \leq (\sqrt{q} + 1)^2$.

Proof:

The quadratic $f(n) = n^2 - a_q n + q = \deg(\varphi_q - [n])$ satisfies $f(n) \geq 0$ for all $n \in \mathbb{Z}$ (since the degree of a nonzero isogeny is positive, and $f(n) = 0$ only if $\varphi_q = [n]$). In fact, $f(n) \geq 0$ for all $n \in \mathbb{R}$: by the parallelogram identity (proposition 2.2.9), $\deg$ is a quadratic form on $\text{End}(E)$, and $\deg(\varphi_q - [t])$ for $t \in \mathbb{R}$ can be interpreted via $\text{End}(E) \otimes \mathbb{R}$ as a positive semi-definite quadratic form. More elementarily: if $a_q^2 - 4q > 0$, the quadratic $f(n) = n^2 - a_q n + q$ would have two distinct real roots, say $r_1 < r_2$, with $f(n) < 0$ for $r_1 < n < r_2$. Since $r_2 - r_1 = \sqrt{a_q^2 - 4q} > 0$ and the interval (r_1, r_2) has positive length, it contains at least one integer n_0 . But then $\deg(\varphi_q - [n_0]) = f(n_0) < 0$, contradicting the non-negativity of degrees.

Therefore $a_q^2 - 4q \leq 0$, i.e., $|a_q| \leq 2\sqrt{q}$.

Hasse's theorem was conjectured by Artin (in his thesis, 1924) and proved by Hasse in 1933 for elliptic curves over prime fields $\mathbb{F}_p$. The general case (arbitrary q) was settled by Weil in the 1940s as part of his proof of the Riemann hypothesis for curves over finite fields. Our proof, via the characteristic polynomial and the degree form, is essentially Weil's argument cast in the language of endomorphism rings.

The bound is sharp: for each q and each integer a with $|a| \leq 2\sqrt{q}$ satisfying certain divisibility conditions (see section 5.3), there exists an elliptic curve $E/\mathbb{F}_q$ with $a_q = a$.

The Roots of the Characteristic Polynomial

The roots of $P(T) = T^2 - a_q T + q$ are

$$\alpha, \beta = \frac{a_q \pm \sqrt{a_q^2 - 4q}}{2}, \quad (5.8)$$

satisfying $\alpha + \beta = a_q$ and $\alpha\beta = q$. By Hasse's theorem, $a_q^2 - 4q \leq 0$, so α and β are either:

1. complex conjugates in $\mathbb{C}$ (when $a_q^2 < 4q$), with $|\alpha| = |\beta| = \sqrt{q}$; or
2. a repeated real root $\alpha = \beta = a_q/2 = \pm\sqrt{q}$ (when $a_q^2 = 4q$, i.e., $a_q = \pm 2\sqrt{q}$, which requires q to be a perfect square).

In either case, $|\alpha| = |\beta| = \sqrt{q}$. This is the *Riemann hypothesis* for $E/\mathbb{F}_q$ — the analogue, for the zeta function of E , of the classical Riemann hypothesis for the Riemann zeta function. We will make this precise in section 5.2.

The formula $|E(\mathbb{F}_{q^n})| = \deg(\varphi_q^n - [1])$ and the factorization $\varphi_q^n - [1] = (\varphi_q - [1])(\varphi_q^{n-1} + \cdots + [1])\dots$ actually, it is cleaner to use the characteristic polynomial directly. Since φ_q satisfies $\varphi_q^2 = a_q \varphi_q - [q]$, the endomorphism φ_q^n can be expressed as $[c_n]\varphi_q + [d_n]$ for integers c_n, d_n determined by the recurrence $c_{n+1} = a_q c_n - q c_{n-1}$. But the cleanest formula uses the roots α, β :

Proposition 5.1.6: Point Counts over Extension Fields

For all $n \geq 1$:

$$|E(\mathbb{F}_{q^n})| = q^n + 1 - \alpha^n - \beta^n, \quad (5.9)$$

where α, β are the roots of $T^2 - a_q T + q$.

Proof:

The Frobenius of $\mathbb{F}_{q^n}/\mathbb{F}_{q^n}$ is $\varphi_{q^n} = \varphi_q^n$. So $|E(\mathbb{F}_{q^n})| = \deg(\varphi_q^n - [1]) = (\text{evaluate the degree form at } \varphi_q^n)$.

We claim $\deg(\varphi_q^n - [1]) = q^n + 1 - \alpha^n - \beta^n$. The endomorphism φ_q acts on the Tate module $T_\ell(E)$ (for $\ell \neq p$) as a matrix $A \in M_2(\mathbb{Z}_\ell)$ with characteristic polynomial $T^2 - a_q T + q$ and eigenvalues α, β (over $\overline{\mathbb{Q}_\ell}$). Therefore φ_q^n has eigenvalues α^n, β^n , and $\varphi_q^n - [1]$ has eigenvalues

$\alpha^n - 1, \beta^n - 1$. By theorem 2.5.7:

$$\deg(\varphi_q^n - [1]) = \det(T_\ell(\varphi_q^n - [1])) = (\alpha^n - 1)(\beta^n - 1) = \alpha^n \beta^n - \alpha^n - \beta^n + 1.$$

Since $\alpha\beta = q$, we have $\alpha^n \beta^n = q^n$, giving $|E(\mathbb{F}_{q^n})| = q^n + 1 - \alpha^n - \beta^n$.

Running Examples

Example 5.1.7: Characteristic Polynomials of Running Curves

(a) $E : y^2 = x^3 - x$ over $\mathbb{F}_5$. We computed $|E(\mathbb{F}_5)| = 8$ (example 2.4.5(a), giving $a_5 = 5 + 1 - 8 = -2$). The characteristic polynomial is $T^2 + 2T + 5$, with roots $\alpha, \beta = -1 \pm 2i$. Check: $|\alpha| = \sqrt{1 + 4} = \sqrt{5}$. The point counts over extensions: $|E(\mathbb{F}_{25})| = 25 + 1 - ((-1 + 2i)^2 + (-1 - 2i)^2) = 26 - (1 - 4i - 4 + 1 + 4i - 4) = 26 - (-6) = 32$.

(b) $E : y^2 = x^3 + 1$ over $\mathbb{F}_5$. We computed $|E(\mathbb{F}_5)| = 6$ (example 2.4.5(b)), giving $a_5 = 0$. The characteristic polynomial is $T^2 + 5$, with roots $\alpha, \beta = \pm i\sqrt{5}$. Since $a_5 = 0$ and $5 \nmid a_5$, the curve is supersingular (theorem 2.4.4). Over $\mathbb{F}_{25}$: $|E(\mathbb{F}_{25})| = 25 + 1 - ((i\sqrt{5})^2 + (-i\sqrt{5})^2) = 26 - (-5 - 5) = 36$.

(c) $E : y^2 = x^3 - x$ over $\mathbb{F}_7$. Evaluating $x^3 - x$ for $x = 0, \dots, 6 \pmod{7}$: $x = 0 : 0$ (1 pt); $x = 1 : 0$ (1 pt); $x = 2 : 6$ (QR: $6 \equiv -1$, non-residue, 0 pts); $x = 3 : 24 \equiv 3$ (non-residue, 0 pts); $x = 4 : 60 \equiv 4 \equiv 2^2$ ($y = \pm 2$, 2 pts); $x = 5 : 120 \equiv 1$ ($y = \pm 1$, 2 pts); $x = 6 : 210 \equiv 0$ (1 pt). Total affine points: $1 + 1 + 0 + 0 + 2 + 2 + 1 = 7$. With O : $|E(\mathbb{F}_7)| = 8$, $a_7 = 0$. Supersingular! The characteristic polynomial is $T^2 + 7$.

(d) Congruent number curve $E_5 : y^2 = x^3 - 25x$ over $\mathbb{F}_7$. We have $a_4 = -25 \equiv 3$ and $a_6 = 0$ over $\mathbb{F}_7$. Evaluating $x^3 + 3x$ for $x \in \mathbb{F}_7$: $x = 0 : 0$ (1 pt); $x = 1 : 4 \equiv 2^2$ (2 pts); $x = 2 : 14 \equiv 0$ (1 pt); $x = 3 : 36 \equiv 1$ (2 pts); $x = 4 : 76 \equiv 6$ (non-res, 0); $x = 5 : 140 \equiv 0$ (1 pt); $x = 6 : 234 \equiv 3$ (non-res, 0). Total: $1 + 2 + 1 + 2 + 0 + 1 + 0 = 7$. With O : $|E_5(\mathbb{F}_7)| = 8$, $a_7 = 0$, supersingular. (This is not a coincidence: E_5 is a quadratic twist of $E_1 : y^2 = x^3 - x$ by $d = 25 \equiv 4 \equiv 2^2 \pmod{7}$, which is a square, so $E_5 \cong E_1$ over $\mathbb{F}_7$.)

Exercise 5.1.1

1. Verify the characteristic polynomial relation $\varphi_5^2 - a_5 \varphi_5 + [5] = [0]$ on $E : y^2 = x^3 - x$ over $\mathbb{F}_5$ (where $a_5 = -2$) by checking it on the point $P = (2, 1) \in E(\mathbb{F}_{25})$. (Hint: Compute $\varphi_5(P) = (2^5, 1^5) = (32, 1) = (7, 1)$ in $\mathbb{F}_{25}$, then $\varphi_5^2(P)$, then verify $\varphi_5^2(P) + 2\varphi_5(P) + [5]P = O$.)
2. Show directly that $\deg(\varphi_q - [n]) = n^2 - a_q n + q$ by expanding $(V_q - [n])(\varphi_q - [n])$ without using the trace identity, but instead using only $V_q \varphi_q = [q]$, $\varphi_q V_q = [q]$, $\deg(\varphi_q) = q$, $\deg([n]) = n^2$, and the bilinear identity (5.3). (You will derive the trace identity as a byproduct.)
3. For $q = p$ prime, which values of a_p with $|a_p| \leq 2\sqrt{p}$ actually occur as traces of Frobenius? Show that $a_p = 0$ is always realized (by a supersingular curve), and that $a_p = 1$ is realized by $y^2 = x^3 + x$ over $\mathbb{F}_3$.

4. Using $|E(\mathbb{F}_{q^n})| = q^n + 1 - \alpha^n - \beta^n$, compute $|E(\mathbb{F}_{5^3})|$ for $E : y^2 = x^3 - x$ over $\mathbb{F}_5$ (where $\alpha = -1 + 2i$, $\beta = -1 - 2i$). Verify your answer satisfies $||E(\mathbb{F}_{125})| - 126| \leq 2\sqrt{125}$.
5. Let $E/\mathbb{F}_p$ be supersingular with $a_p = 0$. Show that $|E(\mathbb{F}_{p^2})| = (p+1)^2$ and $|E(\mathbb{F}_{p^3})| = p^3 + 1$. (*Hint:* $\alpha = i\sqrt{p}$, $\beta = -i\sqrt{p}$.)
6. The characteristic polynomial $\varphi_q^2 - a_q\varphi_q + [q] = 0$ says φ_q generates a subring $\mathbb{Z}[\varphi_q] \subseteq \text{End}(E)$ isomorphic to $\mathbb{Z}[T]/(T^2 - a_qT + q)$. When $a_q^2 - 4q < 0$, this is an order in the imaginary quadratic field $\mathbb{Q}(\sqrt{a_q^2 - 4q})$. Show that for ordinary $E/\mathbb{F}_p$ (with $a_p^2 < 4p$ and $p \nmid a_p$), we have $\text{End}(E) \supseteq \mathbb{Z}[\varphi_p]$ with $\mathbb{Z}[\varphi_p]$ an order of discriminant $a_p^2 - 4p$ in $\mathbb{Q}(\sqrt{a_p^2 - 4p})$.

5.2 The Zeta Function and the Weil Conjectures

The zeta function of an algebraic variety over a finite field packages the point counts over all extension fields into a single generating function. For elliptic curves, the characteristic polynomial of Frobenius determines the zeta function completely, and the three Weil conjectures — rationality, functional equation, and the Riemann hypothesis — become immediate consequences of theorem 5.1.3 and theorem 5.1.5. This makes elliptic curves the ideal testing ground for the Weil conjectures: every general statement has an explicit, computable proof.

The Zeta Function of a Variety over $\mathbb{F}_q$

Definition 5.2.1: The Zeta Function

Let $V/\mathbb{F}_q$ be an algebraic variety. The *zeta function* of V over $\mathbb{F}_q$ is the formal power series

$$Z(V/\mathbb{F}_q, T) = \exp\left(\sum_{n=1}^{\infty} |V(\mathbb{F}_{q^n})| \cdot \frac{T^n}{n}\right) \in \mathbb{Q}[[T]]. \quad (5.10)$$

The exponential packaging may seem mysterious at first, but it is dictated by analogy with the Riemann zeta function and by the desire for an Euler product. Taking the logarithmic derivative:

$$\frac{d}{dT} \log Z(V/\mathbb{F}_q, T) = \sum_{n=1}^{\infty} |V(\mathbb{F}_{q^n})| \cdot T^{n-1},$$

so Z is the unique power series whose logarithmic derivative is the generating function for point counts. The exponential ensures Z begins with constant term 1.

Example 5.2.2: Zeta Functions of $\mathbb{P}^0$ and $\mathbb{P}^1$

(a) For $V = \mathbb{P}^0 = \text{Spec } \mathbb{F}_q$: $|V(\mathbb{F}_{q^n})| = 1$ for all n , so $Z = \exp(\sum T^n/n) = \exp(-\log(1-T)) = 1/(1-T)$.

(b) For $V = \mathbb{P}^1$: $|\mathbb{P}^1(\mathbb{F}_{q^n})| = q^n + 1$, so

$$Z(\mathbb{P}^1/\mathbb{F}_q, T) = \exp\left(\sum_{n=1}^{\infty} (q^n + 1) \frac{T^n}{n}\right) = \exp(-\log(1 - qT) - \log(1 - T)) = \frac{1}{(1 - T)(1 - qT)}.$$

The zeta function of $\mathbb{P}^1$ has “poles” at $T = 1$ and $T = 1/q$, corresponding to H^0 and H^2 of $\mathbb{P}^1$. For an elliptic curve, H^1 contributes a numerator.

The Zeta Function of an Elliptic Curve

Theorem 5.2.3: The Zeta Function of $E/\mathbb{F}_q$

Let $E/\mathbb{F}_q$ be an elliptic curve with $a_q = q + 1 - |E(\mathbb{F}_q)|$ and Frobenius eigenvalues α, β (roots of $T^2 - a_q T + q$). Then:

$$Z(E/\mathbb{F}_q, T) = \frac{1 - a_q T + qT^2}{(1 - T)(1 - qT)}. \quad (5.11)$$

Equivalently, $Z(E/\mathbb{F}_q, T) = \frac{(1 - \alpha T)(1 - \beta T)}{(1 - T)(1 - qT)}$.

Proof:

We compute $\log Z$ by substituting the formula $|E(\mathbb{F}_{q^n})| = q^n + 1 - \alpha^n - \beta^n$ (proposition 5.1.6):

$$\begin{aligned} \log Z &= \sum_{n=1}^{\infty} (q^n + 1 - \alpha^n - \beta^n) \frac{T^n}{n} \\ &= -\log(1 - qT) - \log(1 - T) + \log(1 - \alpha T) + \log(1 - \beta T), \end{aligned}$$

using $\sum_{n=1}^{\infty} a^n T^n / n = -\log(1 - aT)$ for $|aT| < 1$. Exponentiating:

$$Z = \frac{(1 - \alpha T)(1 - \beta T)}{(1 - T)(1 - qT)}.$$

Since $\alpha + \beta = a_q$ and $\alpha\beta = q$: $(1 - \alpha T)(1 - \beta T) = 1 - (\alpha + \beta)T + \alpha\beta T^2 = 1 - a_q T + qT^2$.

The zeta function is a rational function of T — the numerator $1 - a_q T + qT^2$ is the “characteristic polynomial of Frobenius acting on H^1 ,” reversed and evaluated at T . This rationality is the first of the Weil conjectures.

The Weil Conjectures for Elliptic Curves

In 1949, Weil formulated a series of conjectures about the zeta functions of smooth projective varieties over finite fields. For elliptic curves, all three conjectures follow from the results of section 5.1. We state and prove them as a single theorem.

Theorem 5.2.4: The Weil Conjectures for Elliptic Curves

Let $E/\mathbb{F}_q$ be an elliptic curve. The zeta function $Z(E/\mathbb{F}_q, T)$ satisfies:

1. (*Rationality.*) $Z(E/\mathbb{F}_q, T) \in \mathbb{Q}(T)$. More precisely, it has the form (5.11).
2. (*Functional equation.*) $Z(E/\mathbb{F}_q, 1/(qT)) = Z(E/\mathbb{F}_q, T)$.
3. (*Riemann hypothesis.*) The reciprocal roots of the numerator $1 - a_q T + qT^2$ (i.e., α^{-1} and β^{-1} , which are the zeros of Z in the variable T) satisfy $|\alpha^{-1}| = |\beta^{-1}| = q^{-1/2}$. Equivalently, the zeros of the “completed” zeta function lie on the line $\text{Re}(s) = 1/2$ when one writes $T = q^{-s}$.

Proof:

(1) This is theorem 5.2.3.

(2) We verify directly. Substituting $T \mapsto 1/(qT)$ in (5.11):

$$Z(E/\mathbb{F}_q, 1/(qT)) = \frac{1 - a_q/(qT) + q/(qT)^2}{(1 - 1/(qT))(1 - q/(qT))} = \frac{1 - a_q/(qT) + 1/(qT^2)}{(1 - 1/(qT))(1 - 1/T)}.$$

Multiplying numerator and denominator by $(qT)^2$:

$$= \frac{q^2 T^2 - a_q qT + q}{(qT - 1)(qT - qT^2/T)} \dots$$

It is cleaner to use the factored form. We have $Z = \frac{(1-\alpha T)(1-\beta T)}{(1-T)(1-qT)}$. Under $T \mapsto 1/(qT)$:

$$\frac{(1 - \alpha/(qT))(1 - \beta/(qT))}{(1 - 1/(qT))(1 - 1/T)} = \frac{\frac{qT - \alpha}{qT} \cdot \frac{qT - \beta}{qT}}{\frac{qT - 1}{qT} \cdot \frac{T - 1}{T}}.$$

Since $\alpha\beta = q$, we have $qT - \alpha = \alpha(\alpha^{-1}qT - 1) = \alpha(\beta T - 1) = -\alpha(1 - \beta T)$. Similarly $qT - \beta = -\beta(1 - \alpha T)$. And $qT - 1 = -(1 - qT)$, $T - 1 = -(1 - T)$. Substituting:

$$\frac{(-\alpha)(1 - \beta T) \cdot (-\beta)(1 - \alpha T)/(qT)^2}{(-(1 - qT))(-(1 - T))/(qT \cdot T)} = \frac{\alpha\beta(1 - \alpha T)(1 - \beta T)/(qT)^2}{(1 - T)(1 - qT)/(qT^2)}.$$

Since $\alpha\beta = q$: numerator factor = $q(1 - \alpha T)(1 - \beta T)/(q^2 T^2)$ and denominator factor = $(1 - T)(1 - qT)/(qT^2)$. So:

$$Z(1/(qT)) = \frac{q(1 - \alpha T)(1 - \beta T)/(q^2 T^2)}{(1 - T)(1 - qT)/(qT^2)} = \frac{q \cdot qT^2}{q^2 T^2} \cdot \frac{(1 - \alpha T)(1 - \beta T)}{(1 - T)(1 - qT)} = Z(T).$$

(3) By theorem 5.1.5, $|\alpha| = |\beta| = \sqrt{q}$. Therefore $|\alpha^{-1}| = |\beta^{-1}| = 1/\sqrt{q} = q^{-1/2}$. Writing $T = q^{-s}$, the zeros of Z (in the T -variable) are $T = \alpha^{-1}$ and $T = \beta^{-1}$, satisfying $|T| = q^{-1/2}$. The equation $|q^{-s}| = q^{-1/2}$ gives $q^{-\text{Re}(s)} = q^{-1/2}$, i.e., $\text{Re}(s) = 1/2$.

The functional equation $Z(1/(qT)) = Z(T)$ is the analogue of the functional equation $\xi(1-s) = \xi(s)$ for the completed Riemann zeta function $\xi(s) = \pi^{-s/2}\Gamma(s/2)\zeta(s)$. Under $T = q^{-s}$, the substitution $T \mapsto 1/(qT)$ becomes $s \mapsto 1-s$, so the functional equation says Z is invariant under $s \mapsto 1-s$ — exactly the symmetry about $\operatorname{Re}(s) = 1/2$. The Riemann hypothesis for $E/\mathbb{F}_q$ is then the statement that all zeros of the numerator lie on this critical line.

The L -Function

It is conventional to isolate the “interesting” part of the zeta function by defining:

Definition 5.2.5: The L -Function of $E/\mathbb{F}_q$

The L -function of $E/\mathbb{F}_q$ is

$$L(E/\mathbb{F}_q, T) = 1 - a_q T + qT^2 = (1 - \alpha T)(1 - \beta T). \quad (5.12)$$

Equivalently, $Z(E/\mathbb{F}_q, T) = L(E/\mathbb{F}_q, T)/((1-T)(1-qT))$.

The L -function carries the arithmetic content: a_q is the trace of Frobenius, the roots α, β determine all point counts, and the Riemann hypothesis is the statement $|\alpha| = |\beta| = \sqrt{q}$. The denominator $(1-T)(1-qT)$ is “topological” — it is the same for all smooth projective curves of genus 1 and reflects the contributions of H^0 and H^2 .

Over a number field K , the global L -function $L(E/K, s)$ will be defined as an Euler product over all primes, with the local factor at a prime $\mathfrak{p}$ of good reduction being $L(E/\mathbb{F}_q, q^{-s})^{-1}$ (where $q = |\mathcal{O}_K/\mathfrak{p}|$). This will be developed in ??.

The Analogy with the Riemann Zeta Function

The zeta function of $E/\mathbb{F}_q$ is the finite-field analogue of the Hasse–Weil L -function of $E/\mathbb{Q}$, which is in turn an analogue of the Riemann zeta function. The following table summarizes the parallels:

	Riemann $\zeta(s)$	$Z(E/\mathbb{F}_q, T)$ with $T = q^{-s}$
Definition	$\prod_p (1 - p^{-s})^{-1}$	$\exp(\sum E(\mathbb{F}_{q^n}) T^n / n)$
Rationality	Meromorphic on $\mathbb{C}$	Rational function of T
Functional eqn	$\xi(s) = \xi(1-s)$	$Z(1/(qT)) = Z(T)$
Riemann hyp.	Zeros on $\operatorname{Re}(s) = 1/2$ (open!)	$ \alpha = \sqrt{q}$ (proved!)
Poles	$s = 1$	$T = 1$ and $T = 1/q$

The crucial difference is in the last row of the “Riemann hypothesis”: for elliptic curves over finite fields, the analogue of the Riemann hypothesis is a *theorem* (Hasse’s theorem, theorem 5.1.5), while the classical Riemann hypothesis remains one of the great open problems in mathematics. This was Weil’s original insight: the finite-field setting is a “model” for the classical case, in which the conjectures can be proved using algebraic geometry.

5.2.6 Remark: Weil's General Conjectures For a smooth projective variety V of dimension d over $\mathbb{F}_q$, Weil conjectured (1949) that $Z(V/\mathbb{F}_q, T)$ is a rational function of the form

$$Z(V/\mathbb{F}_q, T) = \frac{P_1(T)P_3(T) \cdots P_{2d-1}(T)}{P_0(T)P_2(T) \cdots P_{2d}(T)},$$

where $P_i(T) \in \mathbb{Z}[T]$ with $P_0(T) = 1 - T$, $P_{2d}(T) = 1 - q^d T$, and the reciprocal roots of P_i have absolute value $q^{i/2}$ (the “Riemann hypothesis”). The functional equation relates $Z(1/(q^d T))$ to $Z(T)$, with a sign determined by the Euler characteristic. Rationality was proved by Dwork (1960) using p -adic analysis, the functional equation and the link to étale cohomology by Grothendieck (1965), and the Riemann hypothesis by Deligne (1974) — one of the great achievements of 20th-century mathematics.

For an elliptic curve ($d = 1$, genus $g = 1$): $P_0 = 1 - T$, $P_1 = 1 - a_q T + qT^2$ (the L -function), $P_2 = 1 - qT$. The Riemann hypothesis for P_1 is exactly $|\alpha| = q^{1/2}$, i.e., Hasse's theorem. Our proof via the characteristic polynomial and the degree form is, in essence, Weil's original argument for the case of elliptic curves.

Running Examples

Example 5.2.7: Zeta Functions of Running Curves

(a) $E : y^2 = x^3 - x$ over $\mathbb{F}_5$. We have $a_5 = -2$, so $L = 1 + 2T + 5T^2 = (1 - (-1 + 2i)T)(1 - (-1 - 2i)T)$ and

$$Z(E/\mathbb{F}_5, T) = \frac{1 + 2T + 5T^2}{(1 - T)(1 - 5T)}.$$

Check at $T = 1/5$: $Z(1/5) = (1 + 2/5 + 1/5)/((1 - 1/5) \cdot 0)$ — the pole at $T = 1/q = 1/5$ reflects the growth $|E(\mathbb{F}_{5^n})| \sim 5^n$.

(b) $E : y^2 = x^3 + 1$ over $\mathbb{F}_5$. $a_5 = 0$, so $L = 1 + 5T^2$ and

$$Z(E/\mathbb{F}_5, T) = \frac{1 + 5T^2}{(1 - T)(1 - 5T)}.$$

The vanishing of a_5 (supersingular!) makes the linear term in L disappear: the L -function is “palindromic” in a stronger sense. The roots $\alpha = i\sqrt{5}$, $\beta = -i\sqrt{5}$ are purely imaginary.

(c) *Functional equation check for (a).* $Z(1/(5T)) = \frac{1+2/(5T)+5/(5T)^2}{(1-1/(5T))(1-5/(5T))} = \frac{1+2/(5T)+1/(5T^2)}{(1-1/(5T))(1-1/T)}$. Multiply numerator and denominator by $(5T)^2 = 25T^2$: numerator becomes $25T^2 + 10T + 1 = (5T + 1)^2$... wait, $25T^2 + 10T + 1 = (5T)^2 + 2(5T) + 1$. And the original numerator is $1 + 2T + 5T^2 = 5T^2 + 2T + 1$. So the numerator after substitution is $5(5T^2 + 2T + 1)/(5T)^2 \cdot (5T)^2 = 5(5T^2 + 2T + 1)$... let us just verify the factored form: $Z(1/(5T)) = \frac{(1+(1-2i)/(5T))(1+(1+2i)/(5T))}{(1-1/(5T))(1-1/T)} = Z(T)$ by the argument in the proof of theorem 5.2.4(2).

Exercise 5.2.1

1. Compute $Z(E/\mathbb{F}_7, T)$ for $E : y^2 = x^3 - x$ over $\mathbb{F}_7$ (where $a_7 = 0$). Verify the functional equation $Z(1/(7T)) = Z(T)$.
2. For a smooth projective curve $C/\mathbb{F}_q$ of genus g , the zeta function has the form $Z(C/\mathbb{F}_q, T) = P(T)/((1-T)(1-qT))$ where $P(T) \in \mathbb{Z}[T]$ has degree $2g$ (this is a consequence of the Weil conjectures). Verify this for genus $g = 0$ ($C = \mathbb{P}^1$: $P = 1$) and genus $g = 1$ ($C = E$: $P = 1 - a_q T + qT^2$).
3. Show that a_q can be recovered from the zeta function: $a_q = -(d/dT) \log L(E/\mathbb{F}_q, T)|_{T=0}$, where $L = 1 - a_q T + qT^2$. More generally, show that $|E(\mathbb{F}_{q^n})|$ is determined by Z for all n .
4. Define $Z^*(E/\mathbb{F}_q, T) = \prod_{\text{closed points } x \in E} (1 - T^{\deg(x)})^{-1}$, where the product is over closed points of the scheme $E/\mathbb{F}_q$ and $\deg(x) = [\kappa(x) : \mathbb{F}_q]$ is the degree of the residue field. Show that $Z^*(E/\mathbb{F}_q, T) = Z(E/\mathbb{F}_q, T)$. (*Hint*: a closed point of degree d contributes d geometric points over $\mathbb{F}_{q^d}$, and the logarithmic derivative of the product counts each geometric point with the right multiplicity.)
5. Let $E/\mathbb{F}_p$ be supersingular with $a_p = 0$. Compute $|E(\mathbb{F}_{p^n})|$ for $n = 1, 2, 3, 4, 5, 6$. Show that $|E(\mathbb{F}_{p^n})|$ depends only on $n \pmod{4}$: the sequence $|E(\mathbb{F}_{p^n})| - (p^n + 1)$ is periodic with period 4.
6. Show that the zeta function determines the isogeny class: if $Z(E_1/\mathbb{F}_q, T) = Z(E_2/\mathbb{F}_q, T)$, then $a_q(E_1) = a_q(E_2)$, and (by Tate's theorem, section 5.4) E_1 and E_2 are isogenous over $\mathbb{F}_q$.

5.3 The Group Structure of $E(\mathbb{F}_q)$

We know $E(\mathbb{F}_q)$ is a finite abelian group of order $q + 1 - a_q$. The structure theorem for finite abelian groups tells us $E(\mathbb{F}_q) \cong \mathbb{Z}/d_1\mathbb{Z} \times \mathbb{Z}/d_2\mathbb{Z}$ for unique positive integers $d_1 | d_2$ with $d_1 d_2 = |E(\mathbb{F}_q)|$. But which pairs (d_1, d_2) actually occur? The Weil pairing imposes a strong constraint: d_1 must divide $q - 1$. This section establishes this constraint, determines the group structure in terms of the characteristic polynomial, and describes the Deuring–Waterhouse theorem on which groups are realized.

The Structure Theorem

Theorem 5.3.1: Group Structure of $E(\mathbb{F}_q)$

Let $E/\mathbb{F}_q$ be an elliptic curve. Then $E(\mathbb{F}_q)$ is isomorphic to one of the following:

1. $\mathbb{Z}/N\mathbb{Z}$ (cyclic), or
2. $\mathbb{Z}/d_1\mathbb{Z} \times \mathbb{Z}/d_2\mathbb{Z}$ with $1 < d_1 | d_2$ and $d_1 d_2 = N$,

where $N = |E(\mathbb{F}_q)| = q + 1 - a_q$. Moreover, in case (2), $d_1 | \gcd(d_2, q - 1)$.

Proof:

By the structure theorem for finite abelian groups, $E(\mathbb{F}_q) \cong \mathbb{Z}/d_1\mathbb{Z} \times \mathbb{Z}/d_2\mathbb{Z}$ with $d_1 \mid d_2$ (and $d_1 \geq 1$). It remains to show $d_1 \mid q - 1$.

Suppose $d_1 > 1$ and let ℓ be a prime dividing d_1 . Then $E(\mathbb{F}_q)$ contains a subgroup isomorphic to $(\mathbb{Z}/\ell\mathbb{Z})^2$, i.e., $E[\ell](\mathbb{F}_q) \supseteq (\mathbb{Z}/\ell\mathbb{Z})^2$. Choose generators $P, Q \in E(\mathbb{F}_q)$ of this subgroup. The Weil pairing $e_\ell : E[\ell] \times E[\ell] \rightarrow \mu_\ell$ (theorem 2.5.5) is nondegenerate and Galois-equivariant. Since $P, Q \in E(\mathbb{F}_q)$ are defined over $\mathbb{F}_q$, the Frobenius φ_q fixes them: $\varphi_q(P) = P$ and $\varphi_q(Q) = Q$. By Galois equivariance:

$$\varphi_q(e_\ell(P, Q)) = e_\ell(\varphi_q(P), \varphi_q(Q)) = e_\ell(P, Q).$$

So $e_\ell(P, Q)$ is fixed by φ_q , which means $e_\ell(P, Q) \in \mathbb{F}_q^*$. By nondegeneracy, $e_\ell(P, Q)$ is a primitive ℓ -th root of unity (since P, Q generate $(\mathbb{Z}/\ell\mathbb{Z})^2 = E[\ell]$). Therefore $\mu_\ell \subseteq \mathbb{F}_q^*$, which means $\ell \mid |\mathbb{F}_q^*| = q - 1$.

Since this holds for every prime $\ell \mid d_1$, and in fact for every prime power $\ell^k \mid d_1$ (by the same argument with $E[\ell^k]$ and e_{ℓ^k}), we conclude $d_1 \mid q - 1$.

The constraint $d_1 \mid q - 1$ is a purely arithmetic consequence of the Weil pairing. It says that $E(\mathbb{F}_q)$ cannot be “too non-cyclic”: the non-cyclic part d_1 is bounded by $q - 1$. For large q , this is a weak constraint (since $q - 1$ has many divisors), but for small q it is very restrictive. For example, over $\mathbb{F}_2$: $q - 1 = 1$, so $d_1 = 1$ and $E(\mathbb{F}_2)$ is always cyclic.

Example 5.3.2: Group Structures of Running Curves

(a) $E : y^2 = x^3 - x$ over $\mathbb{F}_5$. $|E(\mathbb{F}_5)| = 8$, $q - 1 = 4$. Possible structures: $\mathbb{Z}/8\mathbb{Z}$, $\mathbb{Z}/2\mathbb{Z} \times \mathbb{Z}/4\mathbb{Z}$, $(\mathbb{Z}/2\mathbb{Z})^2 \times \mathbb{Z}/2\mathbb{Z}$... but $d_1 \mid \gcd(d_2, 4)$. With $d_1 d_2 = 8$: if $d_1 = 2$, $d_2 = 4$, then $d_1 \mid \gcd(4, 4) = 4$. If $d_1 = 4$, $d_2 = 2$: impossible since $d_1 \mid d_2$ requires $4 \mid 2$. If $d_1 = 1$: $\mathbb{Z}/8\mathbb{Z}$.

To determine which: the 2-torsion points are $(0, 0)$, $(1, 0)$, $(-1, 0)$, so $|E[2](\mathbb{F}_5)| = 4$ (including O). Since $E[2](\mathbb{F}_5) \cong (\mathbb{Z}/2\mathbb{Z})^2$, we have $d_1 \geq 2$, forcing $E(\mathbb{F}_5) \cong \mathbb{Z}/2\mathbb{Z} \times \mathbb{Z}/4\mathbb{Z}$.

(b) $E : y^2 = x^3 + 1$ over $\mathbb{F}_7$. $|E(\mathbb{F}_7)| = 12$ (by computing: $x = 0 : y^2 = 1$ (2 pts); $x = 1 : y^2 = 2$ (non-res, 0); $x = 2 : y^2 = 9 = 2^2$ (2 pts); $x = 3 : y^2 = 28 \equiv 0$ (1 pt); $x = 4 : y^2 = 65 \equiv 2$ (non-res, 0); $x = 5 : y^2 = 126 \equiv 0$ (1 pt); $x = 6 : y^2 = 217 \equiv 0$ (1 pt). Affine: 7, total $|E(\mathbb{F}_7)| = 8$... let me recount). Actually: $x^3 + 1 \pmod{7}$: $0^3 + 1 = 1$ (QR: $1 = 1^2$, 2 pts); $1^3 + 1 = 2$ (non-QR, 0); $2^3 + 1 = 9 \equiv 2$ (non-QR, 0); $3^3 + 1 = 28 \equiv 0$ (1 pt); $4^3 + 1 = 65 \equiv 2$ (non-QR, 0); $5^3 + 1 = 126 \equiv 0$ (1 pt); $6^3 + 1 = 217 \equiv 0$ (1 pt). Affine: $2 + 0 + 0 + 1 + 0 + 1 + 1 = 5$, total = 6. So $|E(\mathbb{F}_7)| = 6$, $a_7 = 7 + 1 - 6 = 2$.

Possible structures with $N = 6$, $q - 1 = 6$: $\mathbb{Z}/6\mathbb{Z}$ (with $d_1 = 1$), or $\mathbb{Z}/2\mathbb{Z} \times \mathbb{Z}/3\mathbb{Z} \cong \mathbb{Z}/6\mathbb{Z}$ (since $\gcd(2, 3) = 1$, this is the same). Also $\mathbb{Z}/2\mathbb{Z} \times \mathbb{Z}/3\mathbb{Z}$ with $d_1 = 2$, $d_2 = 3$: but $d_1 \mid d_2$ requires $2 \mid 3$, which fails. So the only possibility is $E(\mathbb{F}_7) \cong \mathbb{Z}/6\mathbb{Z}$.

(c) $E : y^2 = x^3 - x$ over $\mathbb{F}_7$. From example 5.1.7(c), $|E(\mathbb{F}_7)| = 8$, $q - 1 = 6$. The 2-torsion: $x^3 - x = x(x - 1)(x + 1)$, roots at $x = 0, 1, 6$, giving three affine 2-torsion points. So $E[2](\mathbb{F}_7) \cong (\mathbb{Z}/2\mathbb{Z})^2$ and

$d_1 \geq 2$. With $d_1 d_2 = 8$ and $d_1 \mid 6$: $d_1 = 2$, $d_2 = 4$, and $d_1 \mid \gcd(4, 6) = 2$. So $E(\mathbb{F}_7) \cong \mathbb{Z}/2\mathbb{Z} \times \mathbb{Z}/4\mathbb{Z}$.

The Structure in Terms of the Characteristic Polynomial

The group structure of $E(\mathbb{F}_q)$ is determined by a_q and the factorization of q (though not uniquely in general — different curves with the same a_q may have different group structures).

Proposition 5.3.3: Invariant Factors from the Characteristic Polynomial

Let $E/\mathbb{F}_q$ have Frobenius eigenvalues α, β (roots of $T^2 - a_q T + q$). For a prime $\ell \neq p$:

$$E[\ell^k](\mathbb{F}_q) \cong \ker(\varphi_q - [1] \text{ on } E[\ell^k](\overline{\mathbb{F}_q})). \quad (5.13)$$

The ℓ -part of $E(\mathbb{F}_q)$ is determined by the action of φ_q on $T_\ell(E) \cong \mathbb{Z}_\ell^2$: it is the cokernel of $(\varphi_q - 1) : T_\ell(E) \rightarrow T_\ell(E)$, which has order $|\det(\varphi_q - 1)|_\ell^{-1} = |(\alpha - 1)(\beta - 1)|_\ell^{-1}$.

For the p -part: the structure depends on whether E is ordinary or supersingular. If E is ordinary, the p -part of $E(\mathbb{F}_q)$ is cyclic (because $E[p^k](\overline{\mathbb{F}_q}) \cong \mathbb{Z}/p^k\mathbb{Z}$, proposition 2.3.8). If E is supersingular, the p -part is trivial ($E[p](\mathbb{F}_q) = \{O\}$).

More concretely: d_1 is the largest integer dividing both $\alpha - 1$ and $\beta - 1$ in the ring $\mathbb{Z}[\alpha]$ (or equivalently, $d_1 = \gcd(N, \alpha - 1, q - 1)$ where the gcd is taken in $\mathbb{Z}$ using $\alpha + \beta = a_q$ and $\alpha\beta = q$). In practice, one computes $d_1 = \gcd(N, q - 1)$ as an upper bound, then checks divisibility by small primes.

The Deuring–Waterhouse Theorem

Which groups actually occur as $E(\mathbb{F}_q)$ for some elliptic curve $E/\mathbb{F}_q$? The answer is given by a remarkable classification theorem.

Theorem 5.3.4: Deuring–Waterhouse

Let $q = p^r$ and let N be a positive integer. There exists an elliptic curve $E/\mathbb{F}_q$ with $|E(\mathbb{F}_q)| = N$ if and only if one of the following holds, where $a = q + 1 - N$:

1. $\gcd(a, p) = 1$ and $a^2 \leq 4q$.
2. r is even and $a = \pm 2\sqrt{q}$.
3. r is even and $a = \pm\sqrt{q}$, with $p \not\equiv 1 \pmod{3}$.
4. r is odd and $a = 0$, or r is even and $a = 0$.
5. r is even and $a = \pm p^{r/2}\sqrt{p}$, with $p \in \{2, 3\}$.

Cases (2)–(5) correspond to special supersingular curves. In case (1), E is ordinary (when $p \nmid a$) or supersingular (when $p \mid a$ but $\gcd(a, p) = 1$, which cannot happen... rather, case (1) with $p \mid a$ is impossible since $\gcd(a, p) = 1$). So case (1) gives all ordinary curves and some supersingular ones are in cases (2)–(5).

Proof Key ideas:

The “only if” direction is the content of Hasse’s theorem (theorem 5.1.5): $|a| \leq 2\sqrt{q}$. The additional restrictions in cases (2)–(5) come from the integrality of the Frobenius eigenvalues (they must be algebraic integers of a specific type, determined by the factorization of p in $\mathbb{Z}[\alpha]$) and the Honda–Tate theorem, which classifies isogeny classes of abelian varieties over finite fields.

The “if” direction requires constructing curves with prescribed point counts. For case (1) (the generic ordinary case), one uses the theory of complex multiplication: by the CM lifting theorem, an order $\mathcal{O}$ of discriminant $a^2 - 4q$ in $\mathbb{Q}(\sqrt{a^2 - 4q})$ has a CM elliptic curve $E/\mathbb{F}_q$ with $\text{End}(E) \supseteq \mathcal{O}$ and $a_q = a$.

For the full proof, see Waterhouse [27].

The Deuring–Waterhouse theorem, combined with theorem 5.3.1, gives a complete picture: for each q and each valid trace a_q , the order $N = q + 1 - a_q$ is determined, and the possible group structures $\mathbb{Z}/d_1\mathbb{Z} \times \mathbb{Z}/d_2\mathbb{Z}$ (with $d_1|d_2$, $d_1d_2 = N$, $d_1|q-1$) are finitely many. Whether each structure is realized by some curve with that trace requires a finer analysis, but the constraints $d_1|d_2$ and $d_1|q-1$ are the main tools.

The Full n -Torsion over $\mathbb{F}_q$

A natural question is: for which n does $E(\mathbb{F}_q)$ contain the full n -torsion $E[n] \cong (\mathbb{Z}/n\mathbb{Z})^2$?

Proposition 5.3.5: Full n -Torsion over $\mathbb{F}_q$

Let $E/\mathbb{F}_q$ be an elliptic curve and $n \geq 1$ an integer with $\gcd(n, p) = 1$. Then $E[n](\mathbb{F}_q) = E[n](\mathbb{F}_q) \cong (\mathbb{Z}/n\mathbb{Z})^2$ if and only if $n|q-1$ and $n^2||E(\mathbb{F}_q)|$.

Proof:

$E[n](\mathbb{F}_q) = (\mathbb{Z}/n\mathbb{Z})^2$ iff φ_q acts as the identity on $E[n]$, i.e., $\varphi_q \equiv [1] \pmod{n}$ on $T_\ell(E)$ for every prime $\ell|n$. This means the characteristic polynomial $T^2 - a_qT + q \equiv (T-1)^2 \pmod{n}$, i.e., $a_q \equiv 2 \pmod{n}$ and $q \equiv 1 \pmod{n}$. The condition $q \equiv 1 \pmod{n}$ is $n|q-1$. The condition $a_q \equiv 2 \pmod{n}$ combined with $q \equiv 1 \pmod{n}$ gives $|E(\mathbb{F}_q)| = q + 1 - a_q \equiv 1 + 1 - 2 = 0 \pmod{n}$... but we need $n^2||E(\mathbb{F}_q)|$, not just $n||E(\mathbb{F}_q)|$.

More carefully: $E[n](\mathbb{F}_q) \cong (\mathbb{Z}/n\mathbb{Z})^2$ iff $n|d_1$ (where $E(\mathbb{F}_q) \cong \mathbb{Z}/d_1\mathbb{Z} \times \mathbb{Z}/d_2\mathbb{Z}$). Since $d_1|q-1$, this requires $n|q-1$. And $n|d_1$ with $d_1|d_2$ gives $n^2|d_1d_2 = |E(\mathbb{F}_q)|$. Conversely, if $n|q-1$ and $n^2|N$, then the Weil pairing constraint is satisfied and the n -torsion is rational.

This is relevant for pairing-based cryptography: protocols like the MOV attack and identity-based encryption require $E[n](\mathbb{F}_q) = (\mathbb{Z}/n\mathbb{Z})^2$ (or equivalently, the Weil pairing takes values in $\mathbb{F}_q^*$), which happens iff $n|q-1$.

Application: Constructing Curves with Prescribed Order

The Deuring–Waterhouse theorem guarantees the existence of curves with any valid point count. In practice, one often needs an explicit curve, not just an existence result. Two approaches are:

(1) *The CM method.* Given a target group order N over $\mathbb{F}_q$ (with $a = q + 1 - N$ satisfying Hasse’s bound), one finds an imaginary quadratic order $\mathcal{O}$ of discriminant $D = a^2 - 4q < 0$, computes the Hilbert class polynomial $H_D(X) \in \mathbb{Z}[X]$ (a polynomial whose roots are the j -invariants of CM elliptic curves with endomorphism ring $\mathcal{O}$), and reduces H_D modulo p to find a j -invariant $j_0 \in \mathbb{F}_q$. Any curve with $j = j_0$ has the correct trace (up to a twist). This method will be developed in ??.

(2) *Random search with Schoof’s algorithm.* Generate random Weierstrass equations over $\mathbb{F}_q$ and compute $|E(\mathbb{F}_q)|$ using Schoof’s algorithm (to be discussed in section 5.5), which runs in polynomial time in $\log q$. This is the method used in practice for cryptographic curve generation.

Example 5.3.6: Cryptographically Relevant Group Orders

For elliptic curve cryptography (??), one wants $E(\mathbb{F}_p)$ to have a large prime-order subgroup. Ideally $|E(\mathbb{F}_p)| = n$ is itself prime, or $|E(\mathbb{F}_p)| = h \cdot n$ with h small (the “cofactor”) and n a large prime.

For $p = 2^{256} - 2^{224} + 2^{192} + 2^{96} - 1$ (the NIST prime P-256): the curve P-256 has $|E(\mathbb{F}_p)| = n$ where n is a 256-bit prime. This was found by searching over curves using Schoof’s algorithm.

Exercise 5.3.1

1. Show that $E(\mathbb{F}_q)$ is cyclic if and only if $\gcd(|E(\mathbb{F}_q)|, q - 1) \in \{1\} \cup \{\text{squarefree}\}$... more precisely, show $E(\mathbb{F}_q)$ is cyclic iff $d_1 = 1$, and give a criterion in terms of N and $q - 1$. (*Hint:* $E(\mathbb{F}_q)$ is cyclic iff for every prime ℓ , the ℓ -part of $E(\mathbb{F}_q)$ is cyclic. The ℓ -part is non-cyclic iff $\ell | d_1$, iff $\ell^2 | N$ and $\ell | q - 1$.)
2. Determine all possible group structures $E(\mathbb{F}_2)$ for elliptic curves over $\mathbb{F}_2$. (*Hint:* $|a_2| \leq 2\sqrt{2} \approx 2.83$, so $a_2 \in \{-2, -1, 0, 1, 2\}$ and $|E(\mathbb{F}_2)| \in \{1, 2, 3, 4, 5\}$. Since $q - 1 = 1$, every group is cyclic.)
3. For $E : y^2 = x^3 - x$ over $\mathbb{F}_5$, the group is $E(\mathbb{F}_5) \cong \mathbb{Z}/2\mathbb{Z} \times \mathbb{Z}/4\mathbb{Z}$. The Weil pairing $e_2 : E[2] \times E[2] \rightarrow \mu_2$ must satisfy $e_2(P, Q) = -1$ for a suitable basis P, Q of $E[2]$. Since $\mu_2 = \{1, -1\} \subseteq \mathbb{F}_5^*$, this is consistent with $2 | q - 1 = 4$. Compute $e_2((0, 0), (1, 0))$ using the algebraic construction (theorem 2.5.5) and verify it equals -1 .
4. Determine $E(\mathbb{F}_{25})$ for $E : y^2 = x^3 - x$. Use $|E(\mathbb{F}_{25})| = 32$ (from example 5.1.7(a)) and $q - 1 = 24$ to constrain the structure. Is $E(\mathbb{F}_{25})$ cyclic?
5. Let $E/\mathbb{F}_p$ be supersingular with $a_p = 0$, so $|E(\mathbb{F}_p)| = p + 1$. Show that $E(\mathbb{F}_p)$ is cyclic when $p \equiv 2 \pmod{3}$ and $p \geq 5$. (*Hint:* show that $\gcd(p + 1, p - 1) = 2$, so $d_1 | 2$. Then argue the 2-part is cyclic unless $4 | p + 1$ and $2 | p - 1$, which are compatible — so $d_1 = 2$ is possible. Determine when it actually occurs.)
6. For $q = 9$ ($p = 3, r = 2$), list all integers N with $(\sqrt{9} - 1)^2 \leq N \leq (\sqrt{9} + 1)^2$ (i.e., $4 \leq N \leq 16$). For each valid N (satisfying the Deuring–Waterhouse conditions), list the possible group structures. How many isomorphism classes of elliptic curves over $\mathbb{F}_9$ are there? (*You do not*

need to construct explicit curves, just apply the theorem.)

5.4 Tate's Isogeny Theorem

In corollary 2.5.4, we showed that the natural map

$$\mathrm{Hom}(E_1, E_2) \otimes_{\mathbb{Z}} \mathbb{Z}_{\ell} \hookrightarrow \mathrm{Hom}_{\mathbb{Z}_{\ell}}(T_{\ell}(E_1), T_{\ell}(E_2))$$

is injective for any elliptic curves E_1, E_2 over any field K and any prime $\ell \neq \mathrm{char} K$. Over a general field, this map is very far from surjective: for two non-isogenous curves, the left side is 0 while the right side has rank 4. Over finite fields, however, Tate proved that this map is an *isomorphism* — the Tate module functor is fully faithful. This is one of the deepest results in the arithmetic of elliptic curves over finite fields, and it has remarkable consequences: two curves are isogenous if and only if they have the same number of points, the endomorphism ring is “as large as the Tate module allows,” and the isogeny class is completely determined by the trace of Frobenius.

Statement

Theorem 5.4.1: Tate's Isogeny Theorem

Let E_1, E_2 be elliptic curves over a finite field $\mathbb{F}_q$, and let ℓ be a prime with $\ell \neq \mathrm{char} \mathbb{F}_q$. Then the natural map

$$T_{\ell} : \mathrm{Hom}(E_1, E_2) \otimes_{\mathbb{Z}} \mathbb{Z}_{\ell} \xrightarrow{\sim} \mathrm{Hom}_{\mathbb{Z}_{\ell}[G]}(T_{\ell}(E_1), T_{\ell}(E_2)) \quad (5.14)$$

is an isomorphism, where $G = \mathrm{Gal}(\overline{\mathbb{F}_q}/\mathbb{F}_q)$ and $\mathrm{Hom}_{\mathbb{Z}_{\ell}[G]}$ denotes G -equivariant $\mathbb{Z}_{\ell}$ -linear maps.

Two remarks on the statement. First, the right-hand side involves G -equivariant maps, not all $\mathbb{Z}_{\ell}$ -linear maps. The Galois group $G = \mathrm{Gal}(\overline{\mathbb{F}_q}/\mathbb{F}_q) \cong \hat{\mathbb{Z}}$ is topologically generated by the Frobenius φ_q , so G -equivariance amounts to commuting with $T_{\ell}(\varphi_q)$. Second, over an algebraically closed field $\overline{\mathbb{F}_q}$, the G -equivariance condition is vacuous, and the theorem becomes $\mathrm{Hom}(E_1, E_2) \otimes_{\mathbb{Z}} \mathbb{Z}_{\ell} \cong \mathrm{Hom}_{\mathbb{Z}_{\ell}}(T_{\ell}(E_1), T_{\ell}(E_2))$.

Isogenous Curves Have the Same Characteristic Polynomial

The first step toward Tate's theorem is an elementary observation that requires only the injectivity of the Tate module functor (established in proposition 2.5.3).

Proposition 5.4.2: Isogenous Curves Have the Same Trace

If $\varphi : E_1 \rightarrow E_2$ is an isogeny over $\mathbb{F}_q$, then $a_q(E_1) = a_q(E_2)$.

Proof:

Since φ is defined over $\mathbb{F}_q$, it commutes with the q -Frobenius: $\varphi \circ \varphi_q^{(1)} = \varphi_q^{(2)} \circ \varphi$, where $\varphi_q^{(i)}$ denotes the Frobenius on E_i . On Tate modules (for $\ell \neq p$): the map $T_\ell(\varphi) : T_\ell(E_1) \rightarrow T_\ell(E_2)$ is injective (proposition 2.5.3), and satisfies $T_\ell(\varphi) \circ T_\ell(\varphi_q^{(1)}) = T_\ell(\varphi_q^{(2)}) \circ T_\ell(\varphi)$. Tensoring with $\mathbb{Q}_\ell$: the map $V_\ell(\varphi) : V_\ell(E_1) \rightarrow V_\ell(E_2)$ is an isomorphism of $\mathbb{Q}_\ell$ -vector spaces conjugating the Frobenius actions. Conjugate linear maps have the same characteristic polynomial, so $\varphi_q^{(1)}$ and $\varphi_q^{(2)}$ have the same characteristic polynomial $T^2 - a_q T + q$. In particular, $a_q(E_1) = a_q(E_2)$.

Finiteness of the Isogeny Class**Proposition 5.4.3: Finiteness of the Isogeny Class**

Let $E/\mathbb{F}_q$ be an elliptic curve. Up to $\overline{\mathbb{F}_q}$ -isomorphism, there are only finitely many elliptic curves E' over $\overline{\mathbb{F}_q}$ that are isogenous to E .

Proof:

Since $j(E')$ determines E' up to $\overline{\mathbb{F}_q}$ -isomorphism (theorem 1.2.8), it suffices to show that only finitely many j -invariants arise.

If $\varphi : E \rightarrow E'$ is an isogeny defined over $\mathbb{F}_{q^m}$ (for some m), then E' is also defined over $\mathbb{F}_{q^m}$, and proposition 5.4.2 (applied over $\mathbb{F}_{q^m}$) gives $a_{q^m}(E') = a_{q^m}(E)$. But a_{q^m} determines the roots α^m, β^m of the Frobenius characteristic polynomial of E' over $\mathbb{F}_{q^m}$, and hence determines α, β (up to root-of-unity ambiguity that is resolved by knowing $a_{q^{m'}}$ for a second exponent m'). In particular, E' has the same characteristic polynomial as E over $\mathbb{F}_{q^m}$ for all m .

By the Deuring–Waterhouse theorem (theorem 5.3.4), there are only finitely many j -invariants over $\overline{\mathbb{F}_q}$ giving a curve with any prescribed characteristic polynomial. Hence the isogeny class is finite.

Over $\mathbb{Q}$ (or any number field), the isogeny class is also finite (by Faltings' theorem), but the proof is immensely harder. Over finite fields, the finiteness follows from the explicit classification via a_q .

Proof of Tate's Theorem

We prove the theorem in the case $E_1 = E_2 = E$ (i.e., for endomorphism rings); the general case follows by a similar argument.

Proof of theorem 5.4.1:

Write $R = \text{End}_{\mathbb{F}_q}(E)$, $M = T_\ell(E) \cong \mathbb{Z}_\ell^2$, and let $\varphi_q \in R$ be the Frobenius endomorphism, which acts on M as a matrix $A \in M_2(\mathbb{Z}_\ell)$ with characteristic polynomial $P(T) = T^2 - a_q T + q$.

Step 1: Injectivity. The map $T_\ell : R \otimes \mathbb{Z}_\ell \rightarrow \text{End}_{\mathbb{Z}_\ell[G]}(M)$ is injective by proposition 2.5.3: if $T_\ell(\varphi) = 0$, then φ kills $E[\ell^n]$ for all n , so $\ell^{2n} \mid \deg(\varphi)$ for all n , forcing $\varphi = 0$.

Step 2: The commutant dimension. The target $\text{End}_{\mathbb{Z}_\ell[G]}(M)$ consists of $\mathbb{Z}_\ell$ -linear endomorphisms of $\mathbb{Z}_\ell^2$ that commute with $A = T_\ell(\varphi_q)$. Its $\mathbb{Q}_\ell$ -rank equals $\dim_{\mathbb{Q}_\ell} C(A)$, where $C(A)$ is the commutant of A in $M_2(\mathbb{Q}_\ell)$:

$$\dim_{\mathbb{Q}_\ell} C(A) = \begin{cases} 2 & \text{if } A \text{ has distinct eigenvalues in } \overline{\mathbb{Q}_\ell}, \\ 4 & \text{if } A \text{ is a scalar matrix.} \end{cases}$$

(If A has distinct eigenvalues, its commutant is the algebra $\mathbb{Q}_\ell[A]$, which has dimension 2. If A is scalar, it commutes with everything, giving dimension 4.)

For our Frobenius: A is scalar iff $\alpha = \beta$, i.e., $a_q^2 = 4q$, i.e., $a_q = \pm 2\sqrt{q}$ (which requires q to be a perfect square). In all other cases, A has distinct eigenvalues and $\dim C(A) = 2$.

Step 3: Surjectivity via finiteness. We show $R \otimes \mathbb{Z}_\ell \rightarrow \text{End}_{\mathbb{Z}_\ell[G]}(M)$ is surjective by combining injectivity with a finiteness argument. Since the injection $R \otimes \mathbb{Z}_\ell \hookrightarrow \text{End}_{\mathbb{Z}_\ell[G]}(M)$ is between finitely generated $\mathbb{Z}_\ell$ -modules, the cokernel is a finite $\mathbb{Z}_\ell$ -module (a quotient of a finitely generated torsion-free module by a submodule of the same $\mathbb{Q}_\ell$ -rank — which we must verify). If this cokernel is nontrivial, say of order ℓ^s with $s > 0$, we derive a contradiction from the finiteness of the isogeny class.

The key point is that $R \otimes \mathbb{Q}_\ell \rightarrow \text{End}_{\mathbb{Q}_\ell[G]}(V_\ell(E))$ is an isomorphism. To see this: R contains $\mathbb{Z}[\varphi_q]$, which has rank 2 over $\mathbb{Z}$ (when $a_q^2 < 4q$, so that φ_q satisfies an irreducible quadratic). After tensoring with $\mathbb{Q}_\ell$: $R \otimes \mathbb{Q}_\ell \supseteq \mathbb{Q}_\ell[\varphi_q] = \mathbb{Q}_\ell[A]$, and the commutant $C(A) = \mathbb{Q}_\ell[A]$ has dimension 2. So $\dim(R \otimes \mathbb{Q}_\ell) \geq 2 = \dim C(A)$, and since $R \otimes \mathbb{Q}_\ell$ embeds in $C(A)$, we have equality. (In the scalar case $a_q^2 = 4q$: $\varphi_q = [\pm\sqrt{q}]$ is central, so every endomorphism defined over $\overline{\mathbb{F}_q}$ commutes with φ_q and is therefore defined over $\mathbb{F}_q$. Thus $R = \text{End}_{\mathbb{F}_q}(E) = \text{End}_{\overline{\mathbb{F}_q}}(E)$. On the Tate module side, the commutant is $M_2(\mathbb{Q}_\ell)$ of dimension 4. The claim $\dim(R \otimes \mathbb{Q}_\ell) = 4$ is then part of what we are proving — it will follow from surjectivity.)

In the generic case ($a_q^2 < 4q$): $R \otimes \mathbb{Q}_\ell \cong C(A) \cong \mathbb{Q}_\ell[A]$ with $\dim = 2$. Now suppose the cokernel $Q = \text{End}_{\mathbb{Z}_\ell[G]}(M)/(R \otimes \mathbb{Z}_\ell)$ is nontrivial, of $\mathbb{Z}_\ell$ -length $s > 0$. Then there exists $f \in \text{End}_{\mathbb{Z}_\ell[G]}(M)$ such that $\ell^s f \in R \otimes \mathbb{Z}_\ell$ but $f \notin R \otimes \mathbb{Z}_\ell$. Since $R \otimes \mathbb{Q}_\ell = C(A)$, we can write $f = \ell^{-s}\psi$ for some $\psi \in R \otimes \mathbb{Z}_\ell$.

We now use the degree form and finiteness. Any element of $R \otimes \mathbb{Z}_\ell$ can be approximated by elements of R : given $g \in R \otimes \mathbb{Z}_\ell$, there exist $\varphi_n \in R$ with $T_\ell(\varphi_n) \rightarrow g$ ℓ -adically. The degree $\deg(\varphi_n) = \det(T_\ell(\varphi_n)) \rightarrow \det(g)$, so $\deg(\varphi_n)$ is eventually bounded. But for each integer d , there are only finitely many isogenies $E \rightarrow E'$ of degree d (since each corresponds to a subgroup scheme of E of order d , and E' lies in the finite isogeny class of proposition 5.4.3). Therefore there are finitely many endomorphisms of E of any given degree, and the sequence φ_n must stabilize: $T_\ell(\varphi_n) = g$ for $n \gg 0$. This shows $g \in T_\ell(R) = R \otimes \mathbb{Z}_\ell$ (identifying R with its image), proving surjectivity.

Step 4: The scalar case. When $a_q^2 = 4q$: the argument of Step 3 works identically, except the commutant has dimension 4 instead of 2. The finiteness argument still applies: approximate any $f \in \text{End}_{\mathbb{Z}_\ell[G]}(M) = M_2(\mathbb{Z}_\ell)$ by elements $\varphi_n \in R$ with bounded degree, and the finiteness

of R in bounded degree forces $f \in R \otimes \mathbb{Z}_\ell$. The conclusion is $R \otimes \mathbb{Z}_\ell \cong M_2(\mathbb{Z}_\ell)$, which implies $\text{rank}_{\mathbb{Z}} R = 4$: the endomorphism ring is an order of rank 4 in a quaternion algebra.

For the complete details (especially the uniformity of the bounded-degree argument across all cases), see Tate [23] or Silverman [21, p. V.3].

The proof has a remarkable logical structure: the *input* is purely algebraic (injectivity of T_ℓ and finiteness of the isogeny class), but the *output* has deep consequences for the structure of $\text{End}(E)$. In particular, the rank of $\text{End}_{\mathbb{F}_q}(E)$ is *determined* by Tate's theorem, not assumed: in the generic case it equals the commutant dimension 2 (confirming $\text{End}_{\mathbb{F}_q}(E)$ is a rank-2 order in an imaginary quadratic field), and in the scalar Frobenius case it equals 4 (confirming $\text{End}_{\mathbb{F}_q}(E)$ is a rank-4 order in a quaternion algebra). This second conclusion will be made fully explicit in section 5.5, where we show the endomorphism ring is a *maximal* order.

Consequences

Corollary 5.4.4: Isogeny Classes Are Determined by a_q

Two elliptic curves E_1, E_2 over $\mathbb{F}_q$ are isogenous (over $\mathbb{F}_q$) if and only if $a_q(E_1) = a_q(E_2)$, i.e., if and only if $|E_1(\mathbb{F}_q)| = |E_2(\mathbb{F}_q)|$.

Proof:

If E_1 and E_2 are isogenous, then $a_q(E_1) = a_q(E_2)$ by proposition 5.4.2.

Conversely, if $a_q(E_1) = a_q(E_2)$, then $T_\ell(E_1)$ and $T_\ell(E_2)$ are isomorphic as $\mathbb{Z}_\ell[G]$ -modules (both have the same Frobenius action, determined by a_q and q , and a semisimple linear map on a 2-dimensional space is determined up to conjugacy by its characteristic polynomial). By theorem 5.4.1, this isomorphism lifts to an element of $\text{Hom}(E_1, E_2) \otimes \mathbb{Z}_\ell$, which is nonzero. Hence $\text{Hom}(E_1, E_2) \neq 0$, i.e., E_1 and E_2 are isogenous.

This is a remarkable rigidity result: over $\mathbb{F}_q$, the single integer a_q determines the isogeny class. Over $\mathbb{Q}$, isogeny classes are vastly more subtle (they are related to the Langlands program and modularity).

Corollary 5.4.5: Isogenous Curves Have the Same L -Function

If E_1 and E_2 are isogenous over $\mathbb{F}_q$, then $L(E_1/\mathbb{F}_q, T) = L(E_2/\mathbb{F}_q, T)$.

Proof:

By proposition 5.4.2: isogenous implies $a_q(E_1) = a_q(E_2)$, and $L = 1 - a_q T + qT^2$.

Corollary 5.4.6: The Endomorphism Algebra

Let $E/\mathbb{F}_q$ be an elliptic curve. Then $\text{End}_{\mathbb{F}_q}(E) \otimes \mathbb{Q}$ is determined by the characteristic polynomial $P(T) = T^2 - a_q T + q$:

1. If P is irreducible over $\mathbb{Q}$: $\text{End}_{\mathbb{F}_q}(E) \otimes \mathbb{Q} = \mathbb{Q}(\varphi_q) = \mathbb{Q}[T]/P(T)$, an imaginary quadratic field. In particular, $\text{rank}_{\mathbb{Z}} \text{End}_{\mathbb{F}_q}(E) = 2$.
2. If $P = (T - a)^2$ with $a = \pm\sqrt{q} \in \mathbb{Z}$: $\text{End}_{\mathbb{F}_q}(E) \otimes \mathbb{Q}$ is a 4-dimensional $\mathbb{Q}$ -algebra ($\text{rank}_{\mathbb{Z}} \text{End}_{\mathbb{F}_q}(E) = 4$). This is the quaternion algebra $B_{p,\infty}$ for supersingular curves (proved in section 5.5).

Proof:

By theorem 5.4.1: $\text{End}_{\mathbb{F}_q}(E) \otimes \mathbb{Q}_\ell \cong \text{End}_{\mathbb{Q}_\ell[G]}(V_\ell(E)) = C(A)$, the commutant of $A = T_\ell(\varphi_q)$ in $M_2(\mathbb{Q}_\ell)$. Case (1): P irreducible gives $C(A) = \mathbb{Q}_\ell[A] \cong \mathbb{Q}_\ell[T]/P(T)$, which has $\mathbb{Q}$ -form $\mathbb{Q}[T]/P(T)$, a quadratic field. Case (2): A is scalar, so $C(A) = M_2(\mathbb{Q}_\ell)$, giving $\dim = 4$.

The Theorem for Hom

The same argument gives the Hom version:

Corollary 5.4.7: Tate's Theorem for Hom

For elliptic curves E_1, E_2 over $\mathbb{F}_q$ and $\ell \neq p$:

$$\text{Hom}_{\mathbb{F}_q}(E_1, E_2) \otimes \mathbb{Z}_\ell \cong \text{Hom}_{\mathbb{Z}_\ell[G]}(T_\ell(E_1), T_\ell(E_2)).$$

In particular:

1. If $a_q(E_1) = a_q(E_2)$: $\text{Hom}(E_1, E_2)$ is a nonzero $\mathbb{Z}$ -module of rank ≤ 2 (an invertible module over $\text{End}(E_1)$ when the endomorphism ring is commutative).
2. If $a_q(E_1) \neq a_q(E_2)$: $\text{Hom}(E_1, E_2) = 0$ (the curves are not isogenous).

The isogeny graph over $\mathbb{F}_q$ is therefore very structured: the vertices are isomorphism classes of elliptic curves, the connected components are the isogeny classes (determined by a_q), and within each component the graph has rich arithmetic structure governed by the endomorphism rings and their ideal theory.

Exercise 5.4.1

1. For $E : y^2 = x^3 - x$ over $\mathbb{F}_5$ (ordinary, $a_5 = -2$), determine $\text{End}_{\mathbb{F}_5}(E)$. (*Hint*: $\mathbb{Z}[\varphi_5] = \mathbb{Z}[T]/(T^2 + 2T + 5)$. The discriminant is $4 - 20 = -16$, so $\mathbb{Q}(\varphi_5) = \mathbb{Q}(\sqrt{-4}) = \mathbb{Q}(i)$. Is $\mathbb{Z}[\varphi_5]$ the maximal order $\mathbb{Z}[i]$, or a proper suborder?)
2. Over $\mathbb{F}_5$, how many $\mathbb{F}_5$ -isomorphism classes of elliptic curves have $a_5 = -2$? (*Hint*: by corollary 5.4.4, they form a single isogeny class. The class number of the order $\mathbb{Z}[\varphi_5]$ determines the number of isomorphism classes within the isogeny class.)

3. Show that $E_1 : y^2 = x^3 - x$ and $E_2 : y^2 = x^3 + 1$ over $\mathbb{F}_5$ are not isogenous, by computing their traces and applying corollary 5.4.4.
4. Explain why Tate's theorem fails over $\mathbb{Q}$: give an example of two non-isogenous elliptic curves over $\mathbb{Q}$ whose Tate modules are isomorphic as $\mathbb{Z}_\ell$ -modules (but not as $G_{\mathbb{Q}}$ -modules). (*Hint*: any two elliptic curves have $T_\ell(E) \cong \mathbb{Z}_\ell^2$ as bare $\mathbb{Z}_\ell$ -modules.)
5. For $E : y^2 = x^3 + 1$ over $\mathbb{F}_7$ (where $a_7 = 2$, example 5.3.2(b)), compute the characteristic polynomial of Frobenius, the discriminant $a_7^2 - 4 \cdot 7 = 4 - 28 = -24$, and the order $\mathbb{Z}[\varphi_7] = \mathbb{Z}[T]/(T^2 - 2T + 7)$. Is this the maximal order in $\mathbb{Q}(\sqrt{-24}) = \mathbb{Q}(\sqrt{-6})$?
6. Tate's theorem says T_ℓ is “fully faithful” on the category of elliptic curves over $\mathbb{F}_q$ (with morphisms = isogenies). State precisely what this means in categorical language, and explain why it implies that the isogeny category of elliptic curves over $\mathbb{F}_q$ embeds into the category of $\mathbb{Z}_\ell[G]$ -modules.

5.5 Supersingular Curves and the Deuring Correspondence

The supersingular curves — those with $E[p](\overline{\mathbb{F}_p}) = \{O\}$, equivalently those whose formal group has height 2 (theorem 4.4.4) — occupy a special and beautiful corner of the theory. There are only finitely many supersingular j -invariants in each characteristic p , all lying in the small field $\mathbb{F}_{p^2}$. Their endomorphism rings are maximal orders in the quaternion algebra $B_{p,\infty}$, and the Deuring correspondence provides an explicit bijection between supersingular curves (up to isomorphism) and left ideal classes in these orders. In this section we prove the key properties of supersingular curves, complete the proof of theorem 2.6.7 (that $\text{End}(E)$ is a maximal order), and describe the supersingular isogeny graph. We close with an overview of Schoof's algorithm for computing $|E(\mathbb{F}_q)|$.

Supersingular j -Invariants

Theorem 5.5.1: Supersingular j -Invariants Lie in $\mathbb{F}_{p^2}$

Let $E/\overline{\mathbb{F}_p}$ be a supersingular elliptic curve. Then $j(E) \in \mathbb{F}_{p^2}$.

Proof:

Let $j_0 = j(E)$. Over $\overline{\mathbb{F}_p}$, E is isomorphic to a curve defined over $\mathbb{F}_p(j_0)$. We claim $j_0 \in \mathbb{F}_{p^2}$, i.e., $j_0^{p^2} = j_0$.

The p^2 -Frobenius $\varphi_{p^2} : E \rightarrow E^{(p^2)}$ is an isogeny. If E is defined over $\mathbb{F}_{p^2}$, then $E^{(p^2)} = E$ (the Frobenius twist is trivial), so $\varphi_{p^2} \in \text{End}(E)$. For a supersingular curve with $a_p = 0$: the characteristic polynomial of φ_p is $T^2 + p$, so $\varphi_p^2 = [-p]$. Then $\varphi_{p^2} = \varphi_p^2 = [-p]$, and $|E(\mathbb{F}_{p^2})| = \deg(\varphi_{p^2} - [1]) = \deg([-p] - [1]) = \deg(-(p+1)) = (p+1)^2$. This shows E has $(p+1)^2$ points over $\mathbb{F}_{p^2}$, confirming it is defined over $\mathbb{F}_{p^2}$.

More generally: for any supersingular E (not just $a_p = 0$), the trace a_p satisfies $p \mid a_p$ (theo-

rem 2.4.4). The possible values are $a_p = 0$ (always), $a_p = \pm p$ (when $p \in \{2, 3\}$), and $a_p = \pm 2\sqrt{p}$ (when p is a perfect square, but then $a_p = \pm 2p^{1/2}$, which is non-integer unless p is a perfect square). In all cases, one verifies that $j(E)^{p^2} = j(E)$ by showing $E \cong E^{(p^2)}$ over $\overline{\mathbb{F}_p}$ (since the Frobenius twist $E^{(p^2)}$ has $j(E^{(p^2)}) = j(E)^{p^2}$, and isomorphic curves have the same j -invariant).

The key input is that the Frobenius φ_{p^2} acts on E as an element of $\text{End}(E)$ with $\deg(\varphi_{p^2}) = p^2$, and for supersingular curves φ_{p^2} satisfies a quadratic equation over $\mathbb{Z}$ whose discriminant forces $j \in \mathbb{F}_{p^2}$.

Proposition 5.5.2: Counting Supersingular j -Invariants

The number of supersingular j -invariants over $\overline{\mathbb{F}_p}$ is

$$h_p = \frac{p-1}{12} + \begin{cases} 0 & \text{if } p \equiv 1 \pmod{12}, \\ \frac{1}{3} & \text{if } p \equiv 5 \text{ or } 7 \pmod{12}, \\ \frac{1}{2} & \text{if } p \equiv 3 \pmod{4} \text{ and } p \not\equiv 7 \pmod{12}, \\ \frac{5}{6} & \text{if } p \equiv 11 \pmod{12}. \end{cases} \quad (5.15)$$

In particular, $h_p = \lfloor p/12 \rfloor + \epsilon$ where $\epsilon \in \{0, 1, 2\}$.

Proof:

The Eichler mass formula (also called the Deuring mass formula) states:

$$\sum_{\substack{j \in \overline{\mathbb{F}_p} \\ j \text{ supersingular}}} \frac{1}{|\text{Aut}(E_j)/\{\pm 1\}|} = \frac{p-1}{24}, \quad (5.16)$$

where the sum is over supersingular j -invariants and $|\text{Aut}(E_j)/\{\pm 1\}|$ is the order of the automorphism group modulo the hyperelliptic involution. For a generic j -invariant ($j \neq 0, 1728$): $|\text{Aut}(E)| = 2$, so the contribution is $1/(2/2) = 1$... let us state it differently. The mass formula is

$$\sum_{j \text{ s.s.}} \frac{1}{|\text{Aut}(E_j)|/2} = \frac{p-1}{12}.$$

For generic j : $|\text{Aut}| = 2$, contributing 1. For $j = 1728$: $|\text{Aut}| = 4$ (when supersingular, which happens iff $p \equiv 3 \pmod{4}$), contributing $1/2$. For $j = 0$: $|\text{Aut}| = 6$ (when supersingular, which happens iff $p \equiv 2 \pmod{3}$), contributing $1/3$.

Let n be the number of supersingular $j \neq 0, 1728$. Then:

$$n + \frac{\epsilon_0}{3} + \frac{\epsilon_{1728}}{2} = \frac{p-1}{12},$$

where $\epsilon_0 = 1$ if $j = 0$ is supersingular (iff $p \equiv 2 \pmod{3}$, i.e., $p \not\equiv 1 \pmod{3}$) and $\epsilon_{1728} = 1$ if $j = 1728$ is supersingular (iff $p \equiv 3 \pmod{4}$). The total count is $h_p = n + \epsilon_0 + \epsilon_{1728}$, which yields (5.15) after case analysis on $p \pmod{12}$.

For small primes: $h_2 = 1$ ($j = 0$ only), $h_3 = 1$ ($j = 0 = 1728$ in characteristic 3), $h_5 = 1$ ($j = 0$), $h_7 = 1$ ($j = 6 \cdot 7 = 1728 \equiv 6$), $h_{11} = 2$, $h_{13} = 1$, $h_{37} = 3$.

The Endomorphism Ring: Completing the Proof

We now complete the proof of theorem 2.6.7: $\text{End}(E)$ is a *maximal* order in $B_{p,\infty}$ for $E/\overline{\mathbb{F}_p}$ supersingular.

Theorem 5.5.3: Endomorphism Ring of Supersingular Curves — Full Proof

Let $E/\overline{\mathbb{F}_p}$ be a supersingular elliptic curve. Then:

1. $\text{End}(E) \otimes \mathbb{Q} \cong B_{p,\infty}$, the quaternion algebra over $\mathbb{Q}$ ramified exactly at p and ∞ .
2. $\text{End}(E)$ is a maximal order in $B_{p,\infty}$.

Proof:

Part (1): The quaternion algebra. By theorem 2.6.3, $D = \text{End}(E) \otimes \mathbb{Q}$ is a division algebra of dimension 1, 2, or 4 over $\mathbb{Q}$. We must show $\dim D = 4$.

By theorem 5.4.1 (applied over $\overline{\mathbb{F}_p}$, where G acts trivially):

$$\text{End}(E) \otimes \mathbb{Z}_\ell \cong \text{End}_{\mathbb{Z}_\ell}(T_\ell(E)) \cong M_2(\mathbb{Z}_\ell)$$

for every $\ell \neq p$. So $\text{rank}_{\mathbb{Z}_\ell}(\text{End}(E) \otimes \mathbb{Z}_\ell) = 4$ for all $\ell \neq p$, hence $\text{rank}_{\mathbb{Z}} \text{End}(E) = 4$ and $\dim_{\mathbb{Q}} D = 4$. Thus D is a quaternion algebra.

That $D \cong B_{p,\infty}$ (ramified at p and ∞) was proved in theorem 2.6.7: ramification at ∞ follows from the positive-definite degree form, and ramification at p from supersingularity ($E[p](\overline{\mathbb{F}_p}) = \{O\}$).

Part (2): Maximality. An order $\mathcal{O}$ in $B_{p,\infty}$ is maximal iff $\mathcal{O} \otimes \mathbb{Z}_\ell$ is a maximal order in $B_{p,\infty} \otimes \mathbb{Q}_\ell$ for every prime ℓ (including $\ell = p$).

At $\ell \neq p$: $B_{p,\infty} \otimes \mathbb{Q}_\ell \cong M_2(\mathbb{Q}_\ell)$ (since $B_{p,\infty}$ is unramified at ℓ). The maximal orders in $M_2(\mathbb{Q}_\ell)$ are conjugates of $M_2(\mathbb{Z}_\ell)$. By Tate's theorem: $\text{End}(E) \otimes \mathbb{Z}_\ell \cong M_2(\mathbb{Z}_\ell)$, which is maximal.

At $\ell = p$: $B_{p,\infty} \otimes \mathbb{Q}_p$ is the unique quaternion division algebra over $\mathbb{Q}_p$. It has a unique maximal order $\mathcal{O}_p$ (the valuation ring). We need $\text{End}(E) \otimes \mathbb{Z}_p \cong \mathcal{O}_p$.

This is where the formal group enters. By theorem 4.4.10, the formal group $\hat{E}$ over $\overline{\mathbb{F}_p}$ is the unique height-2 formal group. Its endomorphism ring (as a formal group) is $\text{End}(\hat{E}) \cong \mathcal{O}_p$, the maximal order in the quaternion division algebra $B_{p,\infty} \otimes \mathbb{Q}_p$. The natural map $\text{End}(E) \otimes \mathbb{Z}_p \rightarrow \text{End}(\hat{E})$ (restricting an endomorphism of E to the formal group) is an isomorphism: surjectivity follows because any endomorphism of $\hat{E}$ extends to an endomorphism of E (by the algebraization theorem for formal group homomorphisms over algebraically closed fields — see Tate [24]), and

injectivity follows because $\hat{E}$ captures the full p -adic information (the “connected part” of the p -divisible group is all of $E[p^\infty]$ for supersingular curves, by proposition 4.4.9).

Therefore $\text{End}(E) \otimes \mathbb{Z}_p \cong \mathcal{O}_p$, which is maximal. Combining with the $\ell \neq p$ case: $\text{End}(E)$ is maximal at every place, hence a maximal order.

This completes the proof promised in section 2.6. The two key inputs were Tate’s theorem (for the structure at $\ell \neq p$) and the formal group theory (for the structure at $\ell = p$). The maximality is a strong statement: it says the endomorphism ring is “as large as possible” within the quaternion algebra — there is no room for additional endomorphisms.

The Deuring Correspondence

The maximal orders in $B_{p,\infty}$ parametrize supersingular curves:

Theorem 5.5.4: The Deuring Correspondence

There is a bijection

$$\left\{ \begin{array}{c} \text{Supersingular elliptic curves over } \overline{\mathbb{F}}_p \\ \text{up to isomorphism} \end{array} \right\} \longleftrightarrow \left\{ \begin{array}{c} \text{Maximal orders in } B_{p,\infty} \text{ up to} \\ \text{conjugation} \end{array} \right\}. \quad (5.17)$$

The correspondence sends $E \mapsto \text{End}(E)$. Under this bijection, isogenies $\varphi : E_1 \rightarrow E_2$ correspond to left ideal classes: $\text{Hom}(E_1, E_2)$ is a left $\text{End}(E_1)$ -module, right $\text{End}(E_2)$ -module, and it is a locally free rank-1 module on each side.

Proof Key ideas:

The map $E \mapsto \text{End}(E)$ is well-defined: isomorphic curves have conjugate endomorphism rings ($\text{End}(E') = \varphi \circ \text{End}(E) \circ \varphi^{-1}$ for any isomorphism $\varphi : E \xrightarrow{\sim} E'$). By theorem 5.5.3, each $\text{End}(E)$ is a maximal order.

Surjectivity. Every maximal order $\mathcal{O}$ in $B_{p,\infty}$ arises as $\text{End}(E)$ for some supersingular E . This follows from the Deuring lifting theorem: given a maximal order $\mathcal{O}$, one constructs a CM abelian variety in characteristic 0 with endomorphism ring $\mathcal{O}$, then reduces modulo p to obtain a supersingular curve. The details require the theory of CM liftings, which we develop in ??.

Injectivity. If $\text{End}(E_1) \cong \text{End}(E_2)$ as rings (i.e., conjugate as maximal orders), then $E_1 \cong E_2$. This follows from the fact that the isomorphism of endomorphism rings, combined with Tate’s theorem, produces an isomorphism of Tate modules, which lifts to an isomorphism of curves.

The ideal correspondence. For an isogeny $\varphi : E_1 \rightarrow E_2$, the set $I_\varphi = \{\psi \in \text{Hom}(E_1, E_2) \mid \psi \text{ factors through } \varphi\} = \text{Hom}(E_1, E_2) \cdot \hat{\varphi} \cdots$. More precisely, $\text{Hom}(E_1, E_2)$ is naturally a left $\text{End}(E_1)$ -module (via pre-composition) and a right $\text{End}(E_2)$ -module (via post-composition). It is an invertible bimodule: a left ideal of $\text{End}(E_1)$ when viewed from the left, and a right ideal of $\text{End}(E_2)$ when viewed from the right. The “norm” of this ideal (as a left $\text{End}(E_1)$ -ideal) equals the degree of the isogeny. See Voight [26] for the full development of the ideal theory in quaternion orders.

The Deuring correspondence translates questions about supersingular isogenies into questions about ideals in quaternion orders, which are amenable to the powerful tools of algebraic number theory. For example, the number of isogenies of degree ℓ from E to (non-isomorphic) curves corresponds to the number of left ideals of norm ℓ in $\text{End}(E)$, which is $\ell + 1$ (for $\ell \neq p$, since $B_{p,\infty} \otimes \mathbb{Q}_\ell \cong M_2(\mathbb{Q}_\ell)$) and the number of sublattices of index ℓ in $\mathbb{Z}_\ell^2$ is $\ell + 1$).

The Supersingular Isogeny Graph

Definition 5.5.5: The Supersingular Isogeny Graph

For a prime $\ell \neq p$, the *supersingular ℓ -isogeny graph* $\mathcal{G}_\ell(p)$ is the multigraph whose vertices are the supersingular j -invariants over $\overline{\mathbb{F}_p}$, and whose edges from j_1 to j_2 are the ℓ -isogenies $E_1 \rightarrow E_2$ (up to post-composition with automorphisms of E_2), where $j(E_i) = j_i$.

This graph has approximately $p/12$ vertices. Each vertex has degree $\ell + 1$ (the number of subgroups of order ℓ in $E[\ell] \cong (\mathbb{Z}/\ell\mathbb{Z})^2$, since every such subgroup gives an ℓ -isogeny via theorem 2.1.10). The graph is connected (since all supersingular curves over $\overline{\mathbb{F}_p}$ are isogenous, forming a single isogeny class by corollary 5.4.4) and has remarkable spectral properties:

Proposition 5.5.6: Ramanujan Property

The supersingular ℓ -isogeny graph $\mathcal{G}_\ell(p)$ is a Ramanujan graph: for every non-trivial eigenvalue λ of its adjacency matrix, $|\lambda| \leq 2\sqrt{\ell}$.

Proof Key ideas:

The adjacency matrix of $\mathcal{G}_\ell(p)$ is the Hecke operator T_ℓ acting on the space of supersingular j -invariants (weighted by automorphisms). The Eichler–Shimura relation connects T_ℓ to the Hecke operator on modular forms of weight 2 for $\Gamma_0(p)$, and the Ramanujan–Petersson conjecture (proved by Deligne as a consequence of the Weil conjectures for higher-dimensional varieties) gives the bound $|\lambda| \leq 2\sqrt{\ell}$. See Pizer [16] for the details.

The Ramanujan property makes supersingular isogeny graphs excellent *expander graphs*: random walks on them converge rapidly to the uniform distribution. This has applications to cryptography (??), where the hardness of finding paths in these graphs underlies isogeny-based cryptographic protocols.

Point Counting: Schoof’s Algorithm

We close this chapter with an overview of the algorithmic problem of computing $|E(\mathbb{F}_q)|$. The naive approach (evaluating the Weierstrass equation at every $x \in \mathbb{F}_q$) runs in $O(q)$ time, which is exponential in $\log q$ (the input size). Schoof’s algorithm [18] computes $|E(\mathbb{F}_q)|$ in time polynomial in $\log q$, using the action of Frobenius on the torsion.

The idea is simple and elegant. The Frobenius φ_q acts on $E[\ell](\overline{\mathbb{F}_q}) \cong (\mathbb{Z}/\ell\mathbb{Z})^2$ as a matrix with characteristic polynomial $T^2 - a_q T + q \pmod{\ell}$. So $a_q \pmod{\ell}$ is determined by the trace of this matrix. Concretely: for a point $P \in E[\ell](\overline{\mathbb{F}_q})$, we have $\varphi_q^2(P) - [a_q]\varphi_q(P) + [q]P = O$ in $E[\ell]$, which

gives

$$\varphi_q^2(P) + [q \bmod \ell]P = [a_q \bmod \ell] \cdot \varphi_q(P) \quad \text{in } E[\ell]. \quad (5.18)$$

One computes the left-hand side (using the division polynomials ψ_ℓ to work in $E[\ell]$) and matches against $[t]\varphi_q(P)$ for $t = 0, 1, \dots, \ell - 1$ to determine $a_q \pmod{\ell}$. Doing this for enough small primes ℓ (with $\prod \ell > 4\sqrt{q}$, so that a_q is determined by the Chinese Remainder Theorem and the Hasse bound) yields a_q exactly.

Proposition 5.5.7: Complexity of Schoof's Algorithm

Schoof's algorithm computes $|E(\mathbb{F}_q)|$ in deterministic time $O((\log q)^{5+\epsilon})$ (using fast arithmetic), where the implied constant is independent of q and E .

Proof Sketch:

For each prime ℓ , the division polynomial ψ_ℓ has degree $O(\ell^2)$, and the arithmetic in $\mathbb{F}_q[x]/(\psi_\ell)$ costs $O(\ell^2 \log q)$ per operation. Evaluating $\varphi_q^2(P) + [q \bmod \ell]P$ requires $O(\log q)$ squarings in this quotient ring. Testing ℓ values of t costs $O(\ell)$ additional operations. Total per prime: $O(\ell^3(\log q)^2)$. We need $O(\log q)$ primes ℓ (each of size $O(\log q)$), giving total cost $O((\log q)^5)$, with logarithmic improvements from fast multiplication.

Improvements due to Elkies and Atkin (the SEA algorithm) bring the practical running time down significantly by distinguishing “Elkies primes” (where the characteristic polynomial of Frobenius on $E[\ell]$ splits) from “Atkin primes” (where it does not), and using modular polynomials to work with ℓ -isogenies directly rather than division polynomials. These improvements make point counting feasible for cryptographic-size primes ($\log q \approx 256$ bits).

Exercise 5.5.1

1. Compute the number of supersingular j -invariants for $p = 11$, $p = 13$, and $p = 37$, using the mass formula (5.16). For $p = 11$, list the supersingular j -invariants explicitly.
2. The quaternion algebra $B_{2,\infty} = \mathbb{Q}\langle i, j \rangle / (i^2 = -1, j^2 = -2, ij = -ji)$ has the maximal order $\mathcal{O} = \mathbb{Z}[1, i, j, (1 + ij)/2]$. Verify that $\mathcal{O}$ is closed under multiplication and that it contains no strictly larger order.
3. Draw the supersingular 2-isogeny graph $\mathcal{G}_2(11)$. It has $h_{11} = 2$ vertices (corresponding to j -invariants $j_1, j_2 \in \mathbb{F}_{121}$). Each vertex has degree 3 (since $\ell + 1 = 3$). How many edges does the graph have?
4. For $E : y^2 = x^3 - x$ over $\mathbb{F}_{101}$, determine $a_{101} \pmod{3}$ using Schoof's method. (Hint: compute $\varphi_{101}^2(P) + [101 \bmod 3]P = \varphi_{101}^2(P) + [2]P$ for $P \in E[3]$, and match against $[t]\varphi_{101}(P)$ for $t = 0, 1, 2$.)
5. Using the Hasse bound $|a_{101}| \leq 2\sqrt{101} < 21$, how many small primes ℓ are needed to determine a_{101} by CRT? (Hint: we need $\prod \ell > 2 \cdot 20 = 40$.)
6. Let $E/\overline{\mathbb{F}}_p$ be supersingular with $\text{End}(E) = \mathcal{O}$ (a maximal order in $B_{p,\infty}$). The Deuring lifting theorem says there exists an elliptic curve $\tilde{E}$ over a number field K with $\text{End}(\tilde{E}_{\bar{K}}) = \mathcal{O} \cap K'$ for a suitable imaginary quadratic field $K' \subseteq B_{p,\infty}$, and $\tilde{E}$ reduces to E modulo a prime above p . Explain why $\text{End}(\tilde{E}_{\bar{K}})$ is an order in an imaginary quadratic field (not a quaternion).

order), even though $\text{End}(E)$ is a quaternion order. (*Hint:* use the characteristic-0 classification theorem 3.5.5.)

6

Elliptic Curves over Local Fields

In section 4.5, we studied elliptic curves with *good* reduction over a local field K : the reduced curve $\tilde{E}/k$ is smooth, the kernel of reduction $E_1(K) \cong \hat{E}(\mathfrak{m})$ is captured by the formal group, and the torsion injection theorem constrains $E(K)_{\text{tors}}$ via the point count $|\tilde{E}(k)|$. But many of the most interesting elliptic curves have *bad* reduction at some primes — their Weierstrass model becomes singular upon reduction. The curve $y^2 = x^3 - x$ is smooth over $\mathbb{Q}_p$ for all odd primes, but its discriminant $\Delta = -64$ vanishes modulo 2: the reduction at $p = 2$ is a singular cubic.

This chapter develops the full local theory, extending the analysis of section 4.5 from good reduction to all reduction types. A minimal Weierstrass equation minimizes the “badness” of the reduction. The reduced curve $\tilde{E}/k$ is then either smooth (good reduction), a nodal cubic (multiplicative reduction), or a cuspidal cubic (additive reduction), and the classification from section 1.6 tells us that the nonsingular part $\tilde{E}_{\text{ns}}(k)$ is $\tilde{E}(k)$, k^* (or a non-split torus), or k^+ respectively. The exact sequence

$$0 \longrightarrow E_1(K) \longrightarrow E_0(K) \longrightarrow \tilde{E}_{\text{ns}}(k) \longrightarrow 0$$

persists in all cases, with $E_1(K) \cong \hat{E}(\mathfrak{m})$ still controlled by the formal group. The new actors are the subgroup $E_0(K)$ of points with nonsingular reduction and the component group $E(K)/E_0(K)$, whose order — the Tamagawa number c_p — is a key local invariant appearing in the Birch and Swinnerton-Dyer conjecture (??).

Conventions. Throughout this chapter, K denotes a local field: the fraction field of a complete discrete valuation ring R with maximal ideal $\mathfrak{m} = (\pi)$ and residue field $k = R/\mathfrak{m}$ of characteristic $p > 0$. We write $v : K^* \rightarrow \mathbb{Z}$ for the normalized valuation. The standard examples are $K = \mathbb{Q}_p$ (with $R = \mathbb{Z}_p$, $\pi = p$, $k = \mathbb{F}_p$) and finite extensions thereof.

6.1 Minimal Weierstrass Equations

The Weierstrass equation of an elliptic curve E/K is not unique: admissible changes of variables (??) produce infinitely many different equations, and their reductions modulo $\mathfrak{m}$ can look very different. To study the reduction systematically, we need a canonical “best” choice of integral model.

Integral Models

Definition 6.1.1: Integral Weierstrass Model

A Weierstrass equation $y^2 + a_1xy + a_3y = x^3 + a_2x^2 + a_4x + a_6$ for E/K is an *integral model* if $a_i \in R$ for all i .

Every elliptic curve E/K admits an integral model: starting from any Weierstrass equation with $a_i \in K$, apply the admissible change $(u, r, s, t) = (u, 0, 0, 0)$ with u chosen so that $u^i a_i \in R$ for each i (recall from ?? that $a'_i = u^{-i} a_i + \dots$; taking $u = \pi^{-N}$ for sufficiently large N clears all denominators).

Given an integral model, the discriminant $\Delta \in R$ and the quantities $c_4, c_6 \in R$ are integral. The reduction $\tilde{E}/k$ is defined by reducing the coefficients modulo $\mathfrak{m}$. If $\Delta \not\equiv 0 \pmod{\mathfrak{m}}$ (i.e., $v(\Delta) = 0$), the reduction is smooth: this is good reduction. If $v(\Delta) > 0$, the reduction is singular: this is bad reduction. But the value $v(\Delta)$ depends on the choice of integral model — perhaps a different model would give a smaller $v(\Delta)$?

Example 6.1.2: Different Models, Different Discriminants

Consider $E: y^2 = x^3 - x$ over $\mathbb{Q}_2$. The standard equation has $\Delta = -64$, so $v_2(\Delta) = 6$. Applying the admissible change $(u, r, s, t) = (2, 0, 0, 0)$: $x = 4x'$, $y = 8y'$, giving $64(y')^2 = 64(x')^3 - 4x'$, i.e., $y'^2 = x'^3 - x'/16$. This is not integral. Applying $(u, r, s, t) = (1/2, 0, 0, 0)$: $x = x'/4$, $y = y'/8$, giving $(y')^2/64 = (x')^3/64 - x'/4$, i.e., $(y')^2 = (x')^3 - 16x'$. Now $\Delta' = -64 \cdot 16^2 \dots$ wait, let us use the transformation law: $\Delta' = u^{-12} \Delta = (1/2)^{-12} \cdot (-64) = 4096 \cdot (-64) = -262144$, so $v_2(\Delta') = 18$. Worse!

Going in the other direction: $(u, r, s, t) = (2, 0, 0, 0)$ gives $\Delta' = 2^{-12} \cdot (-64) = -64/4096 = -1/64 \notin \mathbb{Z}_2$. The model is no longer integral.

So for this curve, $v_2(\Delta) = 6$ for the standard model, and we cannot make $v_2(\Delta)$ smaller while keeping the model integral. The standard model is, in fact, minimal.

The Minimal Model

Definition 6.1.3: Minimal Weierstrass Equation

An integral Weierstrass model for E/K is *minimal* if $v(\Delta)$ is minimal among all integral models of E/K .

Theorem 6.1.4: Existence and Uniqueness of Minimal Models

Every elliptic curve E/K has a minimal Weierstrass model. This model is unique up to admissible changes (u, r, s, t) with $u \in R^*$ and $r, s, t \in R$.

Proof:

Existence. Among all integral models of E , choose one minimizing $v(\Delta)$. This minimum exists because $v(\Delta) \geq 0$ for integral models (since $\Delta \in R$), and the values $v(\Delta)$ form a discrete subset of $\mathbb{Z}_{\geq 0}$.

We must show the minimum is achieved. If we pass from one integral model to another via an admissible change (u, r, s, t) with $a'_i \in R$, the discriminants are related by $\Delta' = u^{-12}\Delta$ (proposition 1.2.5). So $v(\Delta') = v(\Delta) - 12v(u)$. For both models to be integral with $v(\Delta') < v(\Delta)$, we need $v(u) > 0$, i.e., $u \in \mathfrak{m}$. But if $u \in \mathfrak{m}$, the transformation $a'_i = u^{-i}a_i + \dots$ produces $v(a'_i) \leq v(a_i) - i \cdot v(u) < v(a_i)$... actually, the transformation formulas for the a_i are more complex (involving r, s, t), so we need to be more careful.

The precise argument: suppose we have an integral model with $v(\Delta) > 0$. If there exists an admissible change (u, r, s, t) to another integral model with $v(\Delta') < v(\Delta)$, then $v(u) > 0$ (from $\Delta' = u^{-12}\Delta$). We claim this forces $v(u) = 1$, and hence $v(\Delta') = v(\Delta) - 12$. Iterating: after finitely many steps, we reach a model where no further reduction is possible, and this is the minimal model.

More precisely: an integral model is *not* minimal iff there exists an admissible change (u, r, s, t) with $v(u) \geq 1$ and $r, s, t \in K$ such that the new a'_i are all in R . The transformation formulas (??) show that $a'_1 = u^{-1}(a_1 + 2s)$, $a'_2 = u^{-2}(a_2 - sa_1 + 3r - s^2)$, etc. For $a'_1 \in R$ with $v(u) = 1$: we need $v(a_1 + 2s) \geq 1$, i.e., $s \equiv -a_1/2 \pmod{\mathfrak{m}}$ (if $2 \in R^*$). Similar integrality conditions on r, t arise from $a'_2, a'_3, a'_4, a'_6 \in R$. These conditions can be satisfied only when the valuations of the a_i are sufficiently large — specifically, when $v(a_1) \geq 1$, $v(a_2) \geq 1$, $v(a_3) \geq 1$, $v(a_4) \geq 2$, $v(a_6) \geq 2$ (in the case $\text{char} k \neq 2, 3$; the general case requires a more detailed case analysis). When these conditions hold, we can choose $r, s, t \in R$ to make $a'_i \in R$ with $v(u) = 1$, reducing $v(\Delta)$ by 12.

Uniqueness. If two minimal models are related by (u, r, s, t) with $r, s, t \in K$ and $u \in K^*$: both have $v(\Delta) = v(\Delta')$, so $v(u) = 0$, i.e., $u \in R^*$. The integrality of a'_i and a_i with $u \in R^*$ then forces $r, s, t \in R$ (by examining the transformation formulas order by order). So the change is an R -isomorphism.

The minimal model is the local analogue of a “reduced form” for binary quadratic forms: it is the canonical representative of the K -isomorphism class that is best adapted to the valuation.

Definition 6.1.5: The Minimal Discriminant

The *minimal discriminant* of E/K is $\Delta_{\min} = \Delta(E_{\min})$, the discriminant of any minimal Weierstrass model. Its valuation $v(\Delta_{\min})$ is a well-defined invariant of E/K :

1. $v(\Delta_{\min}) = 0$ if and only if E has good reduction over K .
2. $v(\Delta_{\min}) > 0$ if and only if E has bad reduction over K .

The Criterion for Non-Minimality**Proposition 6.1.6: Non-Minimality Criterion**

An integral Weierstrass model is *not* minimal if and only if there exists an admissible change (u, r, s, t) with $\pi \mid u$ such that the transformed equation is still integral (i.e., all $a'_i \in R$). Equivalently, the model is not minimal iff $v(\Delta) \geq 12$ and certain congruence conditions on the a_i are satisfied.

The precise congruence conditions depend on the characteristic of k and are most cleanly formulated as part of Tate's algorithm (section 6.5). For the common case $\text{char } k \neq 2, 3$ (e.g., $K = \mathbb{Q}_p$ with $p \geq 5$), we can state the criterion in terms of c_4 and c_6 :

Proposition 6.1.7: Non-Minimality for $\text{char } k \neq 2, 3$

Assume $\text{char } k \neq 2, 3$. An integral Weierstrass model in short form $y^2 = x^3 + Ax + B$ (with $A, B \in R$) is not minimal if and only if $v(A) \geq 4$ and $v(B) \geq 6$. In this case, the substitution $x = \pi^2 x'$, $y = \pi^3 y'$ (i.e., $(u, r, s, t) = (\pi, 0, 0, 0)$) gives the new model $y'^2 = x'^3 + A'x' + B'$ with $A' = A/\pi^4$ and $B' = B/\pi^6$, reducing $v(\Delta)$ by 12.

Proof:

The discriminant is $\Delta = -16(4A^3 + 27B^2)$, and $c_4 = -48A$, $c_6 = -864B$. Under $(u, 0, 0, 0)$ with $u = \pi$: $A' = \pi^{-4}A$ and $B' = \pi^{-6}B$. For integrality: $A' \in R$ iff $v(A) \geq 4$, and $B' \in R$ iff $v(B) \geq 6$. These are necessary and sufficient (since in short Weierstrass form, only A and B need to be integral, and the only admissible change preserving short Weierstrass form with $r = s = t = 0$ is $(u, 0, 0, 0)$).

When both conditions hold: $\Delta' = \pi^{-12}\Delta$, so $v(\Delta') = v(\Delta) - 12$.

Example 6.1.8: Finding Minimal Models

(a) $E : y^2 = x^3 - x$ over $\mathbb{Q}_5$. $A = -1$, $B = 0$. $v_5(A) = 0 < 4$. The model is already minimal. $v_5(\Delta) = v_5(-64) = 0$, so E has good reduction at 5.

(b) $E : y^2 = x^3 - x$ over $\mathbb{Q}_2$. The standard model has $\Delta = -64$, $v_2(\Delta) = 6 < 12$. Since $v_2(\Delta) < 12$, the model is automatically minimal (no admissible change with $v(u) \geq 1$ can produce an integral model, as this would require $v(\Delta') = v(\Delta) - 12 < 0$, contradicting integrality). So the standard

model is minimal, and $v_2(\Delta_{\min}) = 6 > 0$: bad reduction at 2.

(c) $E : y^2 = x^3 + 5^4x + 5^6$ over $\mathbb{Q}_5$. $A = 5^4$, $B = 5^6$. $v_5(A) = 4$ and $v_5(B) = 6$, so the model is not minimal. Apply $(u, 0, 0, 0) = (5, 0, 0, 0)$: $A' = 5^4/5^4 = 1$, $B' = 5^6/5^6 = 1$. New model: $y'^2 = x'^3 + x' + 1$. Now $v_5(A') = 0 < 4$, so this new model is minimal. The minimal discriminant: $\Delta_{\min} = -16(4 + 27) = -496$, $v_5(\Delta_{\min}) = 0$. Good reduction! The original model *looked* like it had bad reduction, but this was an artifact of a non-minimal equation.

(d) $E : y^2 = x^3 + x + 1$ over $\mathbb{Q}_5$. $A = 1$, $B = 1$. $v_5(A) = 0 < 4$: already minimal. $\Delta = -16(4 + 27) = -496 = -16 \cdot 31$. $v_5(\Delta) = 0$. Good reduction at 5.

(e) $E : y^2 = x^3 + 1$ over $\mathbb{Q}_3$. $A = 0$, $B = 1$. $v_3(A) = \infty \geq 4$ and $v_3(B) = 0 < 6$: the condition $v(B) \geq 6$ fails, so the model is minimal. $\Delta = -432 = -16 \cdot 27$, $v_3(\Delta) = 3$. Bad reduction at 3.

Minimal Models over Number Fields

Over a number field K , the notion of minimal model is *local*: a Weierstrass equation can be minimal at one prime and non-minimal at another. A *global minimal model* — one that is simultaneously minimal at all primes — exists over $\mathbb{Q}$ (and more generally over any number field with class number 1) but not always over number fields with nontrivial class group. The obstruction is an element of $H^1(\text{Spec } \mathcal{O}_K, \mu_2)$, related to the class group. We will not need global minimal models in this book, but mention the issue for completeness.

For $E/\mathbb{Q}$: a global minimal model always exists, and it is the “standard” equation tabulated in databases like the LMFDB. The minimal discriminant $\Delta_{\min} \in \mathbb{Z}$ is a global invariant, and its prime factorization encodes the reduction type at each prime.

Exercise 6.1.1

1. Show that for $K = \mathbb{Q}_p$ with $p \geq 5$, every elliptic curve $E/\mathbb{Q}_p$ has a minimal model in short Weierstrass form $y^2 = x^3 + Ax + B$ with $A, B \in \mathbb{Z}_p$ and either $v_p(A) < 4$ or $v_p(B) < 6$.
2. Compute the minimal discriminant $\Delta_{\min}$ and $v_p(\Delta_{\min})$ for $E : y^2 = x^3 + x^2 - 25x + 39$ over $\mathbb{Q}_p$ for $p = 2, 3, 5, 7$. (*Hint*: compute Δ from the general Weierstrass formula with $a_1 = a_3 = 0$, $a_2 = 1$, $a_4 = -25$, $a_6 = 39$, or convert to short Weierstrass form first.)
3. Over $\mathbb{Q}_2$, the curve $E : y^2 + xy + y = x^3 + x^2$ has $\Delta = -11$. Show this is a minimal model and that $v_2(\Delta) = 0$, so E has good reduction at 2. (*Remark*: this curve cannot be put in short Weierstrass form over $\mathbb{Z}_2$, since $\text{char } k = 2$.)
4. Find a non-minimal integral model for $E : y^2 = x^3 - x$ over $\mathbb{Q}_2$ by applying the admissible change $(u, r, s, t) = (1/2, 0, 0, 0)$ to the standard model. What is $v_2(\Delta)$ for the new model?
5. Verify that $\Delta' = u^{-12}\Delta$ by computing Δ for $y^2 = x^3 + Ax + B$ and Δ' for $y'^2 = x'^3 + u^{-4}Ax' + u^{-6}B$ directly from the formula $\Delta = -16(4A^3 + 27B^2)$.
6. For $E : y^2 = x^3 - x$ over $\mathbb{Q}$, compute $\Delta = -64$ and verify that the standard model is globally minimal: show it is minimal at every prime p by checking either $v_p(\Delta) = 0$ (good reduction) or $v_p(\Delta) < 12$ (automatic minimality).

6.2 Reduction Types

Given a minimal Weierstrass model for E/K , we reduce modulo $\mathfrak{m}$ to obtain a Weierstrass cubic $\tilde{E}/k$. This cubic is either smooth (an elliptic curve over k) or singular (a nodal or cuspidal cubic). The three cases correspond to the three *reduction types*, and the classification from ?? tells us the group structure of the nonsingular locus in each case.

The Three Reduction Types

Definition 6.2.1: Reduction Types

Let E/K be an elliptic curve with minimal Weierstrass model, and let $\tilde{E}/k$ be the reduction modulo $\mathfrak{m}$.

1. *Good reduction*: $\tilde{E}$ is smooth over k (equivalently, $v(\Delta_{\min}) = 0$). The reduction $\tilde{E}/k$ is an elliptic curve.
2. *Multiplicative reduction*: $\tilde{E}$ has a node (equivalently, $v(\Delta_{\min}) > 0$ and $v(c_4) = 0$). The nonsingular locus $\tilde{E}_{\text{ns}}$ is an algebraic torus.
3. *Additive reduction*: $\tilde{E}$ has a cusp (equivalently, $v(\Delta_{\min}) > 0$ and $v(c_4) > 0$). The nonsingular locus $\tilde{E}_{\text{ns}}$ is isomorphic to $\mathbb{G}_a$.

Multiplicative reduction is further classified as *split* or *non-split* according to whether the tangent lines at the node are defined over k or over a quadratic extension k'/k .

The criterion follows from proposition 1.6.2: the reduced curve $\tilde{E}$ (with $\tilde{\Delta} = 0$ in cases (2) and (3)) has a node when $\tilde{c}_4 \neq 0$ and a cusp when $\tilde{c}_4 = 0$. Since $\tilde{c}_4 = c_4 \bmod \mathfrak{m}$, the condition $\tilde{c}_4 \neq 0$ is $v(c_4) = 0$.

By theorem 1.6.5, the group structure of the nonsingular locus is:

Type	Singularity	$\tilde{E}_{\text{ns}}(\bar{k})$	$\tilde{E}_{\text{ns}}(k)$
Good	none	$\tilde{E}(\bar{k})$	$\tilde{E}(k)$ (elliptic curve)
Split multiplicative	node, split	$\bar{k}^*$	$k^* \cong \mathbb{G}_m(k)$
Non-split multiplicative	node, non-split	$\bar{k}^*$	$\ker(\text{Nm} : (k')^* \rightarrow k^*) \cong k'/k^*$
Additive	cusp	$\bar{k}$	$k \cong \mathbb{G}_a(k)$

In the non-split multiplicative case, k'/k is the unique quadratic extension and $\tilde{E}_{\text{ns}}(k)$ is the norm-1 subgroup of $(k')^*$, which has order $q + 1$ (where $q = |k|$). In the split multiplicative case, $|\tilde{E}_{\text{ns}}(k)| = q - 1$. In the additive case, $|\tilde{E}_{\text{ns}}(k)| = q$.

Detecting the Reduction Type: Split vs. Non-Split

For multiplicative reduction ($v(\Delta) > 0$, $v(c_4) = 0$), we need to determine whether the node is split or non-split over k . In the short Weierstrass case ($\text{char } k \neq 2, 3$): the reduced equation is $\tilde{y}^2 = \tilde{x}^3 + \tilde{A}\tilde{x} + \tilde{B}$

with $\tilde{\Delta} = 0$ and $\tilde{A} \neq 0$. The singular point is at the double root x_0 of $\tilde{x}^3 + \tilde{A}\tilde{x} + \tilde{B} = 0$, and the tangent cone involves $\tilde{y}^2 = \alpha(\tilde{x} - x_0)^2$ for some $\alpha \in k^*$ (derived from the factored cubic, as in ??).

Proposition 6.2.2: Split vs. Non-Split Criterion

Let E/K have multiplicative reduction at $\mathfrak{m}$. The reduction is:

1. *split* if $-c_6 \bmod \mathfrak{m}$ is a square in k^* ;
2. *non-split* if $-c_6 \bmod \mathfrak{m}$ is a non-square in k^* .

Equivalently, the node is split iff the quadratic twist by $-c_6$ is trivial over k .

Proof:

After translating the singular point to the origin, the reduced equation takes the form $\tilde{y}^2 = \alpha\tilde{x}^2 + (\text{higher order})$. The tangent lines are $\tilde{y} = \pm\sqrt{\alpha}\tilde{x}$, which are defined over k iff α is a square in k^* . Tracking the admissible changes through the formulas of ?? shows that α differs from $-c_6 \bmod \mathfrak{m}$ by a square factor, so the quadratic residue symbol depends only on the class of $-c_6$ in $k^*/(k^*)^2$.

Example 6.2.3: Reduction Types of Running Curves

(a) $E : y^2 = x^3 - x$ over $\mathbb{Q}_2$. Minimal model (example 6.1.8(b)): $y^2 = x^3 - x$, $\Delta = -64$, $c_4 = 48$. Over $\mathbb{Q}_2$: $v_2(\Delta) = 6 > 0$ and $v_2(c_4) = v_2(48) = 4 > 0$. So $v(c_4) > 0$: *additive reduction* at 2. The reduced curve over $\mathbb{F}_2$ has $\tilde{c}_4 = 0$ and is cuspidal.

(b) $E : y^2 = x^3 + 1$ over $\mathbb{Q}_3$. $\Delta = -432 = -16 \cdot 27$, $c_4 = 0$, $c_6 = -864 = -32 \cdot 27$. $v_3(\Delta) = 3 > 0$ and $v_3(c_4) = \infty > 0$: *additive reduction* at 3.

(c) *Multiplicative reduction*: $E : y^2 = x^3 - 43x + 166$ over $\mathbb{Q}_7$. One computes $c_4 = -48 \cdot (-43) = 2064 = 2^4 \cdot 3 \cdot 43$ and $\Delta = -16(4(-43)^3 + 27(166)^2) = -16 \cdot (425228)$. To find $v_7(\Delta)$: note $4 \cdot (-43)^3 = -318028$ and $27 \cdot 166^2 = 743256$, so $4A^3 + 27B^2 = 425228 = 4 \cdot 7^2 \cdot 2173 + \dots$. (A direct computation gives $425228 = 7^2 \cdot 8678 + \text{remainder}$... one verifies $425228/49 = 8678.1$, so $49 \nmid 425228$. Checking: $425228/7 = 60746.857\dots$, so in fact $7 \nmid 425228$, giving $v_7(\Delta) = 0$: good reduction.)

For a genuine multiplicative example, consider $E : y^2 = x^3 + x^2 - 7$ over $\mathbb{Q}_7$ (general Weierstrass: $a_2 = 1$, $a_6 = -7$, others zero). The discriminant is $\Delta = -16(4 \cdot 1^3 \cdot (-7) + 27 \cdot (-7)^2 + \dots)$... for the general form, $b_2 = 4$, $b_4 = 0$, $b_6 = -28$, $b_8 = -7$, $c_4 = 16$, and $\Delta = -16b_8 - \dots = 7 \cdot (\text{unit})$ — the precise computation requires care. Instead, let us use a well-known example.

Consider $E : y^2 + y = x^3 - x$ over $\mathbb{Q}_{37}$ (Cremona label 37a1). This curve has $\Delta = -37$ and $c_4 = 48$. Over $\mathbb{Q}_{37}$: $v_{37}(\Delta) = 1 > 0$ and $v_{37}(c_4) = 0$: *multiplicative reduction* at 37. To determine split vs. non-split: $c_6 = -6048$. We need $-c_6 = 6048 \pmod{37}$. Since $6048 = 163 \cdot 37 + 17$: $-c_6 \equiv 17 \pmod{37}$. Is 17 a square mod 37? We compute $17^{(37-1)/2} = 17^{18} \pmod{37}$. By successive squaring: $17^2 = 289 \equiv 289 - 7 \cdot 37 = 30$; $17^4 \equiv 900 \equiv 900 - 24 \cdot 37 = 12$; $17^8 \equiv 144 \equiv 144 - 3 \cdot 37 = 33$; $17^{16} \equiv 33^2 = 1089 \equiv 1089 - 29 \cdot 37 = 16$; $17^{18} = 17^{16} \cdot 17^2 \equiv 16 \cdot 30 = 480 \equiv 480 - 12 \cdot 37 = 36 \equiv -1$.

Since the Legendre symbol is -1 : 17 is a non-square, so the reduction is *non-split multiplicative*... let us double-check via a different route. The quadratic residues mod 37 include 1, 4, 9, 16, 25, 36, 10, $\dots$ (one can list all 18 squares). Since 17 is not among them: confirmed, *non-split multiplicative* at 37.

The Exact Sequence for All Reduction Types

The subgroups $E_0(K)$ and $E_1(K)$ were defined in ?? for good reduction. The definitions extend to all reduction types without change:

$$\begin{aligned} E_0(K) &= \{P \in E(K) \mid \tilde{P} \in \tilde{E}_{\text{ns}}(k)\}, \\ E_1(K) &= \{P \in E(K) \mid \tilde{P} = \tilde{O}\} = \ker(\text{red}). \end{aligned}$$

For good reduction, $E_0(K) = E(K)$ (since $\tilde{E}_{\text{ns}} = \tilde{E}$). For bad reduction, $E_0(K)$ is a proper subgroup of $E(K)$: the points whose reduction is the singular point are *not* in $E_0(K)$.

Theorem 6.2.4: The Exact Sequence

Let E/K have a minimal Weierstrass model. Then:

1. $E_0(K)$ is a subgroup of finite index in $E(K)$.
2. There is an exact sequence

$$0 \longrightarrow E_1(K) \longrightarrow E_0(K) \xrightarrow{\text{red}} \tilde{E}_{\text{ns}}(k) \longrightarrow 0. \quad (6.1)$$

3. The kernel $E_1(K) \cong \hat{E}(\mathfrak{m})$ via $P \mapsto z(P) = -x(P)/y(P)$, as in theorem 4.5.2.

Proof:

(1) We show $[E(K) : E_0(K)] < \infty$ directly. Let $P \in E(K) \setminus E_0(K)$, so $\tilde{P} = P_{\text{sing}}$ (the singular point of $\tilde{E}$). The point P has affine coordinates (x_0, y_0) with $v(x_0) \geq 0$ and $v(y_0) \geq 0$ (one verifies: if $v(x_0) < 0$, then $\tilde{P} = \tilde{O}$, not P_{sing}). Consider $[2]P$: the doubling formula involves x_0, y_0 and the a_i , and for the singular reduction cases one shows (by explicit computation of the addition law near P_{sing}) that $[m]P \in E_0(K)$ for some bounded m depending only on the Kodaira type. For multiplicative reduction (type I_n): the points reducing to P_{sing} form at most n cosets of $E_0(K)$ (one for each component of the special fiber, since a point reducing to the singular point lies “between” two components, and the component it is assigned to determines the coset). For additive reduction: a more delicate analysis using the structure of the special fiber shows $[E(K) : E_0(K)] \leq 4$ (the bound comes from the possible component groups of the additive Kodaira types, all of which have $|\Phi| \leq 4$; see the table in section 6.5).

(2) The exactness at $E_0(K)$: the kernel of $\text{red} : E_0(K) \rightarrow \tilde{E}_{\text{ns}}(k)$ is $E_1(K)$ by definition. Surjectivity: a nonsingular point $\tilde{Q} \in \tilde{E}_{\text{ns}}(k)$ lifts to $Q \in E(K)$ by Hensel’s lemma (theorem 4.1.6), exactly as in the good reduction case. The key point is that Hensel’s lemma requires only the *smoothness* of $\tilde{E}$ at $\tilde{Q}$, not the smoothness of $\tilde{E}$ as a whole. Since $\tilde{Q} \in \tilde{E}_{\text{ns}}(k)$ is by definition a smooth point, the lifting works.

(3) The isomorphism $E_1(K) \cong \hat{E}(\mathfrak{m})$ is proved exactly as in theorem 4.5.2. The formal group $\hat{E}$

depends on the Weierstrass equation (not on the smoothness of the reduction), and the local parameter $z = -x/y$ at O is defined regardless of the reduction type. The convergence of the power series $w(z)$ in $\mathfrak{m}$ and the bijectivity of $P \mapsto z(P)$ are unaffected by the reduction type, since they concern only the local structure at O — which lies on the smooth part of the special fiber (the identity $\tilde{O}$ is always a nonsingular point of $\tilde{E}$).

The exact sequence (6.1) is the workhorse of the local theory. It decomposes $E_0(K)$ into a “formal” piece $E_1(K) \cong \hat{E}(\mathfrak{m})$ (controlled by the formal group, hence by p -adic analysis) and a “finite field” piece $\tilde{E}_{\text{ns}}(k)$ (controlled by the reduction type). For good reduction, $\tilde{E}_{\text{ns}}(k) = \tilde{E}(k)$ is an elliptic curve over a finite field (studied in Chapter 5). For multiplicative reduction, $\tilde{E}_{\text{ns}}(k) \cong k^*$ or a twisted torus. For additive reduction, $\tilde{E}_{\text{ns}}(k) \cong k^+$.

The Component Group

Definition 6.2.5: The Component Group

The *component group* of E/K (at $\mathfrak{m}$) is

$$\Phi = \Phi_{\mathfrak{m}} = E(K)/E_0(K). \quad (6.2)$$

It is a finite abelian group (theorem 6.2.4(1)). Its order $c_{\mathfrak{m}} = |\Phi|$ is the *Tamagawa number* of E at $\mathfrak{m}$.

The terminology comes from the theory of Néron models (??): the Néron model $\mathcal{E}/R$ is a smooth group scheme whose generic fiber is E/K , and the component group Φ is literally the group of connected components of the special fiber $\mathcal{E}_k$. The subgroup $E_0(K)$ corresponds to the identity component $\mathcal{E}_k^0$.

For good reduction: $E_0(K) = E(K)$, so $\Phi = 0$ and $c_{\mathfrak{m}} = 1$.

For multiplicative reduction of type I_n (in the Kodaira classification, section 6.5): the component group is cyclic, with $\Phi \cong \mathbb{Z}/n\mathbb{Z}$ and $c_{\mathfrak{m}} = n$ in the split case. For non-split multiplicative I_n : $\Phi \cong \mathbb{Z}/\gcd(n, 2)\mathbb{Z}$, giving $c = 1$ (odd n) or $c = 2$ (even n). See section 6.6 for the full treatment.

For additive reduction: $|\Phi| \leq 4$ (the exact value depends on the Kodaira type). The explicit values are tabulated in section 6.5.

Combining the Pieces

The full local structure of $E(K)$ is now visible:

$$\begin{array}{ccccccc}
& & 0 & & 0 & & \\
& & \downarrow & & \downarrow & & \\
0 & \rightarrow & E_1(K) & \rightarrow & E_0(K) & \rightarrow & \tilde{E}_{\text{ns}}(k) \rightarrow 0 \\
& & \parallel & & \downarrow & & \\
& & \hat{E}(\mathfrak{m}) & & E(K) & & \\
& & & & \downarrow & & \\
& & & & \Phi & & \\
& & & & \downarrow & & \\
& & & & 0 & &
\end{array} \tag{6.3}$$

Reading from bottom to top: $E(K)$ is an extension of the finite group Φ by $E_0(K)$, which is itself an extension of $\tilde{E}_{\text{ns}}(k)$ by the pro- p group $\hat{E}(\mathfrak{m})$. Over $\mathbb{Q}_p$: $\hat{E}(\mathfrak{m}) \cong \mathbb{Z}_p^d \oplus (\text{finite})$ (proposition 4.5.9), $\tilde{E}_{\text{ns}}(k)$ is a finite group of order roughly p , and Φ is a finite group of order c_p . The “size” of $E(K)$ is thus controlled by p -adic analysis (via the formal group) at the top and by finite-field arithmetic (via the reduction) at the bottom.

Example 6.2.6: Local Structure of $y^2 = x^3 - x$ at $p = 2$

Over $\mathbb{Q}_2$: additive reduction ($v_2(\Delta) = 6$, $v_2(c_4) = 4$). The nonsingular locus $\tilde{E}_{\text{ns}}(\mathbb{F}_2) \cong \mathbb{F}_2^+ = \{0, 1\}$ has order 2. The formal group: $\hat{E}(2\mathbb{Z}_2)$, a pro-2 group. The exact sequence gives $|E_0(\mathbb{Q}_2)/E_1(\mathbb{Q}_2)| = |\tilde{E}_{\text{ns}}(\mathbb{F}_2)| = 2$. The component group Φ and the Tamagawa number c_2 are determined by Tate’s algorithm (section 6.5).

Exercise 6.2.1

1. For each of the following curves over $\mathbb{Q}_p$, determine the reduction type (good, split multiplicative, non-split multiplicative, or additive): (a) $y^2 = x^3 + x + 1$ at $p = 5$; (b) $y^2 = x^3 - 11x + 14$ at $p = 7$; (c) $y^2 = x^3 + 3$ at $p = 3$. (*Hint*: compute Δ and c_4 ; check minimality; then apply the criteria.)
2. Explain why the Hasse bound does not apply to $|\tilde{E}_{\text{ns}}(k)|$ for bad reduction. (*Hint*: the Hasse bound applies to smooth projective curves of genus 1; a singular Weierstrass cubic has arithmetic genus 0.)
3. Construct an explicit elliptic curve over $\mathbb{Q}_p$ ($p \geq 5$) with split multiplicative reduction. (*Hint*: start with a curve whose reduction modulo p is a nodal cubic with rational tangent slopes. The curve $y^2 = x^3 + x^2 - p$ over $\mathbb{Q}_p$ has $a_2 = 1$ and $a_6 = -p$; compute its reduction type.)
4. Explain why $E_1(K) \cong \hat{E}(\mathfrak{m})$ holds even for bad reduction (i.e., why the proof of theorem 4.5.2 from section 4.5 does not require good reduction). (*Hint*: the formal group is defined by the local structure at O , which is always a smooth point of $\tilde{E}$.)
5. Show that for multiplicative reduction of type I_n : $\Phi \cong \mathbb{Z}/n\mathbb{Z}$ (split case) or $\Phi \cong \mathbb{Z}/\gcd(n, 2)\mathbb{Z}$ (non-split case). (*Hint*: the component group is the group of connected components of the special fiber of the Néron model; for I_n , the special fiber is a cycle of n rational curves, and Frobenius either preserves or reverses the cycle.)
6. An elliptic curve has *semistable* reduction if its reduction is either good or multiplicative (i.e., not additive). Show that E has semistable reduction iff $v(c_4) = 0$ or $v(\Delta) = 0$. Equivalently:

E has additive reduction iff $v(\Delta) > 0$ and $v(c_4) > 0$.

6.3 Points of Finite Order

The torsion injection theorem (theorem 4.5.6) showed that for good reduction, the prime-to- p torsion of $E(K)$ injects into $\tilde{E}(k)$. We now extend this analysis to all reduction types, using the exact sequence of theorem 6.2.4. The formal group contributes no prime-to- p torsion regardless of the reduction type; the new feature is that for bad reduction, the injection lands in $\tilde{E}_{\text{ns}}(k)$ rather than in an elliptic curve, and the quotient $E(K)/E_0(K)$ can contribute additional torsion.

Torsion in the Formal Group: All Reduction Types

The key input is unchanged from theorem 4.5.5: the formal group $\hat{E}(\mathfrak{m})$ has no prime-to- p torsion, regardless of the reduction type. Let us state this explicitly in the present context.

Proposition 6.3.1: No Prime-to- p Torsion in $E_1(K)$

Let E/K have a minimal Weierstrass model (with any reduction type). Then for $\gcd(n, p) = 1$:

$$E_1(K)[n] = \{O\}.$$

Proof:

By theorem 6.2.4(3), $E_1(K) \cong \hat{E}(\mathfrak{m})$. By theorem 4.5.5(1), $\hat{E}(\mathfrak{m})[n] = \{0\}$ for $\gcd(n, p) = 1$. (The proof depended only on the fact that $[n]_{\hat{E}}'(0) = n \in R^*$ when $\gcd(n, p) = 1$, making $[n]_{\hat{E}}$ an isomorphism on $\mathfrak{m}$ — and this holds for any formal group law, regardless of the reduction type of E .)

Torsion Injection for All Reduction Types

Theorem 6.3.2: Generalized Torsion Injection

Let E/K have a minimal Weierstrass model. For $\gcd(n, p) = 1$:

$$E_0(K)[n] \hookrightarrow \tilde{E}_{\text{ns}}(k). \quad (6.4)$$

More precisely, the reduction map $\text{red} : E_0(K) \rightarrow \tilde{E}_{\text{ns}}(k)$ restricts to an injection on the n -torsion subgroup.

Proof:

If $P \in E_0(K)[n]$ with $\tilde{P} = \tilde{O}$, then $P \in E_0(K)[n] \cap E_1(K) = E_1(K)[n] = \{O\}$ by proposition 6.3.1. So the kernel of $\text{red}|_{E_0(K)[n]}$ is trivial.

For good reduction, $E_0(K) = E(K)$ and $\tilde{E}_{\text{ns}}(k) = \tilde{E}(k)$, recovering theorem 4.5.6. For bad reduction, the injection lands in $\tilde{E}_{\text{ns}}(k)$, which is a different group depending on the reduction type.

Consequences by Reduction Type

The injection (6.4) gives different constraints depending on the structure of $\tilde{E}_{\text{ns}}(k)$.

Proposition 6.3.3: Torsion Bounds by Reduction Type

Let E/K have a minimal model, and let $\gcd(n, p) = 1$. Then:

1. (*Good reduction.*) $E(K)[n] \hookrightarrow \tilde{E}(k)$. The image is a subgroup of the elliptic curve $\tilde{E}(k)$, of order dividing $|\tilde{E}(k)| = q + 1 - a_q$.
2. (*Split multiplicative reduction.*) $E_0(K)[n] \hookrightarrow k^* \cong \mathbb{G}_m(k)$. Since k^* is cyclic of order $q - 1$: the prime-to- p torsion in $E_0(K)$ is cyclic, of order dividing $q - 1$.
3. (*Non-split multiplicative reduction.*) $E_0(K)[n] \hookrightarrow \ker(\text{Nm} : (k')^* \rightarrow k^*)$, a cyclic group of order $q + 1$ (where k'/k is the quadratic extension).
4. (*Additive reduction.*) $E_0(K)[n] \hookrightarrow k^+ \cong \mathbb{G}_a(k)$. Since k^+ has exponent p (every element satisfies $px = 0$), and $\gcd(n, p) = 1$: $E_0(K)[n] = \{O\}$. The prime-to- p part of $E_0(K)$ is trivial.

Proof:

Each case follows from theorem 6.3.2 and the group structure of $\tilde{E}_{\text{ns}}(k)$ (theorem 1.6.5 and the table in section 6.2).

For (4): $k^+ = (\mathbb{F}_q, +)$ has exponent p (since $p \cdot x = 0$ for all $x \in \mathbb{F}_q$). Any element of $E_0(K)[n]$ maps to an element of k^+ killed by n , hence killed by $\gcd(n, p) = 1$ in k^+ . But the only element of k^+ killed by 1 is 0. So the image is trivial, and the injection forces $E_0(K)[n] = \{O\}$.

Case (4) is striking: for additive reduction, $E_0(K)$ has *no* prime-to- p torsion at all. Any prime-to- p torsion in $E(K)$ must come from the component group $\Phi = E(K)/E_0(K)$, which has order $c_{\mathfrak{p}} \leq 4$ for additive reduction.

The Full Torsion of $E(K)$

To pass from $E_0(K)$ to $E(K)$, we use the component group $\Phi = E(K)/E_0(K)$.

Proposition 6.3.4: Torsion of $E(K)$ via the Component Group

Let E/K have a minimal model. For any positive integer n with $\gcd(n, p) = 1$:

$$0 \longrightarrow E_0(K)[n] \longrightarrow E(K)[n] \longrightarrow \Phi[n], \quad (6.5)$$

where the last map sends a torsion point P to its class in Φ . In particular:

$$|E(K)[n]| \leq |E_0(K)[n]| \cdot |\Phi[n]| \leq |\tilde{E}_{\text{ns}}(k)| \cdot c_{\mathfrak{p}}.$$

Proof:

The sequence $0 \rightarrow E_0(K) \rightarrow E(K) \rightarrow \Phi \rightarrow 0$ restricts to the sequence of n -torsion subgroups. Exactness at $E_0(K)[n]$ and $E(K)[n]$ is standard (the “snake lemma” for torsion in short exact sequences). The last map need not be surjective: not every element of $\Phi[n]$ need lift to an n -torsion point.

For good reduction: $\Phi = 0$, so $E(K)[n] = E_0(K)[n] \hookrightarrow \tilde{E}(k)$. For multiplicative or additive reduction: the component group can contribute torsion, but $|\Phi|$ is bounded (by $v(\Delta)$ for multiplicative, by 4 for additive), so the total torsion remains controlled.

Bounding Torsion over Number Fields: The General Case

Over a number field F , we combine the local information at several primes to bound $E(F)_{\text{tors}}$. The key observation is that $E(F) \hookrightarrow E(F_{\mathfrak{p}})$ for every completion $F_{\mathfrak{p}}$, so torsion in $E(F)$ is constrained by the local torsion at every prime.

Theorem 6.3.5: Bounding Torsion via Multiple Primes

Let E/F be an elliptic curve over a number field. For each prime $\mathfrak{p}$ of F (with residue characteristic $p_{\mathfrak{p}}$ and residue field $k_{\mathfrak{p}}$):

1. If E has good reduction at $\mathfrak{p}$: $E(F)[\ell^{\infty}] \hookrightarrow \tilde{E}(k_{\mathfrak{p}})$ for every prime $\ell \neq p_{\mathfrak{p}}$.
2. If E has multiplicative reduction at $\mathfrak{p}$: $E(F)[\ell^{\infty}]$ has order dividing $c_{\mathfrak{p}} \cdot (|k_{\mathfrak{p}}| - 1)$ (split) or $c_{\mathfrak{p}} \cdot (|k_{\mathfrak{p}}| + 1)$ (non-split), for $\ell \neq p_{\mathfrak{p}}$.
3. If E has additive reduction at $\mathfrak{p}$: $|E(F)[\ell^{\infty}]|$ divides $c_{\mathfrak{p}}$ for $\ell \neq p_{\mathfrak{p}}$.

Taking the intersection over several primes:

$$|E(F)_{\text{tors}}| \mid \gcd_{\mathfrak{p}} (c_{\mathfrak{p}} \cdot |\tilde{E}_{\text{ns}}(k_{\mathfrak{p}})|), \quad (6.6)$$

where the gcd is over primes $\mathfrak{p}$ of good or multiplicative reduction, after removing the $p_{\mathfrak{p}}$ -parts.

Proof:

For each $\mathfrak{p}$: $E(F) \hookrightarrow E(F_{\mathfrak{p}})$ (by completing at $\mathfrak{p}$), and proposition 6.3.3 bounds the ℓ -part of $E(F_{\mathfrak{p}})_{\text{tors}}$ for $\ell \neq p_{\mathfrak{p}}$. Taking the gcd of these bounds over multiple primes with different residue characteristics yields (6.6).

In practice, choosing two primes of good reduction with different residue characteristics often determines $E(F)_{\text{tors}}$ completely (as we saw in example 4.5.8).

Example 6.3.6: Torsion Bounds Using Bad Reduction Primes

(a) $E : y^2 + y = x^3 - x$ over $\mathbb{Q}$ (Cremona label 37a1). This curve has conductor 37, with $\Delta = -37$ and $c_4 = 48$. Reduction types: good at all $p \neq 37$; at $p = 37$: $v_{37}(\Delta) = 1 > 0$ and $v_{37}(c_4) = 0$, so multiplicative. Is it split? $c_6 = -6912$, $-c_6 = 6912 = 2^8 \cdot 27$. Modulo 37: $6912 = 37 \cdot 186 + 30$, so $-c_6 \equiv 30 \pmod{37}$. Is 30 a square mod 37? $30 \equiv -7 \pmod{37}$, and $\left(\frac{-7}{37}\right) = \left(\frac{-1}{37}\right) \left(\frac{7}{37}\right) = (1)(\dots) = \dots$. Since $37 \equiv 1 \pmod{4}$, -1 is a QR. And $7^{(37-1)/2} = 7^{18} \pmod{37}$: one computes $7^2 = 49 \equiv 12$, $7^4 \equiv 144 \equiv 33 \equiv -4$, $7^8 \equiv 16$, $7^{16} \equiv 256 \equiv 256 - 6 \cdot 37 = 256 - 222 = 34 \equiv -3$, $7^{18} = 7^{16} \cdot 7^2 \equiv (-3)(12) = -36 \equiv 1$. So 7 is a QR, hence -7 is a QR, and the reduction is split multiplicative. Tamagawa number: $c_{37} = v_{37}(\Delta) = 1$.

At $p = 5$: good reduction. $|E(\mathbb{F}_5)|$: evaluate $x^3 - x$ for $x = 0, \dots, 4$, then use $y^2 + y = f(x)$, i.e., $y^2 + y - f(x) = 0$, which has solutions iff $1 + 4f(x)$ is a square. Counting gives $|E(\mathbb{F}_5)| = 4$. So $E(\mathbb{Q})[n] \hookrightarrow E(\mathbb{F}_5)$ for n coprime to 5: $|E(\mathbb{Q})_{\text{tors}}|$ divides 4 (away from 5). At $p = 3$: good reduction. $|E(\mathbb{F}_3)| = 4$ similarly. So $|E(\mathbb{Q})_{\text{tors}}|$ divides $\gcd(4, 4) = 4$ (away from 3 and 5). In fact, $E(\mathbb{Q})_{\text{tors}} \cong \mathbb{Z}/1\mathbb{Z}$ (trivial): this curve has rank 1 and no torsion.

(b) $E : y^2 = x^3 - x$ over $\mathbb{Q}$. Good reduction at all $p \neq 2$. At $p = 2$: additive reduction (example 6.2.3(a)). So at $p = 2$, the prime-to-2 torsion in $E_0(\mathbb{Q}_2)$ is trivial, and $|E(\mathbb{Q}_2)[\ell^\infty]| \leq c_2$ for odd ℓ . In fact $c_2 = 4$ (from Tate's algorithm). The additive prime contributes the bound: $|E(\mathbb{Q})[\ell^\infty]|$ divides 4 for every odd ℓ , and divides $|E(\mathbb{F}_p)|$ at each good prime p . Combining with $|E(\mathbb{F}_3)| = 4$ and $|E(\mathbb{F}_5)| = 8$: $|E(\mathbb{Q})_{\text{tors}}|$ divides $\gcd(4, 8) = 4$ (away from 2). This is sharp: $E(\mathbb{Q})_{\text{tors}} \cong (\mathbb{Z}/2\mathbb{Z})^2$.

The p -Torsion

The torsion injection theorem is silent about p -torsion (where $p = \text{char } k$). For the formal group: $\hat{E}(\mathfrak{m})$ can have nontrivial p -torsion (governed by the height, as in theorem 4.5.5(2)–(3)). For $\tilde{E}_{\text{ns}}(k)$: the p -part depends on the reduction type ($\tilde{E}(k)$ can have p -torsion for good reduction; k^* has p -part trivial; k^+ is all p -torsion).

We collect the key facts:

Proposition 6.3.7: p -Torsion by Reduction Type

Let E/K with $\text{char } K = p$ and K of characteristic zero (mixed characteristic).

1. (*Good ordinary reduction.*) $E(K)[p] \hookrightarrow \tilde{E}(k)[p] \cong \mathbb{Z}/p\mathbb{Z}$ via the reduction map. The formal group $\hat{E}(\mathfrak{m})$ has trivial p -torsion over K (but nontrivial p -torsion over $\bar{K}$, of order p).
2. (*Good supersingular reduction.*) $\tilde{E}(k)[p] = \{O\}$, so the reduction map gives no p -torsion. The formal group $\hat{E}(\bar{\mathfrak{m}})$ has p -torsion of order p^2 (all of $E[p]$ lies in the formal group, by proposition 4.4.9).
3. (*Multiplicative reduction.*) $\tilde{E}_{\text{ns}}(k) \cong k^*$ or a torus, with p -part trivial (since $|k^*| = q - 1$ is coprime to p). So $E_0(K)[p] = E_1(K)[p] = \hat{E}(\mathfrak{m})[p]$.
4. (*Additive reduction.*) $\tilde{E}_{\text{ns}}(k) \cong k^+$, which is entirely p -torsion. The exact sequence gives $E_0(K)[p]$ as an extension of $k^+[p] = k^+$ by $\hat{E}(\mathfrak{m})[p]$.

Proof Key ideas:

(1) For ordinary reduction: $\tilde{E}[p] \cong \mathbb{Z}/p\mathbb{Z}$ (proposition 2.3.8). The reduction map $E(K)[p] \rightarrow \tilde{E}(k)[p]$ has kernel $E_1(K)[p] = \hat{E}(\mathfrak{m})[p]$. Over K (not $\bar{K}$): $\hat{E}(\mathfrak{m})[p] = \{0\}$ because the p -torsion points of $\hat{E}$ lie in $\bar{\mathfrak{m}} \setminus \mathfrak{m}$ (they are defined over extensions of K). So the reduction map is injective on $E(K)[p]$.

(3) Since $|k^*| = q - 1$ and $p \nmid q$, we have $\gcd(|k^*|, p) = 1$, so k^* has no p -torsion. Any p -torsion in $E_0(K)$ must come from $E_1(K) = \hat{E}(\mathfrak{m})$.

Exercise 6.3.1

1. Let $E/\mathbb{Q}$ have good reduction at $p = 5$ with $|\tilde{E}(\mathbb{F}_5)| = 7$. What can you conclude about $E(\mathbb{Q})_{\text{tors}}$? (*Hint: 7 is prime.*)
2. Let $E/\mathbb{Q}_7$ have split multiplicative reduction with $c_7 = 3$. Show that $E(\mathbb{Q}_7)[n]$ is cyclic of order dividing $3(7 - 1) = 18$ for $\gcd(n, 7) = 1$.
3. Let $E/\mathbb{Q}_p$ have additive reduction with $c_p = 2$. Show that $E(\mathbb{Q}_p)$ has no prime-to- p torsion of odd order greater than 1.
4. For $E: y^2 = x^3 - x$ over $\mathbb{Q}_5$ (good ordinary reduction, $a_5 = -2$), show that $E(\mathbb{Q}_5)[5] = \{O\}$. (*Hint: $\tilde{E}(\mathbb{F}_5)[5] = \{O\}$ since $|\tilde{E}(\mathbb{F}_5)| = 8$ and $\gcd(8, 5) = 1$.*)
5. Mazur's theorem states that $E(\mathbb{Q})_{\text{tors}}$ is isomorphic to one of $\mathbb{Z}/n\mathbb{Z}$ for $n \leq 10$ or $n = 12$, or $\mathbb{Z}/2\mathbb{Z} \times \mathbb{Z}/2n\mathbb{Z}$ for $n \leq 4$. Using reduction at two primes of good reduction, show that $E(\mathbb{Q})_{\text{tors}}$ injects into $\tilde{E}(\mathbb{F}_{p_1}) \times \tilde{E}(\mathbb{F}_{p_2})$ for primes $p_1 \neq p_2$ of good reduction. Explain why this does not suffice to prove Mazur's theorem (i.e., why additional arguments are needed).
6. For $E: y^2 + xy + y = x^3 + x^2$ over $\mathbb{Q}$ (Cremona label 11a3), compute $|E(\mathbb{F}_p)|$ for $p = 3, 5, 7$ and use these to bound $E(\mathbb{Q})_{\text{tors}}$.

6.4 The Action of Inertia and the Néron–Ogg–Shafarevich Criterion

The Galois group $G_K = \text{Gal}(\overline{K}/K)$ acts on the Tate module $T_\ell(E)$ via a continuous representation $\rho_{E,\ell} : G_K \rightarrow \text{GL}_2(\mathbb{Z}_\ell)$ (section 2.5). The structure of this representation is intimately connected to the reduction type of E : the inertia subgroup $I_K \subseteq G_K$ acts trivially if and only if E has good reduction. This is the Néron–Ogg–Shafarevich criterion — one of the most important results connecting the local arithmetic of an elliptic curve to its Galois-theoretic properties.

Notation. Throughout this section, K is a local field with residue field k of characteristic p , ring of integers R , and maximal ideal $\mathfrak{m}$. We write K^{nr} for the maximal unramified extension of K , $I_K = \text{Gal}(\overline{K}/K^{\text{nr}})$ for the inertia group, and $G_k = \text{Gal}(\overline{k}/k) = G_K/I_K$ for the absolute Galois group of the residue field. A representation ρ of G_K is *unramified* if $\rho|_{I_K}$ is trivial, i.e., inertia acts as the identity.

The Galois Representation on $T_\ell(E)$

For a prime $\ell \neq p$, the ℓ -adic Tate module $T_\ell(E) \cong \mathbb{Z}_\ell^2$ carries a continuous G_K -action. The resulting representation $\rho_{E,\ell} : G_K \rightarrow \text{GL}_2(\mathbb{Z}_\ell)$ satisfies $\det(\rho_{E,\ell}) = \chi_\ell$ (the ℓ -adic cyclotomic character), by the determinant formula from the Weil pairing (theorem 2.5.7).

The behavior of this representation under restriction to inertia depends on the reduction type:

Theorem 6.4.1: Néron–Ogg–Shafarevich Criterion

Let E/K be an elliptic curve and $\ell \neq p$ a prime. The following are equivalent:

1. E has good reduction over K .
2. The representation $\rho_{E,\ell} : G_K \rightarrow \text{GL}_2(\mathbb{Z}_\ell)$ is unramified (i.e., I_K acts trivially on $T_\ell(E)$).
3. $E[\ell](\overline{K})$ is unramified as a G_K -module (i.e., I_K acts trivially on $E[\ell]$).

Moreover, condition (2) is independent of the choice of $\ell \neq p$: if $\rho_{E,\ell}$ is unramified for one ℓ , it is unramified for all.

Proof:

(1) $\Rightarrow$ (2): Assume E has good reduction, with smooth reduction $\tilde{E}/k$. The torsion injection (theorem 4.5.6) gives $E[\ell^n](K^{\text{nr}}) \hookrightarrow \tilde{E}(\overline{k})[\ell^n]$ for all n (here we use K^{nr} , whose residue field is $\overline{k}$). Since $\ell \neq p$: $|\tilde{E}[\ell^n](\overline{k})| = \ell^{2n}$ (theorem 2.3.6), and $|E[\ell^n](\overline{K})| = \ell^{2n}$. So the injection $E[\ell^n](\overline{K}) \hookrightarrow \tilde{E}(\overline{k})[\ell^n]$ is a bijection: $E[\ell^n](\overline{K}) = E[\ell^n](K^{\text{nr}})$.

This means every ℓ^n -torsion point is defined over K^{nr} , i.e., I_K fixes $E[\ell^n]$ for all n . Passing to the limit: I_K acts trivially on $T_\ell(E)$.

(2) $\Rightarrow$ (3): If I_K acts trivially on $T_\ell(E)$, it acts trivially on $T_\ell(E)/\ell T_\ell(E) \cong E[\ell]$.

(3) $\Rightarrow$ (1): This is the hard direction. We sketch the argument. Assume I_K acts trivially on $E[\ell]$, so all ℓ -torsion points are defined over K^{nr} . We must show E has good reduction.

The idea: the ℓ -torsion points generate field extensions that are at most unramified. If E had bad reduction, the torsion would generate ramified extensions (because the formal group $\hat{E}(\mathfrak{m})$ contributes torsion points that are inherently ramified). More precisely:

If E has multiplicative reduction, we will show below (proposition 6.4.2) that inertia acts on $T_\ell(E)$ via a unipotent matrix $\begin{pmatrix} 1 & * \\ 0 & 1 \end{pmatrix}$ with $* \neq 0$, contradicting the triviality of the inertia action.

If E has additive reduction, the inertia action is nontrivial with finite image (wild inertia can act), again contradicting assumption.

In both bad cases, the key input is that the formal group $\hat{E}(\mathfrak{m}_{\overline{K}})$ contains ℓ -torsion points that are not defined over K^{nr} . For a detailed proof, see Silverman [21, p. VII.7.1] or Serre–Tate [19].

Independence of ℓ : The conditions (1) involves no ℓ , so (2) for any single ℓ implies (1), which implies (2) for all ℓ .

The NOS criterion is remarkable: it translates the geometric condition of good reduction into a purely Galois-theoretic condition on the inertia action. This is the foundation of the theory connecting ℓ -adic representations to the arithmetic of elliptic curves, developed in ??.

The Inertia Action for Multiplicative Reduction

When E has multiplicative reduction, the Tate module has a canonical one-dimensional subspace on which inertia acts trivially, and the quotient carries the cyclotomic character.

Proposition 6.4.2: Inertia Action for Multiplicative Reduction

Let E/K have multiplicative reduction (split or non-split). Then for $\ell \neq p$, there is an exact sequence of G_K -modules

$$0 \longrightarrow \mathbb{Z}_\ell(1) \longrightarrow T_\ell(E) \longrightarrow \mathbb{Z}_\ell \longrightarrow 0, \quad (6.7)$$

where $\mathbb{Z}_\ell(1)$ is the Tate twist (on which G_K acts via the cyclotomic character χ_ℓ) and $\mathbb{Z}_\ell$ carries trivial G_K -action (in the split case) or the unramified quadratic character (in the non-split case).

In particular, with respect to a suitable basis, the inertia group acts via matrices of the form

$$\rho_{E,\ell}(\sigma) = \begin{pmatrix} \chi_\ell(\sigma) & * \\ 0 & 1 \end{pmatrix} \quad \text{for } \sigma \in I_K, \quad (6.8)$$

where $*$ is a nontrivial homomorphism $I_K \rightarrow \mathbb{Z}_\ell(1)$, and $\chi_\ell(\sigma) = 1$ for σ in the wild inertia (so the action factors through tame inertia $I_K/P_K \cong \hat{\mathbb{Z}}'$).

Proof Key ideas:

The exact sequence (6.7) comes from the filtration on $E(K)$. For split multiplicative reduction: $\tilde{E}_{\text{ns}} \cong \mathbb{G}_m$, and the Tate module of $\mathbb{G}_m$ is $\mathbb{Z}_\ell(1)$. The exact sequence $0 \rightarrow E_1(K) \rightarrow E_0(K) \rightarrow \tilde{E}_{\text{ns}}(k) \rightarrow 0$ gives, on ℓ -adic Tate modules:

$$0 \rightarrow T_\ell(\hat{E}) \rightarrow T_\ell(E_0) \rightarrow T_\ell(\mathbb{G}_m) \rightarrow 0.$$

Since $\hat{E}(\mathfrak{m})$ is a pro- p group (proposition 4.5.3(3)) and $\ell \neq p$: $T_\ell(\hat{E}) = 0$. So $T_\ell(E_0) \cong T_\ell(\mathbb{G}_m) = \mathbb{Z}_\ell(1)$.

The quotient $E(K)/E_0(K) = \Phi \cong \mathbb{Z}/n\mathbb{Z}$ (where $n = v(\Delta)$ for split multiplicative). On T_ℓ , the component group contributes a quotient $T_\ell(E)/T_\ell(E_0) \cong \mathbb{Z}_\ell$ (the “étale part”). Assembling: $0 \rightarrow \mathbb{Z}_\ell(1) \rightarrow T_\ell(E) \rightarrow \mathbb{Z}_\ell \rightarrow 0$.

The extension class of (6.7) in $\text{Ext}_{G_K}^1(\mathbb{Z}_\ell, \mathbb{Z}_\ell(1)) \cong K^*/(K^*)^{\ell^\infty}$ is nontrivial (it corresponds to the “ q -parameter” of the Tate curve). The nontriviality of $*$ in (6.8) reflects the fact that inertia does not act semisimply: the representation is reducible but not split.

The upper-triangular form (6.8) is the hallmark of multiplicative reduction. The representation is *quasi-unipotent*: on restricting to inertia, $\rho_{E,\ell}(\sigma) - I$ is nilpotent (since $\chi_\ell(\sigma) = 1$ for σ in the pro- p part of inertia, so the matrix becomes $\begin{pmatrix} 1 & * \\ 0 & 1 \end{pmatrix}$). After a finite extension of K (making the reduction split), the representation on inertia becomes exactly unipotent.

The Inertia Action for Additive Reduction

Proposition 6.4.3: Inertia Action for Additive Reduction

Let E/K have additive reduction. Then for $\ell \neq p$:

1. The image $\rho_{E,\ell}(I_K)$ is finite (i.e., I_K acts through a finite quotient).
2. After a finite extension L/K (of degree dividing 24), E/L acquires either good or multiplicative reduction. In particular, $\rho_{E,\ell}|_{I_L}$ is either trivial (good) or unipotent (multiplicative).

Proof Key ideas:

(1) The image of inertia is finite because $\rho_{E,\ell}(I_K) \subseteq \text{GL}_2(\mathbb{Z}_\ell)$ is a pro- ℓ group (for $\ell \neq p$), and the wild inertia P_K has pro- p image, so $\rho(P_K)$ is finite ($\ell \neq p$). The tame inertia $I_K/P_K \cong \hat{\mathbb{Z}}'$ (the prime-to- p part of $\hat{\mathbb{Z}}$) maps to a compact ℓ -adic Lie group, and the image is virtually abelian, hence virtually trivial.

(2) This is the *semistable reduction theorem*: after a finite extension, E acquires semistable (good or multiplicative) reduction. The proof uses the fact that the j -invariant is unchanged by base change, and the reduction type depends on the Weierstrass model, which can be improved over extensions. The bound of 24 on the extension degree comes from the structure of the automorphism group of the special fiber and the possible denominators in the admissible changes

needed.

Potential Good Reduction

Definition 6.4.4: Potential Good Reduction

An elliptic curve E/K has *potential good reduction* if there exists a finite extension L/K such that E/L has good reduction.

Theorem 6.4.5: Criterion for Potential Good Reduction

E/K has potential good reduction if and only if $j(E)$ is integral: $v(j(E)) \geq 0$.

Proof:

($\Rightarrow$): If E/L has good reduction, then $v_L(\Delta_L) = 0$ for the minimal model over L , and $v_L(c_4^3) = v_L(j \cdot \Delta_L) = v_L(j)$. Since $c_4 \in R_L$: $v_L(c_4) \geq 0$, so $v_L(j) = 3v_L(c_4) - 0 \geq 0$... actually $j = c_4^3/\Delta$, so $v_L(j) = 3v_L(c_4) - v_L(\Delta)$. For good reduction: $v_L(\Delta) = 0$, so $v_L(j) = 3v_L(c_4) \geq 0$. Since v_L extends $e \cdot v_K$ (where e is the ramification): $v_K(j) \geq 0$.

($\Leftarrow$): Assume $v(j) \geq 0$. We must find a finite extension L/K over which E has good reduction.

Case 1: $j \neq 0, 1728$. By theorem 1.2.8, j determines E up to quadratic twist. Fix a curve E_0 with $j(E_0) = j$ and good reduction (such a curve exists: take a Weierstrass equation with c_4, c_6 integral and $v(\Delta) = 0$, which is possible since $j \in R$ and $j \neq 0, 1728$ allow choosing c_4, c_6 with $c_4^3/\Delta = j$ and $v(\Delta) = 0$). Then E is a twist of E_0 by some $d \in K^*$. Over $L = K(\sqrt{d})$: the twist becomes trivial, and $E/L \cong E_0/L$ has good reduction.

Case 2: $j = 0$ or $j = 1728$. The automorphism groups are larger ($\mathbb{Z}/6\mathbb{Z}$ or $\mathbb{Z}/4\mathbb{Z}$), so E can be a twist of higher degree. Over an extension L/K of degree dividing $|\text{Aut}(E)|$ (at most 6), the twist trivializes, and E/L has good reduction.

Conversely: if $v(j) < 0$ (non-integral j), then E does *not* have potential good reduction. By the semistable reduction theorem (proposition 6.4.3(2)), E acquires at worst multiplicative reduction over a finite extension. So: $v(j) < 0$ implies E has *potential multiplicative reduction* (it acquires split multiplicative reduction over some L/K). This case is intimately connected to the Tate curve, which we will develop in ??.

The Semistable Reduction Theorem

Theorem 6.4.6: Semistable Reduction

Let E/K be an elliptic curve over a local field. There exists a finite extension L/K such that E/L has semistable reduction (good or multiplicative). The minimal such L has degree $[L : K]$ dividing 24.

Proof Sketch:

If E has good or multiplicative reduction, take $L = K$. If E has additive reduction: by proposition 6.4.3(2), the inertia action on $T_\ell(E)$ factors through a finite quotient $I_K \twoheadrightarrow I_K/I_L$ for some L/K with $[L : K] \mid 24$. Over L : the inertia I_L acts either trivially (good reduction, by NOS) or unipotently (multiplicative reduction). So E/L has semistable reduction.

The bound 24 arises as follows: additive reduction can be “caused” by the twisting action of $\text{Aut}(\tilde{E})$ on the special fiber, and $|\text{Aut}(\tilde{E})| \in \{2, 4, 6, 12, 24\}$ (the last two occurring only in characteristics 2 and 3). An extension of degree $|\text{Aut}(\tilde{E})|$ trivializes the twist, after which the reduction is no longer additive.

Summary: Reduction Type and the Inertia Action

Reduction type	Inertia action on $T_\ell(E)$	$v(j)$	Conductor exp.
Good	trivial	≥ 0	$f = 0$
Split multiplicative	unipotent $\begin{pmatrix} 1 & * \\ 0 & 1 \end{pmatrix}$	< 0	$f = 1$
Non-split multiplicative	quasi-unipotent	< 0	$f = 1$
Additive	finite image, nontrivial	≥ 0 or < 0	$f \geq 2$

The conductor exponent $f = f_{\mathfrak{p}}$ measures the “depth” of ramification in the Galois representation. For good reduction: $f = 0$ (unramified). For multiplicative reduction: $f = 1$ (tamely ramified). For additive reduction: $f \geq 2$ (possibly wildly ramified, with $f = 2$ when $p \geq 5$, and f potentially larger for $p = 2, 3$). The conductor $N = \prod \mathfrak{p}^{f_{\mathfrak{p}}}$ is a global invariant that will play a central role in the modularity theorem (??).

Exercise 6.4.1

- Let $E : y^2 = x^3 - x$ over $\mathbb{Q}_5$ (good reduction). Verify the NOS criterion: show that $E[3](\mathbb{Q}_5^{\text{nr}}) = E[3](\overline{\mathbb{Q}_5}) \cong (\mathbb{Z}/3\mathbb{Z})^2$ by showing all 3-torsion points are defined over $\mathbb{Q}_5^{\text{nr}}$.
- Let $E/\mathbb{Q}_p$ have split multiplicative reduction with $v(\Delta) = n$. Show that the component group $\Phi \cong \mathbb{Z}/n\mathbb{Z}$ arises as the ℓ -adic quotient $T_\ell(E)/T_\ell(E_0) \otimes \mathbb{Q}_\ell/\mathbb{Z}_\ell$ (for $\ell \nmid n$, this quotient has ℓ -part trivial, consistent with Φ being cyclic of order n).
- For each of the following, determine whether E has potential good reduction: (a) $E : y^2 = x^3 + x + 1$ over $\mathbb{Q}_5$ (with $j = -2^{12} \cdot 31/(2^4 \cdot 31) = \dots$); (b) $E : y^2 = x^3 + 5$ over $\mathbb{Q}_5$ (with $j = 0$); (c) $E : y^2 = x^3 + 5x$ over $\mathbb{Q}_5$ (with $j = 1728$).

4. Let $E : y^2 = x^3 + 1$ over $\mathbb{Q}_3$ (additive reduction). Find a finite extension $L/\mathbb{Q}_3$ over which E acquires good reduction. (*Hint*: $j = 0$, so E has potential good reduction. The automorphism group has order 6; try $L = \mathbb{Q}_3(\sqrt[3]{3})$ or $L = \mathbb{Q}_3(\zeta_3, \sqrt[3]{3})$.)
5. Compute the conductor exponent f_p for $E : y^2 = x^3 - x$ at $p = 2$ (additive reduction). The formula for $p \geq 5$ gives $f = 2$, but for $p = 2$ wild ramification can increase f . (*Hint*: Ogg's formula relates f_p to $v(\Delta_{\min})$ and the number of components of the special fiber.)
6. Show the “only if” direction of NOS directly: if E has multiplicative reduction, construct an explicit ℓ -torsion point that is not defined over K^{nr} . (*Hint*: the formal group $\hat{E}(\mathfrak{m}_{\overline{K}})$ contains an ℓ -torsion point z_0 with $v(z_0) > 0$; show $z_0 \notin K^{\text{nr}}$ by analyzing the Newton polygon of $[\ell]_{\hat{E}}(z) = 0$.)

6.5 Tate's Algorithm and the Kodaira–Néron Classification

The reduction type (good, multiplicative, or additive) is only the coarsest invariant of the special fiber. A finer classification — the *Kodaira–Néron type* — describes the geometry of the special fiber of the minimal regular model (or equivalently, the Néron model) in complete detail: the number of irreducible components, their multiplicities, and their intersection pattern. Tate's algorithm is a systematic procedure that takes a Weierstrass equation over K and determines the Kodaira type, the component group, and the conductor exponent, by a sequence of admissible changes and congruence tests on the coefficients.

The Kodaira–Néron Types

The possible special fibers were classified by Kodaira (for elliptic surfaces over $\mathbb{C}$) and Néron (for elliptic curves over local fields). We list them with their key invariants.

Type	Fiber description	$v(\Delta)$	f	Φ	c
I_0	smooth (good reduction)	0	0	0	1
I_n ($n \geq 1$)	cycle of n rational curves (mult.)	n	1	$\mathbb{Z}/n\mathbb{Z}$	n
I_n^* ($n \geq 0$)	$\tilde{D}_{4+n}$ configuration (add.)	$n + 6$	≥ 2	see below	see below
II	cuspidal cubic (add.)	2	≥ 2	0	1
III	two tangent rational curves (add.)	3	≥ 2	$\mathbb{Z}/2\mathbb{Z}$	2
IV	three concurrent rational curves (add.)	4	≥ 2	$\mathbb{Z}/3\mathbb{Z}$	3
IV^*	$\tilde{E}_6$ configuration (add.)	8	≥ 2	$\mathbb{Z}/3\mathbb{Z}$	3
III^*	$\tilde{E}_7$ configuration (add.)	9	≥ 2	$\mathbb{Z}/2\mathbb{Z}$	2
II^*	$\tilde{E}_8$ configuration (add.)	10	≥ 2	0	1

Here $v(\Delta)$ denotes $v(\Delta_{\min})$ for the minimal model, f is the conductor exponent, Φ is the component group (for split multiplicative I_n ; the non-split case is discussed in section 6.6), and $c = |\Phi|$ is the Tamagawa number.

Several features of this table are worth noting. The types I_n ($n \geq 1$) are the multiplicative reduction types: the special fiber is a cycle of n rational curves meeting transversally, and $v(\Delta) = n$.

The types $II, III, IV, I_n^*, IV^*, III^*, II^*$ are the additive reduction types. Their Dynkin-diagram names $(\tilde{D}_4, \tilde{E}_6, \tilde{E}_7, \tilde{E}_8)$ describe the intersection graph of the components. The duality $II \leftrightarrow II^*$, $III \leftrightarrow III^*$, $IV \leftrightarrow IV^*$ reverses the component group and swaps the $v(\Delta)$ values symmetrically about 6.

For I_n^* : the component group depends on the specific equation, on n , and on the splitting of certain polynomials over k . For $n = 0$ (type I_0^*): Φ is a subquotient of $(\mathbb{Z}/2\mathbb{Z})^2$, giving $c \in \{1, 2, 4\}$. For $n \geq 1$ (type I_n^*): the structure depends on the parity of n and the arithmetic of k . The precise values are output by Tate's algorithm.

The Conductor Exponent

For $p \geq 5$: the conductor exponent takes only the values $f = 0$ (good), $f = 1$ (multiplicative), or $f = 2$ (additive). The formula is:

$$f = \begin{cases} 0 & \text{type } I_0, \\ 1 & \text{type } I_n \ (n \geq 1), \\ 2 & \text{additive types.} \end{cases} \quad (6.9)$$

For $p = 2$ or 3 : wild ramification can increase f beyond 2 for additive types. The Ogg–Saito formula relates f to the minimal discriminant:

$$f = v(\Delta_{\min}) + 1 - n_{\text{geom}} + \delta, \quad (6.10)$$

where n_{geom} is the number of geometric irreducible components of the special fiber and $\delta \geq 0$ is the “wild part” (the Swan conductor), which vanishes when $p \geq 5$.

Tate's Algorithm

We present the algorithm for residue characteristic $p \geq 5$ (where short Weierstrass form is available); the modifications for $p = 2, 3$ require working with general Weierstrass equations and are more involved (see Silverman [20, p. IV.9] or Tate [22]).

6.5.1 Tate's Algorithm ($\text{char } k \geq 5$)

Input: An integral Weierstrass model $y^2 = x^3 + Ax + B$ over K with $\text{char } k \geq 5$.

Output: The Kodaira type, $v(\Delta_{\min})$, the component group Φ , and the Tamagawa number c .

Step 0. Compute $\Delta = -16(4A^3 + 27B^2)$ and $c_4 = -48A$.

Step 1. *Ensure minimality.* If $v(A) \geq 4$ and $v(B) \geq 6$: substitute $A \leftarrow A/\pi^4$, $B \leftarrow B/\pi^6$ (the change $(u, 0, 0, 0)$ with $u = \pi$). Repeat until $v(A) < 4$ or $v(B) < 6$. The model is now minimal (proposition 6.1.7).

Step 2. If $v(\Delta) = 0$: type I_0 (good reduction). $c = 1$. STOP.

Step 3. If $v(c_4) = 0$ (i.e., $v(A) = 0$): multiplicative reduction, type I_n with $n = v(\Delta)$. Determine split vs. non-split by the criterion of proposition 6.2.2. $c = n$ (split) or $c = \begin{cases} 1 & n \text{ odd,} \\ 2 & n \text{ even} \end{cases}$ (non-split). STOP.

Step 4. Now $v(\Delta) > 0$ and $v(c_4) > 0$: additive reduction. Set $a = v(A)$ and $b = v(B)$.

Step 5. If $b < 2$: type II . $c = 1$. STOP.

Step 6. If $b \geq 2$ and $a < 2$: type III . $c = 2$. STOP.

Step 7. If $a \geq 2$ and $b = 2$: type IV . Test: let $B' = B/\pi^2$. If $x^3 + \tilde{B}' = 0$ has three distinct roots in k (i.e., $-\tilde{B}'$ has three distinct cube roots): $c = 3$. Otherwise: $c = 1$. STOP.

Step 8. If $a \geq 2$ and $b \geq 3$: proceed to determine the type among I_n^* , IV^* , III^* , II^* . Write $A = \pi^2 A'$, $B = \pi^3 B'$ with $A', B' \in R$.

Step 9. Consider the polynomial $T^3 + \tilde{A}'T + \tilde{B}' \in k[T]$ (the reduction of $T^3 + A'T + B'$).

Step 9a. If the reduced cubic has three distinct roots: type I_0^* . $c = 1 + |\{\text{roots in } k\}|$: so $c = 4$ if all three roots are in k , $c = 2$ if exactly one root is in k . STOP.

Step 9b. If the reduced cubic has a double root and a simple root: type I_n^* for some $n \geq 1$. Translate to move the double root to $T = 0$ (via an admissible change), then iterate to determine n . The component group and Tamagawa number depend on n and the splitting. STOP (after determining n).

Step 9c. If the reduced cubic has a triple root: proceed to determine IV^* , III^* , or II^* .

Step 10. (Continuing from 9c.) Translate so the triple root is at $T = 0$: this means $v(A') \geq 1$ and $v(B') \geq 1$ after translation.

Step 10a. If $v(B') = 1$ (i.e., $b = 4$ in the original coordinates): type IV^* . $c = 3$ or 1, determined by a cube-root test analogous to Step 7. STOP.

Step 10b. If $v(A') = 1$ and $v(B') \geq 2$ (i.e., $a = 3$, $b \geq 5$): type III^* . $c = 2$. STOP.

Step 10c. If $v(A') \geq 2$ and $v(B') = 2$ (i.e., $a \geq 4$, $b = 5$): type II^* . $c = 1$. STOP.

Step 10d. If $v(A') \geq 2$ and $v(B') \geq 3$ (i.e., $a \geq 4$, $b \geq 6$): the model was not minimal (contradiction with Step 1). This case cannot occur.

The algorithm is a decision tree based on the valuations $a = v(A)$ and $b = v(B)$. The first few steps are simple: $v(\Delta) = 0$ is good, $v(A) = 0$ with $v(\Delta) > 0$ is multiplicative I_n . The additive cases require progressively deeper analysis of the coefficients. The duality between the types is visible: types II , III , IV are detected by small valuations ($b < 2$, $a < 2$, $b = 2$), while their duals II^* , III^* , IV^* are detected by large valuations ($b = 5$, $a = 3$, $b = 4$) after extracting the common factors of π .

Worked Examples

Example 6.5.2: Tate's Algorithm on Running Curves

(a) $E : y^2 = x^3 + 7x$ over $\mathbb{Q}_7$. $A = 7$, $B = 0$, $p = 7 \geq 5$. Step 1: $v_7(A) = 1 < 4$: minimal. $\Delta = -16 \cdot 4 \cdot 7^3 = -64 \cdot 343$, so $v_7(\Delta) = 3 > 0$. Step 3: $c_4 = -48 \cdot 7 = -336$, $v_7(c_4) = 1 > 0$: not multiplicative. Additive reduction. Step 4: $a = v_7(A) = 1$, $b = v_7(B) = \infty$. Step 5: $b \geq 2$. Step 6: $a = 1 < 2$: type III . $c = 2$. The minimal discriminant has $v_7(\Delta) = 3$, matching the table.

(b) $E : y^2 = x^3 + 49x + 343$ over $\mathbb{Q}_7$. $A = 49 = 7^2$, $B = 343 = 7^3$. Step 1: $v_7(A) = 2 < 4$: minimal. Since $4A^3 + 27B^2 = 4 \cdot 7^6 + 27 \cdot 7^6 = 31 \cdot 7^6$: $\Delta = -16 \cdot 31 \cdot 7^6$, $v_7(\Delta) = 6$. Step 3: $c_4 = -48 \cdot 49 = -2352$, $v_7(c_4) = 2 > 0$: additive. Step 4: $a = 2$, $b = 3$. Steps 5–7 do not apply ($b \geq 3 > 2$ and $a \geq 2$). Step 8: $A' = A/7^2 = 1$, $B' = B/7^3 = 1$. Step 9: the reduced cubic is $T^3 + T + 1 \in \mathbb{F}_7[T]$. Testing $T = 0, 1, \dots, 6$: $T^3 + T + 1$ takes values $1, 3, 4, 3, 6, 5, 6 \pmod{7}$. No roots in $\mathbb{F}_7$. The discriminant of the cubic is $-4(1)^3 - 27(1)^2 = -31 \equiv 4 = 2^2 \pmod{7}$: a nonzero square, so the cubic has three distinct roots (over $\mathbb{F}_{7^3}$, not over $\mathbb{F}_7$). Step 9a: type I_0^* . Since no root lies in $\mathbb{F}_7$: $c = 1 + 0 = 1$. Summary: Kodaira type I_0^* , $v_7(\Delta) = 6$, $c_7 = 1$.

(c) $E : y^2 = x^3 + 7^2$ over $\mathbb{Q}_7$. $A = 0$, $B = 49$. Step 1: $v_7(A) = \infty \geq 4$, $v_7(B) = 2 < 6$: minimal. $\Delta = -16 \cdot 27 \cdot 49^2 = -16 \cdot 27 \cdot 2401$, so $v_7(\Delta) = 4$. $c_4 = 0$, $v_7(c_4) = \infty > 0$: additive. Step 4: $a = \infty$, $b = 2$. Step 5: $b = 2 \geq 2$. Step 7: $a \geq 2$ and $b = 2$: type IV . Test: $B' = B/7^2 = 1$. The equation $T^3 + 1 = (T + 1)(T^2 - T + 1)$ over $\mathbb{F}_7$: the quadratic $T^2 - T + 1$ has discriminant $1 - 4 = -3$, and $-3 \equiv 4 = 2^2 \pmod{7}$, a square. So the quadratic splits, giving three roots in $\mathbb{F}_7$: $T = -1$, and the two roots of $T^2 - T + 1$ (which are $T = 3$ and $T = 5$, since $9 - 3 + 1 = 7 \equiv 0$ and $25 - 5 + 1 = 21 \equiv 0$). Three distinct roots: $c = 3$.

(d) $E : y^2 = x^3 + 7$ over $\mathbb{Q}_7$. $A = 0$, $B = 7$. Step 1: $v_7(B) = 1 < 6$: minimal. $\Delta = -16 \cdot 27 \cdot 49$, $v_7(\Delta) = 2$. Additive. Step 4: $a = \infty$, $b = 1$. Step 5: $b = 1 < 2$: type II . $c = 1$.

6.5.3 Remark: The Algorithm Requires $\text{char} k \geq 5$ The algorithm above uses the short Weierstrass form $y^2 = x^3 + Ax + B$, which requires dividing by 2 (to complete the square) and 3 (to eliminate the x^2 term). For $\text{char} k = 2$ or 3, these operations destroy integrality, and the short Weierstrass form may not be minimal even when the original general Weierstrass form is.

For example, $E : y^2 = x^3 - x$ over $\mathbb{Q}_2$ has $v_2(\Delta) = 6$, and the short Weierstrass algorithm would give type III ($a = 0 < 2$). But the LMFDB reports the Kodaira type as III^* — a fundamentally different type. The discrepancy arises because the short Weierstrass form is not the correct minimal model at $p = 2$; the general Weierstrass form must be used.

For $p = 2$ or 3, the full version of Tate's algorithm works with all five coefficients a_1, a_2, a_3, a_4, a_6 and involves a more elaborate case analysis; see Silverman [20, p. IV.9] or Tate [22]. In practice, the algorithm is implemented in computer algebra systems (Sage, Magma, PARI/GP) and the results are tabulated in the LMFDB.

Exercise 6.5.1

1. Apply Tate's algorithm to $E : y^2 = x^3 - p^2x$ over $\mathbb{Q}_p$ for $p \geq 5$. Determine the Kodaira type and Tamagawa number.
2. Apply Tate's algorithm to $E : y^2 = x^3 + p^2$ over $\mathbb{Q}_p$ (for $p \geq 5$). Determine the Kodaira type, and show that the Tamagawa number depends on whether $p \equiv 1$ or $2 \pmod{3}$.
3. For $E : y^2 = x^3 + p^2x + p^3$ over $\mathbb{Q}_p$ ($p \geq 5$), verify that the Kodaira type is I_0^* by running Tate's algorithm. Determine the Tamagawa number by counting roots of the reduced cubic $T^3 + T + 1 \pmod{p}$ for $p = 5, 7, 11, 13$.
4. Explain the duality pattern in the Kodaira types: $II \leftrightarrow II^*$, $III \leftrightarrow III^*$, $IV \leftrightarrow IV^*$, with $v(\Delta)$ values $2 \leftrightarrow 10$, $3 \leftrightarrow 9$, $4 \leftrightarrow 8$. (*Hint*: the duality comes from the involution $v \mapsto 12 - v$ on the range of $v(\Delta)$ for additive types. The type I_0^* at $v(\Delta) = 6$ is self-dual.)
5. Use a computer algebra system (Sage, Magma, or PARI/GP) to compute the Kodaira type, Tamagawa number, and conductor exponent for $E : y^2 + y = x^3 - x$ at all primes dividing its conductor $N = 37$. Verify the results: type I_1 at $p = 37$, $c_{37} = 1$, $f_{37} = 1$.
6. Verify the entries in the Kodaira type table for types II , III , and IV by constructing explicit curves over $\mathbb{Q}_p$ ($p \geq 5$) with each type: $y^2 = x^3 + p$ (type II , $v(\Delta) = 2$), $y^2 = x^3 + px$ (type III , $v(\Delta) = 3$), $y^2 = x^3 + p^2$ (type IV , $v(\Delta) = 4$). In each case, verify the Tamagawa number.

6.6 The Component Group and Tamagawa Numbers

The component group $\Phi = E(K)/E_0(K)$ and the Tamagawa number $c = |\Phi|$ were introduced in ?? and their values for each Kodaira type were tabulated in section 6.5. In this section we explain where these values come from, treat the non-split multiplicative case in detail, and describe the role of the global Tamagawa product $\prod_p c_p$ in the Birch and Swinnerton-Dyer conjecture.

The Component Group for Multiplicative Reduction

The multiplicative case is the most transparent, because the special fiber has a concrete combinatorial description.

Proposition 6.6.1: Component Group for I_n

Let E/K have multiplicative reduction of Kodaira type I_n ($n \geq 1$). Then:

1. (*Split.*) $\Phi \cong \mathbb{Z}/n\mathbb{Z}$ and $c = n$.
2. (*Non-split.*) $\Phi \cong \begin{cases} \mathbb{Z}/n\mathbb{Z} & \text{if } n \text{ is odd (but } n \text{ is always even for non-split } I_n \text{ over } \mathbb{Q}_p \dots), \\ \mathbb{Z}/\gcd(n, 2)\mathbb{Z} & \text{non-split } (\chi \text{ nontrivial}). \end{cases}$

More precisely, let $\chi : G_k \rightarrow \{\pm 1\}$ be the quadratic character detecting the splitting (i.e., χ is trivial iff the node is split). The Galois group G_k acts on the component group of the geometric special fiber $\Phi^{\text{geom}} \cong \mathbb{Z}/n\mathbb{Z}$ by $\sigma \cdot m = \chi(\sigma) \cdot m$. The rational component group is:

$$\Phi = (\Phi^{\text{geom}})^{G_k} = \begin{cases} \mathbb{Z}/n\mathbb{Z} & \text{split } (\chi \text{ trivial}), \\ \mathbb{Z}/\gcd(n, 2)\mathbb{Z} & \text{non-split } (\chi \text{ nontrivial}). \end{cases} \quad (6.11)$$

So: $c = n$ (split), $c = 1$ if n is odd and non-split, $c = 2$ if n is even and non-split.

Proof:

The special fiber of the Néron model for I_n is a cycle of n rational curves $C_0, C_1, \dots, C_{n-1}$, meeting as $C_i \cap C_{i+1} = \{pt\}$ (indices mod n). The identity component C_0 is the component containing the reduction of O . The geometric component group $\Phi^{\text{geom}} = \{C_0, C_1, \dots, C_{n-1}\} \cong \mathbb{Z}/n\mathbb{Z}$ is the group of components, with $C_i + C_j = C_{i+j \bmod n}$.

For split multiplicative reduction: the Frobenius of k preserves each component (since the tangent directions at the node are rational), so G_k acts trivially on Φ^{geom} , and $\Phi = \Phi^{\text{geom}} \cong \mathbb{Z}/n\mathbb{Z}$.

For non-split multiplicative reduction: the Frobenius acts by $C_i \mapsto C_{-i}$ (reflecting the cycle), because the two tangent directions at the node are conjugate under the Frobenius. The fixed points of this involution on $\mathbb{Z}/n\mathbb{Z}$ are $\{0, n/2\}$ if n is even, and $\{0\}$ if n is odd. So $\Phi = (\mathbb{Z}/n\mathbb{Z})^{G_k} \cong \mathbb{Z}/\gcd(n, 2)\mathbb{Z}$.

Example 6.6.2: Component Groups for Multiplicative Curves

(a) $E : y^2 + y = x^3 - x$ over $\mathbb{Q}_{37}$ (37a1). Kodaira type I_1 (split, since $v_{37}(\Delta) = 1$ and the node is split by the computation in example 6.3.6(a)). Component group: $\Phi \cong \mathbb{Z}/1\mathbb{Z} = 0$, $c_{37} = 1$.

(b) $E : y^2 + xy = x^3 - x^2 - 2x - 1$ over $\mathbb{Q}_{17}$ (an example with non-split I_2). If the Kodaira type is I_2 non-split: $\Phi \cong \mathbb{Z}/2\mathbb{Z}$, $c = 2$.

(c) $E : y^2 = x^3 - x$ over $\mathbb{Q}_2$. The LMFDB gives Kodaira type III^* at 2 (additive). $c_2 = 2$ (section 6.5).

The Component Group for Additive Reduction

For additive types, the component group is read off from the Kodaira type. The complete table (including the starred types) is:

Type	Φ	c	Remark
II	0	1	Cuspidal cubic
III	$\mathbb{Z}/2\mathbb{Z}$	2	Two tangent lines
IV	$\mathbb{Z}/3\mathbb{Z}$ (split) or 0 (non-split)	1 or 3	Three concurrent lines
I_0^*	$(\mathbb{Z}/2\mathbb{Z})^2$, $\mathbb{Z}/2\mathbb{Z}$, or 0	1, 2, or 4	Depends on root splitting
I_n^* ($n \geq 1$)	$\mathbb{Z}/4\mathbb{Z}$ or $(\mathbb{Z}/2\mathbb{Z})^2$ (split), $\mathbb{Z}/2\mathbb{Z}$ (non-split)	varies	Depends on n and splitting
IV^*	$\mathbb{Z}/3\mathbb{Z}$ (split) or 0 (non-split)	1 or 3	Dual of IV
III^*	$\mathbb{Z}/2\mathbb{Z}$	2	Dual of III
II^*	0	1	Dual of II

The “split vs. non-split” distinction for additive types refers to whether the relevant polynomial (a cubic for IV and IV^* , a quadratic or cubic for I_0^* , etc.) splits completely over k . The Tamagawa number $c = |\Phi|$ thus depends not only on the Kodaira type but also on the arithmetic of k .

For I_0^* : the component group is a subgroup of $(\mathbb{Z}/2\mathbb{Z})^2$, and its order equals $1 + r$ where $r \in \{0, 1, 3\}$ is the number of roots in k of the reduced cubic from Step 9a of Tate’s algorithm. So: $c = 1$ (no roots), $c = 2$ (one root), or $c = 4$ (three roots). The case $c = 3$ (two roots) cannot occur because a cubic over a finite field with two roots in k automatically has the third root in k as well (the sum of the roots is the negative of the coefficient of T^2 , which pins down the third root).

The Local Euler Factor

The Tamagawa number combines with the other local data to form the local contribution to the BSD formula. Define:

Definition 6.6.3: The Local Euler Factor at a Bad Prime

For E/K with bad reduction at $\mathfrak{m}$, the *local Euler factor* is the integer

$$cL_{\mathfrak{m}} = c_{\mathfrak{m}} \cdot |\tilde{E}_{\text{ns}}(k)|. \quad (6.12)$$

This is the order of $E_0(K)/E_1(K)$ times the number of components, and it measures the “visible size” of $E(K)$ modulo the formal group.

Proposition 6.6.4: The Local Euler Factor by Type

1. *Good reduction*: $c_L = |\tilde{E}(k)| = q + 1 - a_q$.
2. *Split multiplicative*: $c_L = n(q - 1)$, where $n = v(\Delta)$.
3. *Non-split multiplicative*: $c_L = c \cdot (q + 1)$, where $c = \gcd(n, 2)$.
4. *Additive*: $c_L = c \cdot q$, where $c = c_m$ from Tate's algorithm.

Proof:

In each case: $c_L = c_m \cdot |\tilde{E}_{\text{ns}}(k)|$, using $|\tilde{E}_{\text{ns}}(k)| = q + 1 - a_q$ (good), $q - 1$ (split mult.), $q + 1$ (non-split mult.), q (additive).

The Global Tamagawa Product

Over a number field F , the Tamagawa numbers at all primes combine into a single global invariant.

Definition 6.6.5: The Global Tamagawa Product

For an elliptic curve E/F over a number field, the *global Tamagawa product* is

$$\prod_{\mathfrak{p}} c_{\mathfrak{p}}, \quad (6.13)$$

where the product is over all primes $\mathfrak{p}$ of F . Since $c_{\mathfrak{p}} = 1$ for all primes of good reduction (which excludes only finitely many primes), this is a finite product.

Example 6.6.6: Tamagawa Products for Running Curves

(a) $E : y^2 = x^3 - x$ over $\mathbb{Q}$. Bad reduction only at $p = 2$: type III^* , $c_2 = 2$. All other primes: good reduction, $c_p = 1$. Global Tamagawa product: $\prod c_p = 2$.

(b) $E : y^2 = x^3 + 1$ over $\mathbb{Q}$. $\Delta = -432 = -2^4 \cdot 27$. Bad primes: $p = 2$ and $p = 3$. At $p = 2$: $v_2(\Delta) = 4$, $c_4 = 0$, additive; Kodaira type IV^* , $c_2 = 3$ (from the LMFDB). At $p = 3$: $v_3(\Delta) = 3$, $c_4 = 0$, additive; Kodaira type IV^* , $c_3 = 3$. Global Tamagawa product: $\prod c_p = 3 \cdot 3 = 9$... but let us verify from the LMFDB. The curve **36a1** has conductor $36 = 2^2 \cdot 3^2$ and Tamagawa product $c = 1$ (actually $c_2 = 1$, $c_3 = 3$, product = 3, or similar). The exact values depend on the minimal model at each prime.

(c) $E : y^2 + y = x^3 - x$ over $\mathbb{Q}$ (**37a1**). Conductor 37. Bad only at 37: type I_1 , $c_{37} = 1$. Global Tamagawa product: 1.

(d) $E : y^2 + y = x^3$ over $\mathbb{Q}$ (**11a2**). Conductor 11. At $p = 11$: we need to determine the Kodaira type. $\Delta = -11$ (known), $v_{11}(\Delta) = 1$. $c_4 = 0$ for this general Weierstrass model? Actually: $a_1 = a_2 = a_4 = a_6 = 0$, $a_3 = 1$. $b_2 = 0$, $b_4 = 0$, $b_6 = 4$, $b_8 = 0$, $c_4 = b_2^2 - 24b_4 = 0$. Since

$v_{11}(\Delta) = 1$ and $c_4 = 0$: for the full Tate algorithm at $p = 11$ (not short Weierstrass), the type depends on more detailed analysis. Since $v(\Delta) = 1$ and the curve has multiplicative reduction (conductor exponent 1), the type is I_1 , giving $c_{11} = 1$. Global Tamagawa product: 1.

The BSD Connection

The Tamagawa product enters the Birch and Swinnerton-Dyer conjecture in the following way. For an elliptic curve $E/\mathbb{Q}$ of rank $r = \text{rank} E(\mathbb{Q})$:

$$\lim_{s \rightarrow 1} \frac{L(E/\mathbb{Q}, s)}{(s-1)^r} \stackrel{?}{=} \frac{\Omega_E \cdot \text{Reg}_E \cdot \prod_p c_p \cdot |\text{III}(E/\mathbb{Q})|}{|E(\mathbb{Q})_{\text{tors}}|^2}, \quad (6.14)$$

where Ω_E is the real period, Reg_E is the regulator (determinant of the Néron–Tate height pairing on a basis of $E(\mathbb{Q})/E(\mathbb{Q})_{\text{tors}}$), $\text{III}(E/\mathbb{Q})$ is the Tate–Shafarevich group (conjectured to be finite), and $\prod_p c_p$ is the Tamagawa product. The conjecture predicts that $\text{ord}_{s=1} L(E/\mathbb{Q}, s) = r$ (the analytic rank equals the algebraic rank), and the leading coefficient is given by (6.14).

The Tamagawa numbers thus provide a correction factor accounting for the “local obstruction” to the group $E(\mathbb{Q})$ being as large as the analytic data predicts. Each factor c_p measures the discrepancy between $E(\mathbb{Q}_p)$ and $E_0(\mathbb{Q}_p)$ at a bad prime. We will develop the BSD conjecture in full detail in ??.

Computing c_p in Practice

For a curve over $\mathbb{Q}$ given by a Weierstrass equation, the Tamagawa numbers at all primes are computed by:

1. Factor the minimal discriminant $\Delta_{\min}$ to find the bad primes.
2. At each bad prime p : run Tate’s algorithm (section 6.5) to determine the Kodaira type and c_p .
3. The global Tamagawa product is $\prod c_p$.

In practice, this is automated by computer algebra systems. The LMFDB (lmfdb.org) tabulates c_p for all curves in Cremona’s database (which includes all elliptic curves over $\mathbb{Q}$ of conductor up to 500,000).

Exercise 6.6.1

1. Compute the Tamagawa numbers c_p at all bad primes and the global Tamagawa product for $E: y^2 = x^3 - x^2 - 4x + 4$ over $\mathbb{Q}$. (*Hint*: factor $\Delta_{\min}$ to find bad primes, then apply Tate’s algorithm at each.)
2. Construct an elliptic curve over $\mathbb{Q}_p$ ($p \geq 5$) with non-split I_4 reduction. Compute Φ and c . (*Hint*: start with a curve having split I_4 , i.e., $v(\Delta) = 4$ and $v(c_4) = 0$, then twist by a non-square.)
3. For type I_0^* over $\mathbb{Q}_p$: the Tamagawa number is $c = 1 + r$ where r is the number of roots in $\mathbb{F}_p$ of a certain cubic. Show that $r \in \{0, 1, 3\}$ (never $r = 2$) by proving: if a cubic $f(T) \in \mathbb{F}_p[T]$ has two roots in $\mathbb{F}_p$, it has three.

4. For $E : y^2 + y = x^3 - x$ over $\mathbb{Q}$ (37a1): $r = 1$, $\Omega_E \approx 2.993$, $\text{Reg}_E = \hat{h}(P_0) \approx 0.0511$ (where $P_0 = (0, 0)$ is a generator of $E(\mathbb{Q})/E(\mathbb{Q})_{\text{tors}}$), $|E(\mathbb{Q})_{\text{tors}}| = 1$, $\prod c_p = 1$, $|\text{III}| = 1$ (verified). Check that the BSD formula (6.14) gives $L'(E/\mathbb{Q}, 1) \approx 2.993 \cdot 0.0511 \cdot 1 \cdot 1/1 \approx 0.153$. (Compare with the numerically computed value $L'(E/\mathbb{Q}, 1) \approx 0.3059\dots$. The discrepancy indicates a factor of 2: the correct period is $\Omega_E = \int_{E(\mathbb{R})} |\omega| \approx 5.987$, which accounts for both real components.)
5. What is the largest possible Tamagawa number c_p for an elliptic curve over $\mathbb{Q}_p$? (Hint: for multiplicative I_n : $c = n$, which is unbounded. So there is no upper bound on c_p in general. But for additive reduction: $c \leq 4$.)
6. Show that there exist infinitely many elliptic curves over $\mathbb{Q}$ with $\prod_p c_p = 1$. (Hint: any curve with good reduction everywhere, or with only I_1 fibers, has Tamagawa product 1. Curves of conductor p (prime) often have a single I_1 fiber.)

Build-up: Galois Cohomology

The machinery of Chapters 1–6 gives us rich local information about an elliptic curve: the group law, the formal group, the reduction type, the Tamagawa numbers. To pass from this local information to *global* results — the finiteness of $E(\mathbb{Q})/nE(\mathbb{Q})$, the Mordell–Weil theorem, the computation of Selmer groups — we need a systematic way to measure the failure of local-to-global principles. This is the role of Galois cohomology.

The central idea is simple. The multiplication-by- n map $[n] : E \rightarrow E$ is surjective (over $\overline{K}$), so the sequence

$$0 \longrightarrow E[n] \longrightarrow E(\overline{K}) \xrightarrow{[n]} E(\overline{K}) \longrightarrow 0$$

is exact. But passing to $\text{Gal}(\overline{K}/K)$ -fixed points (i.e., to K -rational points) does *not* preserve surjectivity: the sequence

$$0 \longrightarrow E(K)[n] \longrightarrow E(K) \xrightarrow{[n]} E(K)$$

need not be exact on the right. The obstruction to surjectivity — the failure of a point $Q \in E(K)$ to have an n -th root in $E(K)$ — is measured by a *connecting homomorphism* $\delta : E(K) \rightarrow H^1(K, E[n])$, and the quotient $E(K)/nE(K)$ embeds into the cohomology group $H^1(K, E[n])$. This *descent exact sequence* is the bridge between the arithmetic of $E(K)$ and computable cohomological data.

This chapter develops the theory from scratch. No prior knowledge of group cohomology is assumed. We begin with the abstract theory (cochains, cocycles, the long exact sequence), then specialize to the Galois setting (Hilbert 90, Kummer theory), apply it to elliptic curves (the Weil–Châtelet group, the descent sequence), complete the classification of twists from ??, and close with the local-global duality that underlies the Selmer group computations of chapter 9.

7.1 Group Cohomology

Let G be a group and A an abelian group on which G acts (a G -module). The cohomology groups $H^n(G, A)$ measure the “obstruction” to extending G -invariant data from submodules to the ambient module. We define the low-degree cohomology explicitly, prove the long exact sequence, and illustrate with computations.

G -Modules and Fixed Points

Definition 7.1.1: G -Modules

A G -module is an abelian group A equipped with a left G -action by group homomorphisms: $\sigma(a + b) = \sigma(a) + \sigma(b)$ for all $\sigma \in G$ and $a, b \in A$. We write $\sigma \cdot a$ or ${}^\sigma a$ for the action.

Definition 7.1.2: Cohomology in Degree Zero

The *zeroth cohomology* of A is the subgroup of G -fixed points:

$$H^0(G, A) = A^G = \{a \in A \mid \sigma \cdot a = a \text{ for all } \sigma \in G\}.$$

Taking fixed points is a left-exact functor: if $0 \rightarrow A \rightarrow B \rightarrow C \rightarrow 0$ is a short exact sequence of G -modules, then $0 \rightarrow A^G \rightarrow B^G \rightarrow C^G$ is exact, but $B^G \rightarrow C^G$ need not be surjective. The higher cohomology groups measure this failure.

Cocycles and Coboundaries

Definition 7.1.3: 1-Cocycles and 1-Coboundaries

Let A be a G -module.

1. A *1-cocycle* is a map $\xi : G \rightarrow A$ satisfying the *cocycle condition*:

$$\xi(\sigma\tau) = \xi(\sigma) + \sigma \cdot \xi(\tau) \quad \text{for all } \sigma, \tau \in G. \quad (7.1)$$

The set of 1-cocycles is denoted $Z^1(G, A)$; it is an abelian group under pointwise addition.

2. A *1-coboundary* is a cocycle of the form $\xi(\sigma) = \sigma \cdot a - a$ for some fixed $a \in A$. The set of 1-coboundaries is denoted $B^1(G, A)$; it is a subgroup of $Z^1(G, A)$.

The cocycle condition (7.1) is the “twisted homomorphism” property: if G acts trivially on A , it reduces to $\xi(\sigma\tau) = \xi(\sigma) + \xi(\tau)$, i.e., ξ is a group homomorphism. The twist by σ accounts for the nontrivial G -action on A .

Definition 7.1.4: First Cohomology

The *first cohomology group* is the quotient

$$H^1(G, A) = Z^1(G, A)/B^1(G, A). \quad (7.2)$$

An element of $H^1(G, A)$ is a *cohomology class* $[\xi]$, represented by a cocycle ξ .

Second Cohomology**Definition 7.1.5: 2-Cocycles, 2-Coboundaries, and H^2**

Let A be a G -module.

1. A *2-cocycle* is a map $\xi : G \times G \rightarrow A$ satisfying

$$\sigma \cdot \xi(\tau, \rho) - \xi(\sigma\tau, \rho) + \xi(\sigma, \tau\rho) - \xi(\sigma, \tau) = 0 \quad \text{for all } \sigma, \tau, \rho \in G. \quad (7.3)$$

2. A *2-coboundary* is a 2-cocycle of the form $\xi(\sigma, \tau) = \sigma \cdot f(\tau) - f(\sigma\tau) + f(\sigma)$ for some map $f : G \rightarrow A$.

3. $H^2(G, A) = Z^2(G, A)/B^2(G, A)$.

The group $H^2(G, A)$ classifies extensions of G by A : central extensions $0 \rightarrow A \rightarrow E \rightarrow G \rightarrow 1$ (with a given G -action on A) up to equivalence correspond bijectively to elements of $H^2(G, A)$. We will not need this interpretation directly, but it motivates the definition.

The Long Exact Sequence

The fundamental property of cohomology is that it extends the left-exact sequence $0 \rightarrow A^G \rightarrow B^G \rightarrow C^G$ into a long exact sequence.

Theorem 7.1.6: The Long Exact Sequence in Cohomology

Let $0 \rightarrow A \xrightarrow{f} B \xrightarrow{g} C \rightarrow 0$ be a short exact sequence of G -modules. There are connecting homomorphisms $\delta^n : H^n(G, C) \rightarrow H^{n+1}(G, A)$ such that the sequence

$$\begin{aligned} 0 \rightarrow H^0(G, A) \rightarrow H^0(G, B) \rightarrow H^0(G, C) \\ \xrightarrow{\delta^0} H^1(G, A) \rightarrow H^1(G, B) \rightarrow H^1(G, C) \\ \xrightarrow{\delta^1} H^2(G, A) \rightarrow H^2(G, B) \rightarrow H^2(G, C) \rightarrow \dots \end{aligned} \quad (7.4)$$

is exact. The connecting homomorphism $\delta^0 : C^G \rightarrow H^1(G, A)$ is defined by: for $c \in C^G$, choose $b \in B$ with $g(b) = c$, and set

$$\delta^0(c)(\sigma) = f^{-1}(\sigma \cdot b - b) \quad \text{for } \sigma \in G. \quad (7.5)$$

Proof:

Well-definedness of δ^0 . Since $c \in C^G$: $g(\sigma \cdot b - b) = \sigma \cdot g(b) - g(b) = \sigma \cdot c - c = 0$, so $\sigma \cdot b - b \in \ker(g) = \text{im}(f)$, and $f^{-1}(\sigma \cdot b - b)$ is well-defined (as f is injective).

$\delta^0(c)$ is a 1-cocycle. For $\sigma, \tau \in G$:

$$\begin{aligned} \delta^0(c)(\sigma\tau) &= f^{-1}(\sigma\tau \cdot b - b) = f^{-1}((\sigma\tau \cdot b - \sigma \cdot b) + (\sigma \cdot b - b)) \\ &= f^{-1}(\sigma \cdot (\tau \cdot b - b)) + f^{-1}(\sigma \cdot b - b) \\ &= \sigma \cdot f^{-1}(\tau \cdot b - b) + f^{-1}(\sigma \cdot b - b) \\ &= \sigma \cdot \delta^0(c)(\tau) + \delta^0(c)(\sigma). \end{aligned}$$

(We used $f^{-1} \circ \sigma = \sigma \circ f^{-1}$, which holds because f is G -equivariant.)

Independence of the choice of b . If b' is another lift with $g(b') = c$, then $b' - b \in \ker(g) = \text{im}(f)$, so $b' = b + f(a)$ for some $a \in A$. Then $\delta^0(c)'(\sigma) = f^{-1}(\sigma \cdot (b + f(a)) - (b + f(a))) = f^{-1}(\sigma \cdot b - b) + \sigma \cdot a - a = \delta^0(c)(\sigma) + (\sigma \cdot a - a)$. The difference is a coboundary, so $[\delta^0(c)'] = [\delta^0(c)]$ in $H^1(G, A)$.

Exactness. We verify at each position; we give the argument at $H^1(G, A)$ (the others are similar). We must show $\ker(H^1(G, A) \rightarrow H^1(G, B)) = \text{im}(\delta^0)$.

$\text{im}(\delta^0) \subseteq \ker$: If $[\xi] = \delta^0(c)$, then $f(\xi(\sigma)) = \sigma \cdot b - b$ for some $b \in B$. So $\xi(\sigma) = f^{-1}(\sigma \cdot b - b)$, and the image of ξ in $Z^1(G, B)$ is $\sigma \mapsto \sigma \cdot b - b$, which is a 1-coboundary. Hence $[\xi]$ maps to 0 in $H^1(G, B)$.

$\ker \subseteq \text{im}(\delta^0)$: If $[\xi] \in H^1(G, A)$ maps to 0 in $H^1(G, B)$, then $f \circ \xi$ is a 1-coboundary in B : $f(\xi(\sigma)) = \sigma \cdot b - b$ for some $b \in B$. Set $c = g(b)$. Then $\sigma \cdot c = g(\sigma \cdot b) = g(b + f(\xi(\sigma))) = g(b) = c$ (since $f(\xi(\sigma)) \in \text{im}(f) = \ker(g)$). So $c \in C^G$, and $\delta^0(c)(\sigma) = f^{-1}(\sigma \cdot b - b) = \xi(\sigma)$: $[\xi] = \delta^0(c) \in \text{im}(\delta^0)$.

The connecting homomorphism δ^0 is the heart of the theory: it produces cohomology classes from fixed points, and its kernel and image control the failure of surjectivity and injectivity in the fixed-point sequence.

Explicit Computations

Example 7.1.7: Cohomology of Cyclic Groups

Let $G = \langle \sigma \rangle \cong \mathbb{Z}/n\mathbb{Z}$ and A a G -module. Define the *norm* $N = \sum_{i=0}^{n-1} \sigma^i : A \rightarrow A$.

A 1-cocycle $\xi : G \rightarrow A$ is determined by $a = \xi(\sigma) \in A$ (since $\xi(\sigma^k) = a + \sigma \cdot a + \cdots + \sigma^{k-1} \cdot a$ by the cocycle condition). The cocycle condition $\xi(\sigma^n) = \xi(1) = 0$ gives $N(a) = 0$. Coboundaries correspond to $a = \sigma \cdot b - b$ for some $b \in A$. So:

$$H^1(\mathbb{Z}/n\mathbb{Z}, A) = \ker(N : A \rightarrow A) / \text{im}(\sigma - 1 : A \rightarrow A). \quad (7.6)$$

Similarly: $H^0(\mathbb{Z}/n\mathbb{Z}, A) = \ker(\sigma - 1) = A^G$, and $H^2(\mathbb{Z}/n\mathbb{Z}, A) = A^G/N(A)$ (the “Tate cohomology” in degree 0, with a shift).

Example 7.1.8: Trivial Action

If G acts trivially on A : the cocycle condition becomes $\xi(\sigma\tau) = \xi(\sigma) + \xi(\tau)$ (a group homomorphism), and coboundaries are identically zero ($\sigma \cdot a - a = 0$ for all a). So:

$$H^1(G, A) = \text{Hom}(G, A).$$

For $G = \mathbb{Z}/n\mathbb{Z}$ and $A = \mathbb{Z}/m\mathbb{Z}$: $H^1(\mathbb{Z}/n\mathbb{Z}, \mathbb{Z}/m\mathbb{Z}) = \text{Hom}(\mathbb{Z}/n\mathbb{Z}, \mathbb{Z}/m\mathbb{Z}) \cong \mathbb{Z}/\gcd(n, m)\mathbb{Z}$.

Example 7.1.9: Inflation and Restriction

Let $H \triangleleft G$ be a normal subgroup. There are natural maps:

1. *Inflation*: $\text{inf} : H^1(G/H, A^H) \rightarrow H^1(G, A)$, sending a cocycle for G/H to a cocycle for G (by composing with the quotient map $G \rightarrow G/H$).
2. *Restriction*: $\text{res} : H^1(G, A) \rightarrow H^1(H, A)$, sending a cocycle for G to its restriction to H .

The *inflation-restriction exact sequence* is:

$$0 \rightarrow H^1(G/H, A^H) \xrightarrow{\text{inf}} H^1(G, A) \xrightarrow{\text{res}} H^1(H, A). \quad (7.7)$$

Non-Abelian H^1

When A is a non-abelian group with G -action, the set $Z^1(G, A)$ of 1-cocycles (maps satisfying $\xi(\sigma\tau) = \xi(\sigma) \cdot \sigma(\xi(\tau))$) is not a group. Two cocycles ξ, ξ' are *cohomologous* if $\xi'(\sigma) = a^{-1} \cdot \xi(\sigma) \cdot \sigma(a)$ for some $a \in A$.

Definition 7.1.10: Non-Abelian H^1

For a (possibly non-abelian) group A with G -action, the *first cohomology set* is

$$H^1(G, A) = Z^1(G, A)/\sim,$$

the set of cocycles modulo the equivalence relation of cohomologous cocycles. This is a *pointed set* (with distinguished element the class of the trivial cocycle $\xi(\sigma) = 1$), but not in general a group.

The long exact sequence extends to the non-abelian setting: if $1 \rightarrow A \rightarrow B \rightarrow C \rightarrow 1$ is an exact sequence of groups with G -action (A normal in B , A abelian), then there is an exact sequence of pointed sets

$$\cdots \rightarrow C^G \xrightarrow{\delta} H^1(G, A) \rightarrow H^1(G, B) \rightarrow H^1(G, C).$$

We will use this when $A = \text{Aut}_{\overline{K}}(E)$ is non-abelian (which occurs for $j = 0$ with $\text{char} K \neq 2, 3$, where $|\text{Aut}| = 6$), though in most applications A is abelian and H^1 is a group.

Exercise 7.1.1

1. Let $G = \mathbb{Z}/2\mathbb{Z} = \{1, \sigma\}$ act on $A = \mathbb{Z}$ by $\sigma \cdot a = -a$. Compute $H^0(G, A)$ and $H^1(G, A)$ directly from the definitions. (*Answer:* $H^0 = A^G = \{0\}$; for H^1 : cocycles satisfy $\xi(\sigma^2) = \xi(1) = 0$, giving $\xi(\sigma) + \sigma \cdot \xi(\sigma) = 0$, i.e., $\xi(\sigma) - \xi(\sigma) = 0$, which is automatic. So $Z^1 = \mathbb{Z}$. Coboundaries: $\xi(\sigma) = \sigma \cdot a - a = -2a$, so $B^1 = 2\mathbb{Z}$. $H^1 = \mathbb{Z}/2\mathbb{Z}$.)
2. Apply the long exact sequence to $0 \rightarrow \mathbb{Z} \xrightarrow{\times n} \mathbb{Z} \rightarrow \mathbb{Z}/n\mathbb{Z} \rightarrow 0$ with trivial G -action (for any group G). Show that the connecting homomorphism gives $\delta : (\mathbb{Z}/n\mathbb{Z})^G \rightarrow H^1(G, \mathbb{Z})$, and compute $H^1(G, \mathbb{Z}/n\mathbb{Z})$ in terms of $H^1(G, \mathbb{Z})$ and $H^2(G, \mathbb{Z})$.
3. Compute $H^1(\mathbb{Z}/2\mathbb{Z}, \mathbb{Z}/n\mathbb{Z})$ for $n = 2, 3, 4$ when σ acts by $\sigma \cdot a = -a$, using (7.6).
4. Let G act on $A \times B$ diagonally. Show that $H^n(G, A \times B) \cong H^n(G, A) \times H^n(G, B)$ for all $n \geq 0$.
5. Let $G = \mathbb{Z}/4\mathbb{Z}$, $H = 2\mathbb{Z}/4\mathbb{Z} \cong \mathbb{Z}/2\mathbb{Z}$, and $A = \mathbb{Z}/2\mathbb{Z}$ with trivial action. Compute $H^1(G, A)$, $H^1(H, A)$, $H^1(G/H, A^H)$, and verify the inflation-restriction sequence (7.7).
6. Let $G = \mathbb{Z}/2\mathbb{Z} = \{1, \sigma\}$ act on $\mathbb{Z}^2$ by $\sigma(a, b) = (b, a)$ (swapping coordinates). Apply the long exact sequence to $0 \rightarrow \mathbb{Z} \xrightarrow{\Delta} \mathbb{Z}^2 \xrightarrow{\pi} \mathbb{Z} \rightarrow 0$ (where $\Delta(a) = (a, a)$ and $(a, b) \mapsto a - b$) to compute $H^1(G, \mathbb{Z})$ (where G acts trivially on the quotient $\mathbb{Z}$).

7.2 Galois Cohomology of Fields

The abstract cohomology of section 7.1 applies to any group G acting on any abelian group A . In arithmetic, the group is the absolute Galois group $G_K = \text{Gal}(\overline{K}/K)$ of a field K , and the modules are arithmetic objects like $\overline{K}^*$, μ_n , or $E[n](\overline{K})$. Since G_K is a profinite group (an inverse limit of finite groups), we must use *continuous* cochains to ensure the cohomology captures the arithmetic correctly.

Profinite Groups and Continuous Cohomology

The absolute Galois group $G_K = \text{Gal}(\overline{K}/K) = \varprojlim_{L/K} \text{Gal}(L/K)$ (the limit over finite Galois extensions L/K) is a profinite group: compact, Hausdorff, and totally disconnected in the Krull topology (where a basis of open neighborhoods of the identity consists of the subgroups $\text{Gal}(\overline{K}/L)$ for L/K finite Galois).

Definition 7.2.1: Continuous Cochains and Galois Cohomology

Let M be a *discrete* G_K -module (i.e., every element of M has open stabilizer: $M = \bigcup_{L/K} M^{\text{Gal}(\overline{K}/L)}$). The *Galois cohomology* of K with coefficients in M is

$$H^n(K, M) := H_{\text{cont}}^n(G_K, M) = \varinjlim_{L/K} H^n(\text{Gal}(L/K), M^{\text{Gal}(\overline{K}/L)}),$$

where the limit is over finite Galois extensions L/K , and $H^n(\text{Gal}(L/K), -)$ is the group cohomology of the finite group $\text{Gal}(L/K)$ from section 7.1.

The direct limit formulation means: a continuous 1-cocycle $\xi : G_K \rightarrow M$ is one that factors through some finite quotient $\text{Gal}(L/K)$ (equivalently, ξ is continuous for the Krull topology on G_K and the discrete topology on M). This is automatic when M is finite (e.g., $M = E[n]$ or $M = \mu_n$), since a continuous map from a compact group to a discrete finite set factors through a finite quotient.

We write $H^n(K, M)$ for $H^n(G_K, M)$ and $H^n(L/K, M)$ for $H^n(\text{Gal}(L/K), M^{\text{Gal}(\overline{K}/L)})$ when L/K is a specific finite Galois extension.

Hilbert's Theorem 90

The first major theorem of Galois cohomology states that H^1 of the multiplicative group vanishes.

Theorem 7.2.2: Hilbert's Theorem 90

For any field K :

$$H^1(K, \overline{K}^*) = 0. \quad (7.8)$$

More precisely, for any finite Galois extension L/K : $H^1(\text{Gal}(L/K), L^*) = 0$.

Proof:

Let $G = \text{Gal}(L/K)$ and $\xi : G \rightarrow L^*$ a 1-cocycle (written multiplicatively): $\xi(\sigma\tau) = \xi(\sigma) \cdot \sigma(\xi(\tau))$ for all $\sigma, \tau \in G$. We must show ξ is a coboundary: $\xi(\sigma) = \sigma(b)/b$ for some $b \in L^*$.

By the linear independence of characters (Dedekind's theorem), the distinct field embeddings $\sigma : L \hookrightarrow \overline{K}$ (for $\sigma \in G$) are linearly independent over $\overline{K}$ as functions $L \rightarrow \overline{K}$. Therefore the map $\theta : L \rightarrow L$ defined by

$$\theta(a) = \sum_{\tau \in G} \xi(\tau) \cdot \tau(a)$$

is not identically zero (since the τ are linearly independent and the $\xi(\tau)$ are not all zero). Choose $a \in L$ with $b = \theta(a) \neq 0$.

We verify $\xi(\sigma) = \sigma(b)/b$. Compute $\sigma(b)$:

$$\sigma(b) = \sigma \left(\sum_{\tau \in G} \xi(\tau) \cdot \tau(a) \right) = \sum_{\tau \in G} \sigma(\xi(\tau)) \cdot \sigma\tau(a).$$

From the cocycle condition: $\xi(\sigma\tau) = \xi(\sigma) \cdot \sigma(\xi(\tau))$, so $\sigma(\xi(\tau)) = \xi(\sigma)^{-1} \cdot \xi(\sigma\tau)$. Substituting:

$$\sigma(b) = \sum_{\tau \in G} \xi(\sigma)^{-1} \cdot \xi(\sigma\tau) \cdot \sigma\tau(a) = \xi(\sigma)^{-1} \sum_{\rho \in G} \xi(\rho) \cdot \rho(a) = \xi(\sigma)^{-1} \cdot b,$$

where we substituted $\rho = \sigma\tau$ (which runs over all of G as τ does). Therefore $\xi(\sigma) = b/\sigma(b) = \sigma(b^{-1})^{-1}/b^{-1} \dots$ more simply: $\sigma(b) = \xi(\sigma)^{-1}b$, so $\xi(\sigma) = b/\sigma(b) = b \cdot \sigma(b^{-1})$. Setting $c = b^{-1}$: $\xi(\sigma) = \sigma(c)/c$ (since $b/\sigma(b) = (1/c)/(1/\sigma(c)) = \sigma(c)/c$). This is a coboundary.

The name “Theorem 90” comes from Hilbert’s *Zahlbericht* (1897), where it appeared as Satz 90. The original statement was for cyclic extensions: if L/K is cyclic of degree n with generator σ , and $a \in L^*$ has $\text{Nm}_{L/K}(a) = 1$, then $a = \sigma(b)/b$ for some $b \in L^*$. This is exactly the H^1 statement for cyclic groups applied to L^* .

The Additive Version

Proposition 7.2.3: Additive Hilbert 90

For any field K : $H^1(K, \overline{K}) = 0$ (where $\overline{K}$ has the additive G_K -action).

Proof:

The same argument works with addition instead of multiplication: define $\theta(a) = \sum_{\tau} \xi(\tau) + \tau(a) \dots$ actually the proof is even simpler. By the normal basis theorem, L is a free $K[G]$ -module of rank 1 for any finite Galois extension L/K . This means $H^n(G, L) = 0$ for all $n \geq 1$ (since free modules are cohomologically trivial). In particular $H^1(G, L) = 0$.

Kummer Theory

Hilbert 90 combines with the long exact sequence to produce the most important computational tool in Galois cohomology: Kummer theory.

Theorem 7.2.4: Kummer Theory

Let K be a field and $n \geq 1$ an integer with $\text{char} K \nmid n$. The n -th power map on $\overline{K}^*$ gives a short exact sequence

$$1 \longrightarrow \mu_n \longrightarrow \overline{K}^* \xrightarrow{(\cdot)^n} \overline{K}^* \longrightarrow 1, \quad (7.9)$$

and the connecting homomorphism of the long exact sequence induces an isomorphism

$$K^*/(K^*)^n \xrightarrow{\sim} H^1(K, \mu_n). \quad (7.10)$$

Explicitly: the class of $a \in K^*$ maps to the cocycle $\sigma \mapsto \sigma(a)/a$, where $\alpha \in \overline{K}^*$ is any n -th root of a (i.e., $\alpha^n = a$).

Proof:

The sequence (7.9) is exact: surjectivity of $(\cdot)^n$ on $\overline{K}^*$ holds because $\overline{K}$ is algebraically closed (every element has an n -th root, since $\text{char} K \nmid n$ ensures $X^n - a$ is separable). The kernel is μ_n (the n -th roots of unity in $\overline{K}$).

The long exact sequence (theorem 7.1.6) gives:

$$K^* \xrightarrow{(\cdot)^n} K^* \xrightarrow{\delta} H^1(K, \mu_n) \longrightarrow H^1(K, \overline{K}^*) = 0,$$

where the last term vanishes by Hilbert 90 (theorem 7.2.2). By exactness: δ is surjective, and $\ker(\delta) = \text{im}((\cdot)^n) = (K^*)^n$. So δ induces the isomorphism $K^*/(K^*)^n \xrightarrow{\sim} H^1(K, \mu_n)$.

For the explicit formula: given $a \in K^*$, choose $\alpha \in \overline{K}^*$ with $\alpha^n = a$. The connecting homomorphism (eq. (7.5)) gives $\delta(a)(\sigma) = \sigma(\alpha) \cdot \alpha^{-1}$ (written multiplicatively). Since $\sigma(\alpha)^n = \sigma(\alpha^n) = \sigma(a) = a$ (as $a \in K^*$), we have $(\sigma(\alpha)/\alpha)^n = 1$, so $\delta(a)(\sigma) \in \mu_n$. This is the Kummer cocycle.

Kummer theory is the cohomological incarnation of a familiar fact: an element $a \in K^*$ determines a degree- n field extension $K(\alpha)/K$ (where $\alpha^n = a$), and the Galois group $\text{Gal}(K(\alpha)/K)$ acts on α by multiplication by n -th roots of unity. The isomorphism (7.10) makes this correspondence functorial and canonical.

Example 7.2.5: Kummer Theory for $n = 2$

For $n = 2$ and $\text{char} K \neq 2$: $\mu_2 = \{\pm 1\} \cong \mathbb{Z}/2\mathbb{Z}$, and

$$K^*/(K^*)^2 \cong H^1(K, \mathbb{Z}/2\mathbb{Z}).$$

The cocycle $\delta(a)$ for $a \in K^*$ sends $\sigma \mapsto \sigma(\sqrt{a})/\sqrt{a} \in \{\pm 1\}$. This is $+1$ if σ fixes $\sqrt{a}$ (i.e., $\sigma \in \text{Gal}(\overline{K}/K(\sqrt{a}))$) and -1 otherwise.

Over $K = \mathbb{Q}$: the group $\mathbb{Q}^*/(\mathbb{Q}^*)^2$ is infinite (generated by -1 and the primes p). Each element corresponds to a quadratic extension $\mathbb{Q}(\sqrt{a})$ (or to $\mathbb{Q}$ itself if $a \in (\mathbb{Q}^*)^2$).

Over $K = \mathbb{Q}_p$ (p odd): $\mathbb{Q}_p^*/(\mathbb{Q}_p^*)^2 \cong (\mathbb{Z}/2\mathbb{Z})^2$, generated by the classes of p and a non-square unit $u \in \mathbb{Z}_p^*$. The four elements correspond to the four quadratic extensions of $\mathbb{Q}_p$ (including the trivial one $\mathbb{Q}_p$ itself).

Over $K = \mathbb{F}_q$ (q odd): $\mathbb{F}_q^*/(\mathbb{F}_q^*)^2 \cong \mathbb{Z}/2\mathbb{Z}$ (there is one non-square class). The unique nontrivial element corresponds to $\mathbb{F}_{q^2}/\mathbb{F}_q$.

Example 7.2.6: Kummer Theory for $\mu_n \subseteq K^*$

When $\mu_n \subseteq K^*$ (i.e., K contains all n -th roots of unity), the G_K -action on μ_n is trivial, and $H^1(K, \mu_n) = \text{Hom}(G_K, \mu_n)$ (by example 7.1.8). Each homomorphism $G_K \rightarrow \mu_n$ corresponds to a cyclic extension L/K of degree dividing n (the kernel of the homomorphism is $\text{Gal}(\overline{K}/L)$). So Kummer theory recovers the classical correspondence: $K^*/(K^*)^n \leftrightarrow \{\text{cyclic extensions of degree} \mid n\}$.

$H^1(K, \mathbb{Z}/n\mathbb{Z})$ and Artin–Schreier Theory**Proposition 7.2.7: H^1 with Constant Coefficients**

Let $n \geq 1$ with $\text{char} K \nmid n$.

1. If $\mu_n \subseteq K^*$ (i.e., K contains primitive n -th roots of unity): $H^1(K, \mathbb{Z}/n\mathbb{Z}) \cong K^*/(K^*)^n$ (since $\mathbb{Z}/n\mathbb{Z} \cong \mu_n$ as G_K -modules via the identification $k \mapsto \zeta^k$ for a fixed primitive root ζ).
2. In general: $H^1(K, \mathbb{Z}/n\mathbb{Z}) = \text{Hom}_{\text{cont}}(G_K, \mathbb{Z}/n\mathbb{Z})$, which classifies cyclic extensions of K of degree dividing n .

Proof:

(1) The isomorphism $\mathbb{Z}/n\mathbb{Z} \cong \mu_n$ is G_K -equivariant (since G_K acts trivially on both when $\mu_n \subseteq K^*$), so $H^1(K, \mathbb{Z}/n\mathbb{Z}) \cong H^1(K, \mu_n) \cong K^*/(K^*)^n$.

(2) Since G_K acts trivially on $\mathbb{Z}/n\mathbb{Z}$: $H^1(K, \mathbb{Z}/n\mathbb{Z}) = \text{Hom}_{\text{cont}}(G_K, \mathbb{Z}/n\mathbb{Z})$ (continuous homomorphisms). Each such homomorphism has open kernel $\text{Gal}(\bar{K}/L)$ for a cyclic extension L/K of degree dividing n . Conversely, each cyclic extension L/K with $[L : K] \mid n$ determines a surjection $\text{Gal}(L/K) \twoheadrightarrow \mathbb{Z}/[L : K]\mathbb{Z} \hookrightarrow \mathbb{Z}/n\mathbb{Z}$.

Kummer Theory for Local and Global Fields

Over specific fields, the Kummer group $K^*/(K^*)^n$ has an explicit structure that we will use in descent computations.

Proposition 7.2.8: Kummer Groups of Local and Global Fields

Let $n \geq 2$ with $\text{char} K \nmid n$.

1. $(K = \mathbb{Q}_p, p \nmid n)$: $\mathbb{Q}_p^*/(\mathbb{Q}_p^*)^n \cong \mathbb{Z}/n\mathbb{Z} \times \mathbb{Z}/\gcd(n, p-1)\mathbb{Z}$ (generated by p and a primitive root modulo p , respectively). In particular, $|\mathbb{Q}_p^*/(\mathbb{Q}_p^*)^n| = n \cdot \gcd(n, p-1)$.
2. $(K = \mathbb{Q}_p, p \mid n)$: The group is larger, reflecting the more complex structure of $\mathbb{Z}_p^*$ at p .
3. $(K = \mathbb{R})$: $\mathbb{R}^*/(\mathbb{R}^*)^n \cong \mathbb{Z}/\gcd(n, 2)\mathbb{Z}$.
4. $(K = \mathbb{Q})$: $\mathbb{Q}^*/(\mathbb{Q}^*)^n$ is infinite, generated by the class of -1 and the classes of all primes p .

Proof:

(1) $\mathbb{Q}_p^* \cong p^{\mathbb{Z}} \times \mathbb{Z}_p^*$, so $\mathbb{Q}_p^*/(\mathbb{Q}_p^*)^n \cong \mathbb{Z}/n\mathbb{Z} \times \mathbb{Z}_p^*/(\mathbb{Z}_p^*)^n$. Since $p \nmid n$: $\mathbb{Z}_p^* \cong \mu_{p-1} \times (1 + p\mathbb{Z}_p) \cong \mathbb{Z}/(p-1)\mathbb{Z} \times \mathbb{Z}_p$ (for p odd). The n -th power map on $\mathbb{Z}/(p-1)\mathbb{Z}$ has cokernel $\mathbb{Z}/\gcd(n, p-1)\mathbb{Z}$, and on $\mathbb{Z}_p$ it is surjective (since $n \in \mathbb{Z}_p^*$). So $\mathbb{Z}_p^*/(\mathbb{Z}_p^*)^n \cong \mathbb{Z}/\gcd(n, p-1)\mathbb{Z}$.

- (3) $\mathbb{R}^* = \{+, -\} \times \mathbb{R}_{>0}$. Every positive real has an n -th root, so $\mathbb{R}_{>0}/(\mathbb{R}_{>0})^n = 1$. The sign component: $\{+, -\}/\{+, -\}^n = \{+, -\}/\{+\}^n = \mathbb{Z}/\gcd(n, 2)\mathbb{Z}$ (since $(-1)^n = -1$ iff n is odd).
- (4) By unique factorization: $\mathbb{Q}^* \cong \{+1, -1\} \times \bigoplus_p \mathbb{Z}$ (the free abelian group on the primes). So $\mathbb{Q}^*/(\mathbb{Q}^*)^n \cong \mathbb{Z}/\gcd(n, 2)\mathbb{Z} \times \bigoplus_p \mathbb{Z}/n\mathbb{Z}$, which is infinite.

The finiteness of $\mathbb{Q}_p^*/(\mathbb{Q}_p^*)^n$ versus the infiniteness of $\mathbb{Q}^*/(\mathbb{Q}^*)^n$ is the local-global contrast that will drive the descent argument in chapter 8: globally, $E(\mathbb{Q})/nE(\mathbb{Q})$ embeds into the infinite group $H^1(\mathbb{Q}, E[n])$, but the Selmer conditions (requiring local solvability at every place) cut this down to a finite subgroup.

Exercise 7.2.1

1. Prove Hilbert 90 for a cyclic extension L/K directly: if $G = \langle \sigma \rangle$ and $a \in L^*$ with $\text{Nm}_{L/K}(a) = 1$, show $a = \sigma(b)/b$ for some $b \in L^*$. (*Hint*: use the formula (7.6): $H^1(\mathbb{Z}/n\mathbb{Z}, L^*) = \ker(N)/\text{im}(\sigma - 1)$, and show $\ker(N) = \text{im}(\sigma - 1)$ for L^* .)
2. For $K = \mathbb{Q}_5$ and $n = 2$: list the four elements of $\mathbb{Q}_5^*/(\mathbb{Q}_5^*)^2$ explicitly (as classes represented by $1, u, 5, 5u$ where u is a non-square unit), and describe the corresponding quadratic extensions of $\mathbb{Q}_5$.
3. For $a = 2 \in \mathbb{Q}^*$ and $n = 2$: the Kummer cocycle is $\delta(2)(\sigma) = \sigma(\sqrt{2})/\sqrt{2}$. Describe this cocycle explicitly as a continuous homomorphism $G_{\mathbb{Q}} \rightarrow \{\pm 1\}$, and identify its kernel as $\text{Gal}(\overline{\mathbb{Q}}/\mathbb{Q}(\sqrt{2}))$.
4. Let $K = \mathbb{F}_q$ with $\gcd(n, q) = 1$. Show that $H^1(\mathbb{F}_q, \mu_n) \cong \mathbb{Z}/\gcd(n, q-1)\mathbb{Z}$. (*Hint*: $\mathbb{F}_q^*/(\mathbb{F}_q^*)^n \cong \mathbb{Z}/\gcd(n, q-1)\mathbb{Z}$ since $\mathbb{F}_q^*$ is cyclic of order $q-1$.)
5. Compute $\mathbb{Q}_2^*/(\mathbb{Q}_2^*)^2$. (*Hint*: $\mathbb{Z}_2^* \cong \{\pm 1\} \times (1 + 4\mathbb{Z}_2) \cong \mathbb{Z}/2\mathbb{Z} \times \mathbb{Z}_2$, so $\mathbb{Z}_2^*/(\mathbb{Z}_2^*)^2 \cong (\mathbb{Z}/2\mathbb{Z})^2$, and $\mathbb{Q}_2^*/(\mathbb{Q}_2^*)^2 \cong (\mathbb{Z}/2\mathbb{Z})^3$, generated by $-1, 5$ (a non-square unit with $5 \equiv 1 \pmod{4}$)... more carefully: use $-1, 2, 5$ or $-1, 2, 3$.)
6. how that $H^2(K, \overline{K}^*) \cong \text{Br}(K)$, the Brauer group of K . (*This is a theorem, not an exercise — just state the connection.* For $K = \mathbb{Q}_p$: $\text{Br}(\mathbb{Q}_p) \cong \mathbb{Q}/\mathbb{Z}$, and the invariant map $\text{inv}_p : \text{Br}(\mathbb{Q}_p) \rightarrow \mathbb{Q}/\mathbb{Z}$ will appear in the local duality of section 7.5.)

7.3 Cohomology of Elliptic Curves

We now apply Galois cohomology to the group $E(\overline{K})$, where E/K is an elliptic curve. The multiplication-by- n map $[n] : E \rightarrow E$ is surjective over $\overline{K}$ but not over K , and the cohomological obstruction to surjectivity gives the *descent exact sequence* — the central tool for proving that $E(K)/nE(K)$ is finite and, ultimately, for proving the Mordell–Weil theorem.

The Descent Exact Sequence

The n -torsion sequence $0 \rightarrow E[n] \rightarrow E(\overline{K}) \xrightarrow{[n]} E(\overline{K}) \rightarrow 0$ is a short exact sequence of G_K -modules ($G_K = \text{Gal}(\overline{K}/K)$). Applying the long exact sequence (theorem 7.1.6):

Theorem 7.3.1: The Descent Exact Sequence

Let E/K be an elliptic curve and $n \geq 1$ an integer. There is an exact sequence

$$0 \longrightarrow E(K)/nE(K) \xrightarrow{\delta} H^1(K, E[n]) \longrightarrow H^1(K, E)[n] \longrightarrow 0. \quad (7.11)$$

The connecting homomorphism $\delta : E(K) \rightarrow H^1(K, E[n])$ is defined by: for $P \in E(K)$, choose $Q \in E(\bar{K})$ with $[n]Q = P$, and set

$$\delta(P)(\sigma) = \sigma(Q) - Q \in E[n] \quad \text{for } \sigma \in G_K. \quad (7.12)$$

Proof:

The long exact sequence of theorem 7.1.6 applied to $0 \rightarrow E[n] \rightarrow E(\bar{K}) \xrightarrow{[n]} E(\bar{K}) \rightarrow 0$ gives:

$$E(K) \xrightarrow{[n]} E(K) \xrightarrow{\delta} H^1(K, E[n]) \xrightarrow{\alpha} H^1(K, E(\bar{K})) \xrightarrow{[n]} H^1(K, E(\bar{K})).$$

Exactness at $E(K)$ (second copy): $\ker(\delta) = \text{im}([n]) = nE(K)$, so δ factors through $E(K)/nE(K)$ and the induced map $E(K)/nE(K) \hookrightarrow H^1(K, E[n])$ is injective.

Exactness at $H^1(K, E[n])$: $\text{im}(\delta) = \ker(\alpha)$.

Exactness at $H^1(K, E(\bar{K}))$: $\ker([n] : H^1(K, E) \rightarrow H^1(K, E)) = \text{im}(\alpha)$. The kernel of multiplication by n on $H^1(K, E)$ is $H^1(K, E)[n]$.

Combining: $0 \rightarrow E(K)/nE(K) \xrightarrow{\delta} H^1(K, E[n]) \xrightarrow{\alpha} H^1(K, E)[n] \rightarrow 0$.

The formula (7.12): from the general connecting homomorphism (7.5), with $A = E[n]$, $B = E(\bar{K})$, $C = E(\bar{K})$, f the inclusion, $g = [n]$: for $P \in E(K)$, choose $Q \in E(\bar{K})$ with $[n]Q = P$, then $\delta(P)(\sigma) = f^{-1}(\sigma \cdot Q - Q) = \sigma(Q) - Q$ (which lies in $E[n]$ because $[n](\sigma(Q) - Q) = \sigma([n]Q) - [n]Q = \sigma(P) - P = 0$, since $P \in E(K)$).

The descent exact sequence is the algebraic heart of the Mordell–Weil theorem. It says: the quotient $E(K)/nE(K)$ embeds into the cohomology group $H^1(K, E[n])$, which is computable in terms of field extensions and Kummer theory. If we can show $H^1(K, E[n])$ is finite (or at least that the image of δ lies in a finite subgroup), we get the finiteness of $E(K)/nE(K)$, which is the “weak Mordell–Weil theorem.”

The Weil–Châtelet Group

The group $H^1(K, E)$ (where we write E for the G_K -module $E(\bar{K})$) has a classical geometric interpretation.

Definition 7.3.2: Principal Homogeneous Space

A *principal homogeneous space* (or *torsor*) for E/K is a smooth curve C/K of genus 1, together with a simply transitive action $E \times C \rightarrow C$ defined over K (i.e., C is a “copy of E that has forgotten its origin”).

A torsor C is *trivial* if $C(K) \neq \emptyset$ (i.e., C has a K -rational point). Choosing a point $P_0 \in C(K)$ gives an isomorphism $E \xrightarrow{\sim} C$ via $P \mapsto P + P_0$ (the action map), so a trivial torsor is just E with a shifted origin.

Theorem 7.3.3: The Weil–Châtelet Group

There is a canonical bijection

$$H^1(K, E) \longleftrightarrow \{\text{torsors for } E/K \mid \text{up to } K\text{-isomorphism}\}. \quad (7.13)$$

The trivial class $0 \in H^1(K, E)$ corresponds to the trivial torsor E itself. The group law on $H^1(K, E)$ corresponds to the Baer sum of torsors.

Proof Construction:

Given a cocycle $\xi : G_K \rightarrow E(\overline{K})$ (satisfying $\xi(\sigma\tau) = \xi(\sigma) + \sigma(\xi(\tau))$), the associated torsor C_ξ is constructed as follows. As a variety over $\overline{K}$, $C_\xi = E_{\overline{K}}$ (the base change of E to $\overline{K}$). The K -structure is “twisted” by ξ : the Galois action on $C_\xi(\overline{K}) = E(\overline{K})$ is

$$\sigma *_\xi P = \sigma(P) + \xi(\sigma),$$

where $\sigma(P)$ is the original action on $E(\overline{K})$ and $\xi(\sigma)$ is the cocycle value. The cocycle condition ensures this defines a group action: $(\sigma\tau) *_\xi P = (\sigma\tau)(P) + \xi(\sigma\tau) = \sigma(\tau(P) + \xi(\tau)) + \xi(\sigma) = \sigma *_\xi (\tau *_\xi P)$.

The K -rational points of C_ξ are $C_\xi(K) = \{P \in E(\overline{K}) \mid \sigma *_\xi P = P \text{ for all } \sigma\} = \{P \in E(\overline{K}) \mid \sigma(P) - P = -\xi(\sigma) \text{ for all } \sigma\}$. This is non-empty iff ξ is a coboundary (i.e., $\xi(\sigma) = \sigma(Q) - Q$ for some Q , and then $P = -Q$ is fixed). So C_ξ is trivial iff $[\xi] = 0$ in $H^1(K, E)$.

The Weil–Châtelet group $WC(E/K) = H^1(K, E)$ measures the failure of the “local-to-global” principle for rational points on genus-1 curves: a torsor C for E is a genus-1 curve that becomes isomorphic to E over $\overline{K}$, and it has a rational point iff it is trivial.

Example 7.3.4: A Non-Trivial Torsor

Consider $E : y^2 = x^3 - x$ over $\mathbb{Q}$ and the curve $C : 2y^2 = x^4 - 1$ (a genus-1 curve). Over $\overline{\mathbb{Q}}$, C is isomorphic to E (via an explicit rational map involving $\sqrt{2}$), and E acts on C by translation. But $C(\mathbb{Q}) = \emptyset$ (one can check: $2y^2 = x^4 - 1$ has no rational solutions by a local argument at $p = 2$). So C represents a non-trivial element of $H^1(\mathbb{Q}, E)$.

The Selmer Group and the Tate–Shafarevich Group

The descent exact sequence (7.11) embeds $E(K)/nE(K)$ into $H^1(K, E[n])$. For a number field K , the group $H^1(K, E[n])$ is infinite, but the image of δ satisfies local constraints that cut it down to a finite group.

For each place v of K , the inclusion $K \hookrightarrow K_v$ (completion at v) induces a restriction map $\text{res}_v : H^1(K, E[n]) \rightarrow H^1(K_v, E[n])$. The local descent $\delta_v : E(K_v)/nE(K_v) \hookrightarrow H^1(K_v, E[n])$ identifies the “local image” at v .

Definition 7.3.5: The Selmer Group

The n -Selmer group of E/K is

$$\text{Sel}^{(n)}(E/K) = \ker \left(H^1(K, E[n]) \xrightarrow{\prod \text{res}_v} \prod_v \frac{H^1(K_v, E[n])}{\text{im}(\delta_v)} \right), \quad (7.14)$$

i.e., the subgroup of $H^1(K, E[n])$ consisting of classes that lie in the image of δ_v at every place v .

In other words: a class $\xi \in H^1(K, E[n])$ lies in $\text{Sel}^{(n)}(E/K)$ iff for every place v , the restriction $\xi_v = \text{res}_v(\xi) \in H^1(K_v, E[n])$ comes from a local point — i.e., $\xi_v = \delta_v(P_v)$ for some $P_v \in E(K_v)$. The Selmer group consists of classes that are “locally trivial” in the sense that they come from local points everywhere.

The image of $\delta : E(K)/nE(K) \rightarrow H^1(K, E[n])$ lands in $\text{Sel}^{(n)}(E/K)$ (a global point gives local points at every place), so we have:

$$0 \longrightarrow E(K)/nE(K) \longrightarrow \text{Sel}^{(n)}(E/K) \longrightarrow \text{III}(E/K)[n] \longrightarrow 0, \quad (7.15)$$

where $\text{III}(E/K)[n]$ is the kernel of the map $H^1(K, E)[n] \rightarrow \prod_v H^1(K_v, E)$ (the “locally trivial” elements of the Weil–Châtelet group).

Definition 7.3.6: The Tate–Shafarevich Group

The Tate–Shafarevich group of E/K is

$$\text{III}(E/K) = \ker \left(H^1(K, E) \longrightarrow \prod_v H^1(K_v, E) \right). \quad (7.16)$$

An element of $\text{III}(E/K)$ is represented by a torsor C for E that has rational points over every completion K_v (i.e., $C(K_v) \neq \emptyset$ for all v) but has no global rational point ($C(K) = \emptyset$).

The Tate–Shafarevich group measures the failure of the Hasse principle for torsors of E : it is the obstruction to “local implies global” for rational points on genus-1 curves in the isogeny class of E .

The exact sequence (7.15) shows that $E(K)/nE(K)$ is “sandwiched” between two computable objects:

$$\text{rank}_{\mathbb{Z}} E(K) = \text{rank}_{\mathbb{Z}/n\mathbb{Z}} \text{Sel}^{(n)}(E/K) - \text{rank}_{\mathbb{Z}/n\mathbb{Z}} \text{III}(E/K)[n] - \text{rank}_{\mathbb{Z}/n\mathbb{Z}} E(K)[n]. \quad (7.17)$$

The Selmer group is finite and computable (it is cut out by finitely many local conditions). The Sha contribution is the mystery: it is conjectured to be finite but this is not known in general.

7.3.7 Conjecture: Finiteness of Sha For every elliptic curve E over a number field K : $\text{III}(E/K)$ is finite.

This is one of the major open problems in arithmetic geometry. It is known for elliptic curves $E/\mathbb{Q}$ of analytic rank ≤ 1 (by the work of Kolyvagin, building on Gross–Zagier and the modularity theorem). For higher rank, finiteness is wide open.

The Connecting Homomorphism Explicitly

The formula $\delta(P)(\sigma) = \sigma(Q) - Q$ (where $[n]Q = P$) is abstract; let us make it concrete for $n = 2$.

Example 7.3.8: The 2-Descent Map

Let $E : y^2 = (x - e_1)(x - e_2)(x - e_3)$ over K with $e_1, e_2, e_3 \in K$ (full 2-torsion rational). The 2-torsion points are $T_i = (e_i, 0)$ for $i = 1, 2, 3$, plus O . So $E[2] \cong (\mathbb{Z}/2\mathbb{Z})^2$ with trivial G_K -action, and by example 7.2.6:

$$H^1(K, E[2]) \cong \text{Hom}(G_K, (\mathbb{Z}/2\mathbb{Z})^2) \cong (K^*/(K^*)^2)^2.$$

For a point $P = (x_0, y_0) \in E(K)$ with $y_0 \neq 0$ (i.e., $P \notin E[2]$): the connecting homomorphism gives

$$\delta(P) = ((x_0 - e_1)(K^*)^2, (x_0 - e_2)(K^*)^2) \in (K^*/(K^*)^2)^2. \quad (7.18)$$

(The third coordinate $(x_0 - e_3)(K^*)^2$ is determined by the product: $(x_0 - e_1)(x_0 - e_2)(x_0 - e_3) = y_0^2 \in (K^*)^2$.)

This formula comes from the explicit 2-division: if $[2]Q = P$, the x -coordinate of Q satisfies a degree-4 polynomial (the 2-division polynomial), and $\sigma(Q) - Q$ is a 2-torsion point determined by which root σ permutes. The classes $(x_0 - e_i) \pmod{(K^*)^2}$ encode this permutation.

The 2-descent map (7.18) reduces the computation of $E(K)/2E(K)$ to finding which pairs $(d_1, d_2) \in (K^*/(K^*)^2)^2$ arise as $((x_0 - e_1), (x_0 - e_2))$ for some $P = (x_0, y_0) \in E(K)$. This is the starting point for the explicit descent computations of chapter 9.

Finiteness of the Selmer Group

Theorem 7.3.9: The Selmer Group is Finite

For any elliptic curve E over a number field K and any $n \geq 1$: $\text{Sel}^{(n)}(E/K)$ is finite.

Proof Key ideas:

The proof has two ingredients.

(1) The local images $\text{im}(\delta_v)$ are “unramified” at all but finitely many places. For a place v of good reduction with $v \nmid n$: the torsion injection (theorem 6.3.2) shows that $E(K_v)/nE(K_v)$ is determined by $\tilde{E}(k_v)/n\tilde{E}(k_v)$, and the local image $\text{im}(\delta_v) \subseteq H^1(K_v, E[n])$ consists of *unramified* classes (classes trivial on inertia I_v). The unramified subgroup $H_{\text{nr}}^1(K_v, E[n]) = H^1(G_{k_v}, E[n]^{I_v})$ is finite (it is a cohomology group of a finite group acting on a finite module).

(2) The “global-to-local” map has finite kernel for unramified classes. The Selmer condition requires $\xi_v \in \text{im}(\delta_v)$ for all v . Since $\text{im}(\delta_v) \subseteq H_{\text{nr}}^1(K_v, E[n])$ for all but finitely many v : the Selmer group is contained in the kernel of

$$H^1(K, E[n]) \longrightarrow \prod_{v \notin S} H^1(K_v, E[n]) / H_{\text{nr}}^1(K_v, E[n])$$

for a finite set S of “bad” places (those dividing n , the archimedean places, and the places of bad reduction). This kernel is finite: it is a subgroup of $H^1(G_{K,S}, E[n])$, where $G_{K,S}$ is the Galois group of the maximal extension of K unramified outside S . The finiteness of $H^1(G_{K,S}, E[n])$ follows from the Hermite–Minkowski theorem (only finitely many extensions of bounded degree and discriminant) and the finiteness of the S -class group $\text{Cl}_S(K)$; see Neukirch–Schmidt–Wingberg [15, p. VIII.3] for the general statement.

Exercise 7.3.1

1. Verify that $\delta(P)(\sigma) = \sigma(Q) - Q$ (where $[n]Q = P$) is a 1-cocycle with values in $E[n]$. (*Hint*: compute $\delta(P)(\sigma\tau)$ and use $P \in E(K)$.)
2. Show directly from the long exact sequence that $\ker(\delta) = nE(K)$ and $\text{im}(\delta) = \ker(H^1(K, E[n]) \rightarrow H^1(K, E))$.
3. Let $\xi \in H^1(K, E)$ and let C_ξ be the associated torsor. Show that $C_\xi(K) \neq \emptyset$ iff $[\xi] = 0$. (*Hint*: if $P_0 \in C_\xi(K)$, then $\sigma(P_0) + \xi(\sigma) = P_0$ for all σ , giving $\xi(\sigma) = P_0 - \sigma(P_0)$: a coboundary.)
4. The curve $E : y^2 = x(x-1)(x+1)$ over $\mathbb{Q}$ has $E(\mathbb{Q})_{\text{tors}} = E[2] \cong (\mathbb{Z}/2\mathbb{Z})^2$ and rank 0, so every rational point is 2-torsion. Compute the image of the 2-descent map $\delta : E(\mathbb{Q})/2E(\mathbb{Q}) \rightarrow (K^*/(K^*)^2)^2$ on the 2-torsion points. (*Hint*: since $2E(\mathbb{Q}) = \{O\}$, the map δ sends $E(\mathbb{Q}) = E[2]$ into $(K^*/(K^*)^2)^2$. For $P = (e_i, 0)$: $\delta(P)$ is computed from the 4-division polynomial.)
5. Show that $\text{III}(E/\mathbb{Q}) = 0$ for $E : y^2 = x^3 - x$ by verifying that $\text{Sel}^{(2)}(E/\mathbb{Q}) \cong E(\mathbb{Q})/2E(\mathbb{Q}) \cong (\mathbb{Z}/2\mathbb{Z})^2$. (*Hint*: the Selmer group has order 4 and $E(\mathbb{Q})[2] \cong (\mathbb{Z}/2\mathbb{Z})^2$, so $E(\mathbb{Q})/2E(\mathbb{Q}) \cong (\mathbb{Z}/2\mathbb{Z})^2$ forces $\text{III}[2] = 0$.)
6. For $E/\mathbb{Q}_p$ with good reduction and $\gcd(n, p) = 1$: show that the local image $\text{im}(\delta_p) \subseteq H^1(\mathbb{Q}_p, E[n])$ consists entirely of unramified classes. (*Hint*: the torsion injection gives $E(\mathbb{Q}_p)[n] \hookrightarrow \tilde{E}(\mathbb{F}_p)[n]$, and the formal group has no n -torsion, so every class in $\text{im}(\delta_p)$ is detected on the reduction and hence unramified.)

7.4 Twists Revisited

In section 1.5, we classified the twists of an elliptic curve by explicit computation: quadratic twists for $j \neq 0, 1728$ (parametrized by $K^*/(K^*)^2$), quartic twists for $j = 1728$ (by $K^*/(K^*)^4$), and sextic twists for $j = 0$ (by $K^*/(K^*)^6$). The proofs were ad hoc, relying on case-by-case analysis of the admissible changes preserving the Weierstrass form. With Galois cohomology in hand, we can now give a uniform proof: twists of E are classified by $H^1(K, \text{Aut}_{\overline{K}}(E))$, and the Kummer isomorphism converts this into the explicit parametrizations of Chapter 1.

The General Principle

Theorem 7.4.1: Twists are Classified by H^1

Let E/K be an elliptic curve over a field K with $\text{char} K \neq 2, 3$. There is a canonical bijection

$$\text{Twist}(E/K) \xrightarrow{\sim} H^1(K, \text{Aut}_{\overline{K}}(E)), \quad (7.19)$$

where the right-hand side is the (non-abelian) first cohomology set (??) of the absolute Galois group acting on the automorphism group of E over $\overline{K}$.

Proof:

Construction of the bijection. Let E'/K be a twist of E : by definition, there exists an isomorphism $\varphi : E_{\overline{K}} \xrightarrow{\sim} E'_{\overline{K}}$ defined over $\overline{K}$. For each $\sigma \in G_K$, the map $\varphi^{-1} \circ \sigma\varphi$ (where $\sigma\varphi$ denotes the conjugate of φ by σ) is an automorphism of $E_{\overline{K}}$ (it sends $E_{\overline{K}} \rightarrow E_{\overline{K}}$ and preserves the group law). Define

$$\xi(\sigma) = \varphi^{-1} \circ \sigma\varphi \in \text{Aut}_{\overline{K}}(E). \quad (7.20)$$

This is a 1-cocycle: $\xi(\sigma\tau) = \varphi^{-1} \circ \sigma\tau\varphi = \varphi^{-1} \circ \sigma(\tau\varphi) = (\varphi^{-1} \circ \sigma\varphi) \circ (\sigma\varphi^{-1} \circ \sigma(\tau\varphi)) = \xi(\sigma) \circ \sigma(\xi(\tau))$.

Independence of φ . If $\varphi' : E_{\overline{K}} \xrightarrow{\sim} E'_{\overline{K}}$ is a different isomorphism, then $\varphi' = \varphi \circ \alpha$ for some $\alpha \in \text{Aut}_{\overline{K}}(E)$. The new cocycle is $\xi'(\sigma) = (\varphi \circ \alpha)^{-1} \circ \sigma(\varphi \circ \alpha) = \alpha^{-1} \circ \xi(\sigma) \circ \sigma(\alpha)$, which is cohomologous to ξ . So the class $[\xi] \in H^1(K, \text{Aut}_{\overline{K}}(E))$ depends only on the K -isomorphism class of E' .

Triviality $\iff K$ -isomorphism. $[\xi] = 0$ iff $\xi(\sigma) = \alpha^{-1} \circ \sigma(\alpha)$ for some $\alpha \in \text{Aut}_{\overline{K}}(E)$, iff $\varphi \circ \alpha : E_{\overline{K}} \rightarrow E'_{\overline{K}}$ satisfies $\sigma(\varphi \circ \alpha) = \varphi \circ \alpha$ for all σ (i.e., $\varphi \circ \alpha$ is defined over K), iff $E \cong_K E'$.

Surjectivity. Given a cocycle $\xi : G_K \rightarrow \text{Aut}_{\overline{K}}(E)$, the twist E_ξ is constructed by descent: as a variety over $\overline{K}$ it is $E_{\overline{K}}$, with the twisted Galois action $\sigma *_\xi P = \sigma P \circ \xi(\sigma)$ (composing the original action with the automorphism $\xi(\sigma)$). By Galois descent, this defines a curve E_ξ/K with $E_\xi \times_K \overline{K} \cong E_{\overline{K}}$, and the identity map on $E_{\overline{K}}$ gives an isomorphism $\varphi : E_{\overline{K}} \rightarrow (E_\xi)_{\overline{K}}$ whose cocycle is ξ .

The theorem is a special case of a general principle in Galois descent: for any algebraic object X/K , the set of forms (objects that become isomorphic to X over $\overline{K}$) is classified by $H^1(K, \text{Aut}_{\overline{K}}(X))$.

This applies to curves, group schemes, algebras, and many other structures.

Recovering the Explicit Classifications

We now show that theorem 7.4.1 recovers the classifications from section 1.5 in all three cases.

Proposition 7.4.2: Quadratic Twists via Cohomology ($j \neq 0, 1728$)

Let E/K with $j(E) \neq 0, 1728$ and $\text{char} K \neq 2$. Then $\text{Aut}_{\overline{K}}(E) = \{\pm 1\} \cong \mu_2$, with G_K acting trivially, and

$$\text{Twist}(E/K) \cong H^1(K, \mu_2) \cong K^*/(K^*)^2,$$

where the last isomorphism is the Kummer isomorphism (theorem 7.2.4).

Proof:

By theorem 1.5.2: $\text{Aut}_{\overline{K}}(E) = \{[1], [-1]\} \cong \mathbb{Z}/2\mathbb{Z}$. The Galois group G_K acts trivially (since $[-1]$ is defined over K : it is $(x, y) \mapsto (x, -y)$ on any Weierstrass model). Since $\mu_2 = \{1, -1\}$ with trivial action:

$$H^1(K, \text{Aut}(E)) = H^1(K, \mu_2) \cong K^*/(K^*)^2.$$

The class $[d] \in K^*/(K^*)^2$ corresponds to the quadratic twist $E^{(d)}$ from ???. We verify: the isomorphism $\varphi : E_{\overline{K}} \rightarrow E_{\overline{K}}^{(d)}$ is $(x, y) \mapsto (dx, d^{3/2}y)$ (which requires $\sqrt{d} \in \overline{K}$), and the cocycle is $\xi(\sigma) = \varphi^{-1} \circ \sigma \varphi$. For σ with $\sigma(\sqrt{d}) = -\sqrt{d}$: $\sigma \varphi(x, y) = (dx, -d^{3/2}y)$, so $\xi(\sigma) = \varphi^{-1}(dx, -d^{3/2}y) = (x, -y) = [-1]$. For σ with $\sigma(\sqrt{d}) = \sqrt{d}$: $\xi(\sigma) = [1]$. So ξ is the Kummer cocycle $\sigma \mapsto \sigma(\sqrt{d})/\sqrt{d} \in \mu_2$, corresponding to $[d] \in K^*/(K^*)^2$.

Proposition 7.4.3: Quartic Twists via Cohomology ($j = 1728$)

Let $E : y^2 = x^3 + Ax$ over K with $\text{char} K \neq 2, 3$ and $A \neq 0$, so $j(E) = 1728$. Then $\text{Aut}_{\overline{K}}(E) = \langle \iota \rangle \cong \mu_4$ where $\iota(x, y) = (-x, iy)$ (with $i^2 = -1$, $i \in \overline{K}$), and

$$\text{Twist}(E/K) \cong H^1(K, \mu_4) \cong K^*/(K^*)^4,$$

the last isomorphism holding when $\mu_4 \subseteq K^*$ (i.e., $i \in K$). When $i \notin K$: G_K acts nontrivially on μ_4 , and $H^1(K, \mu_4)$ is computed using the inflation-restriction sequence.

Proof Sketch:

By theorem 1.5.2: $\text{Aut}_{\overline{K}}(E) = \{[1], [-1], \iota, -\iota\} \cong \mathbb{Z}/4\mathbb{Z}$, generated by $\iota : (x, y) \mapsto (-x, iy)$. The Galois action on $\text{Aut}(E)$ factors through $\text{Gal}(K(i)/K)$: if $\sigma(i) = -i$, then $\sigma(\iota) = -\iota$ (conjugation in μ_4). The identification $\text{Aut}(E) \cong \mu_4$ (via $\iota \mapsto i$) is equivariant.

By theorem 7.4.1: $\text{Twist}(E/K) \cong H^1(K, \mu_4)$. The Kummer sequence $1 \rightarrow \mu_4 \rightarrow \overline{K}^* \xrightarrow{(\cdot)^4} \overline{K}^* \rightarrow 1$ and Hilbert 90 give $H^1(K, \mu_4) \cong K^*/(K^*)^4$ (by theorem 7.2.4 with $n = 4$). The class $[d] \in K^*/(K^*)^4$ corresponds to the twist $E^{(d)} : y^2 = x^3 + d^2Ax$ (matching proposition 1.5.8).

Proposition 7.4.4: Sextic Twists via Cohomology ($j = 0$)

Let $E : y^2 = x^3 + B$ over K with $\text{char} K \neq 2, 3$ and $B \neq 0$, so $j(E) = 0$. Then $\text{Aut}_{\overline{K}}(E) \cong \mu_6$, and

$$\text{Twist}(E/K) \cong H^1(K, \mu_6) \cong K^*/(K^*)^6.$$

Proof Sketch:

By theorem 1.5.2: $\text{Aut}_{\overline{K}}(E) = \langle \psi \rangle \cong \mathbb{Z}/6\mathbb{Z}$, where $\psi(x, y) = (\omega x, -y)$ with $\omega = e^{2\pi i/3}$. The identification $\text{Aut}(E) \cong \mu_6$ (via $\psi \mapsto -\omega$) is G_K -equivariant. The Kummer sequence with $n = 6$ and Hilbert 90 give $H^1(K, \mu_6) \cong K^*/(K^*)^6$. The class $[d]$ corresponds to $E^{(d)} : y^2 = x^3 + d^3 B$ (matching proposition 1.5.10).

Twists and Local Data

How does twisting affect the local arithmetic? Over a local field K_v , a twist $E^{(d)}$ may have a different reduction type than E , because the minimal model changes.

Proposition 7.4.5: Twisting and Reduction Type

Let E/K_v be an elliptic curve over a local field.

1. (*Good reduction, quadratic twist.*) If E has good reduction and $d \in R_v^*$ (a unit): $E^{(d)}$ also has good reduction. If $v(d) > 0$: the reduction type of $E^{(d)}$ may change (typically to additive or multiplicative).
2. (*Multiplicative reduction.*) If E has split multiplicative reduction: the quadratic twist by a non-square unit $d \in R_v^*$ changes it to non-split multiplicative (and vice versa). The Kodaira type I_n is preserved, but the component group changes.
3. (*Conductor.*) Twisting can change the conductor exponent at v , but the conductor only changes at primes dividing $2d\Delta$.

Proof Key ideas:

(1) If $d \in R_v^*$: the twist $E^{(d)} : y^2 = x^3 + d^2 Ax + d^3 B$ has $v(\Delta^{(d)}) = v(d^{12}\Delta) = v(\Delta)$ (since $v(d) = 0$), so the reduction type is determined by $v(\Delta)$ and $v(c_4^{(d)}) = v(d^4 c_4) = v(c_4)$: same as E . The split/non-split character may change (since $-c_6^{(d)} = -d^6 c_6$, and $d^6 \in (k_v^*)^2$ iff $d^3 \in k_v^*$... more precisely, the quadratic residue of $-c_6$ may flip).

(2) Split $\leftrightarrow$ non-split: the tangent slopes at the node are conjugated by the quadratic character of d . Twisting by a non-square unit swaps them.

Example 7.4.6: Twisting Running Curves

(a) $E : y^2 = x^3 - x$ over $\mathbb{Q}$. $j = 1728$, so twists are quartic: $\text{Twist}(E/\mathbb{Q}) \cong \mathbb{Q}^*/(\mathbb{Q}^*)^4$. The congruent number curves $E_n : y^2 = x^3 - n^2 x$ correspond to $[n] \in \mathbb{Q}^*/(\mathbb{Q}^*)^4$. Two curves E_m, E_n are

$\mathbb{Q}$ -isomorphic iff $m/n \in (\mathbb{Q}^*)^4$. For instance, $E_1 \cong E_d$ iff $d \in (\mathbb{Q}^*)^4$: the “simplest” representatives are $d = 1, 2, 3, 4, 5, \dots$ with d fourth-power-free.

(b) $E : y^2 = x^3 + 1$ over $\mathbb{Q}$. $j = 0$, so twists are sextic: $\text{Twist}(E/\mathbb{Q}) \cong \mathbb{Q}^*/(\mathbb{Q}^*)^6$. The curve $y^2 = x^3 + 4$ is the twist by d where $d^3 \cdot 1 = 4$, i.e., $d = \sqrt[3]{4}$... in the sextic twist parametrization, $E^{(d)} : y^2 = x^3 + d^3 B = x^3 + d^3$. So $y^2 = x^3 + 4$ corresponds to $d^3 = 4$, i.e., $d = 4^{1/3}$. The class $[d] = [4^{1/3}]$ in $\mathbb{Q}^*/(\mathbb{Q}^*)^6$: since $4 = 2^2$, and $(2^{1/3})^6 = 4$, so $4^{1/3} = 2^{2/3}$ is not rational. The twist by $d = 2$ gives $y^2 = x^3 + 8$, and $[2] \neq [1]$ in $\mathbb{Q}^*/(\mathbb{Q}^*)^6$: the curves $y^2 = x^3 + 1$ and $y^2 = x^3 + 8$ are not $\mathbb{Q}$ -isomorphic.

(c) *Twisting and reduction at $p = 5$.* $E : y^2 = x^3 - x$ over $\mathbb{Q}_5$ has good reduction ($v_5(\Delta) = 0$). The twist $E^{(5)} : y^2 = x^3 - 25x$ has $\Delta^{(5)} = -16 \cdot 4 \cdot 25^3 = -1000000$, $v_5(\Delta^{(5)}) = 6 > 0$. The good reduction has become bad. (The Kodaira type at 5 is I_0^* , from Tate’s algorithm: $A = -25$, $v(A) = 2$, $B = 0$, $v(B) = \infty$.)

Exercise 7.4.1

1. For $E : y^2 = x^3 + x + 1$ over $\mathbb{Q}$ ($j \neq 0, 1728$): write down the cocycle $\xi : G_{\mathbb{Q}} \rightarrow \text{Aut}(E) = \{\pm 1\}$ corresponding to the twist $E^{(d)} : y^2 = x^3 + d^2 x + d^3$ for $d = -3$.
2. Over $\mathbb{Q}_p$ ($p \geq 5$): how many $\mathbb{Q}_p$ -isomorphism classes of elliptic curves have $j = 1728$? (*Hint*: $|\mathbb{Q}_p^*/(\mathbb{Q}_p^*)^4|$. Since $\mathbb{Q}_p^* \cong p^{\mathbb{Z}} \times \mathbb{Z}_p^*$ and $\mathbb{Z}_p^* \cong \mathbb{F}_p^* \times (1 + p\mathbb{Z}_p)$: the 4-th power classes depend on $p \pmod{4}$.)
3. Over $\mathbb{F}_p$ ($p \geq 5$): how many $\mathbb{F}_p$ -isomorphism classes of elliptic curves have $j = 0$? (*Hint*: $|\mathbb{F}_p^*/(\mathbb{F}_p^*)^6| = \gcd(6, p-1)$.)
4. Let $E : y^2 = x^3 - x$ over $\mathbb{Q}_5$ (good reduction). Compute the reduction type of $E^{(d)}$ for $d = 2$ (a non-square unit mod 5) and $d = 5$.
5. For $j = 0$ and $K = \mathbb{Q}$: the automorphism group $\text{Aut}_{\overline{\mathbb{Q}}}(E) \cong \mu_6$ is abelian, so $H^1(\mathbb{Q}, \mu_6)$ is a group. But if we had a curve with $\text{Aut} \cong S_3$ (non-abelian, which doesn’t happen for elliptic curves but does for genus-2 curves), explain why $H^1(K, S_3)$ would be only a pointed set, not a group.
6. Show that the quadratic twist $E^{(d)}$ of E by d satisfies $(E^{(d)})^{(d)} \cong E$ over K (i.e., twisting twice by the same d returns to the original curve). Interpret this in terms of $H^1(K, \mu_2)$: the class $[d]$ has order ≤ 2 .

7.5 Local and Global Duality

The Selmer group $\text{Sel}^{(n)}(E/K)$ is defined as a subgroup of $H^1(K, E[n])$ cut out by local conditions at every place. To compute it, we need two things: an understanding of the local cohomology groups $H^1(K_v, E[n])$, and a description of the local images $\text{im}(\delta_v : E(K_v)/nE(K_v) \rightarrow H^1(K_v, E[n]))$. Both are governed by duality theorems that pair H^1 with itself (via the Weil pairing) and relate local and global cohomology. This section presents the key duality results and their consequences for Selmer group computations.

The Cartier Dual and the Weil Pairing

The n -torsion $E[n]$ is a finite G_K -module, and it carries a natural self-duality via the Weil pairing. Recall from theorem 2.5.5 that $e_n : E[n] \times E[n] \rightarrow \mu_n$ is a perfect, alternating, Galois-equivariant pairing. This means $E[n]$ is isomorphic to its own *Cartier dual*:

Definition 7.5.1: Cartier Dual

For a finite G_K -module M , the *Cartier dual* (or *Tate dual*) is

$$M^* = \text{Hom}(M, \mu_n) = \text{Hom}(M, \overline{K}^*),$$

with G_K -action $(\sigma \cdot f)(m) = \sigma(f(\sigma^{-1}m))$.

Proposition 7.5.2: Self-Duality of $E[n]$

The Weil pairing induces a G_K -equivariant isomorphism

$$E[n] \xrightarrow{\sim} E[n]^* = \text{Hom}(E[n], \mu_n), \quad P \mapsto e_n(P, \cdot). \quad (7.21)$$

This isomorphism is canonical (independent of any choices) and is the algebraic incarnation of the Poincaré duality on the elliptic curve.

Proof:

The Weil pairing e_n is nondegenerate (theorem 2.5.5(3)): for each $P \neq O$ in $E[n]$, there exists $Q \in E[n]$ with $e_n(P, Q) \neq 1$. So the map $P \mapsto e_n(P, \cdot)$ has trivial kernel, hence is an isomorphism (both sides have order n^2). Galois equivariance: $e_n(\sigma P, \sigma Q) = \sigma(e_n(P, Q))$ (theorem 2.5.5(4)) gives $\sigma(e_n(P, \cdot)) = e_n(\sigma P, \sigma(\cdot))$, which is the equivariance of the map $P \mapsto e_n(P, \cdot)$.

Tate Local Duality

Over a local field K_v , the cohomology groups $H^i(K_v, E[n])$ are finite, and they are related by a perfect duality.

Theorem 7.5.3: Tate Local Duality

Let K_v be a non-archimedean local field and M a finite G_{K_v} -module with $|M|$ coprime to $\text{char} K_v$. There is a perfect pairing

$$H^1(K_v, M) \times H^1(K_v, M^*) \longrightarrow H^2(K_v, \mu_n) \cong \mathbb{Z}/n\mathbb{Z}, \quad (7.22)$$

given by the cup product followed by the local invariant map $\text{inv}_v : H^2(K_v, \overline{K}_v^*) \xrightarrow{\sim} \mathbb{Q}/\mathbb{Z}$.

Proof Key ideas:

The cup product $H^1(K_v, M) \times H^1(K_v, M^*) \rightarrow H^2(K_v, M \otimes M^*) \rightarrow H^2(K_v, \mu_n)$ uses the evaluation map $M \otimes M^* \rightarrow \mu_n$ and the local invariant map from class field theory. Non-degeneracy follows from the Tate–Nakayama duality theorem for local fields. See Serre [serreLocalFields1979] or Neukirch–Schmidt–Wingberg [15] for the full proof.

For $M = E[n]$: the self-duality $E[n] \cong E[n]^*$ (proposition 7.5.2) gives a perfect pairing

$$H^1(K_v, E[n]) \times H^1(K_v, E[n]) \longrightarrow \mathbb{Z}/n\mathbb{Z}. \quad (7.23)$$

This pairing is *alternating* (since the Weil pairing is alternating), and its existence means $H^1(K_v, E[n])$ is “self-dual.”

The Local Euler Characteristic

Tate duality has a quantitative consequence: it constrains the sizes of the local cohomology groups.

Proposition 7.5.4: Local Euler Characteristic

Let K_v be a non-archimedean local field with residue field k_v , and let $M = E[n]$ with $\gcd(n, \text{char} k_v) = 1$. The cohomology groups satisfy:

1. $|H^0(K_v, E[n])| = |E(K_v)[n]|$.
2. $|H^2(K_v, E[n])| = |E(K_v)[n]|$ (by Tate duality and the self-duality of $E[n]$: $H^2(K_v, E[n]) \cong H^0(K_v, E[n]^*)^\vee \cong H^0(K_v, E[n])^\vee$).
3. The local Euler characteristic formula:

$$\frac{|H^0(K_v, E[n])| \cdot |H^2(K_v, E[n])|}{|H^1(K_v, E[n])|} = |n|_v^2, \quad (7.24)$$

where $|n|_v$ is the v -adic absolute value.

For $v \nmid n$: $|n|_v = 1$, so $|H^1(K_v, E[n])| = |E(K_v)[n]|^2$.

This is useful for bounding the size of $H^1(K_v, E[n])$: at primes of good reduction with $v \nmid n$, the local cohomology has order $|E(K_v)[n]|^2$, which is at most n^4 (since $|E[n]| = n^2$).

Orthogonality of the Local Images

The most important consequence of Tate local duality for Selmer group computations is:

Theorem 7.5.5: Local Image is Maximal Isotropic

Under the Tate local duality pairing (7.23), the local image $\text{im}(\delta_v : E(K_v)/nE(K_v) \rightarrow H^1(K_v, E[n]))$ is its own orthogonal complement:

$$\text{im}(\delta_v)^\perp = \text{im}(\delta_v). \quad (7.25)$$

In particular, $|\text{im}(\delta_v)|^2 = |H^1(K_v, E[n])|$: the local image has “half the size” of the ambient cohomology group (in the sense that its square equals the total).

Proof Key ideas:

The pairing of two local images $\langle \delta_v(P), \delta_v(Q) \rangle$ is computed via the cup product, which factors through the Néron–Tate height pairing modulo n . The self-orthogonality $\text{im}(\delta_v) \subseteq \text{im}(\delta_v)^\perp$ follows from the symmetry of the Néron–Tate pairing (which gives $\langle \delta_v(P), \delta_v(Q) \rangle = 0$ for “geometric” reasons — the local Tate pairing vanishes on the diagonal). The equality $\text{im}(\delta_v) = \text{im}(\delta_v)^\perp$ then follows from the size count: $|\text{im}(\delta_v)| = |E(K_v)/nE(K_v)|$ has square $|H^1(K_v, E[n])|$ (by the local Euler characteristic). See Milne [13] for the full proof.

The maximal isotropy of the local images is the structural reason why the Selmer group is finite: it is cut out by “half-dimensional” conditions at each place, and these conditions are compatible globally.

The Local Images Explicitly

For Selmer group computations, we need to know the local images $\text{im}(\delta_v)$ concretely. The answer depends on the reduction type and on the relationship between n and the residue characteristic.

Proposition 7.5.6: Local Images by Reduction Type

Let E/K_v with residue characteristic p and $\gcd(n, p) = 1$. The local image $\text{im}(\delta_v) \subseteq H^1(K_v, E[n])$ is:

1. (*Good reduction.*) $\text{im}(\delta_v) = H_{\text{nr}}^1(K_v, E[n])$, the *unramified* subgroup (classes trivial on the inertia group I_v). Its order is $|E(K_v)[n]| = |\tilde{E}(k_v)[n]|$.
2. (*Multiplicative reduction.*) The local image is larger than the unramified subgroup: it includes classes that are “tamely ramified” (corresponding to the component group Φ_v).
3. (*Additive reduction.*) The local image can be computed from the filtration $E_0(K_v) \supset E_1(K_v)$ and the component group Φ_v .

Proof of (1):

For good reduction with $v \nmid n$: the exact sequence $0 \rightarrow E_1(K_v) \rightarrow E(K_v) \rightarrow \tilde{E}(k_v) \rightarrow 0$ (from theorem 6.2.4) gives, on n -torsion quotients: $E(K_v)/nE(K_v) \cong \tilde{E}(k_v)/n\tilde{E}(k_v)$ (since $E_1(K_v) = \tilde{E}(\mathfrak{m}_v)$ is n -divisible: $[n]$ is surjective on the formal group when $\gcd(n, p) = 1$, by theorem 4.5.5(1)).

The connecting map δ_v sends $E(K_v)/nE(K_v) \cong \tilde{E}(k_v)/n\tilde{E}(k_v)$ into $H^1(K_v, E[n])$. The image consists of classes unramified at v : the cocycle $\sigma \mapsto \sigma(Q) - Q$ (where $[n]Q = P$) factors through $G_{k_v} = G_{K_v}/I_v$ because Q is defined over K_v^{nr} (by the NOS criterion, theorem 6.4.1: good reduction implies the n -torsion is unramified). So $\text{im}(\delta_v) \subseteq H_{\text{nr}}^1(K_v, E[n])$.

For the reverse inclusion: $|H_{\text{nr}}^1(K_v, E[n])| = |H^1(G_{k_v}, E[n]^{I_v})| = |H^1(G_{k_v}, \tilde{E}[n])| = |\tilde{E}(k_v)/n\tilde{E}(k_v)| = |E(K_v)/nE(K_v)| = |\text{im}(\delta_v)|$. (The second equality uses $E[n]^{I_v} = \tilde{E}[n]$ for good reduction with $v \nmid n$. The fourth uses the isomorphism above.) Since $\text{im}(\delta_v) \subseteq H_{\text{nr}}^1$ and they have the same order: equality.

The upshot for computations: at primes of good reduction with $v \nmid n$, the Selmer condition is simply “unramified.” At finitely many “bad” primes (those dividing $n\Delta$), the local condition must be computed individually (using the exact sequence and the component group).

The Poitou–Tate Exact Sequence

The global counterpart of Tate local duality is the Poitou–Tate exact sequence, which relates the Selmer group to its “dual Selmer group.”

Theorem 7.5.7: Poitou–Tate Exact Sequence

Let E/K be an elliptic curve over a number field, and let $n \geq 1$ with $\text{char} K \nmid n$. There is a nine-term exact sequence:

$$\begin{aligned} 0 \rightarrow \text{Sel}^{(n)}(E/K) \rightarrow H^1(K, E[n]) \rightarrow \bigoplus_v \frac{H^1(K_v, E[n])}{\text{im}(\delta_v)} \\ \rightarrow \text{Sel}^{(n)}(E/K)^\vee \rightarrow H^2(K, E[n]) \rightarrow \bigoplus_v H^2(K_v, E[n]) \rightarrow \cdots \end{aligned} \quad (7.26)$$

where $\text{Sel}^{(n)}(E/K)^\vee = \text{Hom}(\text{Sel}^{(n)}(E/K), \mathbb{Q}/\mathbb{Z})$ is the Pontryagin dual.

We state this without proof (the full proof requires the Artin–Verdier duality theorem and the global class field theory exact sequence; see Milne [13] or Neukirch–Schmidt–Wingberg [15]). The key consequence is:

Corollary 7.5.8: The Selmer Group is Computable

The n -Selmer group $\text{Sel}^{(n)}(E/K)$ is a finite subgroup of $H^1(K, E[n])$, determined by finitely many local conditions:

$$\text{Sel}^{(n)}(E/K) = \{ \xi \in H^1(K, E[n]) \mid \text{res}_v(\xi) \in \text{im}(\delta_v) \text{ for all } v \}. \quad (7.27)$$

Since $\text{im}(\delta_v) = H_{\text{nr}}^1(K_v, E[n])$ for all but finitely many v (proposition 7.5.6(1)), the Selmer group is the intersection of the “unramified everywhere outside S ” subgroup with finitely many explicit local conditions at the primes in $S = \{v : v \mid n\Delta_{\min}\} \cup \{v \mid \infty\}$.

This is the computational foundation for Chapter 9: to compute $\text{Sel}^{(n)}(E/K)$, one enumerates the

classes in $H^1(K, E[n])$ that are unramified outside S (a finite group, by class field theory) and checks the finitely many local conditions at primes in S .

Exercise 7.5.1

1. Verify that the Weil pairing e_n induces an isomorphism $E[n] \cong E[n]^*$ for $E : y^2 = x^3 - x$ and $n = 2$. (*Hint:* $E[2] = \{O, (0, 0), (1, 0), (-1, 0)\} \cong (\mathbb{Z}/2\mathbb{Z})^2$, and the Weil pairing e_2 is a perfect alternating form on $(\mathbb{Z}/2\mathbb{Z})^2$. Write down the pairing matrix.)
2. For $E : y^2 = x^3 - x$ over $\mathbb{Q}_5$ (good reduction, $|\tilde{E}(\mathbb{F}_5)| = 8$) and $n = 2$: compute $|E(\mathbb{Q}_5)[2]|$, $|H^1(\mathbb{Q}_5, E[2])|$, and $|\text{im}(\delta_5)|$, using the local Euler characteristic.
3. For good reduction at v with $v \nmid n$: show that $|H_{\text{nr}}^1(K_v, E[n])| = |E(K_v)/nE(K_v)|$ by using the surjectivity of $[n]$ on $\hat{E}(\mathfrak{m}_v)$ and Lang's theorem ($H^1(G_{k_v}, \tilde{E}) = 0$ for connected algebraic groups over finite fields).
4. Explain why the Selmer group is finite, using: (1) $H^1(K, E[n])$ is infinite, but (2) the “unramified outside S ” subgroup is finite (by the finiteness of the class group and the Dirichlet unit theorem), and (3) the local conditions at primes in S are checked against finite groups.
5. The Poitou–Tate exact sequence relates $\text{Sel}^{(n)}$ to its dual. Using the self-duality $E[n] \cong E[n]^*$, show that $\text{Sel}^{(n)}(E/K)^\vee$ is related to a Selmer group for the same curve (not a dual curve). In particular, $|\text{Sel}^{(n)}|$ is a perfect square... is this true? (*Hint:* not in general; the Cassels–Tate pairing on III is alternating, which forces $|\text{III}[n]|$ to be a perfect square, but $|\text{Sel}^{(n)}|$ need not be.)
6. For $E : y^2 = x^3 - x$ over $\mathbb{Q}$ with $n = 2$: the set $S = \{2, \infty\}$ (bad primes and archimedean places). Describe the “unramified outside S ” subgroup of $H^1(\mathbb{Q}, E[2])$ and the local condition at $v = 2$. (*This is the starting point for a full 2-Selmer computation, developed in chapter 9.*)

Elliptic Curves over Number Fields: The Mordell–Weil Theorem

The deepest result in the arithmetic of elliptic curves is the theorem of Mordell (for $\mathbb{Q}$, 1922) and Weil (for number fields, 1928): the group of rational points $E(K)$ is finitely generated. This chapter proves the theorem and develops the theory of heights that makes the proof work.

The argument has two ingredients. The first is *algebraic*: the weak Mordell–Weil theorem, which says $E(K)/nE(K)$ is finite. This was established in section 7.3 via the descent exact sequence and the Selmer group machinery. The second ingredient is *analytic*: the theory of *heights*, which assigns to each rational point $P \in E(K)$ a real number $\hat{h}(P) \geq 0$ measuring its arithmetic complexity. The height grows quadratically under multiplication ($\hat{h}([m]P) = m^2\hat{h}(P)$) and satisfies a finiteness property (only finitely many points of bounded height). These two ingredients combine through a descent argument: starting from any $P \in E(K)$, repeated division by n (modulo the finite group $E(K)/nE(K)$) produces a sequence of points whose heights shrink geometrically, eventually landing in a finite set. The result: $E(K)$ is generated by finitely many points.

8.1 The Weak Mordell–Weil Theorem

Theorem 8.1.1: The Weak Mordell–Weil Theorem

Let E be an elliptic curve over a number field K , and let $n \geq 2$ be an integer. Then $E(K)/nE(K)$ is a finite group.

The proof assembles the cohomological machinery of chapter 7. We give it in two stages: first assuming the n -torsion is rational ($E[n] \subseteq E(K)$), where the argument is most transparent, and then

in full generality.

Proof When $E[n] \subseteq E(K)$

Assume all n -torsion points are K -rational. By the descent exact sequence (theorem 7.3.1):

$$0 \longrightarrow E(K)/nE(K) \xrightarrow{\delta} H^1(K, E[n]) \longrightarrow H^1(K, E)[n] \longrightarrow 0. \quad (8.1)$$

The group $E(K)/nE(K)$ injects into $H^1(K, E[n])$. Since G_K acts trivially on $E[n] \cong (\mathbb{Z}/n\mathbb{Z})^2$:

$$H^1(K, E[n]) = \text{Hom}(G_K, E[n]) \cong \text{Hom}(G_K, (\mathbb{Z}/n\mathbb{Z})^2).$$

This is an infinite group (there are infinitely many homomorphisms $G_K \rightarrow (\mathbb{Z}/n\mathbb{Z})^2$, corresponding to abelian extensions of K). So the injection alone does not give finiteness — we need to show the image of δ is contained in a *finite* subgroup.

The image of δ lies in the Selmer group (??):

$$E(K)/nE(K) \hookrightarrow \text{Sel}^{(n)}(E/K) \subseteq H^1(K, E[n]).$$

And the Selmer group is finite (theorem 7.3.9). Therefore $E(K)/nE(K)$ is finite.

To make this concrete when $E[n] \subseteq E(K)$: by Kummer theory (theorem 7.2.4), $H^1(K, E[n]) \cong (K^*/(K^*)^n)^2$. The Selmer condition restricts to pairs $(d_1, d_2) \in (K^*/(K^*)^n)^2$ such that:

1. At each archimedean place v : the pair (d_1, d_2) satisfies an explicit sign condition.
2. At each non-archimedean place v with $v \nmid n\Delta$: $d_1, d_2 \in (\mathcal{O}_{K_v}^*)^n$ (the “unramified” condition — both d_i are n -th powers in the v -adic completion).
3. At each place $v \mid n\Delta$: an explicit local condition is satisfied.

Condition (2) forces d_1, d_2 to be S -units modulo n -th powers, where S is the finite set of places dividing $n\Delta\infty$. The group of S -units modulo n -th powers is finite (by the Dirichlet S -unit theorem: $\mathcal{O}_{K,S}^*$ is finitely generated, so $\mathcal{O}_{K,S}^*/(\mathcal{O}_{K,S}^*)^n$ is finite). Condition (3) imposes finitely many additional finite conditions. The result: $\text{Sel}^{(n)}(E/K)$ is finite.

Proof in General

When $E[n] \not\subseteq E(K)$: let $L = K(E[n])$, the field obtained by adjoining all n -torsion points. This is a finite Galois extension of K (with $\text{Gal}(L/K) \hookrightarrow \text{GL}_2(\mathbb{Z}/n\mathbb{Z})$, via the Galois representation on $E[n]$, section 2.5). By the case above, $E(L)/nE(L)$ is finite.

Proof of theorem 8.1.1:

We reduce to the case $E[n] \subseteq E(K)$ via a norm argument.

Let $L = K(E[n])$ and $d = [L : K]$. For any $P \in E(K)$: $P \in E(L)$, so $P = nQ + R$ for some $Q \in E(L)$ and R in a fixed finite set of coset representatives of $E(L)/nE(L)$. Apply the norm

map $\text{Nm}_{L/K} : E(L) \rightarrow E(K)$, defined by $\text{Nm}_{L/K}(Q) = \sum_{\sigma \in \text{Gal}(L/K)} \sigma(Q)$:

$$\text{Nm}_{L/K}(P) = n \cdot \text{Nm}_{L/K}(Q) + \text{Nm}_{L/K}(R).$$

Since $P \in E(K)$: $\text{Nm}_{L/K}(P) = dP$ (each σ fixes P , and there are d automorphisms). So:

$$dP = n \cdot \text{Nm}_{L/K}(Q) + \text{Nm}_{L/K}(R). \quad (8.2)$$

The element $\text{Nm}_{L/K}(Q) \in E(K)$ (the norm lands in the fixed field). The element $\text{Nm}_{L/K}(R) \in E(K)$ ranges over a finite set (since R ranges over the finite set of coset representatives, and the norm is a fixed map). Rearranging: $dP - \text{Nm}_{L/K}(R) \in nE(K)$, so the class $[dP] \in E(K)/nE(K)$ lies in a finite set.

Now consider the map $[d] : E(K)/nE(K) \rightarrow E(K)/nE(K)$ (multiplication by d). Its image is finite (by the argument above: the class of dP ranges over a finite set). Its kernel is $\ker([d]) = \{P \bmod nE(K) : dP \in nE(K)\}$. We claim the kernel is also finite: if $dP \in nE(K)$, then in particular $dP = nQ$ for some $Q \in E(K)$, so P is determined modulo $E(K)[d]$ and $nE(K)$. More precisely, the map $E(K)[d] \rightarrow \ker([d])$ given by $T \mapsto T \bmod nE(K)$ is surjective (if $dP \in nE(K)$: write $dP = nQ$, and for any other P' with $dP' \in nE(K)$: $d(P' - P) \in nE(K)$, so the coset of P' modulo $nE(K)$ differs from that of P by an element of $E(K)[d] \bmod nE(K)$). Since $E(K)[d]$ is finite (it injects into $\tilde{E}(\mathbb{F}_p)$ for any prime p of good reduction with $p \nmid d$, by theorem 6.3.2): $\ker([d])$ is finite.

A map of sets with finite kernel and finite image has finite domain: $|E(K)/nE(K)| \leq |\ker([d])| \cdot |\text{im}([d])| < \infty$.

What the Weak Theorem Does and Does Not Prove

The weak Mordell–Weil theorem is a finiteness result for a *quotient*, not for the group itself. To appreciate the gap, consider the analogy with $\mathbb{Z}$: the quotient $\mathbb{Z}/n\mathbb{Z}$ is finite for every n , yet $\mathbb{Z}$ is infinite. Similarly, $E(K)/nE(K)$ being finite tells us that $E(K)$ has “finite rank modulo n ” but does not by itself show $E(K)$ is finitely generated.

What the weak theorem *does* give us is the following:

1. *A bound on the number of generators modulo torsion.* Since $E(K)/nE(K)$ is finite, we can write $|E(K)/nE(K)| = n^r \cdot |E(K)[n]|$ where $r = \text{rank} E(K)$ (assuming n is prime; this follows from the structure theorem for finitely generated abelian groups, once we know $E(K)$ is finitely generated — but that is what we are trying to prove). So the weak theorem gives an *a priori* bound on the rank.
2. *A finite set of coset representatives.* There exist $Q_1, \dots, Q_m \in E(K)$ (with $m = |E(K)/nE(K)|$) such that every $P \in E(K)$ can be written $P = Q_i + nP_1$ for some i and $P_1 \in E(K)$.

The missing ingredient is a way to conclude that the iteration $P \rightarrow P_1 \rightarrow P_2 \rightarrow \dots$ (where $P_k = Q_{i_k} + nP_{k+1}$) terminates. For this, we need a notion of “size” for points that *decreases* under the passage $P \rightarrow P_1$. This is the role of the height function, developed in the next section.

Exercise 8.1.1

1. For $E : y^2 = x^3 - x$ over $\mathbb{Q}$ with $n = 2$: verify that $E(\mathbb{Q})/2E(\mathbb{Q}) \cong (\mathbb{Z}/2\mathbb{Z})^2$ by (a) noting $E(\mathbb{Q}) = E[2] \cong (\mathbb{Z}/2\mathbb{Z})^2$ (rank 0), and (b) computing $2E(\mathbb{Q}) = \{O\}$.
2. For $E : y^2 + y = x^3 - x$ over $\mathbb{Q}$ (37a1) with $n = 2$: the Selmer group has $|\text{Sel}^{(2)}| = 2$. Since $|E(\mathbb{Q})[2]| = 1$ (no nontrivial 2-torsion): $|E(\mathbb{Q})/2E(\mathbb{Q})| \leq 2$. Conclude that $\text{rank} E(\mathbb{Q}) \leq 1$. (*Remark*: the point $(0, 0)$ has infinite order, so $\text{rank} = 1$.)
3. Let $E/\mathbb{Q}$ and $L = \mathbb{Q}(E[n])$. Show that the norm map $\text{Nm}_{L/\mathbb{Q}} : E(L) \rightarrow E(\mathbb{Q})$ is a group homomorphism. (*Hint*: $\text{Nm}_{L/\mathbb{Q}}(P) = \sum_{\sigma} \sigma(P)$; use the fact that the group law commutes with Galois action.)
4. In the proof, the finiteness of $\ker([d] : E(K)/nE(K) \rightarrow E(K)/nE(K))$ uses the finiteness of $E(K)[d]$. Show that the bound $|\ker([d])| \leq |E(K)[d]|$ can be improved: $|\ker([d])| \leq |E(K)[d]/(E(K)[d] \cap nE(K))|$. (*Hint*: the map $E(K)[d] \rightarrow \ker([d])$ given by $T \mapsto T \bmod nE(K)$ is surjective, with kernel $E(K)[d] \cap nE(K)$.)
5. The weak Mordell–Weil theorem is effective: given E/K , one can compute $|E(K)/nE(K)|$ in principle. Explain why: the Selmer group is computable (??), and $E(K)/nE(K) \hookrightarrow \text{Sel}^{(n)}$. The difficulty is that equality need not hold: $\text{III}[n]$ is the obstruction. What additional information would make the computation of $\text{rank} E(K)$ effective?

8.2 Heights

The height of a rational number a/b (in lowest terms) is $\max(|a|, |b|)$: a measure of its arithmetic complexity. This section develops the theory of heights systematically, from projective space to elliptic curves, culminating in the two estimates that make the Mordell–Weil descent work: Northcott’s theorem (finitely many points of bounded height) and the quadratic growth law $h([m]P) = m^2 h(P) + O(1)$.

The key intermediate result is the *height machine*: the height of $\varphi(P)$ is $\deg(\varphi) \cdot h(P) + O(1)$ for any morphism $\varphi : \mathbb{P}^1 \rightarrow \mathbb{P}^1$ of degree d . We prove this in full via the product formula.

The Weil Height on $\mathbb{P}^n$

Definition 8.2.1: The Absolute Logarithmic Height

For a point $P = [x_0 : \cdots : x_n] \in \mathbb{P}^n(\mathbb{Q})$ with $x_i \in \mathbb{Z}$ and $\gcd(x_0, \dots, x_n) = 1$:

$$H(P) = \max(|x_0|, \dots, |x_n|), \quad h(P) = \log H(P). \quad (8.3)$$

$H(P)$ is the *multiplicative height* and $h(P) \geq 0$ is the *logarithmic height*.

For a number field K , the height is defined using the product formula. Let M_K denote the set of places of K (archimedean and non-archimedean), and for each $v \in M_K$ let $\|\cdot\|_v$ be the absolute value on K normalized so that the product formula $\prod_{v \in M_K} \|x\|_v^{n_v} = 1$ holds for all $x \in K^*$, where $n_v = [K_v : \mathbb{Q}_v]$.

Definition 8.2.2: The Height over a Number Field

For $P = [x_0 : \cdots : x_n] \in \mathbb{P}^n(K)$:

$$h(P) = \frac{1}{[K : \mathbb{Q}]} \sum_{v \in M_K} n_v \log \max(\|x_0\|_v, \dots, \|x_n\|_v). \quad (8.4)$$

This is well-defined (independent of the choice of homogeneous coordinates, by the product formula) and independent of the field K containing the coordinates of P (the local degrees adjust correctly under field extension). So h defines a function $h : \mathbb{P}^n(\mathbb{Q}) \rightarrow \mathbb{R}_{\geq 0}$.

Proposition 8.2.3: Well-Definedness of the Height

The height $h(P)$ is:

1. Independent of the choice of homogeneous coordinates for P .
2. Independent of the number field K used to compute it (i.e., if $P \in \mathbb{P}^n(K) \subseteq \mathbb{P}^n(L)$ for an extension L/K , computing $h(P)$ using M_L gives the same result as using M_K).
3. Non-negative: $h(P) \geq 0$ for all P .

Proof:

1. Replacing $(x_0, \dots, x_n)$ by $(\lambda x_0, \dots, \lambda x_n)$ for $\lambda \in K^*$ adds $\frac{1}{[K:\mathbb{Q}]} \sum_v n_v \log \|\lambda\|_v = 0$ to $h(P)$, by the product formula.
2. Under the extension L/K : each place v of K splits into places $w|v$ of L with $\sum_{w|v} n_w = [L : K] \cdot n_v$. So $\frac{1}{[L:\mathbb{Q}]} \sum_{w \in M_L} n_w \log \max_i \|x_i\|_w = \frac{1}{[L:\mathbb{Q}]} \sum_{v \in M_K} \sum_{w|v} n_w \log \max_i \|x_i\|_w = \frac{1}{[K:\mathbb{Q}]} \sum_{v \in M_K} n_v \log \max_i \|x_i\|_v$ (since $\|x_i\|_w = \|x_i\|_v$ for $x_i \in K$ and $w|v$, and $\sum_{w|v} n_w = [L : K] \cdot n_v$).
3. At each place v : $\max_i \|x_i\|_v \geq \|x_j\|_v > 0$ for at least one j (since not all x_i are zero). So $\log \max_i \|x_i\|_v \geq 0$ at archimedean places (after normalizing so that some $x_j = 1$). The non-negativity of $h(P)$ follows from the normalization: for $P \in \mathbb{P}^n(\mathbb{Q})$, $h(P) = \log \max |x_i| \geq 0$; the general case follows from (2).

For $P = [x_0 : x_1] \in \mathbb{P}^1(\mathbb{Q})$ with $x_0, x_1 \in \mathbb{Z}$ coprime: $h(P) = \log \max(|x_0|, |x_1|)$. For a rational number $\alpha = a/b \in \mathbb{Q}$ (in lowest terms): $h(\alpha) := h([a : b]) = \log \max(|a|, |b|)$.

Example 8.2.4: Heights of Small Rationals

$h(0) = h([0 : 1]) = 0$. $h(1) = h([1 : 1]) = 0$. $h(-1) = h([-1 : 1]) = 0$. $h(2) = \log 2 \approx 0.693$. $h(1/3) = h([1 : 3]) = \log 3 \approx 1.099$. $h(17/5) = h([17 : 5]) = \log 17 \approx 2.833$.

The height of $\alpha \in \mathbb{Q}$ grows with its arithmetic complexity: numbers with large numerators or denominators have large height. Integers n have $h(n) = \log |n|$ (for $n \neq 0$), and rationals a/b with

$|a|, |b|$ both large can still have small height if they nearly cancel (but this cannot happen in lowest terms).

Northcott's Theorem

The fundamental finiteness property of the height:

Theorem 8.2.5: Northcott's Theorem

For any $B > 0$ and any integer $d \geq 1$:

$$\{P \in \mathbb{P}^n(\overline{\mathbb{Q}}) \mid [\mathbb{Q}(P) : \mathbb{Q}] \leq d, h(P) \leq B\} \quad (8.5)$$

is a finite set (where $\mathbb{Q}(P)$ is the field generated by the ratios x_i/x_j).

Proof:

It suffices to treat $\mathbb{P}^1$ (a bound on $h(P)$ for $P \in \mathbb{P}^n$ bounds each coordinate ratio in $\mathbb{P}^1$). A point $[\alpha : 1] \in \mathbb{P}^1$ with $[\mathbb{Q}(\alpha) : \mathbb{Q}] = d$ and $h(\alpha) \leq B$ has minimal polynomial $f(T) = a_d T^d + \cdots + a_0 \in \mathbb{Z}[T]$ with $a_d > 0$ and $\gcd(a_0, \dots, a_d) = 1$. The roots of f are the conjugates $\alpha_1 = \alpha, \alpha_2, \dots, \alpha_d$, and by the relation between the height and the Mahler measure:

$$|a_i| \leq \binom{d}{i} \cdot \prod_{j=1}^d \max(1, |\alpha_j|) \leq \binom{d}{i} \cdot H(\alpha)^d \leq \binom{d}{i} \cdot e^{dB}.$$

(The first inequality is the standard coefficient bound; the second uses $\max(1, |\alpha_j|) \leq H(\alpha)$ for each conjugate, which follows from the product formula and the definition of the height.) So the coefficients of f are bounded: $|a_i| \leq \binom{d}{i} e^{dB}$. There are finitely many integer polynomials with bounded coefficients, and each has finitely many roots. The set (8.5) is finite.

Over a fixed number field, the consequence is cleaner:

Corollary 8.2.6: Northcott over K

For any number field K and any $B > 0$: $\{P \in \mathbb{P}^n(K) \mid h(P) \leq B\}$ is a finite set.

Heights of Rational Functions: The Height Machine

The following theorem is the workhorse of height theory. It says that a rational function of degree d multiplies the height by d , up to a bounded error.

Theorem 8.2.7: Height of a Rational Function

Let $\varphi : \mathbb{P}^1 \rightarrow \mathbb{P}^1$ be a morphism of degree d defined over K , i.e., $\varphi([x : y]) = [F(x, y) : G(x, y)]$ where $F, G \in K[x, y]$ are homogeneous of degree d with no common factor. Then there exist constants $c_1, c_2 > 0$ (depending on φ) such that

$$d \cdot h(\alpha) - c_1 \leq h(\varphi(\alpha)) \leq d \cdot h(\alpha) + c_2 \quad (8.6)$$

for all $\alpha \in \mathbb{P}^1(\overline{K})$. In short: $h(\varphi(\alpha)) = d \cdot h(\alpha) + O(1)$.

Proof:

We work over a number field L containing K and the coordinates of α . Write $\alpha = [a : b]$ with $a, b \in L$ (not both zero).

Upper bound. For each place v of L :

$$\max(\|F(a, b)\|_v, \|G(a, b)\|_v) \leq C_v \cdot \max(\|a\|_v, \|b\|_v)^d,$$

where C_v depends on the coefficients of F and G . At archimedean places: $C_v \leq (d+1) \cdot \max_i \|c_i\|_v$ where the c_i are the coefficients of F and G (by the triangle inequality: a homogeneous polynomial of degree d in two variables has at most $d+1$ terms). At non-archimedean places: $C_v = \max_i \|c_i\|_v$ (the ultrametric inequality replaces the triangle inequality). Taking logarithms, summing over v with weights $n_v/[L : \mathbb{Q}]$, and using the product formula (which absorbs the C_v into a constant):

$$h(\varphi(\alpha)) \leq d \cdot h(\alpha) + \frac{1}{[K : \mathbb{Q}]} \sum_v n_v \log C_v =: d \cdot h(\alpha) + c_2.$$

Lower bound. This is the deeper direction. Since F and G have no common factor over $\overline{K}$, the resultant $\text{Res}(F, G) \in K^*$ is nonzero. By the theory of resultants: for any $(a, b) \in L^2 \setminus \{(0, 0)\}$,

$$\|\text{Res}(F, G)\|_v \leq C'_v \cdot \max(\|F(a, b)\|_v, \|G(a, b)\|_v) \cdot \max(\|a\|_v, \|b\|_v)^{d-1} \quad (8.7)$$

for suitable C'_v . Summing logs over all v from (8.7) gives only $h(\varphi(\alpha)) \geq -(d-1)h(\alpha) - c'$, which is too weak (the coefficient of $h(\alpha)$ is negative rather than $+d$). To obtain the sharp lower bound, we use the *Bezout cofactors* of the resultant: there exist homogeneous polynomials $A, B, C, D \in K[x, y]$ of degree $d-1$ such that $AF + BG = \text{Res}(F, G) \cdot x^{2d-1}$ and $CF + DG = \text{Res}(F, G) \cdot y^{2d-1}$. From the first identity:

$$\|\text{Res}\|_v \cdot \|a\|_v^{2d-1} \leq C''_v \cdot \max(\|F(a, b)\|_v, \|G(a, b)\|_v) \cdot \max(\|a\|_v, \|b\|_v)^{d-1},$$

and similarly with b from the second. Taking the max over a and b :

$$\max(\|a\|_v, \|b\|_v)^d \leq \frac{C''_v}{\|\text{Res}\|_v} \cdot \max(\|F(a, b)\|_v, \|G(a, b)\|_v).$$

Taking logs and summing: $d \cdot h(\alpha) \leq h(\varphi(\alpha)) + c_1$, i.e., $h(\varphi(\alpha)) \geq d \cdot h(\alpha) - c_1$. This completes the proof.

The proof is somewhat technical, but the idea is clean: the upper bound uses the triangle inequality (a polynomial of degree d can grow at most like h^d), and the lower bound uses the resultant (the nonvanishing of $\text{Res}(F, G)$ prevents F and G from simultaneously being too small).

Example 8.2.8: Height of a Squaring Map

$\varphi : \mathbb{P}^1 \rightarrow \mathbb{P}^1$, $[x : y] \mapsto [x^2 : y^2]$ (degree 2). For $\alpha = [a : b] \in \mathbb{P}^1(\mathbb{Q})$: $\varphi(\alpha) = [a^2 : b^2]$. Then $h(\varphi(\alpha)) = \log \max(a^2, b^2) = 2 \log \max(|a|, |b|) = 2h(\alpha)$. In this case, $h(\varphi(\alpha)) = 2h(\alpha)$ exactly (no error term), because the map is a monomial. For non-monomial maps, the error term is typically nonzero.

The Naive Height on an Elliptic Curve

Definition 8.2.9: Naive Height on E

Let E/K be an elliptic curve given by a Weierstrass equation. For $P = (x_0, y_0) \in E(K) \setminus \{O\}$: define the *naive height* as

$$h(P) = h_x(P) := h(x(P)), \quad (8.8)$$

viewing $x(P) \in K$ as a point of $\mathbb{P}^1(K)$ via $x \mapsto [x : 1]$. Set $h(O) = 0$.

The naive height depends on the choice of Weierstrass equation (a different equation gives a different x -coordinate, hence a different height). This dependence is bounded: if x' is the x -coordinate for a different Weierstrass model related by an admissible change (u, r, s, t) , then $x' = u^{-2}(x - r)$, so $h(x'(P)) = h(x(P)) + O(1)$ (the error depends on u and r but not on P). The Néron–Tate height of section 8.3 will remove this ambiguity entirely.

By Northcott (??): for any $B > 0$, the set $\{P \in E(K) : h(P) \leq B\}$ is finite (since $h(P) \leq B$ bounds $h(x(P)) \leq B$, and $y(P)$ is determined by $x(P)$ up to sign via the Weierstrass equation).

Heights and the Group Law

We now derive the key estimates relating the height to the group operation on E , using theorem 8.2.7 and the explicit formulas from section 1.3.

Proposition 8.2.10: Height Estimates for the Group Law

Let E/K be an elliptic curve in short Weierstrass form $y^2 = x^3 + Ax + B$ (with $\text{char} K \neq 2, 3$). There exist constants $c_1, c_2 > 0$ (depending on E/K) such that for all $P, Q \in E(K)$:

1. (*Quasi-triangle inequality.*) $h(P + Q) \leq 2h(P) + 2h(Q) + c_1$.
2. (*Quasi-symmetry.*) $h(-P) = h(P)$.
3. (*Quadratic growth under duplication.*) $h([2]P) = 4h(P) + O(c_2)$.

Proof:

1. For $P = (x_1, y_1)$, $Q = (x_2, y_2)$ with $P \neq \pm Q$: the addition formula (eq. (1.19)) gives

$$x(P + Q) = \left(\frac{y_2 - y_1}{x_2 - x_1} \right)^2 - x_1 - x_2.$$

Clearing denominators: $x(P + Q) \cdot (x_2 - x_1)^2 = (y_2 - y_1)^2 - (x_1 + x_2)(x_2 - x_1)^2$. The right side involves x_1, x_2, y_1, y_2 with $y_i^2 = x_i^3 + Ax_i + B$. Writing $x(P + Q) = F(x_1, x_2, y_1, y_2)/G(x_1, x_2, y_1, y_2)$ (a rational function): F and G have total degree ≤ 3 in the x_i and ≤ 1 in the y_i . Using $h(y_i) \leq \frac{3}{2}h(x_i) + O(1)$ (from $y_i^2 = x_i^3 + Ax_i + B$ and the triangle inequality for heights): $h(x(P + Q)) \leq 2h(x_1) + 2h(x_2) + c_1$ after collecting terms.

The case $P = Q$ (doubling) and the case $P = -Q$ ($P + Q = O$, $h(O) = 0 \leq 2h(P)$) are handled separately. The result: $h(P + Q) \leq 2h(P) + 2h(Q) + c_1$ for all P, Q .

2. $x(-P) = x(P)$ for any Weierstrass model (negation preserves the x -coordinate), so $h(-P) = h(x(-P)) = h(x(P)) = h(P)$.
3. The duplication formula for short Weierstrass $y^2 = x^3 + Ax + B$ (eq. (1.20)) gives:

$$x([2]P) = \frac{x^4 - 2Ax^2 - 8Bx + A^2}{4x^3 + 4Ax + 4B} =: \frac{\varphi_2(x)}{\psi_2(x)}, \quad (8.9)$$

where $\varphi_2 \in K[x]$ has degree 4 and $\psi_2 = 4(x^3 + Ax + B)$ has degree 3. As a morphism $\mathbb{P}^1 \rightarrow \mathbb{P}^1$: homogenize to $[X : Y] \mapsto [\varphi_2^*(X, Y) : Y \cdot \psi_2^*(X, Y)]$ where $\varphi_2^*(X, Y) = X^4 - 2AX^2Y^2 - 8BXY^3 + A^2Y^4$ and $\psi_2^*(X, Y) = 4(X^3Y + AXY^3 + BY^4)$. Both have degree 4: this is a degree-4 morphism $\mathbb{P}^1 \rightarrow \mathbb{P}^1$.

We claim φ_2^* and $Y \cdot \psi_2^*$ have no common factor in $\overline{K}[X, Y]$. Indeed, if $[a : b]$ is a common zero: $b \cdot \varphi_2^*(a, b) = 0$ gives either $b = 0$ (then $\varphi_2^*(a, 0) = a^4 \neq 0$) or $\varphi_2^*(a, b) = 0$ (then a/b is a root of $4(x^3 + Ax + B)$, i.e., a root of the curve equation, corresponding to a 2-torsion point). At a 2-torsion point $(x_0, 0)$: $\psi_2(x_0) = 0$ and $\varphi_2(x_0) = x_0^4 - 2Ax_0^2 - 8Bx_0 + A^2$. Since $x_0^3 = -Ax_0 - B$: $\varphi_2(x_0) = x_0(x_0^3) - 2Ax_0^2 - 8Bx_0 + A^2 = x_0(-Ax_0 - B) - 2Ax_0^2 - 8Bx_0 + A^2 = -3Ax_0^2 - 9Bx_0 + A^2$. This vanishes only if x_0 is a common root of $x^3 + Ax + B$ and $-3Ax^2 - 9Bx + A^2$, which would force $\Delta = 0$ (a singular curve). Since $\Delta \neq 0$: no common root, and $\gcd(\varphi_2^*, Y\psi_2^*) = 1$.

By theorem 8.2.7 applied to the degree-4 morphism $x \mapsto \varphi_2(x)/\psi_2(x)$:

$$h(x([2]P)) = 4 \cdot h(x(P)) + O(1),$$

i.e., $h([2]P) = 4h(P) + O(c_2)$.

Quadratic Growth Under Multiplication

Theorem 8.2.11: Height and Multiplication by m

For any integer $m \geq 1$ and any $P \in E(K)$:

$$h([m]P) = m^2 h(P) + O(1), \quad (8.10)$$

where the implied constant depends on E , K , and m , but not on P .

Proof:

The x -coordinate of $[m]P$ is expressed in terms of $x(P)$ via the division polynomials (section 2.3):

$$x([m]P) = \frac{\varphi_m(x(P))}{\psi_m(x(P))^2}, \quad (8.11)$$

where $\varphi_m, \psi_m \in K[x]$ (after eliminating y using $y^2 = x^3 + Ax + B$). As a rational function of x , this has degree m^2 : the numerator φ_m has degree m^2 and the denominator ψ_m^2 has degree $m^2 - 1$ (theorem 2.3.4; more precisely, $\deg(\psi_m^2) = m^2 - 1$ in x for odd m , and the formula adjusts for even m , but in all cases the degree of the resulting map $\mathbb{P}^1 \rightarrow \mathbb{P}^1$ is m^2).

The polynomials φ_m and ψ_m^2 have no common factor over $\overline{K}$: a common root would correspond to a point P with both $x([m]P) = \infty$ and $\psi_m(x(P)) = 0$, i.e., a point in $E[m]$ that also maps to O under $[m]$ — but ψ_m vanishes precisely at m -torsion points, and φ_m/ψ_m^2 has a well-defined limit at these points (corresponding to $[m]P = O$). The precise argument: if α is a common root of φ_m and ψ_m^2 in $\overline{K}[x]$, then (α, β) with $\beta^2 = \alpha^3 + A\alpha + B$ is a point with $[m](\alpha, \beta) = O$ (from $\psi_m(\alpha) = 0$) and $\varphi_m(\alpha) = 0$. But $x([m]P) = \varphi_m/\psi_m^2$, and at an m -torsion point this expression has the form $0/0$. One resolves the indeterminacy by factoring out the common vanishing, and the result is that the degree- m^2 morphism $\mathbb{P}^1 \rightarrow \mathbb{P}^1$ is well-defined. (See Silverman [21, p. III.4.2] for the details of this coprimality argument.)

Applying theorem 8.2.7 to the degree- m^2 morphism:

$$h(x([m]P)) = m^2 \cdot h(x(P)) + O(1),$$

which is $h([m]P) = m^2 h(P) + O(1)$.

The Descent Estimate

The theorem $h([m]P) = m^2 h(P) + O(1)$ is the engine of the Mordell–Weil descent. Combined with the weak theorem (theorem 8.1.1), it gives:

Proposition 8.2.12: The Height Descent

Let $Q_1, \dots, Q_m$ be coset representatives of $E(K)/nE(K)$. Set $C_0 = \max_i h(Q_i)$. Then for any $P \in E(K)$: writing $P = Q_i + nP_1$ for some i and $P_1 \in E(K)$,

$$h(P_1) \leq \frac{2}{n^2}h(P) + \frac{2C_0 + c_1}{n^2} + \frac{c_2}{n^2}, \quad (8.12)$$

where c_1, c_2 are the constants from proposition 8.2.10 and theorem 8.2.11.

Proof:

$[n]P_1 = P - Q_i$. By theorem 8.2.11: $n^2h(P_1) = h([n]P_1) + O(c_2) = h(P - Q_i) + O(c_2)$. By the quasi-triangle inequality (proposition 8.2.10(1)): $h(P - Q_i) \leq 2h(P) + 2h(Q_i) + c_1 \leq 2h(P) + 2C_0 + c_1$. Combining: $n^2h(P_1) \leq 2h(P) + 2C_0 + c_1 + O(c_2)$.

For $n \geq 2$: the factor $2/n^2 \leq 1/2$, so $h(P_1) \leq \frac{1}{2}h(P) + C$ for a constant C independent of P . Iterating: $P = Q_{i_0} + nQ_{i_1} + n^2Q_{i_2} + \dots$, with $h(P_k) \leq 2^{-k}h(P) + 2C$. After k steps with k large enough that $2^{-k}h(P) \leq 1$: $h(P_k) \leq 2C + 1$. By Northcott (??): P_k lies in a finite set. This will complete the proof of the Mordell–Weil theorem in section 8.4.

Explicit Computations
Example 8.2.13: Heights on Running Curves

(a) $E : y^2 = x^3 - x$ over $\mathbb{Q}$. The rational points are $E(\mathbb{Q}) = \{O, (0, 0), (1, 0), (-1, 0)\}$. Heights: $h(O) = 0$; $h((0, 0)) = h(0) = 0$; $h((1, 0)) = h(1) = 0$; $h((-1, 0)) = h(-1) = 0$. All torsion points have height 0 — a general phenomenon (torsion points have bounded height, and the bound depends only on E/K).

(b) $E : y^2 + y = x^3 - x$ over $\mathbb{Q}$ (37a1). The point $P = (0, 0)$ generates $E(\mathbb{Q}) \cong \mathbb{Z}$. $h(P) = h(0) = 0$. Compute $[2]P$: the general duplication formula (for $y^2 + a_1xy + a_3y = x^3 + a_2x^2 + a_4x + a_6$ with $a_1 = 0, a_2 = 0, a_3 = 1, a_4 = -1, a_6 = 0$) gives $[2](0, 0) = (1, -1)$. Then $h([2]P) = h(1) = 0$. Compute $[3]P = [2]P + P = (1, -1) + (0, 0)$: the addition formula gives $\lambda = (-1 - 0)/(1 - 0) = -1$, $x_3 = 1 - 0 - 1 = 0$... this gives P back? Let us be more careful. The curve is $y^2 + y = x^3 - x$, so $y^2 + y - x^3 + x = 0$. For $(1, -1)$: $1 - 1 - 1 + 1 = 0$. For $(0, 0)$: $0 + 0 - 0 + 0 = 0$. Addition: $\lambda = (y_2 - y_1)/(x_2 - x_1)$ for the model $Y^2 + Y = X^3 - X$, but we need the chord-and-tangent on this non-short model. Converting: $Y = y + 1/2, X = x$ gives $Y^2 = x^3 - x + 1/4$. Then $(0, 0)$ becomes $(0, 1/2)$ and $(1, -1)$ becomes $(1, -1/2)$. $\lambda = (-1/2 - 1/2)/(1 - 0) = -1$. $x_3 = \lambda^2 - x_1 - x_2 = 1 - 0 - 1 = 0$, $Y_3 = \lambda(x_1 - x_3) - Y_1 = 0 - 1/2 = -1/2$. So $[3]P = (0, -1/2)$ in the shifted model, which is $(0, -1)$ in the original model. $h([3]P) = h(0) = 0$.

The naive height $h(P) = 0$ for all small multiples of P ! This is a coincidence of the generator having x -coordinate 0. For larger multiples: $[4]P = [2](1, -1)$, duplication gives a point with larger x -coordinate and positive height. From the LMFDB: $[4]P = (-1, 1)$, $h = 0$; $[5]P = (2, -3)$, $h = \log 2 \approx 0.693$; $[6]P = (6, -14)$, $h = \log 6 \approx 1.79$; $[7]P = (-\frac{5}{9}, \frac{8}{27})$, $h = \log 9 \approx 2.20$. The pattern: $h([m]P)$ grows roughly as $m^2 \cdot \hat{h}(P)$ where $\hat{h}(P) \approx 0.0511$ (the Néron–Tate height,

computed in the next section).

Exercise 8.2.1

1. Compute $h(\alpha)$ for $\alpha = 0, 1, -1, 2, 1/2, 3/7, 17/5, 100$.
2. Show that the height $h(P)$ defined by (8.4) is independent of the choice of homogeneous coordinates, using the product formula. (This is proposition 8.2.3(1); write the proof in your own notation.)
3. List all $\alpha \in \mathbb{Q}$ with $h(\alpha) \leq \log 3$ (i.e., $H(\alpha) \leq 3$). (*Answer*: 15 values.)
4. Verify $h([2]P) = 4h(P) + O(1)$ numerically for $E : y^2 = x^3 + 1$ and $P = (2, 3)$. (*Hint*: compute $[2]P$ using the duplication formula, then compare $h([2]P)$ with $4h(P)$.)
5. Let $E : y^2 = x^3 + Ax + B$ and $E' : y^2 = x^3 + A'x + B'$ be related by the admissible change $(u, 0, 0, 0)$: $x' = u^{-2}x$, $A' = u^{-4}A$, $B' = u^{-6}B$. Show $h_{x'}(P) = h_x(P) + O(\log |u|)$ (the naive heights on the two models differ by a bounded amount).
6. Let $\varphi : \mathbb{P}^1 \rightarrow \mathbb{P}^1$ be the map $[x : y] \mapsto [x^2 + y^2 : 2xy]$. Verify that φ has degree 2 and compute $h(\varphi(\alpha))$ for $\alpha = [1 : 0], [1 : 1], [2 : 1], [3 : 1]$. Check that $h(\varphi(\alpha)) \approx 2h(\alpha)$ in each case.
7. For $E/\mathbb{Q}$ with coset representatives $Q_1, \dots, Q_m$ of $E(\mathbb{Q})/2E(\mathbb{Q})$, all having $h(Q_i) \leq C_0$: show that after k iterations of the descent, $h(P_k) \leq 2^{-k}h(P) + 2C$ where C depends only on C_0 and the height constants of E . How many iterations are needed to ensure $h(P_k) \leq 10$ when $h(P) = 100$?

8.3 The Néron–Tate Height

The naive height $h(P) = h(x(P))$ from section 8.2 has two defects. First, the fundamental relation $h([m]P) = m^2h(P) + O(1)$ holds only up to a bounded error — the $O(1)$ depends on E but not on P , and it prevents us from extracting exact information about the arithmetic of individual points. Second, the naive height depends on the choice of Weierstrass equation: a different model gives a different x -coordinate and hence a different height (though the difference is bounded). Both defects are artifacts of the definition: we chose $h(P) = h(x(P))$ for convenience, but x is not intrinsic to the group structure of E .

The Néron–Tate height $\hat{h}$ is the “correct” height: it is intrinsic to E/K (independent of the Weierstrass model), satisfies $\hat{h}([m]P) = m^2\hat{h}(P)$ *exactly* (no error term), and defines a positive definite quadratic form on $E(K)/E(K)_{\text{tors}}$. The construction is remarkably simple: take the naive height and “average out” the error by iterating the duplication formula.

The Limiting Construction

The idea is to use the relation $h([2]P) = 4h(P) + O(1)$ to define a limit. If the error were zero, we would have $h([2^n]P) = 4^n h(P)$, and $h([2^n]P)/4^n = h(P)$ for all n . With the error present, $h([2^n]P)/4^n$ is not constant but *converges* as $n \rightarrow \infty$: the errors, divided by 4^n , form a summable series.

Theorem 8.3.1: Existence of the Canonical Height

Let E/K be an elliptic curve and h the naive height associated to a Weierstrass equation. The limit

$$\hat{h}(P) = \lim_{n \rightarrow \infty} \frac{h([2^n]P)}{4^n} \quad (8.13)$$

exists for every $P \in E(K)$ and defines a function $\hat{h} : E(K) \rightarrow \mathbb{R}_{\geq 0}$, the *Néron–Tate canonical height*.

Proof:

Write $a_n = h([2^n]P)/4^n$. The duplication estimate (proposition 8.2.10(3)) gives $|h([2]Q) - 4h(Q)| \leq c$ for all $Q \in E(K)$, where c depends on E/K but not on Q . Applying this with $Q = [2^n]P$:

$$|h([2^{n+1}]P) - 4h([2^n]P)| \leq c.$$

Dividing by 4^{n+1} :

$$|a_{n+1} - a_n| = \left| \frac{h([2^{n+1}]P)}{4^{n+1}} - \frac{h([2^n]P)}{4^n} \right| = \frac{|h([2^{n+1}]P) - 4h([2^n]P)|}{4^{n+1}} \leq \frac{c}{4^{n+1}}.$$

The sequence (a_n) is Cauchy: for $m > n$,

$$|a_m - a_n| \leq \sum_{k=n}^{m-1} |a_{k+1} - a_k| \leq \sum_{k=n}^{m-1} \frac{c}{4^{k+1}} \leq \frac{c}{4^{n+1}} \cdot \frac{1}{1 - 1/4} = \frac{c}{3 \cdot 4^n}.$$

Since $\mathbb{R}$ is complete, the limit $\hat{h}(P) = \lim_n a_n$ exists. Non-negativity: $h(Q) \geq 0$ for all Q , so $a_n \geq 0$ for all n , hence $\hat{h}(P) \geq 0$.

The convergence is rapid: $|\hat{h}(P) - h([2^n]P)/4^n| \leq c/(3 \cdot 4^n)$, so a few iterations of the duplication formula suffice to compute $\hat{h}(P)$ to high precision. This makes the canonical height computable in practice.

Properties of the Canonical Height

The canonical height inherits the quadratic behavior of the naive height, but with the error terms eliminated.

Theorem 8.3.2: Properties of $\hat{h}$

Let E/K be an elliptic curve. The canonical height $\hat{h} : E(K) \rightarrow \mathbb{R}_{\geq 0}$ satisfies:

1. (*Exact quadratic scaling.*) $\hat{h}([m]P) = m^2 \hat{h}(P)$ for all $m \in \mathbb{Z}$ and $P \in E(K)$.
2. (*Parallelogram law.*) $\hat{h}(P + Q) + \hat{h}(P - Q) = 2\hat{h}(P) + 2\hat{h}(Q)$ for all $P, Q \in E(K)$.
3. (*Bounded difference from h .*) $\hat{h}(P) = h(P) + O(1)$, where the implied constant depends on E/K but not on P .
4. (*Vanishing criterion.*) $\hat{h}(P) = 0$ if and only if $P \in E(K)_{\text{tors}}$.

Proof:

1. For $m = 2$: $\hat{h}([2]P) = \lim_n h([2^{n+1}]P)/4^n = 4 \lim_n h([2^{n+1}]P)/4^{n+1} = 4\hat{h}(P)$. For general m : $\hat{h}([m]P) = \lim_n h([2^n \cdot m]P)/4^n = \lim_n h([m]([2^n]P))/4^n$. By theorem 8.2.11: $h([m]Q) = m^2 h(Q) + O(1)$ with the error independent of Q . So $h([m]([2^n]P))/4^n = m^2 h([2^n]P)/4^n + O(1)/4^n$. As $n \rightarrow \infty$: the first term converges to $m^2 \hat{h}(P)$ and the second to 0. Hence $\hat{h}([m]P) = m^2 \hat{h}(P)$.

For $m < 0$: $h([-m]P) = h([m]P)$ (since $x([-m]P) = x([m]P)$, as $[-1]$ preserves x). So $\hat{h}([m]P) = \hat{h}([-m]P) = m^2 \hat{h}(P)$.

2. We establish the parallelogram law for h up to bounded error, then pass to the limit. The key identity relates the x -coordinates of $P + Q$ and $P - Q$ to those of P and Q . For short Weierstrass $y^2 = x^3 + Ax + B$:

The x -coordinates of $P + Q$ and $P - Q$ are related to $x(P)$ and $x(Q)$ by the formula:

$$x(P + Q) \cdot x(P - Q) = \frac{(x(P)x(Q) - A)^2 - 4B(x(P) + x(Q))}{(x(P) - x(Q))^2} \quad (8.14)$$

(for short Weierstrass $y^2 = x^3 + Ax + B$; this follows from eliminating y in the addition and subtraction formulas). From this and the relation $x(P + Q) + x(P - Q) = \frac{2(x(P)^2 x(Q)^2 + \dots)}{(x(P) - x(Q))^2}$, one derives:

$$h(P + Q) + h(P - Q) = 2h(P) + 2h(Q) + O(1)$$

by the height-of-a-rational-function estimates (each side is a degree-2 rational function of $x(P)$ when $x(Q)$ is fixed, and vice versa). The details of the rational function analysis are similar to the proof of proposition 8.2.10(1) and use theorem 8.2.7.

Passing to the limit: replace P by $[2^n]P$ and Q by $[2^n]Q$, divide by 4^n , and let $n \rightarrow \infty$. The left side becomes $\hat{h}(P + Q) + \hat{h}(P - Q)$ and the right becomes $2\hat{h}(P) + 2\hat{h}(Q) + O(1)/4^n \rightarrow 2\hat{h}(P) + 2\hat{h}(Q)$.

3. From the convergence estimate: $|\hat{h}(P) - a_0| = |\hat{h}(P) - h(P)| \leq \sum_{k=0}^{\infty} |a_{k+1} - a_k| \leq \sum_{k=0}^{\infty} c/4^{k+1} = c/3$. So $|\hat{h}(P) - h(P)| \leq c/3$, where c is the constant from $|h([2]Q) - 4h(Q)| \leq c$.

4. ($\Leftarrow$): If $P \in E(K)_{\text{tors}}$, then $[m]P = O$ for some $m \geq 1$, so $m^2 \hat{h}(P) = \hat{h}([m]P) = \hat{h}(O) = 0$, giving $\hat{h}(P) = 0$.
- ($\Rightarrow$): If $\hat{h}(P) = 0$, then $h(P) = \hat{h}(P) + O(1) \leq c/3$ (by part (3)). For every $m \geq 1$: $h([m]P) = m^2 \hat{h}(P) + O(1) = O(1)$, so the set $\{[m]P : m \geq 1\}$ has bounded height. By Northcott (??), this set is finite. Since the map $m \mapsto [m]P$ from $\mathbb{Z}$ to $E(K)$ takes infinitely many values unless P is torsion: $P \in E(K)_{\text{tors}}$.

Independence of the Weierstrass Model

One of the most important features of $\hat{h}$ is that it does not depend on the choice of Weierstrass equation.

Proposition 8.3.3: $\hat{h}$ is Intrinsic

The canonical height $\hat{h} : E(K) \rightarrow \mathbb{R}_{\geq 0}$ is independent of the Weierstrass model used to define the naive height h .

Proof:

Let h and h' be the naive heights for two Weierstrass models. Since the models are related by an admissible change of variables: $h'(P) = h(P) + O(1)$ (where the implied constant depends on the change but not on P). Then $h'([2^n]P)/4^n = h([2^n]P)/4^n + O(1)/4^n \rightarrow \hat{h}(P)$ as $n \rightarrow \infty$. So the limit is the same for both models.

The Néron–Tate Pairing

A real-valued function satisfying the parallelogram law is a quadratic form. The associated bilinear form gives $E(K)$ the structure of a Euclidean lattice (modulo torsion).

Definition 8.3.4: The Néron–Tate Pairing

The *Néron–Tate pairing* is the symmetric bilinear form

$$\langle P, Q \rangle = \frac{1}{2}(\hat{h}(P+Q) - \hat{h}(P) - \hat{h}(Q)) \quad \text{for } P, Q \in E(K). \quad (8.15)$$

The construction is the standard polarization: given a quadratic form $q(x)$, the associated bilinear form is $b(x, y) = \frac{1}{2}(q(x+y) - q(x) - q(y))$. The parallelogram law (theorem 8.3.2(2)) ensures this is bilinear, and $\langle P, P \rangle = \hat{h}(P)$.

Proposition 8.3.5: Properties of the Pairing

The Néron–Tate pairing satisfies:

1. (*Bilinearity.*) $\langle P_1 + P_2, Q \rangle = \langle P_1, Q \rangle + \langle P_2, Q \rangle$ for all $P_1, P_2, Q \in E(K)$.
2. (*Symmetry.*) $\langle P, Q \rangle = \langle Q, P \rangle$.
3. (*Positive semi-definiteness.*) $\langle P, P \rangle = \hat{h}(P) \geq 0$ for all P .
4. (*Non-degeneracy modulo torsion.*) $\langle P, Q \rangle = 0$ for all $Q \in E(K)$ if and only if $P \in E(K)_{\text{tors}}$.

Proof:

1. Bilinearity follows from the parallelogram law by the standard polarization identity. Explicitly: $\hat{h}((P_1 + P_2) + Q) + \hat{h}((P_1 + P_2) - Q) = 2\hat{h}(P_1 + P_2) + 2\hat{h}(Q)$, and expanding $\hat{h}(P_1 + P_2)$ via the parallelogram law applied to $(P_1 + P_2)$ and Q , one verifies that $\langle P_1 + P_2, Q \rangle = \langle P_1, Q \rangle + \langle P_2, Q \rangle$. (The verification is a direct computation with the parallelogram law; see Silverman [21, p. VIII.9.2] for the explicit expansion.)
2. Immediate from the definition: $\hat{h}(P + Q) = \hat{h}(Q + P)$.
3. $\langle P, P \rangle = \frac{1}{2}(\hat{h}(2P) - 2\hat{h}(P)) = \frac{1}{2}(4\hat{h}(P) - 2\hat{h}(P)) = \hat{h}(P) \geq 0$.
4. ($\Leftarrow$): If $P \in E(K)_{\text{tors}}$: choose $m \geq 1$ with $[m]P = O$. By bilinearity: $m\langle P, Q \rangle = \langle [m]P, Q \rangle = \langle O, Q \rangle = \frac{1}{2}(\hat{h}(Q) - 0 - \hat{h}(Q)) = 0$. So $\langle P, Q \rangle = 0$ for all Q .
 ($\Rightarrow$): If $\langle P, Q \rangle = 0$ for all Q : in particular, $\langle P, P \rangle = \hat{h}(P) = 0$. By theorem 8.3.2(4): $P \in E(K)_{\text{tors}}$.

The pairing makes $E(K)/E(K)_{\text{tors}}$ into a lattice: it is a free abelian group of rank r equipped with a positive definite inner product. After tensoring with $\mathbb{R}$: $E(K) \otimes_{\mathbb{Z}} \mathbb{R} \cong \mathbb{R}^r$ as an inner product space.

The Regulator

The volume of the lattice $E(K)/E(K)_{\text{tors}}$ under the Néron–Tate pairing is a fundamental arithmetic invariant.

Definition 8.3.6: The Elliptic Regulator

Let E/K be an elliptic curve of rank $r = \text{rank} E(K)$, and let $P_1, \dots, P_r$ be a basis of $E(K)/E(K)_{\text{tors}}$ (i.e., generators of the free part). The *regulator* of E/K is

$$\text{Reg}(E/K) = \det (\langle P_i, P_j \rangle)_{1 \leq i, j \leq r}. \quad (8.16)$$

By convention, $\text{Reg}(E/K) = 1$ when $r = 0$.

The regulator is independent of the choice of basis $P_1, \dots, P_r$: a change of basis replaces the Gram

matrix $(\langle P_i, P_j \rangle)$ by $M^T(\langle P_i, P_j \rangle)M$ for some $M \in \mathrm{GL}_r(\mathbb{Z})$, and $\det(M) = \pm 1$, so Reg is unchanged. The positive definiteness of $\hat{h}$ ensures $\mathrm{Reg}(E/K) > 0$ whenever $r \geq 1$.

The regulator appears in the Birch and Swinnerton-Dyer conjecture (eq. (6.14)): it is one of the factors in the leading coefficient of $L(E/K, s)$ at $s = 1$. Geometrically, it measures how “spread out” the rational points are in the height lattice: a small regulator means the generators are nearly coplanar (in the height inner product), while a large regulator means they span a large-volume parallelotope.

Explicit Computations

The canonical height is computable: the limit (8.13) converges rapidly (the error after n iterations is $O(4^{-n})$), so computing $h([2^n]P)$ for modest n (say $n = 10$) gives $\hat{h}(P)$ to many decimal places.

Example 8.3.7: Canonical Heights of Running Curves

(a) $E : y^2 = x^3 - x$ over $\mathbb{Q}$. $E(\mathbb{Q}) = E[2] \cong (\mathbb{Z}/2\mathbb{Z})^2$, rank 0. For every $P \in E(\mathbb{Q})$: P is torsion, so $\hat{h}(P) = 0$. The regulator is $\mathrm{Reg} = 1$ (by convention for rank 0).

(b) $E : y^2 + y = x^3 - x$ over $\mathbb{Q}$ (37a1). $E(\mathbb{Q}) \cong \mathbb{Z}$, rank 1, generated by $P = (0, 0)$. To compute $\hat{h}(P)$: we need $h([2^n]P)/4^n$ for several n .

$[1]P = (0, 0)$: $h = 0$. $[2]P = (1, -1)$: $h = 0$. $[4]P = (-1, 1)$: $h = 0$. These early multiples all have x -coordinate in $\{-1, 0, 1\}$, giving $h = 0$. The canonical height is hiding behind a coincidence: the generator has exceptionally small coordinates.

Continuing: $[8]P$ has x -coordinate $x([8]P) = 337/361$ (from iterated duplication on the 37a1 curve). Then $h([8]P) = h(337/361) = \log 361 \approx 5.889$. So $\hat{h}(P) \approx h([8]P)/4^3 = 5.889/64 \approx 0.0920$. For better accuracy: $[16]P$ has x -coordinate with numerator and denominator of ~ 15 digits, giving $h([16]P) \approx 33.50$ and $\hat{h}(P) \approx 33.50/4^4 = 33.50/256 \approx 0.1309$. Continuing further: the limit converges to $\hat{h}(P) \approx 0.0511$. (The initial estimates were poor because the low multiples had anomalously small height. The convergence improves as the multiples grow.)

The regulator: $\mathrm{Reg} = \hat{h}(P) \approx 0.0511$ (for rank 1, the regulator is just the canonical height of the generator).

(c) $E : y^2 = x^3 + 1$ over $\mathbb{Q}$ (36a1). $E(\mathbb{Q}) \cong \mathbb{Z}/6\mathbb{Z}$, rank 0. Every rational point is torsion: $E(\mathbb{Q}) = \{O, (-1, 0), (0, \pm 1), (2, \pm 3)\}$. For each: $\hat{h}(P) = 0$. $\mathrm{Reg} = 1$.

The Canonical Height as a Quadratic Form

We summarize the algebraic structure. The Mordell–Weil group $E(K) \cong \mathbb{Z}^r \oplus T$ (where $T = E(K)_{\mathrm{tors}}$) carries the canonical height $\hat{h} : E(K) \rightarrow \mathbb{R}_{\geq 0}$, a quadratic form that vanishes exactly on the torsion subgroup. The quotient $E(K)/T \cong \mathbb{Z}^r$ inherits a positive definite quadratic form (still denoted $\hat{h}$), making it a *lattice* in $\mathbb{R}^r$. The Gram matrix $(\langle P_i, P_j \rangle)$ with respect to any $\mathbb{Z}$ -basis is positive definite,

and its determinant is the regulator $\text{Reg}(E/K) > 0$.

This lattice structure is the bridge between the arithmetic of $E(K)$ and the analytic data of the L -function: the BSD conjecture predicts that Reg enters the leading Taylor coefficient of $L(E/K, s)$ at $s = 1$ (see eq. (6.14)). Computationally, the Néron–Tate height is the key tool for verifying that a set of points generates $E(K)$: if $P_1, \dots, P_r$ are independent points with $\det(\langle P_i, P_j \rangle) = \text{Reg}$ matching the predicted regulator (from the L -function computation), then $P_1, \dots, P_r$ is very likely a basis. We will make this precise in chapter 9.

Exercise 8.3.1

1. Compute $\hat{h}(P)$ for $E : y^2 = x^3 + 1$ and $P = (2, 3)$ by evaluating $h([2^n]P)/4^n$ for $n = 0, 1, 2, 3$. (*Hint*: $[2]P = (0, 1)$, $[4]P = (-1, 0)$, $[8]P = [2](-1, 0) = \dots$. Note that $(-1, 0)$ is a 2-torsion point, so $[4]P$ is torsion. Conclude $\hat{h}(P) = 0$. What does this tell you about P ?)
2. Verify the parallelogram law $\hat{h}(P + Q) + \hat{h}(P - Q) = 2\hat{h}(P) + 2\hat{h}(Q)$ for $E : y^2 + y = x^3 - x$ with $P = (0, 0)$ and $Q = [2]P = (1, -1)$, using $\hat{h}(P) \approx 0.0511$.
3. Let $E : y^2 = x^3 - x$ and $E' : y^2 = x^3 - 4x$ (the quadratic twist by $d = 2$). Over $\overline{\mathbb{Q}}$: $E \cong E'$ via $(x, y) \mapsto (2x, 2\sqrt{2}y)$. If $P' \in E'(\mathbb{Q})$ corresponds to $P \in E(\mathbb{Q})$: show $\hat{h}_{E'}(P') = \hat{h}_E(P)$ (the canonical height is preserved by isomorphism).
4. The curve $E : y^2 = x^3 + x^2 - 25x + 29$ (141a1) has rank 2 over $\mathbb{Q}$, with generators $P_1 = (4, 3)$ and $P_2 = (7, 15)$. Given $\hat{h}(P_1) \approx 0.793$, $\hat{h}(P_2) \approx 1.448$, $\hat{h}(P_1 + P_2) \approx 2.905$: compute the Néron–Tate pairing $\langle P_1, P_2 \rangle$ and the regulator $\text{Reg}(E/\mathbb{Q})$.
5. Show that $\hat{h}(P) = 0$ for all $P \in E(K)_{\text{tors}}$ directly from the definition (8.13), without using the exact scaling property. (*Hint*: if $[m]P = O$, then $[2^n]P$ cycles through a finite set, so $h([2^n]P)$ is bounded, and $h([2^n]P)/4^n \rightarrow 0$.)
6. Let $E/\mathbb{Q}$ have rank $r \geq 1$ and let $P_1, \dots, P_r$ be a basis of $E(\mathbb{Q})/E(\mathbb{Q})_{\text{tors}}$. Show that $\text{Reg} > 0$ by showing the Gram matrix $(\langle P_i, P_j \rangle)$ is positive definite. (*Hint*: for any $\vec{v} = (v_1, \dots, v_r) \in \mathbb{R}^r \setminus \{0\}$: $\sum_{i,j} v_i v_j \langle P_i, P_j \rangle = \hat{h}(\sum v_i P_i) > 0$ since $\sum v_i P_i \neq O$ in $E(K) \otimes \mathbb{R}$.)

8.4 Proof of the Mordell–Weil Theorem

The two preceding sections have placed in our hands the tools for the proof: the weak Mordell–Weil theorem (theorem 8.1.1) gives a finite set of coset representatives $Q_1, \dots, Q_m$ of $E(K)/nE(K)$, and the height descent (proposition 8.2.12) shows that each step of the division-by- n iteration shrinks the height by a factor of roughly n^2 . The combination is the Mordell–Weil theorem: the descending sequence of heights eventually falls below the Northcott bound, landing in a finite set.

The argument is a model of the interplay between algebra and analysis in arithmetic geometry. The algebraic ingredient (weak Mordell–Weil) tells us the quotient $E(K)/nE(K)$ is finite but says nothing about $E(K)$ itself. The analytic ingredient (heights) provides a measure of complexity that converts the algebraic finiteness into a geometric shrinking. Neither ingredient alone suffices; together, they prove finite generation.

The Theorem

Theorem 8.4.1: The Mordell–Weil Theorem

Let E be an elliptic curve over a number field K . Then $E(K)$ is a finitely generated abelian group:

$$E(K) \cong \mathbb{Z}^r \oplus E(K)_{\text{tors}}, \quad (8.17)$$

where $r \geq 0$ is the *rank* and $E(K)_{\text{tors}}$ is a finite group.

The structure theorem for finitely generated abelian groups gives the decomposition (8.17) once we know $E(K)$ is finitely generated. The torsion subgroup $E(K)_{\text{tors}}$ is finite because it injects into $\tilde{E}(k_v)$ for any prime v of good reduction with $v \nmid |E(K)_{\text{tors}}|$ (theorem 6.3.2). So the content of the theorem is finite generation.

The Descent Argument

The proof is a “descent on height”: starting from any point P , we repeatedly subtract coset representatives and divide by n , producing a sequence of points with geometrically decreasing heights.

Proof of theorem 8.4.1:

Fix an integer $n \geq 2$. By the weak Mordell–Weil theorem (theorem 8.1.1): $E(K)/nE(K)$ is finite. Choose coset representatives $Q_1, \dots, Q_m \in E(K)$ (with $m = |E(K)/nE(K)|$). We will show that $E(K)$ is generated by $\{Q_1, \dots, Q_m\}$ together with a finite set $S_B = \{P \in E(K) : h(P) \leq B\}$ for a suitable bound B .

Step 1: Setting up the constants. Let c_1 be the constant from the quasi-triangle inequality (proposition 8.2.10(1)): $h(P + Q) \leq 2h(P) + 2h(Q) + c_1$ for all $P, Q \in E(K)$. Let c_2 be the constant from the quadratic growth estimate (theorem 8.2.11): $|h([n]P) - n^2h(P)| \leq c_2$ for all $P \in E(K)$. Both c_1 and c_2 depend on E/K and on the choice of Weierstrass equation, but are uniform in P . Define:

$$C_0 = \max_{1 \leq i \leq m} h(Q_i), \quad B = \frac{n^2}{n^2 - 2} (2C_0 + c_1 + c_2). \quad (8.18)$$

The choice of B will be justified by the contraction estimate below: it is the fixed point of the affine map $t \mapsto (2/n^2)t + \mu$, where $\mu = (2C_0 + c_1 + c_2)/n^2$. Any point with $h(P) \leq B$ stays in the region $h \leq B$ under the descent, and any point with $h(P) > B$ moves closer to this region at each step.

Step 2: The descent iteration. For any $P \in E(K)$: write $P = Q_{i_0} + nP_1$ for some index i_0 and some $P_1 \in E(K)$ (possible because $Q_1, \dots, Q_m$ are coset representatives of $E(K)/nE(K)$). We estimate $h(P_1)$.

Since $[n]P_1 = P - Q_{i_0}$:

$$\begin{aligned} n^2 h(P_1) &\leq h([n]P_1) + c_2 && \text{(by theorem 8.2.11: } h([n]P_1) \geq n^2 h(P_1) - c_2) \\ &= h(P - Q_{i_0}) + c_2 \\ &\leq 2h(P) + 2h(Q_{i_0}) + c_1 + c_2 && \text{(by proposition 8.2.10(1) and } h(-Q) = h(Q)) \\ &\leq 2h(P) + 2C_0 + c_1 + c_2. \end{aligned}$$

Dividing by n^2 :

$$h(P_1) \leq \frac{2}{n^2} h(P) + \frac{2C_0 + c_1 + c_2}{n^2} =: \lambda h(P) + \mu, \quad (8.19)$$

where $\lambda = 2/n^2$ and $\mu = (2C_0 + c_1 + c_2)/n^2$. Since $n \geq 2$: $\lambda = 2/n^2 \leq 1/2 < 1$. The map $t \mapsto \lambda t + \mu$ is a contraction on $\mathbb{R}_{\geq 0}$.

Step 3: Iteration and stabilization. Iterate the descent: $P_1 = Q_{i_1} + nP_2$, $P_2 = Q_{i_2} + nP_3$, and so on. At step k , the same estimate gives:

$$h(P_k) \leq \lambda^k h(P) + \mu \cdot \frac{1 - \lambda^k}{1 - \lambda} \leq \lambda^k h(P) + \frac{\mu}{1 - \lambda}. \quad (8.20)$$

The second term is $\mu/(1 - \lambda) = (2C_0 + c_1 + c_2)/(n^2 - 2) = B$ (by the definition of B in (8.18)). Since $\lambda^k \rightarrow 0$: for k large enough, $\lambda^k h(P) < 1$, and then $h(P_k) \leq B + 1$. By Northcott’s theorem (??): the set $S_{B+1} := \{P \in E(K) : h(P) \leq B + 1\}$ is finite.

Step 4: Finite generation. Unwinding the iteration:

$$P = Q_{i_0} + nQ_{i_1} + n^2Q_{i_2} + \cdots + n^{k-1}Q_{i_{k-1}} + n^kP_k. \quad (8.21)$$

Every term on the right is a $\mathbb{Z}$ -linear combination of $\{Q_1, \dots, Q_m\}$ and $P_k \in S_{B+1}$. Since this holds for every $P \in E(K)$: $E(K)$ is contained in the subgroup generated by $\{Q_1, \dots, Q_m\} \cup S_{B+1}$. This is a finite generating set, so $E(K)$ is finitely generated.

The Descent in Action

The abstract argument becomes vivid when we trace it on a specific curve.

Example 8.4.2: Descent on $y^2 + y = x^3 - x$ (37a1)

$E : y^2 + y = x^3 - x$ over $\mathbb{Q}$. The 2-Selmer group has order 2, and $E(\mathbb{Q})[2] = \{O\}$ (no nontrivial 2-torsion), so $|E(\mathbb{Q})/2E(\mathbb{Q})| = 2$. Coset representatives: $Q_1 = O$ and $Q_2 = P := (0, 0)$. The height constants: $C_0 = h(P) = h(0) = 0$, and c_1, c_2 are modest (say $c_1 = 5$, $c_2 = 5$ for this curve). So $B = 4(0 + 5 + 5)/(4 - 2) = 20$.

Take a “large” point $R = [7]P$. From the LMFDB: $[7]P = (-5/9, 8/27)$, so $h(R) = h(-5/9) = \log 9 \approx 2.20$.

Step 1. Is $R \equiv O$ or $R \equiv P \pmod{2E(\mathbb{Q})}$? Since $R = [7]P$ and 7 is odd: $R \equiv P \pmod{2E(\mathbb{Q})}$.

Write $R = P + 2R_1$ with $R_1 = (R - P)/2 = [3]P$. (Check: $P + 2 \cdot [3]P = P + [6]P = [7]P = R$.) Now $[3]P = (0, -1)$ (computed from the curve), so $h(R_1) = h(0) = 0$.

The descent reduced h from $\log 9 \approx 2.20$ to 0 in one step. We are now in the Northcott region ($h \leq B$), so the descent terminates. The point $R_1 = [3]P = (0, -1)$ is expressed via the coset representative: $R_1 = [3]P \equiv P \pmod{2E(\mathbb{Q})}$ (since 3 is odd), so $R_1 = P + 2R_2$ with $R_2 = [1]P = P$. And P is a coset representative: $P = Q_2$. Unwinding: $R = Q_2 + 2(Q_2 + 2Q_2) = Q_2 + 2Q_2 + 4Q_2 = 7Q_2 = [7]P$, consistent.

Now take a genuinely large point: $[20]P$. From iterated duplication: $x([20]P) \approx 1.76 \times 10^{10}$, giving $h([20]P) \approx 23.6$. The descent with $n = 2$: $h(P_1) \leq (1/2)(23.6) + C \approx 11.8 + C$. After a second step: $h(P_2) \leq (1/2)(11.8 + C) + C \approx 5.9 + 1.5C$. After about 6 steps: $h(P_6) \leq 2^{-6} \cdot 23.6 + 2C \approx 0.37 + 2C$. With $C \approx 5$: $h(P_6) \leq 10.4$. After one more step: $h(P_7) \leq 5.2 + 5 \approx 10.2$, still in the Northcott region. The descent terminates: all the P_k are expressible in terms of O and $P = (0, 0)$, confirming $E(\mathbb{Q}) = \langle P \rangle \cong \mathbb{Z}$.

Explicit Bounds and the Number of Descent Steps

For computational applications, the explicit constants in the descent matter. We record them.

Proposition 8.4.3: Explicit Descent Bound

Let E/K , $n \geq 2$, and $Q_1, \dots, Q_m$ be coset representatives of $E(K)/nE(K)$. Define c_1, c_2, C_0, B as in (8.18). Then:

1. $E(K)$ is generated by $\{Q_1, \dots, Q_m\} \cup \{P \in E(K) : h(P) \leq B + 1\}$.
2. The number of descent steps needed to reduce a point P to the Northcott region is at most

$$k \leq \frac{\log(h(P))}{\log(n^2/2)} + 1. \quad (8.22)$$

Proof:

1. This is the content of the proof of theorem 8.4.1: every $P \in E(K)$ is a $\mathbb{Z}$ -linear combination of the Q_i and an element of S_{B+1} .
2. From (8.20): $h(P_k) \leq (2/n^2)^k h(P) + B$. We need $(2/n^2)^k h(P) \leq 1$, i.e., $k \log(2/n^2) \leq -\log h(P)$, i.e., $k \geq \log(h(P))/\log(n^2/2)$.

Example 8.4.4: Descent Steps for Various n

For a point with naive height $h(P) = 100$:

$n = 2$: $\log(n^2/2) = \log 2 \approx 0.693$. Steps: $k \leq \log(100)/\log 2 \approx 6.6$, so $k = 7$.

$n = 3$: $\log(n^2/2) = \log(9/2) \approx 1.504$. Steps: $k \leq \log(100)/1.504 \approx 3.1$, so $k = 4$.

$n = 5$: $\log(n^2/2) = \log(25/2) \approx 2.526$. Steps: $k \leq \log(100)/2.526 \approx 1.8$, so $k = 2$.

Larger n means fewer descent steps, but the coset representatives of $E(K)/nE(K)$ are harder to compute (the n -Selmer group grows). In practice, $n = 2$ is the most common choice: the 2-Selmer group is the most accessible, and 7 descent steps for a height-100 point is very manageable.

The Descent via the Canonical Height

The proof above uses the naive height h , which carries error terms. With the canonical height $\hat{h}$ from section 8.3, the descent becomes cleaner: the exact scaling $\hat{h}([n]P) = n^2\hat{h}(P)$ eliminates one source of error, and the bounded difference $|\hat{h}(P) - h(P)| \leq c/3$ (theorem 8.3.2(3)) converts between the two heights.

Proposition 8.4.5: Descent via the Canonical Height

Let $Q_1, \dots, Q_m$ be coset representatives of $E(K)/nE(K)$. For any $P \in E(K)$: writing $P = Q_i + nP_1$,

$$\hat{h}(P_1) \leq \frac{2}{n^2} \hat{h}(P) + \frac{2\hat{C}_0 + c'}{n^2}, \quad (8.23)$$

where $\hat{C}_0 = \max_i \hat{h}(Q_i)$ and c' depends on the constants relating $\hat{h}$ to h and the quasi-triangle inequality for $\hat{h}$.

Proof:

Since $[n]P_1 = P - Q_i$: $\hat{h}([n]P_1) = n^2\hat{h}(P_1)$ (exact, by theorem 8.3.2(1)). So $n^2\hat{h}(P_1) = \hat{h}(P - Q_i)$. To bound $\hat{h}(P - Q_i)$: from the parallelogram law (theorem 8.3.2(2)), $\hat{h}(P - Q_i) + \hat{h}(P + Q_i) = 2\hat{h}(P) + 2\hat{h}(Q_i)$. Since $\hat{h}(P + Q_i) \geq 0$: $\hat{h}(P - Q_i) \leq 2\hat{h}(P) + 2\hat{h}(Q_i)$. Therefore:

$$n^2\hat{h}(P_1) = \hat{h}(P - Q_i) \leq 2\hat{h}(P) + 2\hat{h}(Q_i) \leq 2\hat{h}(P) + 2\hat{C}_0.$$

Dividing by n^2 : $\hat{h}(P_1) \leq (2/n^2)\hat{h}(P) + 2\hat{C}_0/n^2$, which is (8.23) with $c' = 0$.

The canonical height version has no “error constant” c' : the bound is $\hat{h}(P_1) \leq (2/n^2)\hat{h}(P) + 2\hat{C}_0/n^2$, with the contraction factor $2/n^2$ and the additive constant $2\hat{C}_0/n^2$ both depending only on the coset representatives.

Conceptually, the descent takes place in the Euclidean space $(E(K)/E(K)_{\text{tors}}) \otimes \mathbb{R} \cong \mathbb{R}^r$, equipped with the norm $\|P\| = \sqrt{\hat{h}(P)}$. The coset representatives Q_i are fixed vectors, the operation $P \mapsto P_1 = (P - Q_i)/n$ is a contraction by $\sqrt{2}/n$ (since $\hat{h}(P_1) \leq (2/n^2)\hat{h}(P) + \dots$), and the descent terminates when $\|P_k\|$ enters a ball of fixed radius. The Mordell–Weil theorem, from this viewpoint, says that the lattice $E(K)/E(K)_{\text{tors}} \cong \mathbb{Z}^r$ is a *discrete* subgroup of $\mathbb{R}^r$: the contraction eventually reaches a fundamental domain, and the finitely many lattice points in that domain generate the group.

Remarks on the Proof

Non-constructivity. The proof is non-constructive in two places. First, the weak Mordell–Weil theorem produces the coset representatives $Q_1, \dots, Q_m$ via the Selmer group, which is computable in principle (??) but requires checking local conditions at every bad prime and contending with the Sha obstruction ($\text{Sel}^{(n)}$ may be strictly larger than $E(K)/nE(K)$). Second, Northcott’s theorem guarantees that the Northcott set S_{B+1} is finite but does not exhibit its elements: finding all points of bounded height requires an exhaustive search.

Effectivity. Despite the non-constructivity, the proof is *effective* in the following sense: given E/K and a Weierstrass equation, the constants c_1, c_2, C_0, B are all computable, and the generating set $\{Q_1, \dots, Q_m\} \cup S_{B+1}$ is a finite, explicitly bounded set. The practical computation of generators (without enumerating all of S_{B+1}) is the subject of chapter 9.

The role of n . The contraction factor $\lambda = 2/n^2$ improves with larger n : using $n = 3$ gives $\lambda = 2/9 \approx 0.22$ (versus $\lambda = 1/2$ for $n = 2$), and $n = 5$ gives $\lambda = 2/25 = 0.08$. Larger n means fewer descent steps, but the Selmer group $\text{Sel}^{(n)}$ becomes harder to compute (it involves $H^1(K, E[n])$ with the full n -torsion, not just the 2-torsion). In practice, $n = 2$ is nearly universal for computations over $\mathbb{Q}$: the 2-Selmer group involves only quadratic extensions and is well-suited to explicit computation.

Historical Remarks

Mordell proved the theorem for $E/\mathbb{Q}$ in 1922 [14]. His argument was a direct 2-descent: he showed that $E(\mathbb{Q})/2E(\mathbb{Q})$ is finite by relating it to the class group and unit group of a number field (essentially, the Selmer group in modern language), then performed the height descent. The key insight — that the quadratic growth of the “size” of $[2]P$ relative to P gives a contraction — was Mordell’s, and it remains the heart of the proof today.

Weil generalized the theorem to abelian varieties A/K over number fields in his 1928 thesis [28], introducing the abstract height theory and the general descent framework. The move from $E/\mathbb{Q}$ to A/K required replacing Mordell’s explicit 2-descent with the Kummer sequence for $A[n]$, and replacing the naive height (based on the x -coordinate) with a height associated to an ample line bundle on A . The logical structure of the proof — weak Mordell–Weil plus height descent implies finite generation — was established by Weil and has not changed since.

The cohomological interpretation (Selmer groups, III, the descent exact sequence) was developed in the 1950s and 1960s by Selmer, Cassels, and Tate. It clarified the structure of the weak Mordell–Weil step and made the obstruction to computing the rank (III) explicit.

Exercise 8.4.1

- Trace the descent argument for $E : y^2 + y = x^3 - x$ over $\mathbb{Q}$ (37a1) with $n = 2$ and starting point $Q = [5]P = (2, -3)$ where $P = (0, 0)$. Verify that $Q \equiv P \pmod{2E(\mathbb{Q})}$, find P_1 with $Q = P + 2P_1$, and check $h(P_1) < h(Q)$.
- For $E : y^2 = x^3 - x$ over $\mathbb{Q}$ with $n = 2$: $E(\mathbb{Q})/2E(\mathbb{Q}) \cong (\mathbb{Z}/2\mathbb{Z})^2$, with coset representatives $O, (0, 0), (1, 0), (-1, 0)$ (all of height 0). Compute the explicit bound B from (8.18) and verify that S_{B+1} contains only the four torsion points (confirming rank 0).
- Compare the number of descent steps required for $n = 2, 3, 5$ to reduce a point of height $h(P) = 1000$ to the Northcott region.

4. The proof shows $E(K) \subseteq \langle Q_1, \dots, Q_m, S_{B+1} \rangle$. Explain why a subgroup of a finitely generated abelian group that is generated by a finite set is finitely generated. (*This is a standard result in algebra, but state the key step: $\langle \Sigma \rangle = \{ \sum n_i g_i : g_i \in \Sigma, n_i \in \mathbb{Z} \}$ is a quotient of $\mathbb{Z}^{|\Sigma|}$, hence finitely generated.*)
5. Redo the descent of example 8.4.2 using the canonical height $\hat{h}$ instead of the naive height h , and verify that the bounds improve (no error terms).
6. The Mordell–Weil theorem generalizes to abelian varieties: if A/K is an abelian variety of dimension g over a number field, then $A(K) \cong \mathbb{Z}^r \oplus A(K)_{\text{tors}}$. Explain what changes in the proof for $g > 1$: (a) the Kummer sequence uses $A[n] \cong (\mathbb{Z}/n\mathbb{Z})^{2g}$; (b) the height is associated to an ample line bundle $\mathcal{L}$; (c) the Northcott property still holds for ample heights. Why does the argument go through unchanged in structure?

8.5 The Rank, Torsion, and the Structure of $E(K)$

The Mordell–Weil theorem tells us $E(K) \cong \mathbb{Z}^r \oplus E(K)_{\text{tors}}$: a free part of rank r and a finite torsion part. This section examines both pieces. The torsion subgroup $E(K)_{\text{tors}}$ is computable in practice — the torsion injection theorem reduces it to finite-field arithmetic, and Mazur’s theorem constrains the possibilities over $\mathbb{Q}$ to a short explicit list. The rank r is the deeper invariant: it measures the “arithmetic complexity” of the curve, and its computation is the central open problem of the subject.

Computing the Torsion Subgroup

The torsion injection theorem (theorem 6.3.2) provides a practical method for computing $E(K)_{\text{tors}}$: reduce modulo several primes of good reduction and take the intersection of the constraints. Over $\mathbb{Q}$, the method is particularly clean.

Proposition 8.5.1: Computing $E(\mathbb{Q})_{\text{tors}}$ via Reduction

Let $E/\mathbb{Q}$ be an elliptic curve, and let p_1, p_2 be two primes of good reduction with $p_1 \neq p_2$. Then:

$$|E(\mathbb{Q})_{\text{tors}}| \mid \gcd(|\tilde{E}(\mathbb{F}_{p_1})|, |\tilde{E}(\mathbb{F}_{p_2})|), \quad (8.24)$$

where the gcd is taken after removing the p_i -parts (i.e., for each prime ℓ , the ℓ -part of $|E(\mathbb{Q})_{\text{tors}}|$ divides $\gcd(v_\ell(|\tilde{E}(\mathbb{F}_{p_1})|), v_\ell(|\tilde{E}(\mathbb{F}_{p_2})|))$ for $\ell \neq p_1, p_2$).

Proof:

For each p_i : $E(\mathbb{Q})[\ell^\infty] \hookrightarrow \tilde{E}(\mathbb{F}_{p_i})[\ell^\infty]$ for $\ell \neq p_i$ (theorem 6.3.2), so $|E(\mathbb{Q})[\ell^\infty]|$ divides $|\tilde{E}(\mathbb{F}_{p_i})|$ (the ℓ -part). Taking the gcd over $i = 1, 2$ for each $\ell \neq p_1, p_2$ gives (8.24).

In practice, two primes almost always suffice to pin down $E(\mathbb{Q})_{\text{tors}}$ exactly. Once the order is determined, one finds the actual torsion points by searching for rational points of small height (Northcott guarantees finitely many) and checking their order.

Example 8.5.2: Torsion Computation on Running Curves

(a) $E : y^2 = x^3 - x$ over $\mathbb{Q}$. $\Delta = -64$: good reduction at all odd primes.

At $p = 3$: $\tilde{E} : y^2 = x^3 - x$ over $\mathbb{F}_3$. The points: $x = 0 : y^2 = 0, y = 0$ (1 affine pt). $x = 1 : y^2 = 0, y = 0$ (1 pt). $x = 2 : y^2 = 8 - 2 = 6 \equiv 0, y = 0$ (1 pt). Plus O : $|\tilde{E}(\mathbb{F}_3)| = 4$.

At $p = 5$: $a_5 = 5 + 1 - |\tilde{E}(\mathbb{F}_5)|$. Computing: $x = 0 : y^2 = 0$ (1 pt). $x = 1 : y^2 = 0$ (1 pt). $x = 2 : y^2 = 6 \equiv 1, y = \pm 1$ (2 pts). $x = 3 : y^2 = 24 \equiv 4, y = \pm 2$ (2 pts). $x = 4 : y^2 = 60 \equiv 0$ (1 pt). Plus O : $|\tilde{E}(\mathbb{F}_5)| = 8$.

$\gcd(4, 8) = 4$. So $|E(\mathbb{Q})_{\text{tors}}| \mid 4$. The 2-torsion points $(0, 0), (1, 0), (-1, 0)$ are visible, giving $E(\mathbb{Q})[2] \cong (\mathbb{Z}/2\mathbb{Z})^2$ (order 4). So $E(\mathbb{Q})_{\text{tors}} = E[2] \cong (\mathbb{Z}/2\mathbb{Z})^2$.

(b) $E : y^2 + y = x^3 - x$ over $\mathbb{Q}$ (37a1). $\Delta = -37$: good reduction at all $p \neq 37$.

At $p = 2$: $\tilde{E} : y^2 + y = x^3 + x$ over $\mathbb{F}_2$. Points: $x = 0 : y^2 + y = 0, y = 0$ or 1 (2 pts). $x = 1 : y^2 + y = 0$, same (2 pts). Plus O : $|\tilde{E}(\mathbb{F}_2)| = 5$.

At $p = 3$: $\tilde{E} : y^2 + y = x^3 - x$ over $\mathbb{F}_3$. Points: $x = 0 : y^2 + y = 0$ (2 pts). $x = 1 : y^2 + y = 0$ (2 pts). $x = 2 : y^2 + y = 6 - 2 = 4 \equiv 1, y^2 + y - 1 = 0$, disc $= 1 + 4 = 5 \equiv 2$ (non-square in $\mathbb{F}_3$: no pts). Plus O : $|\tilde{E}(\mathbb{F}_3)| = 5$.

$\gcd(5, 5) = 5$. So $|E(\mathbb{Q})_{\text{tors}}| \mid 5$: either 1 or 5. Since 5 is prime, $E(\mathbb{Q})_{\text{tors}} \cong \mathbb{Z}/5\mathbb{Z}$ or trivial. But the point $(0, 0)$ has $2(0, 0) = (1, -1) \neq O$, $3(0, 0) = (0, -1)$, $4(0, 0) = (1, 0)$, $5(0, 0) = \dots$. Let us compute more carefully. On the curve $y^2 + y = x^3 - x$: $P = (0, 0)$, $[2]P = (1, -1)$, $[3]P = (0, -1)$, $[4]P = (1, 0)$, $[5]P = ?$. We need to add $[4]P + P = (1, 0) + (0, 0)$. On this model: $\lambda = (0 - 0)/(0 - 1) = 0$. $x_3 = 0 - 0 - 1 + a_1 \cdot 0 - a_2 = \dots$. The general addition formula for $y^2 + a_1xy + a_3y = x^3 + a_2x^2 + a_4x + a_6$ (with $a_1 = 0, a_3 = 1, a_2 = 0, a_4 = -1, a_6 = 0$): $\lambda = (y_2 - y_1)/(x_2 - x_1) = (0 - 0)/(0 - 1) = 0$, $\nu = y_1 - \lambda x_1 = 0$. Then $x_3 = \lambda^2 + a_1\lambda - a_2 - x_1 - x_2 = 0 - 0 - 1 = -1$ and $y_3 = -(\lambda + a_1)x_3 - \nu - a_3 = 0 - 0 - 1 = -1$. So $[5]P = (-1, -1)$. This is not O , so P has order > 5 . But $|E(\mathbb{Q})_{\text{tors}}| \mid 5$ forces $E(\mathbb{Q})_{\text{tors}} = \{O\}$: trivial. So $P = (0, 0)$ has infinite order, and $E(\mathbb{Q}) \cong \mathbb{Z}$.

(c) $E : y^2 = x^3 + 1$ over $\mathbb{Q}$ (36a1). $\Delta = -432$: good reduction at $p \geq 5$.

At $p = 5$: $|\tilde{E}(\mathbb{F}_5)| = 6$ (compute directly or from $a_5 = 0$). At $p = 7$: $|\tilde{E}(\mathbb{F}_7)| = 12$ (from $a_7 = -4$: $7 + 1 + 4 = 12$). $\gcd(6, 12) = 6$. So $|E(\mathbb{Q})_{\text{tors}}| \mid 6$. The visible torsion: $(-1, 0)$ has order 2; $(0, 1)$ has $(0, 1) + (0, 1) = \dots$: on $y^2 = x^3 + 1$, doubling $(0, 1)$: $\lambda = 0/(2) = 0$, $x_2 = 0$, $y_2 = -1$. So $[2](0, 1) = (0, -1)$, and $[3](0, 1) = (0, -1) + (0, 1)$: $\lambda = (-1 - 1)/(0 - 0)$, undefined (vertical line): $[3](0, 1) = O$. So $(0, 1)$ has order 3. The group $\langle (-1, 0), (0, 1) \rangle$ has order $\text{lcm}(2, 3) = 6$, and $|E(\mathbb{Q})_{\text{tors}}| \mid 6$: therefore $E(\mathbb{Q})_{\text{tors}} \cong \mathbb{Z}/6\mathbb{Z}$, generated by $(2, 3)$ (which has order 6: one verifies $[2](2, 3) = (0, 1)$ of order 3, $[3](2, 3) = (-1, 0)$ of order 2).

The Lutz–Nagell Theorem

For elliptic curves over $\mathbb{Q}$ in short Weierstrass form, torsion points have a particularly rigid structure.

Theorem 8.5.3: The Lutz–Nagell Theorem

Let $E : y^2 = x^3 + Ax + B$ with $A, B \in \mathbb{Z}$ and $\Delta \neq 0$. If $P = (x_0, y_0) \in E(\mathbb{Q})$ is a torsion point of order ≥ 2 , then:

1. $x_0, y_0 \in \mathbb{Z}$ (the coordinates are integers).
2. Either $y_0 = 0$ (and P has order 2) or $y_0^2 \mid \Delta$.

Proof:

1. Suppose $P = (x_0, y_0) \in E(\mathbb{Q})_{\text{tors}}$ with $P \neq O$. We show $x_0, y_0 \in \mathbb{Z}$. By the formal group theory of section 4.5: for each prime p , the kernel of reduction $E_1(\mathbb{Q}_p) = \hat{E}(p\mathbb{Z}_p)$ has no prime-to- p torsion (proposition 6.3.1). So $P \pmod{p}$ is a nonsingular point of $\hat{E}(\mathbb{F}_p)$ for every prime p (at primes of good reduction: $E_1(\mathbb{Q}_p)$ has no ℓ -torsion for $\ell \neq p$, so $P \notin E_1(\mathbb{Q}_p)$, meaning P reduces to a nonsingular point).

Write $x_0 = a/b$ with $\gcd(a, b) = 1$. If $b > 1$: let $p \mid b$. Then $v_p(x_0) < 0$, so $v_p(y_0^2) = v_p(x_0^3 + Ax_0 + B)$. Since $v_p(x_0) < 0$ and $v_p(A), v_p(B) \geq 0$: $v_p(x_0^3) = 3v_p(x_0) < v_p(Ax_0) \leq v_p(x_0)$ and $v_p(x_0^3) < 0 \leq v_p(B)$, so $v_p(y_0^2) = 3v_p(x_0)$. This forces $v_p(y_0) = (3/2)v_p(x_0)$, which is not an integer (since $v_p(x_0)$ is a negative integer). Contradiction. So $b = 1$: $x_0 \in \mathbb{Z}$.

From $y_0^2 = x_0^3 + Ax_0 + B \in \mathbb{Z}$: $y_0^2 \in \mathbb{Z}$. If $y_0 \notin \mathbb{Z}$: write $y_0 = c/d$ with $d > 1$ and $\gcd(c, d) = 1$. Then $c^2/d^2 = x_0^3 + Ax_0 + B \in \mathbb{Z}$, so $d^2 \mid c^2$, contradicting $\gcd(c, d) = 1$. So $y_0 \in \mathbb{Z}$.

2. If $y_0 = 0$: P has order 2 (since $-P = (x_0, -y_0) = (x_0, 0) = P$, so $2P = O$). If $y_0 \neq 0$: consider $[2]P = (x_1, y_1)$. The duplication formula gives $x_1 = (3x_0^2 + A)^2/(4y_0^2) - 2x_0$. By part (1): $x_1 \in \mathbb{Z}$ (since $[2]P$ is also torsion). So $4y_0^2 \mid (3x_0^2 + A)^2$ in $\mathbb{Z}$. Continuing: the y -coordinate of $[2]P$ satisfies $y_1^2 = x_1^3 + Ax_1 + B$, and $y_1 \in \mathbb{Z}$.

The key estimate: from the duplication formula and the integrality of x_1 : $y_0^2 \mid (3x_0^2 + A)^2$, so $y_0 \mid (3x_0^2 + A)$... this requires more care. The standard approach uses the formal group filtration: if P has order m with $p \mid m$, then $[m/p]P \in E_1(\mathbb{Q}_p)$, and the formal group gives $v_p(x([m/p]P)) \leq -2$, etc. The conclusion $y_0^2 \mid \Delta$ follows from the discriminant relation and the integrality constraints. See Silverman [21, p. VIII.7.2] for the complete argument.

The Lutz–Nagell theorem reduces the search for torsion points to a finite computation: list all divisors d of Δ , then for each y_0 with $y_0^2 \mid \Delta$, solve $x_0^3 + Ax_0 + B = y_0^2$ for $x_0 \in \mathbb{Z}$. This is a finite set of cubic equations, each with at most three solutions.

Example 8.5.4: Lutz–Nagell Applied

$E : y^2 = x^3 - x$ over $\mathbb{Q}$. $\Delta = -64$. Divisors of Δ : $y_0^2 \mid 64$ gives $y_0 \in \{0, \pm 1, \pm 2, \pm 4, \pm 8\}$. For $y_0 = 0$: $x^3 - x = 0$, $x(x-1)(x+1) = 0$: $x_0 \in \{0, 1, -1\}$. Three points of order 2: $(0, 0)$, $(1, 0)$, $(-1, 0)$. For $y_0 = \pm 1$: $x^3 - x = 1$, $x^3 - x - 1 = 0$. No integer roots (check $x = 0, \pm 1, \pm 2$). For $y_0 = \pm 2$: $x^3 - x = 4$. $x = 2$: $8 - 2 = 6 \neq 4$. No integer roots. Similarly for $y_0 = \pm 4, \pm 8$: no solutions. So

$$E(\mathbb{Q})_{\text{tors}} = \{O, (0, 0), (1, 0), (-1, 0)\} \cong (\mathbb{Z}/2\mathbb{Z})^2.$$

Mazur’s Theorem

The most remarkable constraint on $E(\mathbb{Q})_{\text{tors}}$ is Mazur’s theorem, which limits the torsion to a short list.

Theorem 8.5.5: Mazur’s Theorem

Let $E/\mathbb{Q}$ be an elliptic curve. Then $E(\mathbb{Q})_{\text{tors}}$ is isomorphic to one of the following 15 groups:

$$\begin{aligned} \mathbb{Z}/n\mathbb{Z} & \text{ for } n = 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 12; \\ \mathbb{Z}/2\mathbb{Z} \times \mathbb{Z}/2n\mathbb{Z} & \text{ for } n = 1, 2, 3, 4. \end{aligned} \tag{8.25}$$

Each of these 15 groups occurs for infinitely many curves $E/\mathbb{Q}$ (up to isomorphism).

The proof of Mazur’s theorem is deep: it requires showing that the modular curve $X_1(N)$ has no rational points (other than cusps) for $N \notin \{1, \dots, 10, 12\}$. We will discuss this in ???. For now, we use the theorem as a black box: given $E/\mathbb{Q}$, the torsion injection theorem reduces $E(\mathbb{Q})_{\text{tors}}$ to a divisor of $\gcd(|\tilde{E}(\mathbb{F}_{p_1})|, |\tilde{E}(\mathbb{F}_{p_2})|)$, and Mazur’s theorem constrains the result to the list (8.25). In practice, two primes almost always determine $E(\mathbb{Q})_{\text{tors}}$ uniquely.

Over number fields $K \neq \mathbb{Q}$: the analogue of Mazur’s theorem is the Merel–Parent theorem, which shows $E(K)_{\text{tors}}$ is bounded in terms of $[K : \mathbb{Q}]$ alone (not depending on E). The explicit bounds improve with ongoing research but are not as clean as over $\mathbb{Q}$.

The Rank

The rank $r = \text{rank}E(K)$ is the most mysterious invariant. It measures the dimension of the “free part” of $E(K)$: the number of independent points of infinite order needed to generate $E(K)$ modulo torsion.

Definition 8.5.6: The Rank

The *rank* (or *Mordell–Weil rank*) of E/K is

$$r = \text{rank}E(K) = \dim_{\mathbb{Q}}(E(K) \otimes_{\mathbb{Z}} \mathbb{Q}). \tag{8.26}$$

Equivalently, r is the largest integer such that $E(K)$ contains r points $P_1, \dots, P_r$ that are linearly independent over $\mathbb{Z}$ (i.e., no nontrivial $\mathbb{Z}$ -linear combination $\sum n_i P_i$ is a torsion point).

Example 8.5.7: Ranks of Running Curves

(a) $E : y^2 = x^3 - x$ over $\mathbb{Q}$. $E(\mathbb{Q}) \cong (\mathbb{Z}/2\mathbb{Z})^2$: rank 0. No points of infinite order.

(b) $E : y^2 + y = x^3 - x$ over $\mathbb{Q}$ (37a1). $E(\mathbb{Q}) \cong \mathbb{Z}$: rank 1, generated by $P = (0, 0)$. The multiples $[n]P$ for $n = 1, 2, 3, \dots$ give an infinite sequence of distinct rational points with heights growing as $n^2 \hat{h}(P)$.

(c) $E : y^2 = x^3 + 1$ over $\mathbb{Q}$ (36a1). $E(\mathbb{Q}) \cong \mathbb{Z}/6\mathbb{Z}$: rank 0.

(d) *Congruent number curves.* $E_n : y^2 = x^3 - n^2x$. The rank depends on n : E_5 has rank 1 (the integer $n = 5$ is a congruent number, corresponding to a right triangle with sides $3/2, 20/3, 41/6$), while E_1 has rank 0. The conjecture that n is a congruent number iff $\text{rank} E_n(\mathbb{Q}) \geq 1$ connects the ancient geometric problem to the modern theory. Tunnell’s theorem (conditional on BSD) gives an explicit criterion: n is congruent iff a certain modular form condition is satisfied.

The Regulator in Examples

The regulator $\text{Reg}(E/K) = \det(\langle P_i, P_j \rangle)$ (??) was defined abstractly in section 8.3. We now compute it on the running curves.

Example 8.5.8: Regulators of Running Curves

(a) $E : y^2 = x^3 - x$ over $\mathbb{Q}$. Rank 0: $\text{Reg} = 1$ (by convention).

(b) $E : y^2 + y = x^3 - x$ over $\mathbb{Q}$ (37a1). Rank 1, generator $P = (0, 0)$. $\text{Reg} = \hat{h}(P) \approx 0.0511$. This is small: the generator has unusually low canonical height. The BSD formula predicts $L'(E, 1) = \Omega \cdot \text{Reg} \cdot c \cdot |\text{III}|/|T|^2 \approx 5.987 \cdot 0.0511 \cdot 1 \cdot 1/1 \approx 0.306$, consistent with the numerically computed value $L'(E, 1) \approx 0.3059$.

(c) *A rank-2 curve.* $E : y^2 = x^3 + x^2 - 25x + 29$ (141a1) has rank 2 with generators $P_1 = (4, 3)$ and $P_2 = (7, 15)$. The canonical heights and pairings: $\hat{h}(P_1) \approx 0.793$, $\hat{h}(P_2) \approx 1.448$, $\langle P_1, P_2 \rangle \approx 0.332$. The Gram matrix and regulator:

$$\text{Reg} = \det \begin{pmatrix} 0.793 & 0.332 \\ 0.332 & 1.448 \end{pmatrix} = 0.793 \cdot 1.448 - 0.332^2 \approx 1.038.$$

What the Mordell–Weil Theorem Does Not Tell Us

The theorem guarantees that $E(K)$ is finitely generated, but it is silent on the most important quantitative question:

What is the rank?

No algorithm is known that is guaranteed to compute $\text{rank} E(K)$ for an arbitrary elliptic curve E over a number field K . The Selmer group computation (??) gives an upper bound on the rank (since $\text{rank} \leq \dim_{\mathbb{F}_n} \text{Sel}^{(n)} - \dim_{\mathbb{F}_n} E(K)[n]$), but the Selmer rank can exceed the Mordell–Weil rank when $\text{III} \neq 0$. Computing the rank exactly requires either (a) finding enough independent points to match the Selmer bound, or (b) showing that the Selmer elements not accounted for by rational points lie in III .

Several deep conjectures address the rank:

8.5.9 The Birch and Swinnerton-Dyer Conjecture (Rank Part) $\text{rank} E(K) = \text{ord}_{s=1} L(E/K, s)$.

The analytic rank $\text{ord}_{s=1} L(E/K, s)$ (the order of vanishing of the L -function at $s = 1$) is computable to arbitrary precision (numerically). If the BSD conjecture is true, it would make the Mordell–Weil rank effectively computable. For elliptic curves over $\mathbb{Q}$: the conjecture is known for analytic rank ≤ 1 (by the work of Gross–Zagier and Kolyvagin, building on the modularity theorem).

Boundedness of the rank. It is not known whether the rank of elliptic curves over $\mathbb{Q}$ is bounded. The current record is a curve of rank ≥ 29 (found by Elkies, 2006). Heuristic arguments (due to Park, Poonen, Voight, and Wood) suggest that the rank is bounded on average and that curves of rank ≥ 22 (say) should be exceedingly rare, but no unconditional bound is known.

Effective bounds on generators. Even when the rank is known, the Mordell–Weil theorem does not give an effective bound on the height of generators. The constants in the descent (proposition 8.4.3) depend on the coset representatives of $E(K)/nE(K)$, whose height is controlled by the Selmer group computation. Conjectures of Szpiro and the ABC conjecture would give effective bounds on the minimal discriminant $\Delta_{\min}$ in terms of the conductor N , and these in turn would yield effective height bounds for generators. We discuss these conjectures and their Diophantine consequences in chapter 10.

Summary: The Arithmetic of $E(K)$

We collect the complete local-to-global picture that the book has built so far.

For an elliptic curve E/K over a number field:

Invariant	Description	Computed via
$E(K)_{\text{tors}}$	finite torsion subgroup	torsion injection + Mazur
$r = \text{rank} E(K)$	free rank	Selmer groups (upper bound)
$\hat{h}(P)$	canonical height	limit of duplication formula
$\text{Reg}(E/K)$	regulator	Gram matrix of $\hat{h}$
c_p	Tamagawa numbers	Tate’s algorithm (Ch 6)
$\text{III}(E/K)$	Tate–Shafarevich group	conjectured finite

The structure theorem $E(K) \cong \mathbb{Z}^r \oplus E(K)_{\text{tors}}$ is established. The torsion is computable. The rank is the central mystery. The next chapter develops the computational tools — descent, saturation, and sieving — that allow us to compute the rank and find generators in practice.

Exercise 8.5.1

1. Apply the Lutz–Nagell theorem to find $E(\mathbb{Q})_{\text{tors}}$ for $E : y^2 = x^3 + 1$. (*Hint:* $\Delta = -432$. Divisors: $y_0^2 | 432$. Check $y_0 = 0, \pm 1, \pm 2, \pm 3, \pm 6$.)

2. For $E : y^2 = x^3 + 3x + 5$ over $\mathbb{Q}$ ($\Delta = -27 \cdot 25 - 16 \cdot 27 = \dots$; compute Δ first): use reduction at $p = 5$ and $p = 7$ to determine $E(\mathbb{Q})_{\text{tors}}$.
3. For each of the 15 groups in Mazur’s list (8.25), find a specific elliptic curve $E/\mathbb{Q}$ realizing that torsion group. (*Hint*: the Kubert parametrization gives explicit families; or consult the LMFDB.)
4. For $E : y^2 + y = x^3 - x$ over $\mathbb{Q}$: $|\text{Sel}^{(2)}| = 2$ and $|E(\mathbb{Q})[2]| = 1$. Compute $\text{rank} E(\mathbb{Q})$ from these data. (*Hint*: the rank formula gives $2^r \leq |\text{Sel}^{(2)}|/|E(\mathbb{Q})[2]| = 2$, so $r \leq 1$. The point $(0, 0)$ has infinite order, so $r = 1$.)
5. Show that 5 is a congruent number by finding a non-torsion point on $E_5 : y^2 = x^3 - 25x$. (*Hint*: try $x = -4$: $y^2 = -64 + 100 = 36$, $y = \pm 6$. Verify $(-4, 6)$ has infinite order.)
6. Compute the regulator of $E : y^2 + y = x^3 - x$ (37a1) by computing $\hat{h}(P)$ for $P = (0, 0)$ using $h([2^n]P)/4^n$ for $n = 4, 5, 6$ and extrapolating.

Computing the Mordell–Weil Group

The Mordell–Weil theorem (theorem 8.4.1) guarantees that $E(K)$ is finitely generated, but says nothing about how to find the rank or the generators. This chapter develops the computational tools that turn the abstract theorem into an algorithm. The central method is *descent*: a systematic procedure for computing (or bounding) the 2-Selmer group, from which the rank is extracted. The full pipeline is: compute the Selmer group (Sections 9.1–9.2) to get an upper bound on the rank; search for rational points (Section 9.4) to get a lower bound; if the bounds match, we know the rank, and the points we found generate $E(K)$ modulo torsion; if they don’t match, the gap is explained by III or resolved by higher descent (Section 9.3). The final step is saturation (Section 9.5): verifying that the points we found generate all of $E(K)$, not just a subgroup of finite index.

All of this is best understood through examples. We will carry out the full computation for several curves, including the running examples $y^2 = x^3 - x$ and $y^2 + y = x^3 - x$, as well as the congruent number curves $y^2 = x^3 - n^2x$.

9.1 The 2-Descent with Full 2-Torsion

The simplest and most transparent form of descent occurs when all three 2-torsion points are rational. In this case, $E[2] \subseteq E(K)$, the Galois action on $E[2]$ is trivial, and the descent map (eq. (7.18)) takes values in $(K^*/(K^*)^2)^2$ — a group we can work with explicitly via Kummer theory.

Setup and the Descent Map

Let K be a number field (we work primarily with $K = \mathbb{Q}$ for concreteness) and let

$$E : y^2 = (x - e_1)(x - e_2)(x - e_3) \quad (9.1)$$

with $e_1, e_2, e_3 \in K$ distinct. The 2-torsion subgroup is $E[2] = \{O, T_1, T_2, T_3\}$ where $T_i = (e_i, 0)$, and $E[2] \cong (\mathbb{Z}/2\mathbb{Z})^2$ with trivial Galois action (since $e_i \in K$). By the descent exact sequence (theorem 7.3.1) and Kummer theory (example 7.2.6):

$$0 \longrightarrow E(K)/2E(K) \xrightarrow{\delta} H^1(K, E[2]) \cong (K^*/(K^*)^2)^2 \longrightarrow H^1(K, E)[2] \longrightarrow 0. \quad (9.2)$$

The connecting homomorphism δ was computed in example 7.3.8: for $P = (x_0, y_0) \in E(K)$ with $y_0 \neq 0$,

$$\delta(P) = ([x_0 - e_1], [x_0 - e_2]) \in (K^*/(K^*)^2)^2, \quad (9.3)$$

where $[a]$ denotes the class of a in $K^*/(K^*)^2$. The third coordinate $[x_0 - e_3]$ is determined by the constraint $[x_0 - e_1] \cdot [x_0 - e_2] \cdot [x_0 - e_3] = [y_0^2] = [1]$ (the identity in $K^*/(K^*)^2$).

For the 2-torsion points: a direct computation using the 4-division polynomial (or the limiting form of the descent map) gives $\delta(T_i) = ([e_i - e_j], [e_i - e_k])$ where $\{i, j, k\} = \{1, 2, 3\}$.

The Homogeneous Spaces

The key question is: which pairs $(d_1, d_2) \in (K^*/(K^*)^2)^2$ lie in the image of δ ? A pair (d_1, d_2) is in the image iff there exists a point $P = (x_0, y_0) \in E(K)$ with $x_0 - e_i = d_i t_i^2$ for some $t_i \in K^*$ (where $d_3 = d_1 d_2$ in $K^*/(K^*)^2$ and $y_0 = d_1 d_2 t_1 t_2 t_3$ up to sign). Substituting into the relation $(x_0 - e_1)(x_0 - e_2)(x_0 - e_3) = y_0^2$: the pair (d_1, d_2) lies in the image of δ iff the following system has a solution in K :

Definition 9.1.1: The Homogeneous Space C_{d_1, d_2}

For $d_1, d_2 \in K^*$, the *homogeneous space* associated to the pair (d_1, d_2) is the curve (or system of equations):

$$C_{d_1, d_2} : \quad d_1 z_1^2 - d_2 z_2^2 = e_2 - e_1, \quad d_1 z_1^2 - d_1 d_2 z_3^2 = e_3 - e_1. \quad (9.4)$$

This is a smooth curve of genus 1 (an intersection of two quadrics in $\mathbb{P}^3$), and it is a torsor for E in the sense of ??.

The pair (d_1, d_2) lies in $\text{im}(\delta)$ if and only if $C_{d_1, d_2}(K) \neq \emptyset$ (the homogeneous space has a rational point). The recovery of P from a solution (z_1, z_2, z_3) : $x_0 = e_1 + d_1 z_1^2$, $y_0 = d_1 d_2 z_1 z_2 z_3$ (up to sign).

The 2-Selmer Group

We cannot check $C_{d_1, d_2}(K) \neq \emptyset$ directly (this is the hard problem), but we *can* check local solubility: $C_{d_1, d_2}(K_v) \neq \emptyset$ at every place v . The Selmer group consists of the locally soluble pairs.

Definition 9.1.2: The 2-Selmer Group (Explicit)

The 2-Selmer group is

$$\text{Sel}^{(2)}(E/K) = \{(d_1, d_2) \in (K^*/(K^*)^2)^2 \mid C_{d_1, d_2}(K_v) \neq \emptyset \text{ for all places } v\}. \quad (9.5)$$

It is a finite subgroup of $(K^*/(K^*)^2)^2$, and the descent sequence gives

$$0 \longrightarrow E(K)/2E(K) \longrightarrow \text{Sel}^{(2)}(E/K) \longrightarrow \text{III}(E/K)[2] \longrightarrow 0.$$

To compute $\text{Sel}^{(2)}$ in practice, we need to know which d_1, d_2 to test and how to check local solubility.

Which d_1, d_2 to test. The pair (d_1, d_2) must represent a class in $(K^*/(K^*)^2)^2$ that is “unramified outside S ” (proposition 7.5.6(1)): at every prime $v \notin S$ (where S is the set of primes dividing 2Δ and the archimedean places), we need $d_i \in (K_v^*)^2$, i.e., $v(d_i)$ is even. Over $\mathbb{Q}$: this means d_1, d_2 are S -units modulo squares, where $S = \{p : p \mid 2\Delta\} \cup \{\infty\}$. The S -units modulo squares form a finite group: if $S = \{\infty, p_1, \dots, p_k\}$, then $\mathbb{Q}_S^*/(\mathbb{Q}_S^*)^2 \cong (\mathbb{Z}/2\mathbb{Z})^{k+1}$ (generated by $-1, p_1, \dots, p_k$). So we test all pairs (d_1, d_2) with d_1, d_2 among the 2^{k+1} squarefree products of $\{-1, p_1, \dots, p_k\}$.

How to check local solubility. For each place v and each pair (d_1, d_2) :

1. *Archimedean* ($v = \infty$): $C_{d_1, d_2}(\mathbb{R}) \neq \emptyset$ iff the system (9.4) has a real solution. This is a sign condition: since $d_1 z_1^2 \geq 0$ (or ≤ 0) depending on the sign of d_1 , the solubility depends on the signs of $d_1, d_2, d_1 d_2$ relative to $e_2 - e_1$ and $e_3 - e_1$.
2. *Non-archimedean* ($v = p$ with $p \nmid 2\Delta$): The curve C_{d_1, d_2} is smooth over $\mathbb{Z}_p$ (since d_1, d_2 are p -adic units and Δ is a p -adic unit), so $C_{d_1, d_2}(\mathbb{Q}_p) \neq \emptyset$ by Hensel’s lemma (lift a smooth $\mathbb{F}_p$ -point). Local solubility is automatic at these primes.
3. *Non-archimedean* ($v = p$ with $p \mid 2\Delta$): Requires an explicit check. For $p \mid \Delta$ (bad reduction): the structure of $C_{d_1, d_2} \pmod{p}$ determines solubility. For $p = 2$: the analysis involves 2-adic squares and is more delicate.

The upshot: the 2-Selmer group is computed by testing $O(2^{2(k+1)})$ pairs (d_1, d_2) (with $k = |S| - 1$) against local conditions at finitely many primes. For curves over $\mathbb{Q}$ with small conductor, $|S|$ is small and the computation is feasible by hand.

Full Worked Example: $y^2 = x^3 - x$

We carry out the complete 2-descent for our running curve $E : y^2 = x^3 - x = x(x-1)(x+1)$ over $\mathbb{Q}$.

Example 9.1.3: 2-Descent on $y^2 = x^3 - x$

$e_1 = 0, e_2 = 1, e_3 = -1, \Delta = -64$. The set $S = \{2, \infty\}$ (the only prime dividing $2\Delta = -128$ is $p = 2$).

Step 1: The S -units modulo squares. $\mathbb{Q}_S^*/(\mathbb{Q}_S^*)^2 = \langle [-1], [2] \rangle \cong (\mathbb{Z}/2\mathbb{Z})^2$. So the d_i range over $\{1, -1, 2, -2\}$ (squarefree representatives). The pairs $(d_1, d_2) \in (\mathbb{Q}_S^*/(\mathbb{Q}_S^*)^2)^2$ form a group of order $4^2 = 16$.

Step 2: The image of δ . The descent map $\delta : E(\mathbb{Q})/2E(\mathbb{Q}) \rightarrow (\mathbb{Q}^*/(\mathbb{Q}^*)^2)^2$ is injective (theorem 7.3.1), and $|E(\mathbb{Q})/2E(\mathbb{Q})| = |E(\mathbb{Q})| = 4$ (since $E(\mathbb{Q}) = E[2]$ and $2E(\mathbb{Q}) = \{O\}$). So $\text{im}(\delta)$ is a subgroup of order 4 inside $(\mathbb{Q}_S^*/(\mathbb{Q}_S^*)^2)^2$, and it is contained in $\text{Sel}^{(2)}$. Rather than computing the exact images of the T_i (which involves careful sign conventions in the descent map at 2-torsion; see Silverman [21, p. X.1.4]), we determine $\text{Sel}^{(2)}$ directly by local solubility and check whether $|\text{Sel}^{(2)}| = 4$. If so: $\text{Sel}^{(2)} = \text{im}(\delta)$ and $\text{III}[2] = 0$.

Step 3: Local solubility. We check each of the 16 pairs (d_1, d_2) for local solubility at $v = \infty$ and $v = 2$.

At $v = \infty$: The system $d_1 z_1^2 - d_2 z_2^2 = 1$ and $d_1 z_1^2 - d_1 d_2 z_3^2 = -1$ (since $e_2 - e_1 = 1$ and $e_3 - e_1 = -1$). For real solubility: we need $d_1 z_1^2 \geq 1$ (from the first equation, if $d_2 > 0$; or $d_1 z_1^2 \leq 1$ if $d_2 < 0$, etc.). Systematic check: the pair (d_1, d_2) is real-soluble iff the signs of $d_1, d_2, d_1 d_2$ allow both equations to have solutions simultaneously. After checking all 16 sign patterns: the real-soluble pairs are those with $d_1 > 0$ (from the second equation: $d_1 z_1^2 = -1 + d_1 d_2 z_3^2$ requires $d_1 d_2 > 0$ or $d_1 > 0$ for a solution with z_1 small). The precise analysis gives 8 real-soluble pairs.

At $v = 2$: This requires checking 2-adic solubility of the system. The analysis involves Hensel’s lemma and explicit lifts modulo 8 (since $\mathbb{Q}_2^*/(\mathbb{Q}_2^*)^2$ has exponent 2 and the quadratic forms modulo 8 determine solubility). After the computation: 4 additional pairs are eliminated.

Step 4: Result. After checking all local conditions: $|\text{Sel}^{(2)}| = 4$. Since $|E(\mathbb{Q})/2E(\mathbb{Q})| = 4$ (all of $E(\mathbb{Q})$ is 2-torsion) and $E(\mathbb{Q})/2E(\mathbb{Q}) \hookrightarrow \text{Sel}^{(2)}$: $\text{Sel}^{(2)} = E(\mathbb{Q})/2E(\mathbb{Q})$ and $\text{III}(E/\mathbb{Q})[2] = 0$.

Conclusion: $|E(\mathbb{Q})/2E(\mathbb{Q})| = 4 = 2^0 \cdot |E(\mathbb{Q})[2]|$. So $\text{rank} E(\mathbb{Q}) = 0$. The Mordell–Weil group is $E(\mathbb{Q}) = E[2] \cong (\mathbb{Z}/2\mathbb{Z})^2$.

Full Worked Example: A Rank-2 Curve

To see the descent detect non-torsion points, we consider a curve with full 2-torsion and positive rank.

Example 9.1.4: 2-Descent on a Rank-2 Curve

$E : y^2 = x(x-1)(x+3) = x^3 + 2x^2 - 3x$ over $\mathbb{Q}$. Here $e_1 = 0, e_2 = 1, e_3 = -3$. $\Delta = -16(e_1 - e_2)^2(e_1 - e_3)^2(e_2 - e_3)^2 = -16 \cdot 1 \cdot 9 \cdot 16 = -2304 = -2^8 \cdot 3^2$. The set $S = \{2, 3, \infty\}$.

Step 1: The S -units modulo squares. $\mathbb{Q}_S^*/(\mathbb{Q}_S^*)^2 = \langle [-1], [2], [3] \rangle \cong (\mathbb{Z}/2\mathbb{Z})^3$, with 8 elements. The group $(\mathbb{Q}_S^*/(\mathbb{Q}_S^*)^2)^2$ has 64 elements.

Step 2: A preliminary observation. Before computing the Selmer group, note that $T_2 = (1, 0) \in 2E(\mathbb{Q})$: the criterion for a 2-torsion point $T_i = (e_i, 0)$ to lie in $2E(\mathbb{Q})$ is that $(e_i - e_j)(e_i - e_k)$ is a square in $\mathbb{Q}^*$ (this follows from the structure of the 4-division polynomial). For T_2 : $(e_2 - e_1)(e_2 - e_3) = 1 \cdot 4 = 4 = 2^2$, a perfect square. So $T_2 \in 2E(\mathbb{Q})$. This means $E(\mathbb{Q})/2E(\mathbb{Q})$ has fewer than 4 classes of torsion points: only O, T_1 , and T_3 represent distinct cosets (with $T_2 \equiv O$).

The quotient $E(\mathbb{Q})/2E(\mathbb{Q})$ has $|E(\mathbb{Q})[2]|/|E(\mathbb{Q})[2] \cap 2E(\mathbb{Q})| \cdot 2^r = (4/2) \cdot 2^r = 2^{r+1}$ elements.

Step 3: Local solubility. We check all 64 pairs (d_1, d_2) for local solubility at $v = \infty, 2, 3$.

At $v = \infty$: The system $d_1 z_1^2 - d_2 z_2^2 = e_2 - e_1 = 1$ and $d_1 z_1^2 - d_1 d_2 z_3^2 = e_3 - e_1 = -3$ must have a real solution. From the first equation: $d_1 z_1^2 = 1 + d_2 z_2^2$, and from the second: $d_1 z_1^2 = -3 + d_1 d_2 z_3^2$. Equating: $1 + d_2 z_2^2 = -3 + d_1 d_2 z_3^2$, i.e., $d_1 d_2 z_3^2 - d_2 z_2^2 = 4$, i.e., $d_2(d_1 z_3^2 - z_2^2) = 4$. For $d_2 > 0$: need $d_1 z_3^2 > z_2^2$ (or $d_1 > 0$), soluble. For $d_2 < 0$: need $d_1 z_3^2 < z_2^2$, soluble if $d_1 < 0$ (take $z_3 = 0$, $z_2^2 = -4/d_2 > 0$) or $d_1 > 0$ with z_2 large. A careful analysis shows that real solubility eliminates roughly half the pairs: 32 pairs survive.

At $v = 3$: The conditions $d_1 z_1^2 - d_2 z_2^2 \equiv 1 \pmod{3}$ and $d_1 z_1^2 - d_1 d_2 z_3^2 \equiv 0 \pmod{3}$ (since $e_3 - e_1 = -3 \equiv 0$). The second equation gives $d_1(z_1^2 - d_2 z_3^2) \equiv 0$, so either $3 \mid d_1$ or $z_1^2 \equiv d_2 z_3^2 \pmod{3}$. Combined with the first equation: the 3-adic analysis eliminates further pairs.

At $v = 2$: The 2-adic analysis (checking lifts modulo 8) eliminates additional pairs.

Step 4: Result. After checking all local conditions: $|\text{Sel}^{(2)}| = 8$.

Step 5: Rank determination. From Step 2: $|E(\mathbb{Q})/2E(\mathbb{Q})| = 2^{r+1}$. The descent sequence gives $|E(\mathbb{Q})/2E(\mathbb{Q})| \leq |\text{Sel}^{(2)}| = 8$, so $2^{r+1} \leq 8$ and $r \leq 2$. If $\text{III}[2] = 0$: $|E(\mathbb{Q})/2E(\mathbb{Q})| = 8$, giving $r = 2$. But let us check: the point $P = (-1, 2)$ lies on E (verify: $(-1)^3 + 2(-1)^2 - 3(-1) = -1 + 2 + 3 = 4 = 2^2$). Since $P \notin E[2]$ (as $y_0 = 2 \neq 0$): P has infinite order, so $r \geq 1$.

To determine whether $r = 1$ or $r = 2$: we would need either a second independent point (suggesting $r = 2$) or a proof that $\text{III}[2] \neq 0$ (forcing $r = 1$). A computation (via the Cassels–Tate pairing or a numerical L -function evaluation) shows $\text{III}[2] = 0$ and $r = 2$ for this curve, with a second generator $Q = (3, 6)$ (verify: $27 + 18 - 9 = 36 = 6^2$). The independence of P and Q : $\hat{h}(P) \approx 0.23$, $\hat{h}(Q) \approx 0.69$, $\langle P, Q \rangle \approx 0.11$, and $\text{Reg} = \hat{h}(P)\hat{h}(Q) - \langle P, Q \rangle^2 \approx 0.15 > 0$.

The worked examples illustrate the general pattern: the 2-Selmer group is a finite, explicitly computable $\mathbb{F}_2$ -vector space; its dimension gives an upper bound on the rank; and searching for rational points gives a lower bound. When the bounds match (i.e., when $\text{III}[2] = 0$), the rank is determined.

The Algorithm in Summary

We collect the full 2-descent procedure into an algorithm.

9.1.5 Algorithm: 2-Descent with Full 2-Torsion **Input:** $E : y^2 = (x-e_1)(x-e_2)(x-e_3)$ over $\mathbb{Q}$ with $e_i \in \mathbb{Q}$.

Output: $\dim_{\mathbb{F}_2} \text{Sel}^{(2)}(E/\mathbb{Q})$, hence a bound $\text{rank} E(\mathbb{Q}) \leq \dim \text{Sel}^{(2)} - \dim E(\mathbb{Q})[2] = \dim \text{Sel}^{(2)} - 2$.

Step 1. Compute Δ and the set $S = \{p : p|2\Delta\} \cup \{\infty\}$.

Step 2. Form the group $\mathbb{Q}_S^*/(\mathbb{Q}_S^*)^2 \cong (\mathbb{Z}/2\mathbb{Z})^{|S|}$, generated by $\{-1\} \cup \{p : p \in S \setminus \{\infty\}\}$.

Step 3. For each pair (d_1, d_2) in $(\mathbb{Q}_S^*/(\mathbb{Q}_S^*)^2)^2$ (there are $2^{2|S|}$ pairs): test local solubility of C_{d_1, d_2} at each place $v \in S$.

1. At $v = \infty$: check real solubility (sign conditions).
2. At $v = p$ for $p|2\Delta$: check p -adic solubility (Hensel's lemma, or explicit lift modulo p^k for small k).

(At $v \notin S$: local solubility is automatic, by Hensel's lemma applied to the smooth reduction.)

Step 4. $\text{Sel}^{(2)} = \{(d_1, d_2) : C_{d_1, d_2} \text{ is locally soluble at all } v\}$. Compute $\dim_{\mathbb{F}_2} \text{Sel}^{(2)}$.

Step 5. Search for rational points on E (by trying small x -values, or by the methods of section 9.4). Each independent point found increases the lower bound on the rank by 1.

Step 6. If $\text{rank}_{\text{lower}} = \dim \text{Sel}^{(2)} - 2$: the rank is determined and $\text{III}[2] = 0$.

Exercise 9.1.1

1. For $E : y^2 = x(x-1)(x+1)$ over $\mathbb{Q}$: verify that $S = \{2, \infty\}$ and $\mathbb{Q}_S^*/(\mathbb{Q}_S^*)^2 \cong (\mathbb{Z}/2\mathbb{Z})^2$. List all 16 pairs (d_1, d_2) and determine which are locally soluble at $v = \infty$.
2. For $E : y^2 = x(x-1)(x+1)$ and the pair $(d_1, d_2) = (1, 2)$: write out the homogeneous space $C_{1,2}$ explicitly. Show that $C_{1,2}(\mathbb{R}) \neq \emptyset$ but $C_{1,2}(\mathbb{Q}_2) = \emptyset$, so $(1, 2) \notin \text{Sel}^{(2)}$.
3. Derive the formula $\delta(T_i) = [(e_i - e_j)(e_i - e_k)], [e_i - e_j]$ for 2-torsion points from the connecting homomorphism (theorem 7.3.1), by choosing a 4-division point Q with $[2]Q = T_i$ and computing $\sigma(Q) - Q$ for $\sigma \in G_K$.
4. Apply the 2-descent to $E : y^2 = x(x-2)(x+4)$ over $\mathbb{Q}$. Determine S , compute $|\text{Sel}^{(2)}|$, and bound the rank. (*Hint:* $\Delta = -16 \cdot 4 \cdot 16 \cdot 36 = -2^{10} \cdot 3^2 \cdot \dots$; $S = \{2, 3, \infty\}$.)
5. Give an example of an elliptic curve $E/\mathbb{Q}$ with full 2-torsion such that $|\text{Sel}^{(2)}| > |E(\mathbb{Q})/2E(\mathbb{Q})|$, i.e., $\text{III}(E/\mathbb{Q})[2] \neq 0$. (*Hint:* the curve $y^2 = x(x-6)(x+6)$ has $\text{Sel}^{(2)}$ of order 16 but $E(\mathbb{Q})/2E(\mathbb{Q})$ of order 4: $|\text{III}[2]| = 4$.)
6. Explain why local solubility at primes $p \nmid 2\Delta$ is automatic. (*Hint:* the curve C_{d_1, d_2} has good reduction at such p , and a smooth curve over $\mathbb{F}_p$ of genus 1 always has an $\mathbb{F}_p$ -point when p is sufficiently large — in fact, by the Weil bound, $|C(\mathbb{F}_p)| \geq p + 1 - 2\sqrt{p} > 0$ for $p \geq 5$. For $p = 3$: check directly.)

9.2 Descent via 2-Isogeny

The full 2-descent of section 9.1 requires all three 2-torsion points to be rational. Many important curves — including the congruent number curves $y^2 = x^3 - n^2x$ — have only *one* rational 2-torsion point. For these curves, the full 2-descent does not apply directly (the Galois action on $E[2]$ is nontrivial, so $H^1(K, E[2]) \not\cong (K^*/(K^*)^2)^2$). Instead, we use the 2-torsion point to construct a 2-isogeny $\varphi : E \rightarrow E'$ and perform descent separately via φ and its dual $\hat{\varphi}$. This “isogeny descent” is more elementary than the full 2-descent: each step involves only a single quadratic condition (rather than a system of two), and the local solubility checks reduce to testing whether certain integers are represented by binary quadratic forms.

Setup: The 2-Isogeny and Its Dual

Let E/K be an elliptic curve with a rational 2-torsion point $T = (e, 0)$. After translation, we may assume $T = (0, 0)$ and write

$$E : y^2 = x^3 + ax^2 + bx = x(x^2 + ax + b), \quad (9.6)$$

where $a, b \in K$ and the discriminant $\Delta_E = 16b^2(a^2 - 4b) \neq 0$. The point $T = (0, 0)$ generates a subgroup $\langle T \rangle = \{O, T\} \cong \mathbb{Z}/2\mathbb{Z}$.

By Vélú’s formula (example 2.1.12), the quotient $E' = E/\langle T \rangle$ is an elliptic curve with a 2-isogeny $\varphi : E \rightarrow E'$ having $\ker(\varphi) = \{O, T\}$. The explicit formulas are:

$$E' : Y^2 = X^3 - 2aX^2 + (a^2 - 4b)X, \quad \varphi(x, y) = \left(\frac{y^2}{x^2}, \frac{y(b - x^2)}{x^2} \right). \quad (9.7)$$

The curve E' also has a rational 2-torsion point $T' = (0, 0)$, and the dual isogeny $\hat{\varphi} : E' \rightarrow E$ (section 2.2) satisfies $\hat{\varphi} \circ \varphi = [2]$ on E and $\varphi \circ \hat{\varphi} = [2]$ on E' .

The Two Descent Exact Sequences

The multiplication-by-2 map $[2] = \hat{\varphi} \circ \varphi$ factors through φ , giving two short exact sequences of G_K -modules:

$$0 \rightarrow \langle T \rangle \rightarrow E \xrightarrow{\varphi} E' \rightarrow 0, \quad 0 \rightarrow \langle T' \rangle \rightarrow E' \xrightarrow{\hat{\varphi}} E \rightarrow 0. \quad (9.8)$$

Taking Galois cohomology produces two descent sequences:

$$0 \rightarrow E'(K)/\varphi(E(K)) \xrightarrow{\delta_\varphi} H^1(K, \langle T \rangle) \rightarrow H^1(K, E)[\varphi] \rightarrow 0, \quad (9.9)$$

$$0 \rightarrow E(K)/\hat{\varphi}(E'(K)) \xrightarrow{\delta_{\hat{\varphi}}} H^1(K, \langle T' \rangle) \rightarrow H^1(K, E')[\hat{\varphi}] \rightarrow 0. \quad (9.10)$$

Since $\langle T \rangle \cong \mathbb{Z}/2\mathbb{Z}$ with trivial Galois action (the point T is K -rational): $H^1(K, \langle T \rangle) \cong K^*/(K^*)^2$ by Kummer theory (theorem 7.2.4). Similarly for $\langle T' \rangle$. So both descent maps land in $K^*/(K^*)^2$.

The Explicit Descent Maps

The connecting homomorphism δ_φ is computed by the standard recipe: for $Q = (X_0, Y_0) \in E'(K)$ with $Q \neq O$ and $X_0 \neq 0$:

$$\delta_\varphi(Q) = [X_0] \in K^*/(K^*)^2. \quad (9.11)$$

For $Q = T' = (0, 0) \in E'(K)$: $\delta_\varphi(T') = [a^2 - 4b] \in K^*/(K^*)^2$ (the X -coordinate of T' is zero, and the formula degenerates; the correct value is the discriminant factor). For $Q = O$: $\delta_\varphi(O) = [1]$.

The geometric meaning: $[X_0] = [d]$ lies in the image of δ_φ iff there exists a point $(x_0, y_0) \in E(K)$ with $\varphi(x_0, y_0) = Q$, i.e., $y_0^2/x_0^2 = X_0$, which means x_0 divides y_0^2 and $X_0 = y_0^2/x_0^2$. The class $[X_0]$ in $K^*/(K^*)^2$ depends only on whether X_0 is a square times a “fundamental discriminant” associated to the isogeny.

Similarly, for the dual: for $P = (x_0, y_0) \in E(K)$ with $P \neq O, T$:

$$\delta_{\hat{\varphi}}(P) = [x_0] \in K^*/(K^*)^2. \quad (9.12)$$

For $P = T = (0, 0)$: $\delta_{\hat{\varphi}}(T) = [b]$.

The Selmer Groups

The φ -Selmer group and $\hat{\varphi}$ -Selmer group are defined by local solubility, exactly as in section 9.1:

$$\text{Sel}^{(\varphi)}(E/K) = \left\{ d \in K^*/(K^*)^2 \mid C_d^{(\varphi)}(K_v) \neq \emptyset \text{ for all } v \right\}, \quad (9.13)$$

where $C_d^{(\varphi)}$ is the homogeneous space (a genus-1 curve) parametrizing preimages of a point with δ_φ -class d . Explicitly, the local condition at a place v asks whether d is the X -coordinate (modulo squares) of a K_v -rational point on E' . This translates to: $dz^2 \in \{X : (X, Y) \in E'(K_v)\}$ for some $z \in K_v^*$, which is equivalent to the solubility of

$$C_d^{(\varphi)} : dw^2 = d^2 - 2ad + a^2 - 4b \quad (9.14)$$

(after simplification; this is a quadratic in w parametrized by d). The Selmer group $\text{Sel}^{(\hat{\varphi})}$ is defined analogously for the dual isogeny.

In practice, the local conditions simplify greatly. Over $\mathbb{Q}$ with $E : y^2 = x(x^2 + ax + b)$:

Proposition 9.2.1: Local Conditions for $\text{Sel}^{(\hat{\varphi})}$

Let $d \in \mathbb{Q}^*$ be squarefree. Then $d \in \text{Sel}^{(\hat{\varphi})}(E/\mathbb{Q})$ iff $d \mid b$ and the equation

$$dz^2 = d^2 + ad + b/d \cdot d = d^2 + ad + b \quad (9.15)$$

is soluble over $\mathbb{Q}_v$ for all places v . Equivalently (since $d \mid b$): write $b = d \cdot e$. Then $d \in \text{Sel}^{(\hat{\varphi})}$ iff the quadratic form $N^2 - aM^2 - eL^2 + (ae - d)M^2L^2/N^2 \dots$ — more precisely, the homogeneous space $C_d^{(\hat{\varphi})} : dw^2 = d^2 + ad + b$ is locally soluble everywhere.

The key simplification: d must divide b (or $2b$, depending on the model), so only finitely many squarefree d need to be tested. For each such d : the local solubility at the archimedean place is a sign condition ($d > 0$ if $d^2 + ad + b > 0$, or appropriate signs), and the local solubility at a prime p is checked by Hensel’s lemma.

The Rank Formula

The two Selmer groups bracket the rank from above. The key relation comes from combining the two descent sequences with the factorization $[2] = \hat{\varphi} \circ \varphi$.

Proposition 9.2.2: Rank Bound via Isogeny Descent

Let E/K be an elliptic curve with a 2-isogeny $\varphi : E \rightarrow E'$ and dual $\hat{\varphi} : E' \rightarrow E$. Then:

$$2^{\text{rank} E(K)} = \frac{|E(K)/2E(K)|}{|E(K)[2]|} = \frac{|\text{im}(\delta_{\hat{\varphi}})| \cdot |\text{im}(\delta_{\varphi})|}{4}, \quad (9.16)$$

where $|\text{im}(\delta_{\hat{\varphi}})| = |E(K)/\hat{\varphi}(E'(K))|$ and $|\text{im}(\delta_{\varphi})| = |E'(K)/\varphi(E(K))|$. In particular:

$$\text{rank} E(K) \leq \dim_{\mathbb{F}_2} \text{Sel}^{(\hat{\varphi})} + \dim_{\mathbb{F}_2} \text{Sel}^{(\varphi)} - 2. \quad (9.17)$$

Proof:

The factorization $[2] = \hat{\varphi} \circ \varphi$ gives $|E(K)/2E(K)| = |E(K)/\hat{\varphi}(E'(K))| \cdot |E'(K)/\varphi(E(K))| / |\text{correction}|$.

More precisely, the exact sequence $E(K) \xrightarrow{\varphi} E'(K) \xrightarrow{\hat{\varphi}} E(K) \rightarrow E(K)/2E(K) \rightarrow 0$ gives (by counting):

$$|E(K)/2E(K)| = \frac{|E(K)/\hat{\varphi}(E'(K))| \cdot |E'(K)/\varphi(E(K))|}{|\hat{\varphi}(E'(K))/2E(K)| \cdot |E'(K)/\varphi(E(K))|} \cdots$$

The clean formula: from the snake lemma applied to $0 \rightarrow \ker \varphi \rightarrow E \xrightarrow{\varphi} E' \rightarrow 0$ and $[2]: |E(K)/2E(K)| \cdot |E(K)[2]| = |E(K)/\hat{\varphi}(E'(K))| \cdot |E'(K)/\varphi(E(K))|$ (this is the “ φ - $\hat{\varphi}$ descent formula,” proved by an alternating product of orders in the exact hexagon; see Silverman [21, p. X.4.9]). Since $|E(K)[2]| = 2$ (the single rational 2-torsion point T) or 4 (if all are rational): the formula (9.16) follows. The bound (9.17) uses $|\text{im}(\delta_{\hat{\varphi}})| \leq |\text{Sel}^{(\hat{\varphi})}|$ and $|\text{im}(\delta_{\varphi})| \leq |\text{Sel}^{(\varphi)}|$.

The Congruent Number Curves

The descent via 2-isogeny is perfectly suited to the congruent number curves $E_n : y^2 = x^3 - n^2x = x(x-n)(x+n)$. Here $a = 0$, $b = -n^2$, and $T = (0, 0)$.

The 2-isogeny $\varphi : E_n \rightarrow E'_n$ has image curve $E'_n : Y^2 = X^3 + 4n^2X$ (from (9.7) with $a = 0$, $b = -n^2$). The dual isogeny $\hat{\varphi} : E'_n \rightarrow E_n$ satisfies $\hat{\varphi} \circ \varphi = [2]$.

Example 9.2.3: Descent on $E_1 : y^2 = x^3 - x$ (Not a Congruent Number)

$n = 1$. $E_1 : y^2 = x(x-1)(x+1)$. $E'_1 : Y^2 = X^3 + 4X$.

The $\hat{\varphi}$ -descent. $\delta_{\hat{\varphi}}(P) = [x_0]$ for $P = (x_0, y_0) \in E_1(\mathbb{Q})$. We need $d \in \text{Sel}^{(\hat{\varphi})}$: squarefree d dividing $b = -1$, so $d \in \{1, -1\}$. Local solubility: $d = 1$: always soluble ($P = O$ maps to $[1]$). $d = -1$: this is $\delta_{\hat{\varphi}}(T) = [b] = [-1]$; corresponds to $T = (0, 0)$, and the homogeneous space is locally soluble (it has a rational point). So $|\text{Sel}^{(\hat{\varphi})}| = 2$.

The φ -descent. $\delta_{\varphi}(Q) = [X_0]$ for $Q = (X_0, Y_0) \in E'_1(\mathbb{Q})$. We need squarefree d dividing $4b' = 4(a^2 - 4b) = 4 \cdot 4 = 16$, i.e., $d \mid 2$: $d \in \{1, -1, 2, -2\}$. Local solubility at $v = \infty$: $E'_1(\mathbb{R})$ contains points with $X_0 > 0$ (e.g., $(2, \sqrt{16}) = (2, 4)$), so $d > 0$ is possible. Also $X_0 < 0$ gives no real points: $X^3 + 4X < 0$ for $X < 0$ but $Y^2 \geq 0$... check: $X = -1$: $-1 - 4 = -5 < 0$. So $E'_1(\mathbb{R})$ has $X_0 \geq 0$ only (plus the component with $X_0 \leq -2$: $X = -3$: $-27 - 12 = -39 < 0$; $X = -4$:

$-64 - 16 = -80 < 0$). In fact $X^3 + 4X = X(X^2 + 4) \geq 0$ iff $X \geq 0$. So $Y^2 = X(X^2 + 4) \geq 0$ requires $X \geq 0$. The archimedean condition: $d > 0$. This eliminates $d = -1, -2$.

At $v = 2$: checking 2-adic solubility for $d = 2$: we need a point on E'_1 with $X_0 \equiv 2 \pmod{(\mathbb{Q}_2^*)^2}$. The point $(2, 4)$: $4^2 = 16 = 2(4 + 4) = 2 \cdot 8 = 16$. So $d = 2$ is locally soluble. We get $|\text{Sel}^{(\varphi)}| = 2$ (only $d = 1$ and $d = 2$, but we need to check whether $d = 2$ is globally soluble; in fact $(2, 4) \in E'_1(\mathbb{Q})$, so $d = 2 \in \text{im}(\delta_\varphi)$).

Rank. $|E_1(\mathbb{Q})[2]| = 4$ (all 2-torsion is rational for E_1). $|\text{im}(\delta_\varphi)| \leq |\text{Sel}^{(\varphi)}| = 2$. $|\text{im}(\delta_\varphi)| \leq |\text{Sel}^{(\varphi)}| = 2$. From (9.16): $2^r = |\text{im}(\delta_\varphi)| \cdot |\text{im}(\delta_\varphi)|/4 \leq 2 \cdot 2/4 = 1$. So $r = 0$: $E_1(\mathbb{Q}) = E_1[2] \cong (\mathbb{Z}/2\mathbb{Z})^2$. The integer 1 is not a congruent number.

Example 9.2.4: Descent on $E_5 : y^2 = x^3 - 25x$ (A Congruent Number)

$n = 5$. $E_5 : y^2 = x(x - 5)(x + 5)$. $a = 0$, $b = -25$. $E'_5 : Y^2 = X^3 + 100X$.

The $\hat{\varphi}$ -descent. The elements of $\text{Sel}^{(\hat{\varphi})}$ are squarefree divisors d of $2b = -50$ (or more precisely, the S -units modulo squares for $S = \{2, 5, \infty\}$) that satisfy local conditions at all places. The squarefree S -units: $d \in \{\pm 1, \pm 2, \pm 5, \pm 10\}$. For each, we check local solubility of the associated homogeneous space.

At $v = \infty$: the x -coordinates of real points on $E_5 : y^2 = x(x^2 - 25)$ satisfy $x \leq -5$ or $0 \leq x \leq 5$ (from the sign analysis of $x(x^2 - 25)$). Since both positive and negative x -values occur, both signs of $d = [x]$ are possible over $\mathbb{R}$. At $v = 2$ and $v = 5$: explicit Hensel lifting checks eliminate some candidates. The result: $|\text{Sel}^{(\hat{\varphi})}| = 4$.

The φ -descent. By a similar computation: $|\text{Sel}^{(\varphi)}| = 4$.

Rank. $|E_5(\mathbb{Q})[2]| = 4$ (all 2-torsion is rational). $2^r \leq |\text{Sel}^{(\hat{\varphi})}| \cdot |\text{Sel}^{(\varphi)}|/4 = 4 \cdot 4/4 = 4$, so $r \leq 2$. If $\text{III}[\varphi] = \text{III}[\hat{\varphi}] = 0$: $r = 2$. But we should check: is there a non-torsion rational point?

The point $P = (-4, 6)$: $(-4)^3 - 25(-4) = -64 + 100 = 36 = 6^2$. So $r \geq 1$. Is there a second independent point? Try $Q = (25/4, 75/8)$: $(25/4)^3 - 25(25/4) = 15625/64 - 625/4 = 15625/64 - 10000/64 = 5625/64$, and $(75/8)^2 = 5625/64$. Independence: $\hat{h}(P)$ and $\hat{h}(Q)$ give a non-degenerate Gram matrix (the regulator is positive). So $r = 2$, $\text{III}[\varphi] = \text{III}[\hat{\varphi}] = 0$, and $E_5(\mathbb{Q}) \cong \mathbb{Z}^2 \oplus (\mathbb{Z}/2\mathbb{Z})^2$.

The congruent number. Since $\text{rank} E_5(\mathbb{Q}) \geq 1$: 5 is a congruent number. The correspondence: a non-torsion point $P = (-4, 6)$ on E_5 gives a right triangle with rational sides and area 5. The triangle: from $P = (x, y) \in E_n(\mathbb{Q})$ with $y \neq 0$, the sides are $a = |x^2 - n^2|/y$, $b = 2nx/y$, $c = (x^2 + n^2)/y$. For $P = (-4, 6)$: $a = |16 - 25|/6 = 3/2$, $b = 2 \cdot 5 \cdot (-4)/6 = -40/6 = -20/3$ (take absolute value: $20/3$), $c = (16 + 25)/6 = 41/6$. Check: $(3/2)^2 + (20/3)^2 = 9/4 + 400/9 = 81/36 + 1600/36 = 1681/36 = (41/6)^2$. Area: $(1/2)(3/2)(20/3) = 5$.

Comparison with Full 2-Descent

The isogeny descent has both advantages and disadvantages relative to the full 2-descent of section 9.1. The advantage: each Selmer group involves only $K^*/(K^*)^2$ (a single quadratic condition), rather than $(K^*/(K^*)^2)^2$ (a pair of quadratic conditions). This makes the local solubility checks simpler and the number of candidates smaller. The disadvantage: the rank bound (9.17) involves *two* Selmer groups (for φ and $\hat{\varphi}$), and the individual Selmer ranks can be larger than necessary. In the worst case, both $\text{III}[\varphi]$ and $\text{III}[\hat{\varphi}]$ are nontrivial, giving a weaker bound than the full 2-Selmer group.

When $E[2] \subseteq E(K)$ (all 2-torsion rational): both methods are available, and the full 2-descent is generally more efficient (it gives a single Selmer group with a sharper rank bound). The isogeny descent is most useful when E has exactly one rational 2-torsion point — the situation for the congruent number curves, and more generally for any curve of the form $y^2 = x^3 + ax^2 + bx$ with $a^2 - 4b$ not a perfect square.

Exercise 9.2.1

1. Apply the φ - $\hat{\varphi}$ descent to $E_7 : y^2 = x^3 - 49x$ to determine whether 7 is a congruent number. (*Hint:* Compute $|\text{Sel}^{(\hat{\varphi})}|$ and $|\text{Sel}^{(\varphi)}|$. Search for a non-torsion point.)
2. For $E : y^2 = x^3 + ax^2 + bx$ with $\varphi : E \rightarrow E'$ as above: verify that $E' : Y^2 = X^3 - 2aX^2 + (a^2 - 4b)X$ has discriminant $\Delta_{E'} = 16(a^2 - 4b)^2 \cdot (-4)(a^2 - 4b + 4 \cdot 2a \cdot \dots) = \dots$. Compute $\Delta_{E'}$ in terms of a, b and verify $\Delta_{E'} \neq 0$ iff $\Delta_E \neq 0$.
3. Using the correspondence between non-torsion points on E_n and right triangles with area n : find the rational right triangle with area 6. (*Hint:* find a non-torsion point on $E_6 : y^2 = x^3 - 36x$.)
4. For $E'_5 : Y^2 = X^3 + 100X$: list all squarefree divisors of 200 and determine which lie in $\text{Sel}^{(\varphi)}$ by checking local solubility at $v = \infty, 2, 5$.
5. Prove the rank formula (9.16) by considering the exact sequence $0 \rightarrow E[\hat{\varphi}]/E[\varphi] \rightarrow E/\varphi(E') \xrightarrow{\hat{\varphi}} E/2E \rightarrow E/\hat{\varphi}(E') \rightarrow 0$ (where $E[\hat{\varphi}]/E[\varphi]$ accounts for the kernel overlap) and counting orders.
6. Show that $n = 1, 2, 3$ are not congruent numbers by proving $\text{rank} E_n(\mathbb{Q}) = 0$ via isogeny descent. (*Hint:* for $n = 1$: done in example 9.2.3. For $n = 2, 3$: the Selmer groups have $|\text{Sel}^{(\hat{\varphi})}| \cdot |\text{Sel}^{(\varphi)}| = 4$, giving $r = 0$.)

9.3 Higher Descent and the Limits of 2-Descent

The 2-descent (in both forms: sections 9.1 and 9.2) computes the 2-Selmer group $\text{Sel}^{(2)}$ and deduces a bound $\text{rank} E(K) \leq \dim_{\mathbb{F}_2} \text{Sel}^{(2)} - \dim_{\mathbb{F}_2} E(K)[2]$. When the Selmer rank equals the Mordell–Weil rank plus the torsion rank (i.e., when $\text{III}(E/K)[2] = 0$), the 2-descent determines the rank exactly. But this need not happen: $\text{III}[2]$ can be nontrivial, inflating the Selmer group beyond what is accounted for by rational points. This section discusses what to do when the 2-descent leaves a gap, the structural constraints that III must satisfy, and the broader landscape of descent methods.

When 2-Descent Fails: Nontrivial $\text{III}[2]$

The Tate–Shafarevich group $\text{III}(E/K)$ (??) measures the failure of the local-to-global principle for torsors of E : an element of III is a genus-1 curve C/K that is a torsor for E (??) with $C(K_v) \neq \emptyset$ for all places v but $C(K) = \emptyset$. Such a C represents an element of $\text{Sel}^{(n)}$ that does not come from a rational point on E .

The simplest examples of nontrivial $\text{III}[2]$ arise from curves of moderate conductor. The phenomenon is not rare: a positive proportion of elliptic curves over $\mathbb{Q}$ (ordered by height) have nontrivial $\text{III}[2]$.

Example 9.3.1: A Curve with $\text{III}[2] \neq 0$

Consider $E : y^2 = x^3 + 6x^2 + x$ over $\mathbb{Q}$ (Cremona label 571a1). This curve has a single rational 2-torsion point at $T = (0, 0)$.

The φ - $\hat{\varphi}$ descent gives $|\text{Sel}^{(\hat{\varphi})}| = 4$ and $|\text{Sel}^{(\varphi)}| = 4$. The rank formula (proposition 9.2.2) gives $2^r \cdot |E(\mathbb{Q})[2]| \leq |\text{Sel}^{(\hat{\varphi})}| \cdot |\text{Sel}^{(\varphi)}| = 16$. With $|E(\mathbb{Q})[2]| = 2$: $2^{r+1} \leq 16$, so $r \leq 3$.

A point search finds no non-torsion rational points of small height. The L -function computation gives $L(E, 1) \neq 0$ (and in fact $L(E, 1) \approx 1.19$), so the BSD conjecture predicts $r = 0$. Since $r = 0$ and $|E(\mathbb{Q})/2E(\mathbb{Q})| = |E(\mathbb{Q})[2]| = 2$: the descent sequence gives $|\text{III}[2]| = |\text{Sel}^{(\hat{\varphi})}| \cdot |\text{Sel}^{(\varphi)}| / (|E(\mathbb{Q})/2E(\mathbb{Q})| \cdot |E(\mathbb{Q})[2]|) = 16/4 = 4$. So $\text{III}(E/\mathbb{Q})[2] \cong (\mathbb{Z}/2\mathbb{Z})^2$: two independent locally soluble torsors that have no global rational point.

The gap between the Selmer rank and the Mordell–Weil rank is measured by $\text{III}[2]$: the larger $\text{III}[2]$ is, the weaker the 2-descent bound. For curves with large $\text{III}[2]$, the 2-descent gives almost no useful information about the rank.

Second Descent (Descent on a 2-Covering)

When the 2-descent produces a Selmer element $(d_1, d_2) \in \text{Sel}^{(2)}$ that might or might not come from a rational point, we can attempt to resolve the ambiguity by performing a *second descent*: a descent on the homogeneous space C_{d_1, d_2} itself.

The idea is as follows. The homogeneous space $C = C_{d_1, d_2}$ from ?? is a genus-1 curve, everywhere locally soluble (since $(d_1, d_2) \in \text{Sel}^{(2)}$). If $C(K) \neq \emptyset$: $(d_1, d_2) \in \text{im}(\delta)$ and corresponds to a rational point on E . If $C(K) = \emptyset$: the pair (d_1, d_2) represents a nontrivial element of $\text{III}[2]$. The second descent attempts to distinguish these two cases by embedding C into a higher-dimensional space and checking whether the resulting variety has rational points.

Concretely: a 2-descent on C (viewed as a torsor for E) produces a 2-covering of C , which is a curve D of genus 1 mapping to C with degree 2. The curve D is a torsor for E of period 4 (it maps to E via a degree-4 isogeny), and it is a curve in $\mathbb{P}^3$ (an intersection of two quadrics). The 2-covering D has a rational point iff C does; but D can be analyzed by a 4-descent, which involves checking local solubility of quadric intersections in $\mathbb{P}^3$. This is computationally harder than the original 2-descent (which involved quadratic forms) but is still feasible.

The second descent was developed systematically by Cassels and Birch–Swinnerton-Dyer in the 1960s, and has been implemented in modern algebra systems (Magma, SageMath). It is the primary tool

for distinguishing genuine Sha elements from “lazy” Selmer elements that happen to correspond to rational points of large height.

Example 9.3.2: Second Descent Resolving III

For the curve $E : y^2 = x^3 + 6x^2 + x$ of example 9.3.1: the two nontrivial elements of $\text{Sel}^{(2)} \setminus \text{im}(\delta)$ correspond to genus-1 curves C_1, C_2 that are everywhere locally soluble but have no rational points. A second descent on C_1 produces a 2-covering $D_1 \subseteq \mathbb{P}^3$, and the 4-Selmer analysis shows D_1 is *not* everywhere locally soluble (it fails at a prime $p|\Delta$). This confirms that $C_1(K) = \emptyset$ and the corresponding Selmer element lies in $\text{III}[2]$.

In general, the second descent reduces $|\text{III}[2]|$ to $|\text{III}[4]/2\text{III}[2]|$: elements of $\text{III}[2]$ that survive the second descent are “deeper” obstructions lying in $\text{III}[4]$.

n -Descent for $n > 2$

The 2-descent is the most common because it involves only quadratic conditions, which are amenable to explicit computation. For $n = 3, 4$, and higher, analogous descent procedures exist, but the homogeneous spaces become more complex.

3-descent. The 3-Selmer group consists of locally soluble torsors of period dividing 3. These torsors are plane cubic curves in $\mathbb{P}^2$: smooth genus-1 curves embedded by a degree-3 line bundle. Computing the 3-Selmer group requires enumerating locally soluble plane cubics up to equivalence — a more involved computation than the quadratic conditions of the 2-descent, but still feasible. The 3-descent is useful when E has a rational 3-torsion point (giving a 3-isogeny $\varphi : E \rightarrow E'$, analogous to the 2-isogeny descent of section 9.2), or when the 2-descent fails to determine the rank and $\text{III}[3]$ is the relevant obstruction.

4-descent. The 4-Selmer group consists of locally soluble torsors of period dividing 4, which are intersections of two quadrics in $\mathbb{P}^3$ (genus-1 curves embedded by a degree-4 line bundle). The 4-descent refines the 2-descent: every element of $\text{Sel}^{(4)}$ maps to an element of $\text{Sel}^{(2)}$ via the natural projection, and the fiber over a given 2-Selmer element detects whether it lifts to a 4-torsion class. The “second descent” of the previous subsection is precisely this: starting from a 2-covering C_{d_1, d_2} (in $\text{Sel}^{(2)}$), we ask whether C lifts to a 4-covering D in $\text{Sel}^{(4)}$. If C is locally soluble but D is not: the element lies in $\text{III}[2] \setminus \text{III}[4]$, and the second descent detects it.

n -descent for general n . The n -Selmer group consists of locally soluble torsors of period dividing n , which are genus-1 curves embedded in $\mathbb{P}^{n-1}$ by a degree- n line bundle. For $n = 5$: the curves live in $\mathbb{P}^4$ (cut out by 5 quadrics). For $n \geq 6$: the computational complexity grows rapidly, and n -descent is rarely performed in practice.

The general pattern: higher n gives a sharper bound on the rank (since $\text{III}[n]$ contains more information than $\text{III}[2]$ for n divisible by 2), but at the cost of increased computational difficulty. In practice, the combination of 2-descent and 4-descent suffices for most curves of moderate conductor.

III in Practice: How Common Is It?

The Tate–Shafarevich group is not merely a theoretical obstruction: it occurs frequently enough that any serious computational approach must contend with it.

Example 9.3.3: The Distribution of $\text{III}[2]$

Among all elliptic curves over $\mathbb{Q}$ of conductor $\leq 100,000$ (from the LMFDB database):

The vast majority have $\text{III}[2] = 0$: for these curves, the 2-descent determines the rank exactly. A smaller but significant fraction have $|\text{III}[2]| = 4$ (the smallest nontrivial case allowed by the Cassels–Tate perfect-square constraint). Curves with $|\text{III}[2]| \geq 16$ are rarer but exist: the curve $y^2 = x^3 - x^2 - 79x + 289$ (conductor 5765) has $|\text{III}[2]| = 16$.

The analytic order of III (predicted by the BSD formula) can be much larger: curves with $|\text{III}| = 9, 25, 49$ (odd perfect squares) require n -descent for odd n to detect. The curve $E : y^2 + y = x^3 - 7$ (conductor 5077, rank 0) has $|\text{III}| = 1$ (despite having the largest conductor among rank-0 curves in its neighborhood). The curve with Cremona label **681b1** has $|\text{III}_{\text{an}}| = 9$ (the analytic order from BSD).

The key message for the practitioner: one should always compute the L -function value $L(E, 1)$ (or $L'(E, 1)$ if the root number is -1) alongside the descent. The L -function gives the analytic rank, which — assuming BSD — determines the Mordell–Weil rank. The descent then determines $|\text{III}[n]|$ as the quotient $|\text{Sel}^{(n)}|/|E(\mathbb{Q})/nE(\mathbb{Q})|$. When the L -function is unavailable or the analytic rank is large (≥ 2), descent is the primary tool.

The Cassels–Tate Pairing

A fundamental structural constraint on III is provided by the Cassels–Tate pairing: a bilinear form on III whose existence forces $|\text{III}[n]|$ to be a perfect square. This is a deep theorem with important computational consequences.

Theorem 9.3.4: The Cassels–Tate Pairing

Let E/K be an elliptic curve over a number field. There exists a non-degenerate alternating bilinear pairing

$$\langle \cdot, \cdot \rangle_{\text{CT}} : \text{III}(E/K) \times \text{III}(E/K) \longrightarrow \mathbb{Q}/\mathbb{Z}. \quad (9.18)$$

In particular: if $\text{III}(E/K)$ is finite, then $|\text{III}(E/K)|$ is a perfect square.

The proof (due to Cassels for $n = 2$ and extended by Tate) uses the Poitou–Tate exact sequence (theorem 7.5.7) and the Weil pairing. The alternating property follows from the alternating nature of the Weil pairing itself. We omit the proof and focus on the consequences.

The perfect-square constraint is a powerful computational tool. If a 2-descent gives $|\text{Sel}^{(2)}| = 2^s$ and a point search gives rank $\geq r$: then $|\text{III}[2]| = 2^{s-2-r}$ (where $|E(K)[2]| = 2^2$ in the full 2-torsion case). Since $|\text{III}[2]|$ must be a perfect square: $s - 2 - r$ must be even. This constrains the rank to have a definite parity.

Example 9.3.5: The Perfect Square Constraint in Action

$E/\mathbb{Q}$ with full 2-torsion and $|\text{Sel}^{(2)}| = 8$. Then $|\text{III}[2]| = 8/(2^r \cdot 4) = 2^{1-r}$. For $|\text{III}[2]|$ to be a perfect square: $1 - r$ must be even and ≥ 0 . So $r = 1$ (giving $|\text{III}[2]| = 1$) or $r = -1$ (impossible). The rank must be 1: the Cassels–Tate pairing has determined the rank without finding a single

rational point!

More generally: if $|E(K)[2]| = 4$ and $|\text{Sel}^{(2)}| = 2^s$: then $r = s - 2 - 2k$ where $k \geq 0$ and $2k$ is the rank of $\text{III}[2]$ as an $\mathbb{F}_2$ -vector space. Since $2k$ must be even (by the perfect square constraint): $r \equiv s \pmod{2}$. The Selmer dimension and the rank have the same parity.

The Parity Conjecture

The Cassels–Tate pairing implies that $\text{rank} E(K)$ and $\dim_{\mathbb{F}_2} \text{Sel}^{(2)} - \dim_{\mathbb{F}_2} E(K)[2]$ have the same parity (assuming III is finite). A much deeper conjecture relates the parity of the rank to the *root number* $w(E/K) \in \{\pm 1\}$, a sign that appears in the functional equation of the L -function.

9.3.6 The Parity Conjecture Let E/K be an elliptic curve over a number field. Then

$$(-1)^{\text{rank} E(K)} = w(E/K), \quad (9.19)$$

where $w(E/K)$ is the root number (the sign of the functional equation $L(E/K, s) = w(E/K) \cdot N^{1-s} \cdot (\text{gamma factors}) \cdot L(E/K, 2-s)$).

The root number $w(E/K)$ is computable from local data: $w(E/K) = \prod_v w_v(E/K_v)$, where the product runs over all places of K and each local factor w_v depends on the reduction type at v . For $E/\mathbb{Q}$: $w_\infty = -1$ (always), and $w_p = +1$ for primes of good reduction, $w_p = -1$ for split multiplicative reduction, $w_p = +1$ for non-split multiplicative reduction, and w_p depends on the conductor exponent for additive reduction.

The parity conjecture is a consequence of the BSD conjecture: if $\text{ord}_{s=1} L(E/K, s) = r$, then the functional equation forces $(-1)^r = w(E/K)$. Conversely, the parity conjecture is known unconditionally in many cases:

Theorem 9.3.7: The Parity Theorem (Nekovář, Dokchitser–Dokchitser)

Let $E/\mathbb{Q}$ be an elliptic curve. Then:

$$(-1)^{\text{rank} E(\mathbb{Q})} = w(E/\mathbb{Q}), \quad (9.20)$$

assuming the finiteness of $\text{III}(E/\mathbb{Q})$.

The proof (due to Nekovář for semistable curves, extended by T. and V. Dokchitser to the general case) uses p -adic heights, Iwasawa theory, and the structure of Selmer groups over cyclotomic towers. The unconditional result (without assuming III finite) is known for all semistable $E/\mathbb{Q}$ by work of the Dokchitser brothers.

Example 9.3.8: Root Number and Rank Parity

(a) $E: y^2 = x^3 - x$ over $\mathbb{Q}$ (conductor $N = 32$). $\Delta = -64$: good reduction at odd primes, additive at $p = 2$. The root number: $w = \prod_v w_v$. $w_\infty = -1$. At $p = 2$: the local root number for additive reduction depends on the conductor exponent ($f_2 = 5$ for this curve) and is $w_2 = -1$. At all other primes: $w_p = +1$ (good reduction). So $w = (-1)(-1) = +1$. Prediction: $\text{rank} E(\mathbb{Q})$ is even. Indeed $\text{rank} = 0$.

(b) $E : y^2 + y = x^3 - x$ over $\mathbb{Q}$ (37a1, conductor $N = 37$). $\Delta_{\min} = -37^3$ (the minimal discriminant of this curve; the standard model has $c_4 = -48$, $c_6 = -216$). The reduction at $p = 37$: since $37 \mid \Delta_{\min}$ and $37 \nmid c_4$, the reduction is multiplicative with Kodaira type I_1 . The tangent directions at the singular point of $\tilde{E}/\mathbb{F}_{37}$ determine whether the reduction is split or non-split: from the LMFDB, the reduction is non-split multiplicative, giving $w_{37} = +1$. At the archimedean place: $w_{\infty} = -1$. At all other primes: $w_p = +1$ (good reduction). So $w = (-1)(+1) = -1$. Prediction: $\text{rank}E(\mathbb{Q})$ is odd. Indeed $\text{rank} = 1$.

(c) $E_5 : y^2 = x^3 - 25x$ over $\mathbb{Q}$ (conductor $N = 400 = 2^4 \cdot 5^2$). The local root numbers at primes of additive reduction (here $p = 2$) are notoriously delicate to compute by hand. From the LMFDB: $w(E_5/\mathbb{Q}) = +1$. Prediction: rank is even. Indeed $\text{rank} = 2$ (example 9.2.4).

The computations above illustrate that while the root number is in principle determined by local data, the local factors at primes of additive reduction (especially $p = 2$) require care. In practice, one uses a computer algebra system or the LMFDB to compute $w(E/\mathbb{Q})$ reliably.

Beyond Descent: Heegner Points

For elliptic curves of analytic rank 1 (i.e., $L(E, 1) = 0$ and $L'(E, 1) \neq 0$), a completely different method produces generators directly, bypassing the descent-and-search pipeline entirely: the *Heegner point* construction.

The idea: choose an imaginary quadratic field $K = \mathbb{Q}(\sqrt{-D})$ such that all primes dividing the conductor N of E split in K (the “Heegner hypothesis”). The theory of complex multiplication on the modular curve $X_0(N)$ produces a point $y_K \in E(K)$. The Gross–Zagier formula relates the canonical height of y_K to $L'(E/K, 1)$: if $L'(E/K, 1) \neq 0$, then $\hat{h}(y_K) > 0$, so y_K has infinite order. By a careful choice of K , the trace $P = \text{Tr}_{K/\mathbb{Q}}(y_K) \in E(\mathbb{Q})$ is a point of infinite order, giving a generator of $E(\mathbb{Q})/E(\mathbb{Q})_{\text{tors}}$.

The Heegner point method is the only known unconditional way to produce generators for rank-1 curves. Combined with Kolyvagin’s Euler system argument: if the analytic rank is 1, then the Mordell–Weil rank is exactly 1, III is finite, and the Heegner point generates $E(\mathbb{Q})$ modulo torsion (up to finite index). This is the content of the Gross–Zagier–Kolyvagin theorem, which we will discuss in ??.

For analytic rank ≥ 2 : no analogue of the Heegner point construction is known, and finding generators remains a computational challenge.

The State of the Art

Combining all available tools, the practical computation of $\text{rank}E(\mathbb{Q})$ for a given curve proceeds as follows.

1. *2-descent*. Compute $\text{Sel}^{(2)}$ (via full 2-descent or isogeny descent). This gives an upper bound $r \leq s - t$ where $s = \dim_{\mathbb{F}_2} \text{Sel}^{(2)}$ and $t = \dim_{\mathbb{F}_2} E(\mathbb{Q})[2]$.
2. *Point search*. Search for rational points of small height (by the methods of section 9.4). Each independent point found raises the lower bound on r .

3. *Parity check.* Compute $w(E/\mathbb{Q})$: if $w = +1$ and the lower bound is odd, or $w = -1$ and the lower bound is even, search harder for another point.
4. *Second descent.* If the Selmer bound and the point-search bound differ by 2 or more: perform a 4-descent (second descent) to refine the Selmer bound.
5. *L-function computation.* Compute $L(E, s)$ numerically at $s = 1$. If $L(E, 1) \neq 0$: the analytic rank is 0, and (by Kolyvagin) the Mordell–Weil rank is 0. If $L'(E, 1) \neq 0$ and $L(E, 1) = 0$: the analytic rank is 1, and (by Gross–Zagier + Kolyvagin) the Mordell–Weil rank is 1.
6. *Heegner points.* When the analytic rank is 1: the Heegner point construction produces an explicit generator, bypassing the search.

For curves of analytic rank ≤ 1 : the combination of the Gross–Zagier formula, Kolyvagin’s theorem, and the modularity theorem (Wiles et al.) gives a complete, unconditional determination of the rank. For analytic rank ≥ 2 : no unconditional method is known, and the rank is determined by the combination of descent bounds and point searches.

In the LMFDB database (which catalogues all elliptic curves over $\mathbb{Q}$ of conductor $\leq 500,000$): the rank is known for every entry. The vast majority are settled by 2-descent plus an L -function computation. A small fraction require 4-descent or more sophisticated methods. The largest verified rank in the database is 4; curves of rank ≥ 5 are rare and their ranks are typically determined by combining 2-descent with extensive point searches.

Exercise 9.3.1

1. Show directly (without the Cassels–Tate pairing) that $|\text{III}(E/\mathbb{Q})[2]|$ is a perfect square when $E[2] \subseteq E(\mathbb{Q})$. (*Hint:* $\text{III}[2]$ is a subgroup of $(K^*/(K^*)^2)^2/\text{im}(\delta)$, and the Weil pairing induces a non-degenerate alternating form on this quotient.)
2. Compute the root number $w(E/\mathbb{Q})$ for $E_5 : y^2 = x^3 - 25x$ (conductor $N = 400 = 2^4 \cdot 5^2$). Use the result to predict the parity of $\text{rank} E_5(\mathbb{Q})$ and verify against example 9.2.4.
3. An elliptic curve over $\mathbb{Q}$ has $|\text{Sel}^{(2)}| = 16$, $|E(\mathbb{Q})[2]| = 4$, and a point search finds 2 independent non-torsion points. What are the possible values of $\text{rank} E(\mathbb{Q})$ and $|\text{III}[2]|$? Use the Cassels–Tate perfect-square constraint to narrow the possibilities.
4. Explain why a second descent (4-descent) can detect that a 2-Selmer element lies in $\text{III}[2]$: if C is the homogeneous space for the element, and $D \rightarrow C$ is a 2-covering, then $C(K) \neq \emptyset$ iff $D(K) \neq \emptyset$. The local solubility of D is a strictly stronger condition than that of C : D may fail to be locally soluble even when C is.
5. For $E : y^2 = x^3 + 1$ over $\mathbb{Q}$ (with $j = 0$): the 3-torsion includes the point $(0, 1)$ of order 3. Explain why a 3-descent might be useful for this curve (even though the 2-descent already shows $\text{rank} = 0$): the 3-Selmer group provides independent information about $\text{III}[3]$.
6. The curve $E : y^2 = x^3 - 7x + 6$ has $L(E, 1) = 0$ and $L'(E, 1) \neq 0$. What does this imply about $\text{rank} E(\mathbb{Q})$ (assuming BSD)? What does it imply unconditionally (via Gross–Zagier and Kolyvagin)?

9.4 Searching for Points and Height Bounds

The descent machinery of Sections 9.1–9.3 provides an upper bound on the rank via the Selmer group. To determine the rank exactly (when the Selmer bound exceeds the known rank), we need to find rational points of infinite order. This is the complementary problem: descent works “from above” (bounding rank by $\dim \text{Sel}^{(n)}$), while point-searching works “from below” (producing independent points that witness rank $\geq r$). When the two bounds meet, the rank is determined.

Finding rational points on an elliptic curve is, in general, a hard problem. No polynomial-time algorithm is known, and the difficulty grows rapidly with the height of the generators. This section develops the practical methods: naive search, sieving, and the theoretical height bounds that tell us how far we need to look.

Naive Search

The most straightforward approach: enumerate rational numbers $x_0 = a/b$ (in lowest terms, with $|a|, |b| \leq N$) and check whether $y_0^2 = x_0^3 + Ax_0 + B$ has a rational solution (i.e., whether the right side is a perfect square in $\mathbb{Q}$).

For $E : y^2 = x^3 + Ax + B$ over $\mathbb{Q}$ with $A, B \in \mathbb{Z}$: writing $x_0 = a/b$ with $\gcd(a, b) = 1$:

$$y_0^2 = \frac{a^3 + Ab^2a + Bb^3}{b^3}, \quad (9.21)$$

which is rational iff $b^3 \mid (a^3 + Ab^2a + Bb^3)$. Since $\gcd(a, b) = 1$: $b^3 \mid Bb^3$ automatically, and $b^3 \mid a^3$ iff $b = \pm 1$. So for $b \neq \pm 1$: we need $b^3 \mid (Ab^2a + Bb^3) = b^2(Aa + Bb)$, i.e., $b \mid (Aa + Bb)$, i.e., $b \mid Aa$, i.e., $b \mid A$ (since $\gcd(a, b) = 1$). For integer Weierstrass models with small A : this severely restricts the possible denominators of x_0 .

Proposition 9.4.1: Denominator Structure for Rational Points

Let $E/\mathbb{Q}$ be given by a minimal Weierstrass equation. Every point $P = (x_0, y_0) \in E(\mathbb{Q}) \setminus \{O\}$ can be written

$$x_0 = \frac{a}{d^2}, \quad y_0 = \frac{c}{d^3}, \quad (9.22)$$

with $a, c, d \in \mathbb{Z}$ and $\gcd(a, d) = \gcd(c, d) = 1$.

This is a nontrivial constraint: the denominator of x_0 is a perfect square, and the denominator of y_0 is a perfect cube (with the same base d). It means that a point of height $h(P) = h(x_0) \leq B$ has $x_0 = a/d^2$ with $|a| \leq e^B d^2$ and $d \leq e^{B/2}$, so the search space is $O(e^{3B/2})$ pairs (a, d) — exponential in B , but much smaller than the naive $O(e^{2B})$ one would get by enumerating arbitrary rationals a/b .

Proof:

At each prime p : $v_p(x_0) = v_p(a) - 2v_p(d)$. If $v_p(x_0) < 0$: P lies in the kernel of reduction $E_1(\mathbb{Q}_p)$, and the formal group structure (section 4.5) gives $v_p(x_0) = -2v_p(z)$ for some uniformizer $z \in p\mathbb{Z}_p$, so $v_p(x_0)$ is even and negative. Writing $v_p(x_0) = -2m$: $v_p(a) = 0$ and $v_p(d) = m$. Similarly $v_p(y_0) = -3m$ gives $v_p(c) = 0$ and $v_p(d) = m$. So $\gcd(a, d) = \gcd(c, d) = 1$.

The denominator structure $x_0 = a/d^2$, $y_0 = c/d^3$ means: searching for points of naive height $h(P) = h(x_0) \leq B$ requires enumerating pairs (a, d) with $|a| \leq e^B d^2$ and $d \leq e^{B/2}$. The search space has size $O(e^{3B/2})$: exponential in B .

Example 9.4.2: Naive Search on $y^2 + y = x^3 - x$ (37a1)

For the minimal model $y^2 + y = x^3 - x$: converting to short Weierstrass, $Y^2 = X^3 - X + 1/4$, or equivalently $(2Y)^2 = 4X^3 - 4X + 1$. But it is simpler to search on the original model: try $x = 0, \pm 1, \pm 2, \dots$ and check whether $y^2 + y - (x^3 - x) = 0$ has a rational root $y = (-1 \pm \sqrt{1 + 4(x^3 - x)})/2$.

$x = 0$: $y^2 + y = 0$, $y(y + 1) = 0$: $y = 0$ or $y = -1$. Points: $(0, 0)$ and $(0, -1)$.

$x = 1$: $y^2 + y = 0$: $y = 0$ or $y = -1$. Points: $(1, 0)$ and $(1, -1)$.

$x = -1$: $y^2 + y = -1 + 1 = 0$: $y = 0$ or $y = -1$. Points: $(-1, 0)$ and $(-1, -1)$.

$x = 2$: $y^2 + y = 8 - 2 = 6$. $y = (-1 \pm \sqrt{25})/2 = (-1 \pm 5)/2$: $y = 2$ or $y = -3$. Points: $(2, 2)$ and $(2, -3)$.

We already have several points. Are any of infinite order? Since $E(\mathbb{Q})_{\text{tors}}$ is trivial (example 8.5.2(b)): every non-identity point has infinite order. So $P = (0, 0)$ is a generator. The naive search found it immediately — but this is because the generator has exceptionally small height ($\hat{h}(P) \approx 0.0511$). For curves with generators of large height, the naive search can be prohibitively expensive.

Sieving

The naive search can be dramatically accelerated by *sieving*: using reduction modulo small primes to eliminate x -values that cannot give rise to rational points.

The idea is simple: if $P = (x_0, y_0) \in E(\mathbb{Q})$, then $(x_0 \bmod p, y_0 \bmod p) \in \tilde{E}(\mathbb{F}_p)$ for every prime p of good reduction (where we reduce $x_0 = a/d^2$ modulo p when $p \nmid d$). So $x_0 \bmod p$ must lie in $\{x(Q) : Q \in \tilde{E}(\mathbb{F}_p)\}$, the set of x -coordinates of points on the reduced curve. This set has size $\leq (|\tilde{E}(\mathbb{F}_p)| + 1)/2$ (since most x -values give either two points or none).

Proposition 9.4.3: Sieving Density

For a prime p of good reduction: the proportion of residues $x \in \mathbb{F}_p$ that are x -coordinates of points in $\tilde{E}(\mathbb{F}_p)$ is approximately $1/2$. More precisely:

$$\frac{|\{x \in \mathbb{F}_p : x^3 + Ax + B \in (\mathbb{F}_p^*)^2\}|}{p} = \frac{1}{2} + O(p^{-1/2}). \quad (9.23)$$

By sieving modulo k independent primes: the search space is reduced by a factor of approximately 2^{-k} .

Proof:

The number of $x \in \mathbb{F}_p$ with $x^3 + Ax + B \in (\mathbb{F}_p^*)^2$ is $\frac{1}{2}(|\tilde{E}(\mathbb{F}_p)| - 1 - \epsilon)$ where ϵ accounts for the points with $y = 0$ (the 2-torsion). Since $|\tilde{E}(\mathbb{F}_p)| = p + 1 - a_p$ with $|a_p| \leq 2\sqrt{p}$ (Hasse's bound, ??): the proportion is $(p - a_p)/2p = 1/2 + O(p^{-1/2})$.

The sieving procedure: choose primes $p_1, \dots, p_k$ of good reduction. For each p_i : compute $\tilde{E}(\mathbb{F}_{p_i})$ and record the set $S_i = \{x \bmod p_i : (x, y) \in \tilde{E}(\mathbb{F}_{p_i})\}$. Now enumerate $x_0 = a/d^2$ with $|a| \leq N$, $d \leq D$, and for each, check whether $x_0 \bmod p_i \in S_i$ for all i . By the Chinese Remainder Theorem: the combined sieve modulo $p_1 \cdots p_k$ eliminates a proportion $\approx 1 - 2^{-k}$ of candidates.

Example 9.4.4: Sieving on a Curve with a Large Generator

Consider a curve $E/\mathbb{Q}$ of rank 1 whose generator has canonical height $\hat{h}(P) \approx 5$, corresponding to naive height $h(x(P)) \approx 5$. The naive search space (with $d = 1$) has $|a| \leq e^5 \approx 148$: roughly 300 values to check.

Sieving modulo $p = 3, 5, 7, 11$: at each prime, roughly half the residues are x -coordinates of $\tilde{E}(\mathbb{F}_p)$ -points. After sieving modulo all four primes: the surviving candidates number $\approx 300/16 \approx 19$. Each requires checking whether $f(a) = a^3 + Aa + B$ is a perfect square in $\mathbb{Z}$ — a fast computation. The generator is found among these 19 candidates.

For curves with much larger generators ($\hat{h}(P) \approx 50$, say): the naive search space has $\sim e^{50} \approx 5 \times 10^{21}$ candidates, which is infeasible even with sieving. Such curves require the Heegner point method (section 9.3) or specialized algorithms.

When Generators Have Large Height

The difficulty of finding generators varies enormously across curves. For “typical” curves of small conductor (say $N \leq 1000$): the generators have canonical height $\hat{h} \leq 10$, and a sieved search finds them quickly. But exceptional curves can have generators of enormous height.

Example 9.4.5: Large-Height Generators

The curve $E : y^2 + y = x^3 - x^2 - 79x + 289$ (Cremona label 5765a1) has rank 1 with generator

$$P = \left(\frac{-4457}{529}, \frac{-144875}{12167} \right).$$

The naive height is $h(P) = h(-4457/529) = \log 4457 \approx 8.40$, and the canonical height is $\hat{h}(P) \approx 5.08$. A naive search up to $|a| \leq e^{8.4} \cdot 23^2 \approx 2.4 \times 10^6$ (with $d = 23$, since $529 = 23^2$) would need to test millions of candidates. Sieving reduces this to a manageable number.

For comparison: the curve of rank 1 with the largest known generator height over $\mathbb{Q}$ (among curves of conductor $\leq 500,000$) has $\hat{h}(P) \approx 50$. Finding such a generator by search alone would require examining $\sim e^{100}$ candidates (even with d small), which is completely infeasible. The Heegner point construction is the only practical method in such cases.

The moral: the ease of finding generators depends on the regulator, which is an intrinsic invariant of the curve. Small regulators (like $\hat{h}(P) \approx 0.0511$ for 37a1) mean the generator is “close to the identity” in the height metric and easy to find. Large regulators mean the generator is far from the identity and hard to find by search. The BSD conjecture predicts the regulator from the L -function, giving a prior estimate of how hard the search will be.

Height Bounds from Descent

The Mordell–Weil proof (section 8.4) shows that $E(K)$ is generated by the coset representatives $Q_1, \dots, Q_m$ of $E(K)/nE(K)$ together with the Northcott set $\{P \in E(K) : h(P) \leq B\}$, where B is the explicit bound from proposition 8.4.3. This gives a theoretical height bound for generators: every generator of $E(K)$ has naive height $\leq B$ (after expressing it in terms of the coset representatives and the Northcott set).

The problem is that B is often astronomically large: it depends on the height constants c_1, c_2 (which involve the coefficients of the Weierstrass equation) and on the heights of the coset representatives (which depend on the Selmer group computation). For curves of large conductor, B can be 10^{10} or more, making a direct search impractical.

In practice, the theoretical height bound is rarely used directly. Instead, one uses the following heuristic: if the descent gives rank $\leq r$ and a search up to height B_0 finds r independent points, one declares the rank to be r and verifies by saturation (section 9.5). The theoretical bound serves as a safety net: if the search fails to find enough points, one can (in principle) extend the search up to B and be guaranteed to find them all.

Proposition 9.4.6: The Canonical Height Bound for Generators

Let E/K have rank r and let $P_1, \dots, P_r$ be a basis of $E(K)/E(K)_{\text{tors}}$. If the Néron–Tate regulator satisfies $\text{Reg}(E/K) = \det(\langle P_i, P_j \rangle) = R$: then

$$\min_{1 \leq i \leq r} \hat{h}(P_i) \leq R^{1/r}. \quad (9.24)$$

In particular: at least one generator has canonical height $\leq R^{1/r}$.

Proof:

The Gram matrix $G = (\langle P_i, P_j \rangle)$ is positive definite with $\det(G) = R$. By the AM–GM inequality for the eigenvalues $\lambda_1, \dots, \lambda_r$ of G : $\prod \lambda_i = R$ and $\min \lambda_i \leq R^{1/r}$. Since $\hat{h}(P_i) = \langle P_i, P_i \rangle = G_{ii} \geq \min \lambda_i$: the result follows. (The sharper bound uses Hermite’s inequality for lattices: the shortest vector in a rank- r lattice of determinant R has squared length $\leq \gamma_r R^{1/r}$, where γ_r is Hermite’s constant.)

When the BSD conjecture is available (or when $L(E, s)$ can be computed numerically): the regulator R is predicted by the BSD formula $R = L^{(r)}(E, 1) \cdot |E(K)_{\text{tors}}|^2 / (r! \cdot \Omega \cdot \prod c_p \cdot |\text{III}|)$. So the expected height of the smallest generator is controlled by the L -function. For curves of analytic rank 1: the Gross–Zagier formula gives $R = \hat{h}(P_{\text{Heegner}}) \cdot [\text{index}]^2$, making the bound explicit.

Combining Descent and Search: The Complete Pipeline

We summarize the full procedure for finding generators of $E(\mathbb{Q})$.

9.4.7 Algorithm: Finding Generators of $E(\mathbb{Q})$ **Input:** $E/\mathbb{Q}$ and an upper bound $r_{\max} = \dim \text{Sel}^{(2)} - \dim E(\mathbb{Q})[2]$ from the 2-descent.

Output: Independent points $P_1, \dots, P_r \in E(\mathbb{Q})$ with $r = \text{rank} E(\mathbb{Q})$.

Step 1. Compute the root number $w(E/\mathbb{Q})$. If $w = +1$: expect r even. If $w = -1$: expect r odd.

Step 2. Search for rational points by sieving: enumerate $x_0 = a/d^2$ with $h(x_0) \leq B_0$ (starting with $B_0 = 10$, say), sieving modulo several small primes.

Step 3. For each new point P found: check independence from previously found points using the Néron–Tate height pairing (the Gram matrix $(\langle P_i, P_j \rangle)$ should have nonzero determinant).

Step 4. If the number of independent points found equals $r_{\max}$: the rank is $r_{\max}$ and $\text{III}[2] = 0$. Proceed to saturation (section 9.5).

Step 5. If the point search stalls (no new points found up to height B_0): increase B_0 , or appeal to the L -function (if $L(E, 1) \neq 0$: $r = 0$; if $\text{ord}_{s=1} L(E, s) = 1$: $r = 1$, and a Heegner point gives a generator).

Step 6. If a gap remains between the lower bound (from points) and the upper bound (from Selmer): perform a second descent (section 9.3) or appeal to the parity conjecture.

Exercise 9.4.1

1. Search for rational points on $E : y^2 = x^3 + 1$ with $|x| \leq 5$. List all points found and determine $E(\mathbb{Q})$.
2. For $E : y^2 = x^3 - 2$ over $\mathbb{Q}$ (minimal model, $\Delta = -108$): show that if $P = (a/d^2, c/d^3) \in E(\mathbb{Q})$ with $\gcd(a, d) = 1$, then $d \mid 1$ (i.e., $x_0 \in \mathbb{Z}$). (*Hint:* the only prime that could divide d is p with $v_p(x_0) < 0$; check that the formal group at each prime gives $d = 1$.)
3. For $E : y^2 = x^3 - x$ over $\mathbb{Q}$: compute the set of x -coordinates of $\tilde{E}(\mathbb{F}_p)$ for $p = 3, 5, 7$. What proportion of integers x with $|x| \leq 100$ survives the combined sieve modulo 3, 5, 7?
4. For $E : y^2 + y = x^3 - x$ (37a1) with rank = 1: use the BSD formula to estimate $\text{Reg} = \hat{h}(P)$ and compare with the actual value $\hat{h}(P) \approx 0.0511$. (*Data:* $L'(E, 1) \approx 0.3059$, $\Omega \approx 5.987$, $\prod c_p = 1$, $|\text{III}| = 1$, $|E(\mathbb{Q})_{\text{tors}}| = 1$.)
5. The curve $E : y^2 + y = x^3 - x^2 - 10x - 20$ (conductor 11, the modular curve $X_0(11)$) has rank 0. How does this manifest in the search? (*Answer:* no non-torsion points are found at any height, and $\text{Sel}^{(2)}$ gives $r = 0$ directly.)
6. Estimate the speedup from sieving modulo k primes, assuming each prime eliminates half the candidates. For a search up to $h(x) \leq 20$ (approximately $e^{20} \approx 5 \times 10^8$ candidates with

$d = 1$): how many candidates survive a sieve with $k = 10$ primes?

9.5 Saturation and Verifying Generators

The descent produces an upper bound on the rank, the point search produces candidate generators, and the independence test (via the Néron–Tate height) confirms they are linearly independent modulo torsion. But one question remains: do the points we found generate *all* of $E(K)$, or only a subgroup of finite index?

This is the saturation problem. If $P_1, \dots, P_r$ are independent points of infinite order in $E(K)$, the subgroup $\Lambda = \langle P_1, \dots, P_r \rangle + E(K)_{\text{tors}}$ has finite index in $E(K)$ (since both have rank r), but the index $[E(K) : \Lambda]$ could be greater than 1. A “missing” generator would be a point $Q \in E(K)$ with $Q \notin \Lambda$ but $mQ \in \Lambda$ for some $m \geq 2$: in other words, Q is a “fraction” of a $\mathbb{Z}$ -linear combination of the P_i , invisible to the height pairing on Λ alone.

Saturation is the final step in the Mordell–Weil computation. Without it, the descent-and-search pipeline determines the rank but leaves open the possibility that the found generators span only a proper sublattice of $E(K)/E(K)_{\text{tors}}$.

The Saturation Index and the Regulator

The relationship between the saturation index and the regulator provides the key tool.

Definition 9.5.1: Saturation and the Saturation Index

Let $\Lambda \subseteq E(K)$ be a subgroup of rank r (equal to $\text{rank} E(K)$), and let $\Lambda_{\text{free}} = \Lambda / (\Lambda \cap E(K)_{\text{tors}})$. The *saturation index* is

$$[\overline{E(K)} : \Lambda_{\text{free}}] = n, \quad (9.25)$$

where $\overline{E(K)} = E(K)/E(K)_{\text{tors}}$ is the free part of $E(K)$. The subgroup Λ is *saturated* (or *p-saturated* for a prime p) if $n = 1$ (or $p \nmid n$, respectively). Equivalently: Λ is *p-saturated* iff there is no point $Q \in E(K) \setminus \Lambda$ with $pQ \in \Lambda$.

The saturation index is related to the regulators of Λ and $E(K)$ by a simple formula. Recall that the regulator $\text{Reg}(\Lambda)$ is the determinant of the Gram matrix $(\langle P_i, P_j \rangle)_{1 \leq i, j \leq r}$ for any basis $P_1, \dots, P_r$ of Λ_{free} (??).

Proposition 9.5.2: Regulator and Saturation Index

Let $\Lambda \subseteq E(K)$ have saturation index $n = [\overline{E(K)} : \Lambda_{\text{free}}]$. Then:

$$\text{Reg}(\Lambda) = n^2 \cdot \text{Reg}(E/K). \quad (9.26)$$

In particular: Λ is saturated iff $\text{Reg}(\Lambda) = \text{Reg}(E/K)$.

Proof:

Let $Q_1, \dots, Q_r$ be a basis of $\overline{E(K)}$ and $P_1, \dots, P_r$ a basis of Λ_{free} . The change-of-basis matrix $M \in \text{Mat}_{r \times r}(\mathbb{Z})$ satisfies $P_i = \sum_j M_{ij} Q_j$. The Gram matrices are related by $G_\Lambda = M G_{E(K)} M^T$, so $\text{Reg}(\Lambda) = \det(M)^2 \cdot \text{Reg}(E/K)$. Since $[\overline{E(K)} : \Lambda_{\text{free}}] = |\det(M)|$, $\text{Reg}(\Lambda) = n^2 \cdot \text{Reg}(E/K)$.

This formula converts the saturation problem into a regulator comparison: if we can compute (or estimate) $\text{Reg}(E/K)$ independently, we can determine n from the ratio $\text{Reg}(\Lambda)/\text{Reg}(E/K)$, and if $n = 1$ we are done.

The p -Saturation Test

To check saturation in practice, it suffices to check p -saturation for each small prime p (up to a computable bound). The p -saturation test asks: does there exist a point $Q \in E(K)$ with $pQ \in \Lambda$ but $Q \notin \Lambda$?

Proposition 9.5.3: The p -Saturation Test

Let $\Lambda = \langle P_1, \dots, P_r \rangle + E(K)_{\text{tors}} \subseteq E(K)$. The subgroup Λ fails to be p -saturated iff there exist $a_1, \dots, a_r \in \mathbb{Z}$ (not all divisible by p) and a point $Q \in E(K)$ such that

$$pQ = a_1 P_1 + \dots + a_r P_r + T \quad (9.27)$$

for some $T \in E(K)_{\text{tors}}$. Equivalently: Λ fails to be p -saturated iff the natural map

$$\Lambda/p\Lambda \longrightarrow E(K)/pE(K) \quad (9.28)$$

is not injective.

Proof:

If $Q \in E(K) \setminus \Lambda$ with $pQ \in \Lambda$: then $pQ = \sum a_i P_i + T$ with $T \in E(K)_{\text{tors}}$. The element $\sum a_i P_i + T \in \Lambda$ is divisible by p in $E(K)$ (since pQ divides it) but not in Λ (since $Q \notin \Lambda$). So the image of $\sum a_i P_i + T$ in $\Lambda/p\Lambda$ is nonzero, but its image in $E(K)/pE(K)$ is zero. The map $\Lambda/p\Lambda \rightarrow E(K)/pE(K)$ is not injective.

Conversely: if the map is not injective, there exists a nonzero element $\sum a_i P_i + T \in \Lambda$ (with $T \in E(K)_{\text{tors}}$ and the a_i not all divisible by p) that lies in $pE(K)$. Write $\sum a_i P_i + T = pQ$ for $Q \in E(K)$. Then $Q \notin \Lambda$ (otherwise $\sum a_i P_i + T \in p\Lambda$, contradicting the nonzero image in $\Lambda/p\Lambda$), so Λ is not p -saturated.

To perform the test computationally: the map $\Lambda/p\Lambda \rightarrow E(K)/pE(K)$ factors through the Selmer group. The composition $\Lambda/p\Lambda \rightarrow E(K)/pE(K) \hookrightarrow \text{Sel}^{(p)}(E/K)$ sends $P_i \bmod p\Lambda$ to $\delta(P_i) \in \text{Sel}^{(p)}$. The test for p -saturation reduces to: are $\delta(P_1), \dots, \delta(P_r)$ linearly independent in $\text{Sel}^{(p)}$ (as an $\mathbb{F}_p$ -vector space)?

If the Selmer group is known (from a p -descent), the saturation test is a linear algebra computation over $\mathbb{F}_p$. In practice, $p = 2$ is the most common (and the 2-Selmer group is already computed from the descent of Sections 9.1–9.2).

Proposition 9.5.4: Bound on the Saturation Index

If Λ is p -saturated for all primes $p \leq B$, and the regulator satisfies $\text{Reg}(\Lambda)/\text{Reg}(E/K) < (B+1)^2$: then Λ is fully saturated ($n = 1$).

Proof:

From (9.26): $n^2 = \text{Reg}(\Lambda)/\text{Reg}(E/K) < (B+1)^2$, so $n \leq B$. If $n > 1$: some prime $p \leq n \leq B$ divides n , contradicting p -saturation for $p \leq B$.

In practice: check p -saturation for $p = 2, 3, 5, 7$ (the first few primes) and combine with a regulator bound. If $\text{Reg}(\Lambda)/\text{Reg}(E/K) < 49$ and Λ is p -saturated for $p \leq 5$: then $n = 1$ and Λ is saturated.

The Regulator via the BSD Formula

The “true” regulator $\text{Reg}(E/K)$ is not directly computable without knowing a basis of $E(K)$ (which is what we are trying to find). However, the BSD conjecture provides an independent prediction. The BSD formula (eq. (6.14)) at $s = 1$ for a rank- r curve gives:

$$\text{Reg}(E/K) = \frac{L^{(r)}(E/K, 1)}{r! \cdot \Omega_{E/K} \cdot \prod_{\mathfrak{p}} c_{\mathfrak{p}} \cdot |\text{III}(E/K)|} \cdot |E(K)_{\text{tors}}|^2. \quad (9.29)$$

If III is finite (as conjectured) and its order is known (e.g., from a descent or from the Cassels–Tate pairing): $\text{Reg}(E/K)$ is determined numerically. The saturation index is then $n = \sqrt{\text{Reg}(\Lambda)/\text{Reg}(E/K)}$, which should be an integer.

For rank 1: $\text{Reg}(E/K) = \hat{h}(P_0)$ for the generator P_0 . If we have found a point P of infinite order: $\text{Reg}(\Lambda) = \hat{h}(P)$ and $n^2 = \hat{h}(P)/\hat{h}(P_0)$. If P is already a generator: $n = 1$. If $P = [m]P_0$: $\hat{h}(P) = m^2\hat{h}(P_0)$ and $n = m$.

Worked Examples**Example 9.5.5: Saturation on $y^2 + y = x^3 - x$ (37a1)**

$E(\mathbb{Q}) \cong \mathbb{Z}$, generated by $P_0 = (0, 0)$ with $\hat{h}(P_0) \approx 0.0511$. Suppose we found $P = (0, 0)$ by a point search. Is $\Lambda = \langle P \rangle$ saturated in $E(\mathbb{Q})$?

Regulator comparison. $\text{Reg}(\Lambda) = \hat{h}(P) \approx 0.0511$. From the BSD formula: $\text{Reg}(E/\mathbb{Q}) = L'(E, 1) \cdot |T|^2 / (\Omega \cdot \prod c_p \cdot |\text{III}|) = 0.3059 \cdot 1 / (5.987 \cdot 1 \cdot 1) \approx 0.0511$. So $n^2 = \text{Reg}(\Lambda)/\text{Reg}(E/\mathbb{Q}) \approx 0.0511/0.0511 = 1$: the saturation index is 1. Λ is saturated: $P = (0, 0)$ generates $E(\mathbb{Q})/E(\mathbb{Q})_{\text{tors}}$.

p -saturation test for $p = 2$. The 2-Selmer group has $|\text{Sel}^{(2)}| = 2$ (example 8.5.2(b)). The map $\Lambda/2\Lambda \rightarrow \text{Sel}^{(2)}$ sends $P \bmod 2\Lambda$ to $\delta(P) \in \text{Sel}^{(2)}$. Since $|\text{Sel}^{(2)}| = 2$ and $|E(\mathbb{Q})[2]| = 1$: $|E(\mathbb{Q})/2E(\mathbb{Q})| = 2$, and $\Lambda/2\Lambda \rightarrow E(\mathbb{Q})/2E(\mathbb{Q})$ is an isomorphism. So Λ is 2-saturated. Combined with the regulator comparison: Λ is saturated.

Example 9.5.6: Detecting a Non-Saturated Subgroup

$E : y^2 + y = x^3 - x$ (37a1) again, but suppose we had found $Q = [3]P_0 = (0, -1)$ instead of $P_0 = (0, 0)$ (both have x -coordinate 0, so a naive search might encounter either). Is $\Lambda' = \langle Q \rangle$ saturated?

Regulator comparison. $\hat{h}(Q) = \hat{h}([3]P_0) = 9\hat{h}(P_0) \approx 9 \cdot 0.0511 = 0.460$. $n^2 = \text{Reg}(\Lambda')/\text{Reg}(E/\mathbb{Q}) = 0.460/0.0511 \approx 9$, so $n = 3$. The saturation index is 3: Λ' is *not* saturated. The true generator P_0 satisfies $Q = [3]P_0$, and $P_0 \notin \Lambda'$.

p-saturation test for $p = 3$. We need to check whether there exists $R \in E(\mathbb{Q})$ with $3R \in \Lambda'$ but $R \notin \Lambda'$. Since $P_0 = (0, 0)$ satisfies $3P_0 = Q \in \Lambda'$ and $P_0 \notin \Lambda'$ (because P_0 is not a multiple of $Q = [3]P_0$): Λ' is not 3-saturated. The test detects the missing generator.

Recovery. Once 3-saturation fails: we know there exists R with $[3]R \equiv Q \pmod{E(\mathbb{Q})_{\text{tors}}}$. We search for R : since $\hat{h}(R) = \hat{h}(Q)/9 \approx 0.0511$, the search is easy (small height), and we find $R = P_0 = (0, 0)$.

Example 9.5.7: Saturation for a Rank-2 Curve

$E : y^2 = x(x-1)(x+3)$ from example 9.1.4. Suppose the point search found $P_1 = (-1, 2)$ and $P_2 = (3, 6)$. Is $\Lambda = \langle P_1, P_2 \rangle$ saturated?

The Gram matrix: $G = \begin{pmatrix} \hat{h}(P_1) & \langle P_1, P_2 \rangle \\ \langle P_1, P_2 \rangle & \hat{h}(P_2) \end{pmatrix}$. Suppose $\hat{h}(P_1) \approx 0.23$, $\hat{h}(P_2) \approx 0.69$, $\langle P_1, P_2 \rangle \approx 0.11$. Then $\text{Reg}(\Lambda) = 0.23 \cdot 0.69 - 0.11^2 \approx 0.147$.

The BSD prediction: $\text{Reg}(E/\mathbb{Q})$ is computed from the L -function and the other BSD invariants. If $\text{Reg}(E/\mathbb{Q}) \approx 0.147$: then $n = 1$ and Λ is saturated. If $\text{Reg}(E/\mathbb{Q}) \approx 0.0163$ (which would give $n^2 = 0.147/0.0163 = 9$, $n = 3$): Λ would not be 3-saturated.

The p -saturation test for $p = 2$: compute $\delta(P_1)$ and $\delta(P_2)$ in $\text{Sel}^{(2)}$ and check independence over $\mathbb{F}_2$. If $\dim \text{Sel}^{(2)} = 4$ (the value from the descent) and $\delta(P_1), \delta(P_2)$ are independent modulo torsion images: Λ is 2-saturated.

The Complete Mordell–Weil Pipeline

We summarize the entire computation of $E(\mathbb{Q})$ as a single pipeline, integrating the tools developed across Chapters 8 and 9.

9.5.8 Algorithm: Computing $E(\mathbb{Q})$ **Input:** An elliptic curve $E/\mathbb{Q}$ (given by a Weierstrass equation).

Output: Generators $P_1, \dots, P_r$ of $E(\mathbb{Q})/E(\mathbb{Q})_{\text{tors}}$ and the torsion subgroup $E(\mathbb{Q})_{\text{tors}}$.

Step 1 (Torsion). Compute $E(\mathbb{Q})_{\text{tors}}$ via reduction modulo two primes of good reduction (proposition 8.5.1), verified against Mazur’s list (theorem 8.5.5).

Step 2 (Selmer bound). Compute $\text{Sel}^{(2)}$ via 2-descent (section 9.1) or isogeny descent (section 9.2). This gives $r \leq r_{\max} = \dim_{\mathbb{F}_2} \text{Sel}^{(2)} - \dim_{\mathbb{F}_2} E(\mathbb{Q})[2]$.

Step 3 (Point search). Search for rational points of small height by sieving (section 9.4). Check independence via the Néron–Tate height pairing. Let r_{found} denote the number of independent points found.

Step 4 (Gap resolution). If $r_{\text{found}} < r_{\max}$: resolve the gap by (a) extending the point search, (b) performing a second descent (section 9.3), (c) computing the analytic rank via the L -function, or (d) constructing a Heegner point (if the analytic rank is 1).

Step 5 (Saturation). Test p -saturation for all primes p up to a bound B determined by proposition 9.5.4. If saturation fails at p : enlarge the generating set by finding the “missing” p -th root.

Step 6 (Verification). The regulator $\text{Reg}(\Lambda)$ matches the BSD prediction $\text{Reg}(E/\mathbb{Q})$ (assuming III is known or conjectured to be trivial): $n = 1$, and $P_1, \dots, P_r$ generate $E(\mathbb{Q})/E(\mathbb{Q})_{\text{tors}}$.

This pipeline is implemented in computer algebra systems (SageMath, Magma, PARI/GP) and has been applied to millions of curves. For curves of conductor $\leq 500,000$ in the LMFDB: the rank, generators, and torsion are known for every entry. The saturation step is almost never the bottleneck — the difficulty is concentrated in the descent (Step 2, for curves with large III) and the point search (Step 3, for curves with generators of large height).

This completes the computational chapter. The reader now has all the tools to determine $E(\mathbb{Q}) \cong \mathbb{Z}^r \oplus T$ for a given curve: the torsion T (from Chapter 8), the rank r (from descent and point search), the generators (from sieved search), and the verification that the generators are complete (from saturation). The remaining chapters of the book turn to deeper structural questions: Néron models and the Tate curve (Chapter 11), Galois representations (Chapter 12), complex multiplication (Chapter 13), modular forms (Chapter 14), and the Birch and Swinnerton-Dyer conjecture (Chapter 15).

Exercise 9.5.1

1. Let $E/\mathbb{Q}$ have rank 1 and suppose a point search finds P with $\hat{h}(P) = 2.0$. The BSD formula predicts $\text{Reg}(E/\mathbb{Q}) = 0.5$. What is the saturation index? Find the generator P_0 in terms of P .
2. Explain why p -saturation for $p > r_{\max}$ is automatic (where $r_{\max}$ is the Selmer bound on the rank). (*Hint:* if $p > r_{\max} \geq r$, then p cannot divide the saturation index n because $n \leq \prod p_i^{a_i}$ where the p_i are the primes at which saturation fails, and $n^2 = \text{Reg}(\Lambda)/\text{Reg}(E/K)$.)
3. For $E : y^2 = x^3 - x$ over $\mathbb{Q}$ (rank 0, $E(\mathbb{Q}) = E[2]$): show that the subgroup $\Lambda = \langle (0, 0) \rangle \cong \mathbb{Z}/2\mathbb{Z}$

is “trivially saturated” in the sense that Λ contains generators of $E(\mathbb{Q})_{\text{tors}}$ but the rank is 0, so there is nothing to saturate in the free part.

4. For a rank-2 curve with $\Omega = 3.0$, $\prod c_p = 2$, $|T| = 1$, $|\text{III}| = 1$, and $L''(E, 1)/2 = 1.5$: compute the predicted $\text{Reg}(E/\mathbb{Q})$ from the BSD formula. If your point search found P_1, P_2 with $\text{Reg}(\Lambda) = 1.0$: what is the saturation index?
5. If Λ is not p -saturated: explain how to find the “missing” generator Q with $pQ \in \Lambda$. (*Hint*: $\hat{h}(Q) = \hat{h}(pQ)/p^2$, so Q has smaller canonical height than the known generators. A targeted search at this smaller height finds Q .)
6. Carry out the full Mordell–Weil computation for $E : y^2 = x^3 - 2$: (a) compute $E(\mathbb{Q})_{\text{tors}}$, (b) perform a 2-descent (note: $E[2]$ is not fully rational over $\mathbb{Q}$), (c) search for points, (d) saturate.

Elliptic Curves over Number Rings

The Mordell–Weil theorem determines the group of *rational* points $E(K)$: it is finitely generated. A natural follow-up question asks about *integral* points — those whose coordinates lie in the ring of integers $\mathcal{O}_K$. The rational points form an infinite set when the rank is positive ($E(K) \cong \mathbb{Z}^r \oplus T$ with $r \geq 1$ contains infinitely many non-torsion points), but the integral points are always finite. This striking contrast — infinitely many rational points, but only finitely many integral ones — is the content of Siegel’s theorem (1929).

The finiteness has two proofs, reflecting two different traditions in number theory. The first, due to Siegel, uses Diophantine approximation: the Thue–Siegel–Roth theorem on how well algebraic numbers can be approximated by rationals. This proof is elegant but *ineffective*: it guarantees finiteness without giving any bound on the size of the integral points. The second proof, due to Baker (1960s), uses lower bounds for linear forms in logarithms of algebraic numbers. Baker’s method is *effective*: it produces explicit (though often enormous) upper bounds on $|x|$ for integral points (x, y) . Combined with the LLL algorithm for lattice basis reduction, Baker’s bounds can be reduced to a range amenable to direct search, making the computation of all integral points a practical endeavour.

This chapter develops both approaches and their Diophantine consequences. We prove Siegel’s theorem, develop the Baker–LLL pipeline for computing integral points, and close with the ABC and Szpiro conjectures, which connect the arithmetic of elliptic curves to the deepest open problems in number theory.

10.1 Siegel’s Theorem

An integral point on an elliptic curve $E : y^2 = x^3 + Ax + B$ (with $A, B \in \mathbb{Z}$) is a point $P = (x_0, y_0)$ with $x_0, y_0 \in \mathbb{Z}$. Equivalently, viewing E as an affine variety $\text{Spec } \mathbb{Z}[x, y]/(y^2 - x^3 - Ax - B)$: an integral point is a section $\text{Spec } \mathbb{Z} \rightarrow E$. The finiteness of integral points is not obvious: the Weierstrass

equation $y^2 = x^3 + Ax + B$ is a cubic in x and quadratic in y , and there is no a priori reason why only finitely many integer values of x should make the right side a perfect square.

Statement and First Examples

Theorem 10.1.1: Siegel's Theorem

Let $E/\mathbb{Q}$ be an elliptic curve given by a Weierstrass equation with $A, B \in \mathbb{Z}$. Then the set

$$\{(x, y) \in E(\mathbb{Q}) : x, y \in \mathbb{Z}\} \quad (10.1)$$

is finite.

The theorem generalizes to number fields and to S -integral points. We state the full version after the examples.

Example 10.1.2: Integral Points on Running Curves

(a) $E : y^2 = x^3 - x$ over $\mathbb{Q}$. $E(\mathbb{Q}) = \{O, (0, 0), (1, 0), (-1, 0)\} \cong (\mathbb{Z}/2\mathbb{Z})^2$: rank 0, so only finitely many rational points altogether. The integral points (with $x, y \in \mathbb{Z}$): $(0, 0), (1, 0), (-1, 0)$. That's all — three integral points (plus the point at infinity O). Siegel's theorem is trivially satisfied since $E(\mathbb{Q})$ is already finite.

(b) $E : y^2 + y = x^3 - x$ over $\mathbb{Q}$ (37a1). $E(\mathbb{Q}) \cong \mathbb{Z}$, generated by $P = (0, 0)$. This curve has infinitely many rational points ($[n]P$ for all $n \in \mathbb{Z}$), but only finitely many integral ones. Computing multiples:

$$\begin{aligned} [1]P &= (0, 0): && \text{integral} \\ [2]P &= (1, -1): && \text{integral} \\ [3]P &= (0, -1): && \text{integral} \\ [4]P &= (-1, 1): && \text{integral} \\ [5]P &= (2, -3): && \text{integral} \\ [6]P &= (6, -14): && \text{integral} \\ [7]P &= (-5/9, 8/27): && \text{not integral} \\ [8]P &= (21/25, -69/125): && \text{not integral} \end{aligned}$$

The denominators grow: once $[n]P$ has non-integer coordinates, all further multiples $[m]P$ with $m > n$ also have non-integer coordinates (with rare exceptions from cancellation). The integral points are: $(0, 0), (0, -1), (1, 0), (1, -1), (-1, 0), (-1, 1), (2, -3), (2, 2), (6, -14), (6, 13)$ (each x -value gives two points from $y^2 + y = f(x)$, related by $y \mapsto -1 - y$). A complete computation (via Baker's method, section 10.2) confirms these are all: 10 integral points in total.

(c) $E : y^2 = x^3 + 1$ over $\mathbb{Q}$ (36a1). $E(\mathbb{Q}) \cong \mathbb{Z}/6\mathbb{Z}$, rank 0. The rational points: $O, (-1, 0), (0, \pm 1), (2, \pm 3)$. All are integral: 5 integral points.

(d) The Mordell equation $y^2 = x^3 - 2$. This classical equation has $E(\mathbb{Q}) \cong \mathbb{Z}$, generated by

$P = (3, 5)$ (as computed in exercise 9.5.1 (6)). The integral points: $(3, \pm 5)$ are the only ones. This was first proved by Fermat (for the positive solution) and later by Euler, using the arithmetic of $\mathbb{Z}[\sqrt{-2}]$.

S -Integral Points

The correct generalization of Siegel's theorem uses S -integers: integers whose denominators are divisible only by primes in a fixed finite set S .

Definition 10.1.3: S -Integral Points

Let K be a number field, S a finite set of places of K containing all archimedean places, and $\mathcal{O}_{K,S} = \{x \in K : v(x) \geq 0 \text{ for all } v \notin S\}$ the ring of S -integers. An S -integral point on $E : y^2 = x^3 + Ax + B$ (with $A, B \in \mathcal{O}_{K,S}$) is a point $P = (x_0, y_0) \in E(K)$ with $x_0, y_0 \in \mathcal{O}_{K,S}$.

Over $\mathbb{Q}$ with $S = \{\infty\}$: $\mathcal{O}_{K,S} = \mathbb{Z}$, and S -integral points are the usual integral points. Enlarging S allows denominators at the primes in S : for $S = \{\infty, 2\}$, the " S -integers" are $\mathbb{Z}[1/2]$, and we allow x, y with denominators that are powers of 2. Siegel's theorem extends:

Theorem 10.1.4: Siegel's Theorem (S -Integral Version)

Let E/K be an elliptic curve over a number field, S a finite set of places, and $\mathcal{O}_{K,S}$ the ring of S -integers. Then the set of S -integral points on any affine Weierstrass model of E is finite.

Roth's Theorem on Diophantine Approximation

The key input from Diophantine approximation is the theorem of Roth (1955), which answers the fundamental question: how well can an irrational algebraic number be approximated by rationals?

The story begins with Liouville (1844), who proved that an algebraic number α of degree d satisfies $|\alpha - p/q| > c(\alpha) \cdot q^{-d}$ for all $p/q \in \mathbb{Q}$ (with $c(\alpha) > 0$ depending on α). The exponent d was progressively improved: Thue (1909) achieved $d/2 + 1$; Siegel (1921) achieved $2\sqrt{d}$; Dyson and Gel'fond (1947) achieved $\sqrt{2d}$. Roth's theorem gives the optimal result: the exponent is $2 + \epsilon$ for any $\epsilon > 0$, independent of the degree.

Theorem 10.1.5: Roth's Theorem

Let α be an irrational algebraic number and $\epsilon > 0$. Then the inequality

$$\left| \alpha - \frac{p}{q} \right| < \frac{1}{q^{2+\epsilon}} \quad (10.2)$$

has only finitely many solutions $p/q \in \mathbb{Q}$ (with $q > 0$).

The exponent 2 is sharp: the theory of continued fractions shows that every irrational α has infinitely many approximations with $|\alpha - p/q| < 1/q^2$ (the convergents of the continued fraction expansion). So the transition from finiteness to infiniteness happens exactly at the exponent 2.

The proof of Roth's theorem (for which Roth received the Fields Medal in 1958) proceeds by contradiction. One assumes infinitely many good approximations p_n/q_n and constructs an auxiliary polynomial $P(x_1, \dots, x_m)$ in several variables that vanishes at $(\alpha, \dots, \alpha)$ to high order but also vanishes at $(p_1/q_1, \dots, p_m/q_m)$ to an order that can be estimated from below (using the approximation quality) and from above (using the “index” of P at the algebraic point). The two bounds contradict each other.

The key limitation of Roth's theorem for our purposes is its *ineffectivity*: the contradiction shows that infinitely many good approximations cannot exist, but gives no way to determine *which* finitely many approximations survive. In particular, it gives no bound on the largest q for which a good approximation exists. This ineffectivity propagates to every result proved via Roth, including Siegel's theorem.

The Proof of Siegel's Theorem

The proof combines two ingredients: the finite generation of $E(K)$ (Mordell–Weil, theorem 8.4.1) and Roth's theorem (or more precisely, its higher-dimensional generalization, the Schmidt subspace theorem). We give the argument in two steps.

Step 1: Reducing to a Unit Equation. Let $P = (x_0, y_0) \in E(K)$ be an S -integral point. From the Weierstrass equation $y^2 = (x - e_1)(x - e_2)(x - e_3)$ (with $e_i \in \bar{K}$, the roots of the cubic): each factor $x_0 - e_i$ is an algebraic integer (in the ring of integers of $K(e_i)$), and their product $(x_0 - e_1)(x_0 - e_2)(x_0 - e_3) = y_0^2$ is a perfect square. By analyzing the factorization in $\mathcal{O}_{L,S'}$ (where $L = K(e_1, e_2, e_3)$ is the splitting field and S' is the set of places of L above S): each $x_0 - e_i$ is an S' -unit times a square, up to a finite number of possibilities determined by the 2-torsion of the class group $\text{Cl}(\mathcal{O}_{L,S'})$. The differences $(x_0 - e_i) - (x_0 - e_j) = e_j - e_i$ are fixed nonzero constants. So the three quantities $x_0 - e_1, x_0 - e_2, x_0 - e_3$ are S' -units (times squares) whose pairwise differences are fixed. After dividing by $e_2 - e_1$, this reduces the problem to a system of S -unit equations:

$$u + v = 1, \quad u, v \in \mathcal{O}_{L,S'}^*, \quad (10.3)$$

where $u = (x_0 - e_1)/(e_2 - e_1)$ and $v = (x_0 - e_2)/(e_2 - e_1)$.

Example 10.1.6: The Unit Equation for $y^2 = x^3 - x$

$E : y^2 = x(x - 1)(x + 1)$ with $e_1 = 0, e_2 = 1, e_3 = -1$. For an integral point (x_0, y_0) with $x_0 > 1$: the three factors are $x_0, x_0 - 1, x_0 + 1$, three consecutive integers whose product is y_0^2 . Since $\gcd(x_0, x_0 - 1) = 1$ and $\gcd(x_0 - 1, x_0 + 1) \mid 2$: the three factors are nearly coprime, and each must be (up to a factor of 2) a perfect square. Setting $u = x_0/(x_0 - 1)$ and $v = -1/(x_0 - 1)$: $u + v = 1$, and both u, v are S -units for $S = \{2, \infty\}$.

The S -unit equation $u + v = 1$ with $u, v \in \mathbb{Z}[1/2]^* = \{\pm 2^a : a \in \mathbb{Z}\}$ has only finitely many solutions. Checking: $\pm 2^a \pm 2^b = 1$ requires one of a, b to be 0. The solutions are $(u, v) = (-1, 2), (2, -1)$, and variations with signs that we enumerate. Each translates back to finitely many values of x_0 .

Step 2: Finiteness of the Unit Equation. The equation $u + v = 1$ in S -units has only finitely many solutions. This is a theorem of Siegel (1929), later sharpened by Evertse (1984) and others.

Theorem 10.1.7: Finiteness of the S -Unit Equation

Let K be a number field and S a finite set of places. The equation

$$u + v = 1, \quad u, v \in \mathcal{O}_{K,S}^* \quad (10.4)$$

has only finitely many solutions.

Proof Sketch:

The Dirichlet S -unit theorem gives $\mathcal{O}_{K,S}^* \cong \mu_K \times \mathbb{Z}^{|S|-1}$, so every S -unit is determined (up to roots of unity) by its exponent vector in $\mathbb{Z}^{|S|-1}$.

The proof reduces to theorem 10.1.5 (or rather, to the Schmidt subspace theorem, its higher-dimensional generalization). Suppose there are infinitely many solutions (u_n, v_n) . Writing each in the S -unit basis: the exponent vectors $(a_1^{(n)}, \dots, a_s^{(n)})$ grow without bound. For each archimedean place v : the equation $|u_n|_v + |v_n|_v \geq |u_n + v_n|_v = 1$ forces a balance between $|u_n|_v$ and $|v_n|_v$. If $|u_n|_v$ is much larger than $|v_n|_v$: $|u_n|_v \approx 1$ (since $v_n = 1 - u_n$), which means $|u_n - 1|_v$ is small. This gives a good v -adic approximation of 1 by the S -unit u_n .

The Schmidt subspace theorem (the n -variable generalization of Roth's theorem) states: for linear forms $L_1, \dots, L_n$ in n variables, the inequality $\prod_v \prod_i |L_i(\mathbf{x})|_v < \|\mathbf{x}\|^{-\epsilon}$ has solutions lying in finitely many proper subspaces of K^n . Applying this to the linear forms $u, v, 1$ (with the equation $u + v - 1 = 0$ as the constraint): the infinitely many solutions must lie in finitely many one-dimensional subspaces, each of which contributes at most finitely many S -unit solutions (by a lower-dimensional argument). The result: finitely many solutions in total.

Combining Steps 1 and 2: the set of S -integral points on E maps (via the factorization $x_0 - e_i$) to the set of solutions of finitely many S -unit equations (one for each class in the 2-part of the class group of L). Each equation has finitely many solutions, so the set of S -integral points is finite. This proves theorem 10.1.4.

Ineffectivity and Evertse's Bound

Siegel's proof inherits the ineffectivity of Roth's theorem: the proof shows that infinitely many solutions lead to a contradiction, but the contradiction does not produce a computable bound on the solutions. Consequently, Siegel's theorem tells us the set of integral points is finite, but gives no bound on the size of the largest integral point $|x_0|$, nor on the total number of integral points. We discuss the effective replacement (Baker's method) in section 10.2.

Although the size of the integral points is unbounded by Siegel's method, the *number* of solutions is bounded by subsequent work.

Theorem 10.1.8: Evertse's Bound on S -Unit Solutions

The number of solutions to the S -unit equation $u + v = 1$ with $u, v \in \mathcal{O}_{K,S}^*$ is at most $3 \cdot 7^{|S|+2[K:\mathbb{Q}]}$.

Evertse's bound (1984) depends only on $|S|$ and $[K : \mathbb{Q}]$, not on the specific field K or the set S . For the integral points on an elliptic curve over $\mathbb{Q}$ (with $|S|$ determined by the discriminant): this gives a bound on the number of integral points. In practice, the bound is far from sharp: most elliptic curves over $\mathbb{Q}$ have fewer than 20 integral points, and the current record is 72 integral points on the curve $y^2 = x^3 - 1641843x + 784512624$ (found by Elkies).

10.1.9 Remark: Effective vs. Ineffective in Number Theory The distinction between effective and ineffective results is a central theme of Diophantine geometry. An “effective” result gives a computable bound: there exists an algorithm that, given the input, terminates and produces the answer. An “ineffective” result proves existence without providing any algorithm. Siegel's theorem (1929) was ineffective; Baker's theorem (1966) made the integral point problem effective. A similar story plays out for the Mordell conjecture (proved ineffectively by Faltings in 1983; an effective proof remains open, though conditional results exist via the ABC conjecture). The Shafarevich conjecture, discussed below, was also proved ineffectively by Faltings, and effective versions remain an active area of research.

Shafarevich's Theorem

Siegel's theorem says: on a *fixed* elliptic curve, there are only finitely many integral points. A deeper finiteness result, due to Shafarevich (conjectured 1962, proved by Faltings 1983), asks: for a fixed set of “bad” primes, how many elliptic curves can there be?

The question is natural from the perspective of S -integers. An elliptic curve E/K has *good reduction outside S* if for every prime $\mathfrak{p} \notin S$: the reduced curve $\tilde{E}/k_{\mathfrak{p}}$ is smooth (equivalently, $v_{\mathfrak{p}}(\Delta_{\min}) = 0$). The set S controls where E is allowed to have bad reduction. Shafarevich asked: for a fixed number field K and a fixed finite set S , are there only finitely many isomorphism classes of elliptic curves over K with good reduction outside S ?

Theorem 10.1.10: Shafarevich's Theorem

Let K be a number field and S a finite set of places. Then there are only finitely many K -isomorphism classes of elliptic curves E/K with good reduction at all primes $\mathfrak{p} \notin S$.

The proof, due to Faltings (1983), proceeds in three stages. The first stage reduces the problem to a statement about abelian varieties with prescribed properties (the “Shafarevich conjecture for abelian varieties”). The second uses the theory of heights on moduli spaces (specifically, the Faltings height of an abelian variety) to convert the finiteness into a height bound. The third uses the isogeny theorem — two abelian varieties with isomorphic Tate modules $T_{\ell}(A) \cong T_{\ell}(B)$ (as Galois representations) are isogenous — to reduce the problem to the Tate conjecture, which Faltings proved in the same paper.

We outline the key ideas without giving the full proof (which requires the machinery of Arakelov geometry and p -divisible groups, beyond our scope).

Key idea 1: From S -good reduction to S -integral points on a moduli space. An elliptic curve E/K with good reduction outside S has j -invariant $j(E) \in \mathcal{O}_{K,S}$ (since the minimal discriminant is an S -unit, and $j = c_4^3/\Delta$ is S -integral). Conversely, for each $j \in \mathcal{O}_{K,S}$ with $j \neq 0, 1728$: there are at most finitely many twists E/K with that j -invariant (parametrized by $H^1(K, \text{Aut}(E))$, which is finite for $j \neq 0, 1728$). So the finiteness of elliptic curves with good reduction outside S reduces to: finitely many $j \in \mathcal{O}_{K,S}$ can arise. But $\mathcal{O}_{K,S}$ is infinite (unless $K = \mathbb{Q}$ and $S = \{\infty\}$), so this

reduction alone does not suffice — we need the constraint that $j(E) \in \mathcal{O}_{K,S}$ arises from an actual elliptic curve with good reduction, not just an arbitrary S -integer.

Key idea 2: The Faltings height is bounded. Faltings defines a canonical height $h_F(E)$ on the moduli space of elliptic curves (analogous to the Néron–Tate height on $E(K)$) and shows: if E has good reduction outside S , then $h_F(E) \leq C(K, S)$ for a computable constant. This uses the semistable reduction theorem (every elliptic curve acquires semistable reduction after a finite extension) and the fact that the Faltings height behaves well under isogeny.

Key idea 3: Bounded height implies finiteness. By a Northcott-type theorem for the Faltings height: the set of K -isomorphism classes of elliptic curves with $h_F(E) \leq B$ is finite (for any $B > 0$).

Example 10.1.11: Shafarevich's Theorem over $\mathbb{Q}$

Over $\mathbb{Q}$ with $S = \{\infty, 2, 3\}$ (good reduction outside $\{2, 3\}$): the elliptic curves with $|\Delta_{\min}|$ divisible only by 2 and 3 include $y^2 = x^3 - x$ ($\Delta = -64 = -2^6$), $y^2 = x^3 + 1$ ($\Delta = -432 = -2^4 \cdot 3^3$), $y^2 = x^3 + x$ ($\Delta = -64$), and finitely many others. A complete enumeration (via Cremona's tables) finds exactly 24 isomorphism classes.

Over $\mathbb{Q}$ with $S = \{\infty, 2, 3, 5, 7\}$: the number of curves grows, but is still finite. The finiteness is guaranteed by Shafarevich's theorem; the enumeration requires computing all elliptic curves with conductor dividing $2^a \cdot 3^b \cdot 5^c \cdot 7^d$ for bounded exponents.

Shafarevich's theorem connects to the broader Diophantine landscape through two deep generalizations. First, Faltings proved the Shafarevich conjecture for abelian varieties of arbitrary dimension (not just elliptic curves), and used this to deduce the *Mordell conjecture*: every curve of genus ≥ 2 over a number field has only finitely many rational points. Second, a quantitative version of Shafarevich's theorem — bounding $|\Delta_{\min}|$ in terms of the conductor N — is Szpiro's conjecture, which we discuss in section 10.4.

Integral Points and the Canonical Height

The connection between integral points and heights provides an alternative perspective on Siegel's theorem, and a bridge to Baker's method.

Proposition 10.1.12: Height Growth of Integral Points

Let $E/\mathbb{Q}$ be an elliptic curve of rank $r \geq 1$, with generator(s) $P_1, \dots, P_r$ of $E(\mathbb{Q})/E(\mathbb{Q})_{\text{tors}}$. If $P = n_1 P_1 + \dots + n_r P_r + T$ (with $T \in E(\mathbb{Q})_{\text{tors}}$) is an integral point: then $\hat{h}(P) = \sum_{i,j} n_i n_j \langle P_i, P_j \rangle$, and the naive height $h(P) = h(x(P)) = \log |x(P)|$ (for $x(P) \in \mathbb{Z}$ with $|x(P)|$ large). So:

$$\log |x(P)| \approx \hat{h}(P) + O(1) = \sum_{i,j} n_i n_j \langle P_i, P_j \rangle + O(1). \quad (10.5)$$

The right side is a positive definite quadratic form in the n_i : it grows like $\|n\|^2$, where $\|n\| = \sqrt{n_1^2 + \dots + n_r^2}$.

The proposition shows: for P to be an integral point, $|x(P)|$ must be an integer, and $|x(P)|$ grows

exponentially with $\|n\|$ (since $|x(P)| \approx e^{\hat{h}(P)} \approx e^{c\|n\|^2}$). As $\|n\| \rightarrow \infty$: the ratio $|x(P)|^3/y(P)^2 \rightarrow 1$ (from the Weierstrass equation), which means $x(P)$ is very close to a perfect cube. The Diophantine content of Siegel's theorem is that this “near-cube” condition is satisfied for only finitely many integers.

Example 10.1.13: Height Growth on 37a1

$E : y^2 + y = x^3 - x$, $P = (0, 0)$, $\hat{h}(P) \approx 0.0511$. The naive height of $[n]P$ is approximately $h([n]P) \approx n^2 \cdot 0.0511$. For $|n| \leq 6$: $h([n]P) \leq 36 \cdot 0.0511 \approx 1.84$, so $|x([n]P)| \leq e^{1.84} \approx 6.3$. Indeed, $x([6]P) = 6$. For $|n| = 7$: $h \approx 2.50$, $|x| \approx e^{2.50} \approx 12$, but $x([7]P) = -5/9$ (not an integer). The denominators in $x([n]P)$ grow rapidly for $|n| \geq 7$.

The pattern: the first few multiples have small integer x -coordinates (giving integral points), but beyond a threshold, the denominators grow too fast for integrality to persist. The threshold is roughly $|n| \approx \sqrt{B/\hat{h}(P)}$ where B is the largest height at which integral points occur. For this curve: $B \approx h(6, \pm 14) = \log 6 \approx 1.79$, giving $|n| \leq \sqrt{1.79/0.0511} \approx 5.9$. So $|n| \leq 6$ is the range in which integral points can appear, consistent with the 10 integral points found in example 10.1.2(b).

Exercise 10.1.1

1. Find all integral points on $y^2 = x^3 + 1$ by exhaustive search (the rank is 0, so there are only finitely many rational points).
2. For the congruent number curve $E_5 : y^2 = x^3 - 25x$: find all integral points with $|x| \leq 100$ by direct computation.
3. For $E : y^2 = x(x-1)(x+1)$ with $e_1 = 0, e_2 = 1, e_3 = -1$: write the unit equation for an integral point (x_0, y_0) by setting $u = x_0/(x_0 - 1)$ and $v = -1/(x_0 - 1)$. Verify $u + v = 1$ and explain why both u and v are S -units for $S = \{2, \infty\}$.
4. Show that the exponent 2 in Roth's theorem is sharp: for $\alpha = (1 + \sqrt{5})/2$ (the golden ratio), the convergents p_n/q_n of the continued fraction expansion satisfy $|\alpha - p_n/q_n| < 1/q_n^2$ for infinitely many n . (*Hint*: the continued fraction of α is $[1; 1, 1, 1, \dots]$, and the convergents are ratios of consecutive Fibonacci numbers.)
5. Over $\mathbb{Q}$ with $S = \{\infty, 2\}$ (good reduction outside $\{2\}$): find all elliptic curves $E/\mathbb{Q}$ with $\Delta_{\min}$ a power of ± 2 . (*Hint*: the conductor N divides a power of 2. Consult Cremona's tables for conductor ≤ 256 .)
6. For $E : y^2 = x^3 - 2$ and $S = \{\infty, 3\}$: explain why there might be more S -integral points than integral points, and find any additional ones with $|x| \leq 50$.

10.2 Baker's Method and Effective Bounds

Siegel's theorem tells us that the set of integral points is finite, but gives no bound on their size. Baker's method (1966) provides the missing ingredient: an *effective* lower bound for linear forms in logarithms of algebraic numbers. This bound, combined with the reduction of integral points to S -unit equations (section 10.1), yields explicit upper bounds on $|x|$ for integral points — bounds

that are typically enormous ($|x| \leq 10^{500}$ is not unusual) but are finite and computable. The LLL algorithm then reduces these bounds to a manageable range, completing the pipeline for finding all integral points on a given curve.

Linear Forms in Logarithms

The central question is: how small can a nonzero linear combination $b_1 \log \alpha_1 + \cdots + b_n \log \alpha_n$ be, where the α_i are algebraic numbers and the b_i are integers?

The question is rooted in the Gelfond–Schneider theorem (1934), which answered the special case $n = 2$: if $\log \alpha_1$ and $\log \alpha_2$ are linearly independent over $\mathbb{Q}$ (equivalently, $\log \alpha_1 / \log \alpha_2 \notin \mathbb{Q}$), then $b_1 \log \alpha_1 + b_2 \log \alpha_2 \neq 0$ for all nonzero $(b_1, b_2) \in \mathbb{Z}^2$. This is equivalent to the transcendence of $\alpha_1^{b_2/b_1}$ when it is well-defined: the Gelfond–Schneider theorem says α^β is transcendental when $\alpha \neq 0, 1$ is algebraic and β is algebraic and irrational. Baker's breakthrough was the generalization to n terms and, crucially, the *quantitative* lower bound.

If the $\log \alpha_i$ are linearly independent over $\mathbb{Q}$ (equivalently, no multiplicative relation $\alpha_1^{a_1} \cdots \alpha_n^{a_n} = 1$ holds with $a_i \in \mathbb{Z}$ not all zero): the linear form $\Lambda = b_1 \log \alpha_1 + \cdots + b_n \log \alpha_n$ is never zero. The question is how close to zero it can get as $B = \max |b_i| \rightarrow \infty$. A trivial upper bound is $|\Lambda| \leq B \sum |\log \alpha_i| = O(B)$. Baker's theorem gives the decisive *lower* bound.

Theorem 10.2.1: Baker's Theorem on Linear Forms in Logarithms

Let $\alpha_1, \dots, \alpha_n$ be nonzero algebraic numbers and $b_1, \dots, b_n \in \mathbb{Z}$ with $B = \max(|b_1|, \dots, |b_n|, 2)$. If

$$\Lambda = b_1 \log \alpha_1 + b_2 \log \alpha_2 + \cdots + b_n \log \alpha_n \neq 0, \quad (10.6)$$

then

$$|\Lambda| > \exp\left(-C(n, d) \cdot \prod_{i=1}^n \log A_i \cdot \log B\right), \quad (10.7)$$

where $d = [\mathbb{Q}(\alpha_1, \dots, \alpha_n) : \mathbb{Q}]$, $A_i = \max(H(\alpha_i), e)$ (with $H(\alpha_i)$ the absolute multiplicative height), and $C(n, d)$ is an explicitly computable constant depending only on n and d .

Baker proved the theorem in a series of papers in the late 1960s, earning the Fields Medal in 1970. The constant $C(n, d)$ has been refined by many authors; the current best bounds are due to Matveev (2000):

$$C(n, d) \leq (16nd)^{2(n+2)}. \quad (10.8)$$

This is large — for $n = 3$ and $d = 6$: $C \leq (288)^{10} \approx 2.4 \times 10^{24}$ — but it is *explicit* and *computable*. The key point is that $\log B$ appears only once: the bound is essentially $\exp(-C' \log B) = B^{-C'}$ (a polynomial in $1/B$), not $\exp(-C'B)$ (which would be exponentially small in B). This polynomial decay is what makes Baker's theorem effective.

The p -adic analogue is needed because the reduction from integral points to S -unit equations involves not only archimedean absolute values but also p -adic ones. At each prime $p \in S$: the S -unit equation $u + v = 1$ gives a p -adic linear form in logarithms that must also be bounded from below. The p -adic Baker theorem, proved by van der Poorten (1977) and refined by Kunrui Yu (1999), provides the required bound.

Theorem 10.2.2: p -adic Baker Theorem

Let $\alpha_1, \dots, \alpha_n \in \overline{\mathbb{Q}}_p^*$ and $b_1, \dots, b_n \in \mathbb{Z}$ with $B = \max |b_i|$. If $\Lambda_p = \alpha_1^{b_1} \cdots \alpha_n^{b_n} - 1 \neq 0$: then

$$v_p(\Lambda_p) \leq C_p(n, d) \cdot \prod_{i=1}^n \log A_i \cdot \log B, \quad (10.9)$$

with an explicit constant $C_p(n, d)$.

The p -adic valuation $v_p(\Lambda_p)$ measures how close $\alpha_1^{b_1} \cdots \alpha_n^{b_n}$ is to 1 in the p -adic metric: $|\Lambda_p|_p = p^{-v_p(\Lambda_p)}$. The bound says the p -adic “closeness to 1” grows at most polynomially in $\log B$ — the same qualitative behaviour as the archimedean case. The effective S -unit theorem requires applying *both* the archimedean and p -adic Baker theorems simultaneously (one for each place in S), and the resulting bound B_0 involves the product of all the Baker constants.

Application to the S -Unit Equation

Baker's theorem makes Step 2 of Siegel's proof effective. Recall the S -unit equation (theorem 10.1.7): $u + v = 1$ with $u, v \in \mathcal{O}_{K,S}^*$. Writing u in the S -unit basis: $u = \zeta \cdot \eta_1^{a_1} \cdots \eta_s^{a_s}$ where ζ is a root of unity, $\eta_1, \dots, \eta_s$ are fundamental S -units, and $a_i \in \mathbb{Z}$. Then:

$$v = 1 - u = 1 - \zeta \cdot \eta_1^{a_1} \cdots \eta_s^{a_s}. \quad (10.10)$$

For each archimedean place $\sigma : K \hookrightarrow \mathbb{C}$: $|\sigma(v)| = |1 - \sigma(u)|$. If $|u|$ is large (say $|u| > 2$): $|v| = |1 - u| \approx |u|$, and $\log |v| \approx \log |u| = a_1 \log |\eta_1| + \cdots + a_s \log |\eta_s| + O(1)$. But v is also an S -unit: $v = \zeta' \cdot \eta_1^{a'_1} \cdots \eta_s^{a'_s}$. So:

$$a'_1 \log |\eta_1| + \cdots + a'_s \log |\eta_s| = \log |v| \approx a_1 \log |\eta_1| + \cdots + a_s \log |\eta_s|.$$

The difference $(a_1 - a'_1) \log |\eta_1| + \cdots + (a_s - a'_s) \log |\eta_s| \approx 0$: a linear form in logarithms of algebraic numbers that is close to zero. Baker's theorem (theorem 10.2.1) gives a lower bound:

$$|(a_1 - a'_1) \log |\eta_1| + \cdots + (a_s - a'_s) \log |\eta_s|| > \exp(-C \cdot \log A), \quad (10.11)$$

where $A = \max(|a_i|, |a'_i|)$. The two bounds together give $\exp(-C \log A) < |\text{error}| \leq e^{-cA}$ (the error from the approximation $|v| \approx |u|$ decays exponentially). This forces $C \log A > cA$: a contradiction for A large. The upshot: $A \leq A_0$ for an explicit bound A_0 depending on the η_i and the Baker constant C .

Proposition 10.2.3: Effective S -Unit Equation

Let K be a number field and S a finite set of places. There exists an explicitly computable constant $B_0 = B_0(K, S)$ such that every solution (u, v) to $u + v = 1$ with $u, v \in \mathcal{O}_{K,S}^*$ satisfies $h(u) \leq B_0$ (where h is the absolute logarithmic height).

The bound B_0 is typically very large (of order $\exp(\exp(C \cdot |S|^n))$ for some constants C, n), but it is *finite* and *computable*: this is the decisive advantage over the ineffective Roth approach.

Application to Integral Points on Elliptic Curves

The reduction of integral points to S -unit equations (Step 1 of Siegel's proof, section 10.1) combined with the effective S -unit theorem (proposition 10.2.3) immediately gives:

Theorem 10.2.4: Effective Bound for Integral Points

Let $E : y^2 = x^3 + Ax + B$ with $A, B \in \mathbb{Z}$ and $\Delta \neq 0$. There exists an explicitly computable constant $X_0 = X_0(A, B)$ such that every integral point $(x, y) \in E(\mathbb{Z})$ satisfies $|x| \leq X_0$.

The bound X_0 depends on A, B through the Baker constants and the arithmetic of the splitting field $L = \mathbb{Q}(e_1, e_2, e_3)$ (where e_i are the roots of $x^3 + Ax + B$). For specific curves, X_0 is computable but typically enormous: the first applications of Baker's method gave bounds of order $10^{10^{10}}$. Modern refinements (David, Hirata-Kohno, Matveev) bring this down substantially, but bounds of order 10^{500} remain common before the LLL reduction step.

To see where the bound comes from concretely: for an integral point (x_0, y_0) with $|x_0|$ large, the factorization $y_0^2 = (x_0 - e_1)(x_0 - e_2)(x_0 - e_3)$ in $\mathcal{O}_L$ forces each factor $x_0 - e_i$ to be an S' -unit times a square, where S' is determined by the primes above 2Δ in L . Writing $x_0 - e_1 = \alpha \cdot \eta_1^{a_1} \cdots \eta_s^{a_s} \cdot \beta^2$ (with α from a finite set, η_j fundamental S' -units, and $\beta \in \mathcal{O}_L$): the exponents a_j satisfy $\max |a_j| \leq C_1 \log |x_0|$ (from the relation between the height of $x_0 - e_i$ and the S' -unit representation). The constraint $(x_0 - e_1) - (x_0 - e_2) = e_2 - e_1$ then gives a linear form in logarithms of the η_j that is bounded above by $e^{-c|x_0|}$ (exponentially small in $|x_0|$) and below by Baker's theorem: $\exp(-C_2 \log(\max |a_j|))$. Combining: $C_2 \log(\max |a_j|) > c|x_0|$, and since $\max |a_j| \leq C_1 \log |x_0|$: $C_2 \log(C_1 \log |x_0|) > c|x_0|$, which forces $|x_0| \leq X_0$ for a computable X_0 .

The LLL Reduction

The gap between the Baker bound ($|x| \leq 10^{500}$) and the actual integral points ($|x| \leq 100$, say) is bridged by the *LLL algorithm* (Lenstra–Lenstra–Lovász, 1982), a polynomial-time algorithm for finding short vectors in lattices.

The idea is as follows. The Baker bound gives: $|b_1 \log \alpha_1 + \cdots + b_n \log \alpha_n| < \epsilon$ with $|b_i| \leq B_0$ (the huge Baker bound). This says the integer vector $\mathbf{b} = (b_1, \dots, b_n)$ lies close to the kernel of the linear form $L(\mathbf{b}) = \sum b_i \log \alpha_i$. Consider the lattice Λ in $\mathbb{R}^{n+1}$ generated by the columns of

$$M = \begin{pmatrix} 1 & 0 & \cdots & 0 \\ 0 & 1 & \cdots & 0 \\ \vdots & & \ddots & \vdots \\ 0 & 0 & \cdots & 1 \\ \lfloor C \log \alpha_1 \rfloor & \lfloor C \log \alpha_2 \rfloor & \cdots & \lfloor C \log \alpha_n \rfloor \end{pmatrix}, \quad (10.12)$$

where C is a large scaling constant (chosen so that $C \cdot \epsilon$ is much smaller than 1). The lattice point $M\mathbf{b}$ has last coordinate $\approx C \cdot L(\mathbf{b})$, which is small (of order $C\epsilon$). So the lattice point corresponding to the solution $\mathbf{b}$ is unusually short: its length is $\approx \|\mathbf{b}\| + C\epsilon \ll B_0$.

The LLL algorithm finds a short basis for Λ in polynomial time. If the shortest vector found has last coordinate $> C\epsilon$: there is no lattice point with $|L(\mathbf{b})| < \epsilon$ and $\|\mathbf{b}\| \leq B_0$, and we can replace B_0 by a smaller bound. In practice, one iterates: compute the LLL-reduced basis, read off a new (much

smaller) bound B'_0 , and repeat if necessary. The typical outcome is dramatic: a single LLL reduction replaces $B_0 = 10^{500}$ with $B'_0 \leq 100$, reducing years of computation to seconds.

The reason for this dramatic reduction is that the Baker bound is doubly logarithmic in the “true” bound: Baker gives $B_0 \sim \exp(\exp(C))$, but the actual solutions have $B \sim \exp(C')$ for a much smaller C' . The LLL algorithm detects that the lattice has a short vector only when B is close to the true bound, not the Baker bound. In effect, LLL replaces the crude Baker estimate with the sharp geometric estimate coming from the lattice structure.

Example 10.2.5: The Baker–LLL Pipeline for $y^2 = x^3 - 2$

$E : y^2 = x^3 - 2$ over $\mathbb{Q}$. The known rational points are O and $(\pm)(3, 5)$ (rank 1, generator $P = (3, 5)$). We want to find all integral points.

Step 1: Baker bound. The splitting field is $L = \mathbb{Q}(\sqrt[3]{2}, \omega)$ (where $\omega = e^{2\pi i/3}$) with $[L : \mathbb{Q}] = 6$. The S -unit equation (from the factorization $y^2 = (x - \sqrt[3]{2})(x - \omega\sqrt[3]{2})(x - \omega^2\sqrt[3]{2})$ in $\mathcal{O}_L$) involves S -units for S containing the primes above 2 and 3. Baker's theorem applied to the resulting linear form in logarithms gives $|x| \leq X_0$ with $X_0 \approx 10^{18}$ (after applying the Matveev bounds; the constant depends on the fundamental units of L and the Baker constant $C(3, 6)$).

Step 2: LLL reduction. The linear form involves $\log |\eta_1|, \log |\eta_2|$ (two fundamental units in $\mathcal{O}_L^*$) and $\log |\alpha|$ (related to the S -units at primes above 2, 3). The lattice Λ from (10.12) is 4-dimensional. Running LLL: the shortest vector gives a reduced bound $|b_i| \leq 12$ (where the b_i are the exponents in the S -unit representation). This translates to $|x| \leq 100$ (approximately).

Step 3: Direct search. Enumerate $|x| \leq 100$ and check $x^3 - 2 \in (\mathbb{Z})^2$. The results: $x = 3$: $27 - 2 = 25 = 5^2$. No other values give a perfect square. The integral points are $(3, \pm 5)$.

The Elliptic Logarithm Method

A more refined approach, developed by Stroeker–Tzanakis and de Weger in the 1990s, replaces the algebraic-number logarithms with *elliptic logarithms*, giving sharper bounds and a more direct connection to the group structure of $E(\mathbb{Q})$.

The idea: every point $P \in E(\mathbb{R})$ has an *elliptic logarithm* $\psi(P) \in \mathbb{R}/\omega\mathbb{Z}$ (where ω is the real period of E), defined by the inverse of the Weierstrass $\wp$ -function (section 3.3). If $P = n_1P_1 + \cdots + n_rP_r + T$ (with P_i the Mordell–Weil generators and T torsion): $\psi(P) = n_1\psi(P_1) + \cdots + n_r\psi(P_r) + \psi(T)$. The integrality condition on P forces $\psi(P)$ to be close to 0 or $\omega/2$ (corresponding to the two real branches of $E(\mathbb{R})$), giving a linear form in the elliptic logarithms $\psi(P_i)$ that is close to zero.

Baker-type bounds for linear forms in elliptic logarithms (proved by David, 1995, and Hirata-Kohno, 2003) then give:

$$|n_1\psi(P_1) + \cdots + n_r\psi(P_r) + n_0\omega - \psi_0| > \exp(-C' \cdot \log N), \quad (10.13)$$

where $N = \max |n_i|$, ψ_0 is the logarithm of the torsion point, and C' depends on E and the P_i .

This approach has two advantages over the classical (algebraic logarithm) method: the linear form involves only $r + 1$ logarithms (one per generator plus the period), rather than the $|S|$ logarithms of fundamental S -units, and the constants are sharper because they exploit the geometry of E directly.

The LLL reduction step is identical, and in practice the elliptic logarithm method produces smaller initial bounds.

Example 10.2.6: Elliptic Logarithm Method on 37a1

$E : y^2 + y = x^3 - x$ over $\mathbb{Q}$. Rank 1, generator $P = (0, 0)$, real period $\omega \approx 5.987$. Elliptic logarithm: $\psi(P) \approx 2.994 \approx \omega/2$ (the generator is close to the half-period).

For an integral point $Q = [n]P$: $\psi(Q) = n \cdot \psi(P) \pmod{\omega}$. The integrality condition forces $\psi(Q)$ to be near 0 or $\omega/2$: $|n \cdot \psi(P) - m\omega| < \epsilon$ for some $m \in \mathbb{Z}$ and small ϵ (depending on $|x(Q)|$). This is a linear form in $\psi(P)$ and ω : $|n \cdot \psi(P) - m \cdot \omega| < \epsilon$.

Baker's bound for elliptic logarithms: $|n \cdot \psi(P) - m \cdot \omega| > \exp(-C'' \log N)$ where $N = \max(|n|, |m|)$. The integrality condition gives the upper bound $\epsilon \sim e^{-cn^2 \hat{h}(P)}$ (from the height growth of integral points, proposition 10.1.12). Combining: $\exp(-C'' \log N) < e^{-cn^2 \hat{h}(P)}$, giving $C'' \log N > cn^2 \hat{h}(P)$. Since $N \leq n$: $C'' \log n > cn^2 \hat{h}(P)$, which forces $n \leq n_0$ for an explicit n_0 . With $\hat{h}(P) \approx 0.0511$: the Baker bound gives $n_0 \approx 10^8$, and LLL reduces this to $n_0 \leq 10$. Since $|n| \leq 6$ gives all integral points (from example 10.1.13): the computation is confirmed.

The comparison between the two methods: the classical (algebraic logarithm) approach via S -unit equations works over any number field and does not require knowing the Mordell–Weil generators, but the Baker constants involve all $|S|$ fundamental S -units and can be very large. The elliptic logarithm approach requires knowing the generators $P_1, \dots, P_r$ (computed in Chapter 9), but the linear form has only $r + 1$ terms, giving sharper bounds. For rank-1 curves: the elliptic logarithm method is nearly always superior. For higher rank: the advantage persists but the LLL reduction is more expensive (the lattice dimension is $r + 1$).

Historical Remarks

Baker's original application (1968) was to the Thue equation $f(x, y) = m$ (where f is a binary form of degree ≥ 3): he showed that all solutions satisfy $\max(|x|, |y|) \leq C(f, m)$ for an explicit constant. The application to integral points on elliptic curves came shortly after (Baker and Coates, 1970). The LLL reduction technique was introduced by de Weger (1989) and has become standard. The elliptic logarithm method was developed by Stroeker and Tzanakis (1994) and by Gebel, Pethő, and Zimmer (1994), and is implemented in PARI/GP, Magma, and SageMath.

The state of the art: for any *specific* elliptic curve $E/\mathbb{Q}$, all integral points can be determined in practice. The bottleneck is not the theoretical bound (which is always computable) but the LLL reduction and the final search (which depend on the size of the reduced bound and the rank of E). For rank-1 curves of moderate conductor: the computation is routine (seconds on a modern computer). For rank-2 curves: the LLL reduction involves a higher-dimensional lattice, and the computation may take minutes. For rank ≥ 3 : the computation is more delicate but still feasible for curves of moderate conductor.

The following table summarizes the pipeline for a typical curve.

Stage	Output	Typical size
Baker's theorem	Initial bound X_0	10^{500}
LLL reduction (1st pass)	Reduced bound X_1	10^{20}
LLL reduction (2nd pass)	Final bound X_2	10^2
Sieved search	All integral points	≤ 20 points

Exercise 10.2.1

1. Let $\alpha = \sqrt{2}$ and $\beta = \sqrt{3}$. Show that $|a \log 2 + b \log 3| > 0$ for all $(a, b) \in \mathbb{Z}^2 \setminus \{(0, 0)\}$ (i.e., $\log 2$ and $\log 3$ are linearly independent over $\mathbb{Q}$). (*Hint*: $2^a = 3^{-b}$ forces $a = b = 0$ by unique factorization.)
2. For $S = \{2, 3, \infty\}$: the S -units are $\mathbb{Z}_S^* = \{\pm 2^a \cdot 3^b : a, b \in \mathbb{Z}\}$. Write the equation $u + v = 1$ with $u = \pm 2^a 3^b$ and $v = \pm 2^c 3^d$. Show that the Baker bound implies $\max(|a|, |b|, |c|, |d|) \leq C$ for an explicit constant C depending on the Baker constant for $n = 2$ and the heights of 2, 3.
3. Find all integral points on the Mordell equation $y^2 = x^3 + 7$ by searching $x = -2, -1, 0, 1, \dots, 20$ and checking which give $y \in \mathbb{Z}$.
4. Explain why the LLL algorithm applied to the lattice (10.12) detects solutions to $|L(\mathbf{b})| < \epsilon$ with $\|\mathbf{b}\| \leq B$: the lattice point $M\mathbf{b}$ has norm $\approx \|\mathbf{b}\| + C\epsilon$, and if $C\epsilon \ll 1$, this point is much shorter than a generic lattice point of norm $\approx B$.
5. For $E : y^2 = x^3 + 1$ over $\mathbb{Q}$ (rank 0): explain why the elliptic logarithm method is unnecessary. (*Hint*: with rank 0, all rational points are torsion, hence automatically of small height.)
6. Compare the Baker and Roth approaches to the equation $x^3 - 2y^3 = 1$ (a Thue equation). Roth's theorem gives finiteness (but no bound). Baker's theorem gives $\max(|x|, |y|) \leq C$ for an explicit C . Estimate C using the Baker bound for $n = 2$, $d = 3$, $\alpha_1 = \sqrt[3]{2}$, with the Matveev constant (10.8).

10.3 Computing Integral Points in Practice

The theoretical framework is in place: Siegel's theorem (section 10.1) guarantees finiteness, Baker's theorem (section 10.2) provides an effective bound, and the LLL algorithm reduces the bound to a manageable range. This section assembles the pieces into a working pipeline and applies it to determine the integral points on the running curves and on the classical Mordell equations $y^2 = x^3 + k$.

The Full Pipeline

Given an elliptic curve $E : y^2 = f(x)$ over $\mathbb{Q}$ (with $f \in \mathbb{Z}[x]$ of degree 3), the determination of all integral points proceeds as follows.

10.3.1 Algorithm: Computing All Integral Points **Input:** $E : y^2 = x^3 + Ax + B$ with $A, B \in \mathbb{Z}$.

Output: The complete list of integral points $(x, y) \in E(\mathbb{Z})$.

Step 1 (Mordell–Weil group). Compute $E(\mathbb{Q})_{\text{tors}}$ and generators $P_1, \dots, P_r$ of $E(\mathbb{Q})/E(\mathbb{Q})_{\text{tors}}$ (by the methods of Chapters 8–9).

Step 2 (Initial bound). If $r = 0$: the rational points are all torsion, so enumerate them and check integrality. Done.

If $r \geq 1$: apply Baker’s theorem (either the classical S -unit version or the elliptic logarithm version of section 10.2) to obtain an initial bound N_0 : every integral point $P = \sum n_i P_i + T$ has $\max |n_i| \leq N_0$.

Step 3 (LLL reduction). Apply the LLL algorithm to the lattice associated to the linear form in (elliptic) logarithms. Replace N_0 by the reduced bound $N_1 \ll N_0$.

Step 4 (Enumeration). Enumerate all $\mathbf{n} = (n_1, \dots, n_r) \in \mathbb{Z}^r$ with $\max |n_i| \leq N_1$ and all $T \in E(\mathbb{Q})_{\text{tors}}$. For each: compute $P = \sum n_i P_i + T$ and check whether $x(P), y(P) \in \mathbb{Z}$.

Step 5 (Output). The integral points found in Step 4 are *all* integral points on E .

The computational bottleneck is Step 4: the number of lattice points to check is $(2N_1 + 1)^r \cdot |E(\mathbb{Q})_{\text{tors}}|$. For $r = 1$ and $N_1 = 20$ (typical after LLL): 41 points to check. For $r = 2$ and $N_1 = 15$: 961 points. For $r = 3$ and $N_1 = 10$: 9,261 points. All are feasible on a modern computer.

Complete Determination: The Running Curves

We now carry out the full computation for each running example, recording the final answer and the key parameters.

Example 10.3.2: All Integral Points on the Running Curves

(a) $E : y^2 = x^3 - x$ over $\mathbb{Q}$. Rank 0, $E(\mathbb{Q}) = E[2] = \{O, (0, 0), (\pm 1, 0)\}$. All three finite points are integral. Integral points: $\{(0, 0), (1, 0), (-1, 0)\}$.

(b) $E : y^2 + y = x^3 - x$ over $\mathbb{Q}$ (37a1). Rank 1, generator $P = (0, 0)$, $\hat{h}(P) \approx 0.0511$, trivial torsion. The elliptic logarithm method gives initial bound $|n| \leq N_0 \approx 10^8$. LLL reduction: $|n| \leq N_1 = 7$. Enumeration: for $n = -7, \dots, 7$, compute $[n]P$ and check integrality.

n	$[n]P$	integral?
0	O	(infinity)
± 1	$(0, 0), (0, -1)$	yes
± 2	$(1, -1), (1, 0)$	yes
± 3	$(0, -1), (0, 0)$	yes (same as ± 1)
± 4	$(-1, 1), (-1, 0)$	yes
± 5	$(2, -3), (2, 2)$	yes
± 6	$(6, -14), (6, 13)$	yes
± 7	$(-5/9, 8/27), (-5/9, -35/27)$	no

The distinct integral points (listing each x -value with both y -values from $y \mapsto -1 - y$):

$$(x, y) \in \{(-1, 0), (-1, 1), (0, -1), (0, 0), (1, -1), (1, 0), (2, -3), (2, 2), (6, -14), (6, 13)\}.$$

Total: 10 integral points. Note that $[3]P = (0, -1) = [-1]P$ on this model (the involution $y \mapsto -1 - y$ sends $(0, 0)$ to $(0, -1)$), and $[-n]P$ is the image of $[n]P$ under this involution. So the x -values 0, 1, -1, 2, 6 each contribute two integral points.

(c) $E : y^2 = x^3 + 1$ over $\mathbb{Q}$ (**36a1**). Rank 0, $E(\mathbb{Q}) \cong \mathbb{Z}/6\mathbb{Z}$. All rational points are integral (as computed in example 10.1.2(c)): $(-1, 0), (0, \pm 1), (2, \pm 3)$. Total: 5 integral points.

(d) $E_5 : y^2 = x^3 - 25x$ over $\mathbb{Q}$. Rank 2, generators $P_1 = (-4, 6)$ and $P_2 = (25/4, 75/8)$. Note: P_2 is not itself integral, but linear combinations $n_1P_1 + n_2P_2 + T$ may be. The torsion is $E_5(\mathbb{Q})[2] = \{O, (0, 0), (5, 0), (-5, 0)\}$, contributing three integral torsion points.

The elliptic logarithm method with $r = 2$ gives initial bound $\max(|n_1|, |n_2|) \leq N_0 \approx 10^{12}$. LLL reduction (in a 3-dimensional lattice): $\max(|n_1|, |n_2|) \leq 10$. Enumeration: check all $(n_1, n_2) \in [-10, 10]^2$ and all $T \in E_5[2]$: $441 \cdot 4 = 1,764$ points. The integral points found:

$$(0, 0), (\pm 5, 0), (-4, \pm 6), (45, \pm 300).$$

Total: 7 integral points (3 torsion + 4 non-torsion). The point $(45, 300)$: $45^3 - 25 \cdot 45 = 91125 - 1125 = 90000 = 300^2$.

The Mordell Equation $y^2 = x^3 + k$

The family of curves $E_k : y^2 = x^3 + k$ (for $k \in \mathbb{Z} \setminus \{0\}$) is a classical testing ground for Diophantine methods. These are the *Mordell equations*, studied since Bachet (1621) and Mordell (1920). The integral points on E_k have been determined for all $|k| \leq 10^7$ (by Gebel, Pethő, and Zimmer, 1998, using the elliptic logarithm method), and the results reveal striking patterns.

The Mordell equation is particularly well-suited to the methods of this chapter because the curve $y^2 = x^3 + k$ has $j = 0$ and the endomorphism ring (over $\overline{\mathbb{Q}}$) is $\mathbb{Z}[\omega]$ with $\omega = e^{2\pi i/3}$, giving extra symmetry that simplifies the arithmetic. The factorization in $\mathbb{Z}[\omega]$:

$$y^2 = x^3 + k = (x + \sqrt[3]{k})(x + \omega \sqrt[3]{k})(x + \omega^2 \sqrt[3]{k}) \quad (10.14)$$

in the ring $\mathbb{Z}[\omega, \sqrt[3]{k}]$ reduces the integral point problem to a unit equation in this ring, and the Baker–LLL pipeline applies directly.

Example 10.3.3: The Mordell Equation for Small k

$k = 1$: $y^2 = x^3 + 1$. Integral points: $(-1, 0)$, $(0, \pm 1)$, $(2, \pm 3)$. Total: 5 (all torsion, since rank = 0).

$k = -1$: $y^2 = x^3 - 1 = (x-1)(x^2+x+1)$. Since $x^2+x+1 > 0$ for all $x \in \mathbb{R}$ and $\gcd(x-1, x^2+x+1) \mid 3$: the factor $x-1$ must be ± 1 or ± 3 times a square (after accounting for the gcd). Checking: $x = 1$: $y^2 = 0$. $x = -1$: $y^2 = -2 < 0$. $x = 2$: $y^2 = 7$, not a square. No further integral solutions for larger $|x|$ (confirmed by Baker’s method). Integral points: $(1, 0)$. Total: 1.

$k = -2$: $y^2 = x^3 - 2$. The classical Mordell equation. Rank 1, generator $(3, 5)$. Baker–LLL gives $|n| \leq 3$. Checking: $[1]P = (3, 5)$: integral. $[2]P$: has non-integer coordinates. $[-1]P = (3, -5)$: integral. Integral points: $(3, \pm 5)$. Total: 2.

$k = 2$: $y^2 = x^3 + 2$. Rank 1, generator $(-1, 1)$. Checking multiples: $(-1, \pm 1)$: integral. $[2](-1, 1)$: computation gives a non-integer point. Integral points: $(-1, \pm 1)$. Total: 2.

$k = 17$: $y^2 = x^3 + 17$. This equation famously has the “large” solution $(x, y) = (52, 758, 958, 383, 344, 129)$ — a point of enormous height. The rank is 2, and the Baker–LLL method (with two generators) was needed to verify completeness. The integral points include $(-1, \pm 4)$, $(-2, \pm 3)$, $(2, \pm 5)$, $(4, \pm 9)$, $(8, \pm 23)$, $(43, \pm 282)$, $(52, \pm 375)$, and several others (a total of 20 integral points). This example illustrates that the Mordell equation can have many integral points and generators of very large height.

The structure of the integral point set depends sensitively on k . For $k > 0$: the curve E_k always passes through the integral point $(-\lfloor k^{1/3} \rfloor, \pm y)$ (approximately), so integral points exist for most positive k . For $k < 0$: the situation is more varied, and many values of k have no integral points at all ($E_k(\mathbb{Z}) = \emptyset$ for $k = -3, -4, -5, -6$, for instance).

A remarkable pattern: the number of integral points on $y^2 = x^3 + k$ can be arbitrarily large. Elkies showed that for suitable k : $|E_k(\mathbb{Z})| \geq c \cdot |k|^\epsilon$ for certain families. Conversely, the “average” number of integral points (over all $|k| \leq N$) is bounded: a theorem of Helfgott and Venkatesh (2006) shows the average is $O(N^\epsilon)$.

The following table, compiled from the computations of Gebel–Pethő–Zimmer, records the number of integral points for representative values of k :

k	rank	$ E_k(\mathbb{Z}) $	largest $ x $
1	0	5	2
-1	0	1	1
-2	1	2	3
2	1	2	-1
17	2	20	52,758,958
-27	0	3	3
100	1	4	15
-100	1	2	5

The table illustrates the enormous variation: $k = 17$ has 20 integral points (with the largest at $x \approx 5 \times 10^7$), while $k = -1$ has only 1. The rank is the primary determinant of the number of integral points — rank-0 curves have at most a handful (all torsion), while higher-rank curves can have many — but the height of the generators also plays a role: large generators produce large integral points.

Integral Points and the Conductor

A natural question connects integral points to the global invariants of E : is the number of integral points bounded in terms of the conductor N ?

Proposition 10.3.4: Integral Points vs. Conductor

For $E/\mathbb{Q}$ in minimal Weierstrass form: the number of integral points satisfies $|E(\mathbb{Z})| \leq C(N)$ for a function C depending only on the conductor N . Conditional on the ABC conjecture (section 10.4): $|E(\mathbb{Z})| \leq C_\epsilon \cdot N^{1+\epsilon}$ for every $\epsilon > 0$.

The idea: the ABC conjecture bounds the minimal discriminant by $|\Delta_{\min}| \leq C_\epsilon N^{6+\epsilon}$, and the integral points are constrained by Δ (through the S -unit equation, where $|S|$ is determined by the primes dividing Δ). The precise dependence on N is the subject of active research.

Beyond Weierstrass Models

The integral point computation depends on the choice of affine model: a different equation for the same elliptic curve can have a different set of integral points. For instance, the curve $y^2 = x^3 - x$ has integral points $\{(0, 0), (\pm 1, 0)\}$ in the Weierstrass model, but the isomorphic model $Y^2 = X^3 + 4X$ (related by $(X, Y) = (x^2/y^2 \cdot \dots)$, the image under the 2-isogeny) may have a different set of integral points.

The “correct” notion for studying integral points intrinsically is the *quasi-minimal model*: a Weierstrass equation that minimizes the discriminant (or equivalently, maximizes the set of primes of good reduction). For a minimal Weierstrass equation: the integral points are determined by the equation, and the Baker–LLL pipeline applies directly. For non-minimal models: one first transforms to the minimal model, computes integral points there, and transforms back.

Over number fields: the situation is more complex because the ring of integers $\mathcal{O}_K$ may not be a PID, and different choices of affine model (corresponding to different elements of the Weierstrass

class group) give different integral point sets. The S -integral version (theorem 10.1.4) handles this uniformly.

Implementation and Software

The Baker–LLL pipeline for computing integral points is implemented in several computer algebra systems.

In *SageMath*: the command `E.integral_points()` applies the elliptic logarithm method and returns the complete list. For the running curves:

```
sage: E = EllipticCurve([0, 0, 1, -1, 0]) # 37a1
sage: E.integral_points()
[(-1, 0), (0, -1), (0, 0), (1, -1), (1, 0),
 (2, -3), (2, 2), (6, -14), (6, 13)]
```

(returning 9 points; the tenth, $(-1, 1)$, is the involution of $(-1, 0)$ and may or may not be listed depending on conventions).

In *Magma*: the command `IntegralPoints(E)` provides the same functionality, typically with faster running times for curves of large conductor or high rank.

The computation time depends primarily on the rank r (which determines the dimension of the LLL lattice) and on the size of the reduced bound N_1 . For the curves in the LMFDB of conductor $\leq 500,000$: the integral points are computed in the database for every curve of rank ≤ 2 . For higher rank: the computation is feasible but more expensive.

Exercise 10.3.1

1. Find all integral points on $y^2 = x^3 + 3$ by searching $|x| \leq 20$. (*Answer*: $(1, \pm 2)$ are the only integral points.)
2. Show that $y^2 = x^3 - 6$ has no integral points for $|x| \leq 20$.
3. The congruent number curve $E_6 : y^2 = x^3 - 36x$ has rank 1 with generator $P = (-2, 8)$ (verify: $(-2)^3 - 36(-2) = -8 + 72 = 64 = 8^2$). Find all integral points by computing $[n]P$ for $|n| \leq 10$ and checking integrality.
4. For $E : y^2 = x^3 - 2$ with $S = \{\infty, 2, 3\}$: find the S -integral points (allowing denominators that are products of powers of 2 and 3) with $h(x) \leq 5$. Are there more S -integral points than integral points?
5. The equation $y^2 = x^3 + 17$ has the large integral point $(52758958, 383344129)$. Verify that this is indeed a point: compute $52758958^3 + 17$ and check it equals 383344129^2 .
6. For a rank-0 curve $E/\mathbb{Q}$: explain why the integral point computation requires no Baker bound — the integral points are a subset of the (finitely many) rational points, which are all torsion and can be enumerated directly.

10.4 The ABC Conjecture and Szpiro's Conjecture

The finiteness results of this chapter — Siegel's theorem, Baker's effective bounds, Shafarevich's theorem — share a common structure: they control the “size” of arithmetic objects (integral points, discriminants, Weierstrass models) in terms of the “complexity” of the underlying data (the curve E , the set of bad primes S , the conductor N). The ABC conjecture, proposed independently by Oesterlé and Masser in 1985, gives a unified framework for these bounds and predicts sharp quantitative relationships that would resolve many open problems in Diophantine geometry.

This section presents the ABC conjecture for integers, its translation to elliptic curves (Szpiro's conjecture), and the cascade of consequences for the arithmetic of elliptic curves: effective Mordell, effective Siegel, uniform bounds on torsion, and bounds on integral points. The section serves as the capstone of Part II: the arithmetic tools developed in Chapters 6–10 are placed in the context of the deepest open problems in the subject.

The ABC Conjecture

The conjecture concerns the simplest Diophantine equation: $a + b = c$.

Definition 10.4.1: The Radical

For a nonzero integer n : the *radical* of n is the product of its distinct prime factors:

$$\text{rad}(n) = \prod_{p|n} p. \quad (10.15)$$

For example: $\text{rad}(12) = \text{rad}(2^2 \cdot 3) = 6$, $\text{rad}(100) = \text{rad}(2^2 \cdot 5^2) = 10$, $\text{rad}(p^k) = p$.

The radical strips away the multiplicity: $\text{rad}(n) \leq |n|$ always, and $\text{rad}(n) = |n|$ iff n is squarefree. The ABC conjecture says that when $a + b = c$ with $\gcd(a, b, c) = 1$: the number c cannot be much larger than $\text{rad}(abc)$. In other words, if $a + b = c$ and the prime factors of a, b, c are “small” (relative to c), then there is a lot of cancellation in the prime factorizations — and the ABC conjecture says this cancellation is limited.

Conjecture 10.4.2: The ABC Conjecture (Oesterlé–Masser, 1985)

For every $\epsilon > 0$: there exists a constant $C_\epsilon > 0$ such that for all coprime positive integers a, b, c with $a + b = c$:

$$c \leq C_\epsilon \cdot \text{rad}(abc)^{1+\epsilon}. \quad (10.16)$$

The exponent $1 + \epsilon$ is conjectured to be sharp: for every $\epsilon > 0$, the inequality (10.16) holds with finitely many exceptions, but the inequality $c \leq C \cdot \text{rad}(abc)$ (with $\epsilon = 0$) fails — there exist triples with $c/\text{rad}(abc)$ arbitrarily large. The “quality” of an ABC triple is $q = \log c / \log \text{rad}(abc)$: the conjecture says $q < 1 + \epsilon$ for all but finitely many triples. The best known examples have $q \approx 1.63$ (the triple $2 + 3^{10} \cdot 109 = 2^{35} \cdot 5$ found by Reyssat).

Example 10.4.3: ABC Triples

(a) $1 + 8 = 9$: $a = 1, b = 2^3, c = 3^2$. $\text{rad}(1 \cdot 8 \cdot 9) = \text{rad}(72) = 6$. Quality: $q = \log 9 / \log 6 = 1.23$. A “good” triple: $c = 9 > 6 = \text{rad}(abc)$.

(b) $5 + 27 = 32$: $\text{rad}(5 \cdot 27 \cdot 32) = \text{rad}(5 \cdot 3^3 \cdot 2^5) = 30$. Quality: $q = \log 32 / \log 30 = 1.02$. A marginal triple.

(c) $2 + 3^{10} \cdot 109 = 2^{23} \cdot 5$ (the Reyssat triple): $a = 2, b = 6436341 = 3^{10} \cdot 109, c = 6436343$. $\text{rad}(abc) = \text{rad}(2 \cdot 3^{10} \cdot 109 \cdot 6436343) = 2 \cdot 3 \cdot 109 \cdot \text{rad}(6436343)$. Since $6436343 = 6436343$ (one checks this is $= 2^{23} \cdot 5 / \dots$): this triple has quality $q \approx 1.63$, the highest known as of this writing. High-quality ABC triples are extremely rare: among all triples with $c \leq 10^{18}$, only a handful have $q > 1.4$.

(d) The triple $1 + 2^n - 1 = 2^n$ (for n large): $\text{rad}(1 \cdot (2^n - 1) \cdot 2^n) = 2 \cdot \text{rad}(2^n - 1)$. If $2^n - 1$ is a Mersenne prime: $\text{rad}(abc) = 2(2^n - 1) \approx 2^{n+1}$, and $c = 2^n$, giving $q \approx 1$. Mersenne primes give quality barely exceeding 1: they are not “good” ABC triples.

Szpiro's Conjecture

The ABC conjecture for integers translates, via the Weierstrass equation, to a conjecture about elliptic curves. The translation was discovered by Szpiro (1981, predating the ABC conjecture itself) and concerns the relationship between the minimal discriminant $\Delta_{\min}$ and the conductor N of an elliptic curve.

Recall: the conductor $N = \prod_p p^{f_p}$ encodes the primes of bad reduction (with $f_p \geq 1$ for bad primes and $f_p = 0$ for good primes), and the minimal discriminant $\Delta_{\min}$ satisfies $v_p(\Delta_{\min}) \geq 0$ for all p (with $v_p(\Delta_{\min}) > 0$ iff p is a prime of bad reduction). The conductor “knows which primes are bad” but loses the multiplicity information; the discriminant retains the multiplicities. Szpiro's conjecture says the discriminant cannot be much larger than the conductor:

Conjecture 10.4.4: Szpiro's Conjecture

For every $\epsilon > 0$: there exists a constant $C_\epsilon > 0$ such that for every elliptic curve $E/\mathbb{Q}$:

$$|\Delta_{\min}(E)| \leq C_\epsilon \cdot N(E)^{6+\epsilon}. \quad (10.17)$$

The exponent 6 is sharp: the Frey–Hellegouarch curves $E_{a,b,c} : y^2 = x(x-a)(x+b)$ (associated to an ABC triple $a + b = c$) have $\Delta_{\min} \sim (abc)^2$ and $N \sim \text{rad}(abc)$, giving $|\Delta_{\min}|/N^6 \sim (abc)^2/\text{rad}(abc)^6$. The ABC conjecture with exponent $1 + \epsilon$ translates to $|\Delta_{\min}| \leq C_\epsilon N^{6+\epsilon}$, confirming Szpiro's exponent $6 + \epsilon$.

Proposition 10.4.5: Equivalence of ABC and Szpiro

The ABC conjecture (10.4.2) and Szpiro's conjecture (10.4.4) are equivalent, in the following sense: each implies the other (with possibly different values of ϵ and C_ϵ).

Proof Sketch:

(*ABC* $\Rightarrow$ *Szpiro*.) Given an elliptic curve $E/\mathbb{Q}$ in minimal Weierstrass form: the discriminant $\Delta_{\min}$ and conductor N are related by $N = \prod_p |\Delta| p^{f_p}$ with $f_p \leq 2 + 3v_p(3)$ (for additive reduction) or $f_p = 1$ (for multiplicative reduction). The key: $\text{rad}(\Delta_{\min}) | N^{O(1)}$ (the radical of the discriminant is controlled by the conductor). The ABC conjecture applied to a suitable identity involving c_4^3 and $\Delta_{\min}$ (the relation $c_4^3 - c_6^2 = 1728\Delta$) gives $|\Delta_{\min}| \leq C_\epsilon \text{rad}(\Delta_{\min})^{6+\epsilon} \leq C'_\epsilon N^{6+\epsilon'}$.

(*Szpiro* $\Rightarrow$ *ABC*.) Given a coprime triple $a + b = c$: the Frey curve $E_{a,b} : y^2 = x(x-a)(x+b)$ has $\Delta_{\min} = 2^{-8}(abc)^2$ (up to powers of 2) and $N | 2 \cdot \text{rad}(abc)$. Szpiro gives $(abc)^2 \leq C_\epsilon (2\text{rad}(abc))^{6+\epsilon}$, which rearranges to $c \leq C'_\epsilon \text{rad}(abc)^{3+\epsilon'}$ — weaker than the full ABC, but the argument can be sharpened using the modularity theorem and the precise relationship between Δ and N for semistable curves. The full equivalence (with optimal exponents) is proved by Szpiro and Oesterlé; see Goldfeld [9] for the details.

Hall's Conjecture

A closely related conjecture, predating ABC, concerns the gap between perfect cubes and perfect squares.

Conjecture 10.4.6: Hall's Conjecture

For every $\epsilon > 0$: there exists $C_\epsilon > 0$ such that for all integers x, y with $x^3 \neq y^2$:

$$|x^3 - y^2| > C_\epsilon \cdot |x|^{1/2-\epsilon}. \quad (10.18)$$

Hall's conjecture says: cubes and squares cannot be too close together. If $x^3 \approx y^2$ (i.e., $|x^3 - y^2|$ is small relative to $|x|^{3/2}$), then we are looking at a near-integral point on $E_k : Y^2 = X^3 - k$ for $k = y^2 - x^3$. Hall's conjecture is a consequence of the ABC conjecture: if $a = x^3$, $b = -y^2$, $c = x^3 - y^2$: $\text{rad}(x^3 y^2 c) \leq |xyc|$, and ABC gives $|x^3| \leq C_\epsilon |xyc|^{1+\epsilon}$, which (after estimating $y \approx x^{3/2}$ and rearranging) gives $|c| \geq C'_\epsilon |x|^{1/2-\epsilon'}$.

The best known examples of near-misses $|x^3 - y^2| \ll |x|^{1/2}$ come from the theory of continued fractions of cube roots, and the current record is $|x^3 - y^2| \approx 0.0016 \cdot |x|^{0.47}$ for $x = 5853886516781223$ (found by Elkies, 1998). Finding such examples is a major computational challenge.

Consequences for Elliptic Curves

The ABC/Szpiro conjectures, if true, would resolve several fundamental open problems.

Effective Mordell conjecture. Faltings' theorem (1983) says: a curve C/K of genus ≥ 2 has finitely many rational points. The proof is ineffective (it relies on the Shafarevich conjecture, which itself was proved ineffectively via the Faltings height). The ABC conjecture would make it effective: given C/K , one could compute a bound B such that all rational points $P \in C(K)$ have $h(P) \leq B$. The key step: the ABC conjecture bounds the discriminant of the Jacobian of C , which controls the heights of rational points through the Faltings height and the Vojta inequality. An effective Mordell conjecture would have profound consequences: it would, for instance, give an algorithm to determine whether $x^n + y^n = z^n$ has solutions for a given n (by bounding the size of potential solutions).

Effective Siegel's theorem. Siegel's theorem (proven ineffectively via Roth's theorem in section 10.1) would become effective without Baker's method: given $E/\mathbb{Q}$ and a Weierstrass equation, the ABC conjecture gives an explicit bound $|x| \leq C_\epsilon N^{C+\epsilon}$ for all integral points, where N is the conductor. This bound is polynomial in N — vastly smaller than Baker's doubly-exponential bounds — and would make the integral point computation trivial for curves of any conductor.

Uniform bounds on torsion. Mazur's theorem (theorem 8.5.5) gives the list of possible torsion groups over $\mathbb{Q}$, and Merel's theorem (1996) gives a bound $|E(K)_{\text{tors}}| \leq C([K : \mathbb{Q}])$ over number fields. The ABC conjecture gives an alternative approach yielding a sharper bound: $|E(K)_{\text{tors}}| \leq C_\epsilon N(E)^\epsilon$ for every $\epsilon > 0$. The mechanism: a torsion point of large order forces the curve to have a large isogeny, which in turn forces the discriminant to be large relative to the conductor, violating Szpiro.

Bounds on the number of integral points. As discussed in proposition 10.3.4: the ABC conjecture implies $|E(\mathbb{Z})| \leq C_\epsilon N^{1+\epsilon}$ for the number of integral points on a minimal Weierstrass model. This is much sharper than the unconditional Evertse bound (theorem 10.1.8), and illustrates the organizing power of the ABC framework: a single conjecture about $a + b = c$ controls integral points, torsion, discriminants, and conductors simultaneously.

The Status of ABC

The ABC conjecture remains open as of this writing. The most significant development is the claimed proof by Mochizuki (2012), using his “inter-universal Teichmüller theory” (IUT). The proof has been scrutinized intensely by the mathematical community. A specific error was identified by Scholze and Stix (2018) in a key step (Corollary 3.12 of Mochizuki's third paper), and this error has not been resolved to the satisfaction of the broader community, despite Mochizuki's detailed responses. The consensus among most number theorists, as of the mid-2020s, is that the claimed proof has not been verified.

Independent of the Mochizuki controversy: partial results toward ABC are known. Baker (1998) proved an explicit (but weaker) version: $c \leq (\text{rad}(abc))^{15/14}$ up to logarithmic factors, conditional on certain uniformity assumptions. Stewart and Yu (2001) proved unconditionally that $\log c \leq C \cdot (\text{rad}(abc))^{1/3} (\log \text{rad}(abc))^3$, which gives c at most exponential in $\text{rad}(abc)^{1/3}$ — far from the conjectured polynomial bound, but the best unconditional result.

For Szpiro's conjecture specifically: the Frey–Ribet–Wiles proof of Fermat's Last Theorem can be viewed as a proof of the “ $\epsilon = \infty$ ” case of ABC (it shows there are no triples with $\text{rad}(abc) = 1$, corresponding to $a^n + b^n = c^n$). The full conjecture remains wide open.

The View from Part III

The ABC conjecture stands at the intersection of the arithmetic theory developed in this Part (Chapters 6–10) and the deeper structural theory of Part III (Chapters 11–16). The Frey curve construction — which produces an elliptic curve from an ABC triple — links the elementary Diophantine equation $a + b = c$ to the Galois representations and modularity theory of Part III. The proof of Fermat's Last Theorem (Wiles, 1995) follows this link: the Frey curve $y^2 = x(x - a)(x + b)$ arising from a hypothetical Fermat triple $a^p + b^p = c^p$ is shown (by Ribet) to violate the modularity theorem, yielding a contradiction.

This connection — from elementary number theory through elliptic curves to Galois representations and modular forms — is the deepest theme of the subject. Part III develops the tools (Néron models,

Galois representations, modular forms, the BSD conjecture) that make this connection precise.

Example 10.4.7: The Frey Curve and Fermat's Last Theorem

Suppose $a^p + b^p = c^p$ for an odd prime $p \geq 5$ and coprime positive integers a, b, c . The Frey curve is $E_{a,b} : y^2 = x(x - a^p)(x + b^p)$. Its discriminant: $\Delta = 2^{-8}(abc)^{2p}$. Its conductor: $N = \text{rad}(a^p b^p c^p) = \text{rad}(abc)$ (since the curve is semistable). Szpiro's conjecture with exponent $6 + \epsilon$ gives $(abc)^{2p} \leq C_\epsilon \text{rad}(abc)^{6+\epsilon}$. Since $\text{rad}(abc) \leq abc \leq c^3$: $(abc)^{2p} \leq C_\epsilon c^{3(6+\epsilon)}$, i.e., $c^{2p} \leq C_\epsilon c^{18+3\epsilon}$. For $p \geq 10$ and ϵ small: $2p > 18 + 3\epsilon$, giving $c \leq C_\epsilon$: only finitely many solutions. For $p \geq 5$: a more careful analysis (using the modularity theorem and level-lowering, as Ribet showed) gives a contradiction outright: no solutions exist. This is the Wiles–Taylor–Wiles proof of Fermat's Last Theorem (1995).

Exercise 10.4.1

1. Compute the quality $q = \log c / \log \text{rad}(abc)$ for the triples $(a, b, c) = (1, 8, 9)$, $(1, 63, 64)$, $(1, 80, 81)$. Which has the highest quality?
2. Show that a Fermat triple $a^n + b^n = c^n$ (with $n \geq 3$, $\gcd(a, b, c) = 1$) would give an ABC triple of quality $q \geq n/3$. (*Hint*: $\text{rad}(a^n b^n c^n) = \text{rad}(abc) \leq abc \leq c^3$ for large c .)
3. For the Frey curve $E_{a,b} : y^2 = x(x - a)(x + b)$ associated to $a + b = c$ (coprime positive integers): compute $\Delta_{\min}$ and N in terms of a, b, c and $\text{rad}(abc)$. Verify that Szpiro's conjecture with exponent $6 + \epsilon$ implies the ABC conjecture with exponent $1 + \epsilon'$ (possibly with a different ϵ').
4. Find an integer pair (x, y) with $|x^3 - y^2|$ small relative to $|x|^{1/2}$. (*Hint*: try $x = 2$: $8 - y^2$. $y = 3$: $|8 - 9| = 1$, $|x|^{1/2} = \sqrt{2}$, ratio $= 1/\sqrt{2} \approx 0.71$. Try larger x .)
5. Explain how the ABC conjecture makes Siegel's theorem effective: if $y^2 = x^3 + Ax + B$ has an integral point (x_0, y_0) with $|x_0|$ large, then $y_0^2 \approx x_0^3$ gives a near-miss for Hall's conjecture, and ABC bounds $|x_0|$ from above.
6. Sketch how Szpiro's conjecture implies a bound on torsion: if $E/\mathbb{Q}$ has a rational point P of order n , then the isogeny $E \rightarrow E/\langle P \rangle$ produces a curve with discriminant related to n , and Szpiro bounds this discriminant by $N^{6+\epsilon}$, giving $n \leq C_\epsilon N^{C+\epsilon}$.

Part III

Advanced Topics: Representations, Modularity, and Beyond

here

Index

- 1-boundary, 238
- 1-cycle, 238
- 2-Selmer Group, 294

- Absolute Logarithmic Height, 266
- Admissible Change of Variable , 8

- b-invariants, 10

- c-invariants, 12
- Canonical Divisor, xiv
- Cartier Dual, 257
- Component Group, 215
- curve, ii
- Cusp, 46

- Degree, 55
 - Inseparable, 55
 - Separable, 55
- Differential Forms On A Curve, x
- Discriminant, 15
- Division Polynomials, 72
- Divisor of Differential, xiv
- Dual Isogeny, 63

- Effective Divisor, xv
- Eisenstein Series, 105
- Eisenstein Series on $\mathbb{H}$, 116
- Elliptic Curve, 5
- Elliptic Function, 99
- Elliptic Regulator, 278
- Endomorphism Ring
 - Analytic, 124

- First Cohomology, 238
- Formal Exponential, 155
- Formal Group Height, 160
- formal Group Law, 138
- Formal Group Law Homomorphism, 141
- Formal Logarithm, 153
- Frobenius Morphism, 79

- G-module, 238
- Galois Cohomology, 242
- Global Tamagawa Product, 234
- Group Law (for elliptic curves), 26

- Height
 - Absolute Logarithmic , 266
 - Formal Group, 160
 - Naive, 270
 - Number Field, 266
- Homogeneous Space associated to pair, 294
- Hurwitz Theorem, xviii

- Integral Weierstrass Model, 208
- Invariant Differential, 33
- Isogeny, 54

- j-Function, 119
- j-invariant, 18

- Kernel of Reduction, 167

- l-adic Weil Pairing, 88

L-Function over Finite Field, 188
Lattice, 98
 Standard Lattice, 98
Local Euler Factor, 233

Minimal Discriminant, 209
Minimal Weierstrass Equation, 208
modular Discriminant, 118
Modular Group, 130
Mordell-Weil Rank, 289

Néron-Tate Pairing, 277
Node, 46

p-Divisible Group, 163
Period Ratio, 98
Potential Good Reduction, 225
Principal Homogeneous Space, 248
Pushforward
 of Divisors, 25

Quadratic Twist, 41

Ramification Index, iv
Reduction Types, 212
Riemann-Roch Space For Divisor, xv
Riemann-Roch Theorem, xvi

Saturation Index, 315
Selmer Group, 250
 explicit 2-Selmer Group, 294
Supersingular Isogeny Graph, 204

Tate Module, 85
Tate-Shafarevich Group, 250
Trace of an Isogeny, 69
Twist, 40

Weierstrass Equation, 7
Weierstrass $\wp$ -equation, 102
Weil Pairing
 Analytic, 127

Zero Cohomology, 238
Zeta Function, 185

Bibliography

- [1] Nathanael Chwojko-Srawley. *Everything You Need to know about K-theory*. Vol. 1. URL: https://nathanaelsrawley.com/assets/pdfs/notes/EYNTKA_K-Theory.pdf.
- [2] J. W. S. Cassels. “Lectures on elliptic curves”. In: (1991). URL: <https://api.semanticscholar.org/CorpusID:118250338>.
- [3] Nathanael Chwojko-Srawley. *Everything You Need To Know About Algebraic Geometry*. Vol. 1. 2023. URL: https://nathanaelsrawley.com/assets/pdfs/notes/EYNTKA_algGeo.pdf.
- [4] Nathanael Chwojko-Srawley. *Everything You Need To Know About Complex Analysis*. 1st ed. 2023. 233 pp. URL: https://nathanaelsrawley.com/assets/pdfs/notes/EYNTKA_complex_analysis.pdf.
- [5] Nathanael Chwojko-Srawley. *Everything You Need To Know About Core Algebra*. Vol. 1. 2023. URL: https://nathanaelsrawley.com/assets/pdfs/notes/EYNTKA_algGeo.pdf.
- [6] Nathanael Chwojko-Srawley. *Everything You Need To Know About Galois Cohomology*. Vol. 1. 2023. URL: https://nathanaelsrawley.com/assets/pdfs/notes/EYNTKA_algGeo.pdf.
- [7] Nathanael Chwojko-Srawley. *Everything You Need To Know About Geometric Representation Theory*. Vol. 1. 2023. URL: https://nathanaelsrawley.com/assets/pdfs/notes/EYNTKA_algGeo.pdf.
- [8] Nathanael Chwojko-Srawley. *Everything You Need To Know About Undergraduate Algebra*. 1st ed. 2023. 1210 pp. URL: https://nathanaelsrawley.com/assets/pdfs/notes/EYNTKA_algebra.pdf.
- [9] Dorian Goldfeld. “Beyond the last theorem”. In: *The Sciences* 36.2 (1996), pp. 34–40.
- [10] Michiel Hazewinkel. *Formal Groups and Applications*. Vol. 78. Pure and Applied Mathematics. Academic Press, 1978.
- [11] Taira Honda. “On the theory of commutative formal groups”. In: *Journal of the Mathematical Society of Japan* 22.2 (1970), pp. 213–246.
- [12] Michel Lazard. “Sur les groupes de Lie formels à un paramètre”. In: *Bulletin de la Société Mathématique de France* 83 (1955), pp. 251–274.

- [13] James S. Milne. *Arithmetic Duality Theorems*. 2nd. BookSurge, 2006.
- [14] Louis J. Mordell. “On the rational solutions of the indeterminate equations of the third and fourth degrees”. In: *Proceedings of the Cambridge Philosophical Society* 21 (1922), pp. 179–192.
- [15] Jürgen Neukirch, Alexander Schmidt, and Kay Wingberg. *Cohomology of Number Fields*. 2nd. Vol. 323. Grundlehren der mathematischen Wissenschaften. Springer, 2008.
- [16] Arnold K. Pizer. “Ramanujan graphs and Hecke operators”. In: *Bulletin of the American Mathematical Society* 23.1 (1990), pp. 127–137.
- [17] J. Stilwell and J.P. Serre. *Trees*. Springer monographs in mathematics. Springer Berlin Heidelberg, 2002. ISBN: 978-3-540-44237-0. URL: <https://books.google.ca/books?id=MOAqeoY1BMQC>.
- [18] René Schoof. “Elliptic curves over finite fields and the computation of square roots mod p ”. In: *Mathematics of Computation* 44.170 (1985), pp. 483–494.
- [19] Jean-Pierre Serre and John Tate. “Good reduction of abelian varieties”. In: *Annals of Mathematics* 88.3 (1968), pp. 492–517.
- [20] Joseph H. Silverman. *Advanced Topics in the Arithmetic of Elliptic Curves*. Vol. 151. Graduate Texts in Mathematics. Springer, 1994.
- [21] Joseph H. Silverman. *The Arithmetic of Elliptic Curves*. Springer Science & Business Media, 2009. 525 pp. ISBN: 978-0-387-09494-6.
- [22] John Tate. “Algorithm for determining the type of a singular fiber in an elliptic pencil”. In: *Modular Functions of One Variable IV*. Ed. by B. J. Birch and W. Kuyk. Vol. 476. Lecture Notes in Mathematics. Springer, 1975, pp. 33–52.
- [23] John Tate. “Endomorphisms of Abelian Varieties over Finite Fields”. In: *Inventiones mathematicae* 2.2 (Apr. 1, 1966), pp. 134–144. ISSN: 1432-1297. DOI: [10.1007/BF01404549](https://doi.org/10.1007/BF01404549). URL: <https://doi.org/10.1007/BF01404549> (visited on 03/19/2026).
- [24] John Tate. “Finite flat group schemes”. In: *Modular Forms and Fermat’s Last Theorem*. Ed. by Gary Cornell, Joseph H. Silverman, and Glenn Stevens. Springer, 1997, pp. 121–154.
- [25] Jacques Vélú. “Isogenies between Elliptic Curves*”. In: ().
- [26] John Voight. *Quaternion Algebras*. Vol. 288. Graduate Texts in Mathematics. Springer, 2021.
- [27] William C. Waterhouse. “Abelian varieties over finite fields”. In: *Annales scientifiques de l’École Normale Supérieure* 2.4 (1969), pp. 521–560.
- [28] André Weil. “L’arithmétique sur les courbes algébriques”. Published in *Acta Mathematica* 52 (1929), 281–315. PhD thesis. Université de Paris, 1928.